

# Cross-Platform

## Data Processing

- Use Cases & Challenges -

Zoi Kaoudi

Jorge Quiané

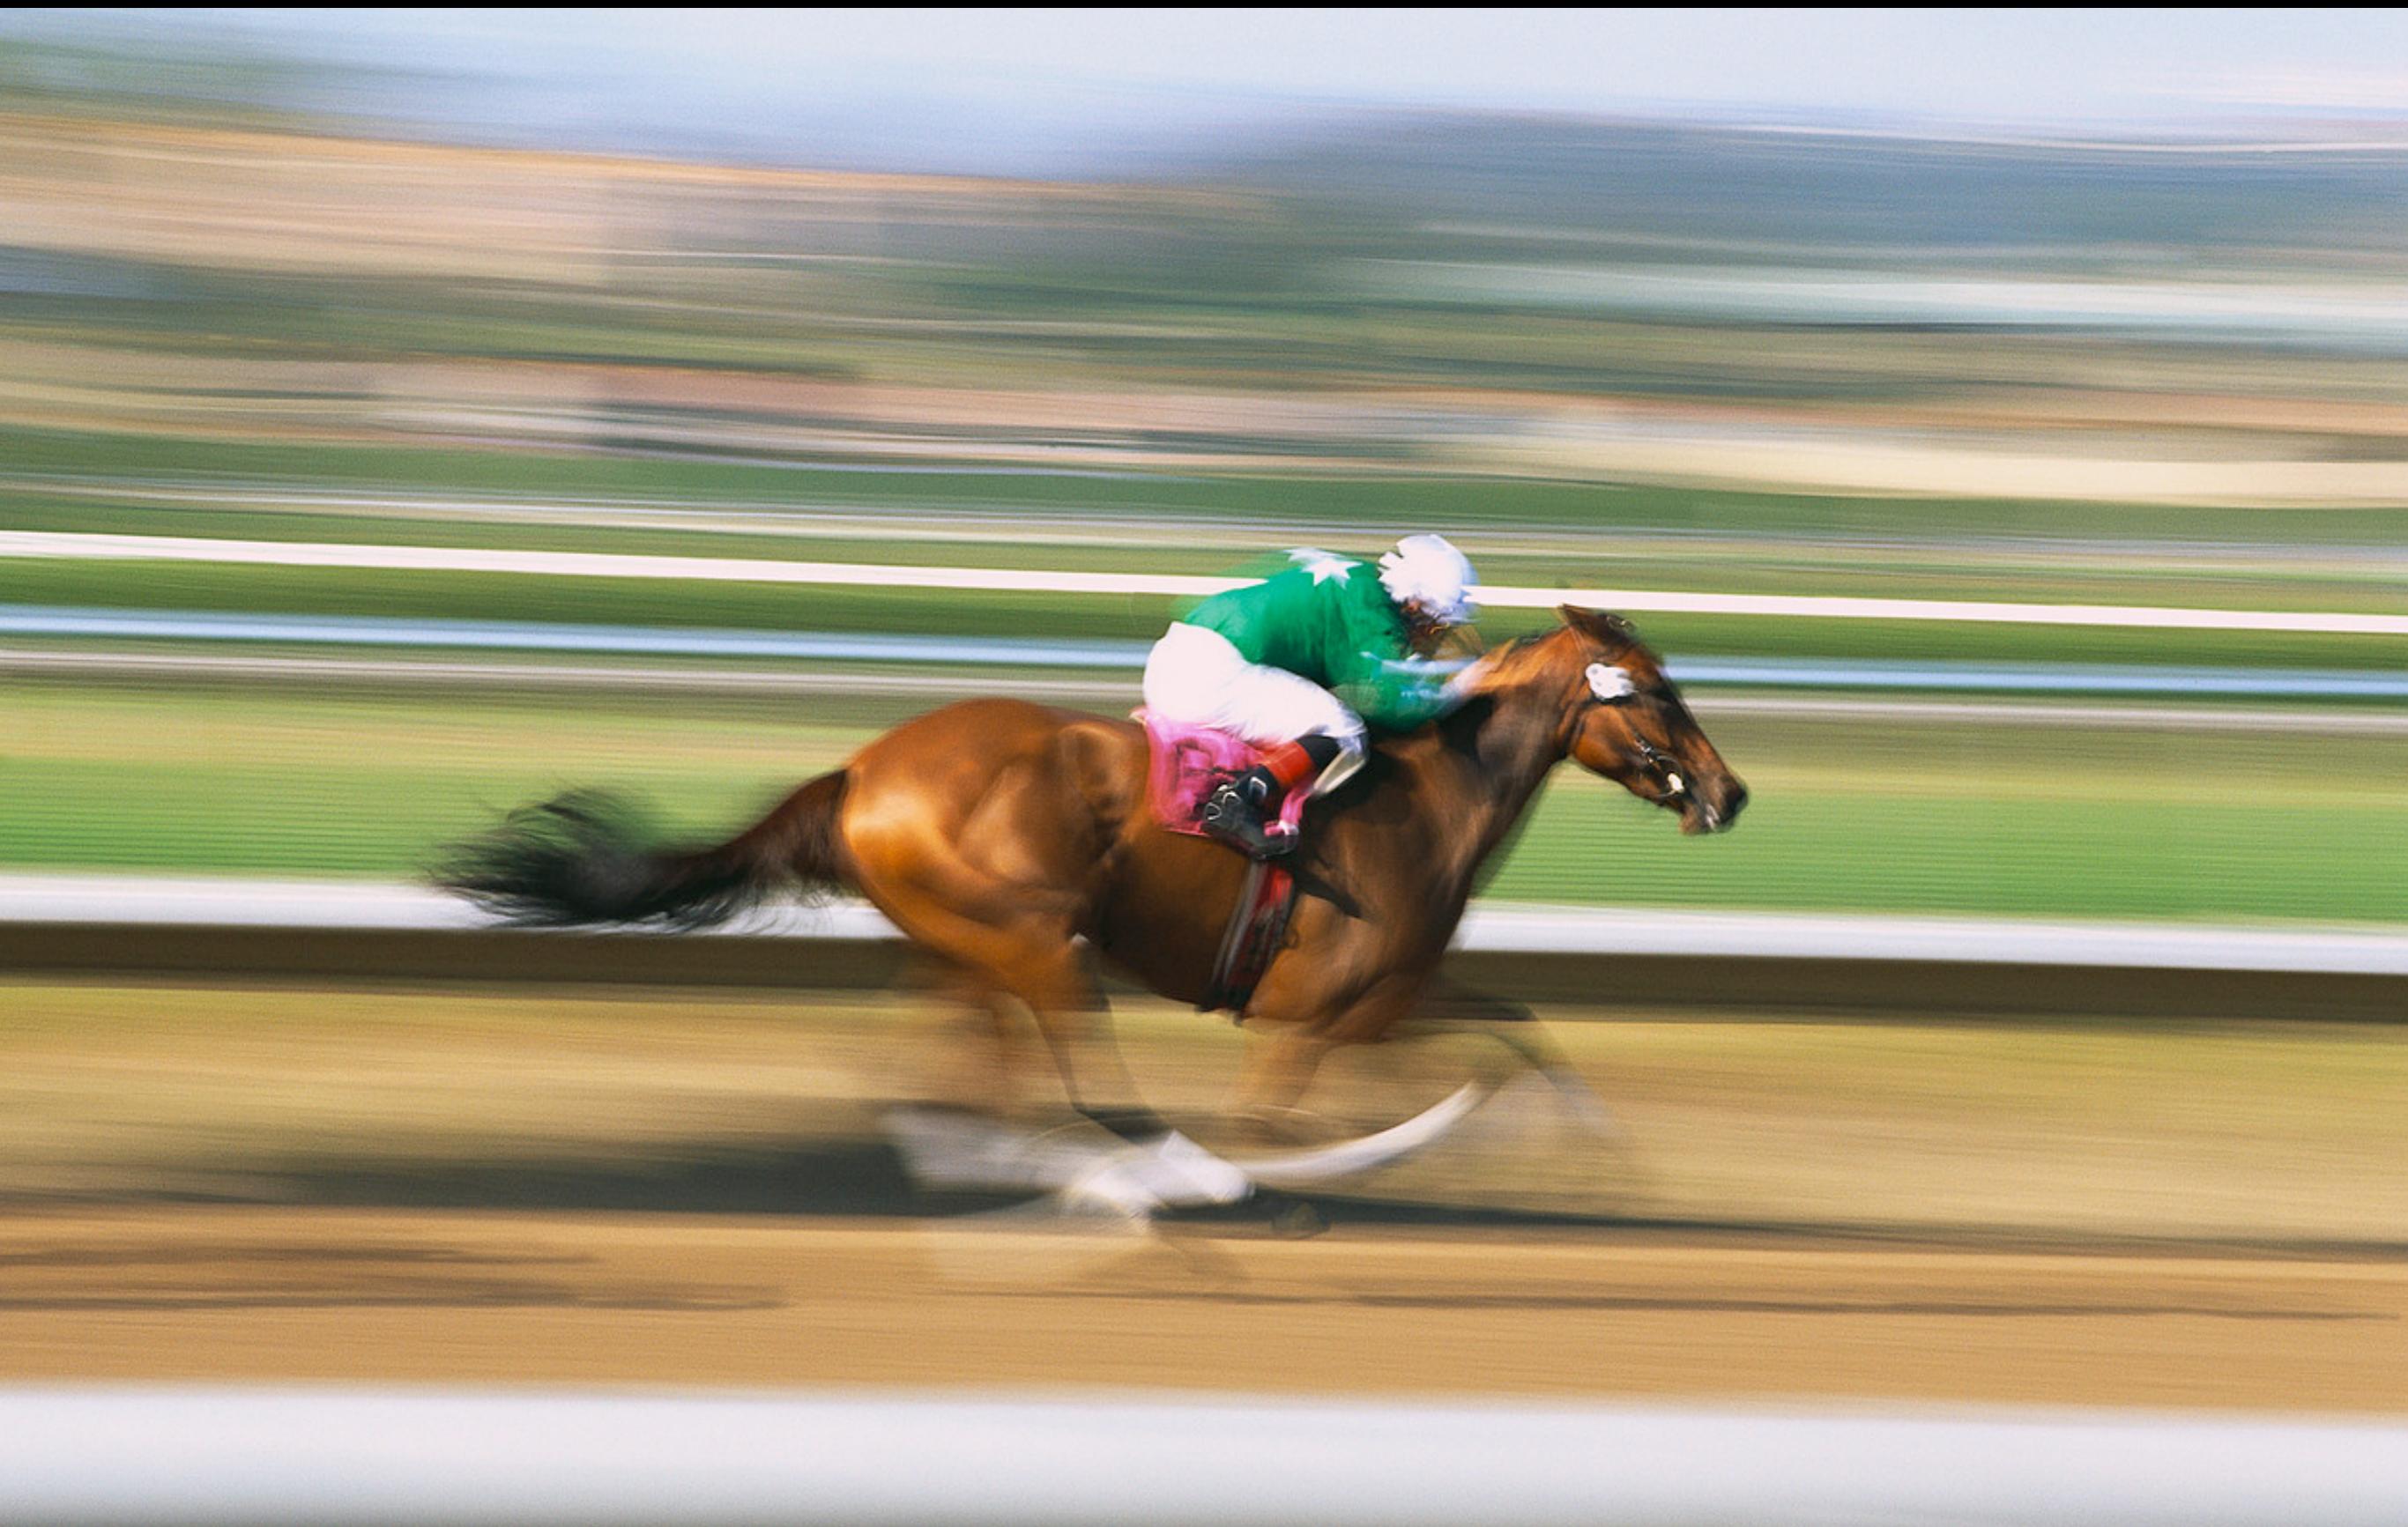
Qatar Computing Research Institute  
(QCRI)



# Nature Inspired

# Going to Race?

# Going to Race?



# Going to the Desert?

# Going to the Desert?



# Going to the Forest?

# Going to the Forest?



# Going to the Everest?

# Going to the Everest?



# Going to Fly?

# Going to Fly?

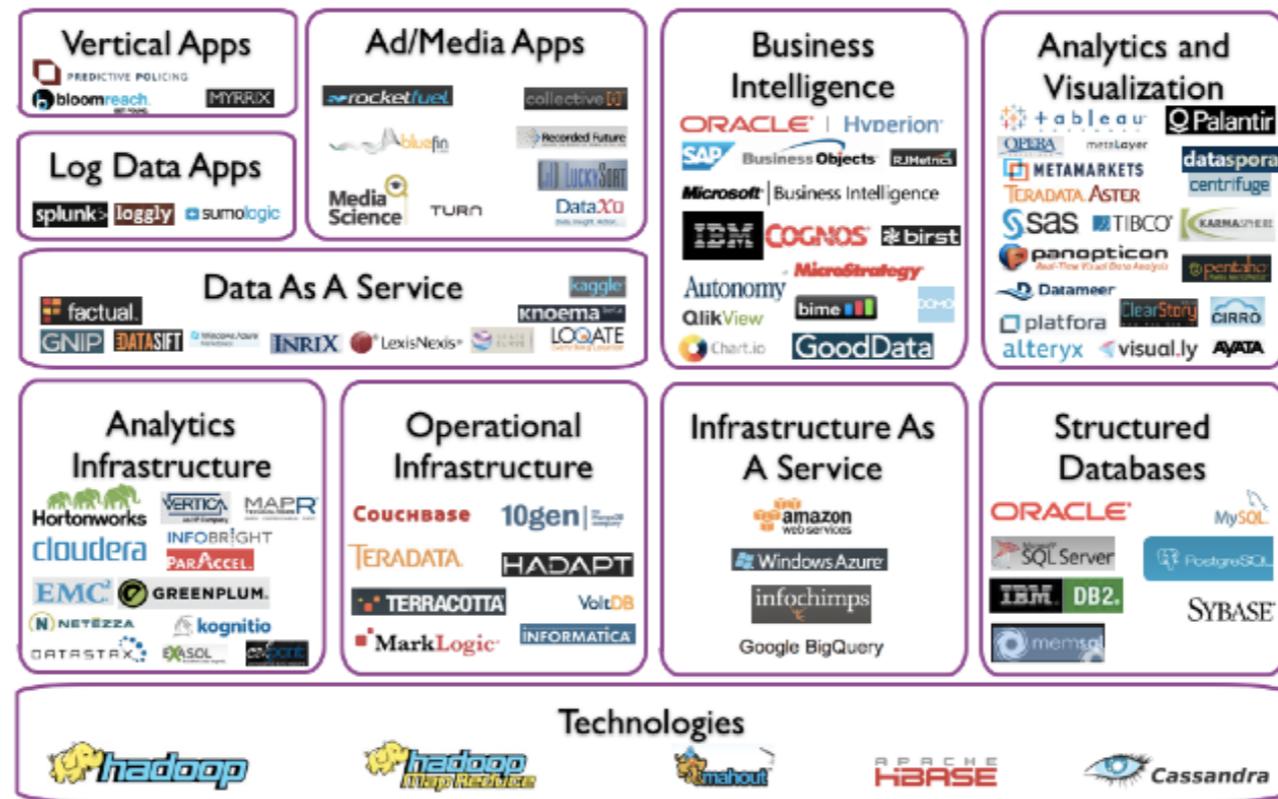


Fly!

# FACT 1

— One Size Does Not Fit All —

# Big Data Landscape 2014

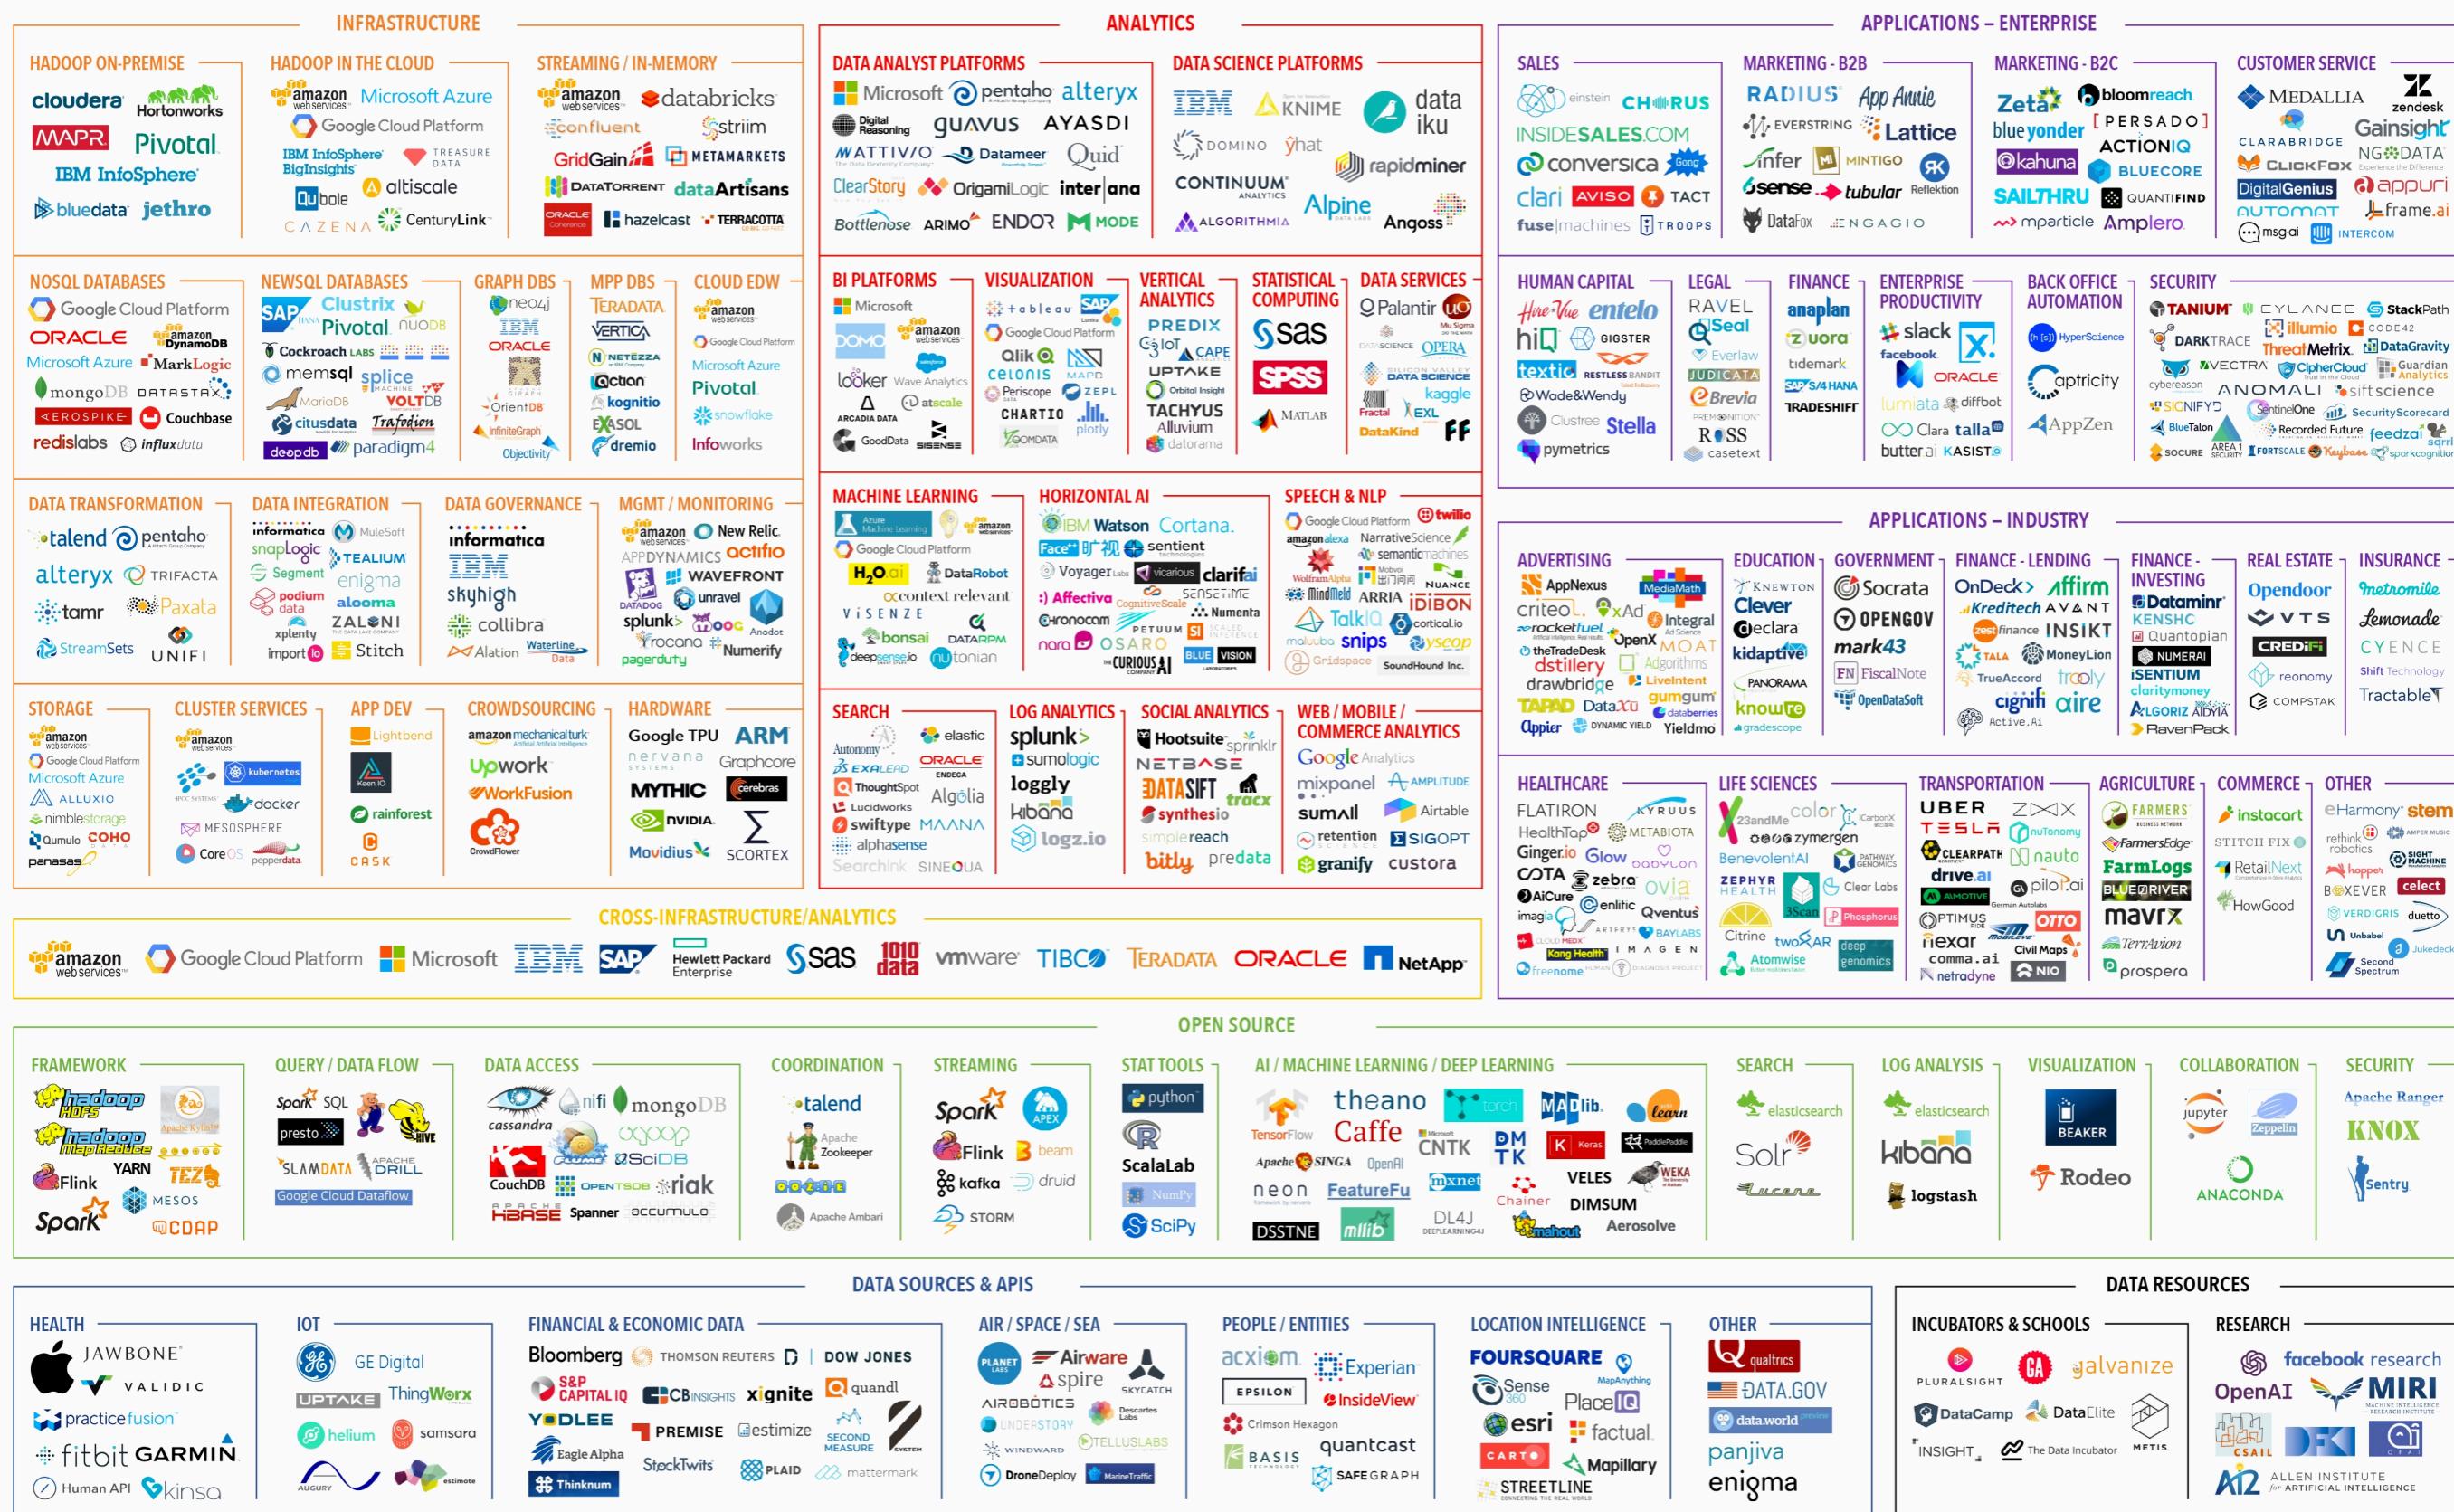


Copyright © 2012 Dave Feinleib

dave@vcDave.com

[blogs.forbes.com/davefeinleib](http://blogs.forbes.com/davefeinleib)

# Big Data Landscape 2017



# FACT 2

## — Zoo of Systems —

# Current Data Analytics

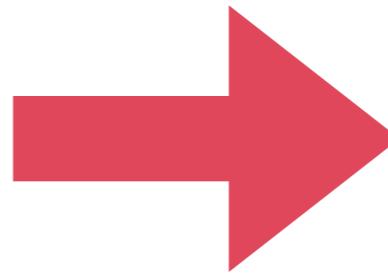
## SPJ Queries

```
SELECT name
FROM employee e
INNER JOIN Person p
ON p.firstname = e.name
```

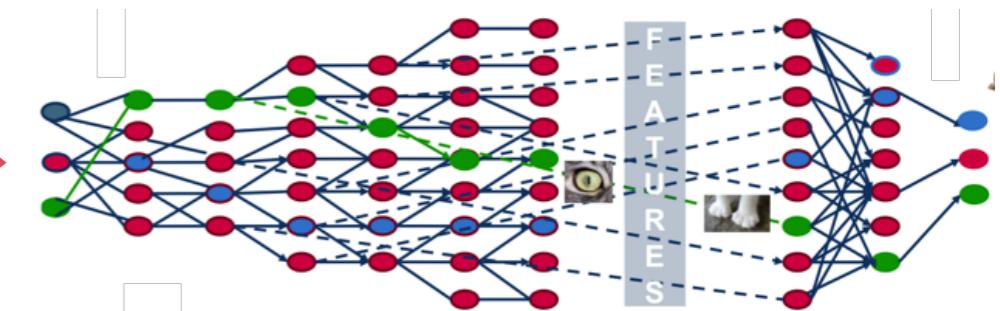
# Current Data Analytics

## SPJ Queries

```
SELECT name  
FROM employee e  
INNER JOIN Person p  
ON p.firstname = e.name
```



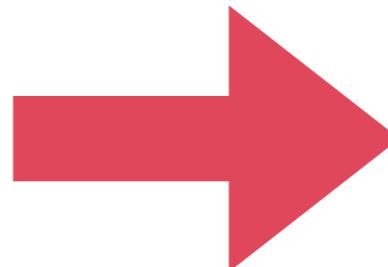
## Deep Learning



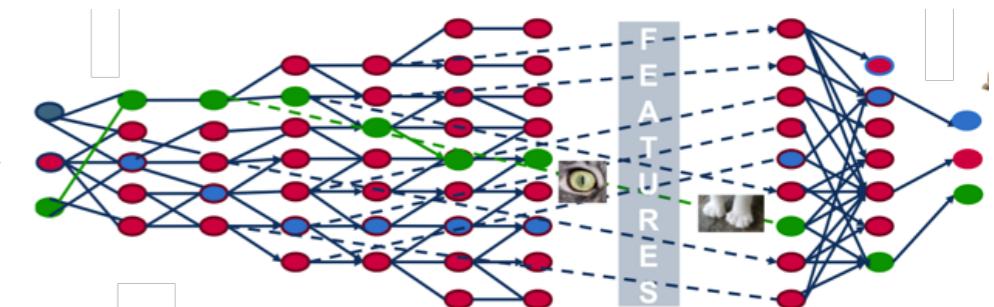
# Current Data Analytics

## SPJ Queries

```
SELECT name  
FROM employee e  
INNER JOIN Person p  
ON p.firstname = e.name
```



## Deep Learning



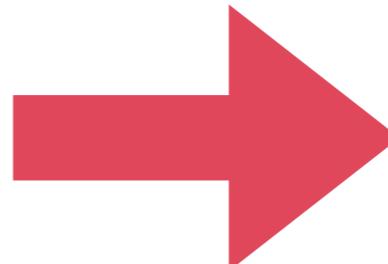
## Relational Data

Activity Code	Activity Name
23	Patching
24	Overlay
25	Crack Sealing

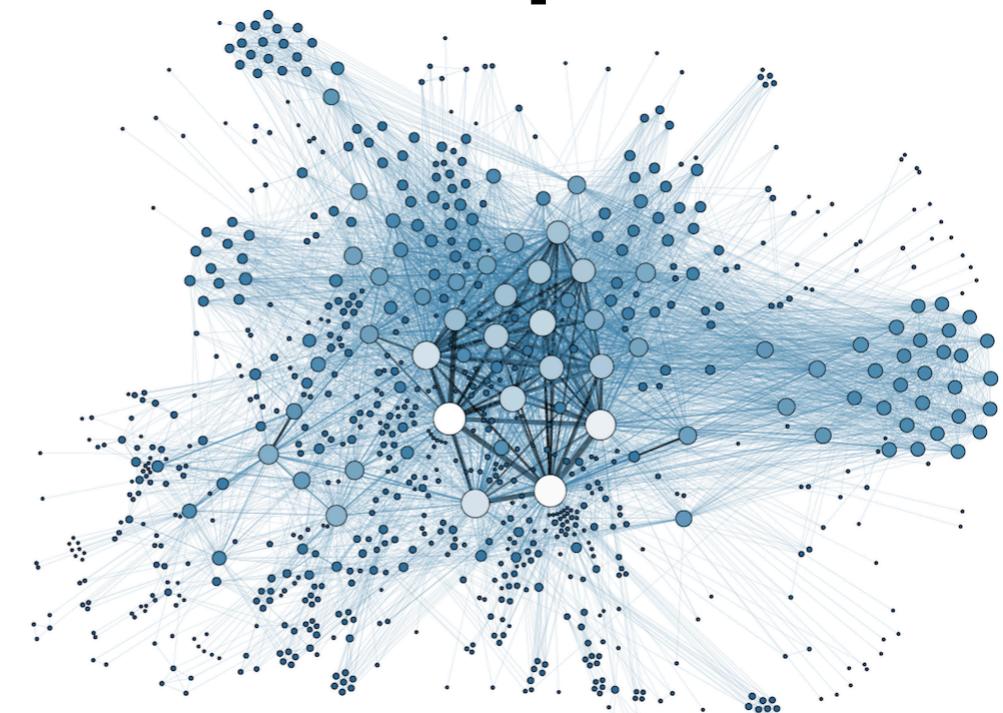
Key = 24

Activity Code	Date	Route No.
24	01/12/01	I-95
24	02/08/01	I-66

Date	Activity Code	Route No.
01/12/01	24	I-95
01/15/01	23	I-495
02/08/01	24	I-66



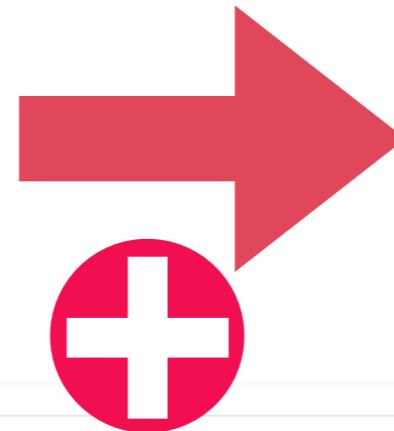
## Graph Data



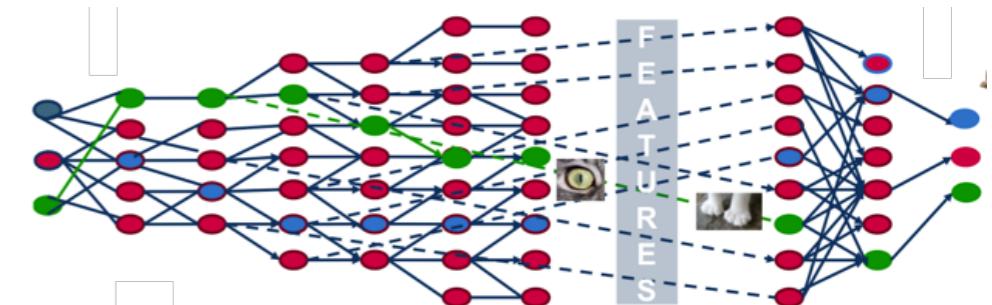
# Current Data Analytics

## SPJ Queries

```
SELECT name  
FROM employee e  
INNER JOIN Person p  
ON p.firstname = e.name
```



## Deep Learning



## — Queries Beyond a Single Platform —

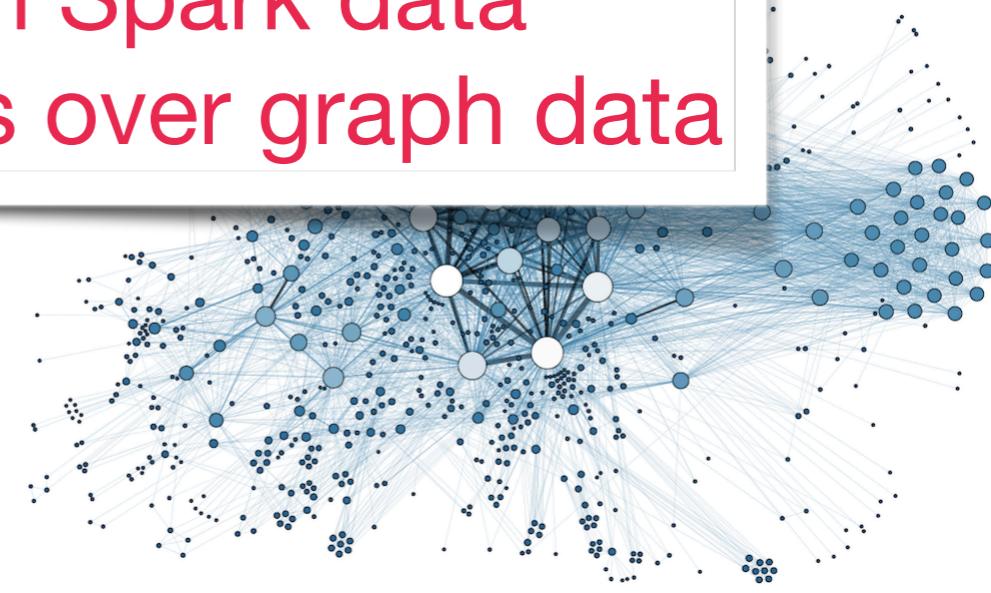
R Q1: unstructured + structured data processing

ta Q2: Join Postgres data with Spark data

Q3: Run machine learning tasks over graph data

Date	Activity Code	Route No.
01/12/01	24	I-95
01/15/01	23	I-495
02/08/01	24	I-66

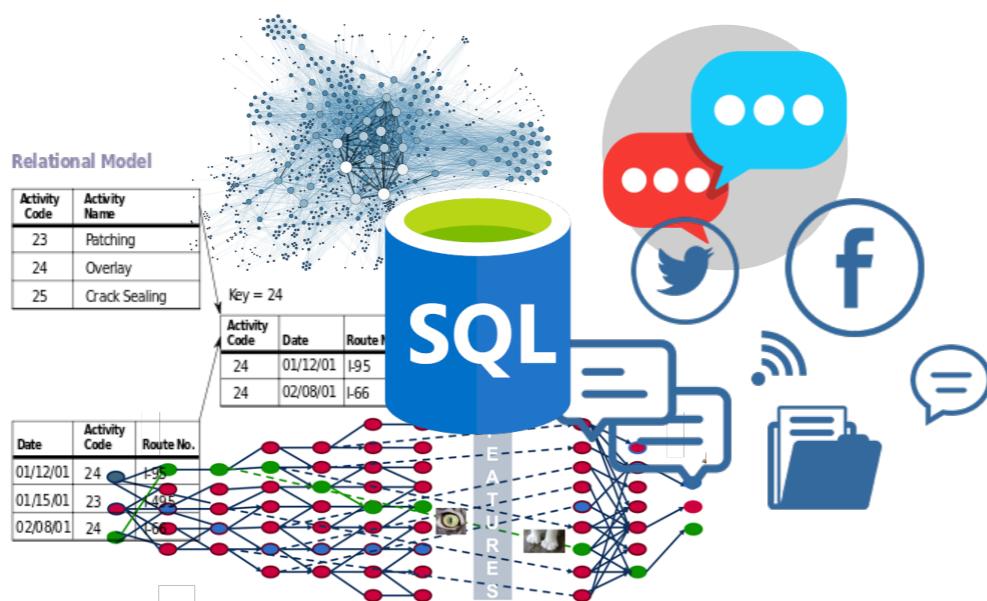
CUSTID	DRIVID	INVOICEINFO
24	01/12/01	I-95
24	02/08/01	I-66



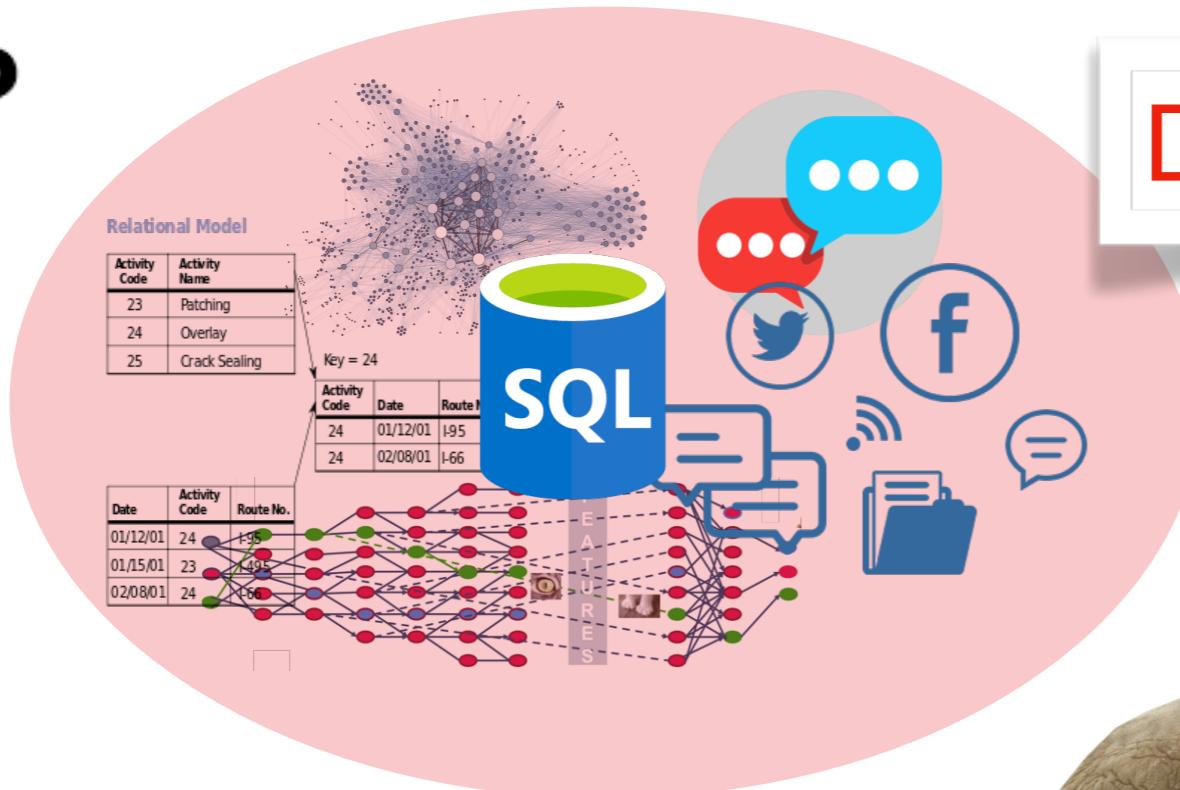
# FACT 3

— Variety of Data Analytics —

# What to Do?



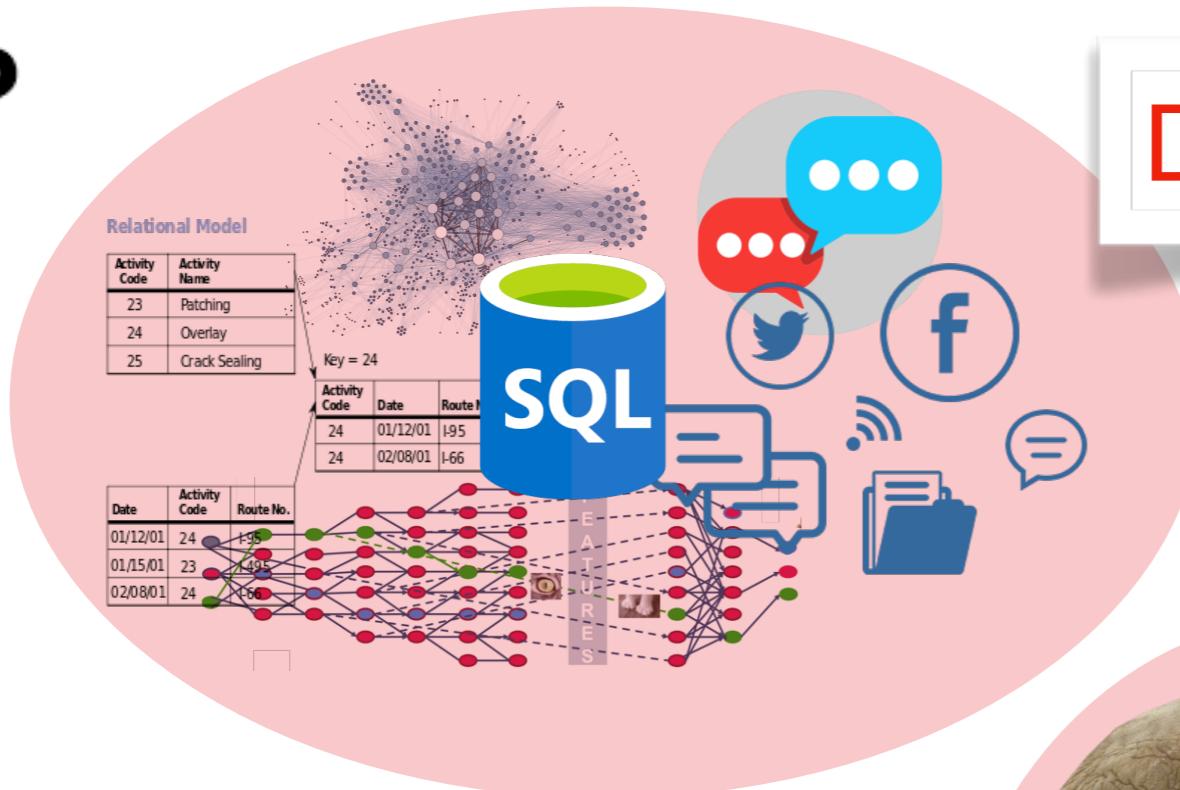
# What to Do?



Diverse Data Analytics



# What to Do?

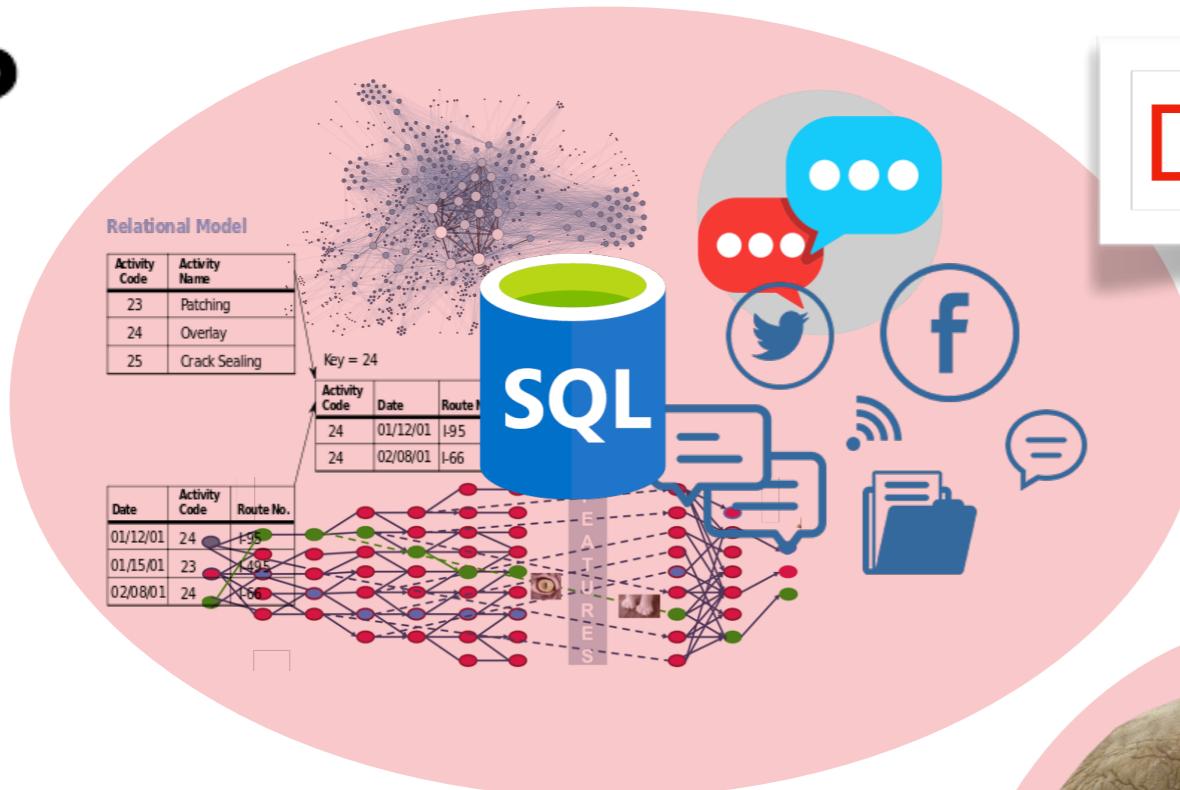


Diverse Data Analytics

One Cannot Fit All

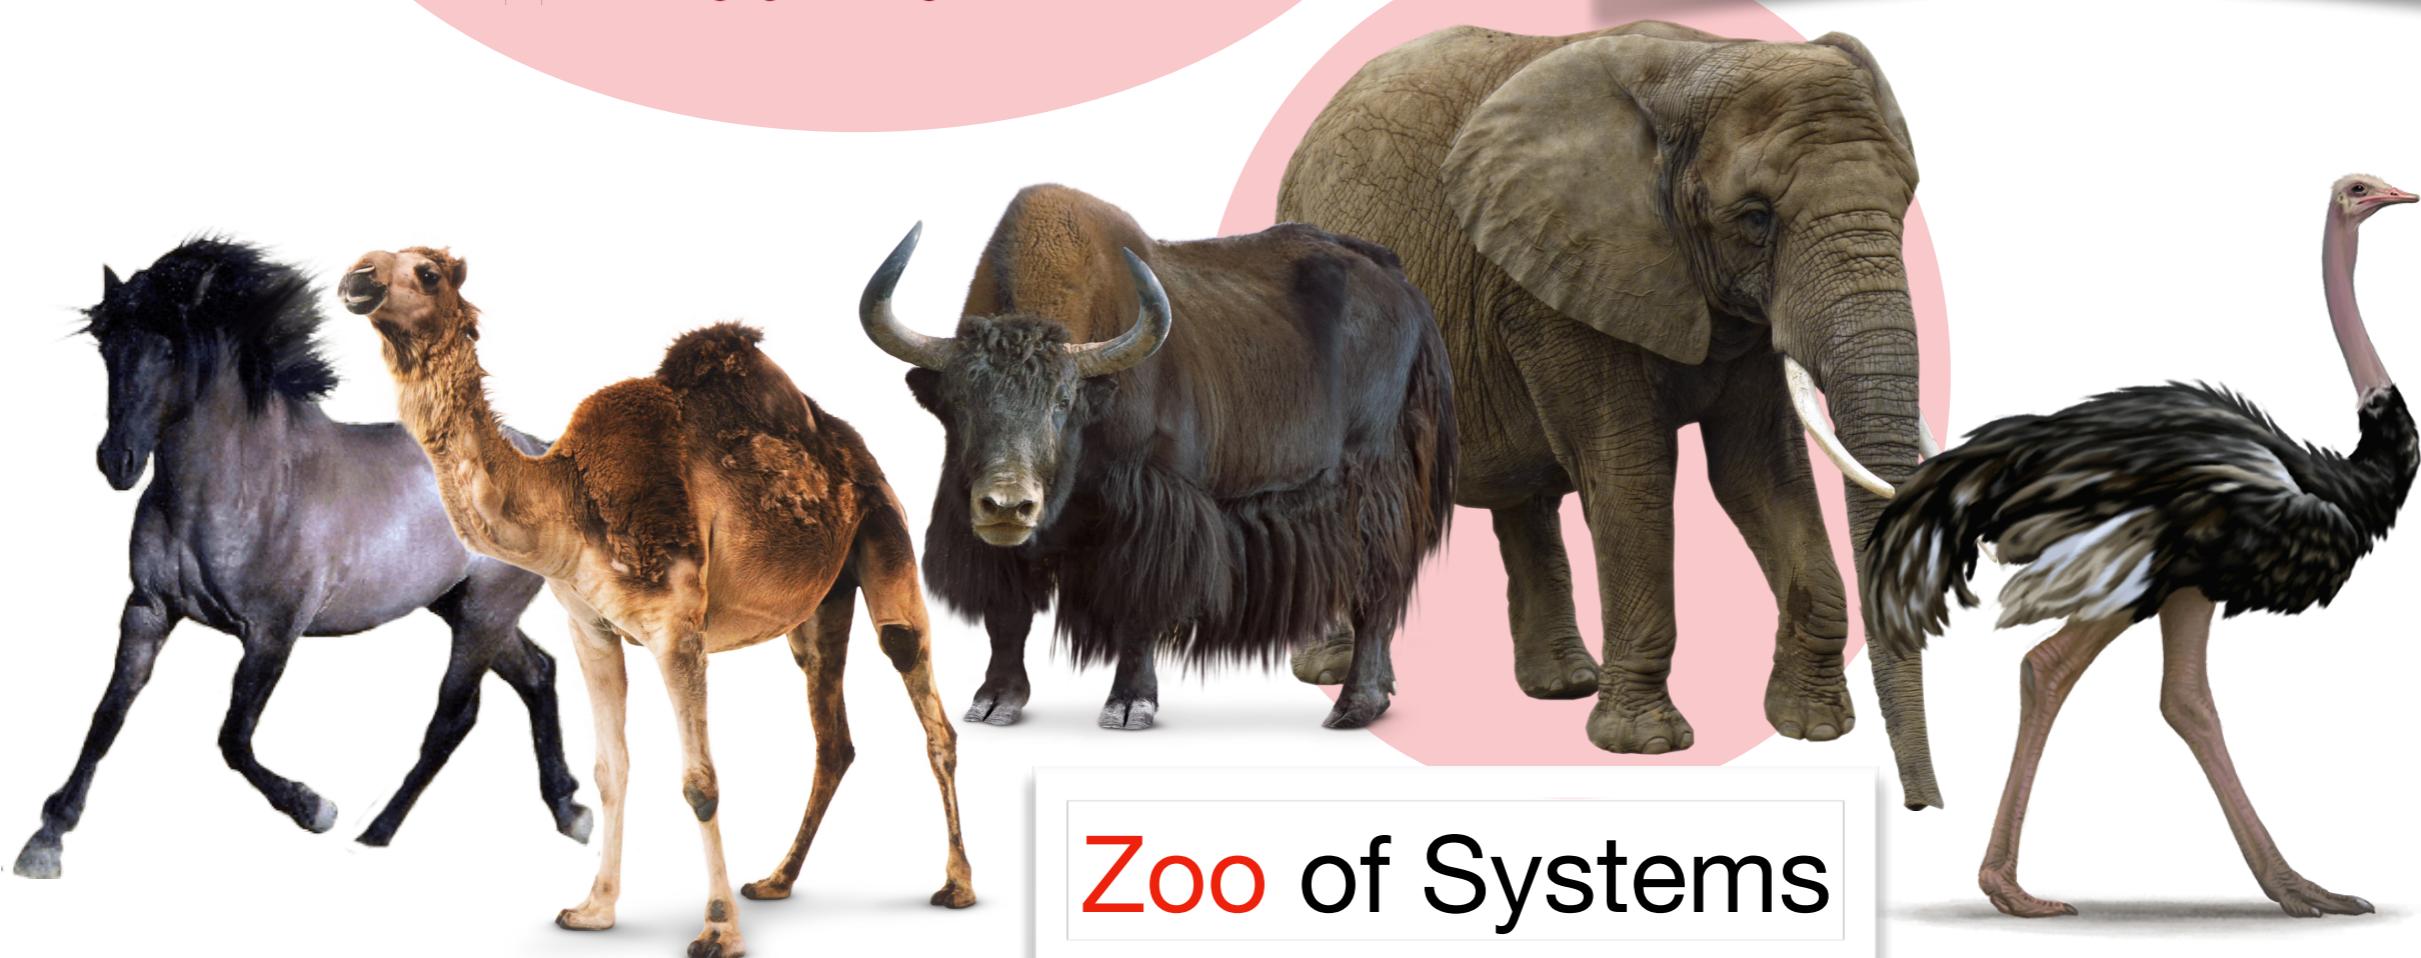


# What to Do?



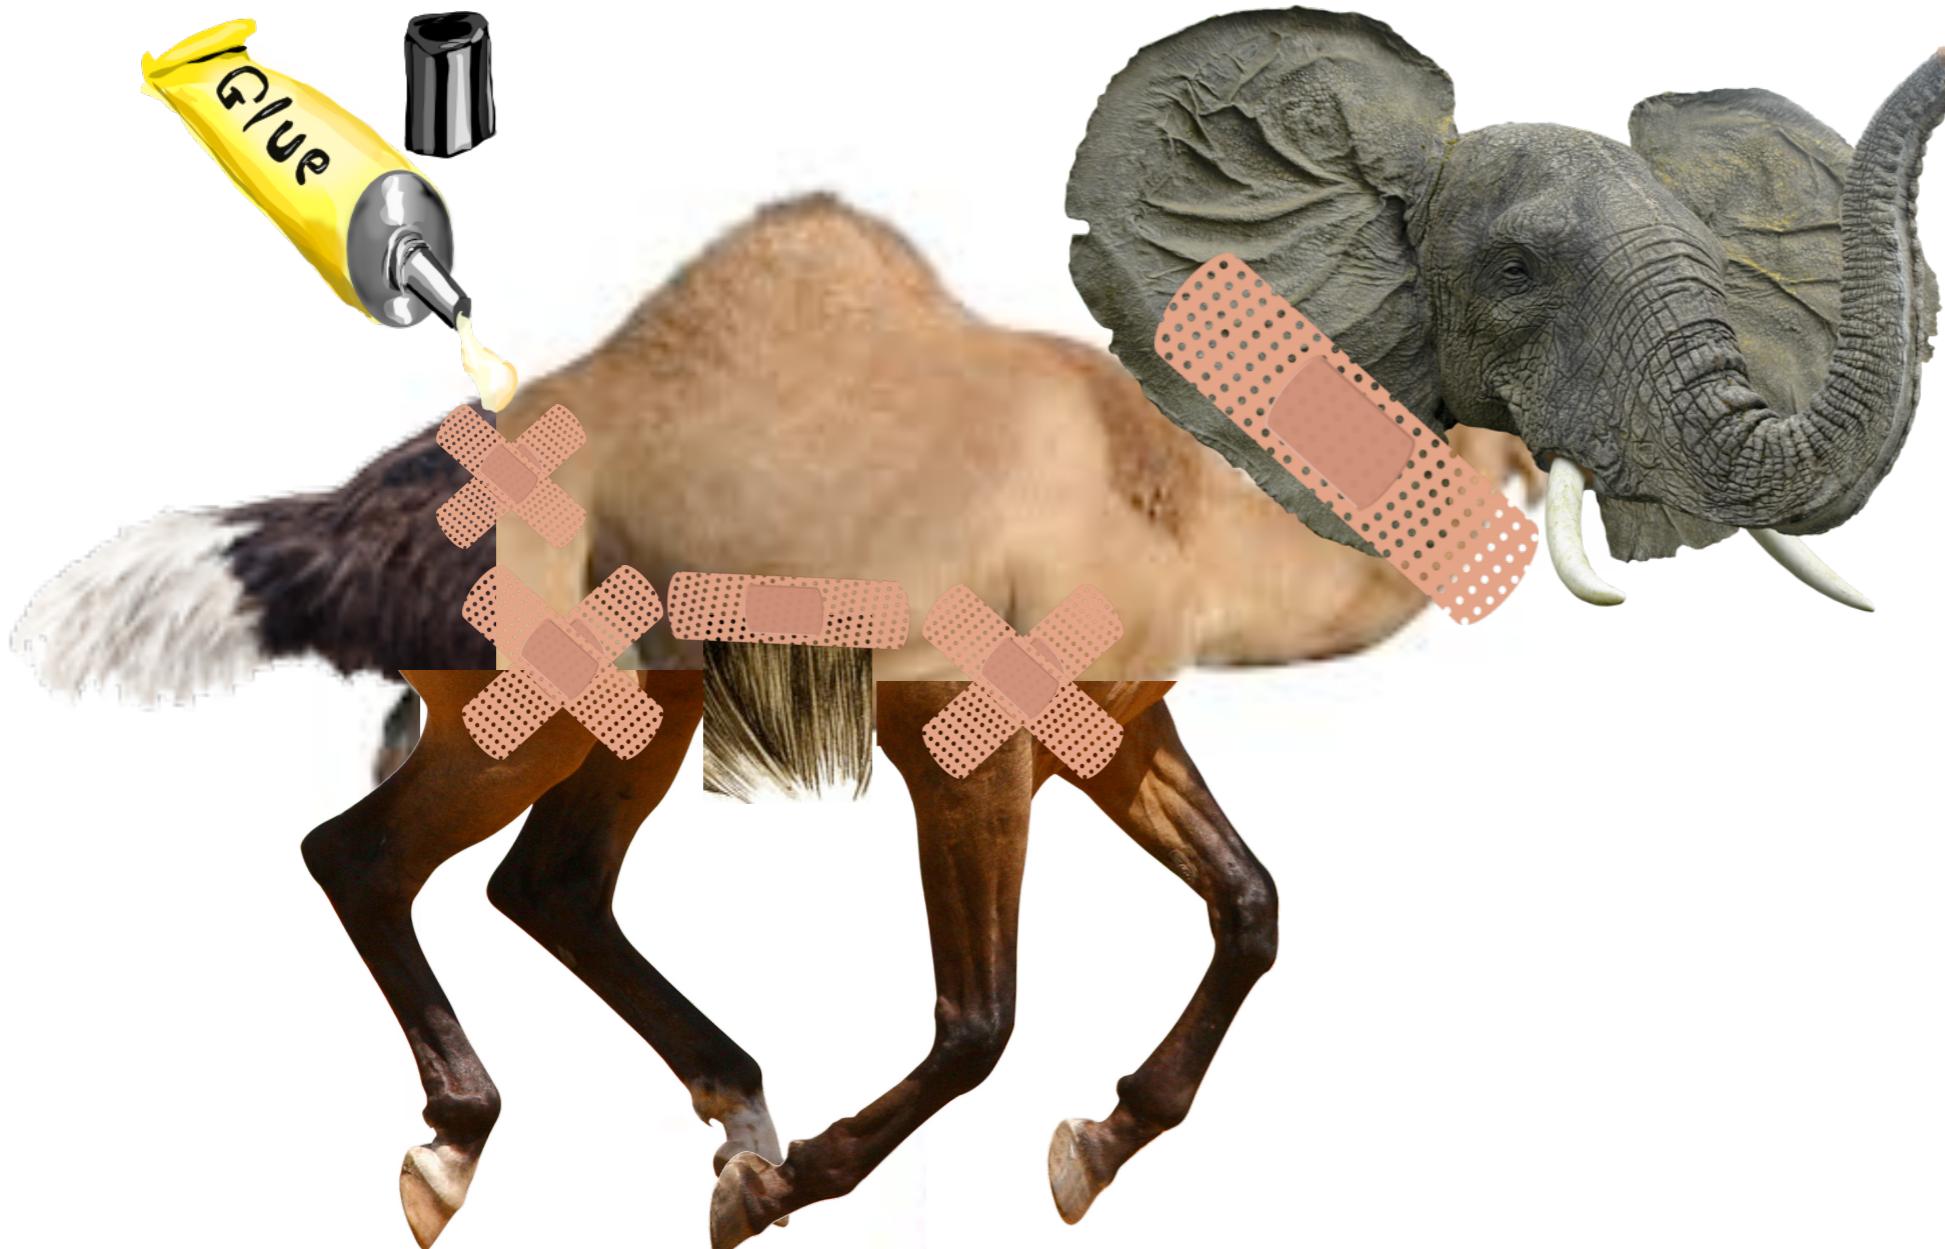
Diverse Data Analytics

One Cannot Fit All

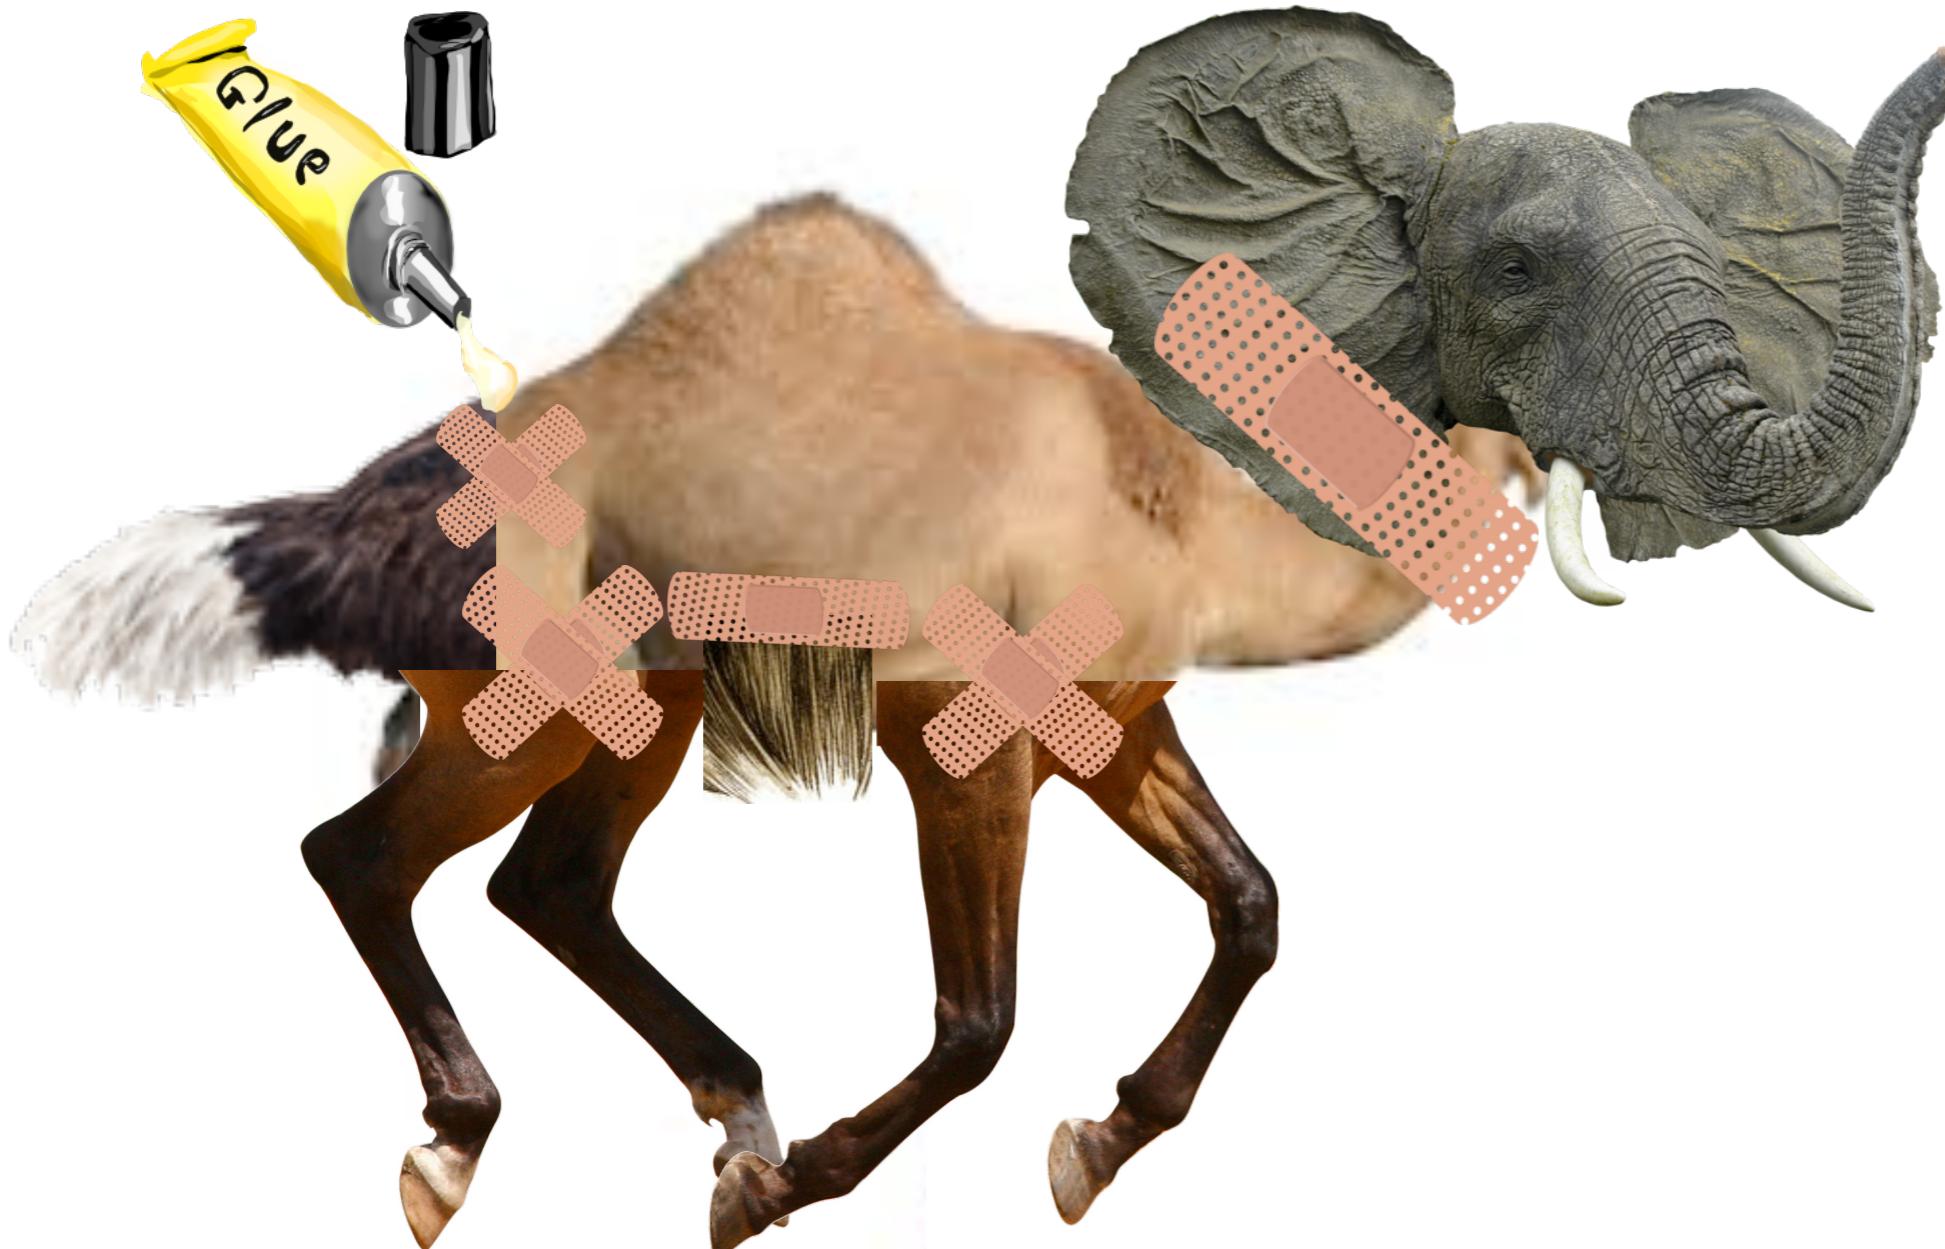


Zoo of Systems

# Hybrid Systems



# Hybrid Systems



Examples: HadoopDB, SystemML, FlumeJava ...

# Hybrid Systems

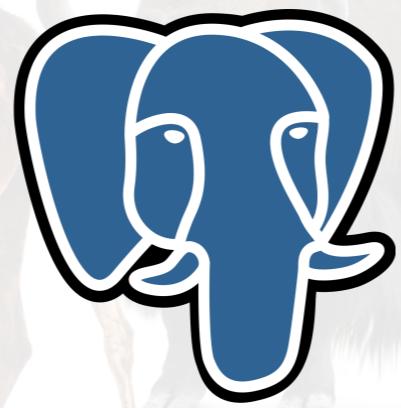


Examples: HadoopDB, SystemML, FlumeJava ...

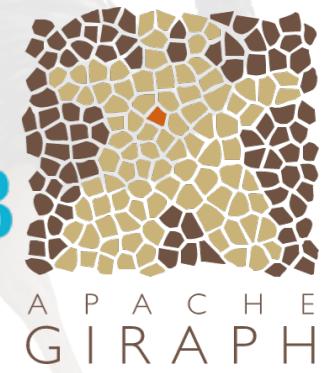
# Taming Them to Live Together



# Taming Them to Live Together



SciDB

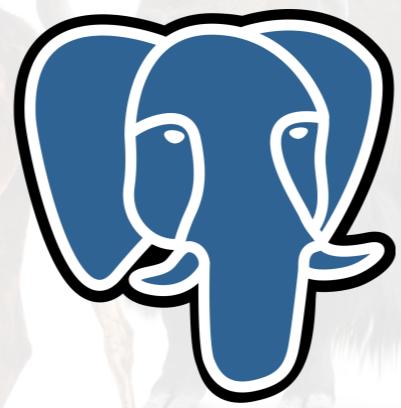


# Taming Them to Live Together

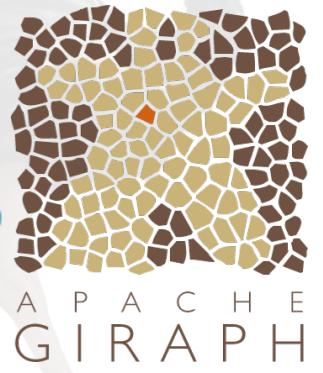


use of multiple data processing platforms for query processing (single run or across several runs)

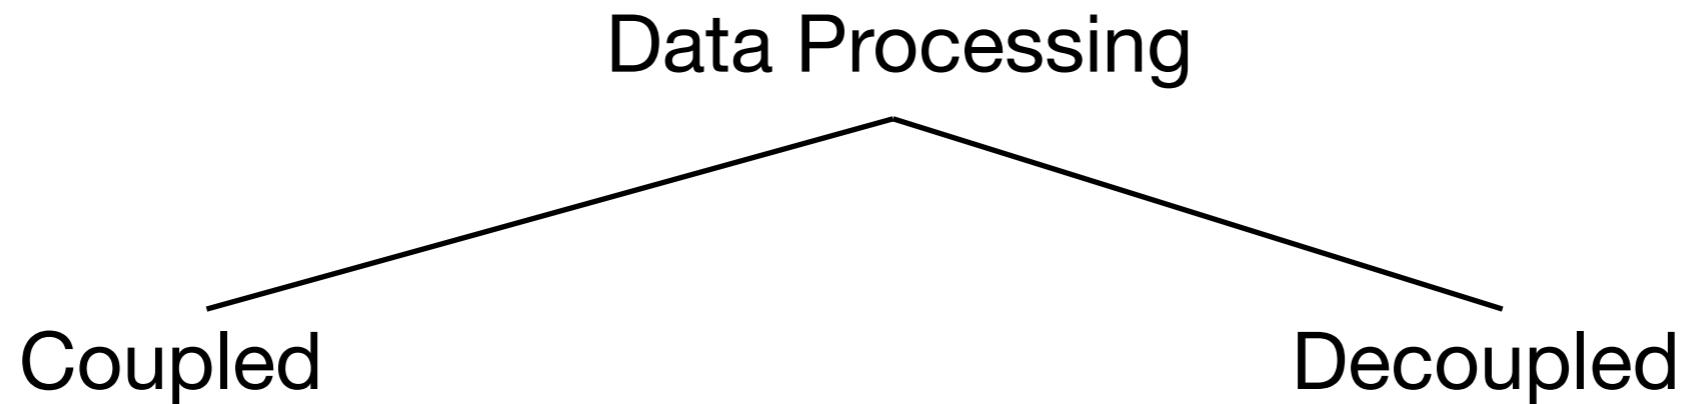
## Cross-Platform Data Processing



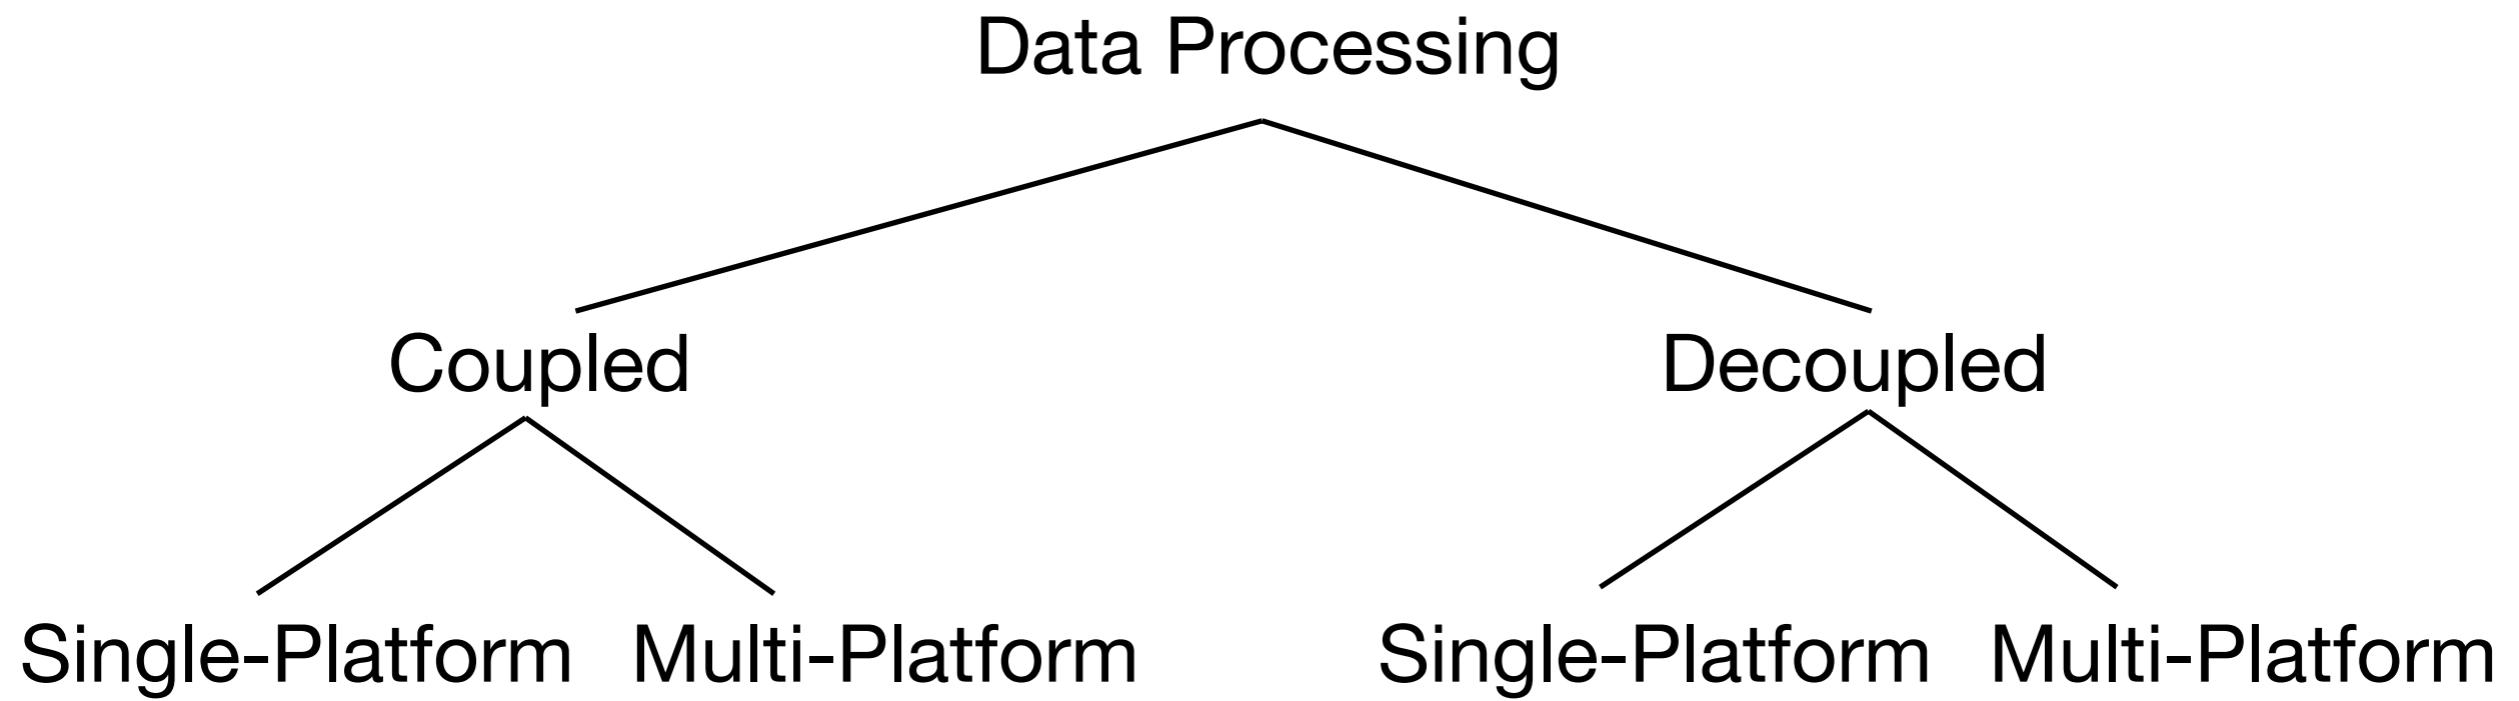
SciDB



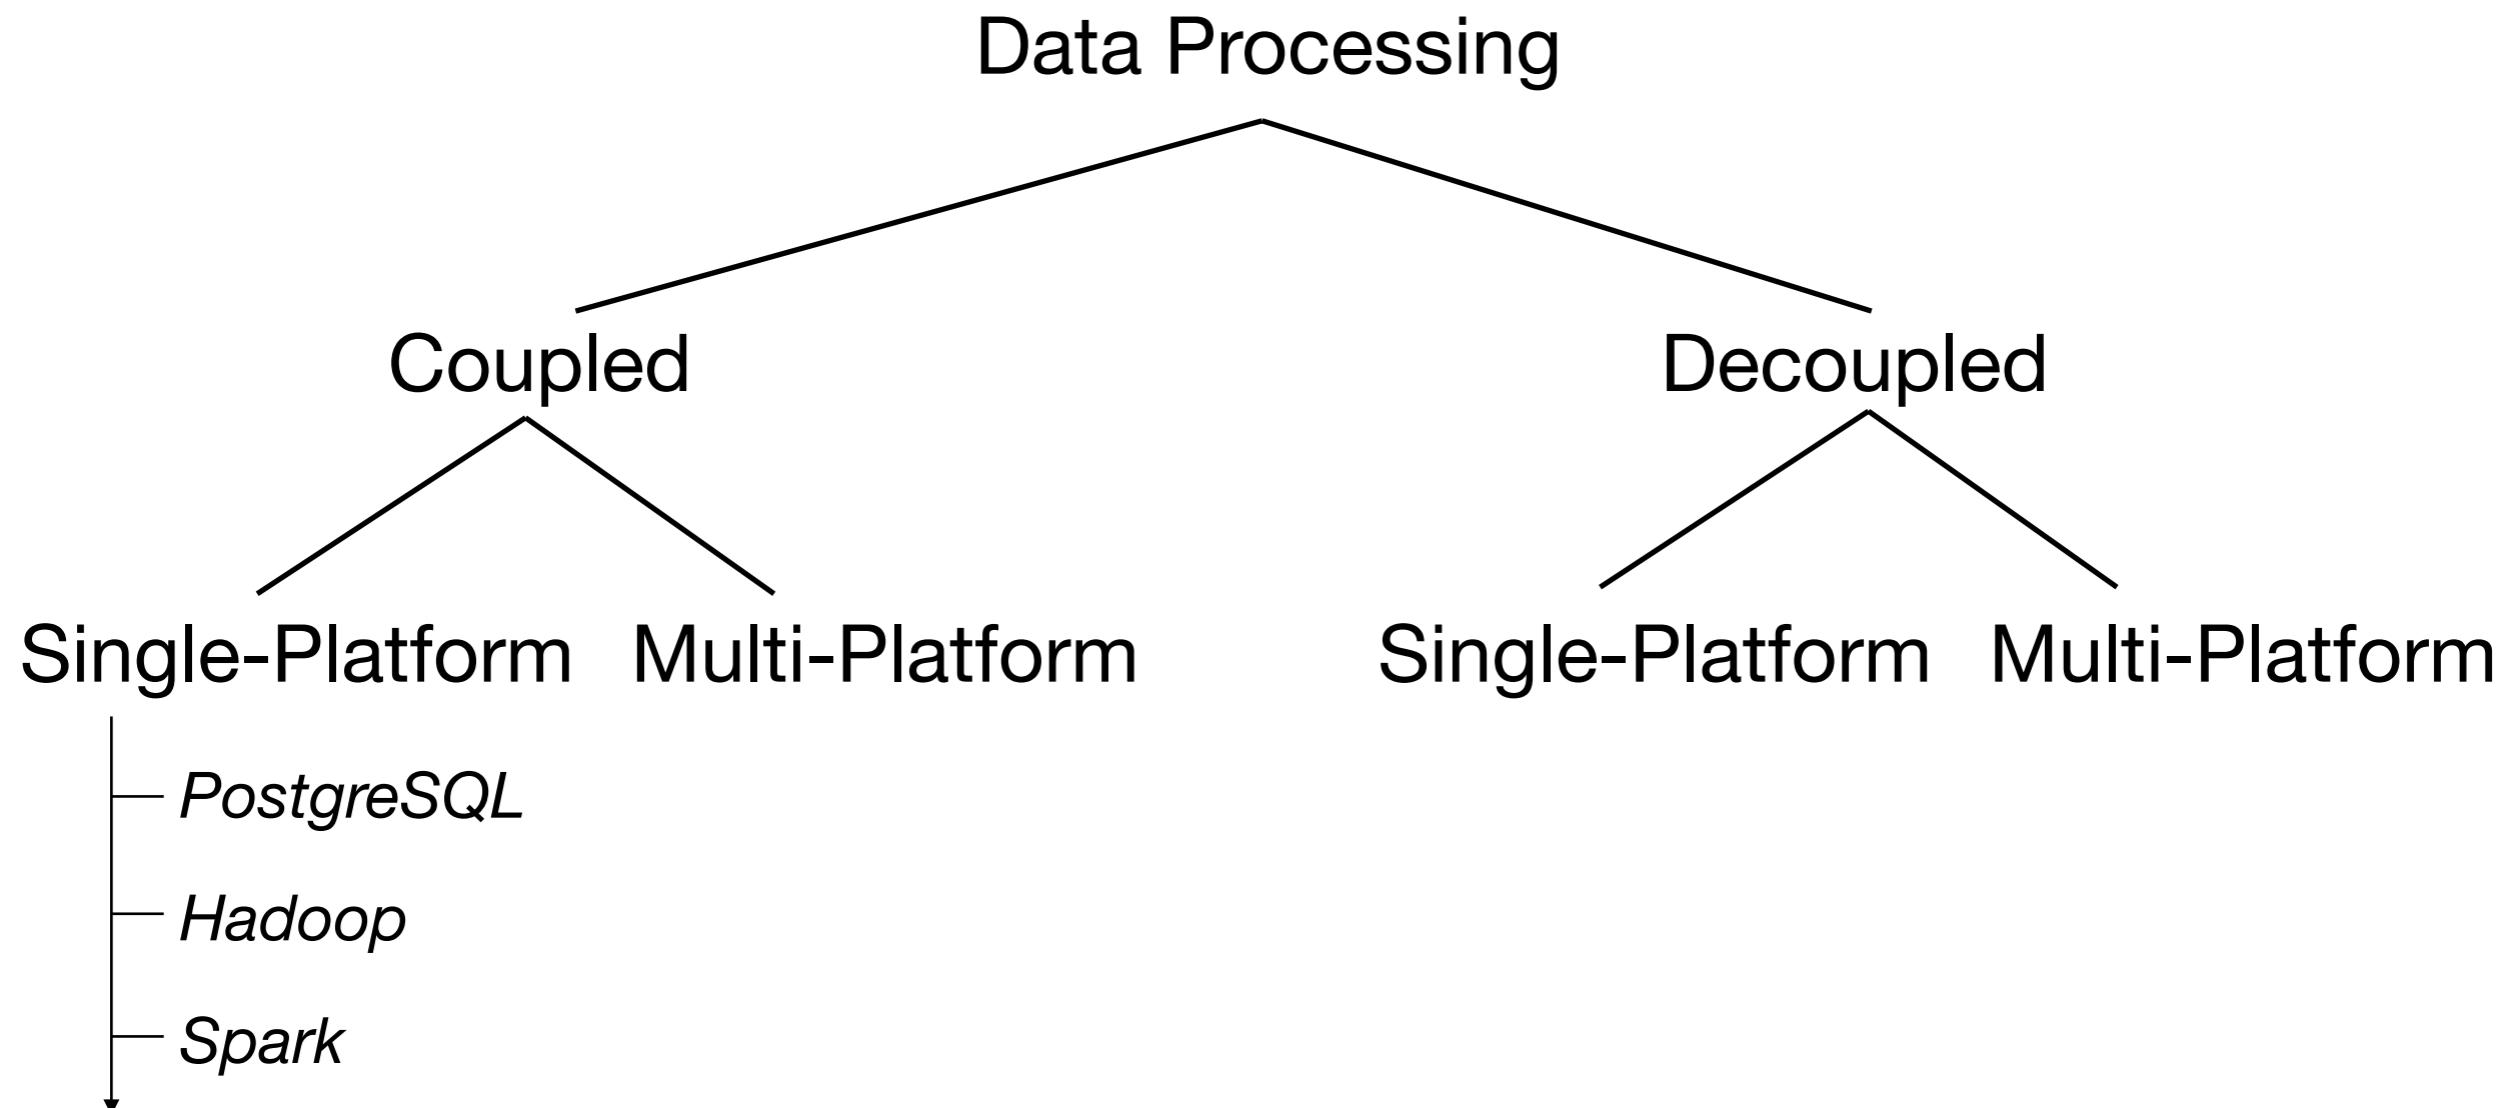
# Data Processing Taxonomy



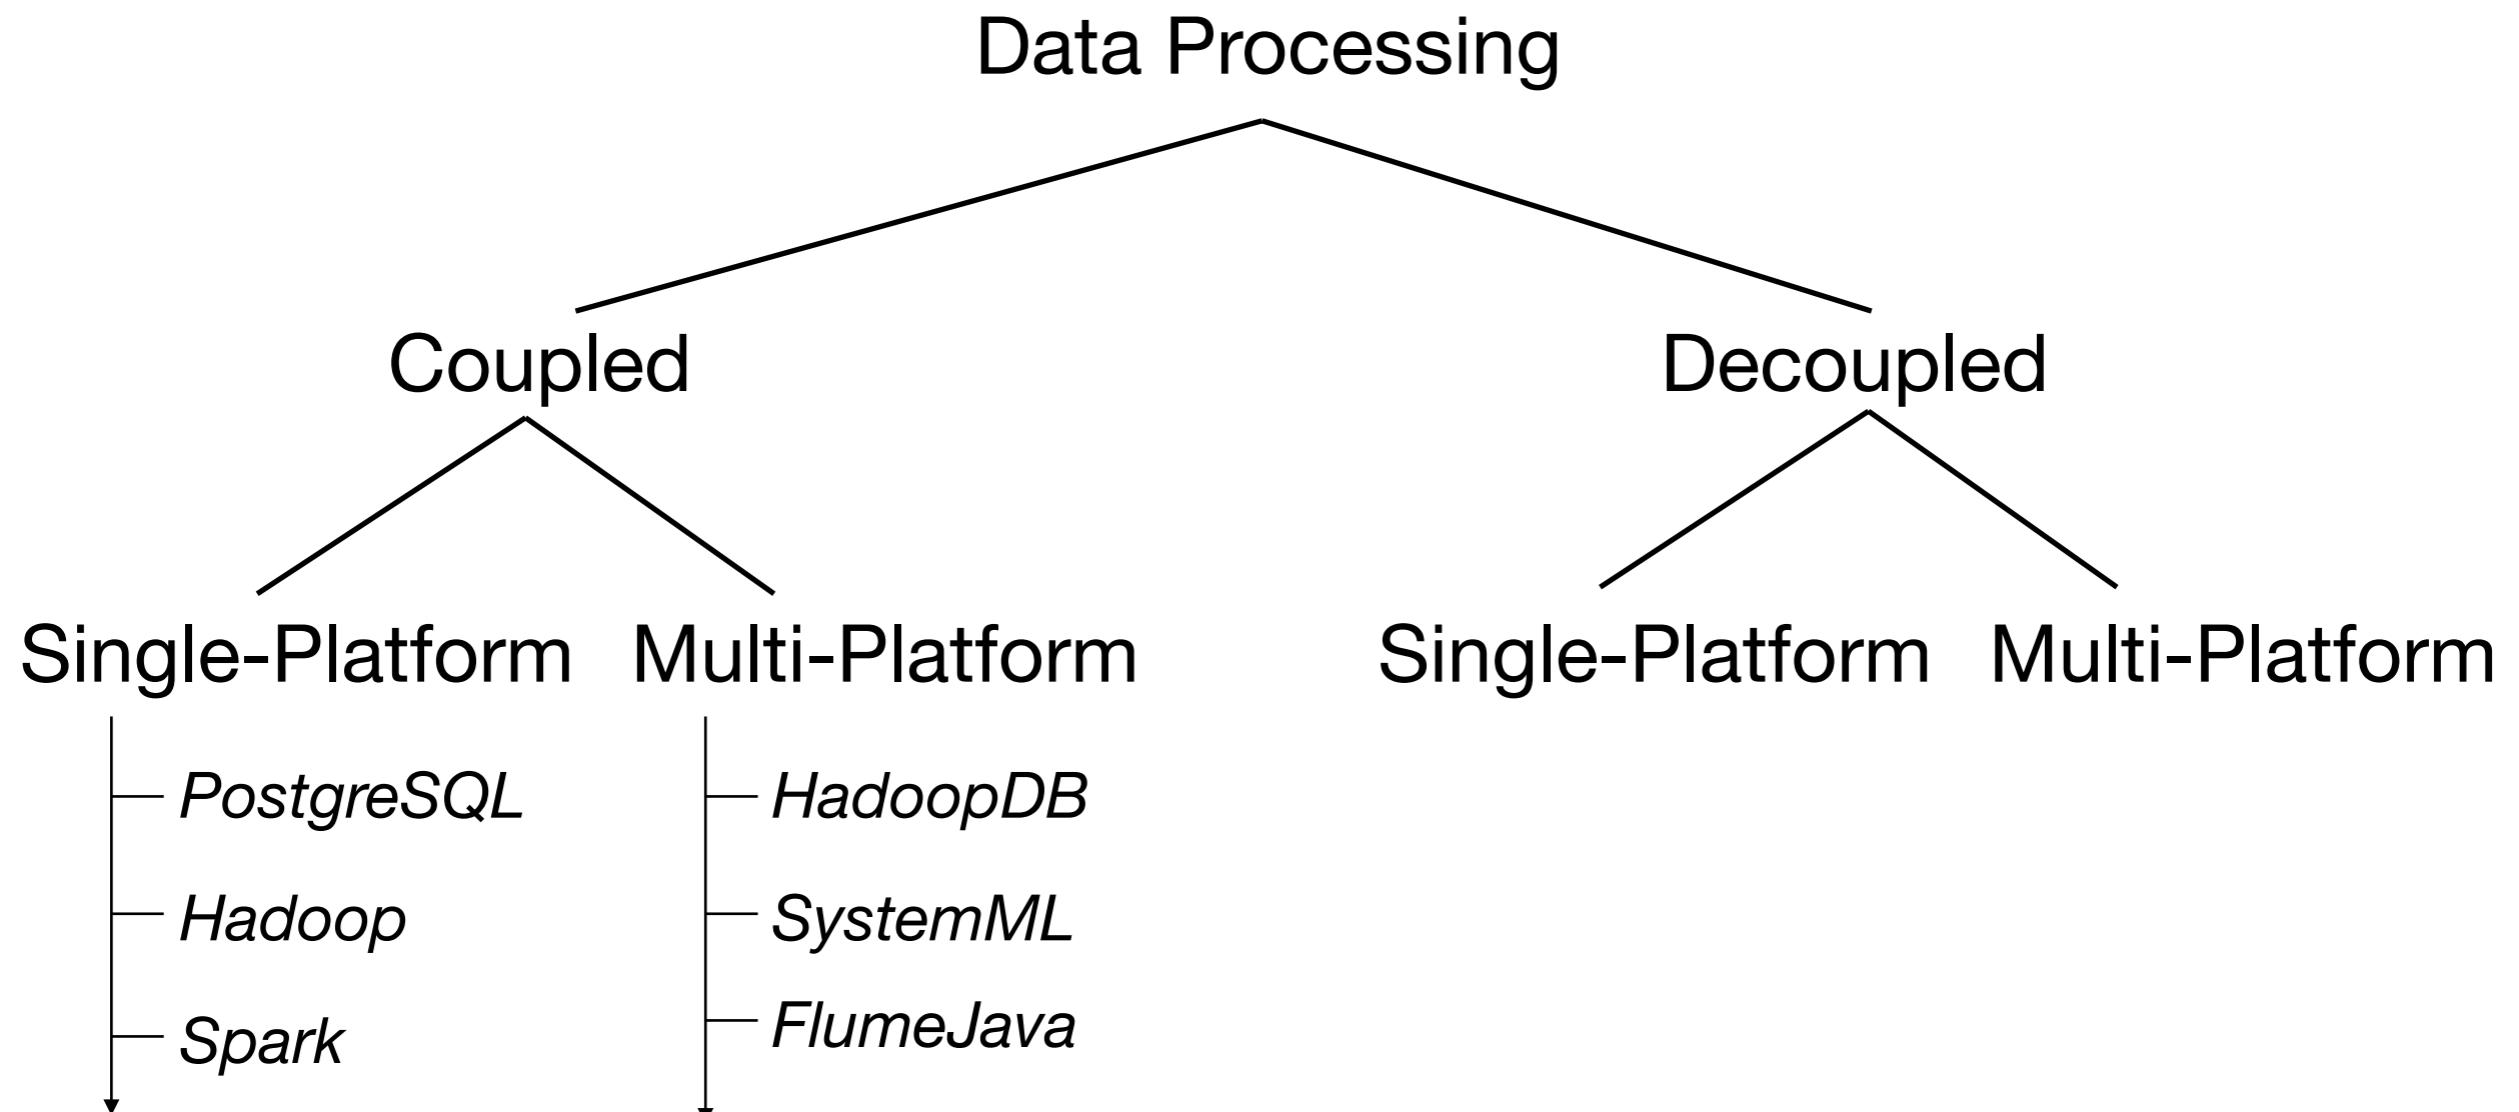
# Data Processing Taxonomy



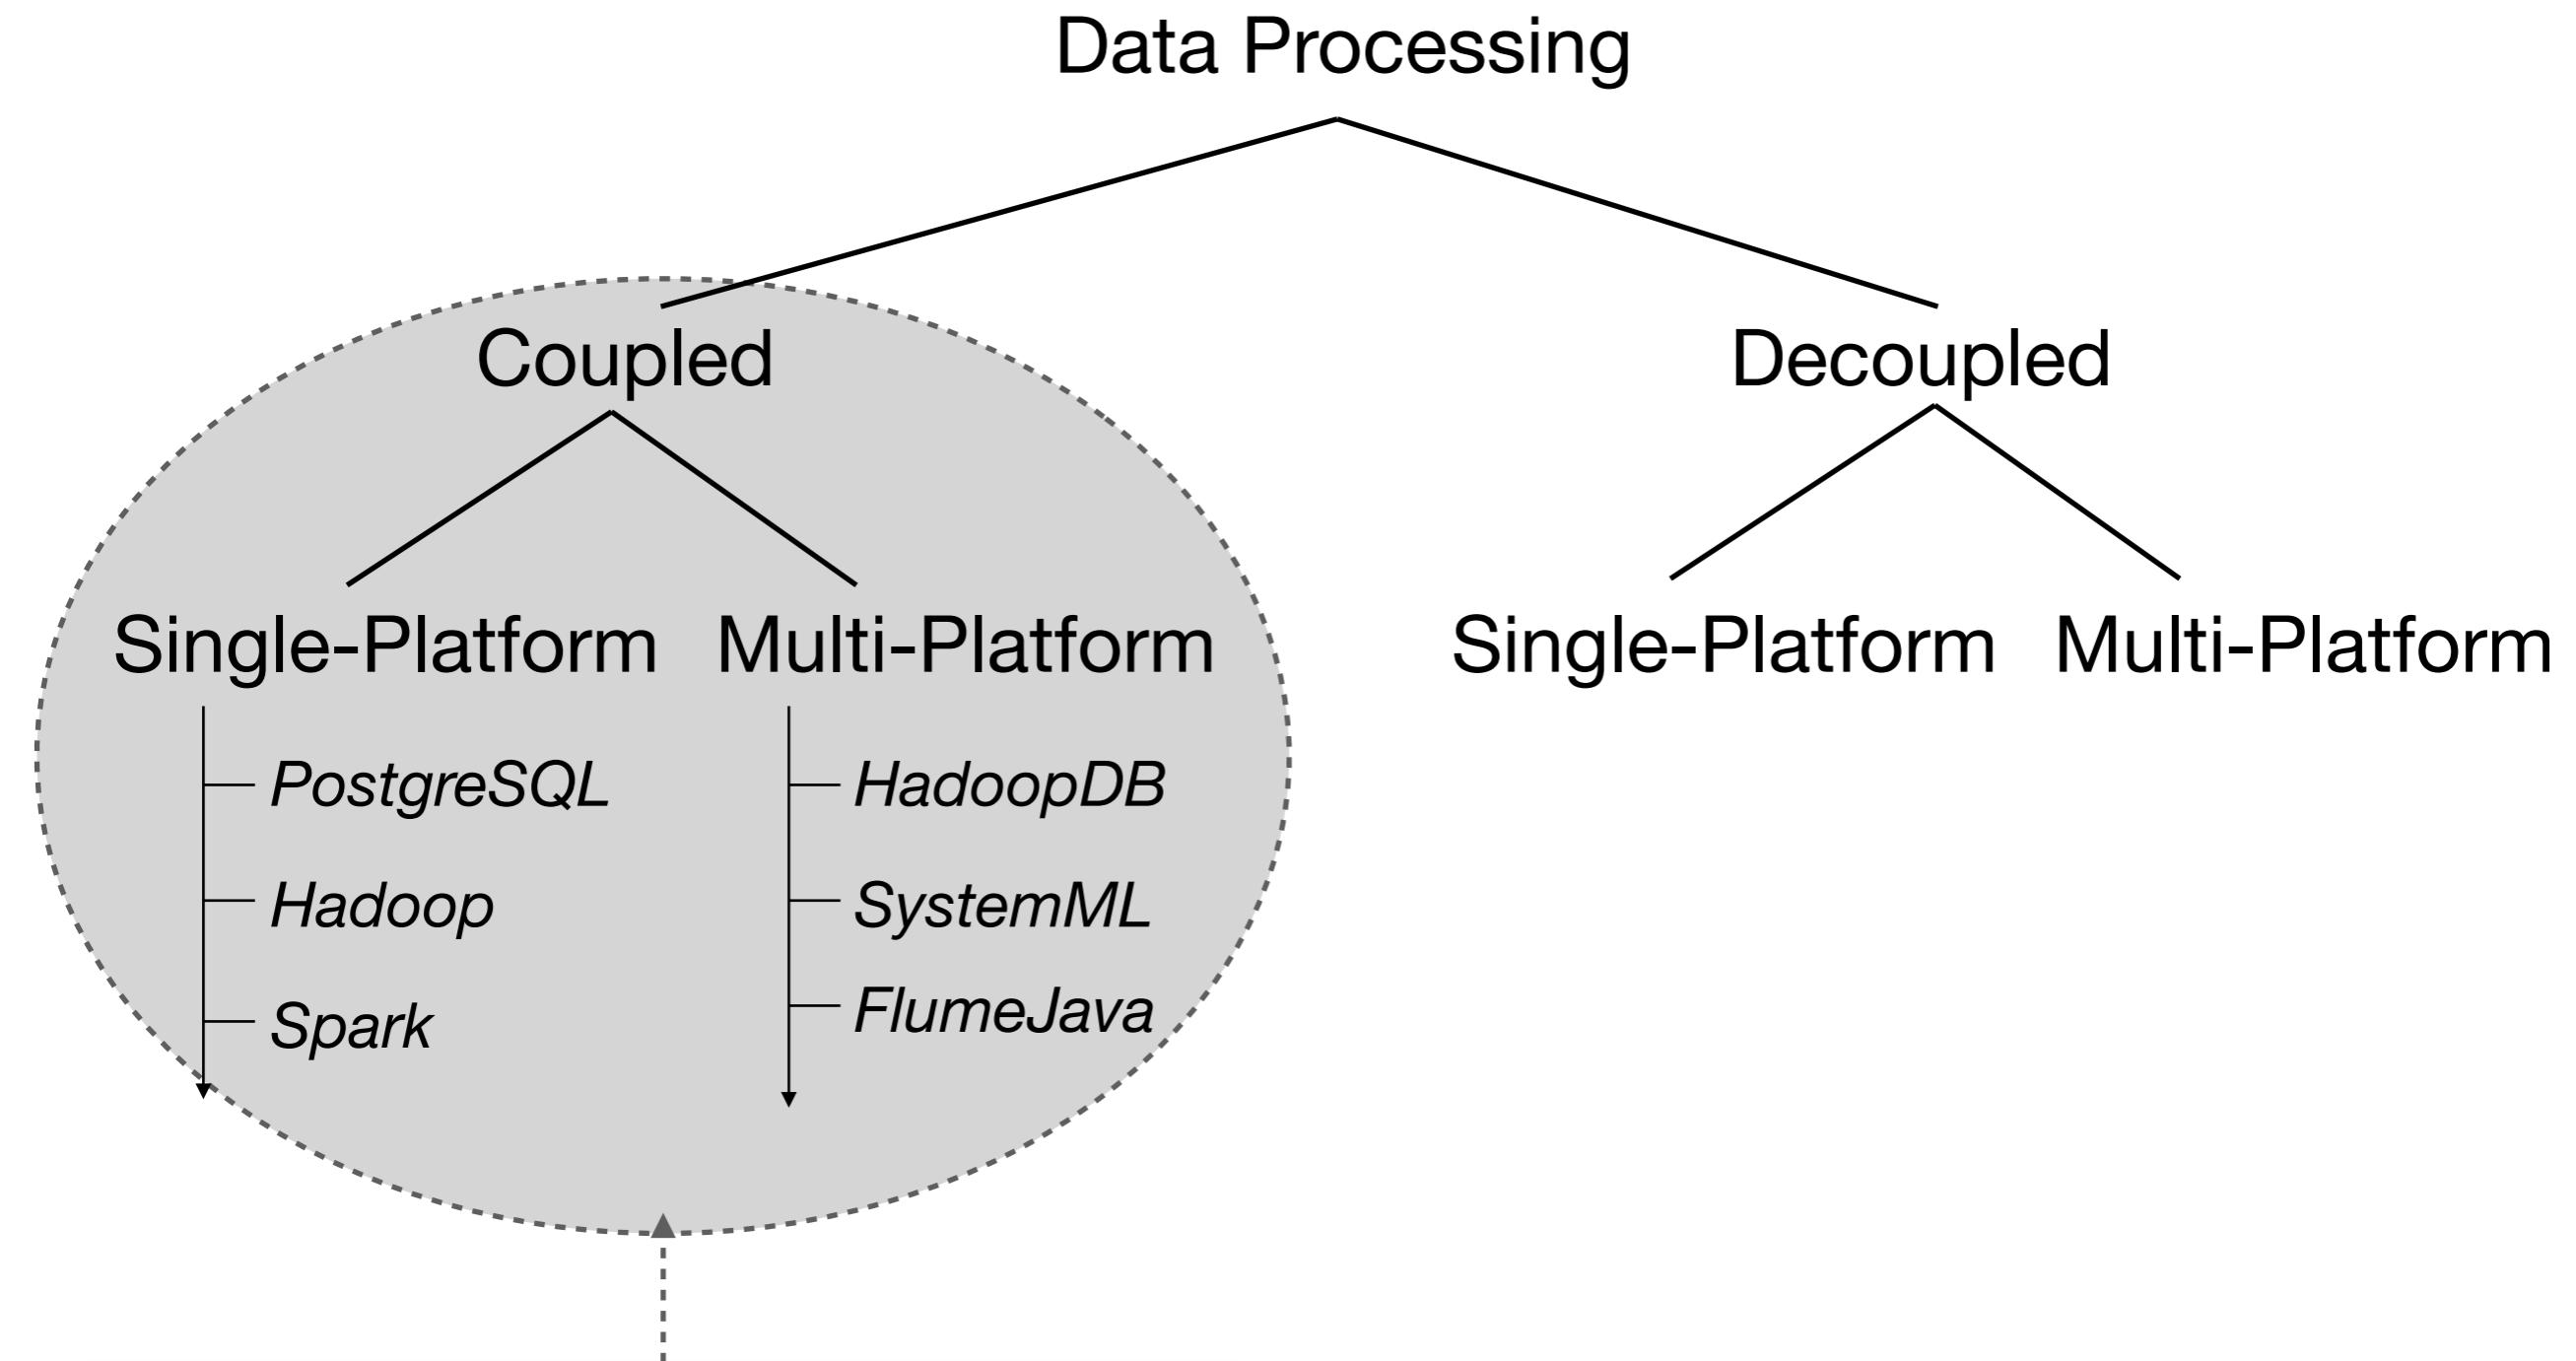
# Data Processing Taxonomy



# Data Processing Taxonomy

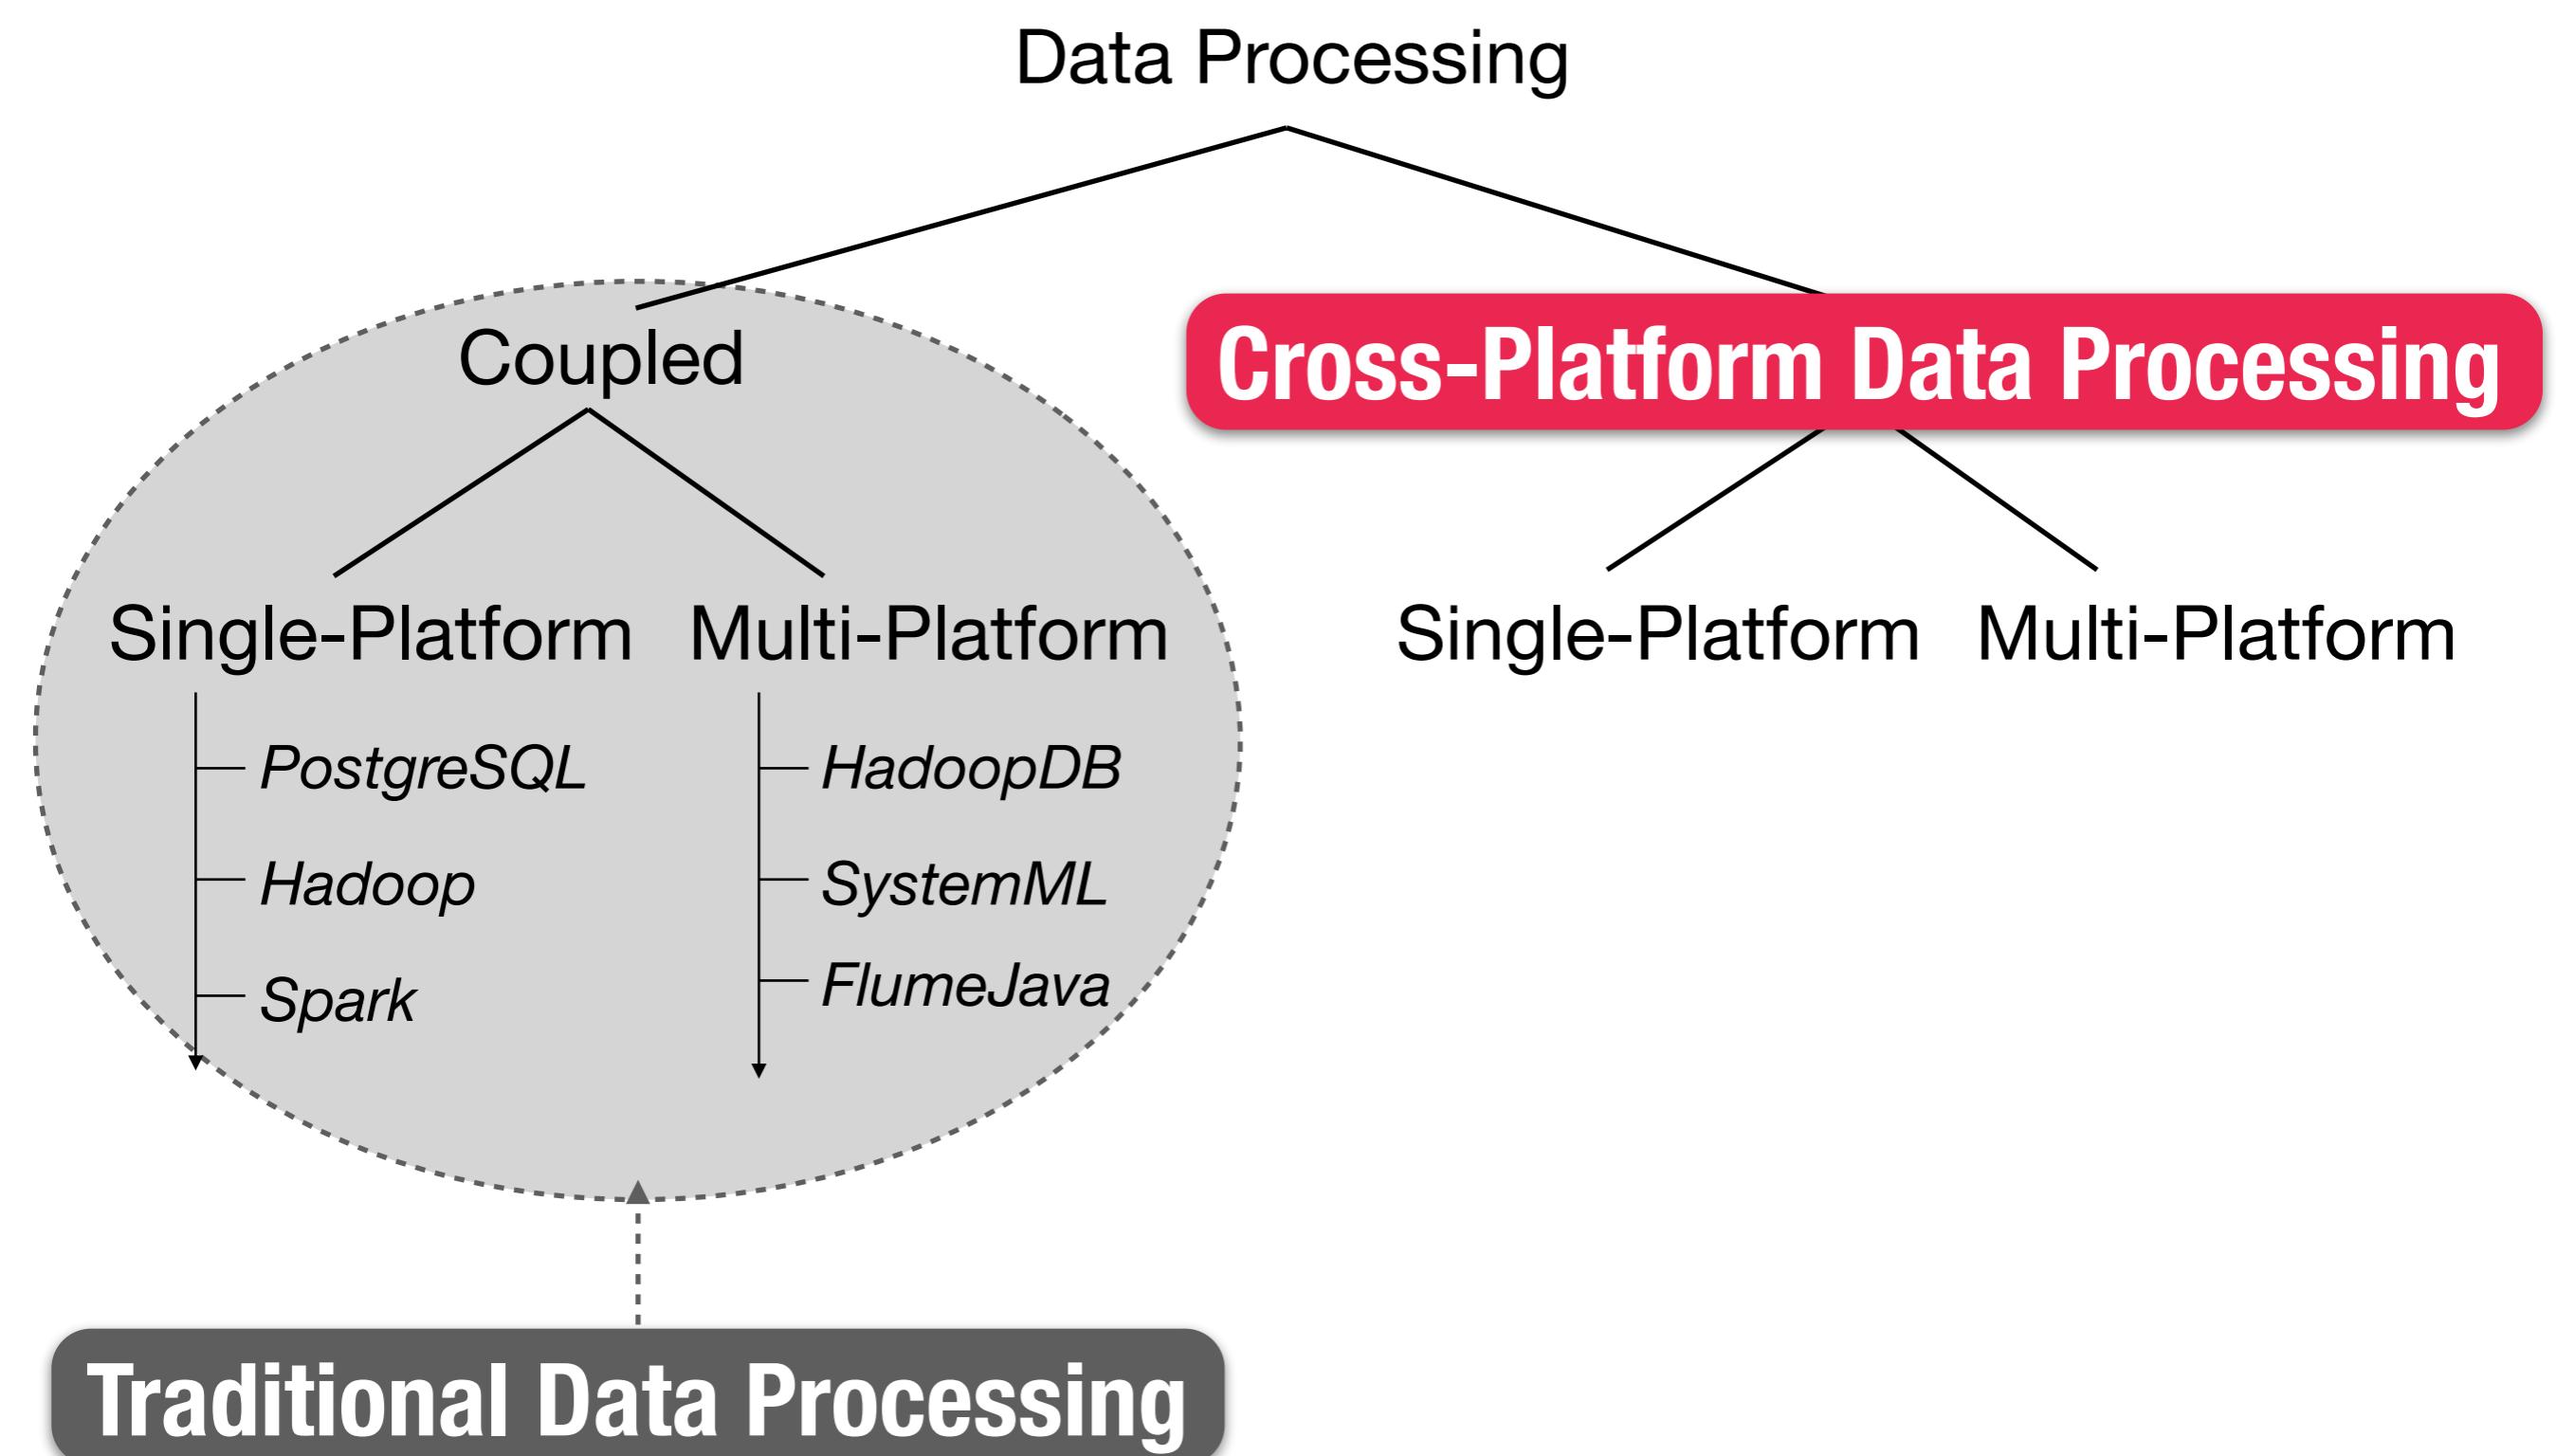


# Data Processing Taxonomy



Traditional Data Processing

# Data Processing Taxonomy



# Federated DBs vs Cross-Platform

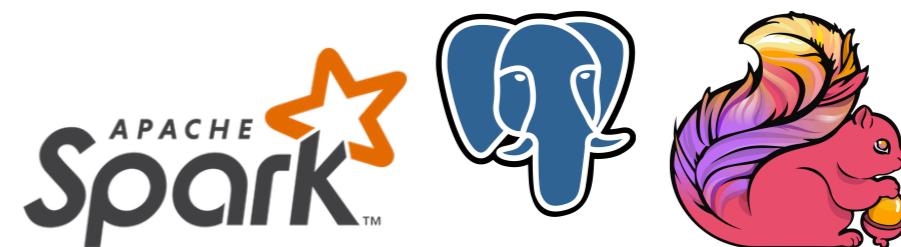
**Single Data Model**

**Federated DB System**



**Disparate Data Models**

**Cross-Platform System**



# Cross-Platform Data Processing

Motivation

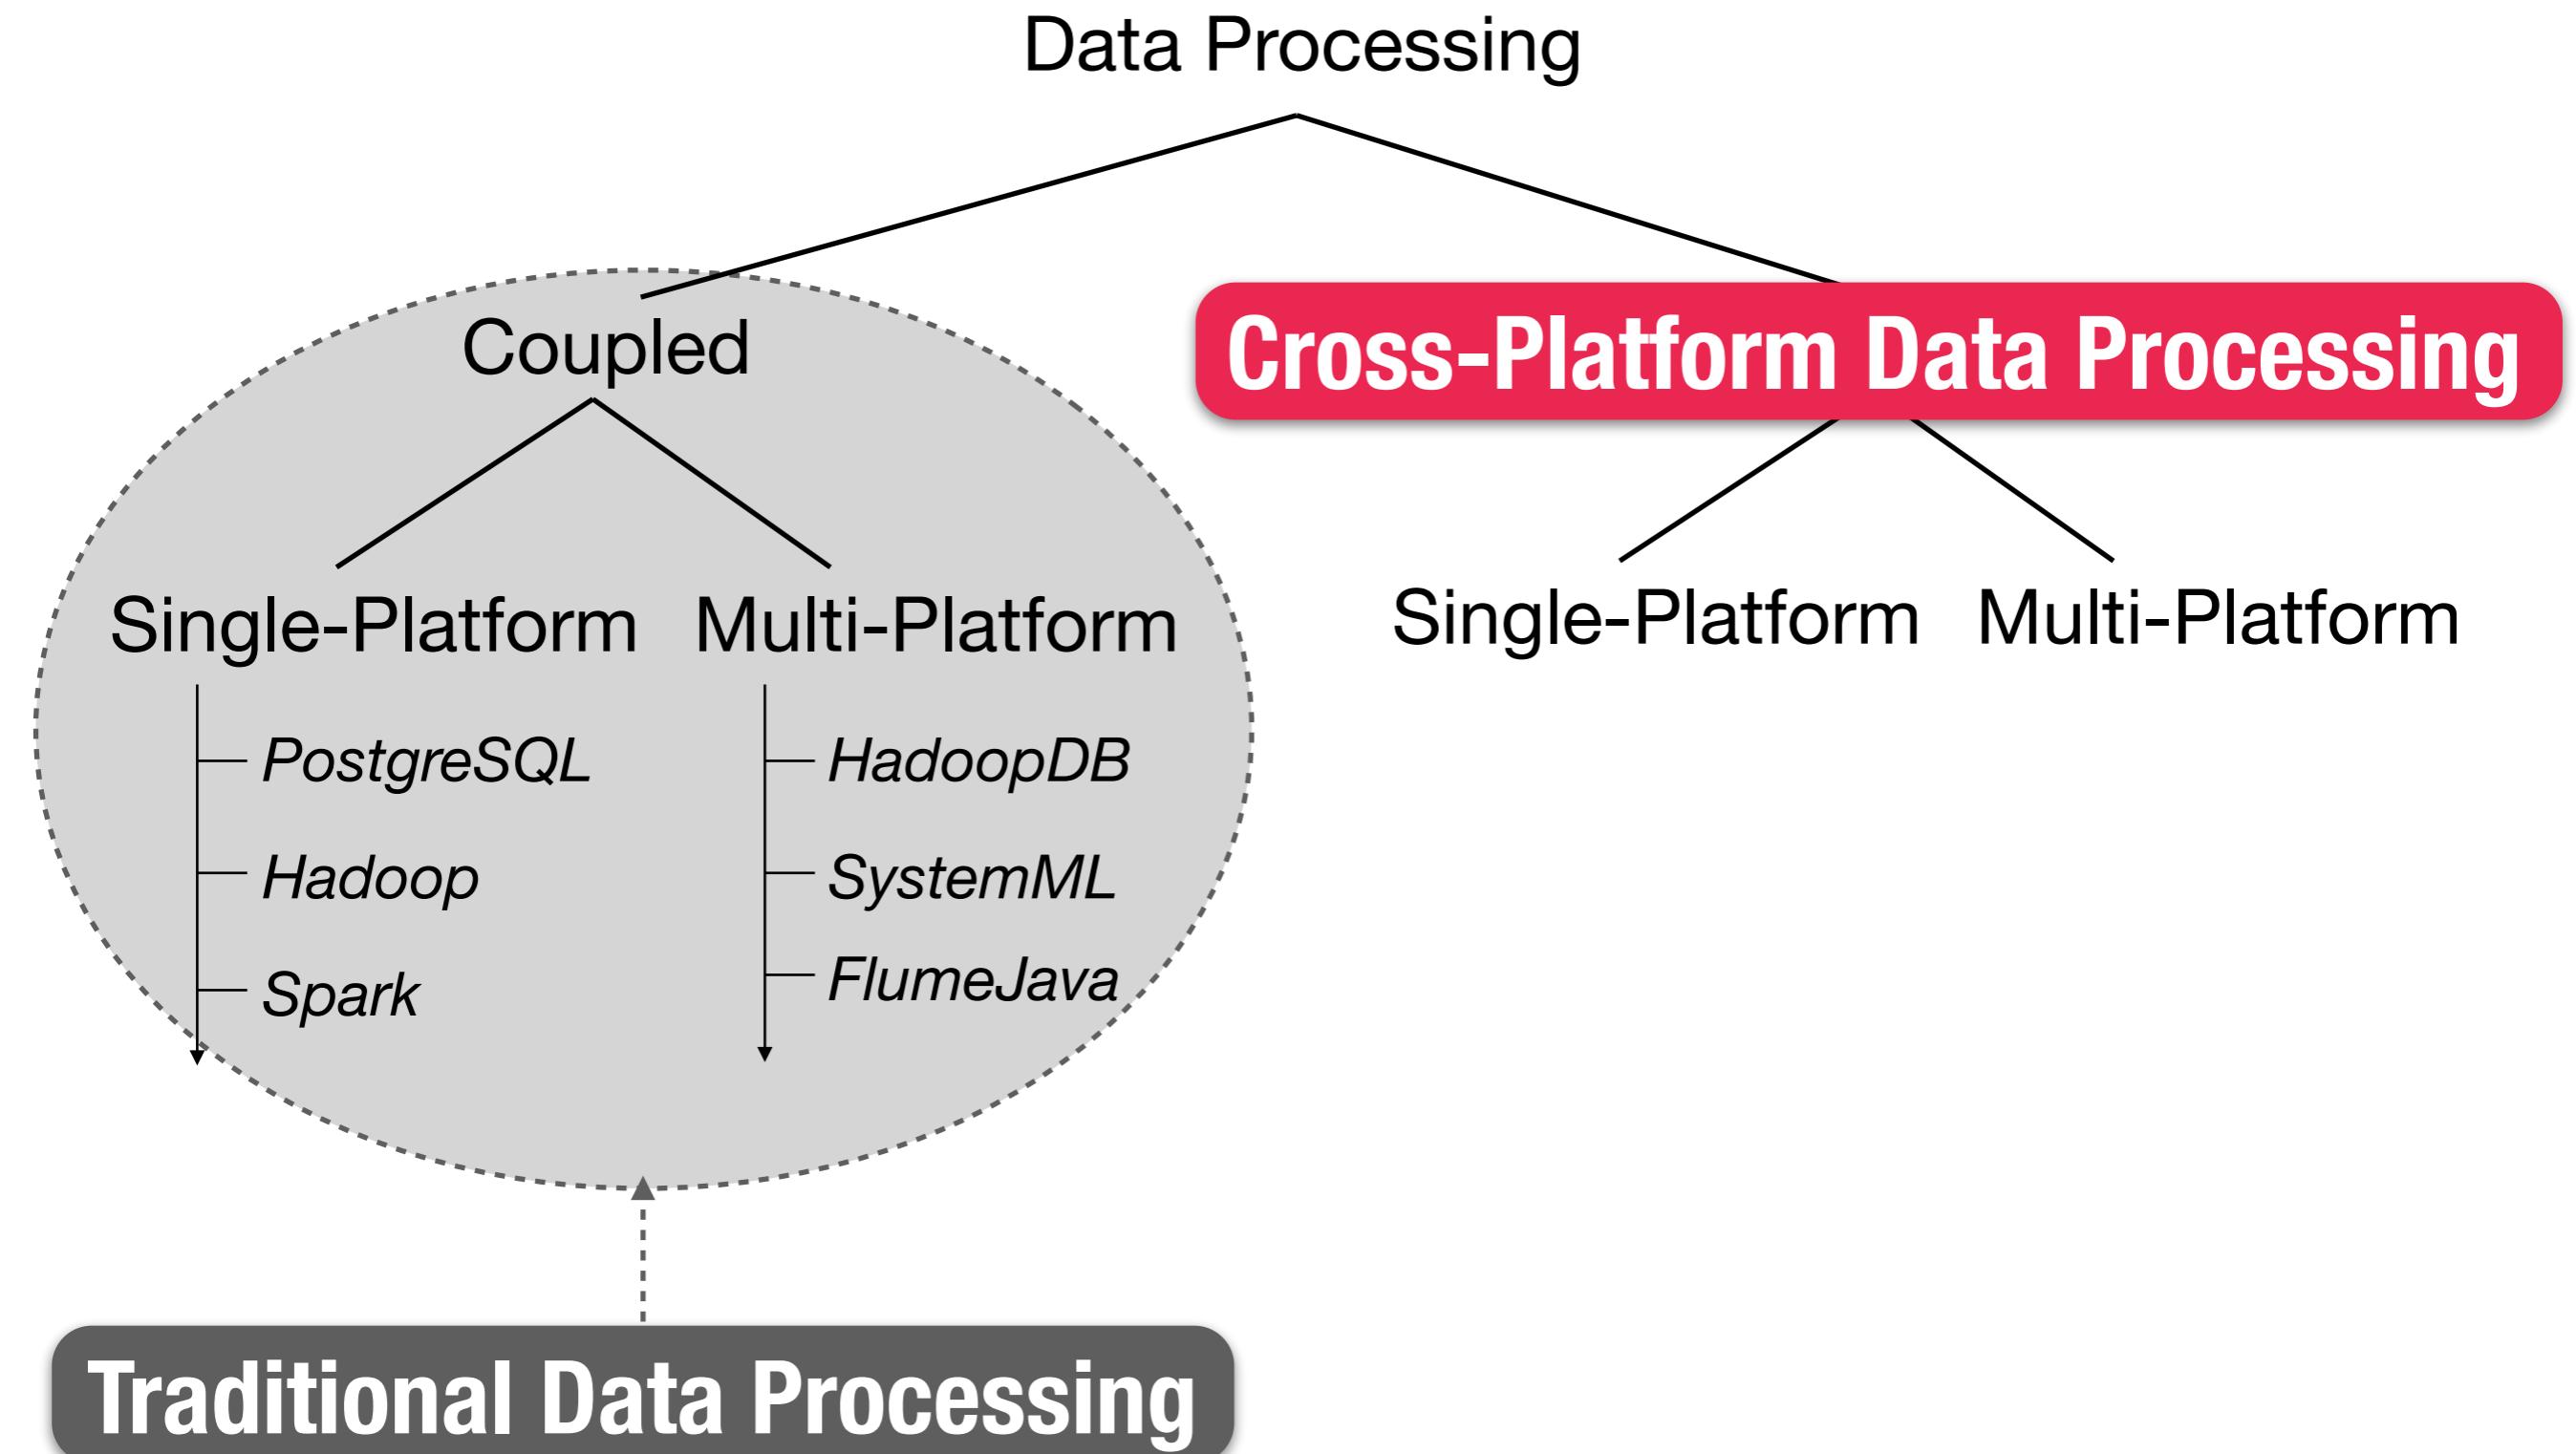
Summary

Use Cases

Current  
Efforts

Challenges

# Data Processing Taxonomy



# Cross-Platform Use Cases

Data Processing

**Cross-Platform Data Processing**

Single-Platform

Multi-Platform

# Cross-Platform Use Cases

Data Processing

## Cross-Platform Data Processing

### Decoupled Single-Platform

processing a query on *any single* data processing platform

Multi-Platform

# Cross-Platform Use Cases

Data Processing

## Cross-Platform Data Processing

Decoupled Single-Platform

Multi-Platform

processing ***one single*** query on  
***several*** data processing platforms...

# Cross-Platform Use Cases

Data Processing

## Cross-Platform Data Processing

Decoupled Single-Platform

Multi-Platform

processing ***one single*** query on  
***several*** data processing platforms...

Opportunistic

... to reduce the total cost of the input query

# Cross-Platform Use Cases

Data Processing

## Cross-Platform Data Processing

Decoupled Single-Platform

Multi-Platform

processing ***one single*** query on  
***several*** data processing platforms...

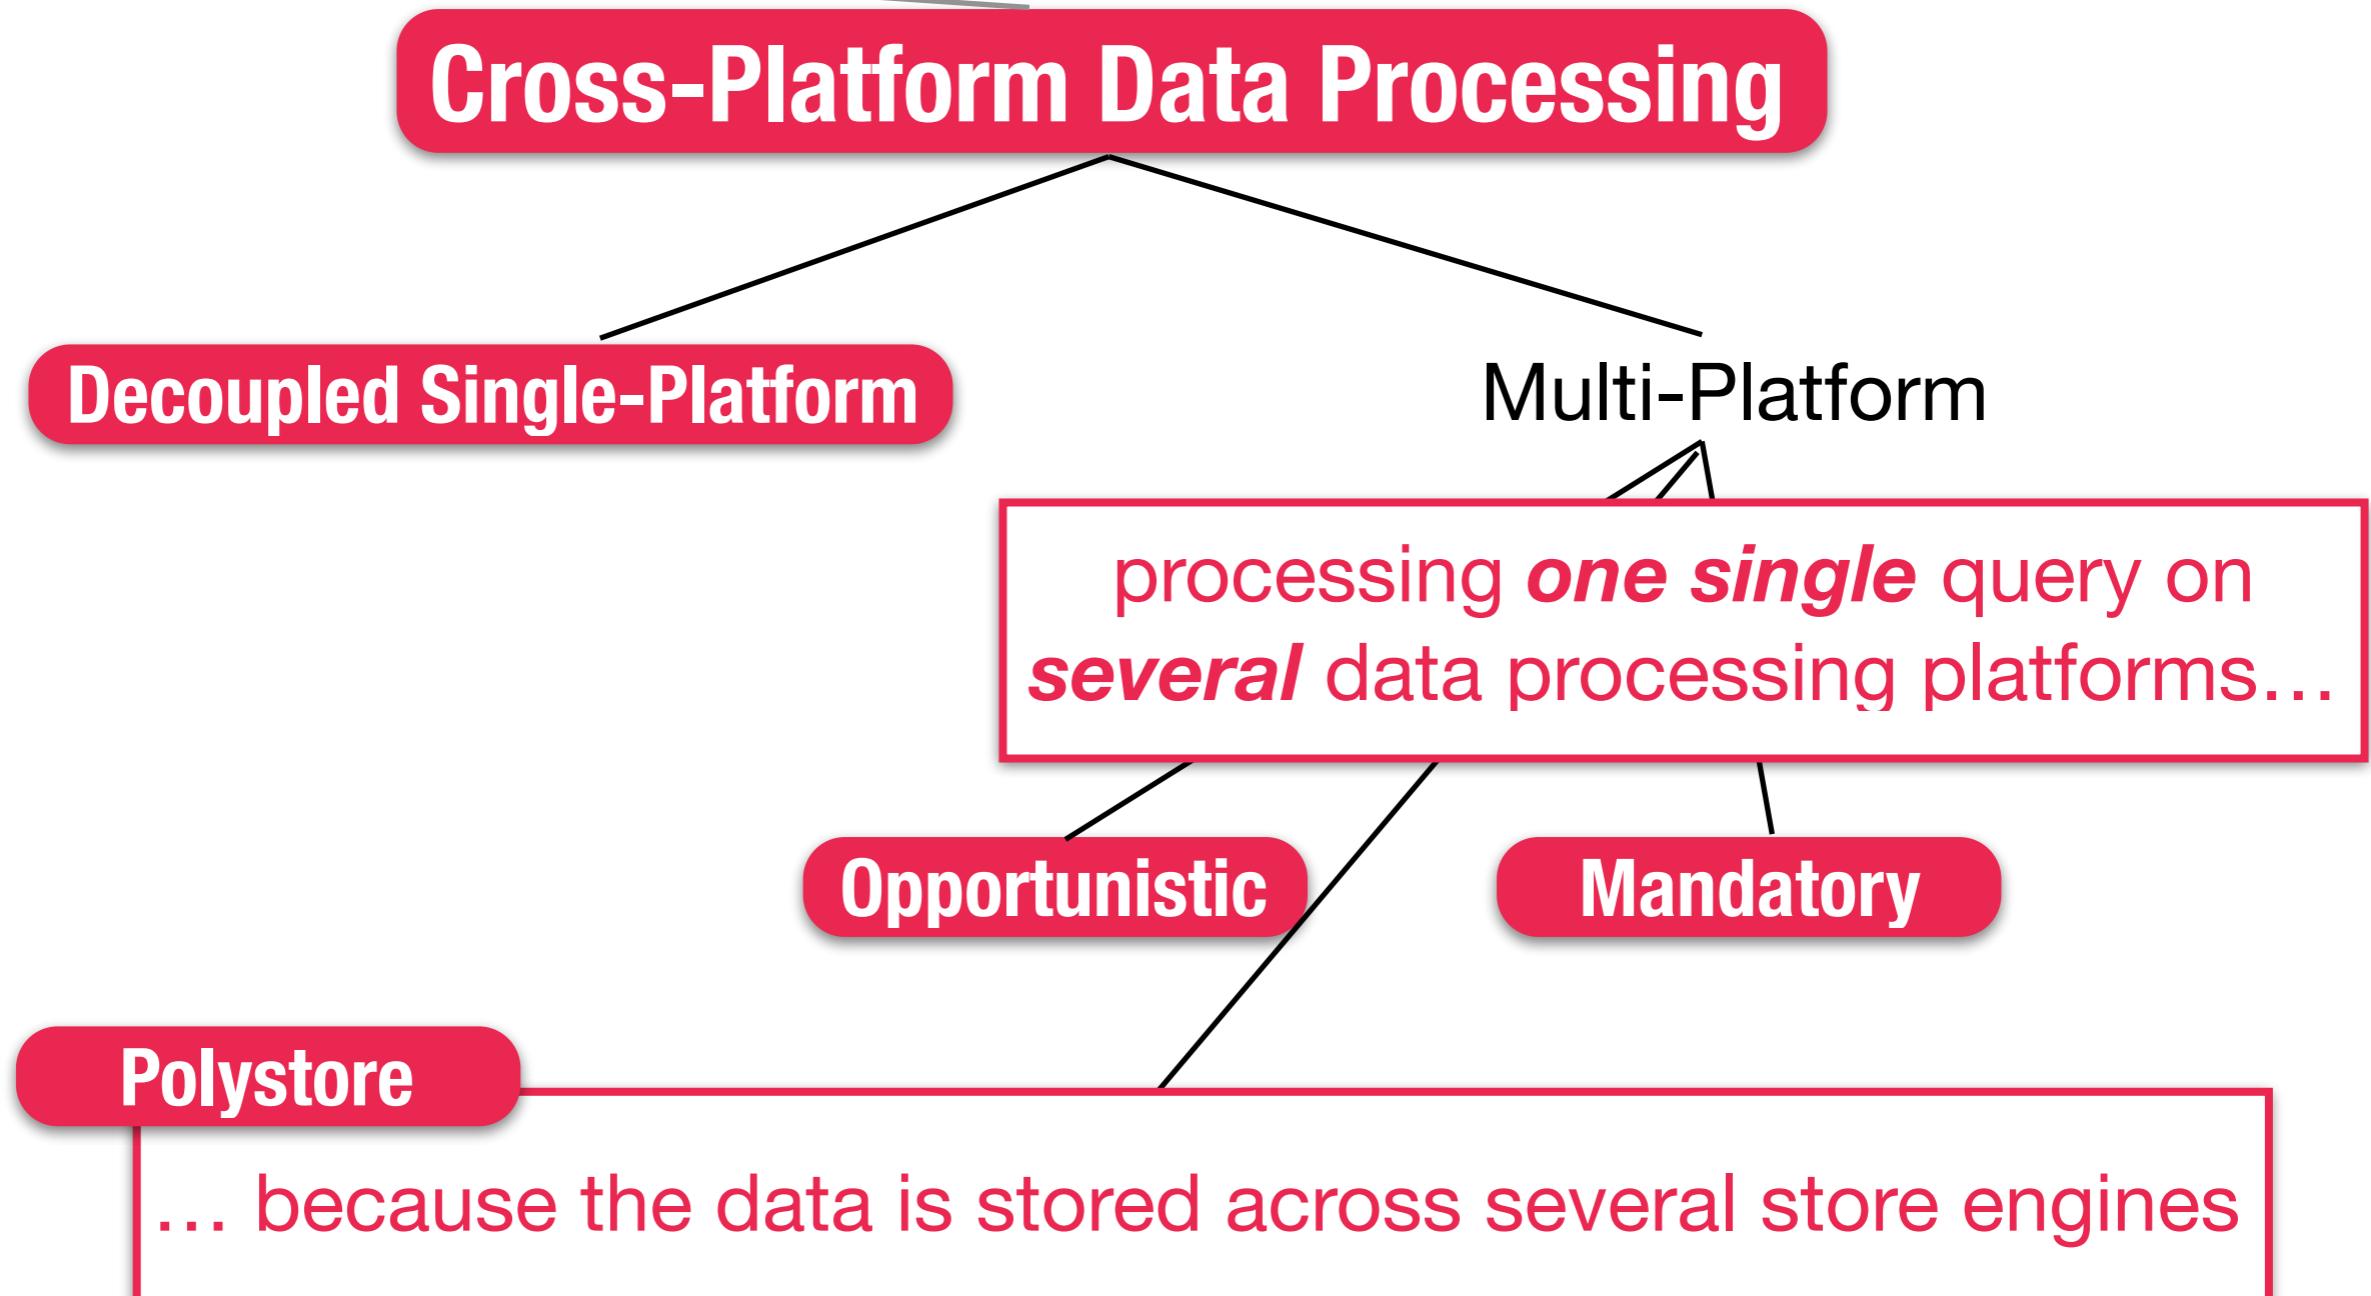
Opportunistic

Mandatory

... because the platform where the data is cannot run it entirely

# Cross-Platform Use Cases

Data Processing



# Cross-Platform Use Cases

Data Processing

**Cross-Platform Data Processing**

**Decoupled Single-Platform**

Multi-Platform

Opportunistic

Mandatory

Polystore

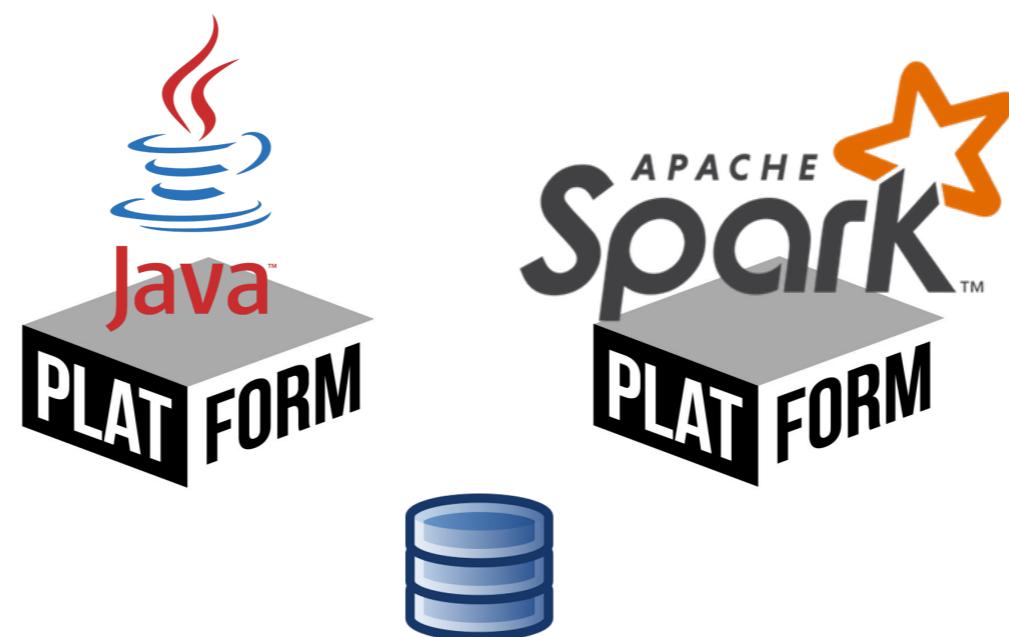
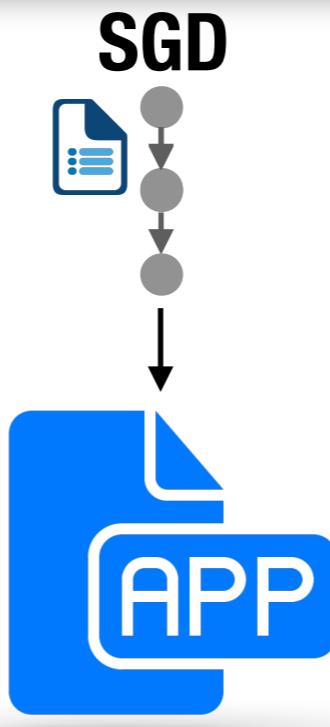
# Decoupled Single-Platform

using *any single* data processing platform to process a query



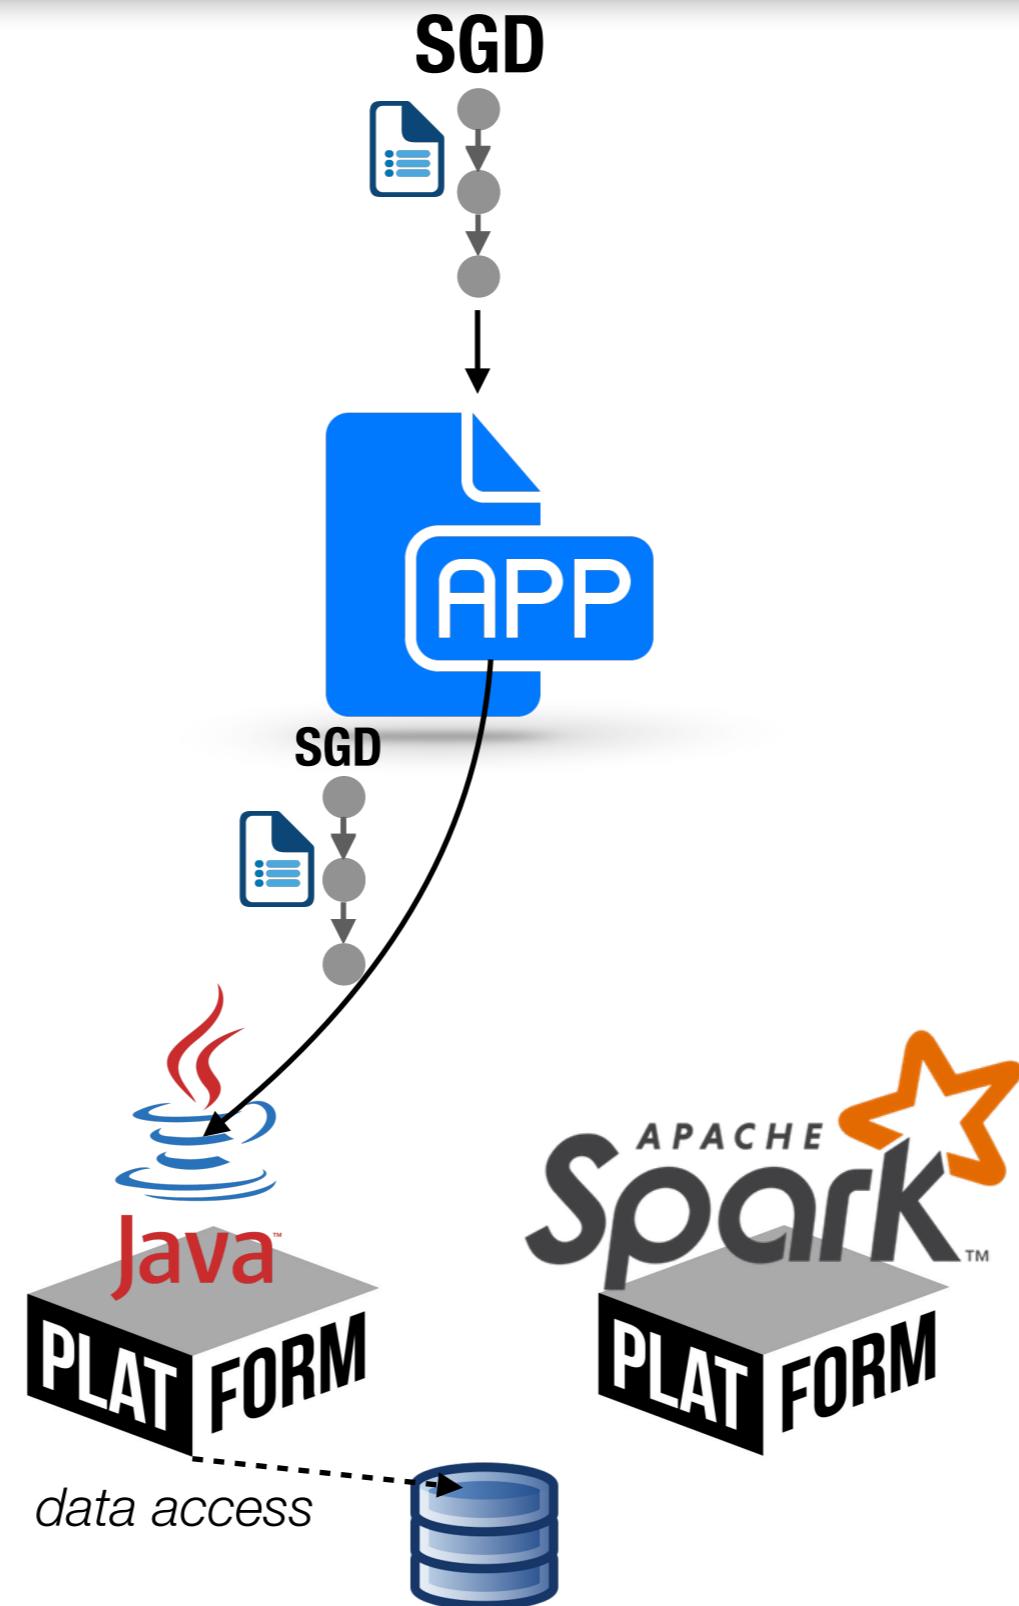
# Decoupled Single-Platform

using *any single* data processing platform to process a query



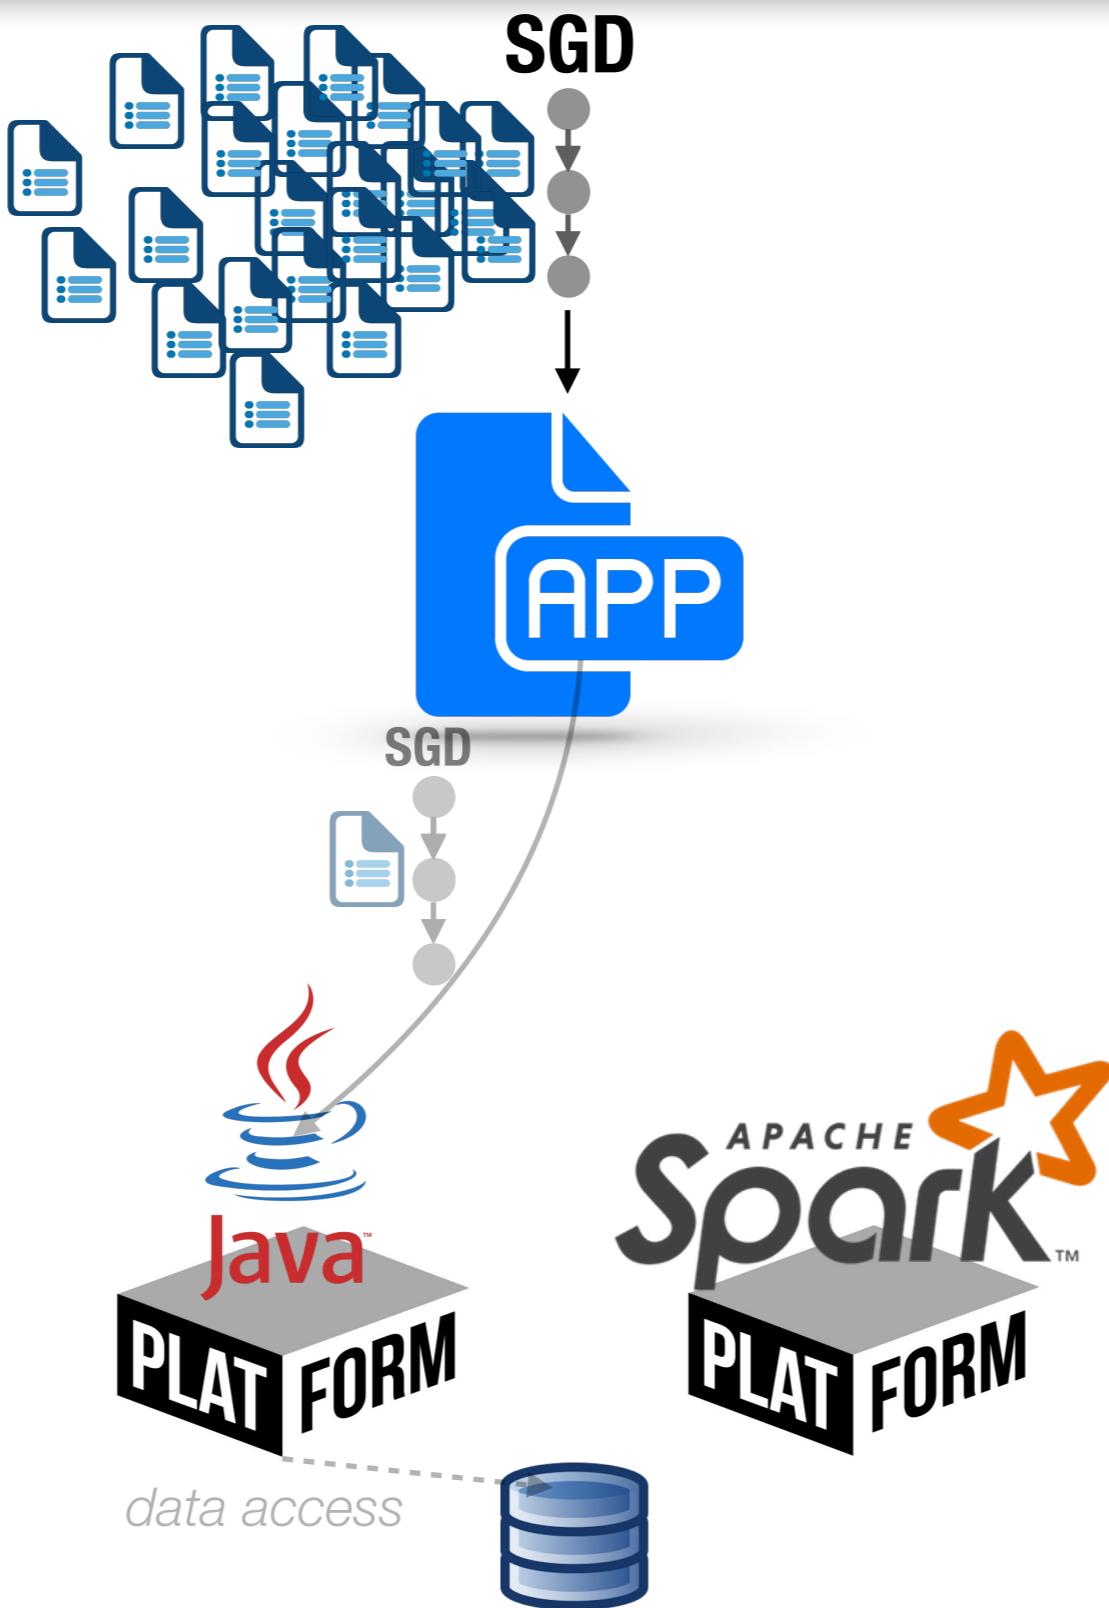
# Decoupled Single-Platform

using *any single* data processing platform to process a query



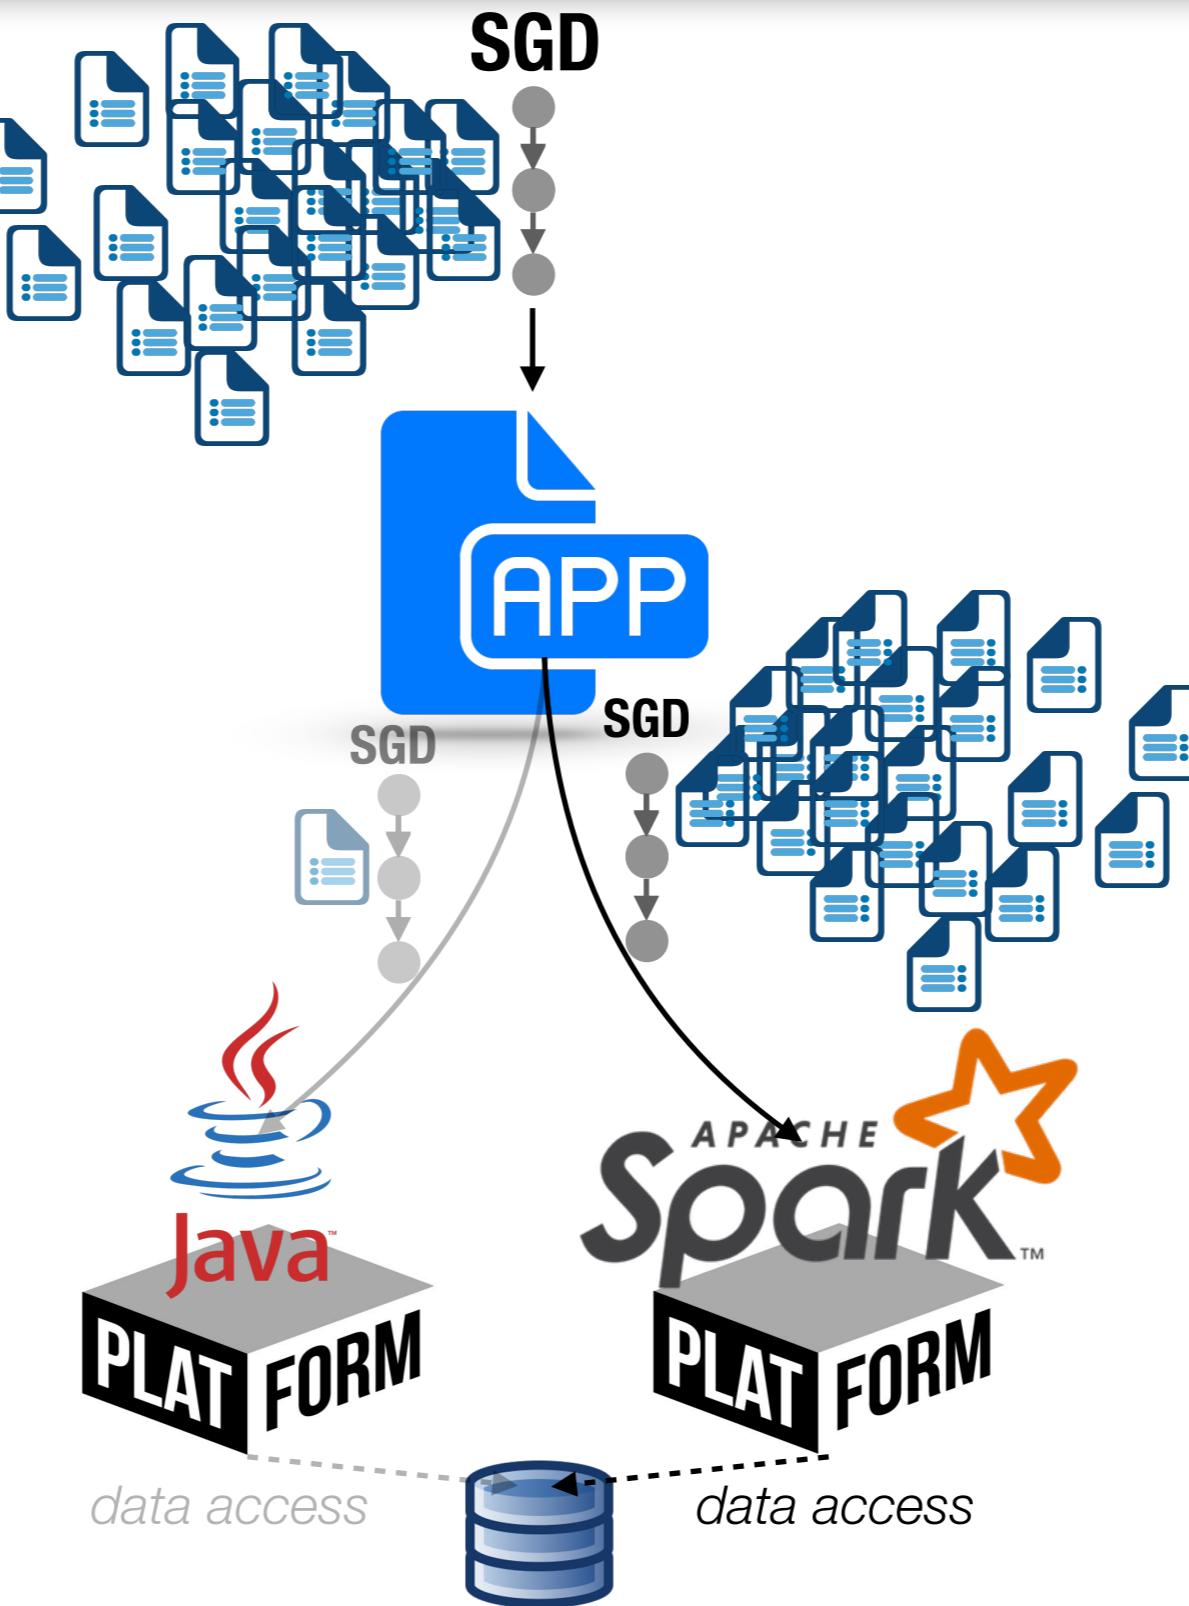
# Decoupled Single-Platform

using *any single* data processing platform to process a query



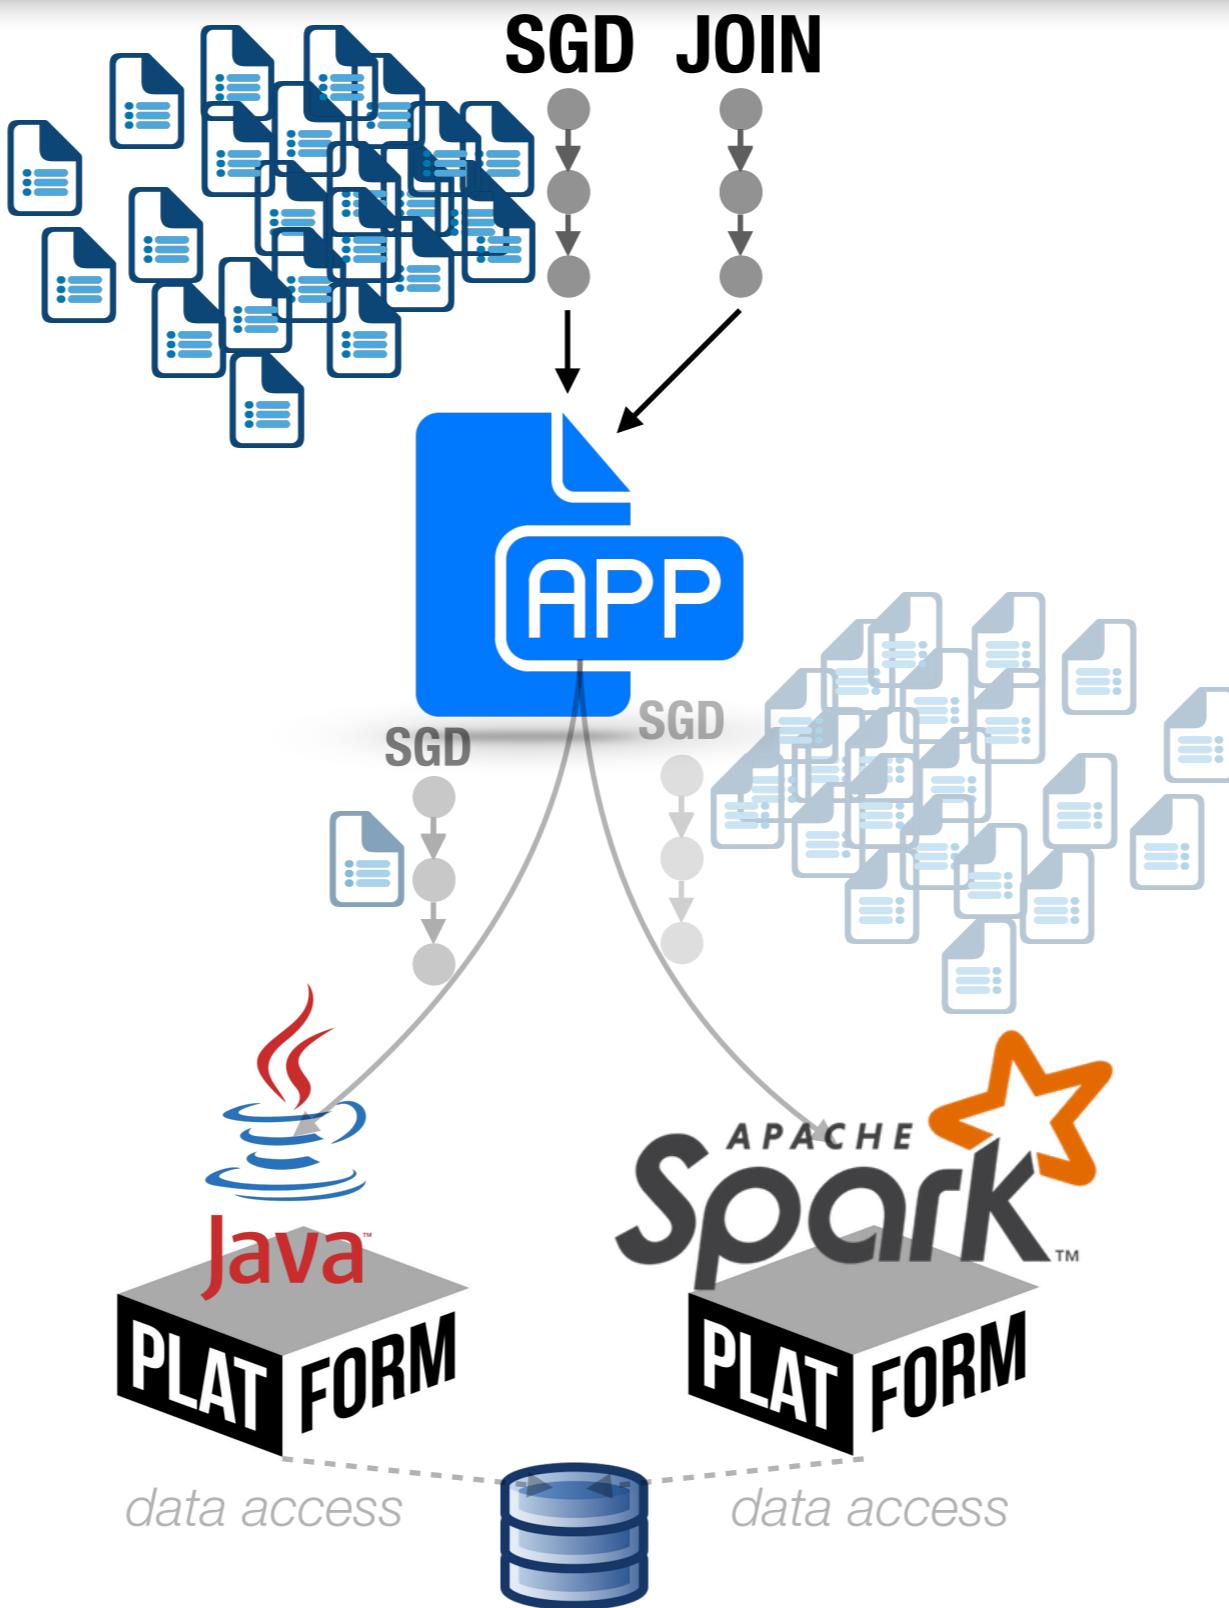
# Decoupled Single-Platform

using *any single* data processing platform to process a query



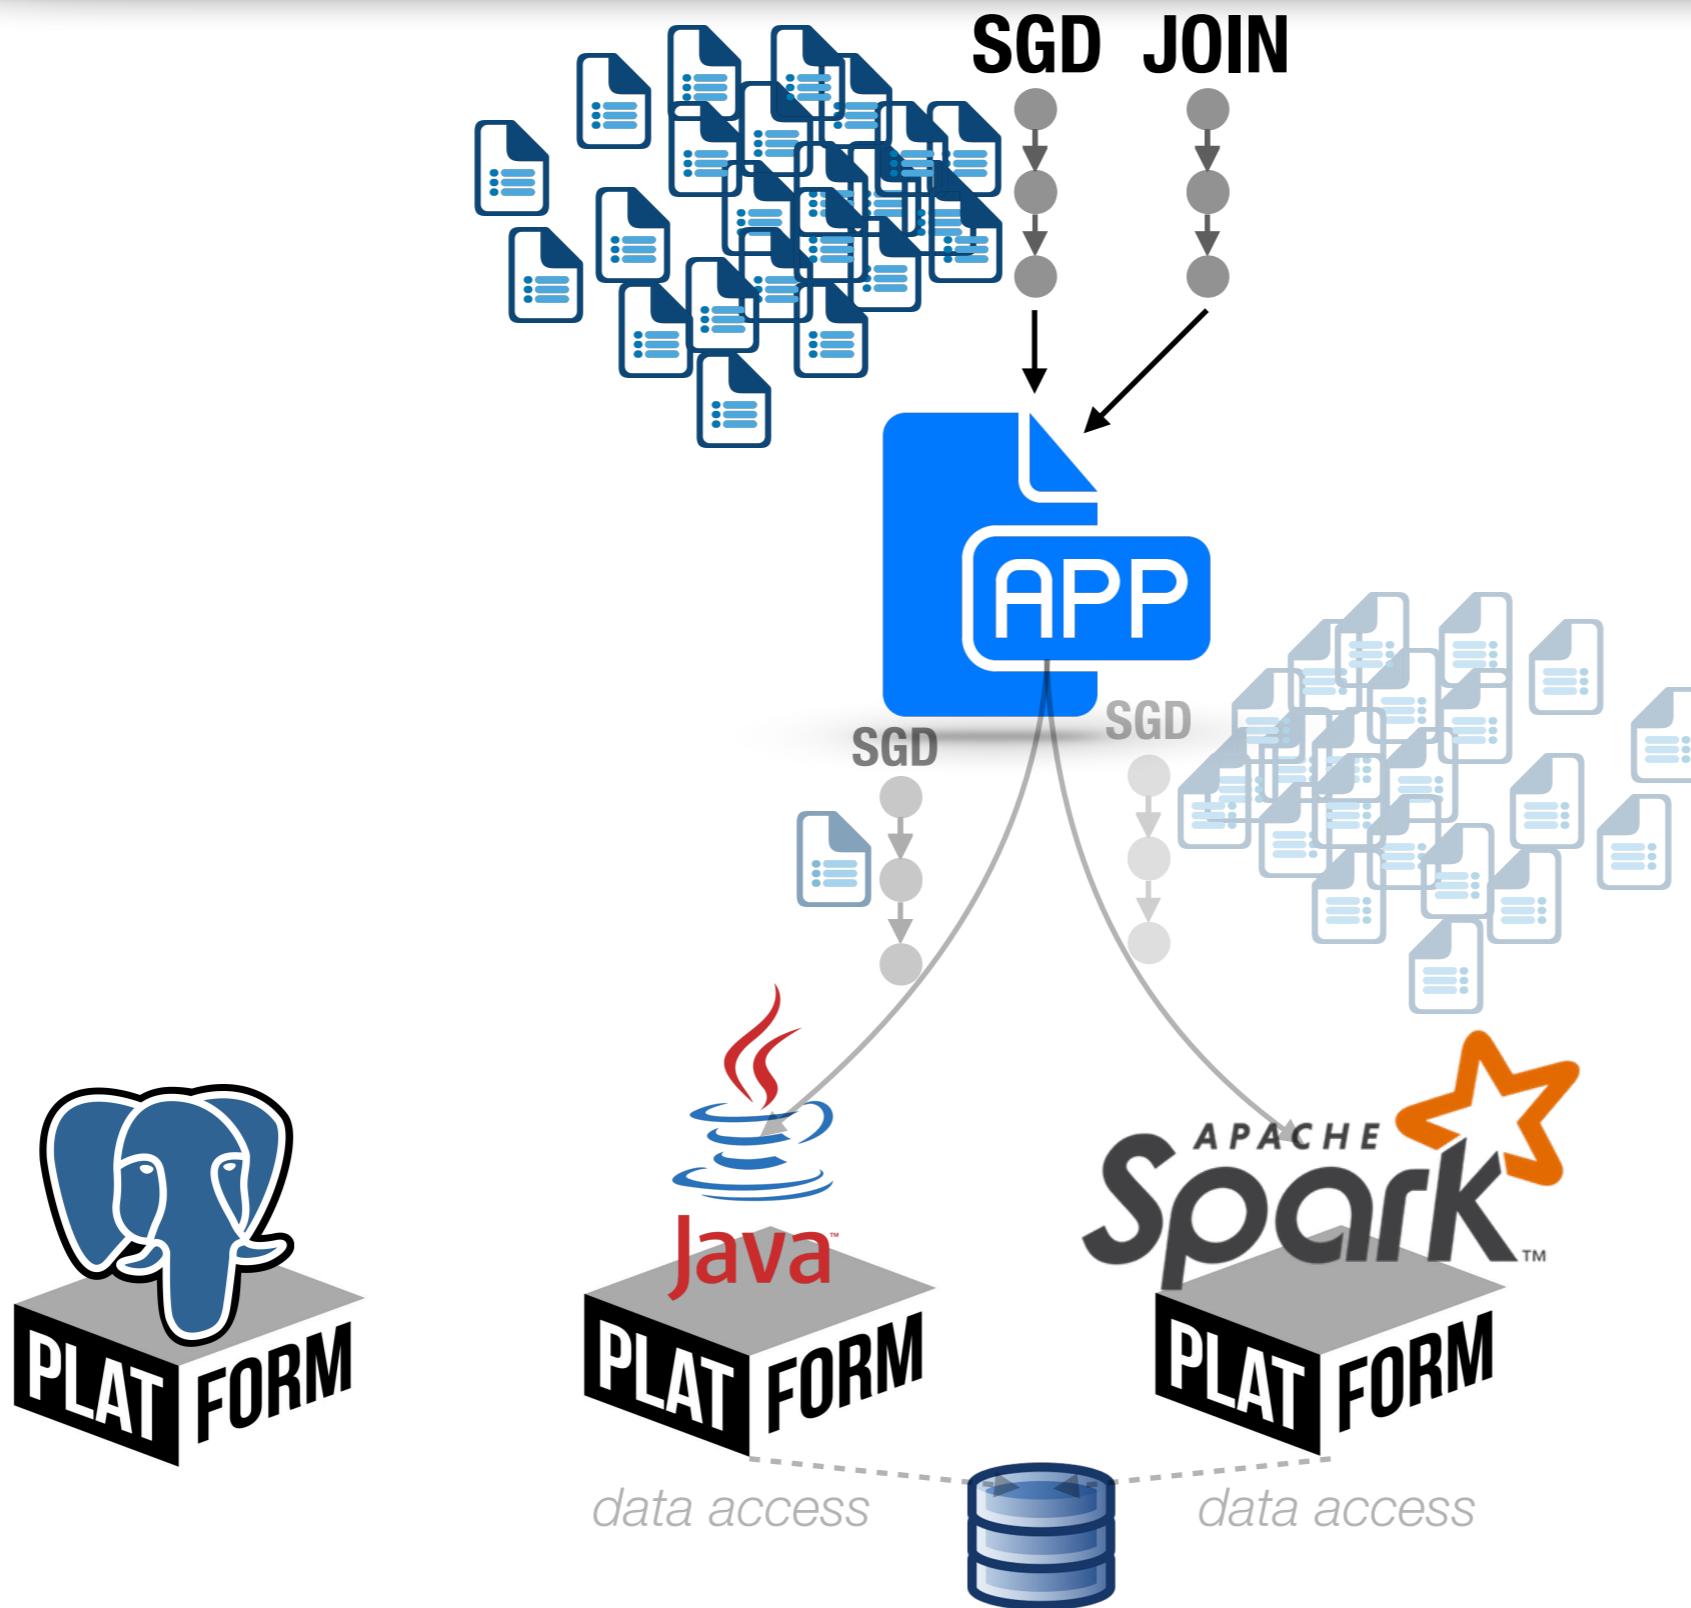
# Decoupled Single-Platform

using *any single* data processing platform to process a query



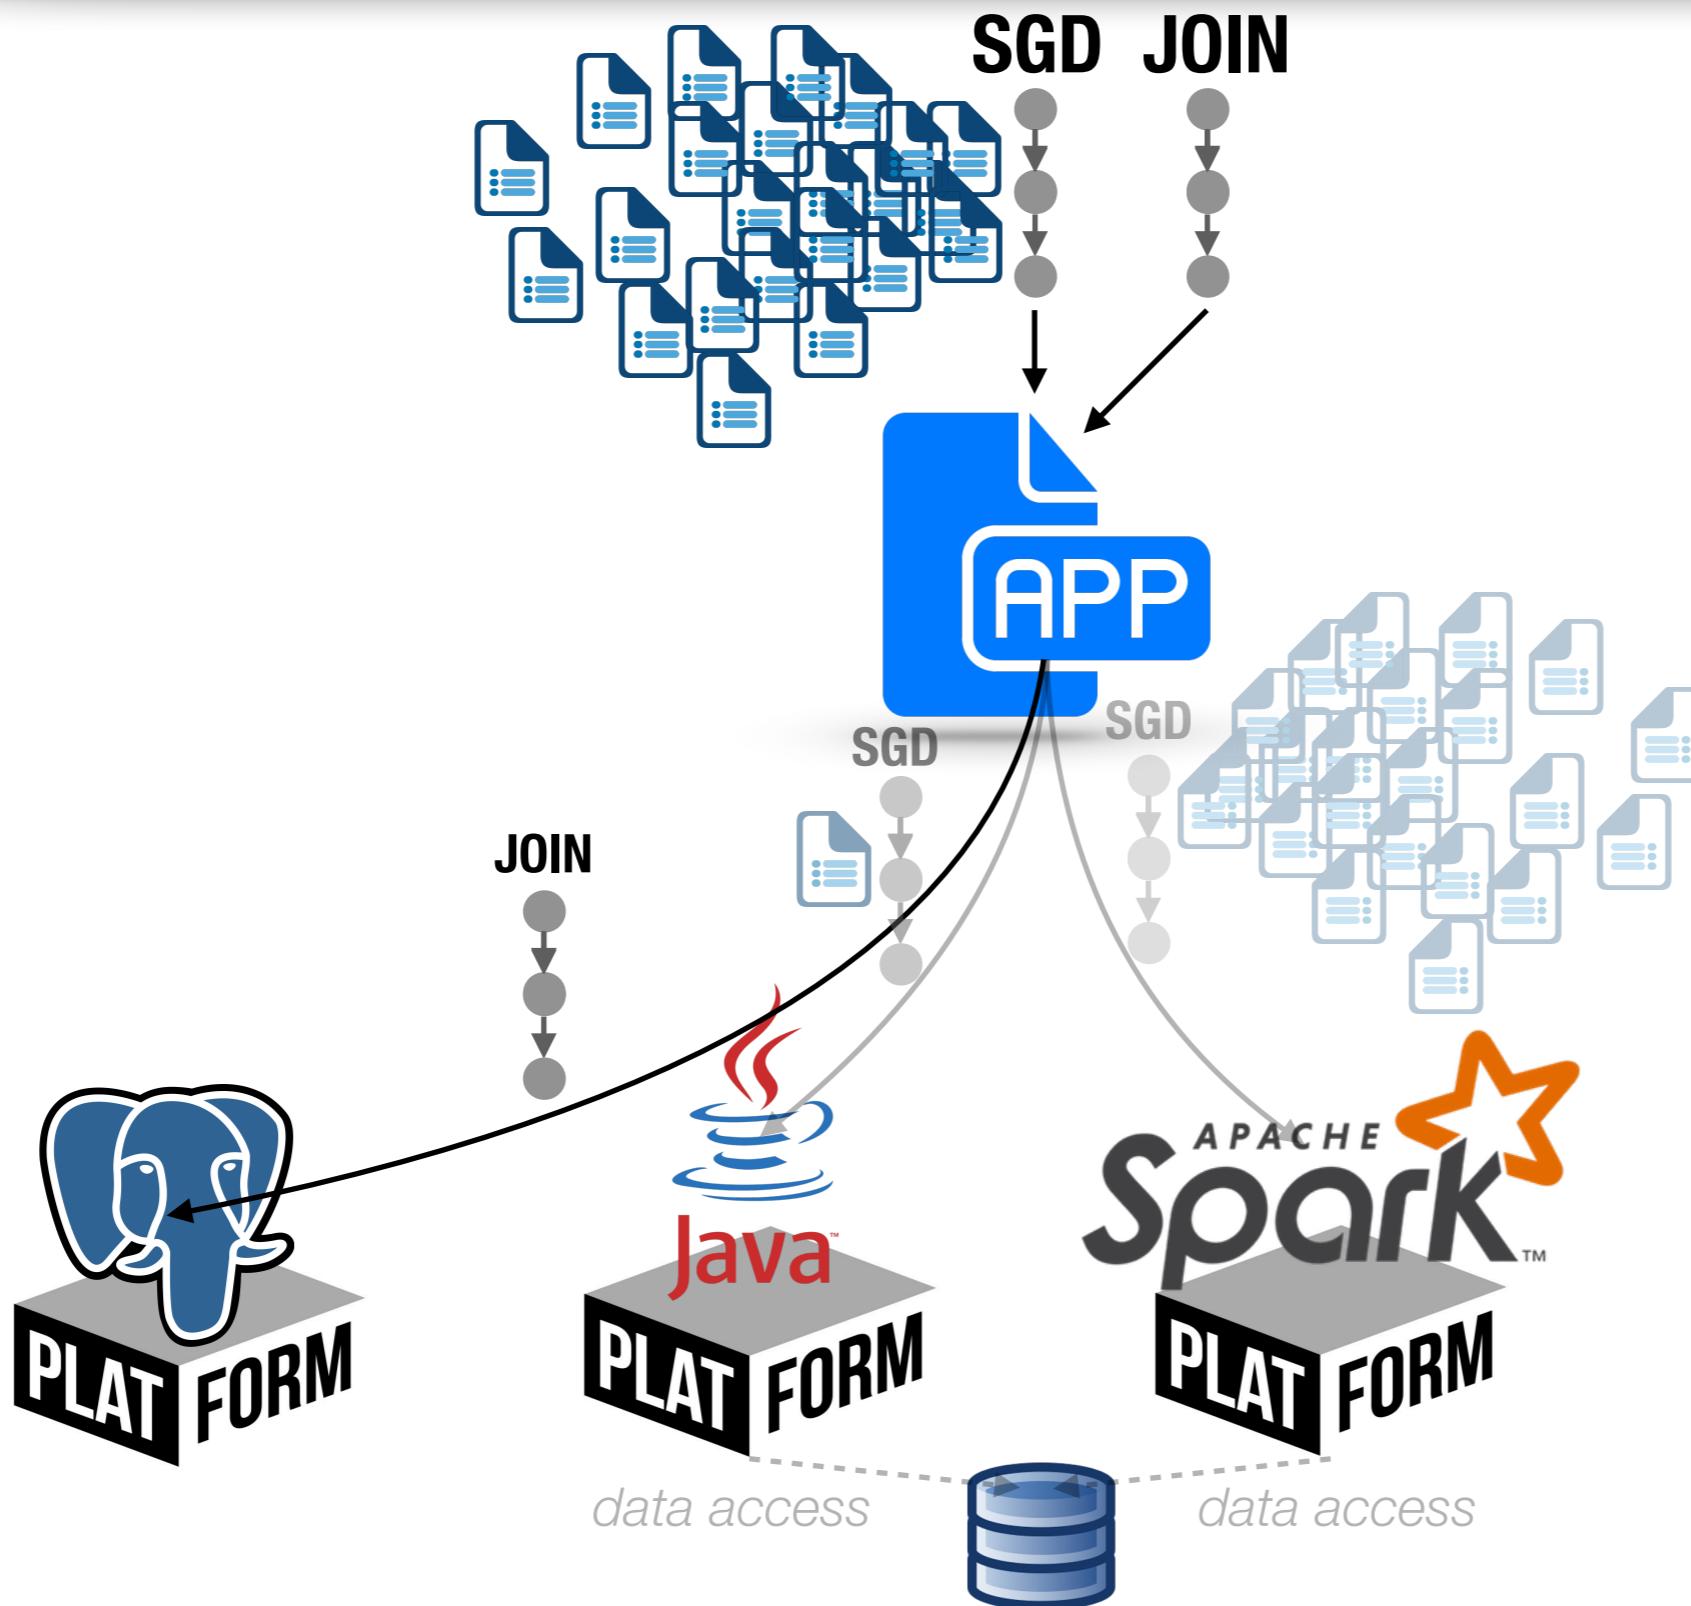
# Decoupled Single-Platform

using *any single* data processing platform to process a query



# Decoupled Single-Platform

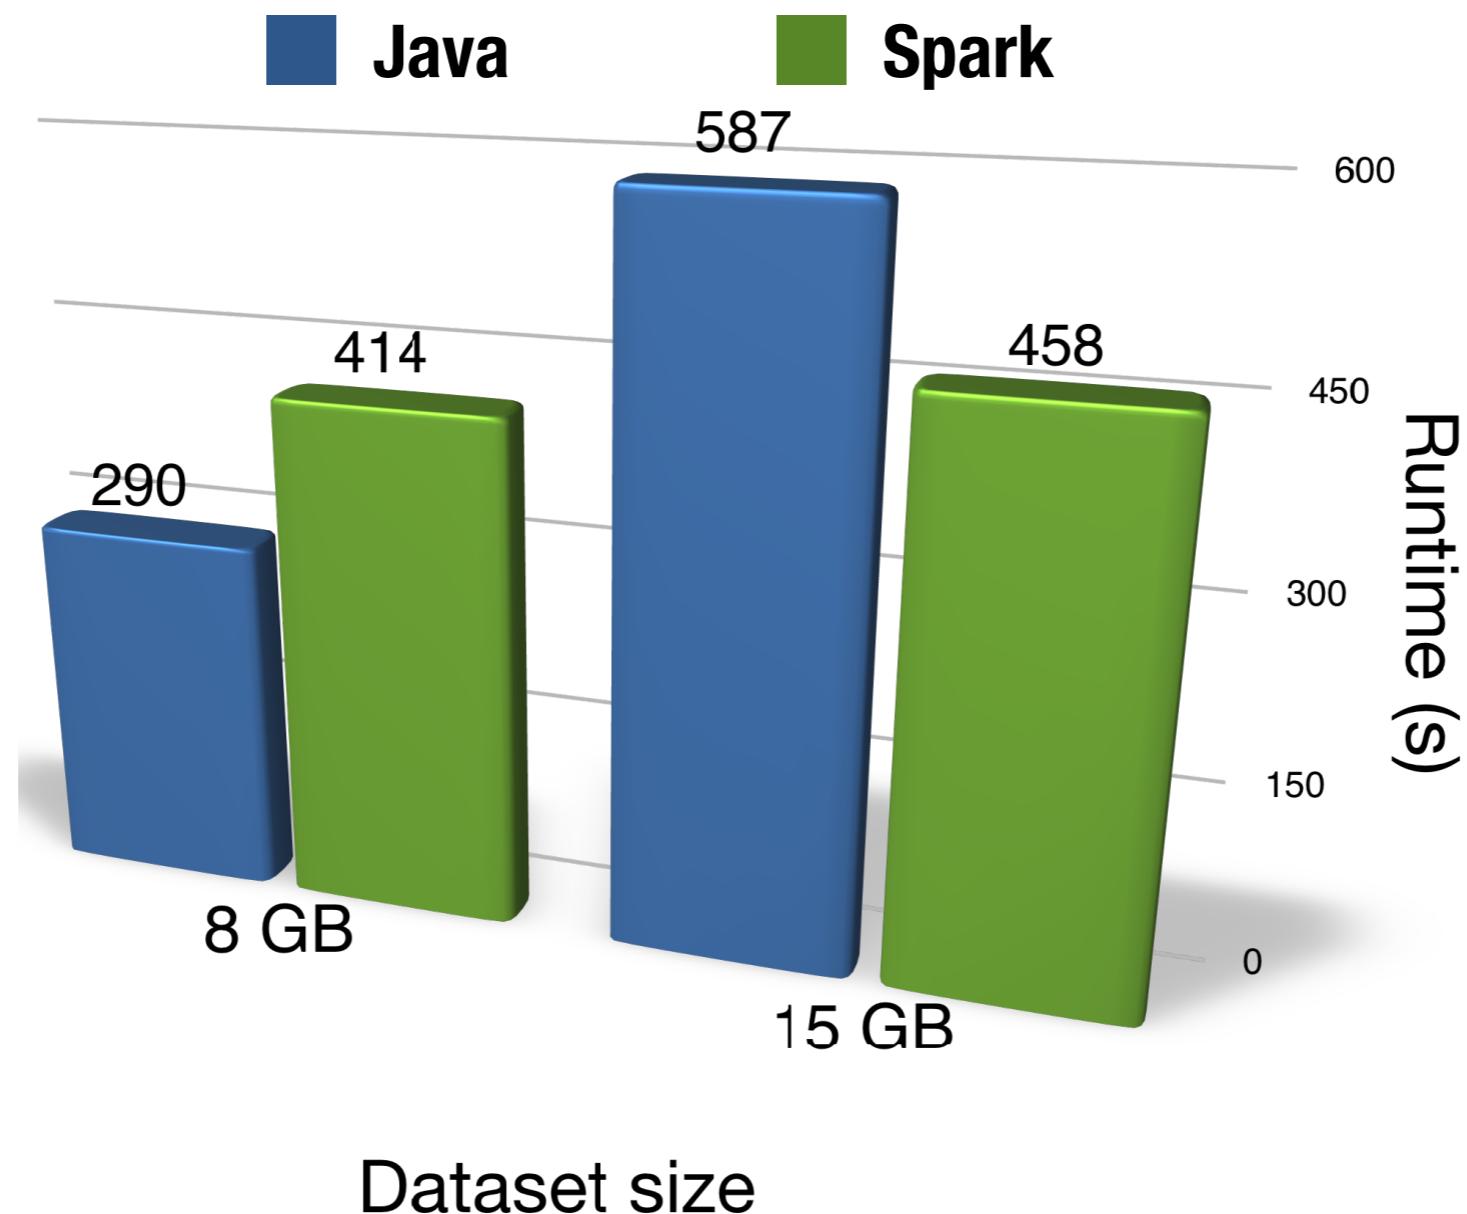
using *any single* data processing platform to process a query



# Decoupled Single-Platform

using *any single* data processing platform to process a query

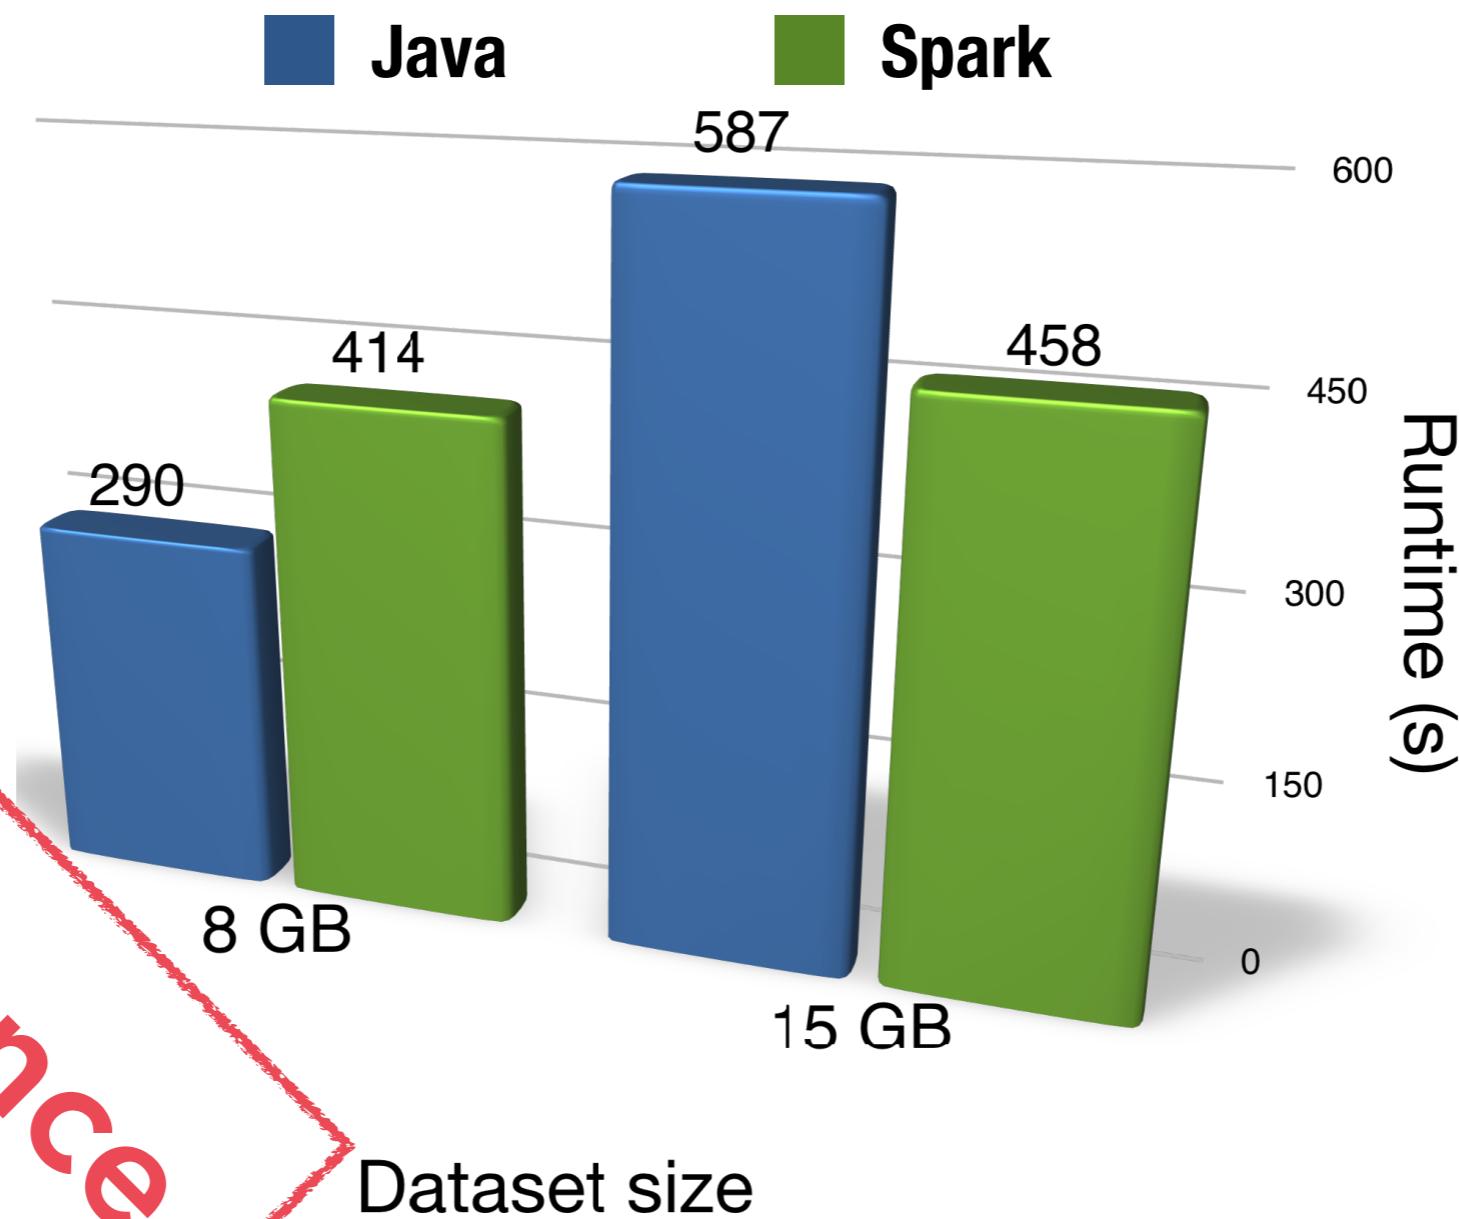
## SGD



# Decoupled Single-Platform

using *any single* data processing platform to process a query

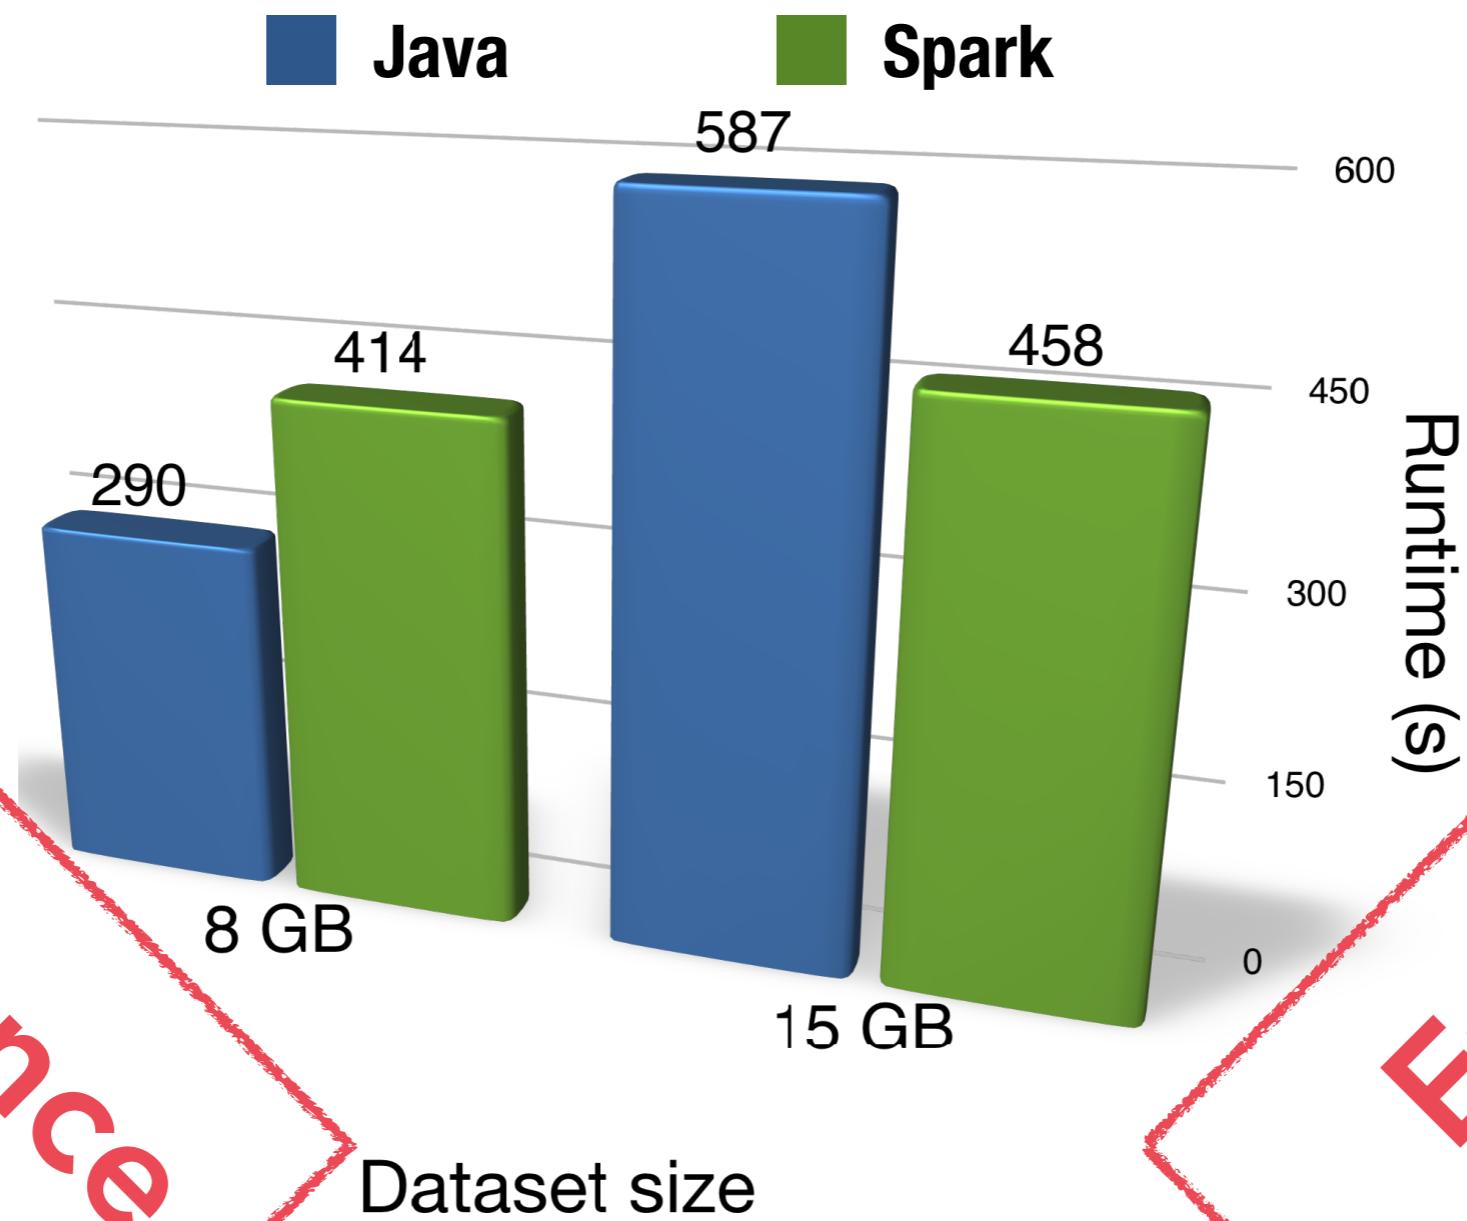
## SGD



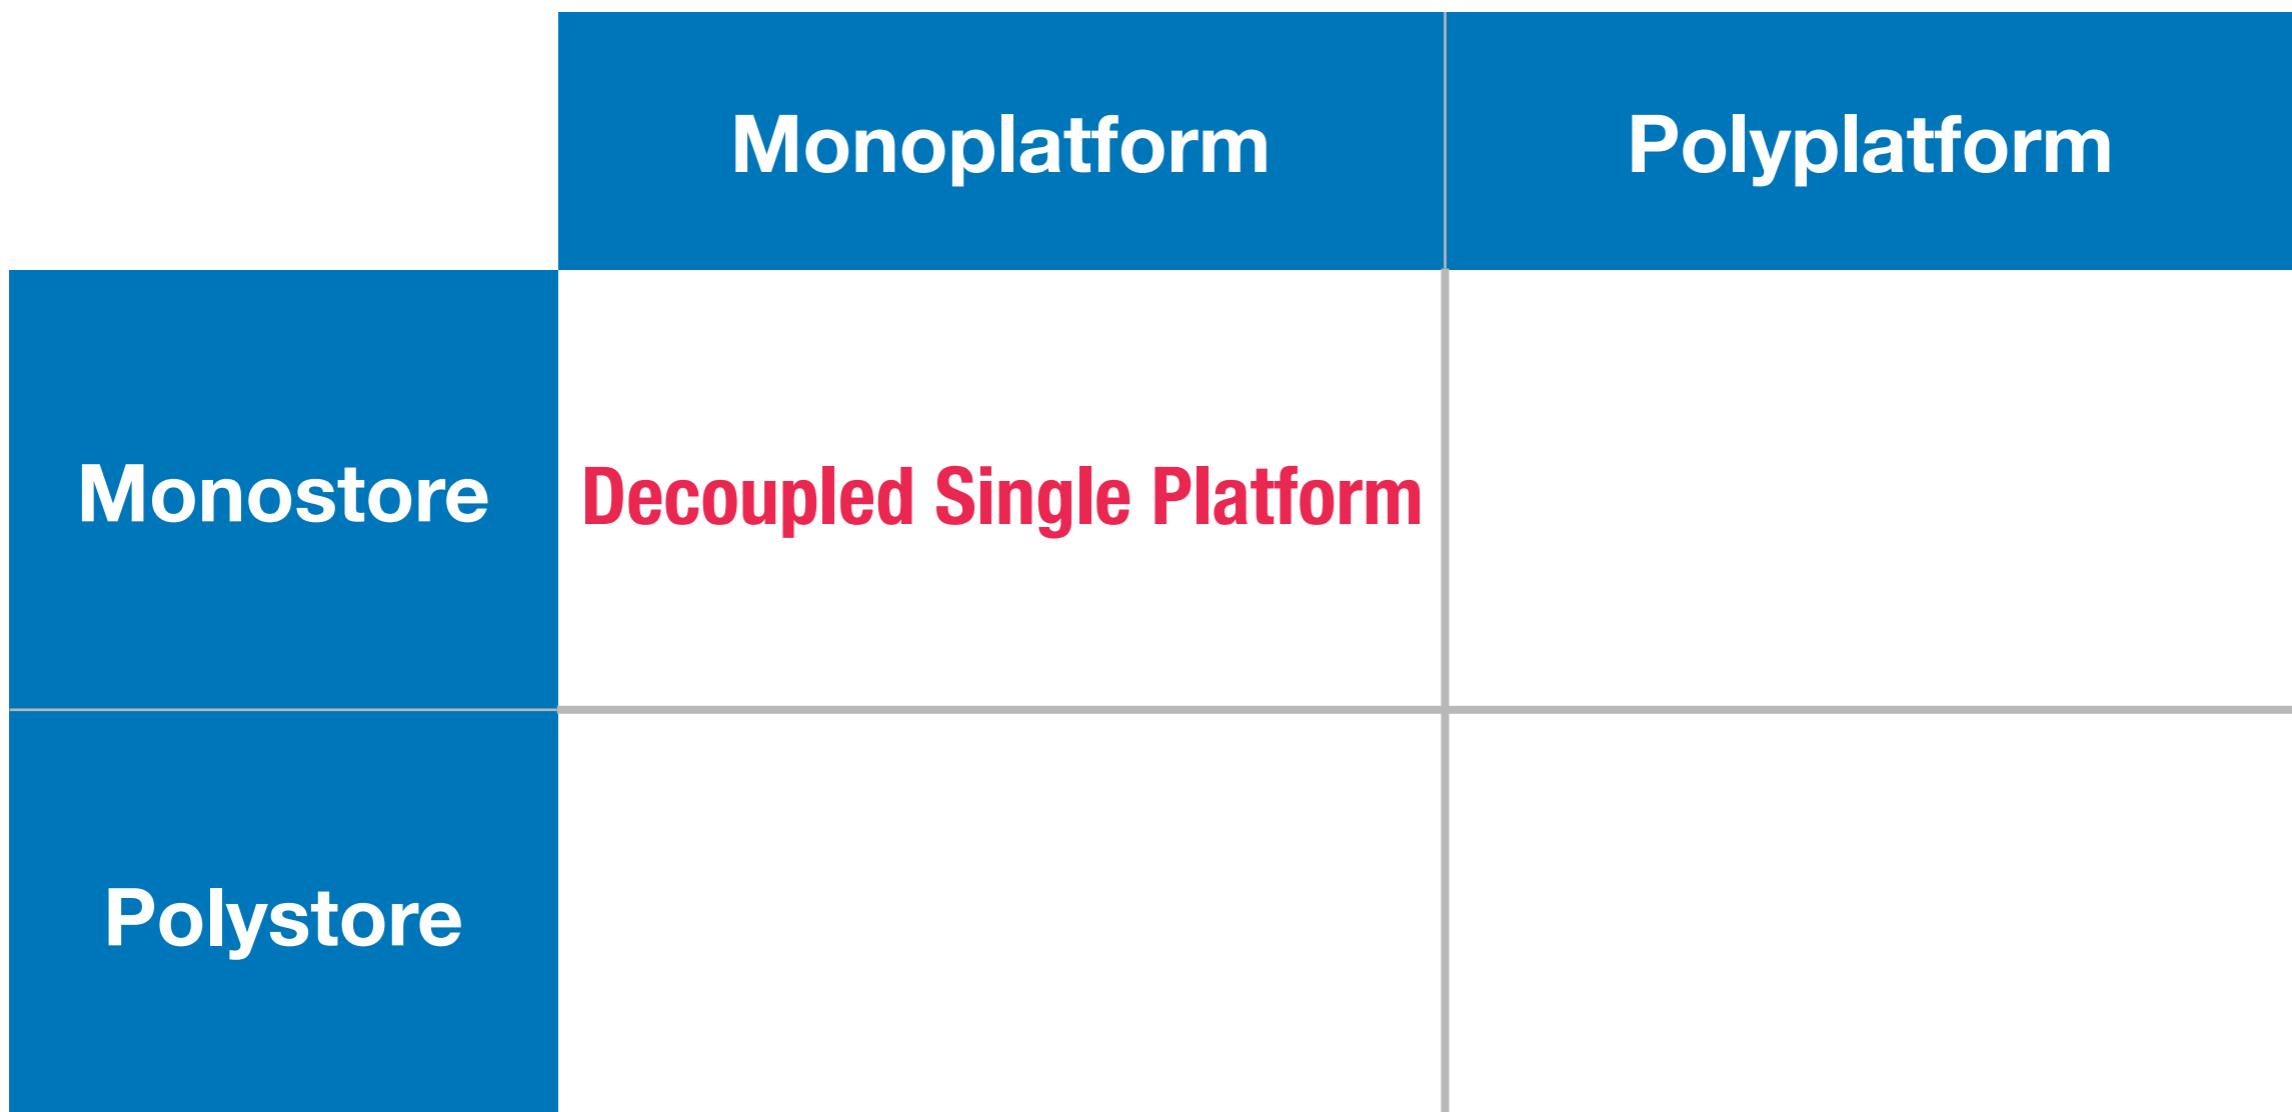
# Decoupled Single-Platform

using *any single* data processing platform to process a query

## SGD



# System-Store Quadrant



# Cross-Platform Use Cases

Data Processing

**Cross-Platform Data Processing**

**Decoupled Single-Platform**

Multi-Platform

**Opportunistic**

**Mandatory**

**Polystore**

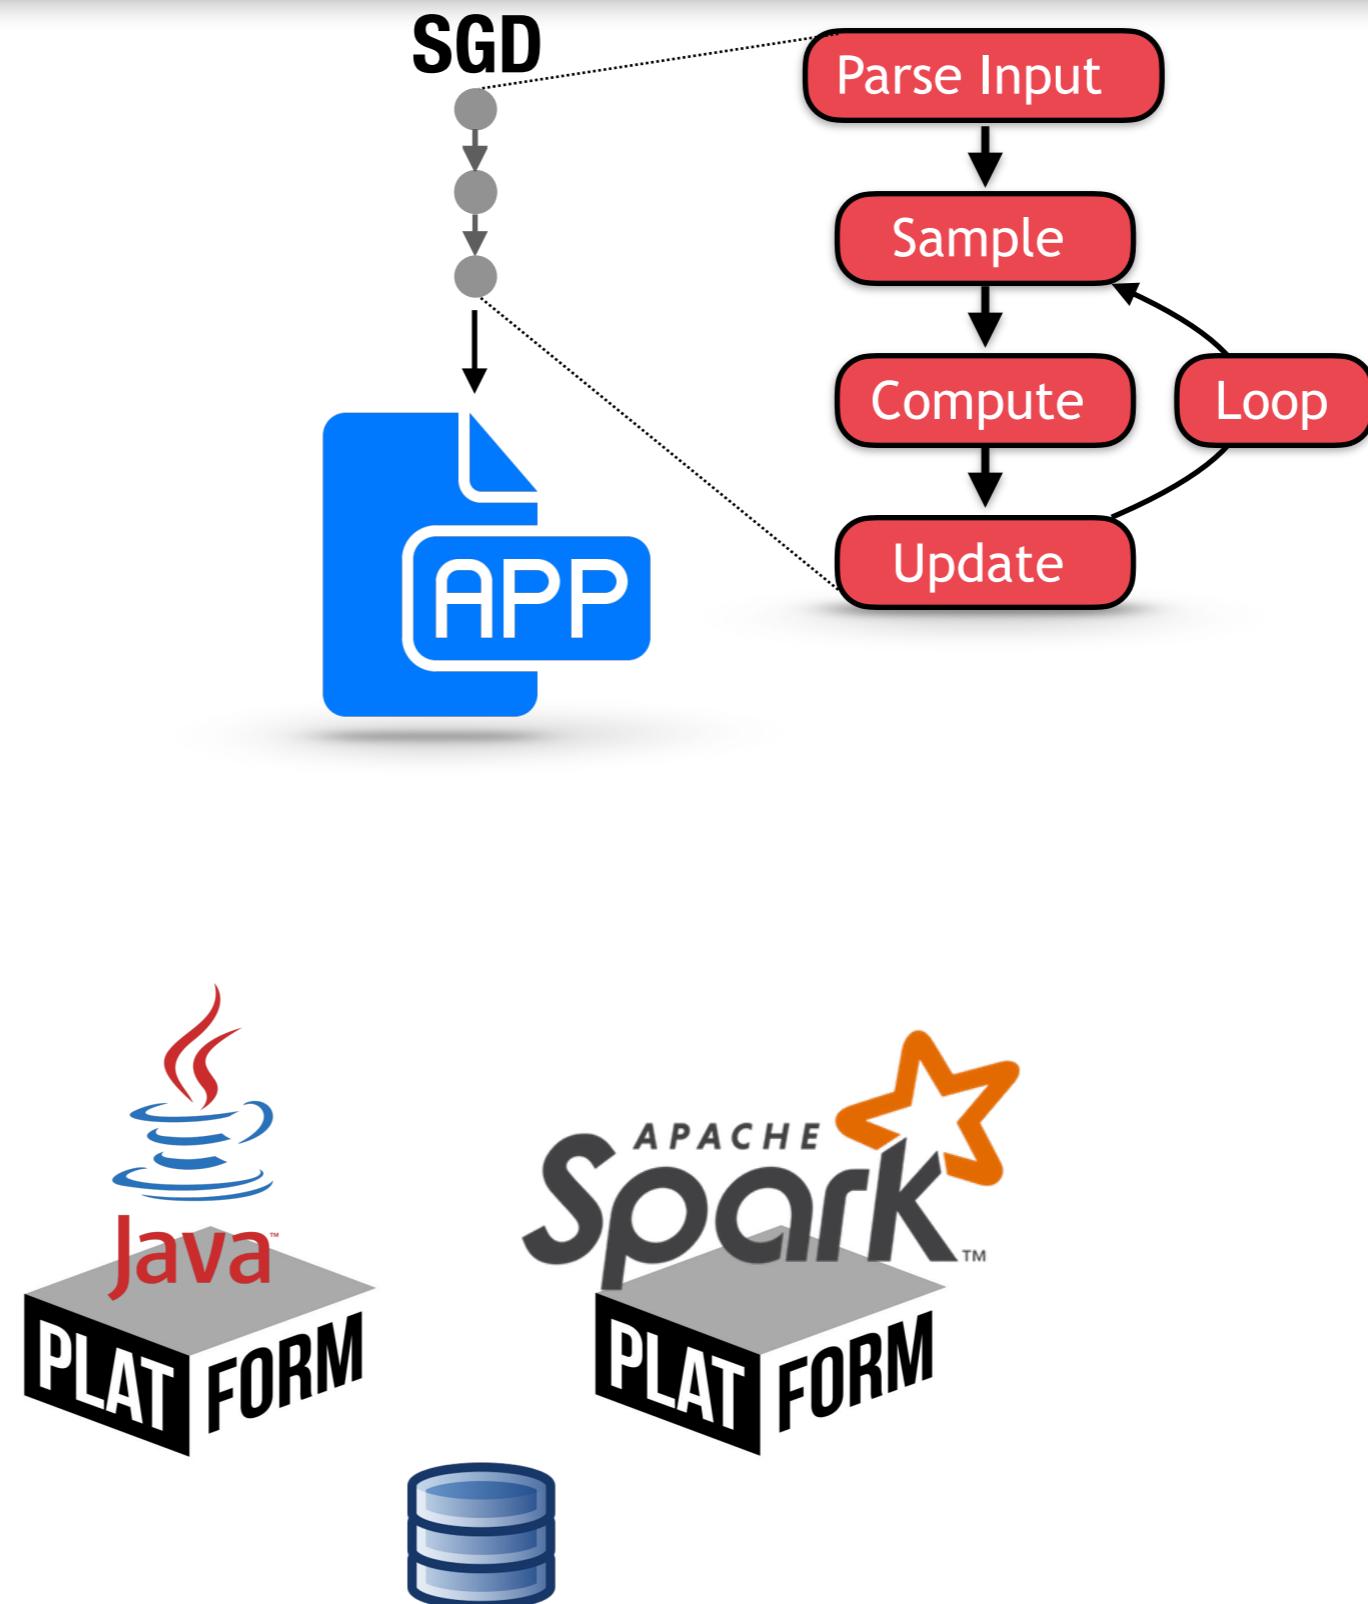
# Opportunistic Cross-Platform

using **several** data processing platforms to reduce the cost of a query



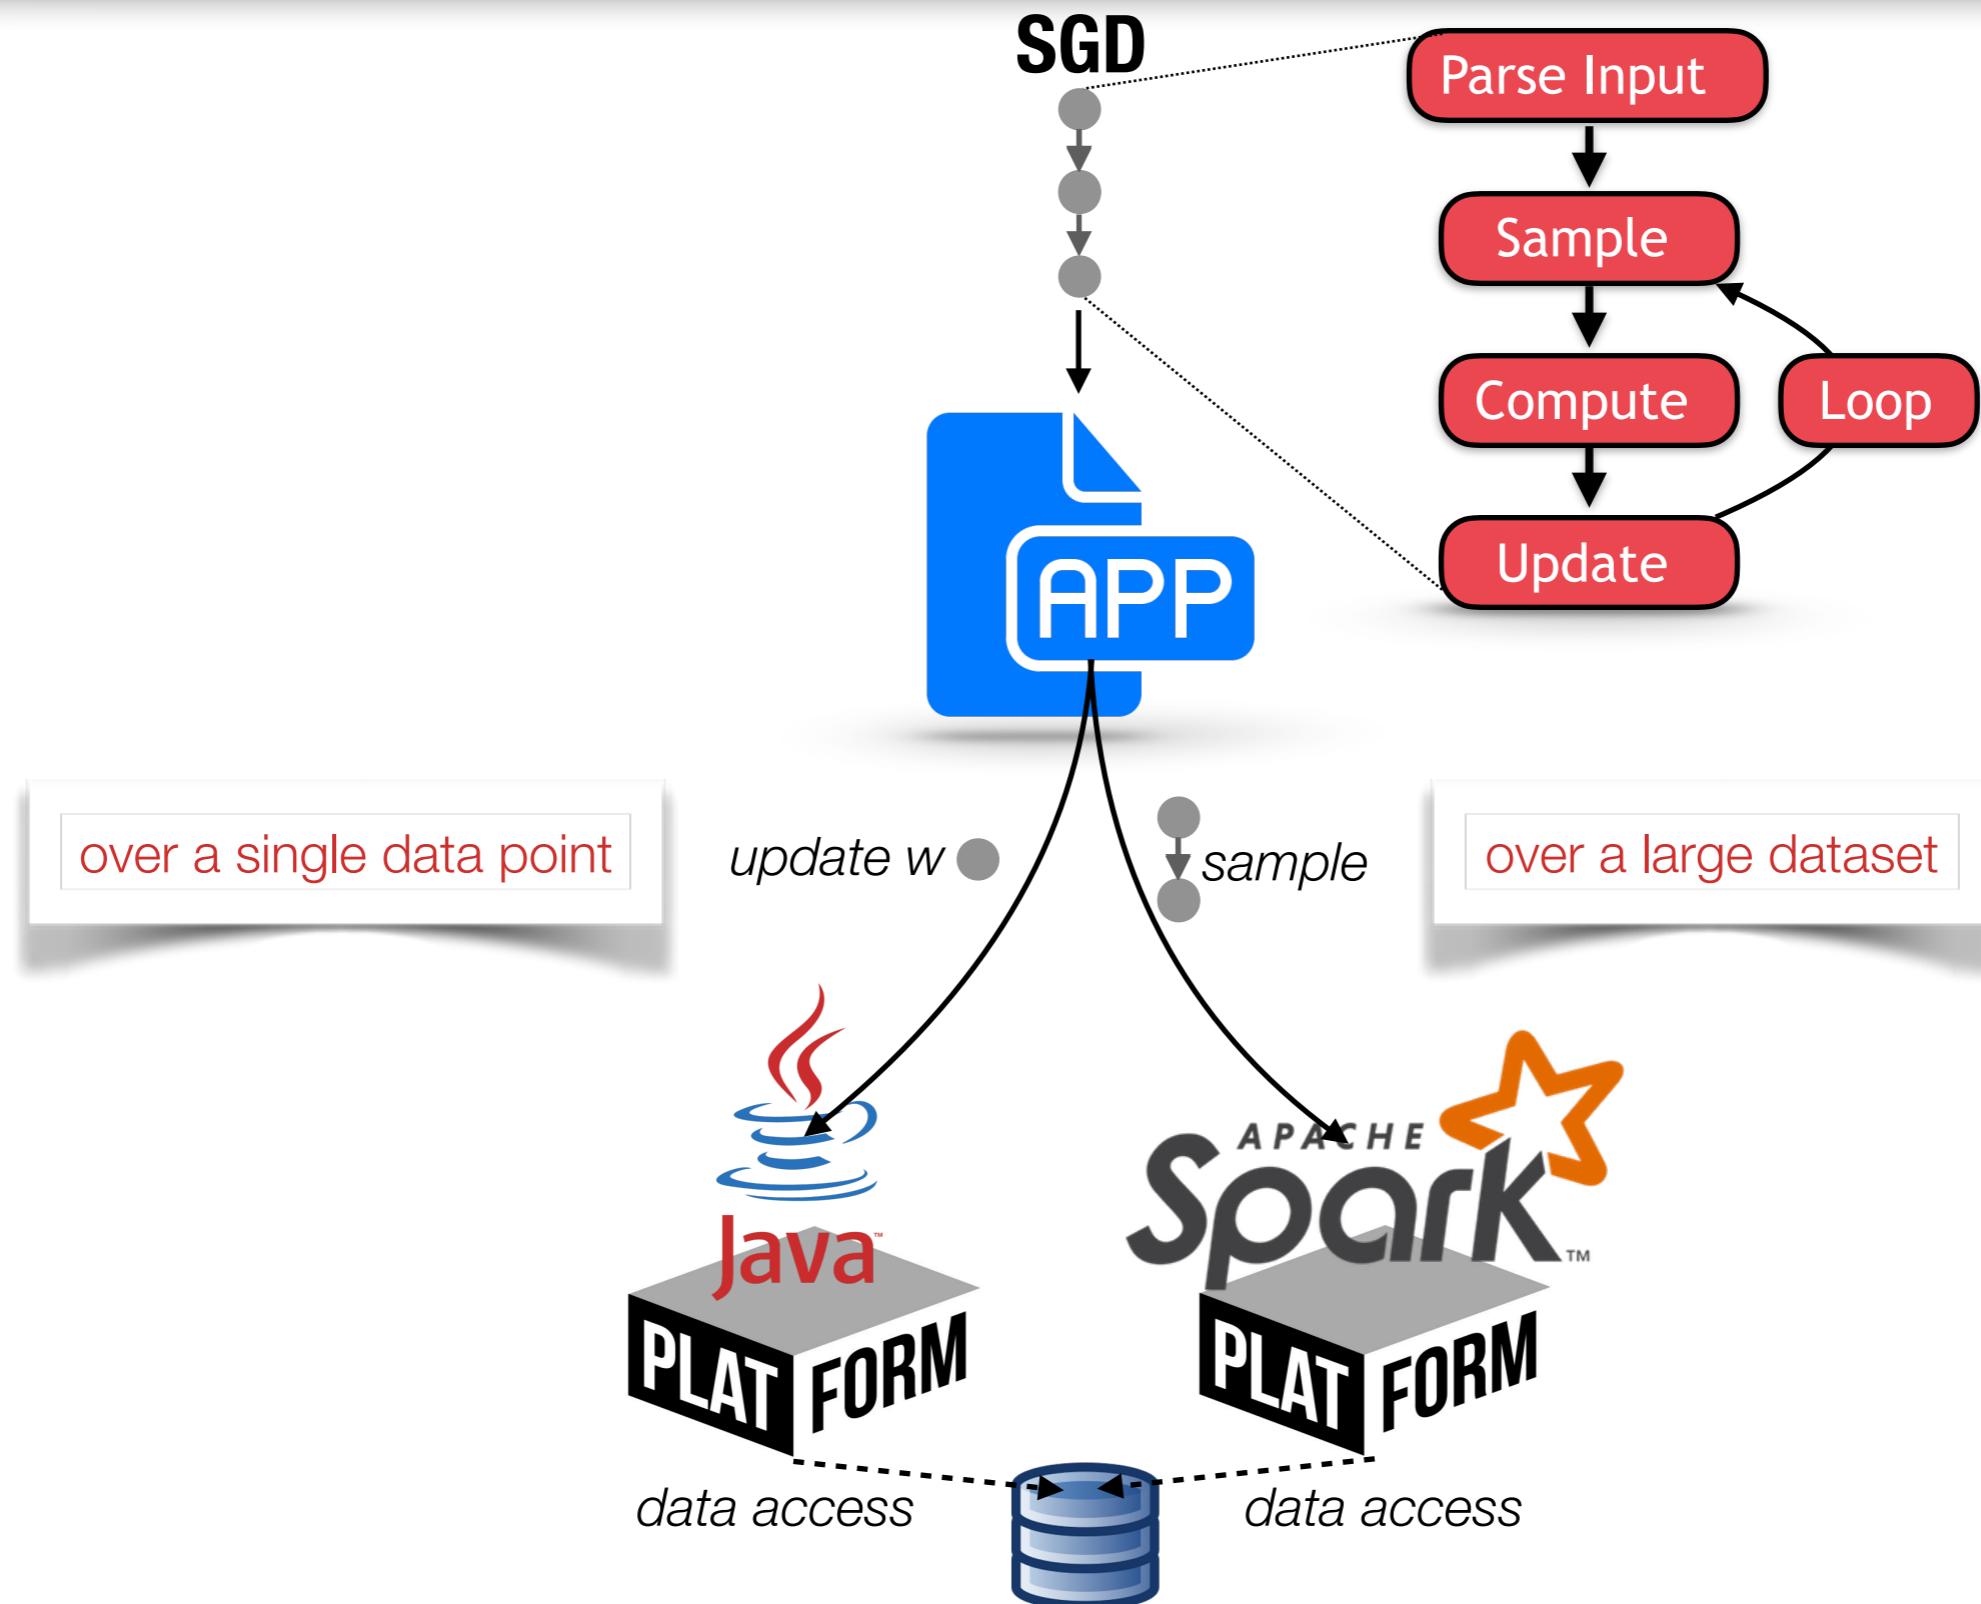
# Opportunistic Cross-Platform

using **several** data processing platforms to reduce the cost of a query



# Opportunistic Cross-Platform

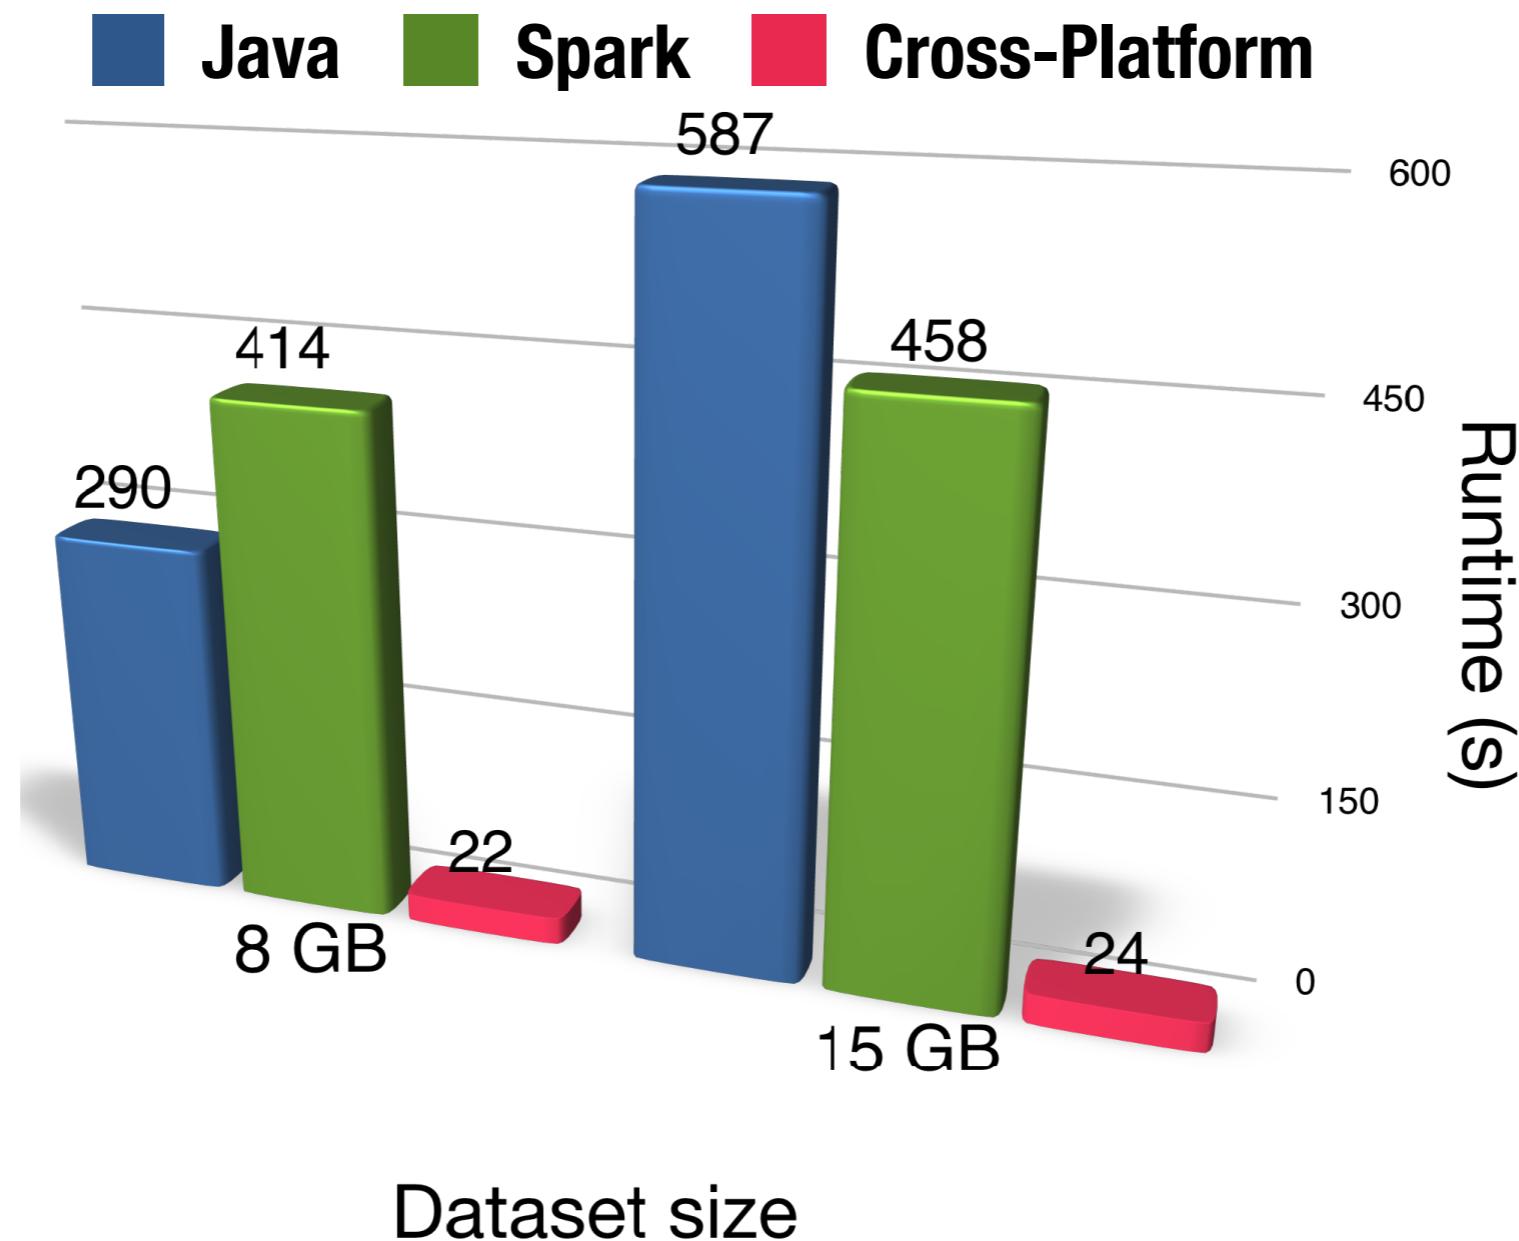
using **several** data processing platforms to reduce the cost of a query



# Opportunistic Cross-Platform

using **several** data processing platforms to reduce the cost of a query

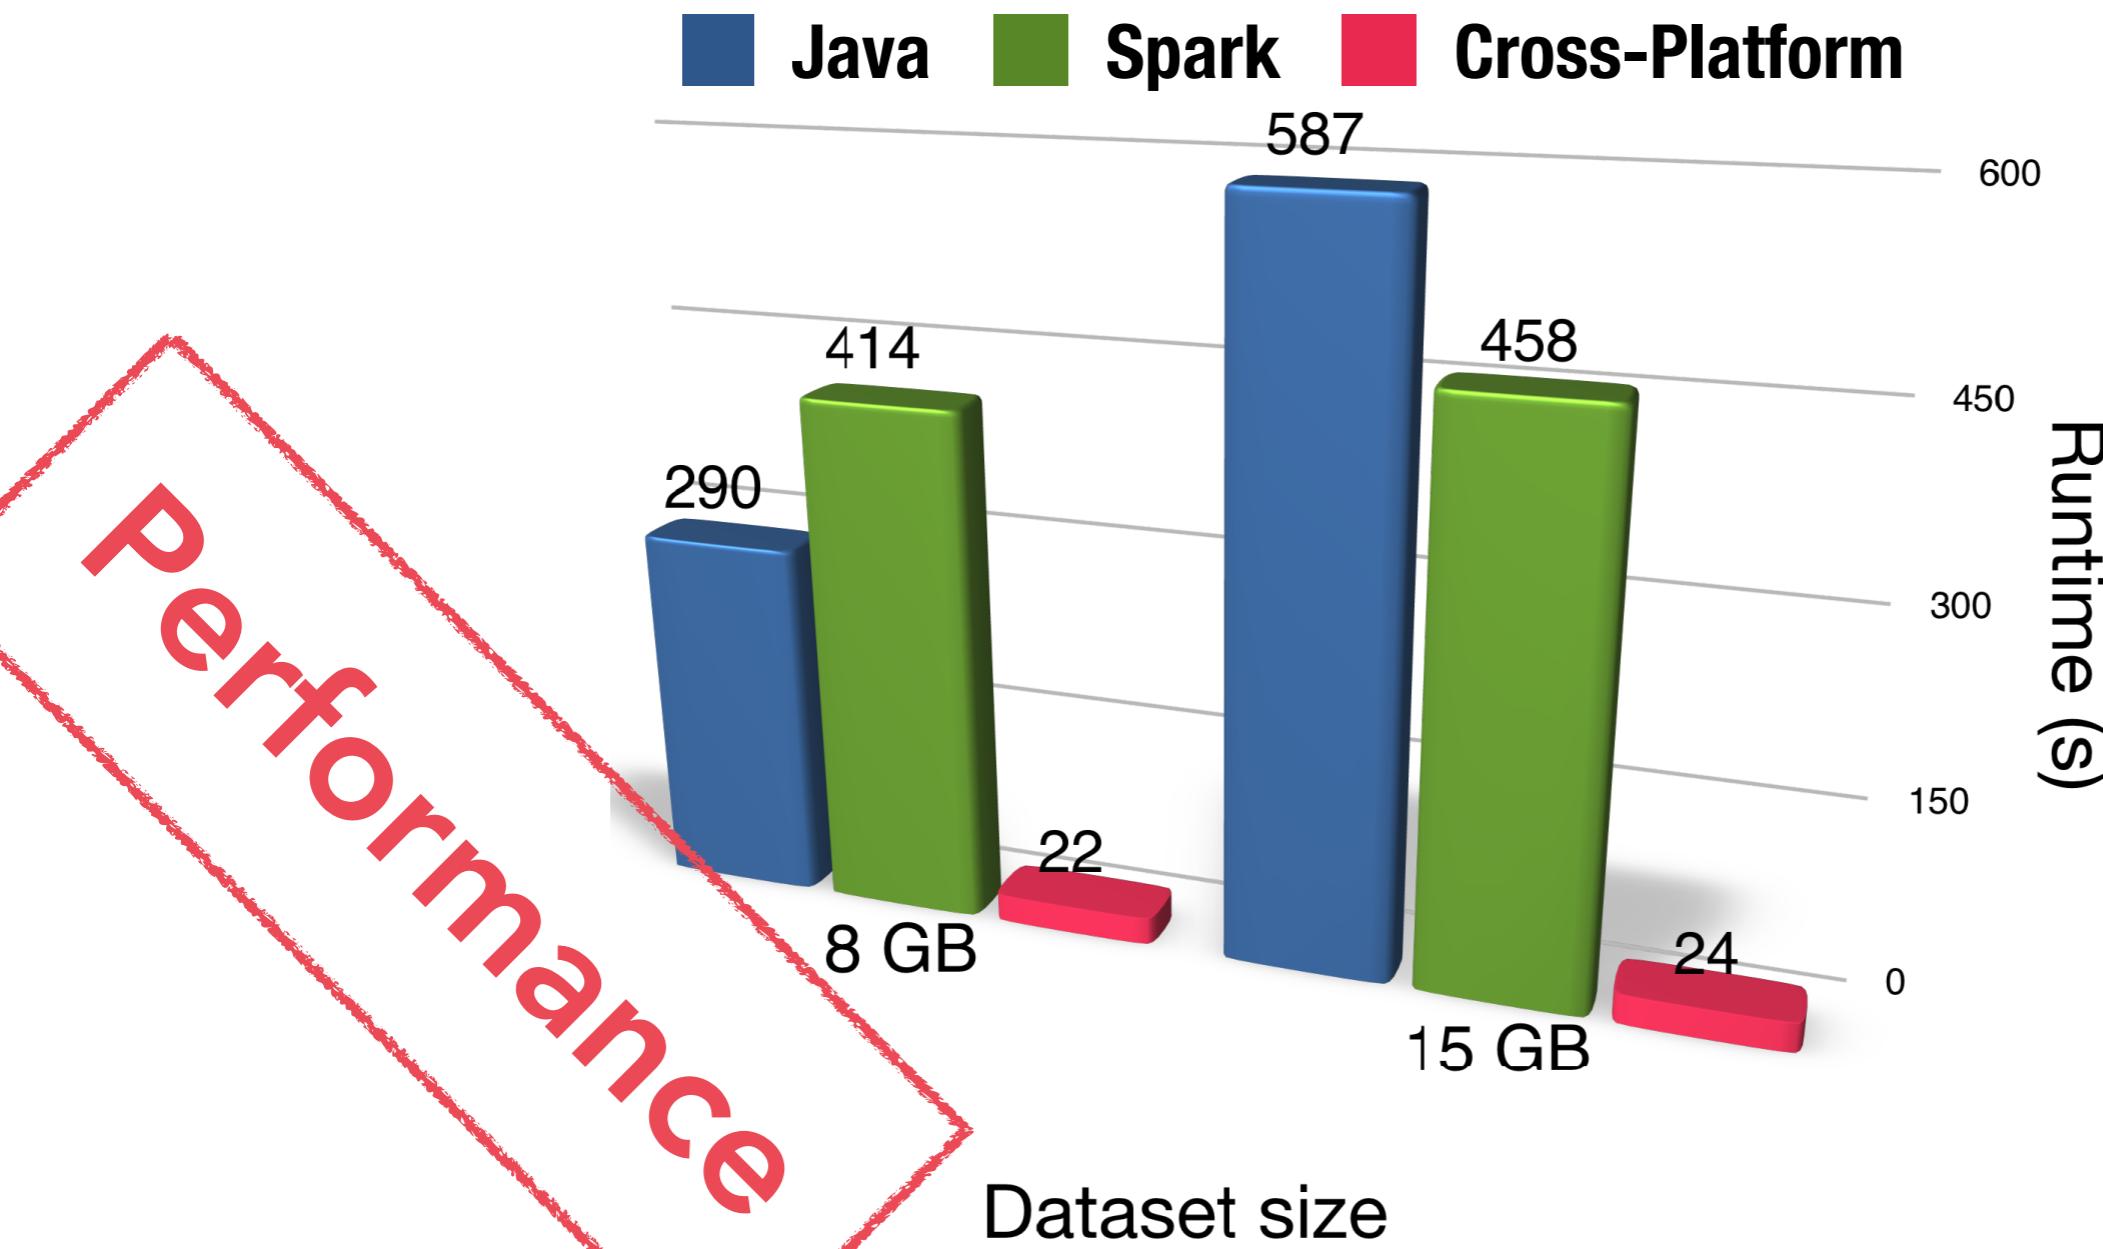
## SGD



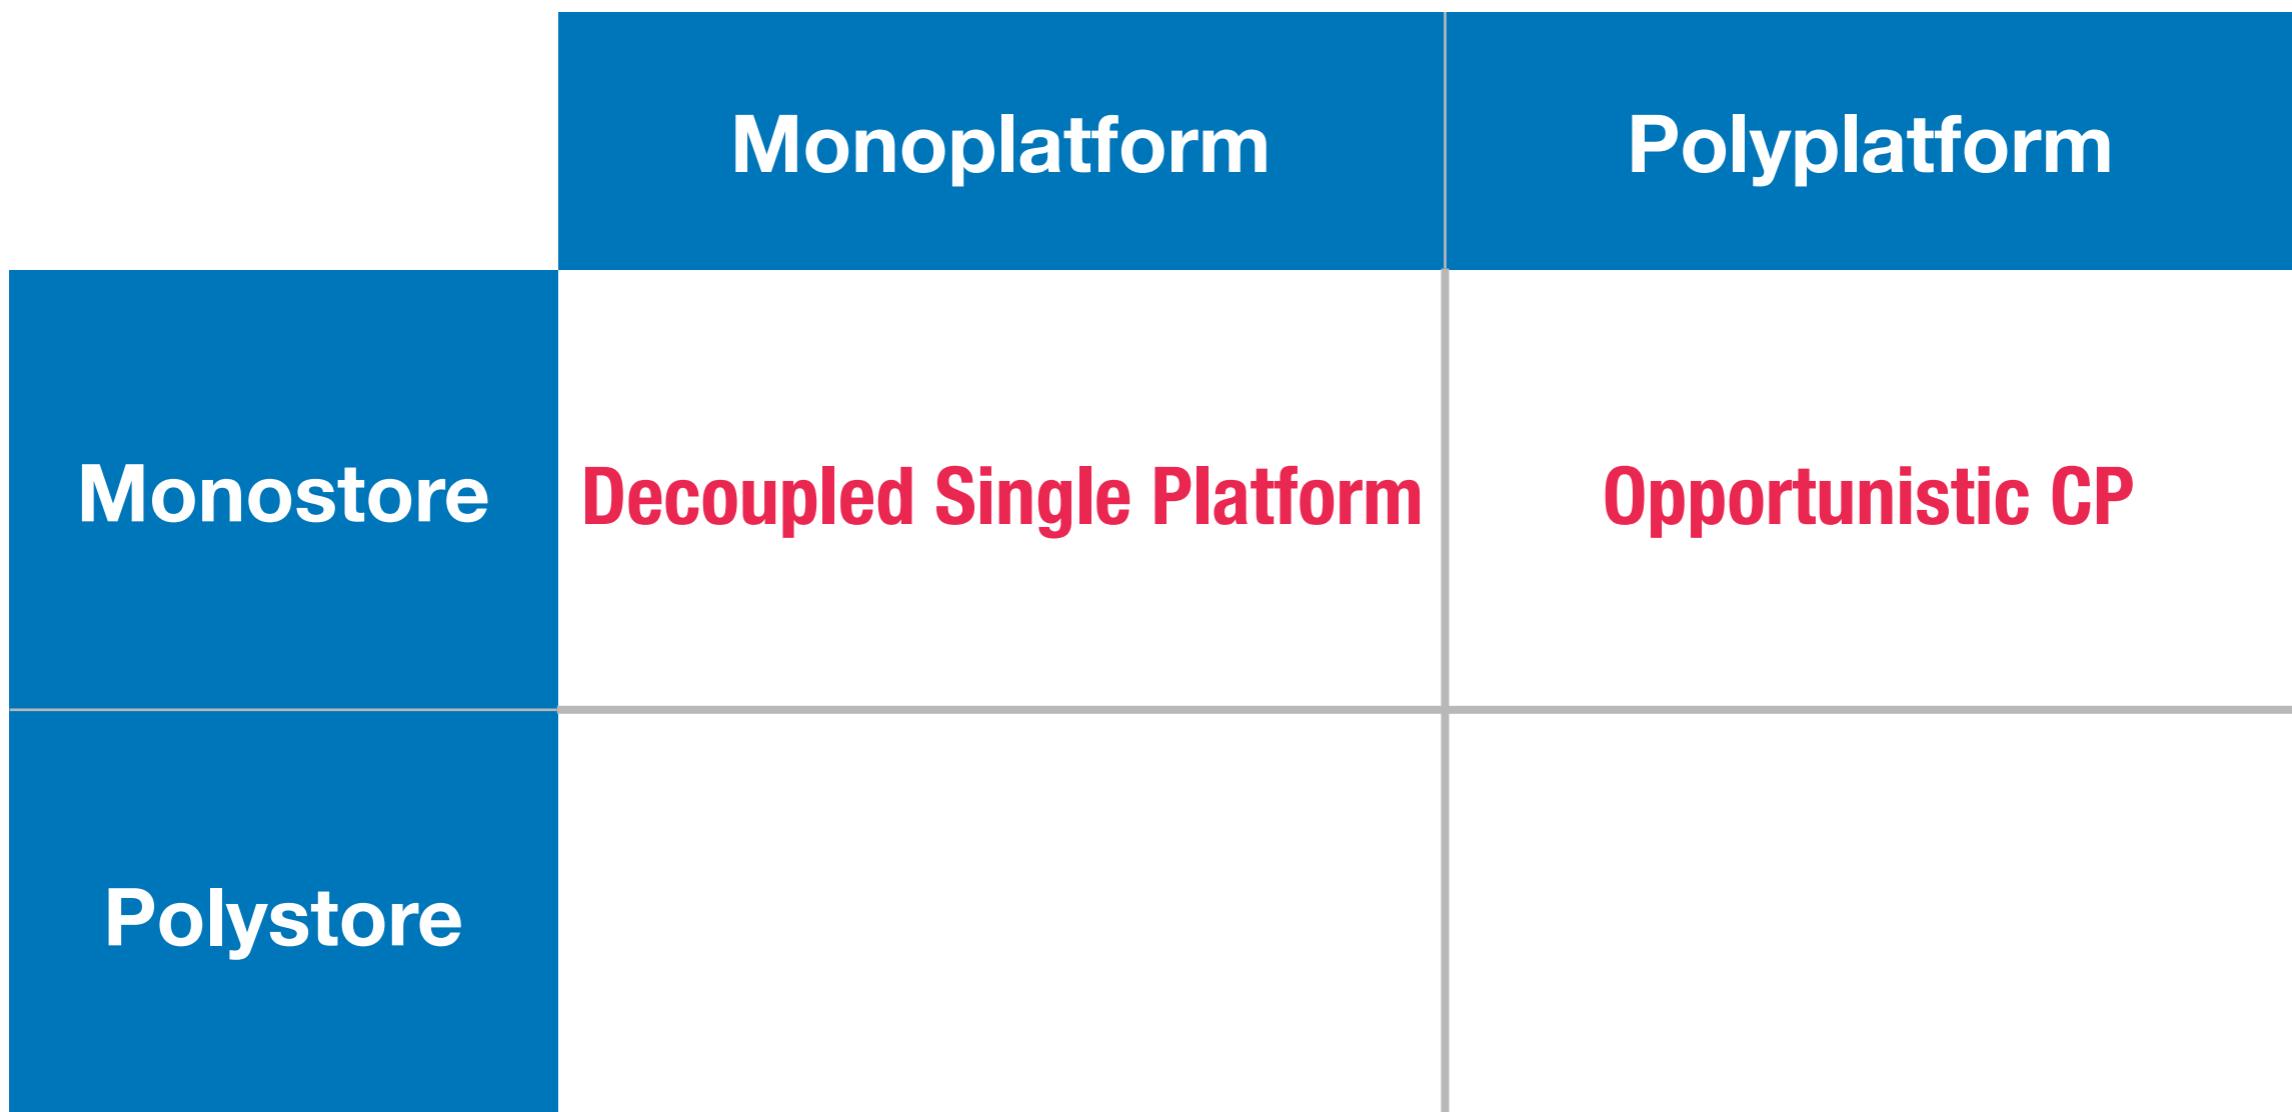
# Opportunistic Cross-Platform

using **several** data processing platforms to reduce the cost of a query

## SGD



# System-Store Quadrant



# Cross-Platform Use Cases

Data Processing

**Cross-Platform Data Processing**

**Decoupled Single-Platform**

Multi-Platform

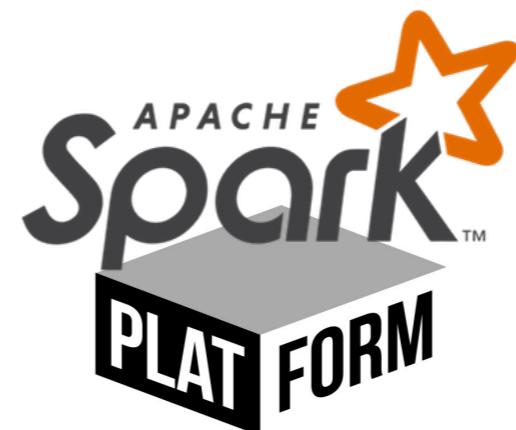
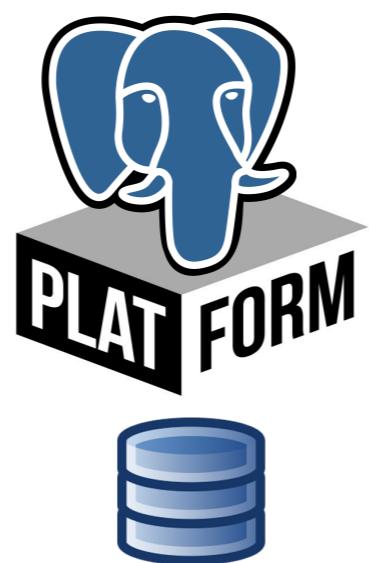
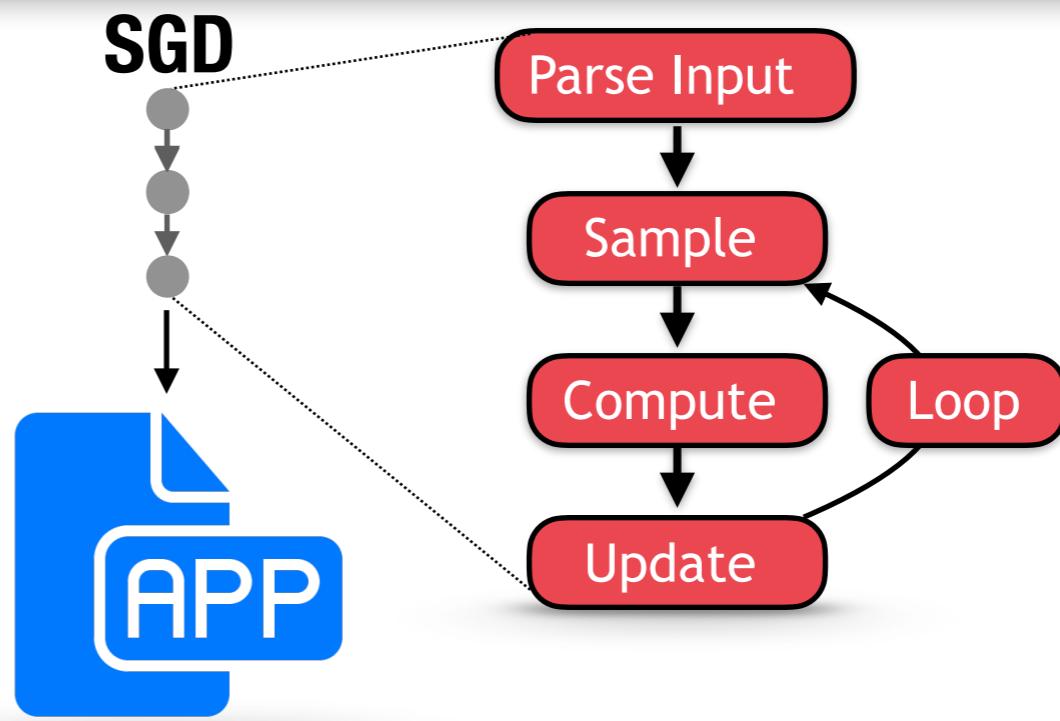
**Opportunistic**

**Mandatory**

**Polystore**

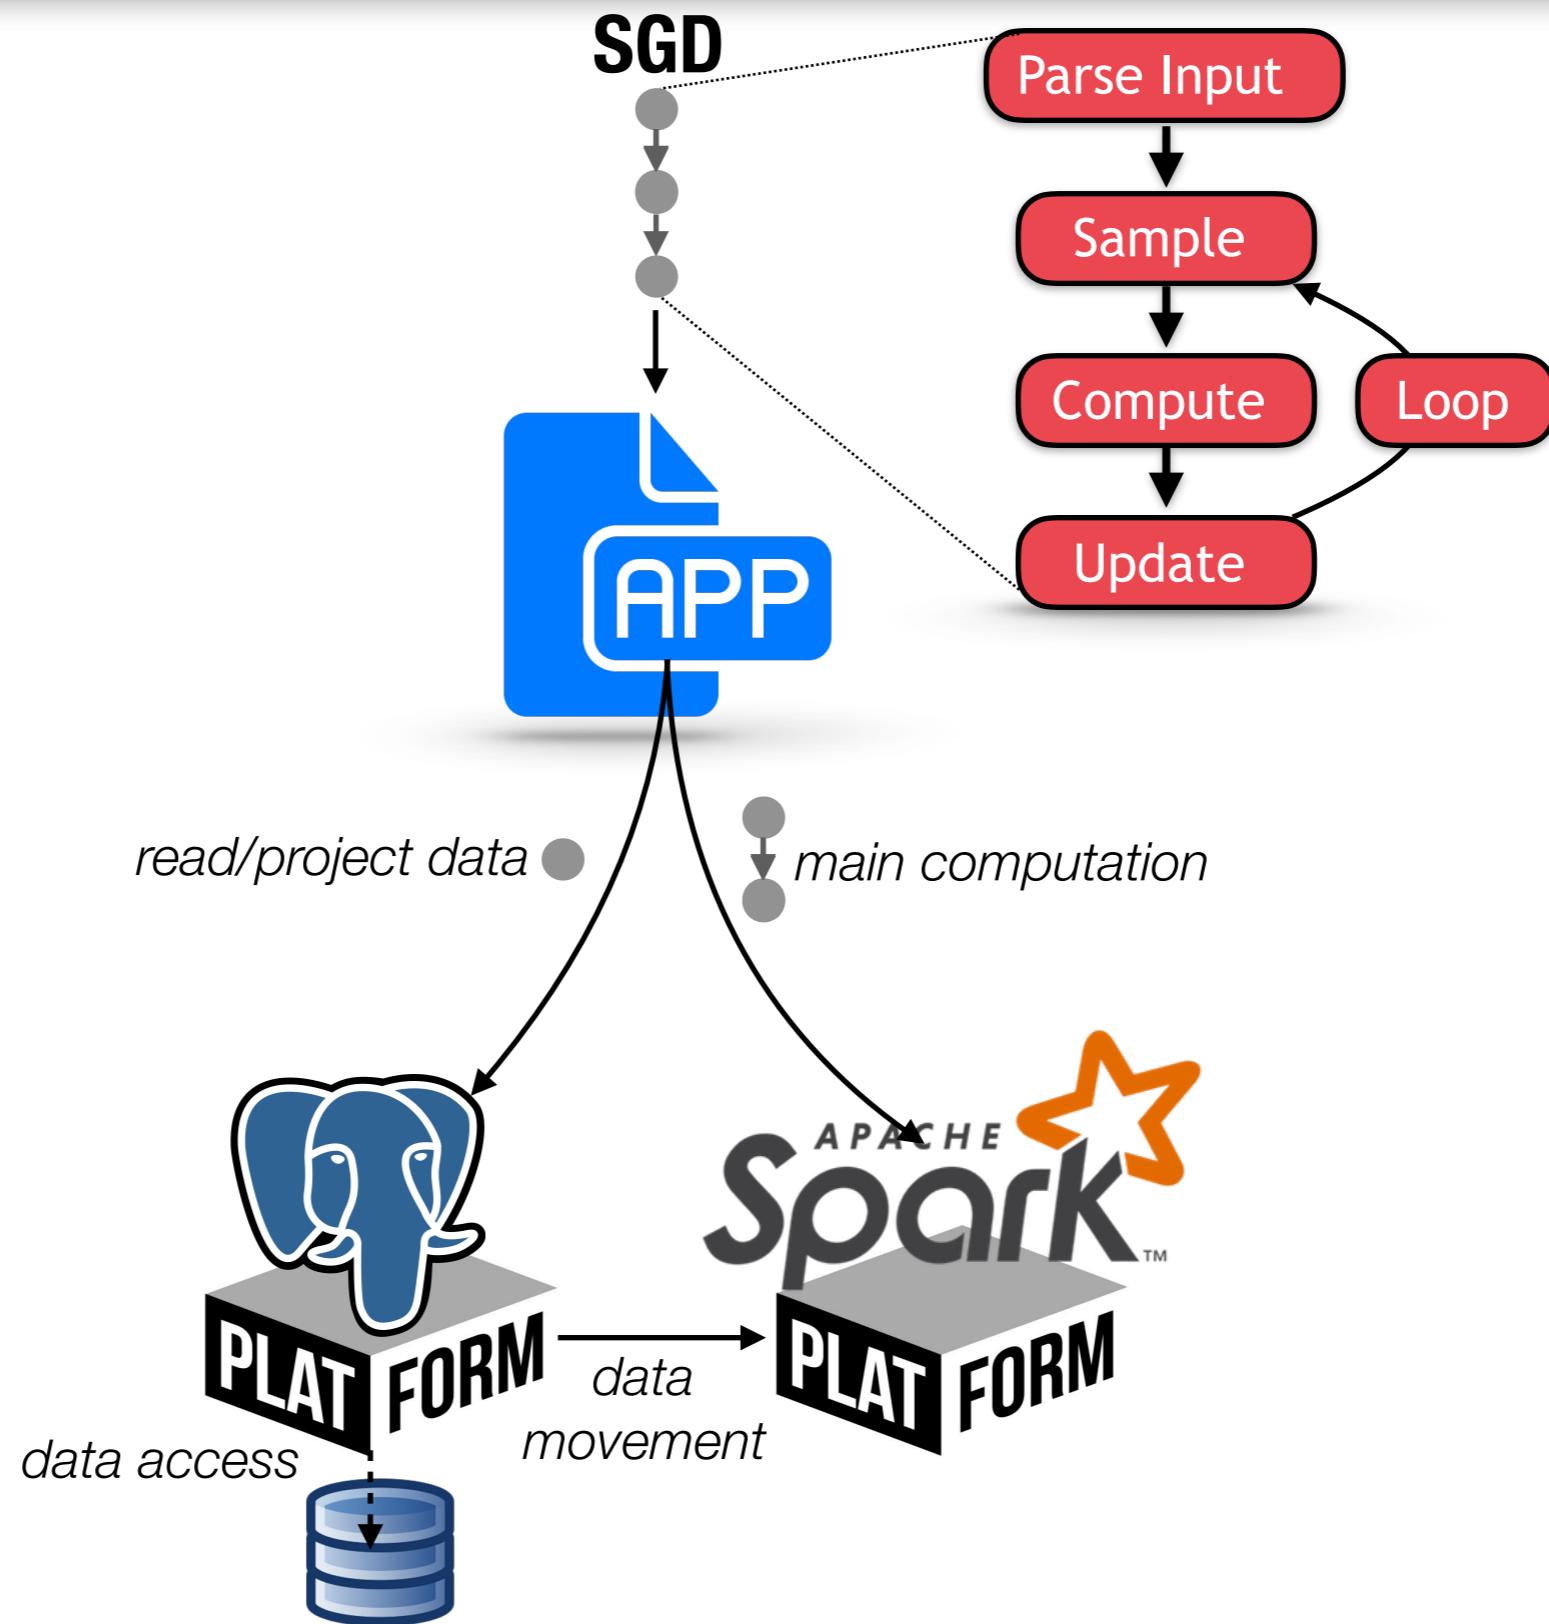
# Mandatory Cross-Platform

using **several** data processing platforms to be able to process a query

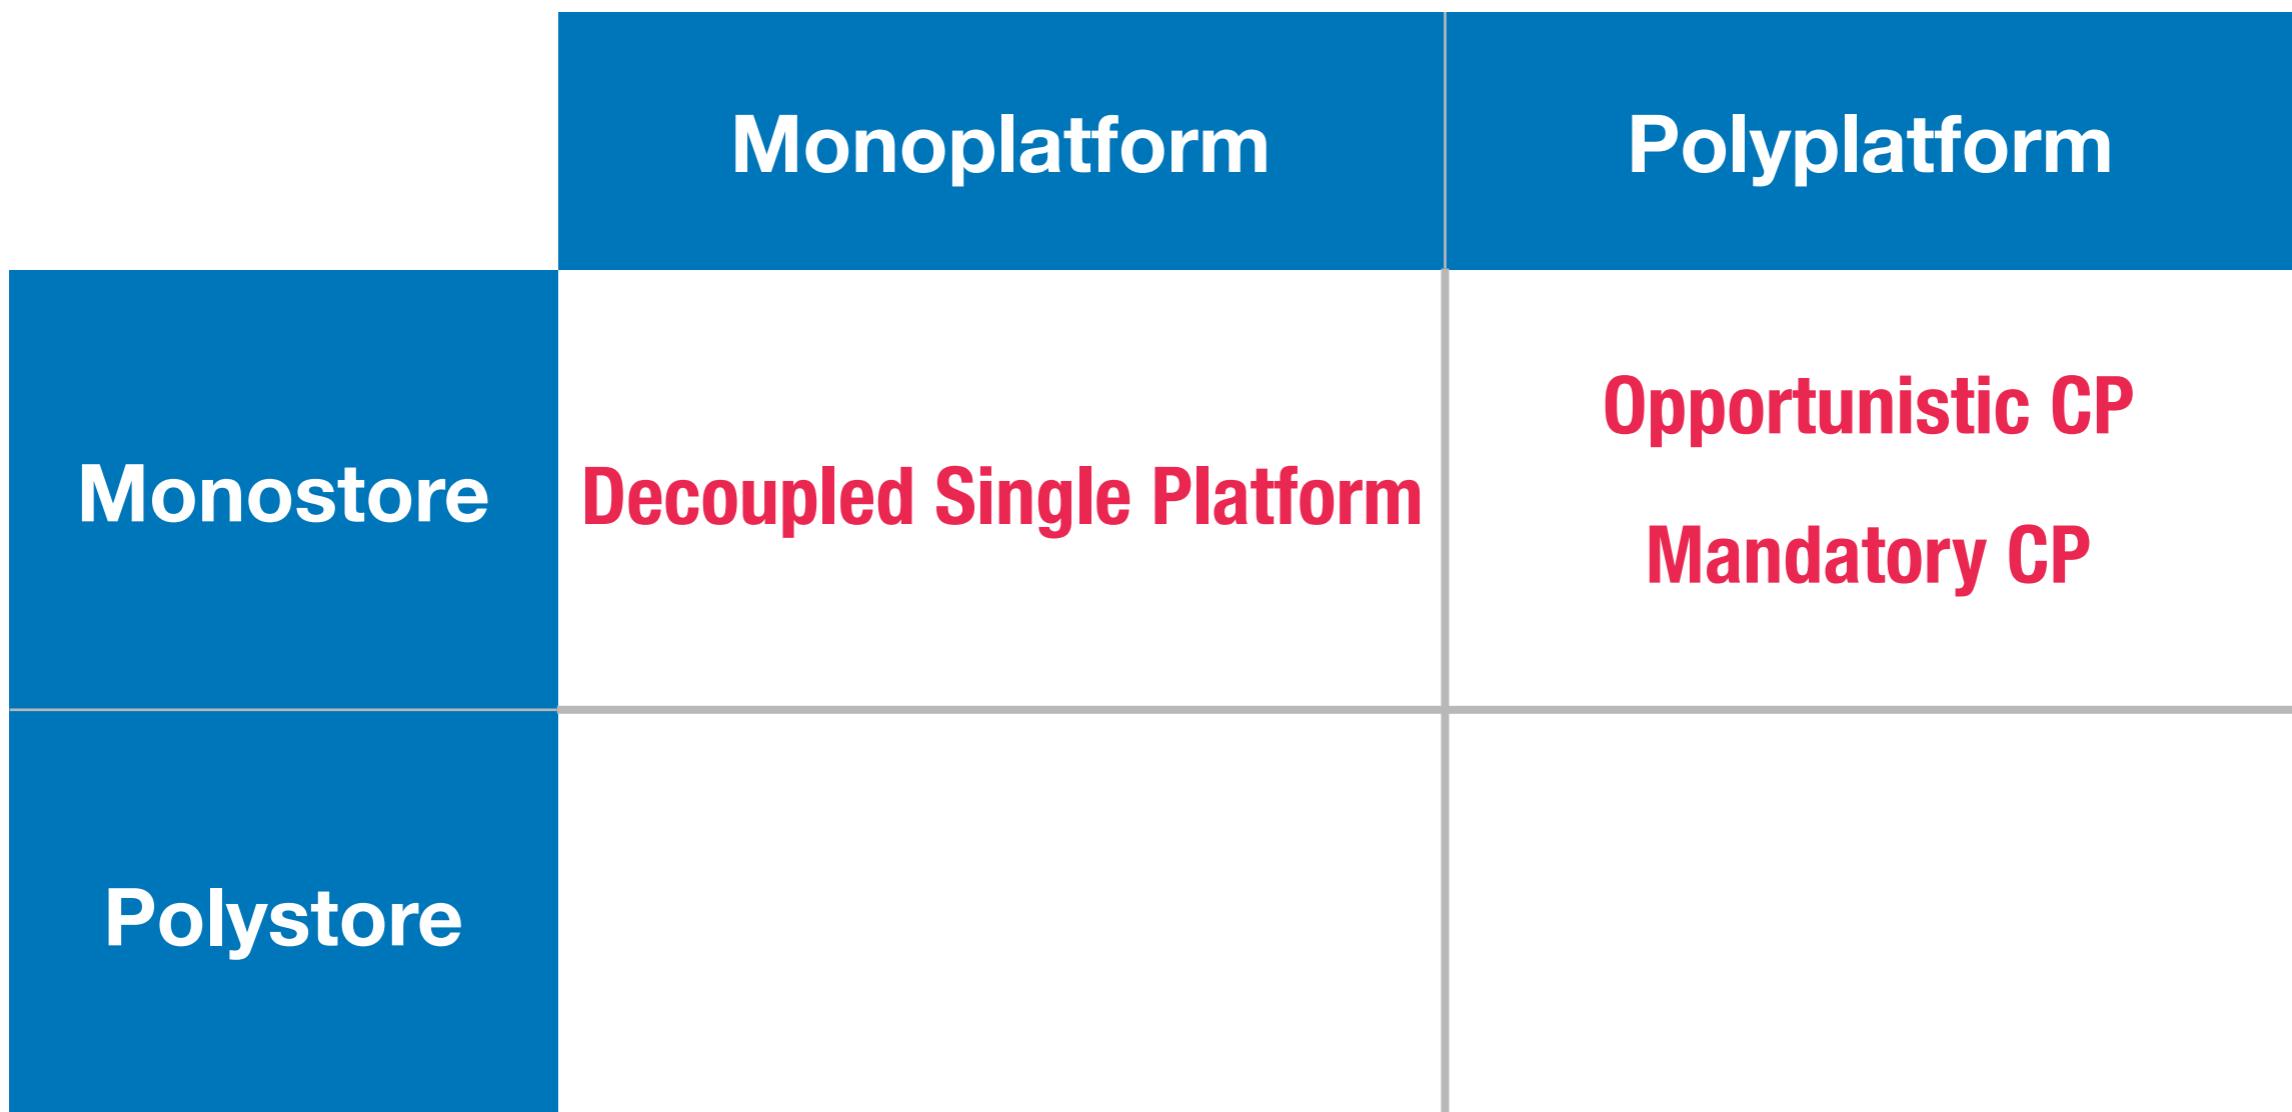


# Mandatory Cross-Platform

using **several** data processing platforms to be able to process a query



# System-Store Quadrant



# Cross-Platform Use Cases

Data Processing

**Cross-Platform Data Processing**

**Decoupled Single-Platform**

Multi-Platform

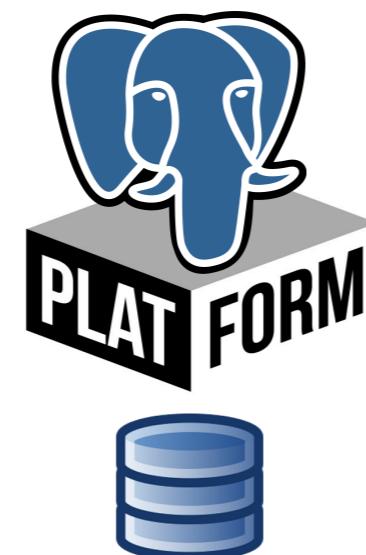
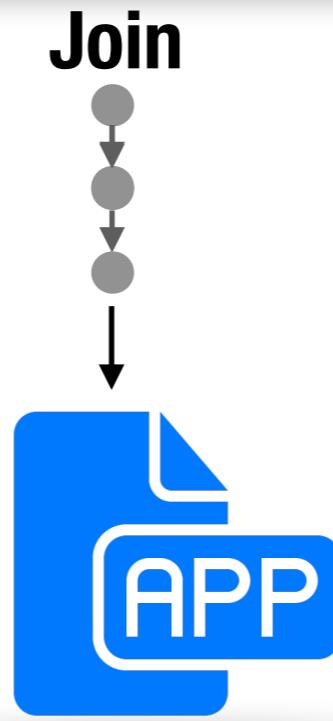
Opportunistic

Mandatory

**Polystore**

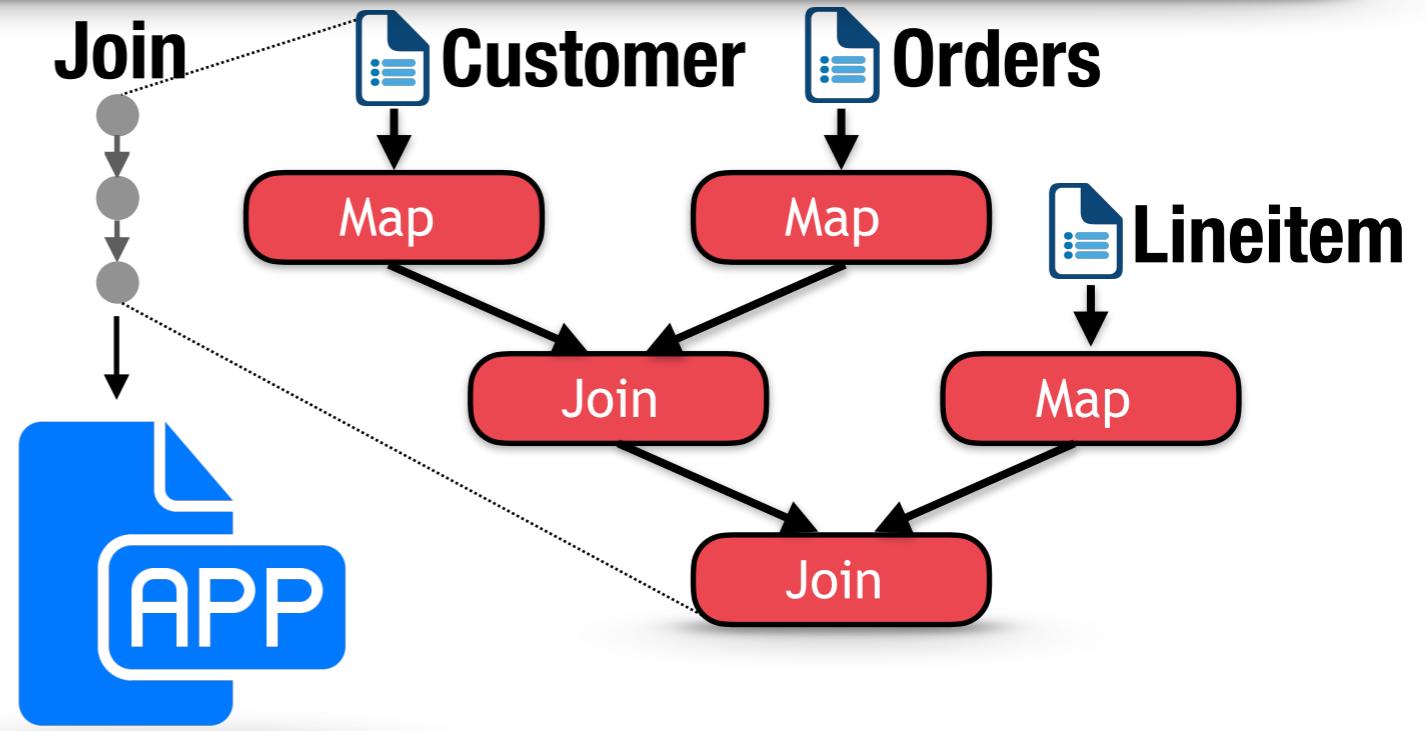
# Polystore

using **several** data processing platforms because data is spread



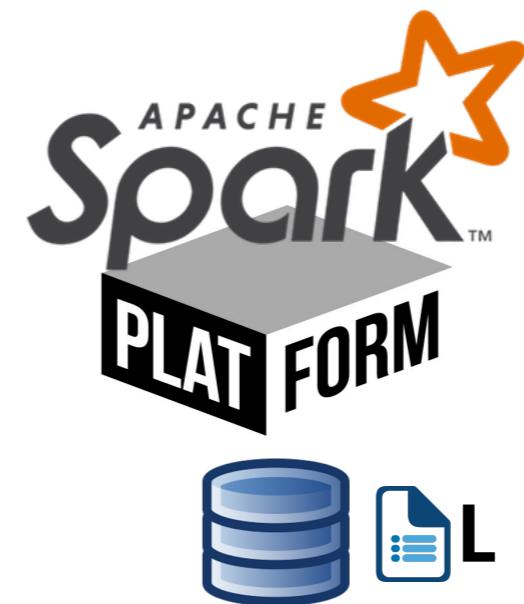
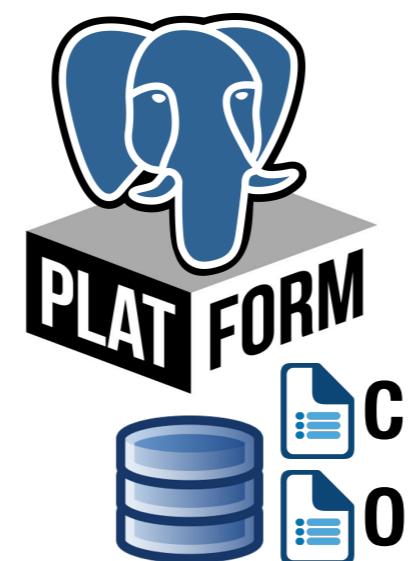
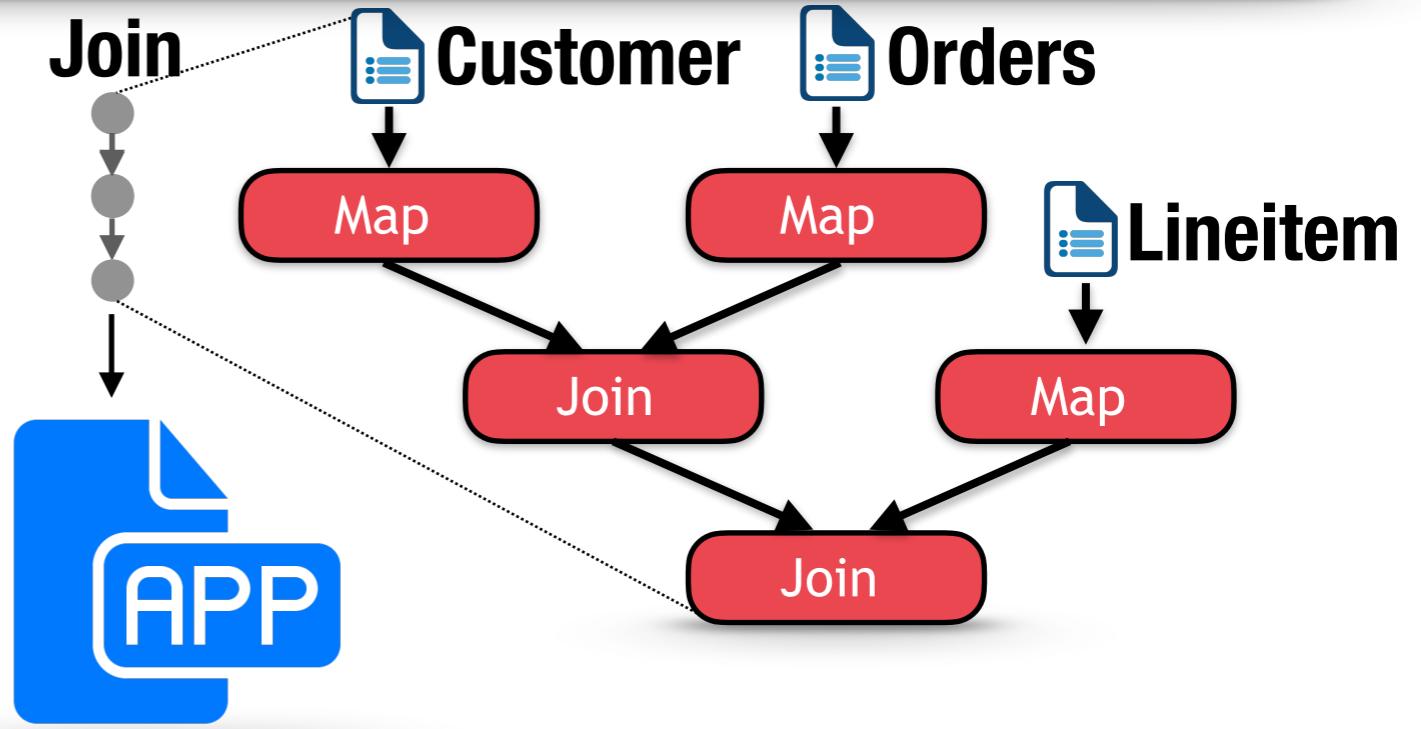
# Polystore

using **several** data processing platforms because data is spread



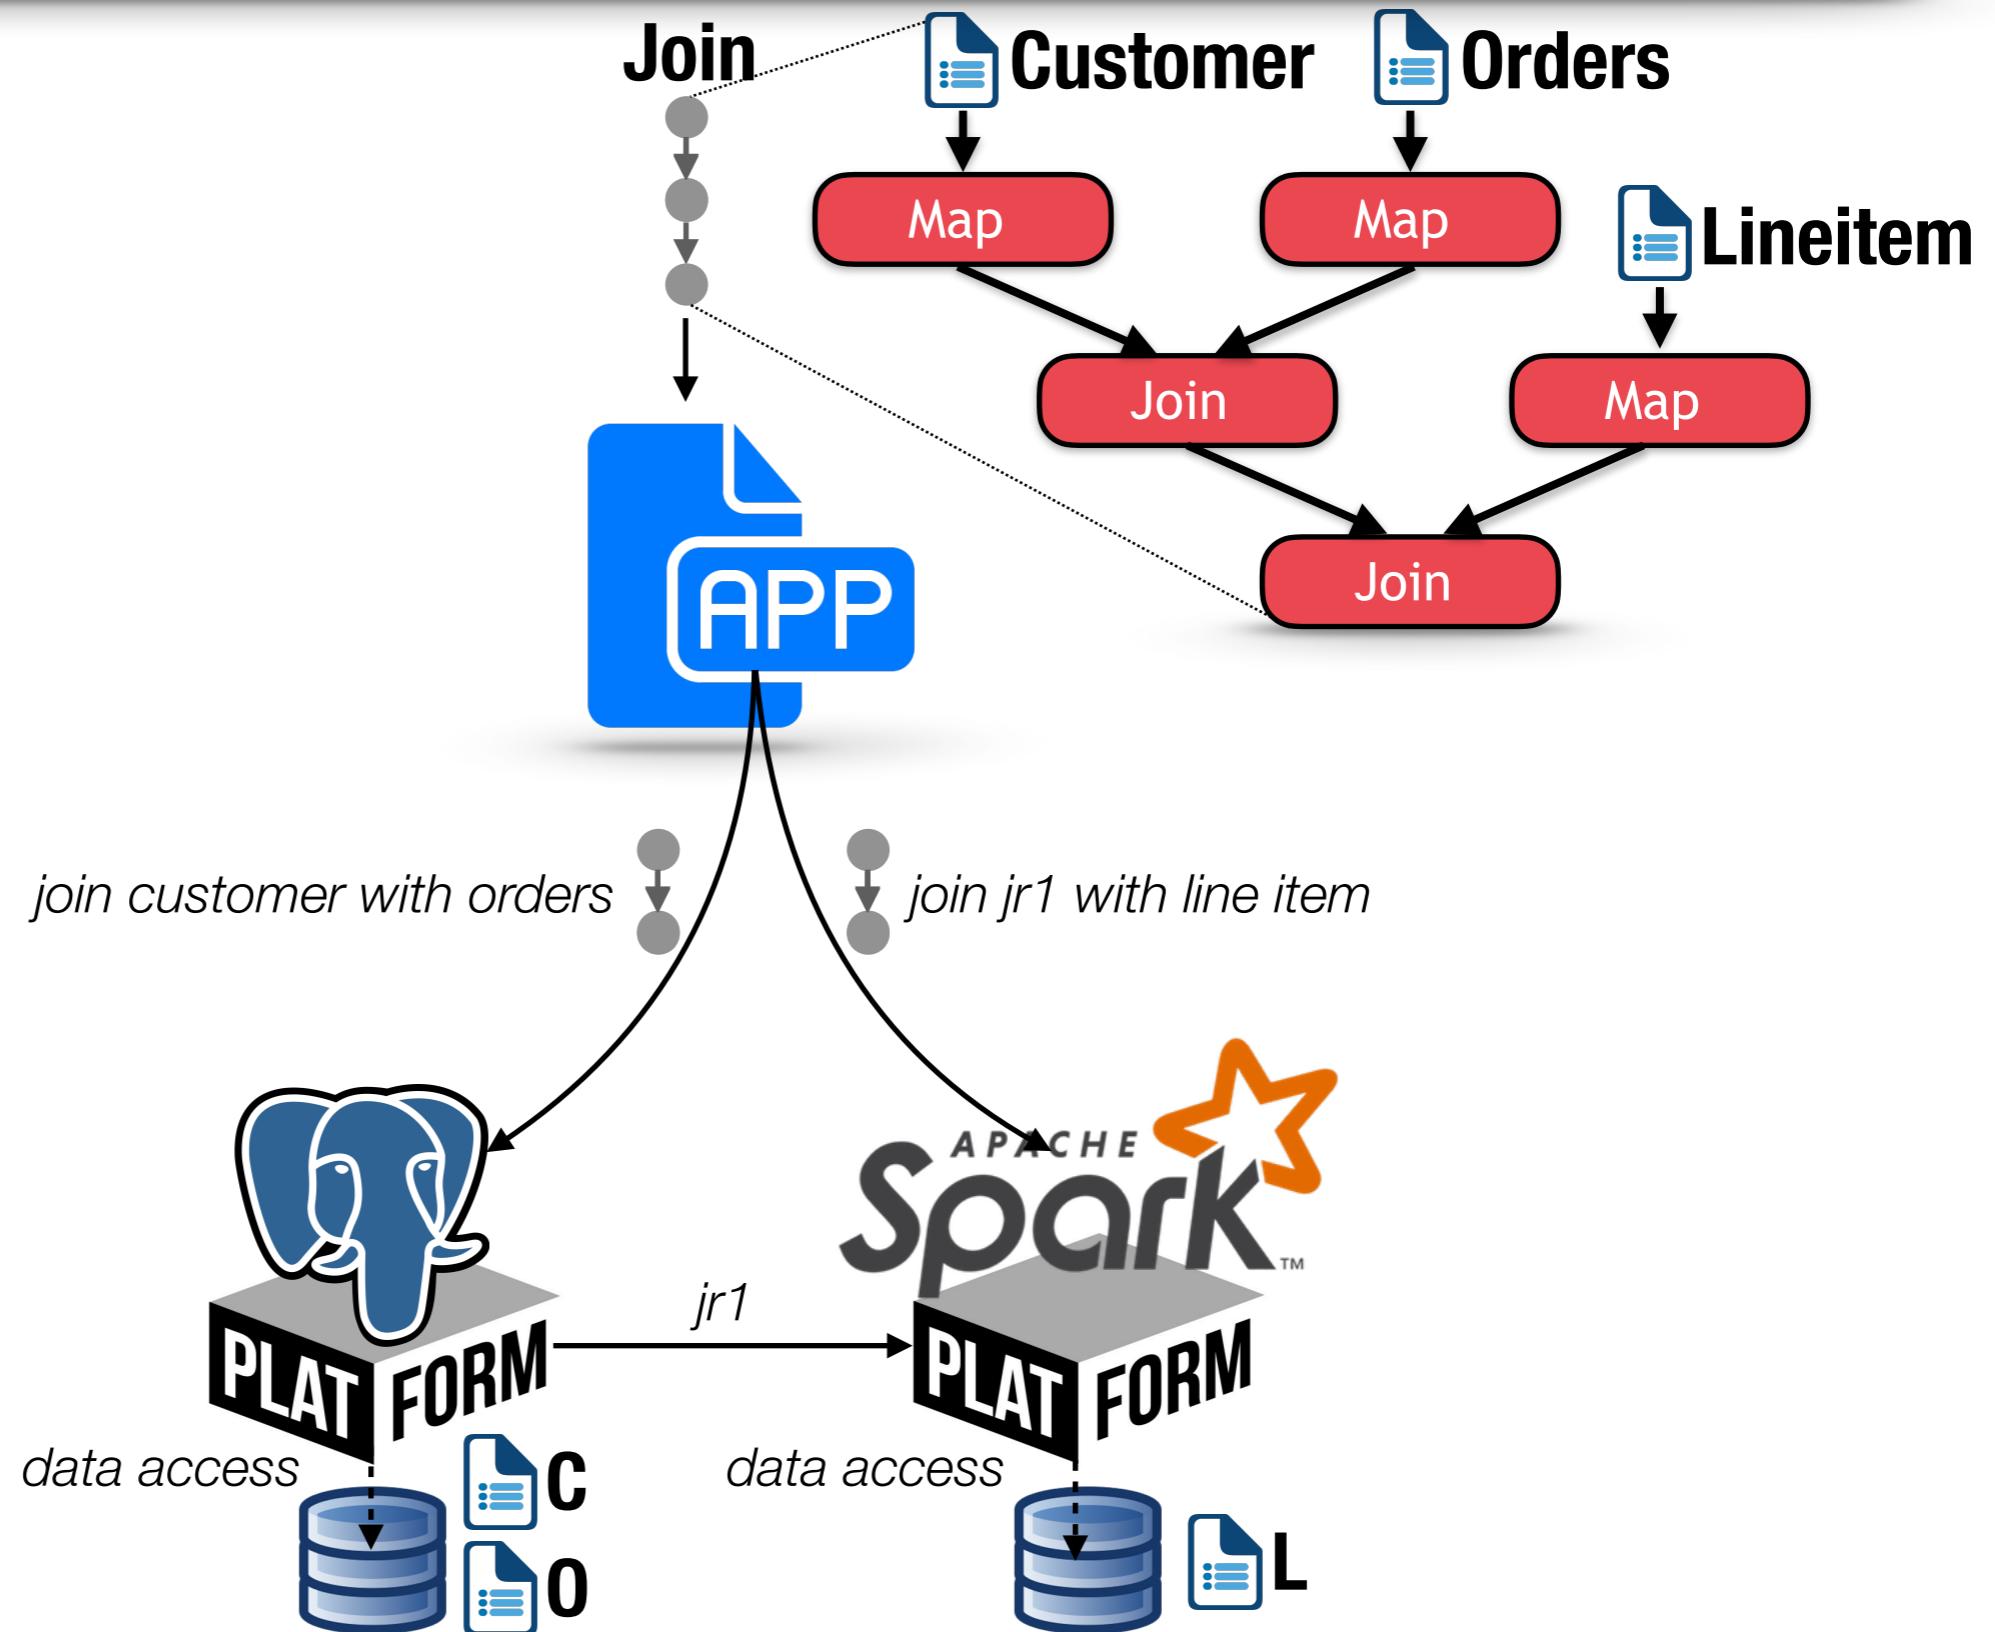
# Polystore

using **several** data processing platforms because data is spread

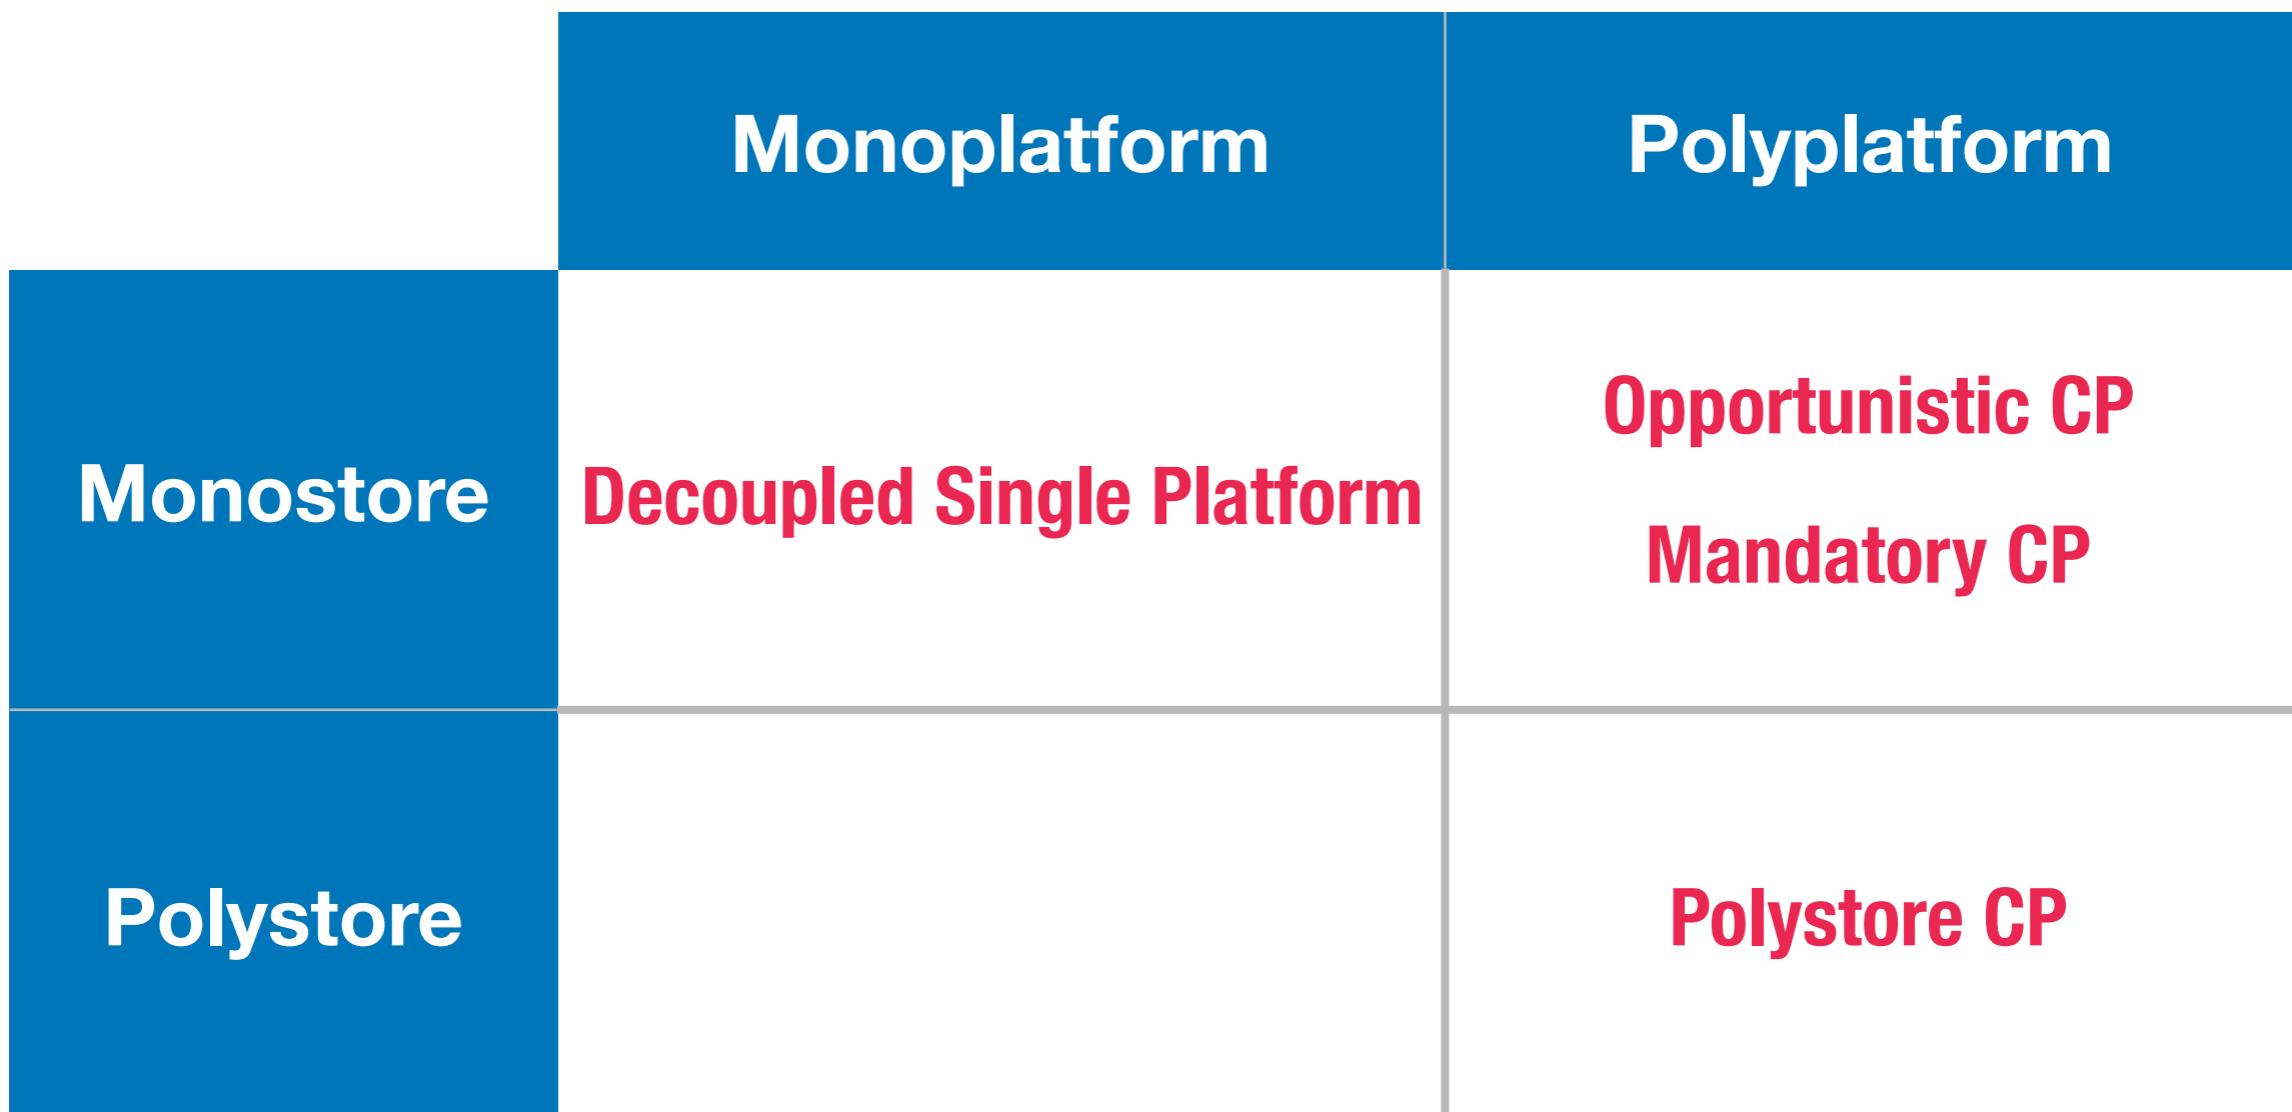


# Polystore

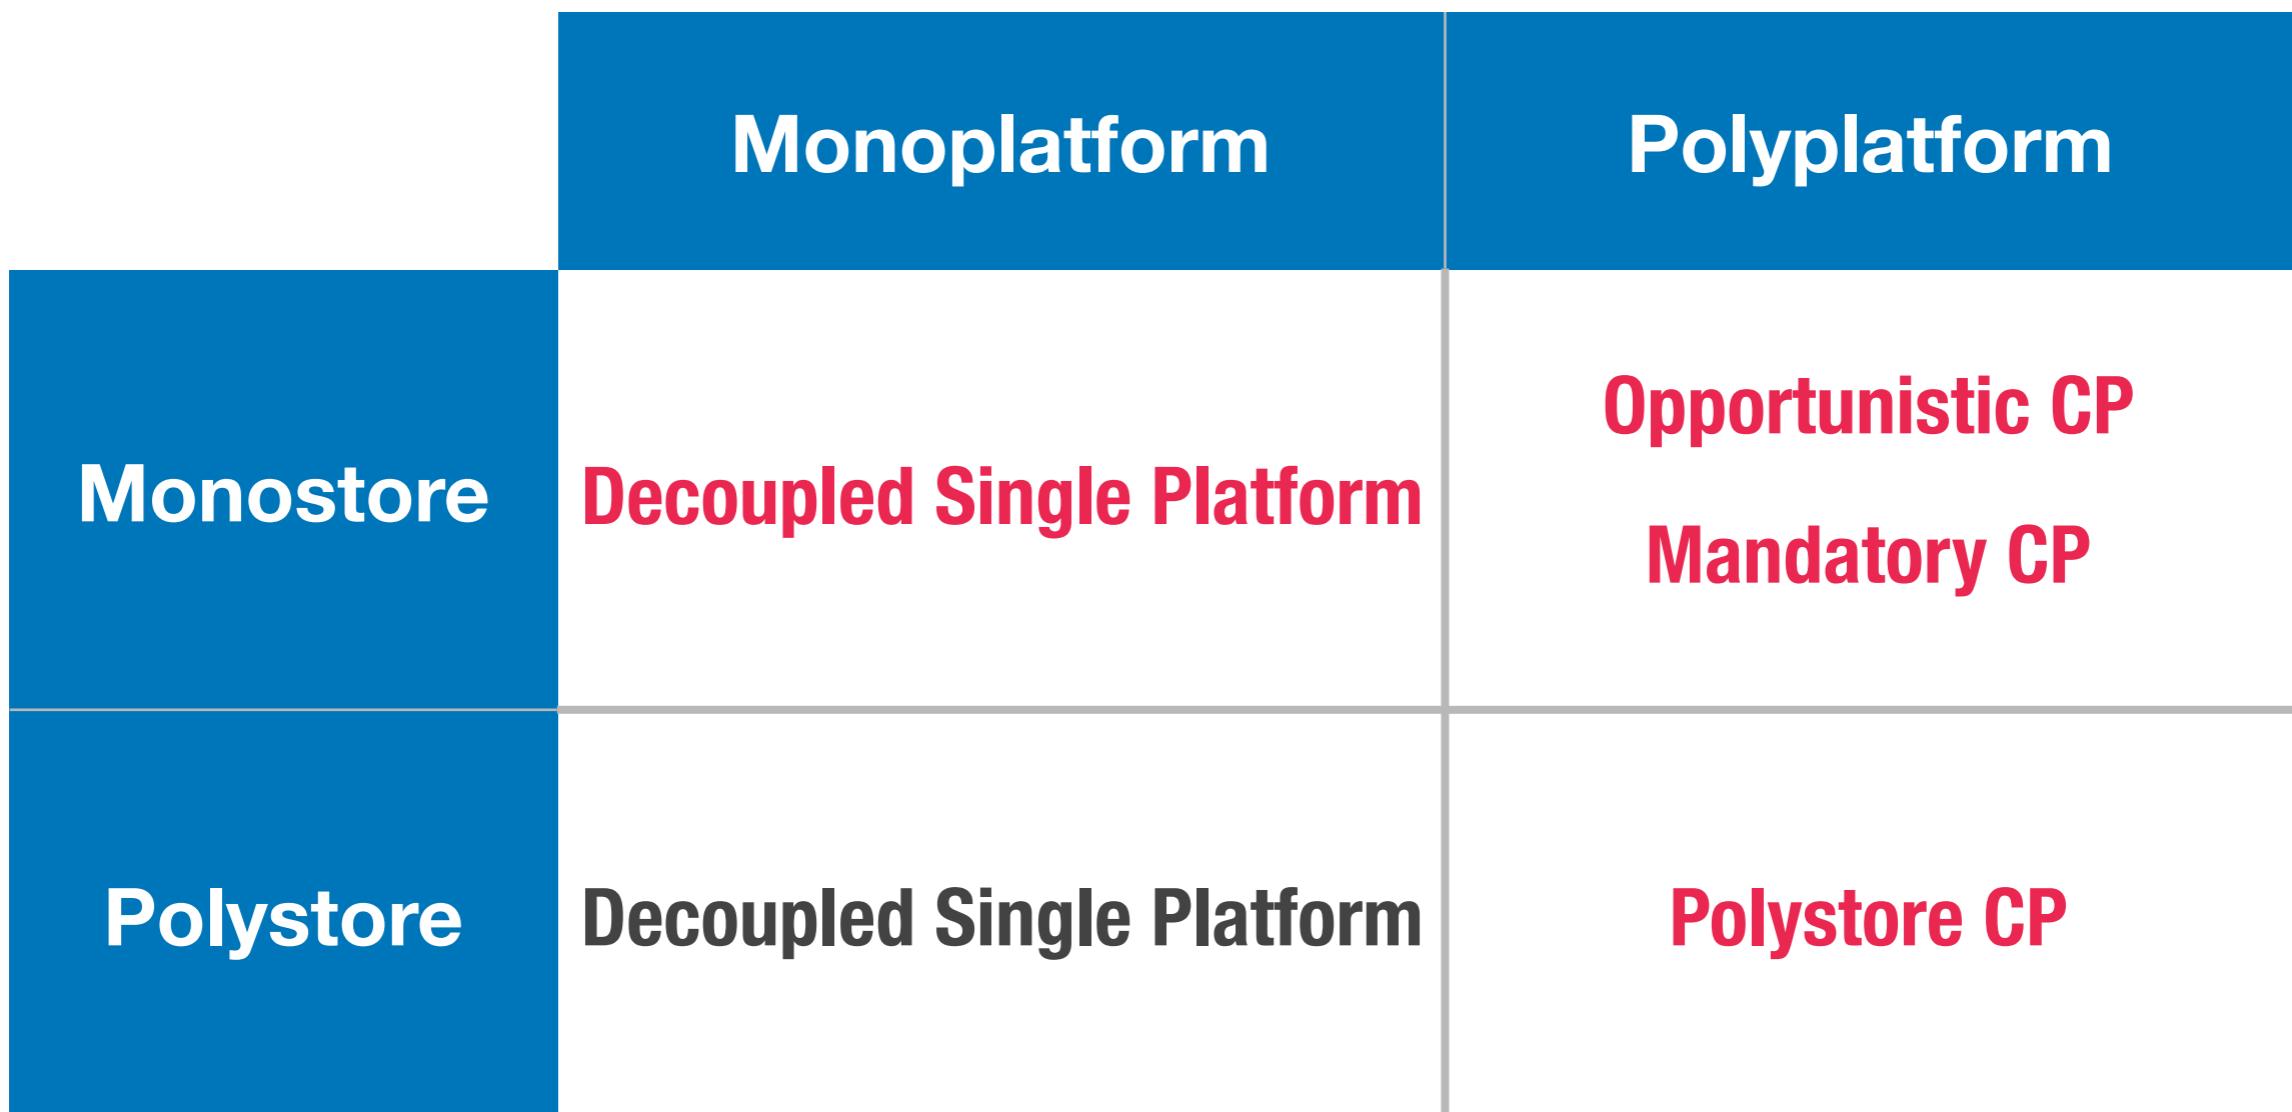
using **several** data processing platforms because data is spread



# System-Store Quadrant



# System-Store Quadrant



# Federated DBs vs Cross-Platform (II)

Polystore

Single Data Model

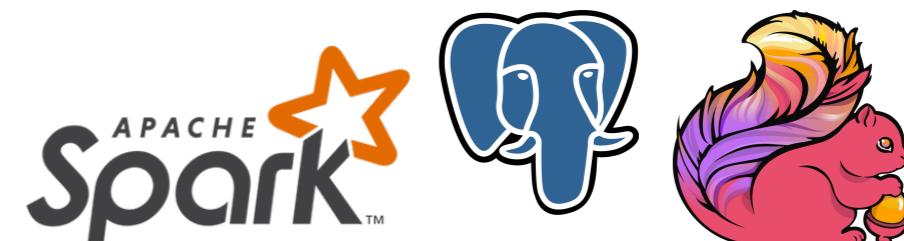
Federated DB System



Decoupled Single Platform  
Opportunistic CP  
Mandatory CP  
Polystore CP

Disparate Data Models

Cross-Platform System



# Cross-Platform Data Processing

Motivation

Summary

Use Cases

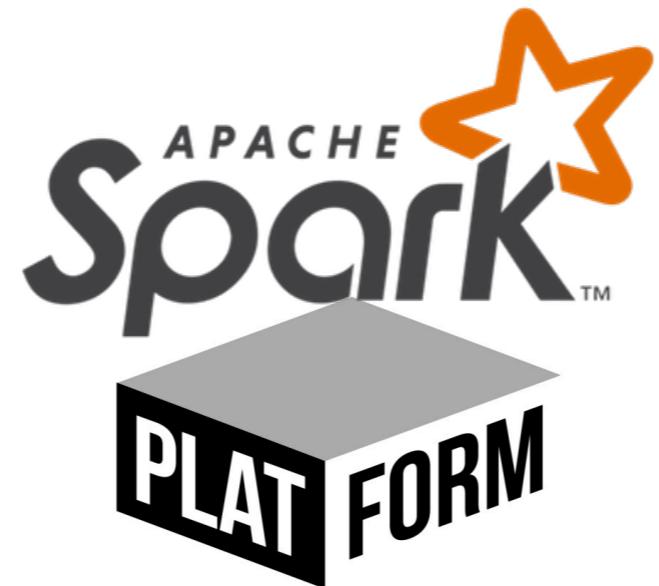
Current  
Efforts

Challenges

# Challenges

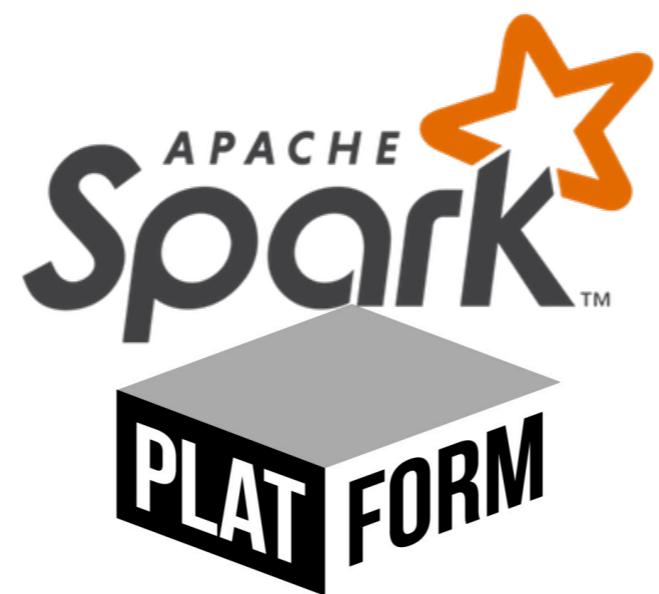


# Challenges



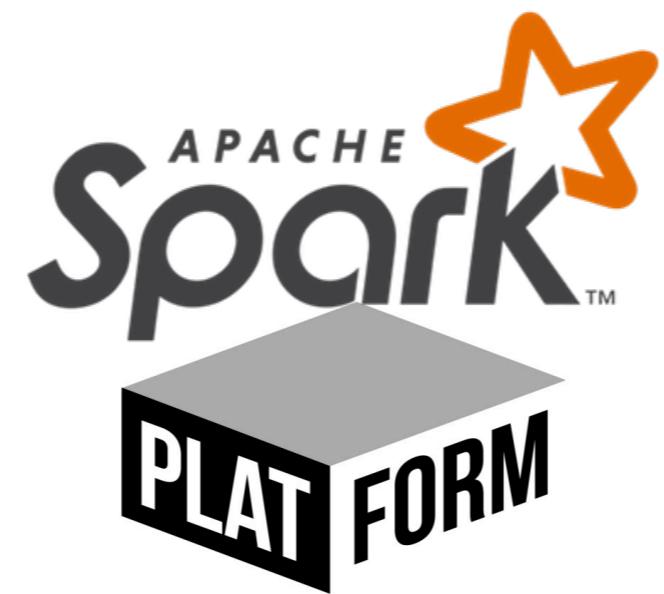
# Challenges

## 1 Decoupling Applications



# Challenges

## 1 Decoupling Applications



# Challenges

## 1 Decoupling Applications



# Challenges

## 1 Decoupling Applications



# Challenges

1 Decoupling Applications



3

Automatic Cross-Platform  
Data Processing



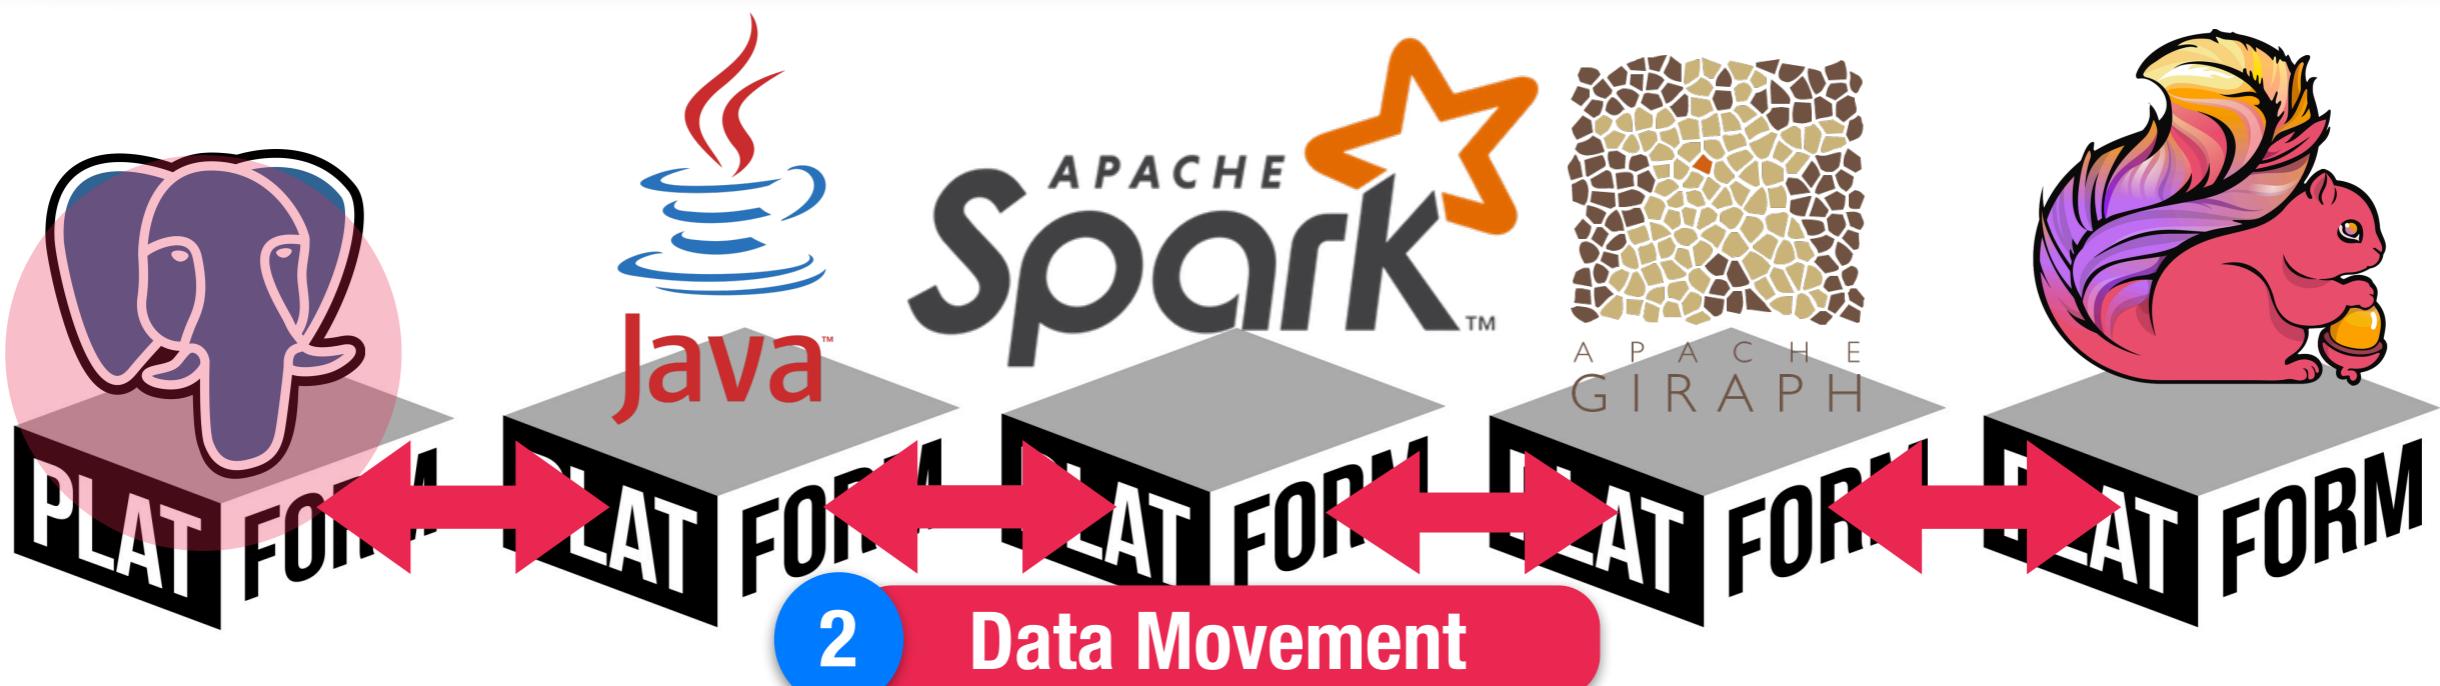
# Challenges

1 Decoupling Applications



3

Automatic Cross-Platform  
Data Processing



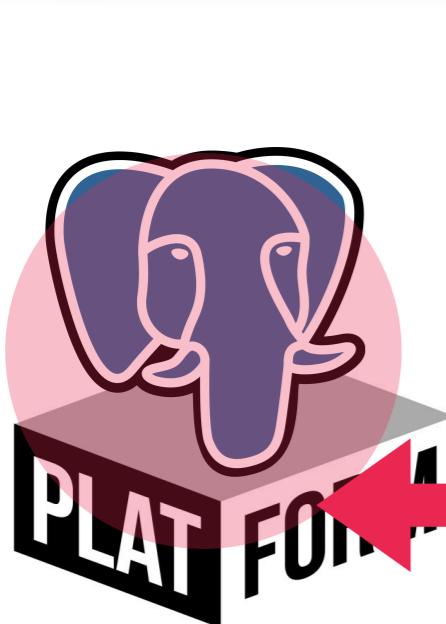
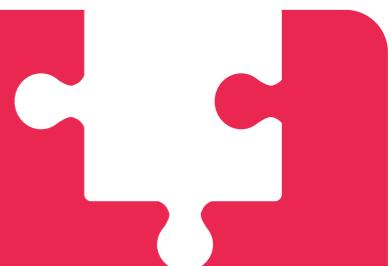
# Challenges

1 Decoupling Applications



3

Automatic Cross-Platform  
Data Processing



2

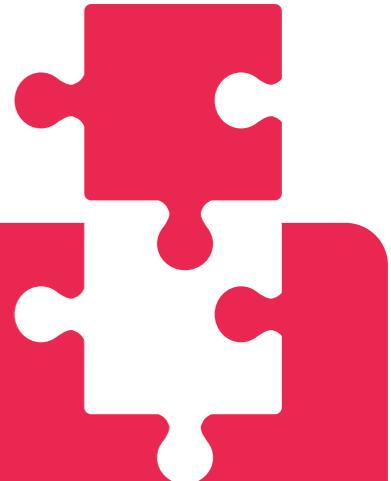
Data Movement

# Challenges

1 Decoupling Applications

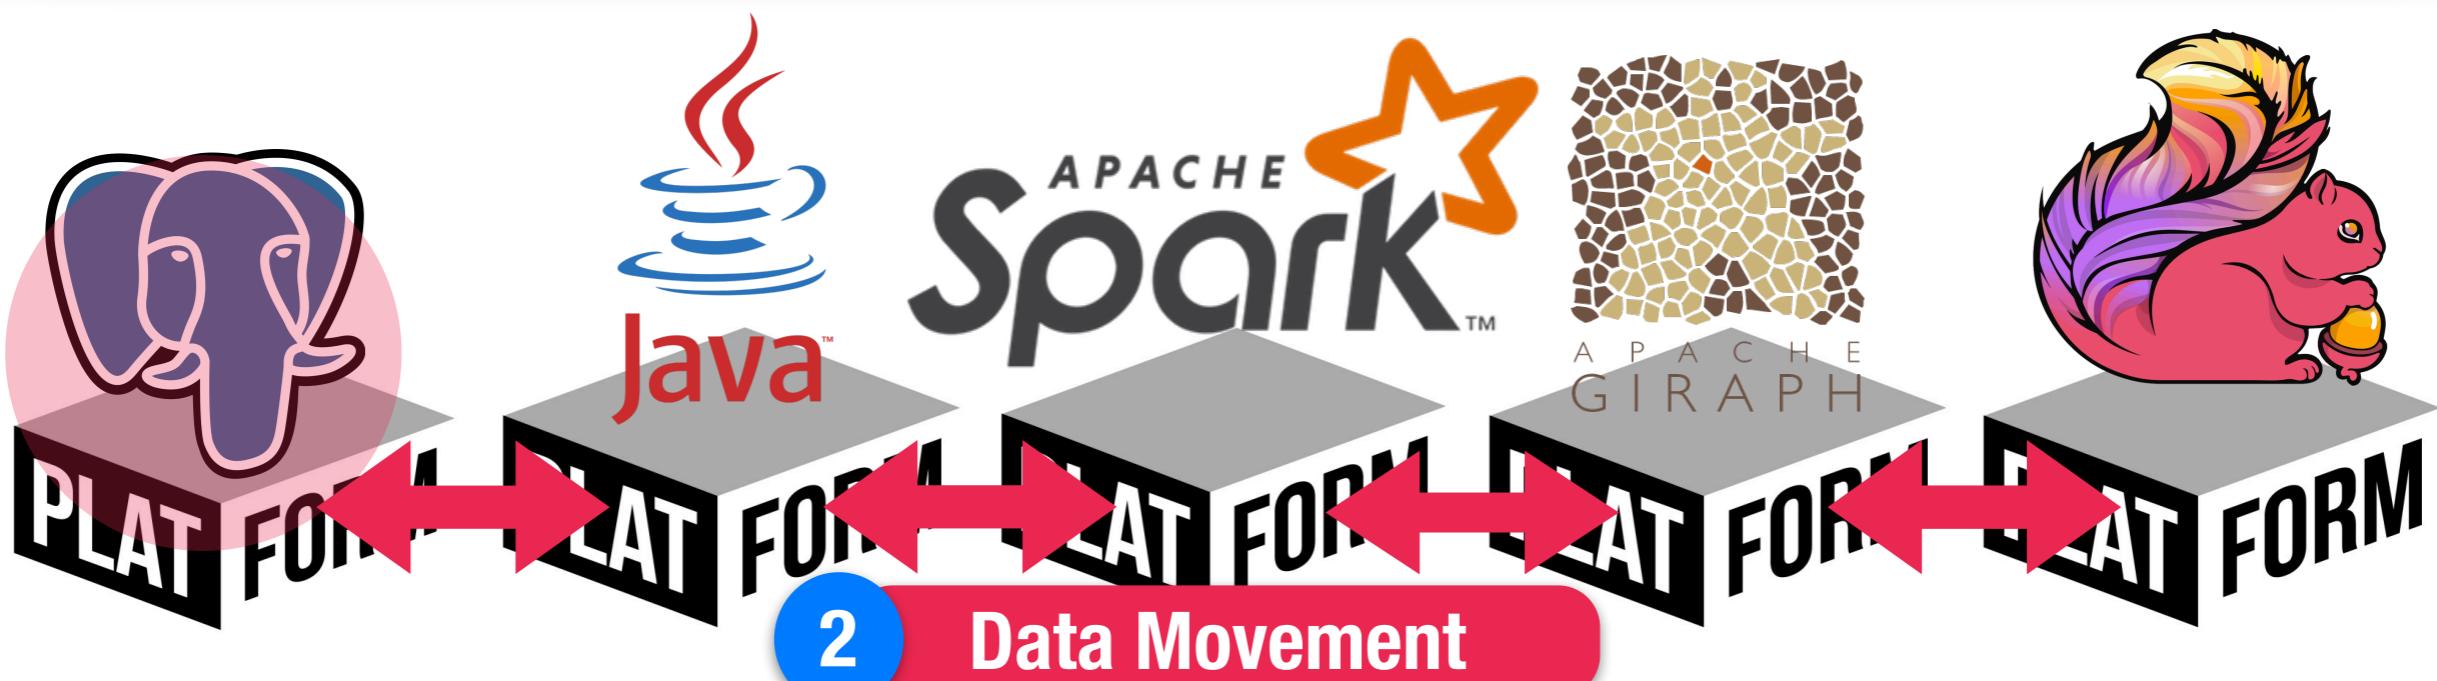


4 Extensibility



3

Automatic Cross-Platform  
Data Processing



# Cross-Platform Data Processing

Summary

Motivation

Use Cases

Current  
Efforts

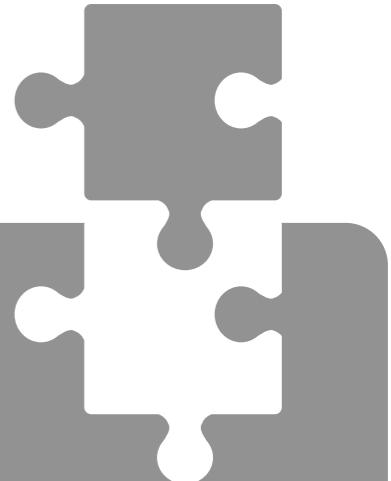
Challenges

# Challenges

1 Decoupling Applications

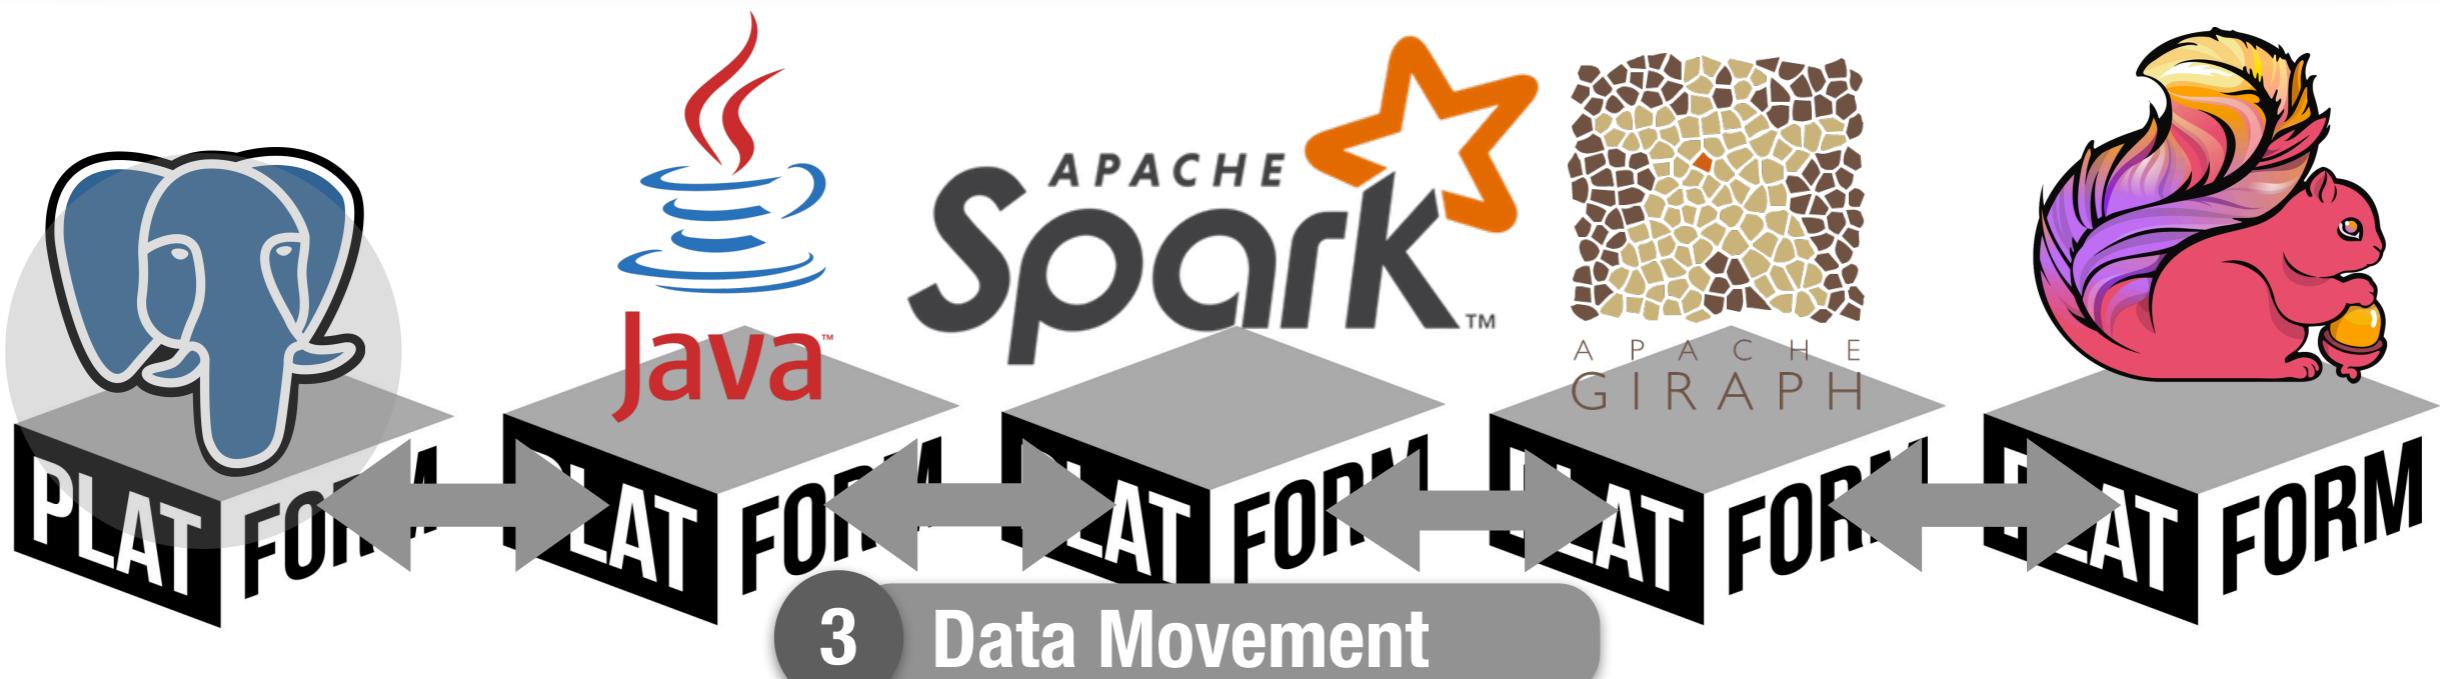


4 Extensibility



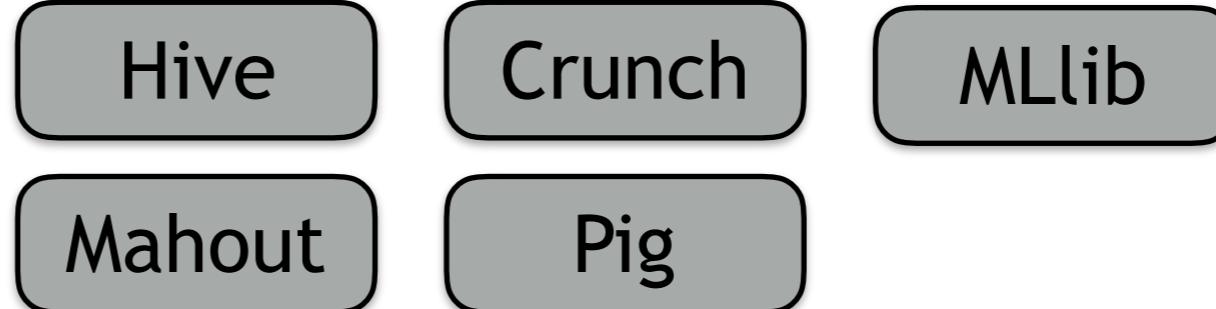
2

Automatic Cross-Platform  
Data Processing

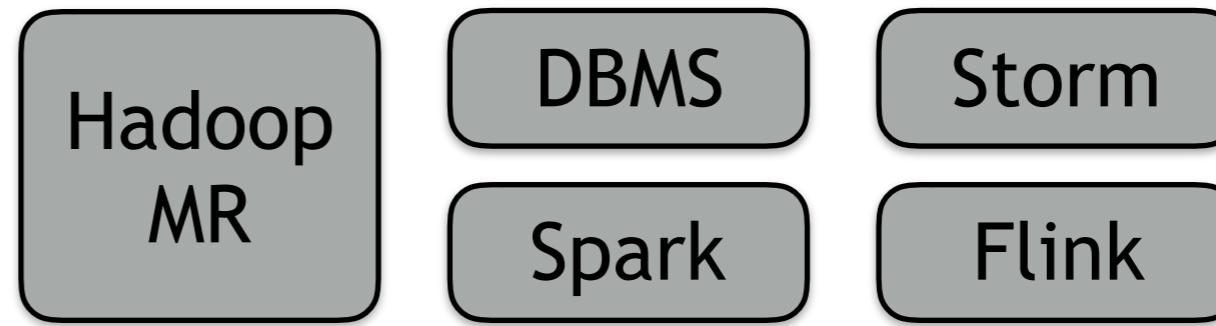


# Decoupling Applications

Applications  
/Frontends



Processing  
Platforms



Storage  
Engines



# Decoupling Applications

Applications  
/Frontends

Hive

Crunch

MLlib

Mahout

Pig



Processing  
Platforms

Hadoop  
MR

DBMS

Storm

Spark

Flink

Storage  
Engines

HDFS

S3

Local FS

# Decoupling Applications

Applications  
/Frontends

Hive

Crunch

MLlib

Mahout

Pig

## Cross-Platform System

Processing  
Platforms

Hadoop  
MR

DBMS

Storm

Spark

Flink

Storage  
Engines

HDFS

S3

Local FS

# Decoupling Applications

Applications  
/Frontends

Hive

Crunch

MLlib

Mahout

Pig

**Musketeer**

**Rheem**

**BigDawg**

**Apache Beam**

**QoX**

**DBMS+**

Processing  
Platforms

Hadoop  
MR

DBMS

Storm

Spark

Flink

Storage  
Engines

HDFS

S3

Local FS

# Decoupling Applications

Applications  
/Frontends

Hive

Crunch

MLlib

Mahout

Pig

**Musketeer**

**Rheem**

**BigDawg**

**Apache Beam**

**QoX**

**DBMS+**

Processing  
Platforms

Hadoop  
MR

DBMS

Storm

Spark

Flink

Storage  
Engines

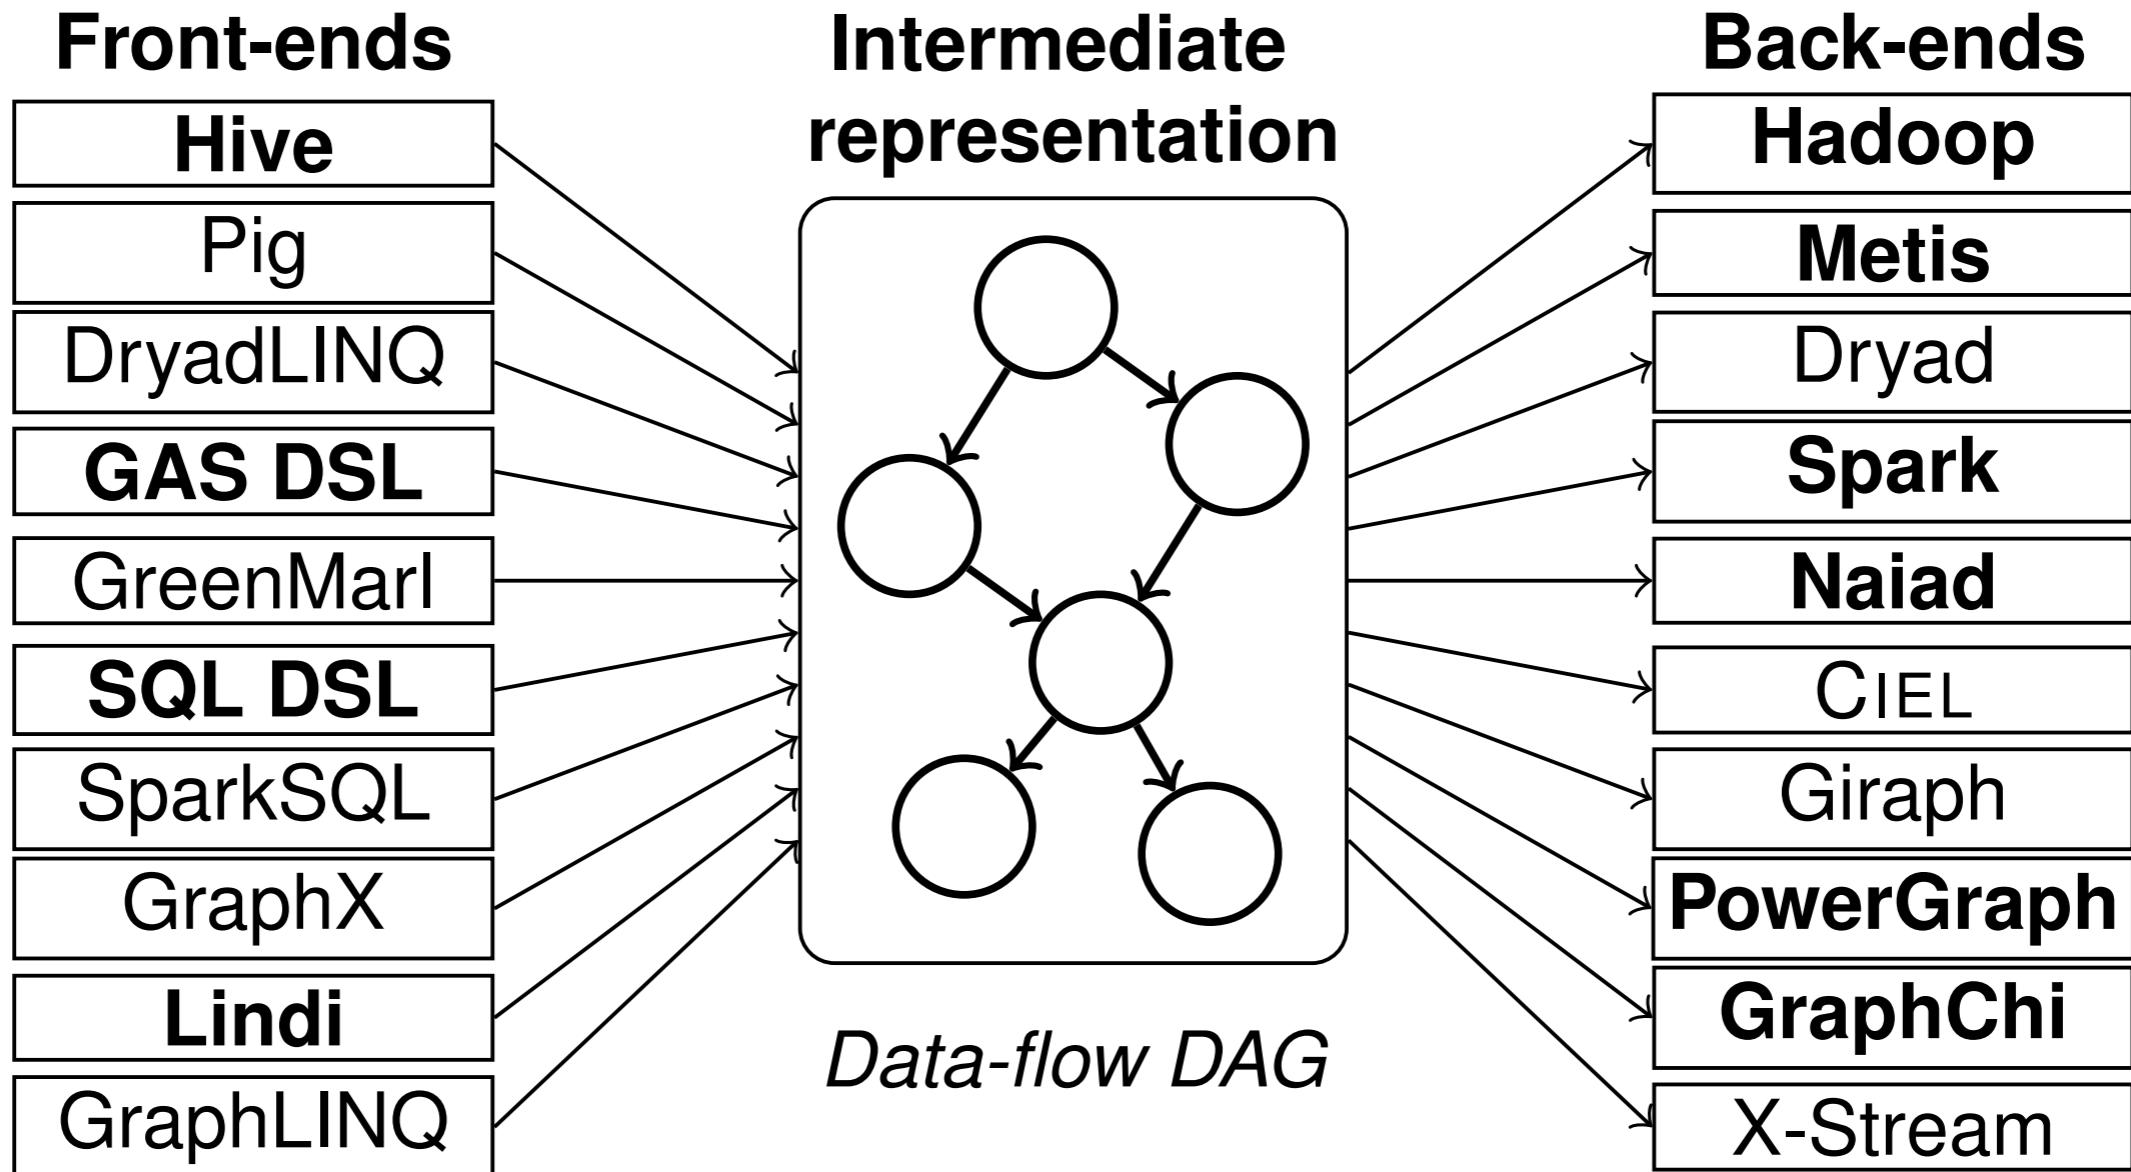
HDFS

S3

Local FS

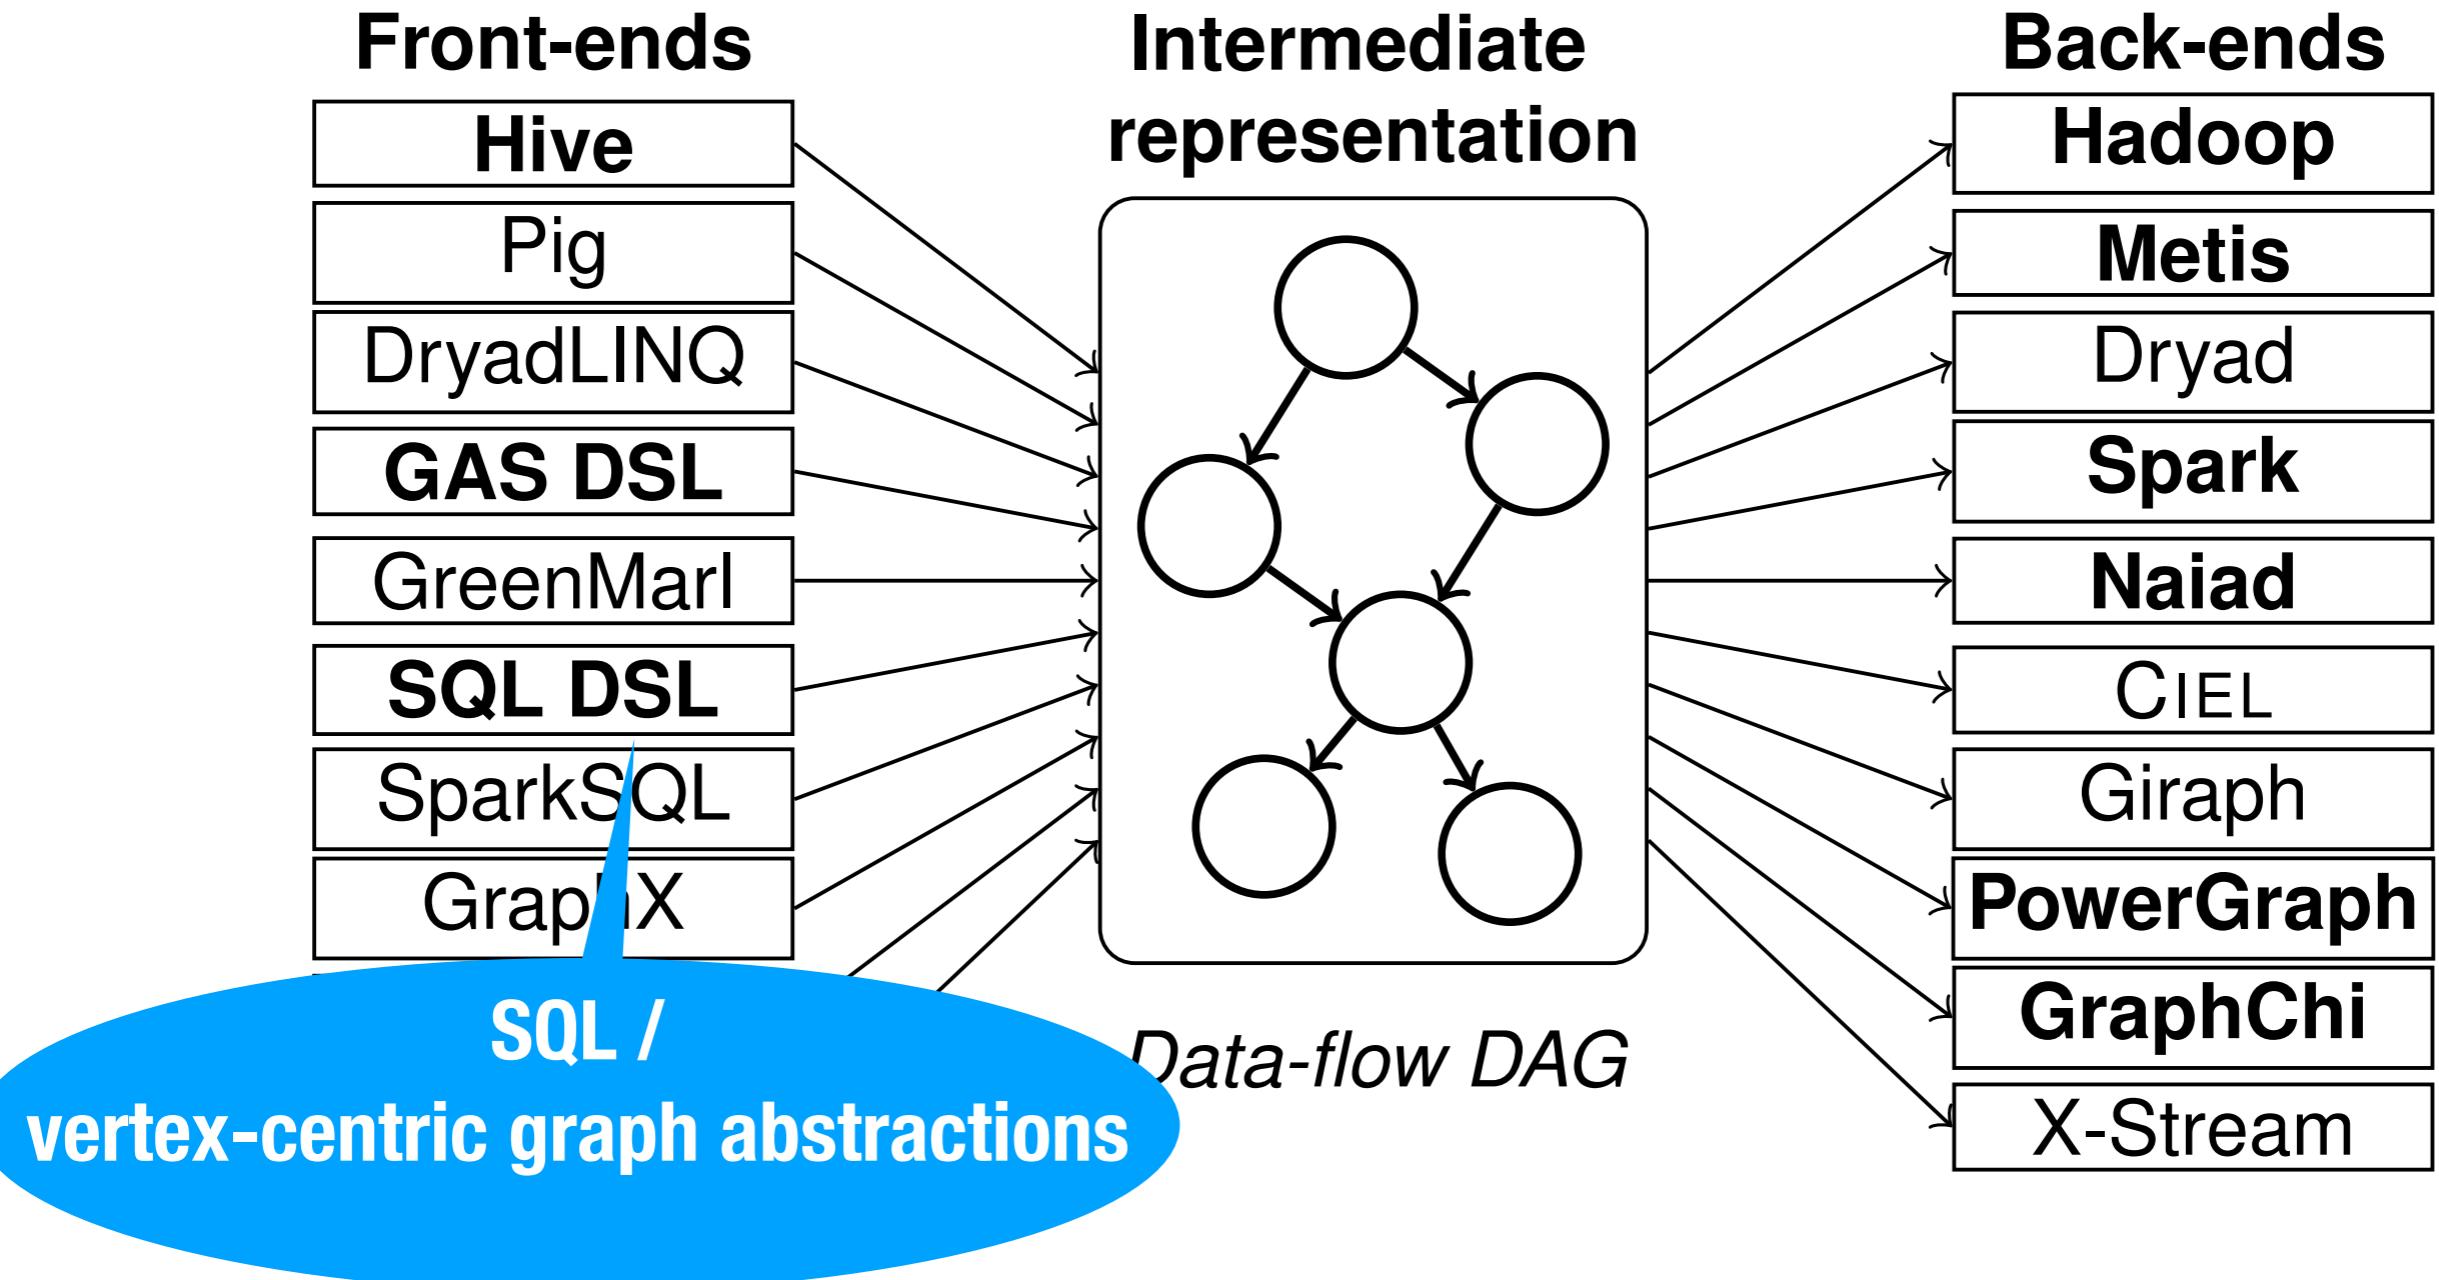
# Musketeer

DAG-based *intermediate representation*



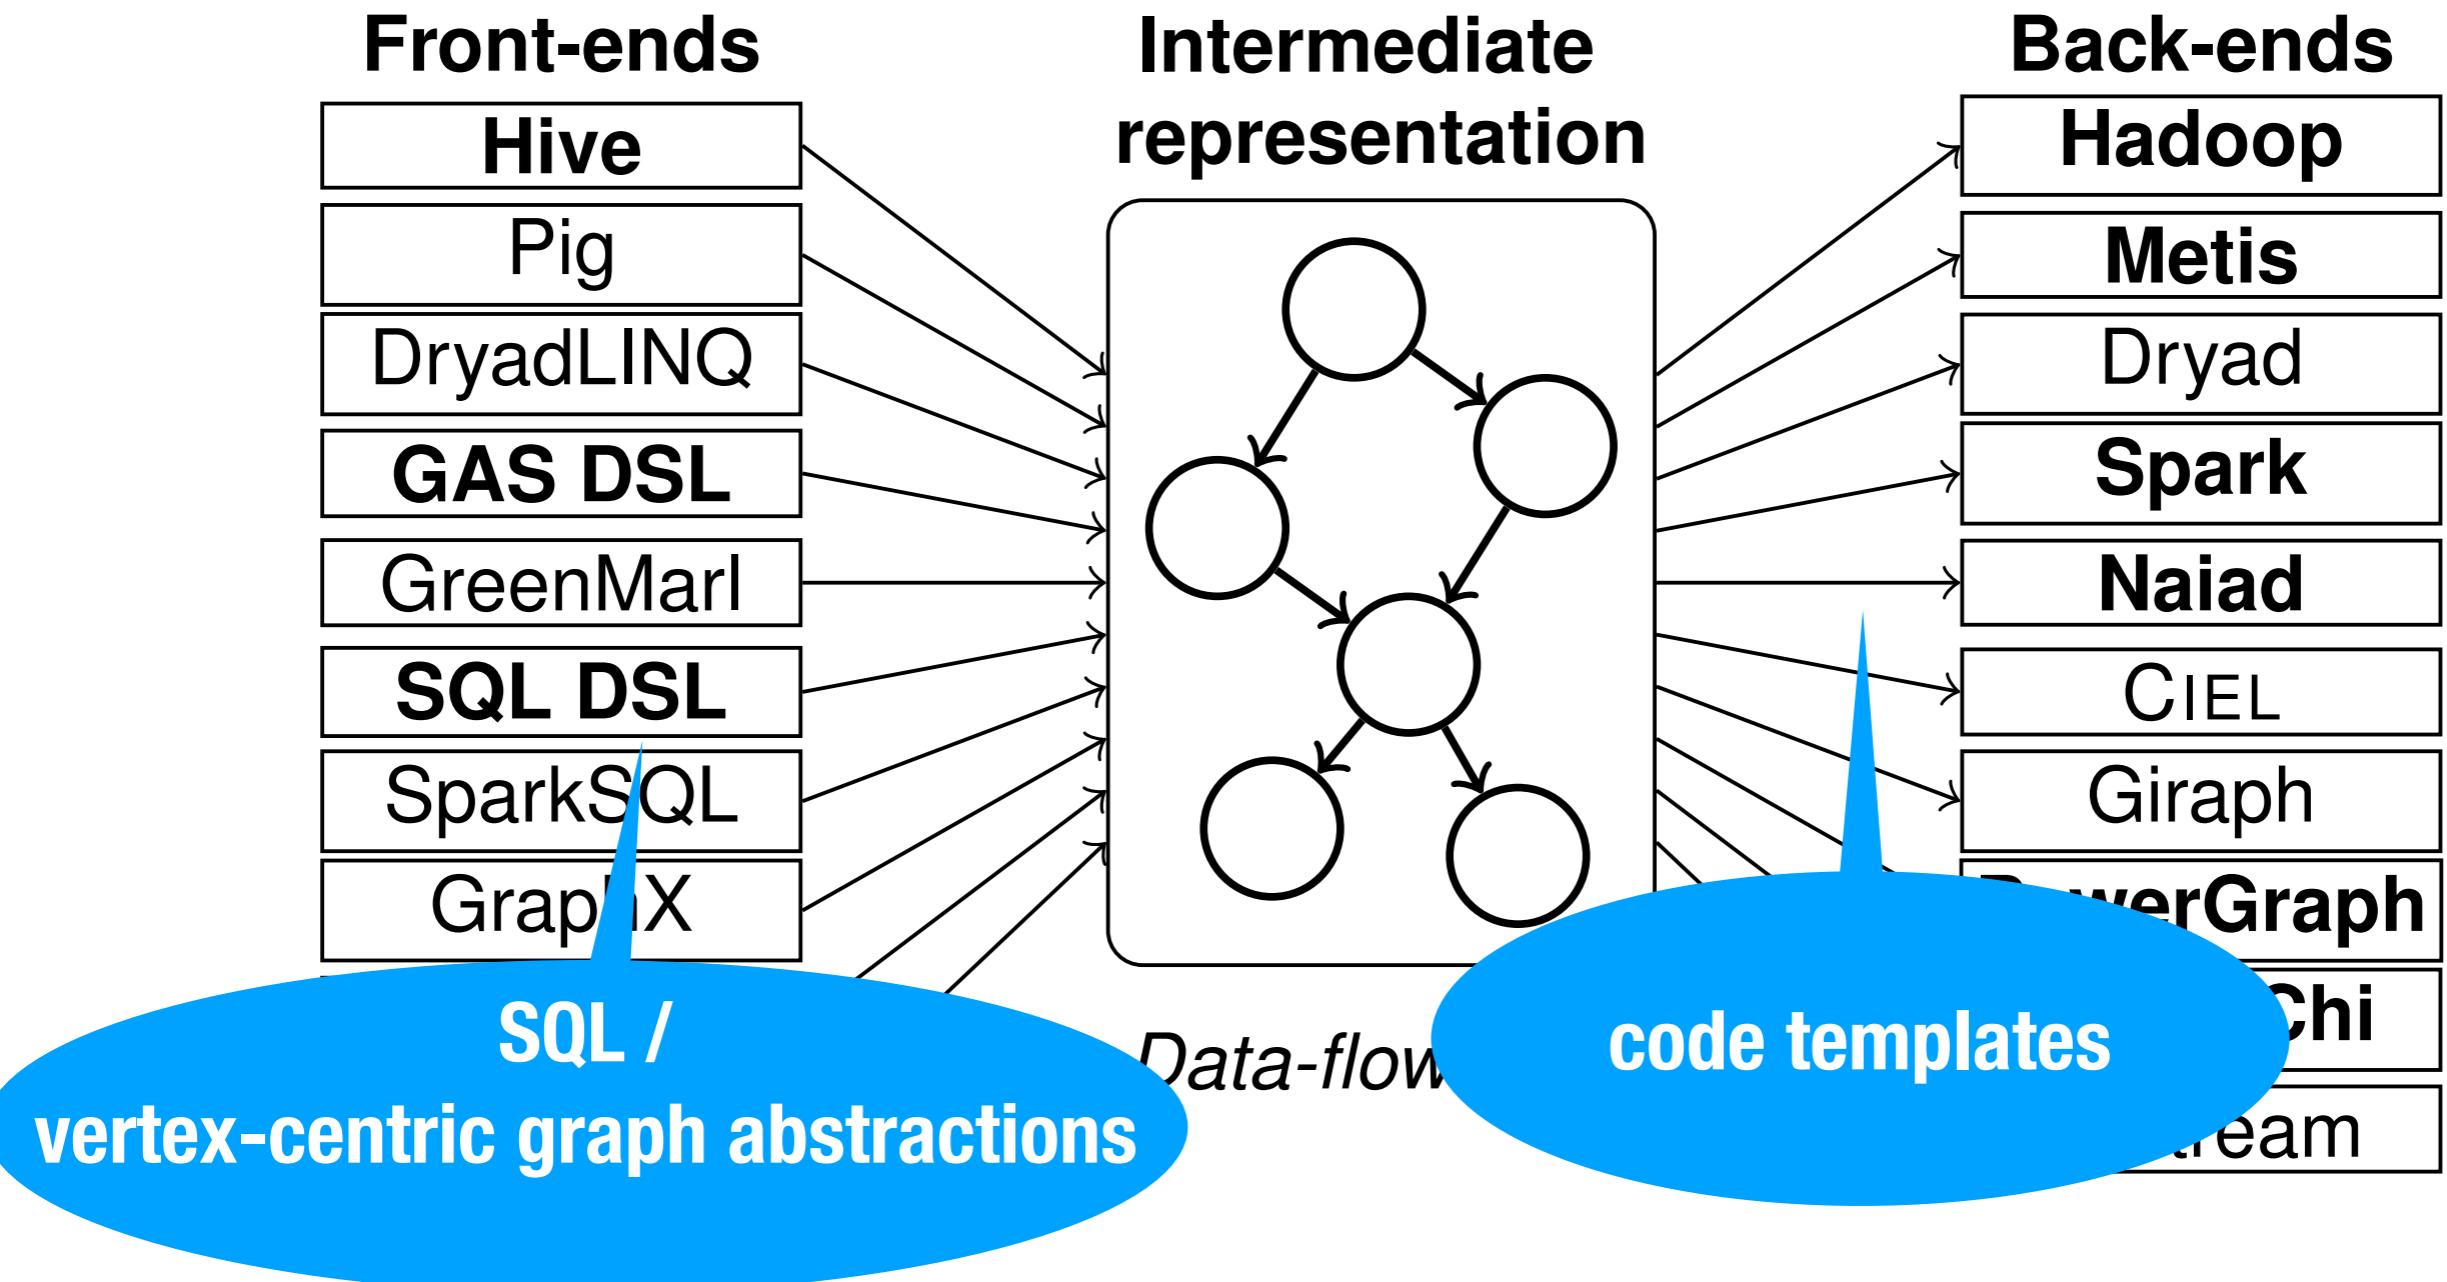
# Musketeer

DAG-based *intermediate representation*



# Musketeer

DAG-based *intermediate representation*



# Decoupling Features

	Processing Model	CP Granularity	Mappings
Musketeer	Dataflow DAG + Loop	Operator	Code templates

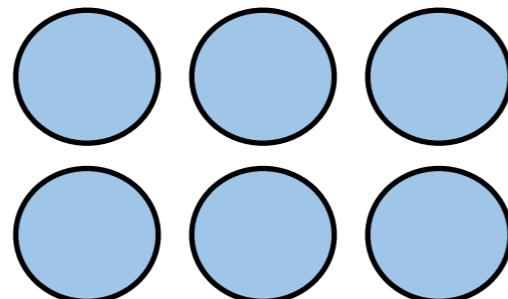
# Rheem & Ires

***2 layers of operators***

platform-agnostic

**Rheem operators (Rheem)**

**Abstract operators (Ires)**



# Rheem & Ires

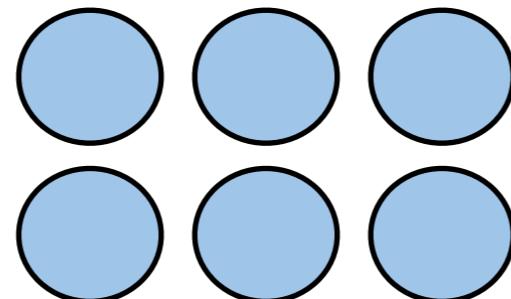
***2 layers of operators***



platform-agnostic

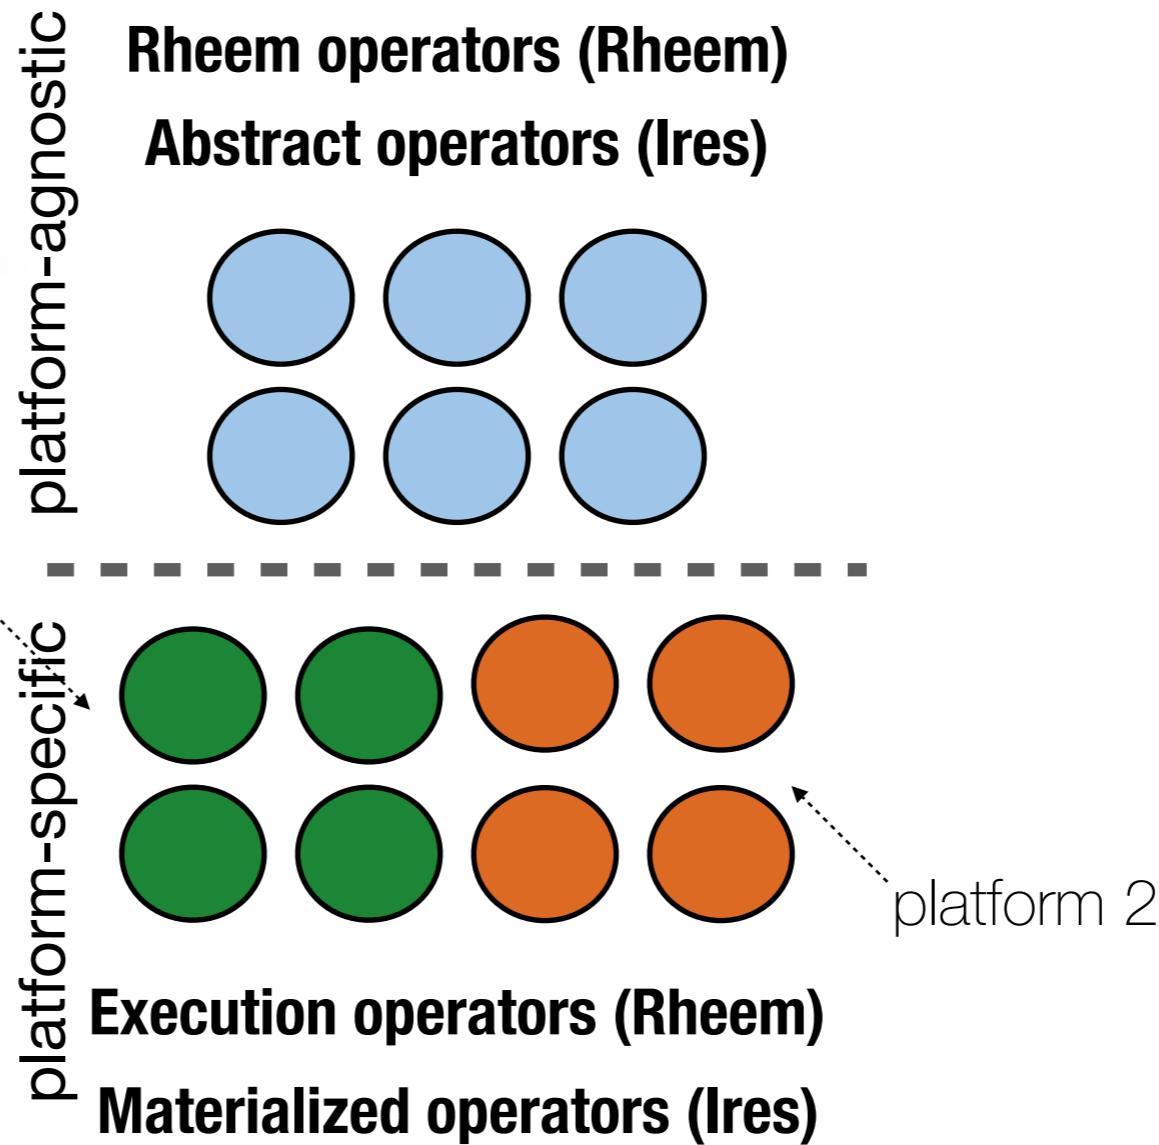
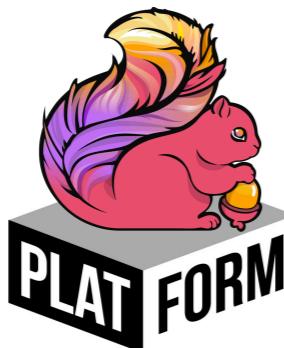
**Rheem operators (Rheem)**

**Abstract operators (Ires)**



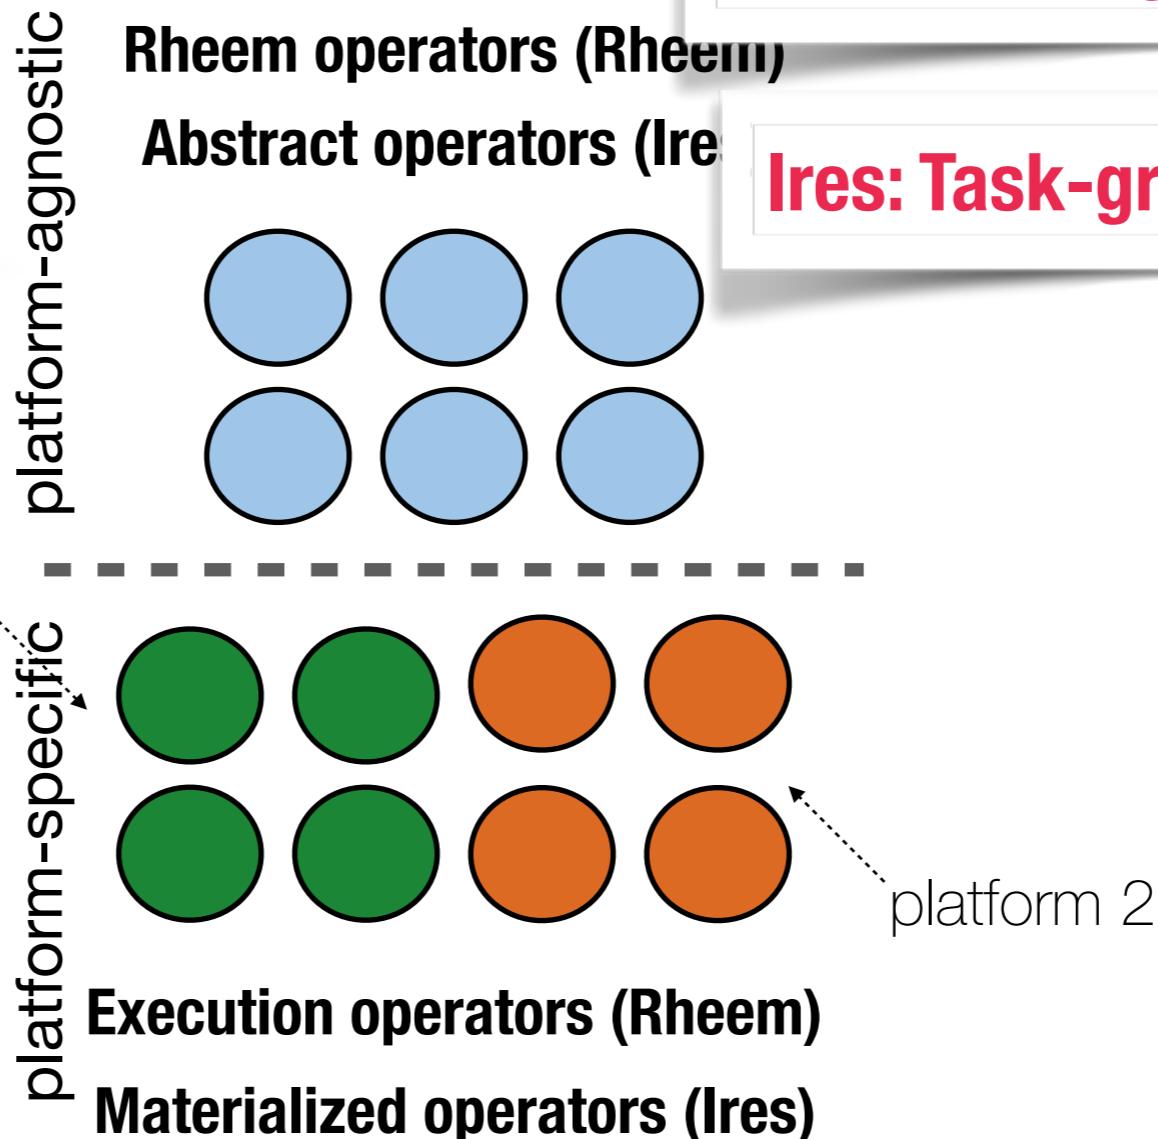
# Rheem & Ires

***2 layers of operators***



# Rheem & Ires

**2 layers of operators**

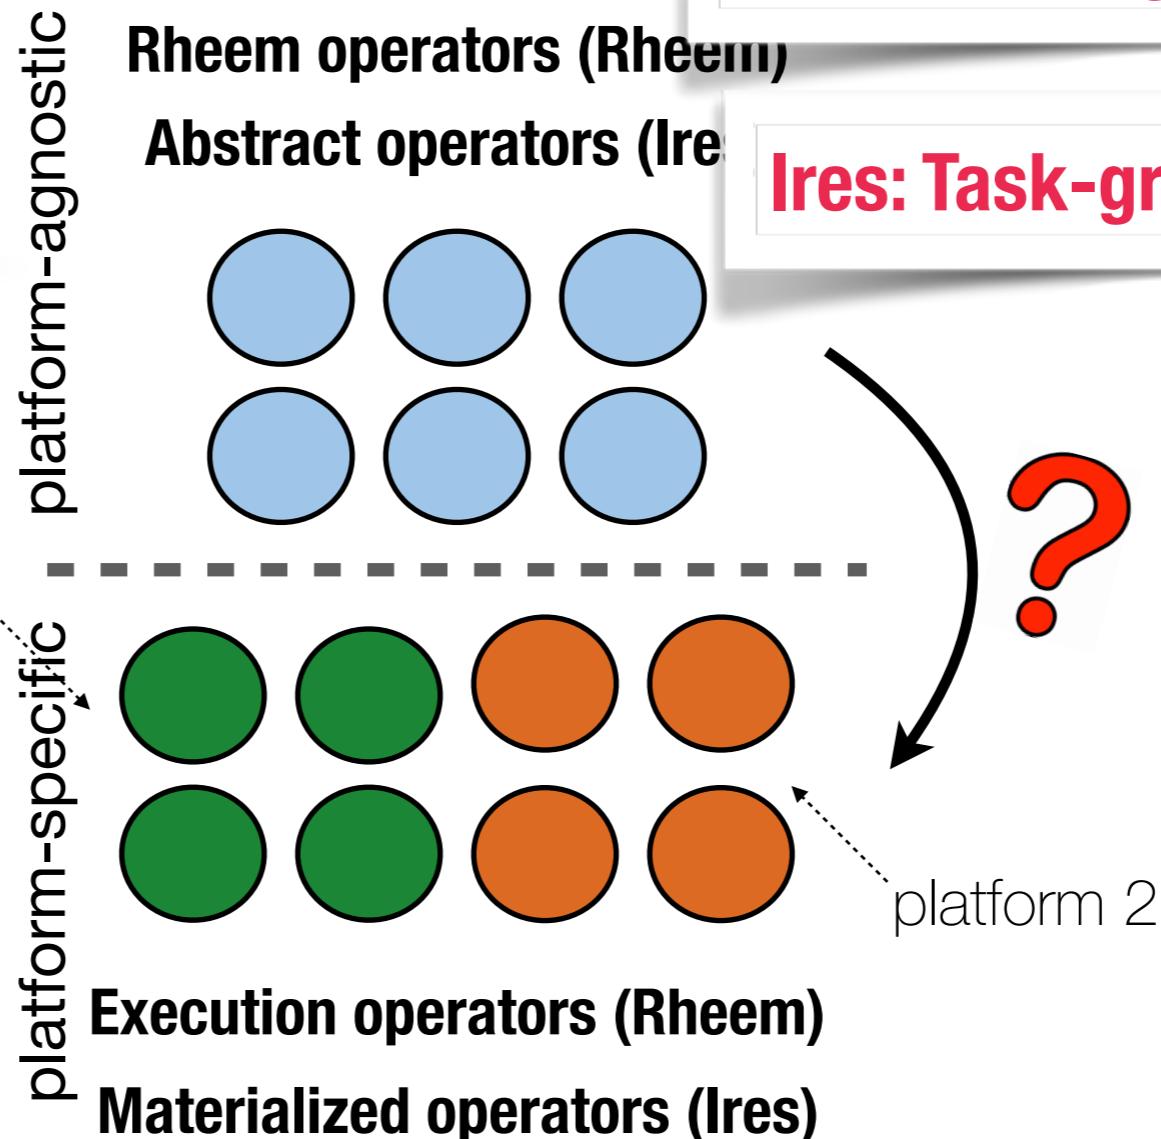


**Rheem: fine-granular operators**

**Ires: Task-granular operators**

# Rheem & Ires

**2 layers of operators**



**Rheem: fine-granular operators**

**Rheem operators (Rheem)**

**Abstract operators (Ires)**

**Ires: Task-granular operators**

# Rheem & Ires

***Graph-based*** mappings

***Metadata-based*** mappings

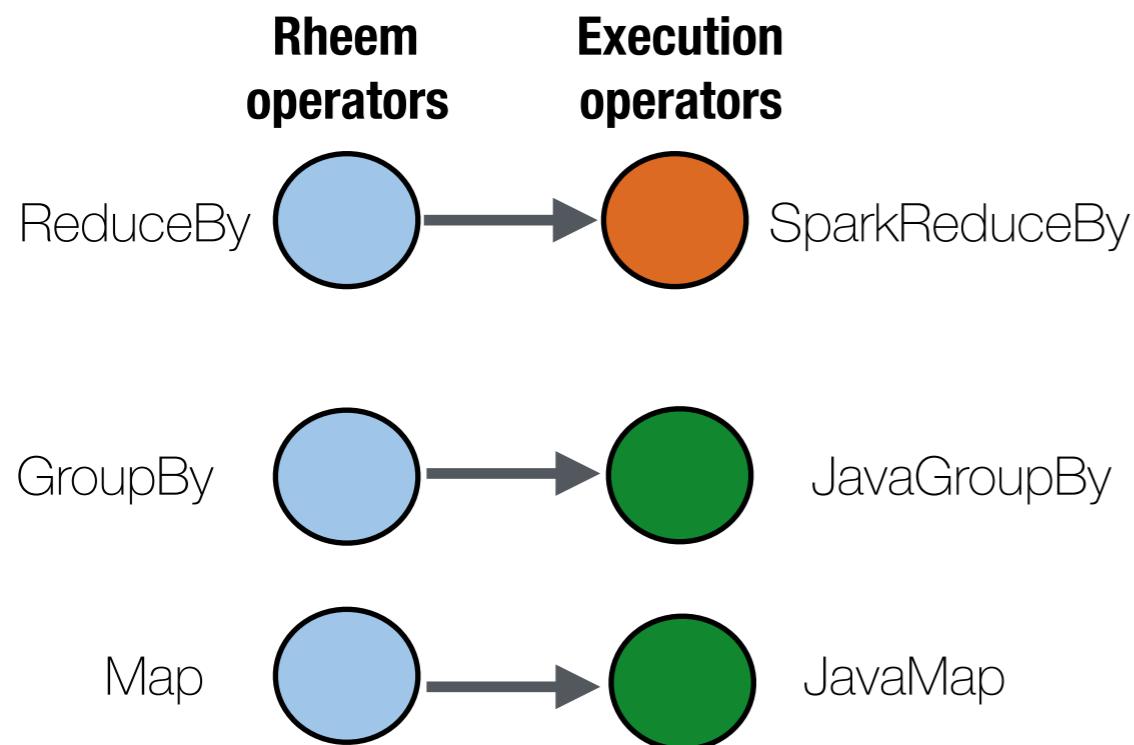
# Rheem

&

# Ires

***Graph-based*** mappings

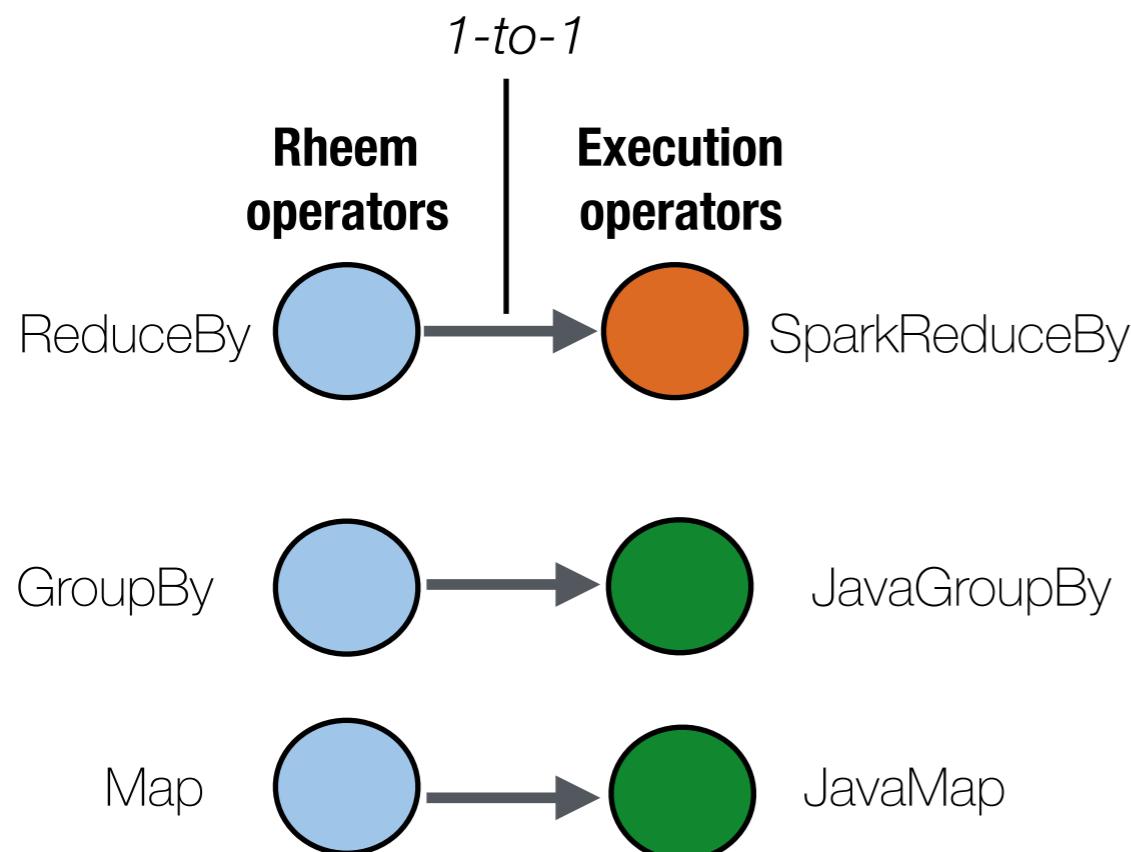
***Metadata-based*** mappings



# Rheem & Ires

**Graph-based** mappings

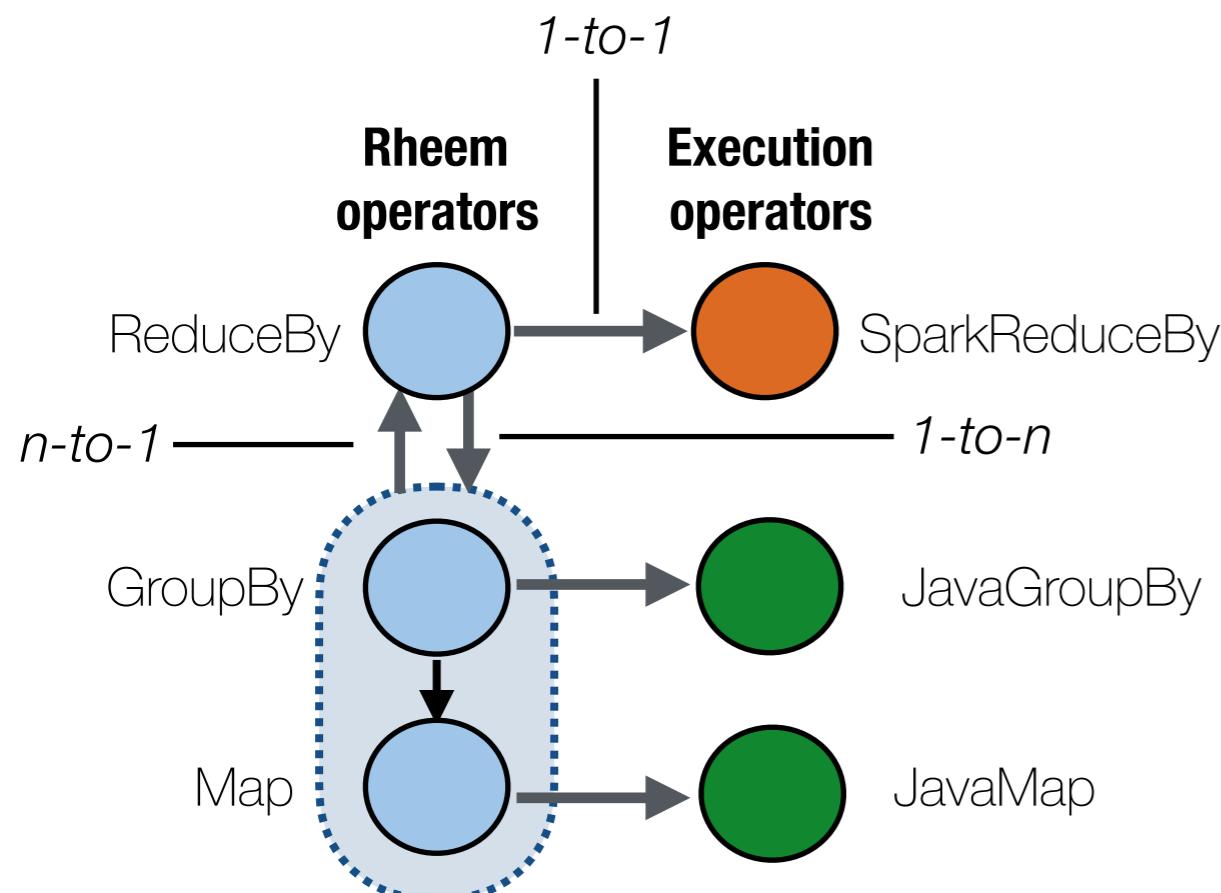
**Metadata-based** mappings



# Rheem & Ires

**Graph-based** mappings

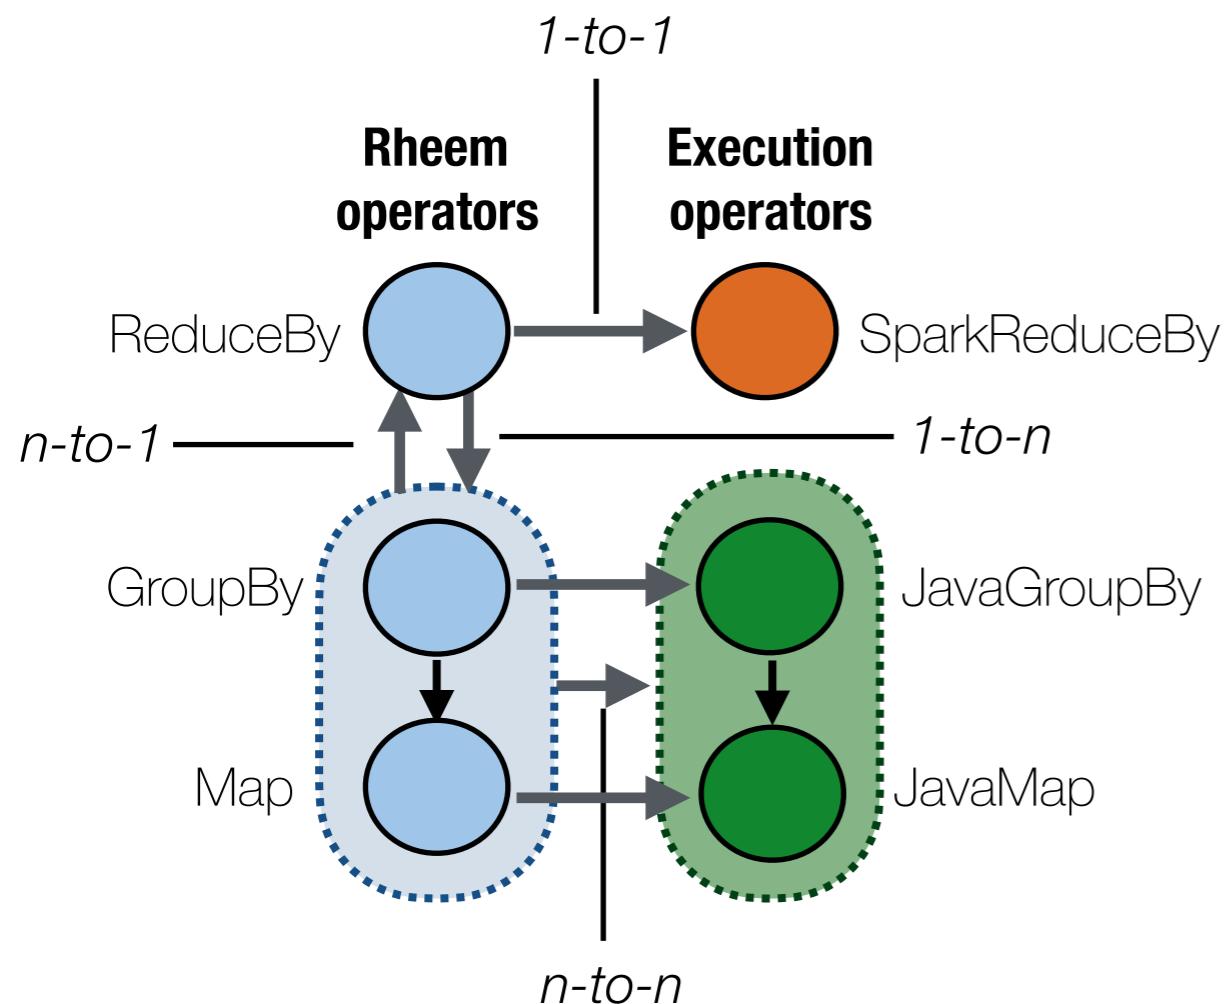
**Metadata-based** mappings



# Rheem & Ires

**Graph-based** mappings

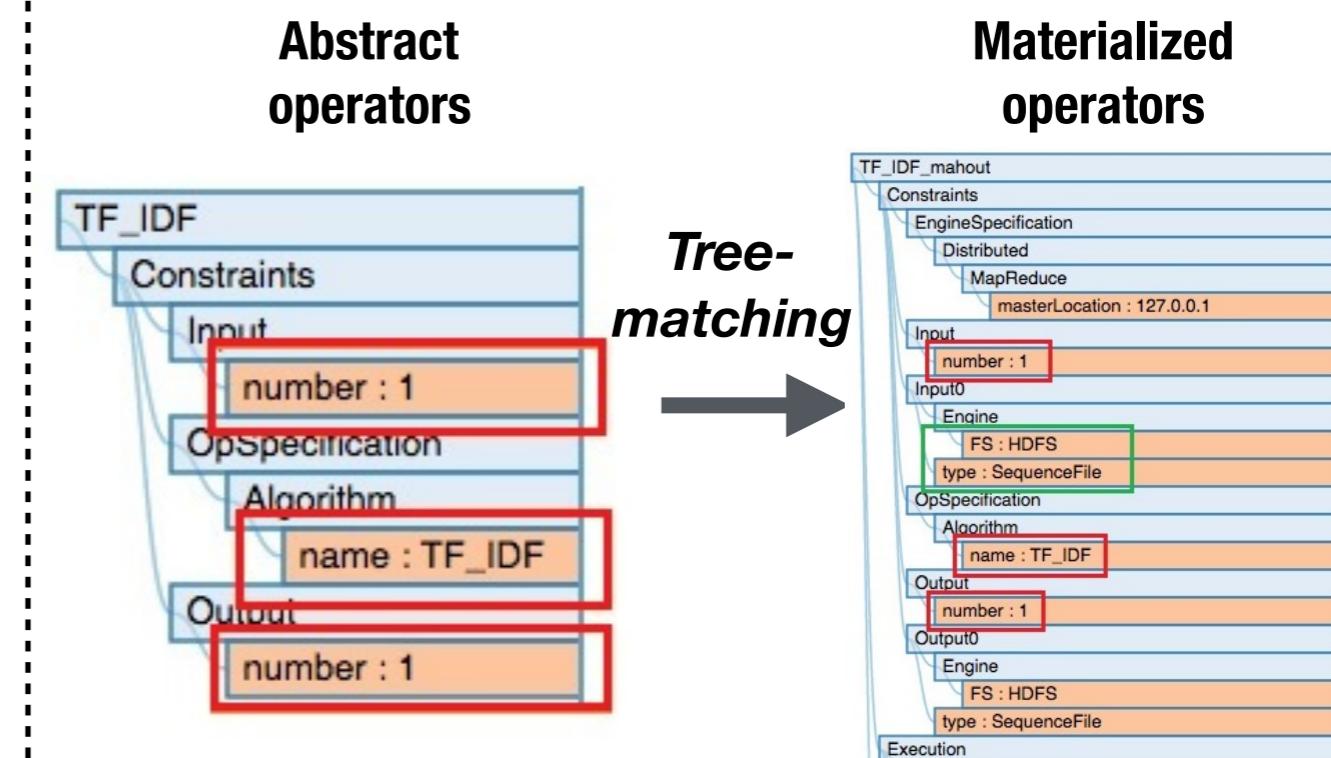
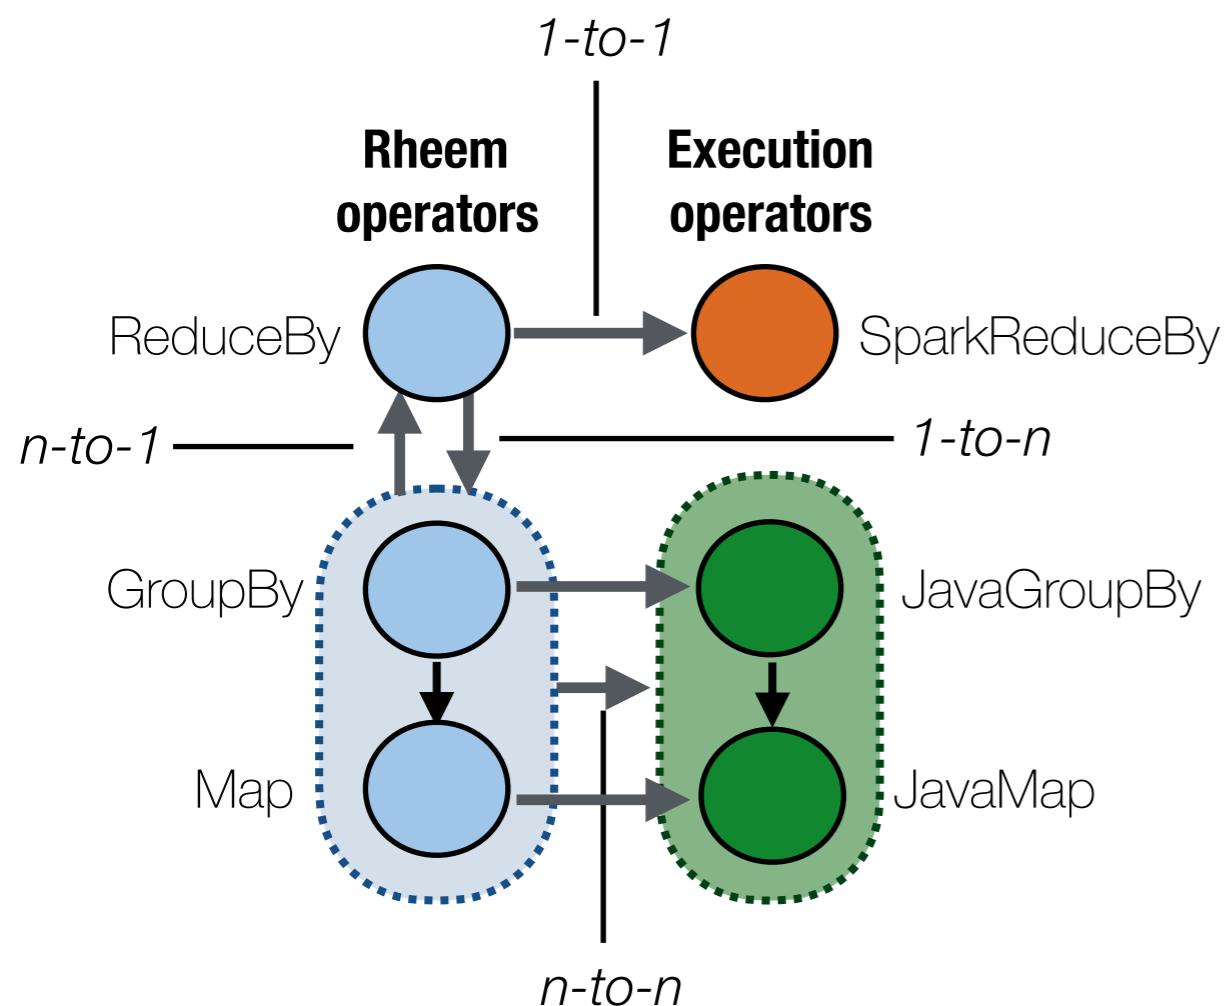
**Metadata-based** mappings



# Rheem & Ires

**Graph-based** mappings

**Metadata-based** mappings

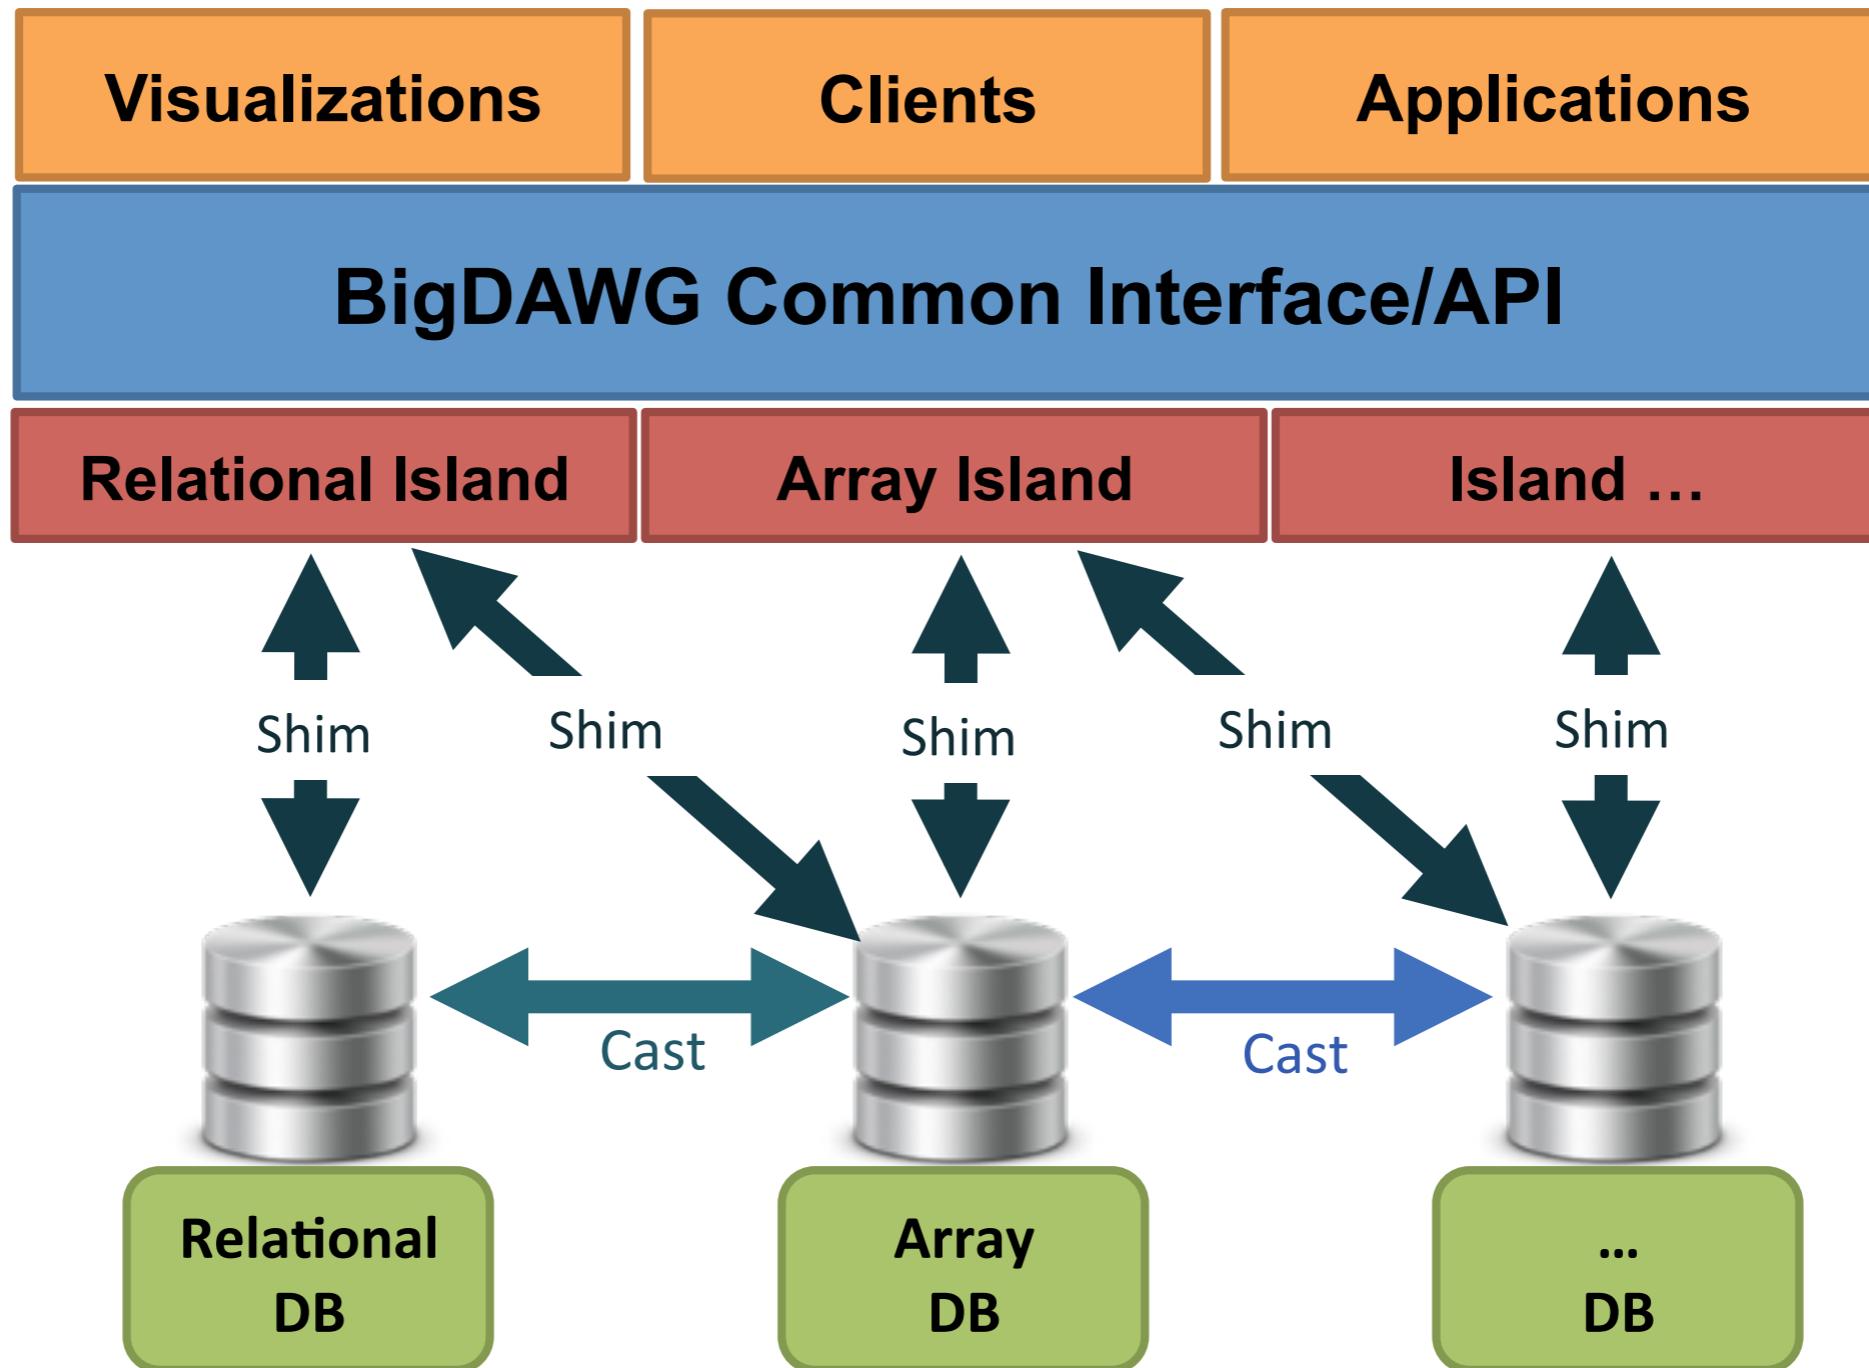


# Decoupling Features

	Processing Model	CP Granularity	Mappings
Musketeer	Dataflow DAG + Loop	Operator	Code templates
Rheem	Dataflow DAG + Loop	Operator	Graph-based
Ires	Workflow graph	Task	Metadata-based

# BigDawg

*Two-tier* decoupling



# Decoupling Features

	Processing Model	CP Granularity	Mappings
<b>Musketeer</b>	Dataflow DAG + Loop	Operator	Code templates
<b>Rheem</b>	Dataflow DAG + Loop	Operator	Graph-based
<b>Ires</b>	Workflow graph	Task	Metadata-based
<b>BigDawg</b>	Query Plan Tree	Subquery	Shims

# Decoupling Features

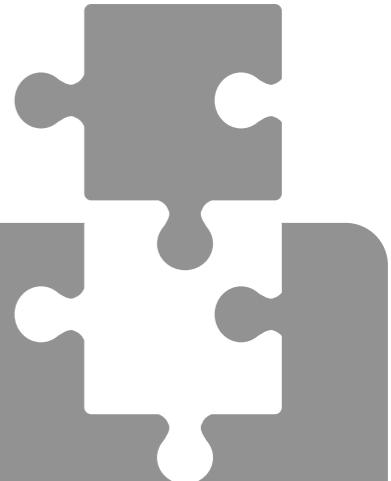
	Processing Model	CP Granularity	Mappings
<b>Musketeer</b>	Dataflow DAG + Loop	Operator	Code templates
<b>Rheem</b>	Dataflow DAG + Loop	Operator	Graph-based
<b>Ires</b>	Workflow graph	Task	Metadata-based
<b>BigDawg</b>	Query Plan Tree	Subquery	Shims
<b>Myria</b>	Extended relational	Operator	Rewrite rules
<b>QoX</b>	Dataflow graph	Task	?
<b>Cyclops</b>	Windowed aggregation	Operator	?
<b>Apache Beam</b>	Dataflow DAG	Pipeline	?
<b>CloudMdsQL</b>	Query Plan Tree	Subquery	

# Challenges

1 Decoupling Applications

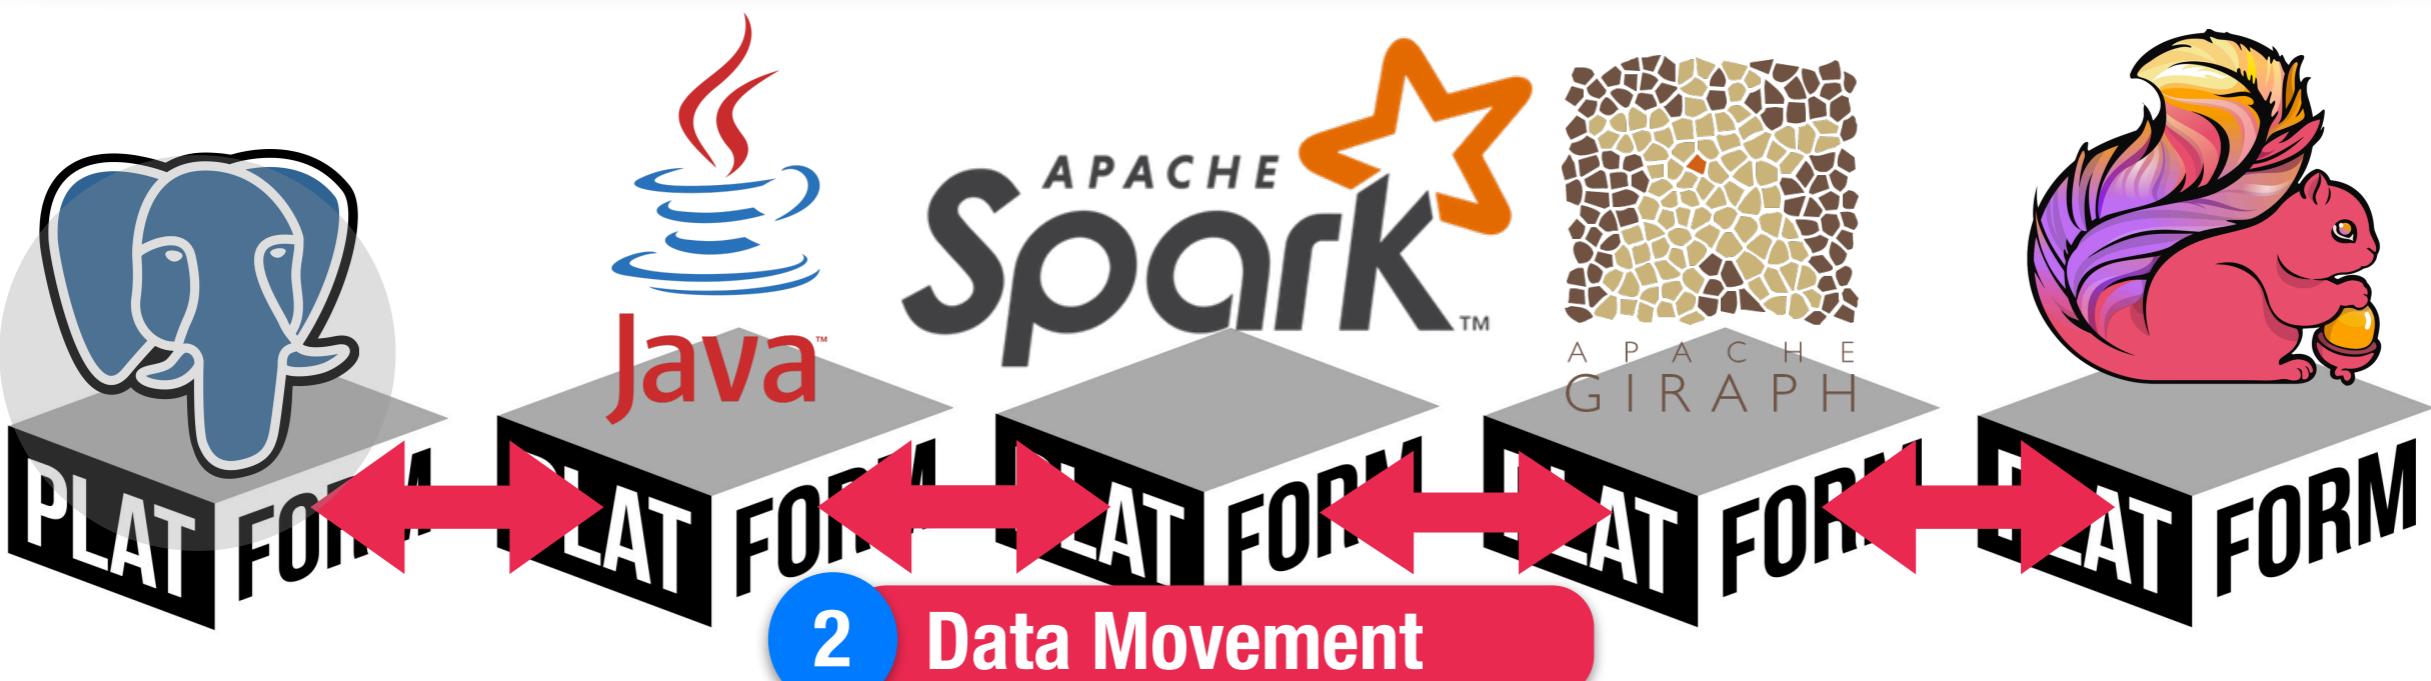


4 Extensibility



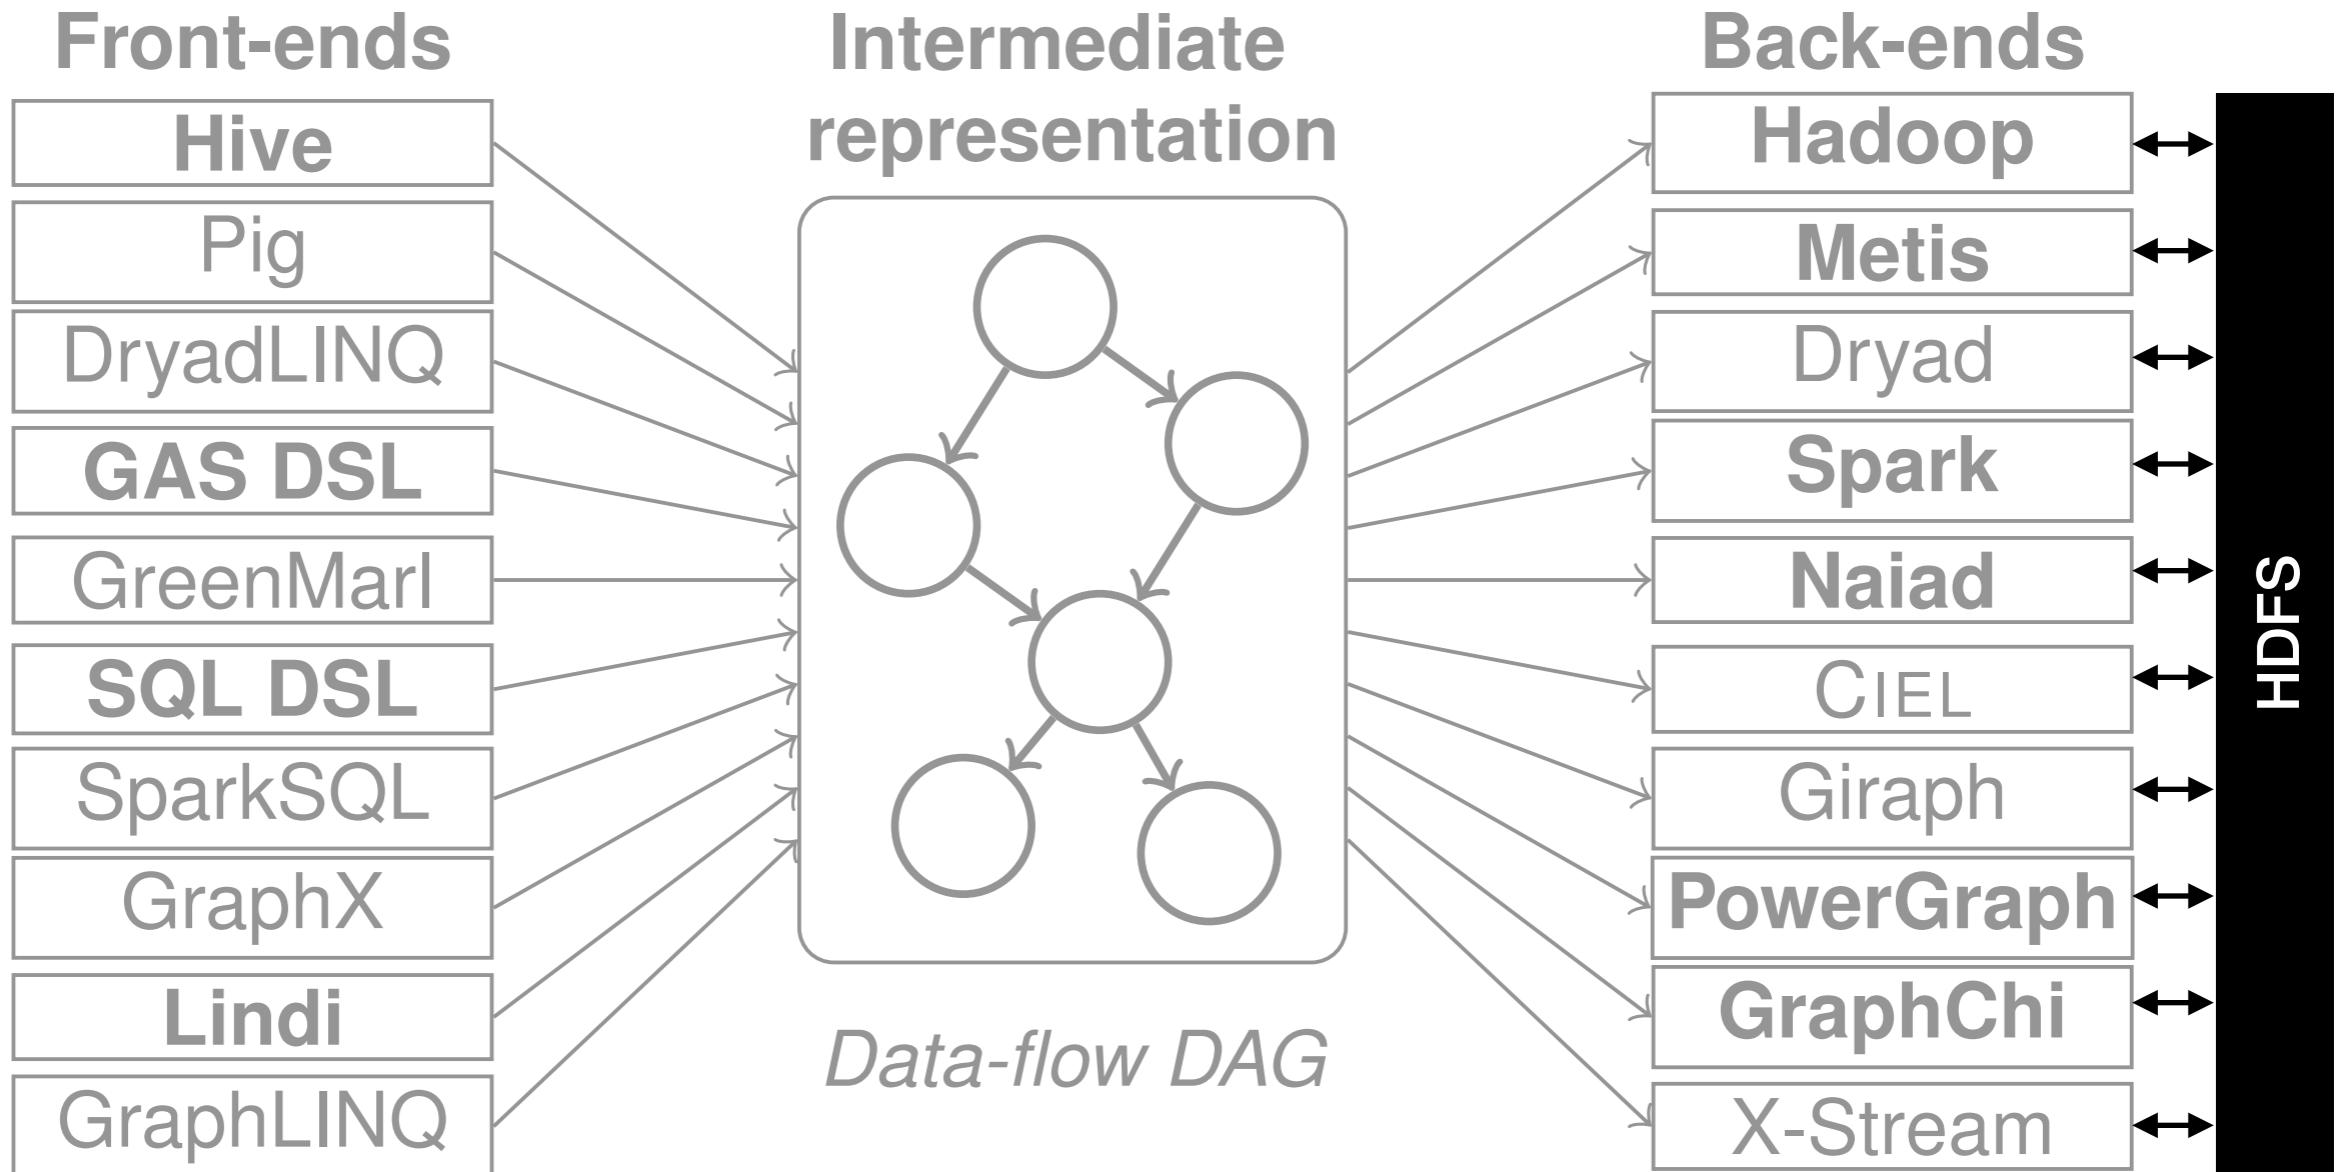
3

Automatic Cross-Platform  
Data Processing



# Musketeer

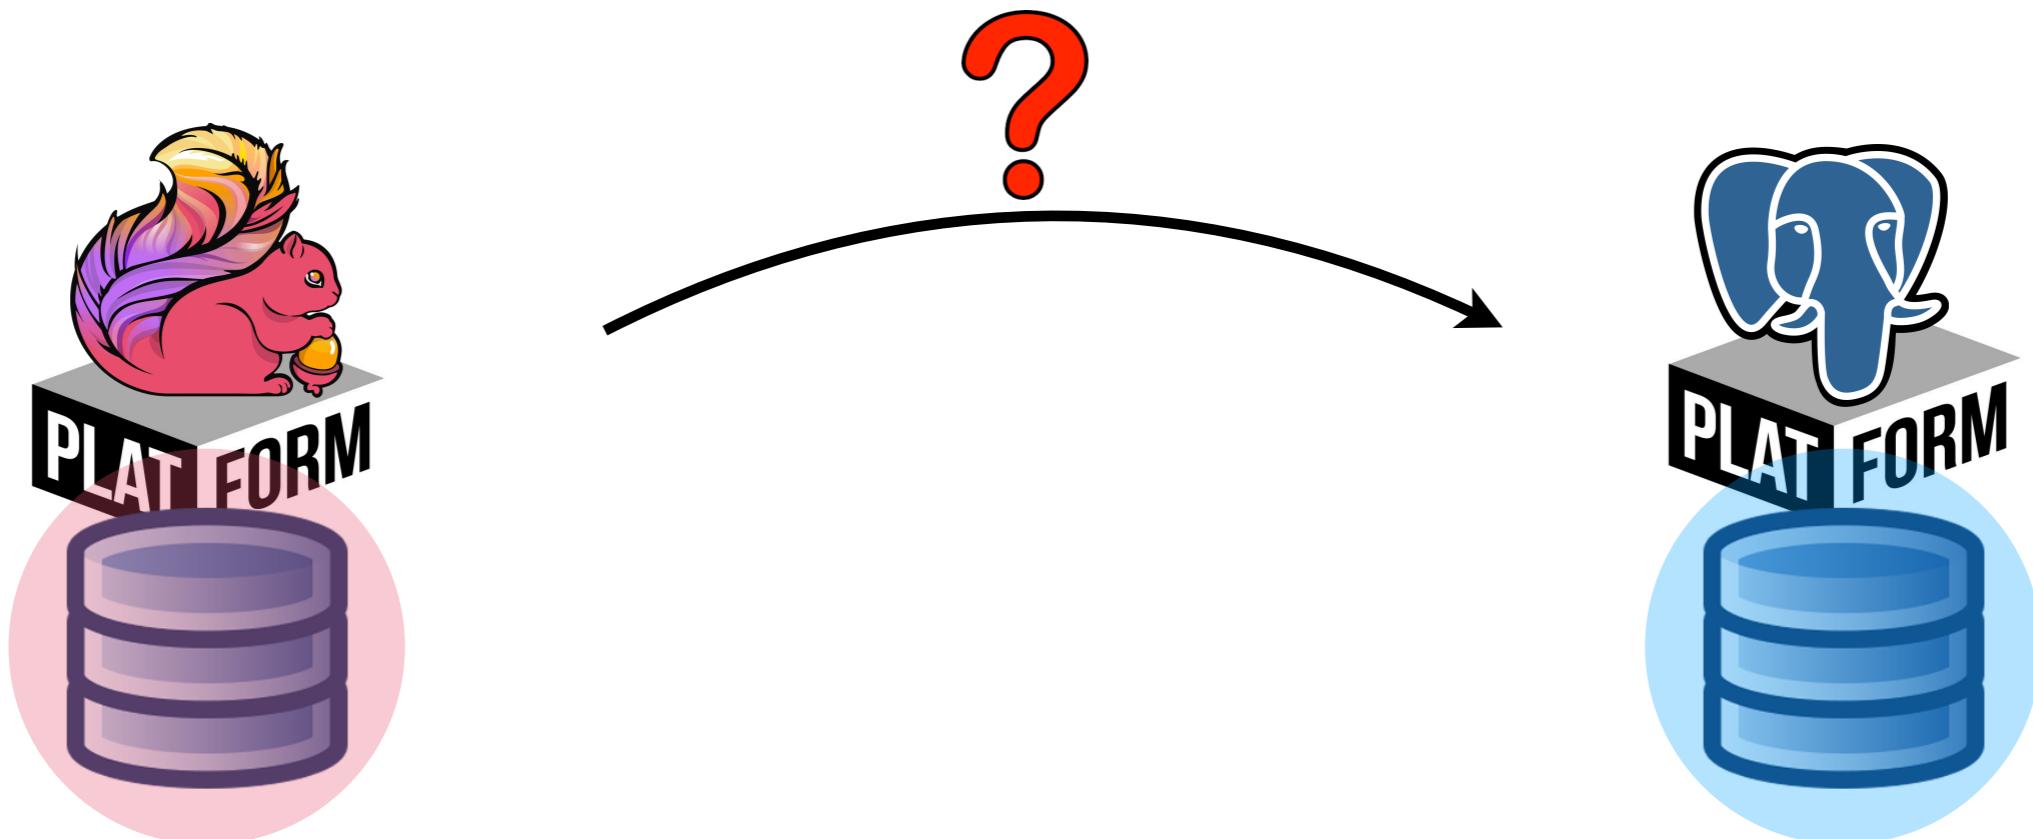
Communication *always* via **HDFS**



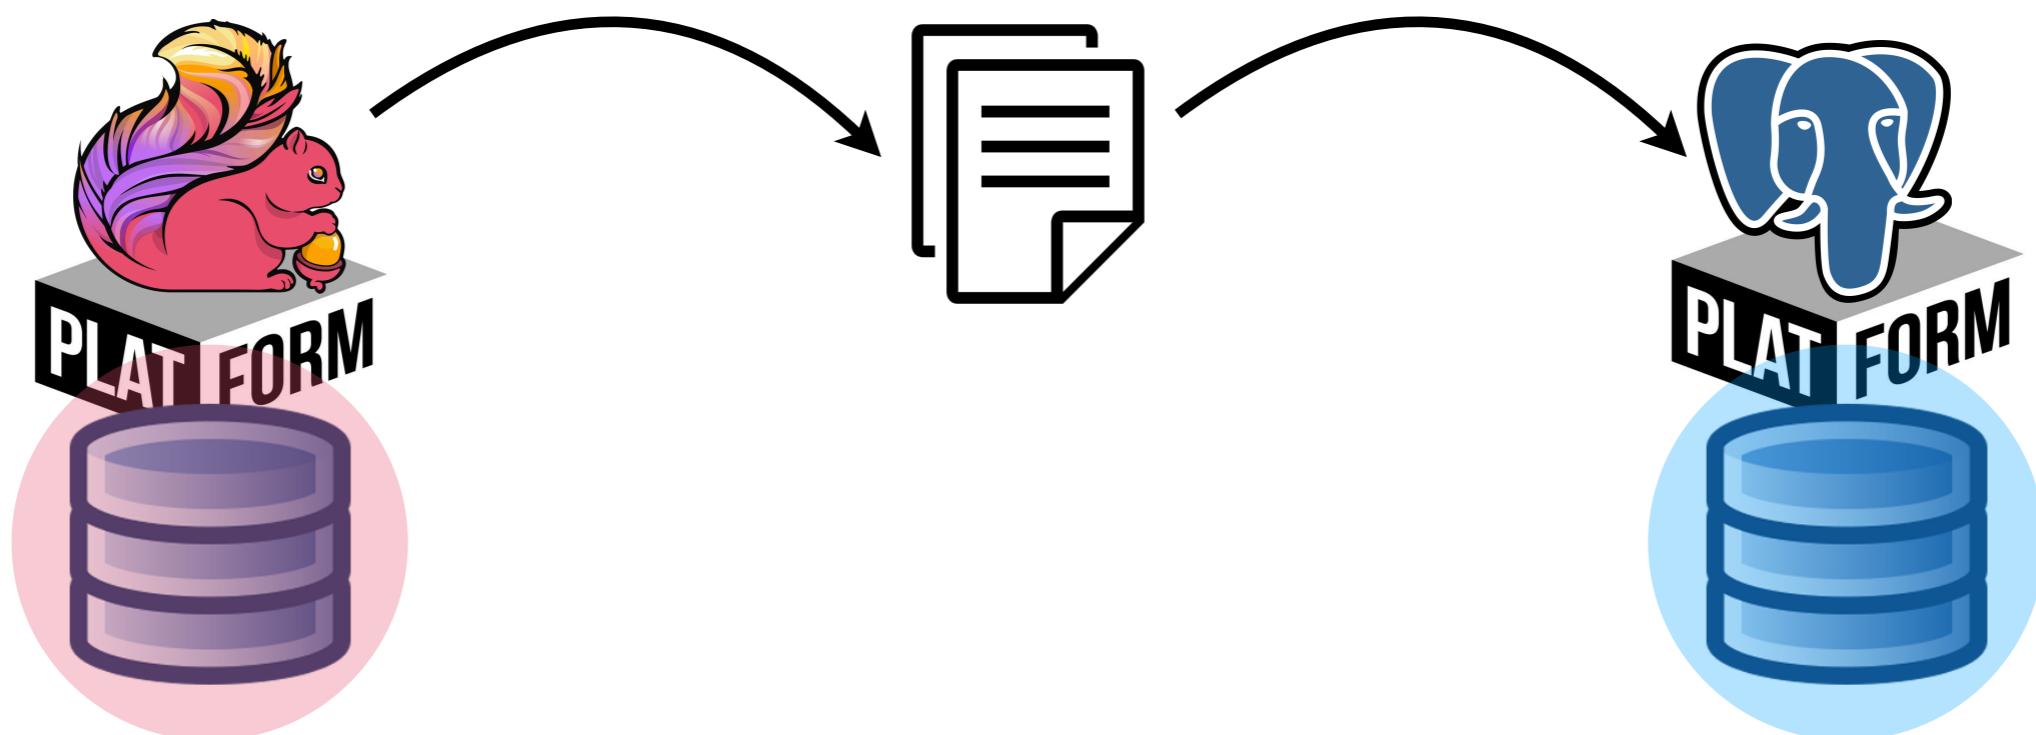
# Data Movement Features

	Connection	Transfer Means
Musketeer	Direct	HDFS Files

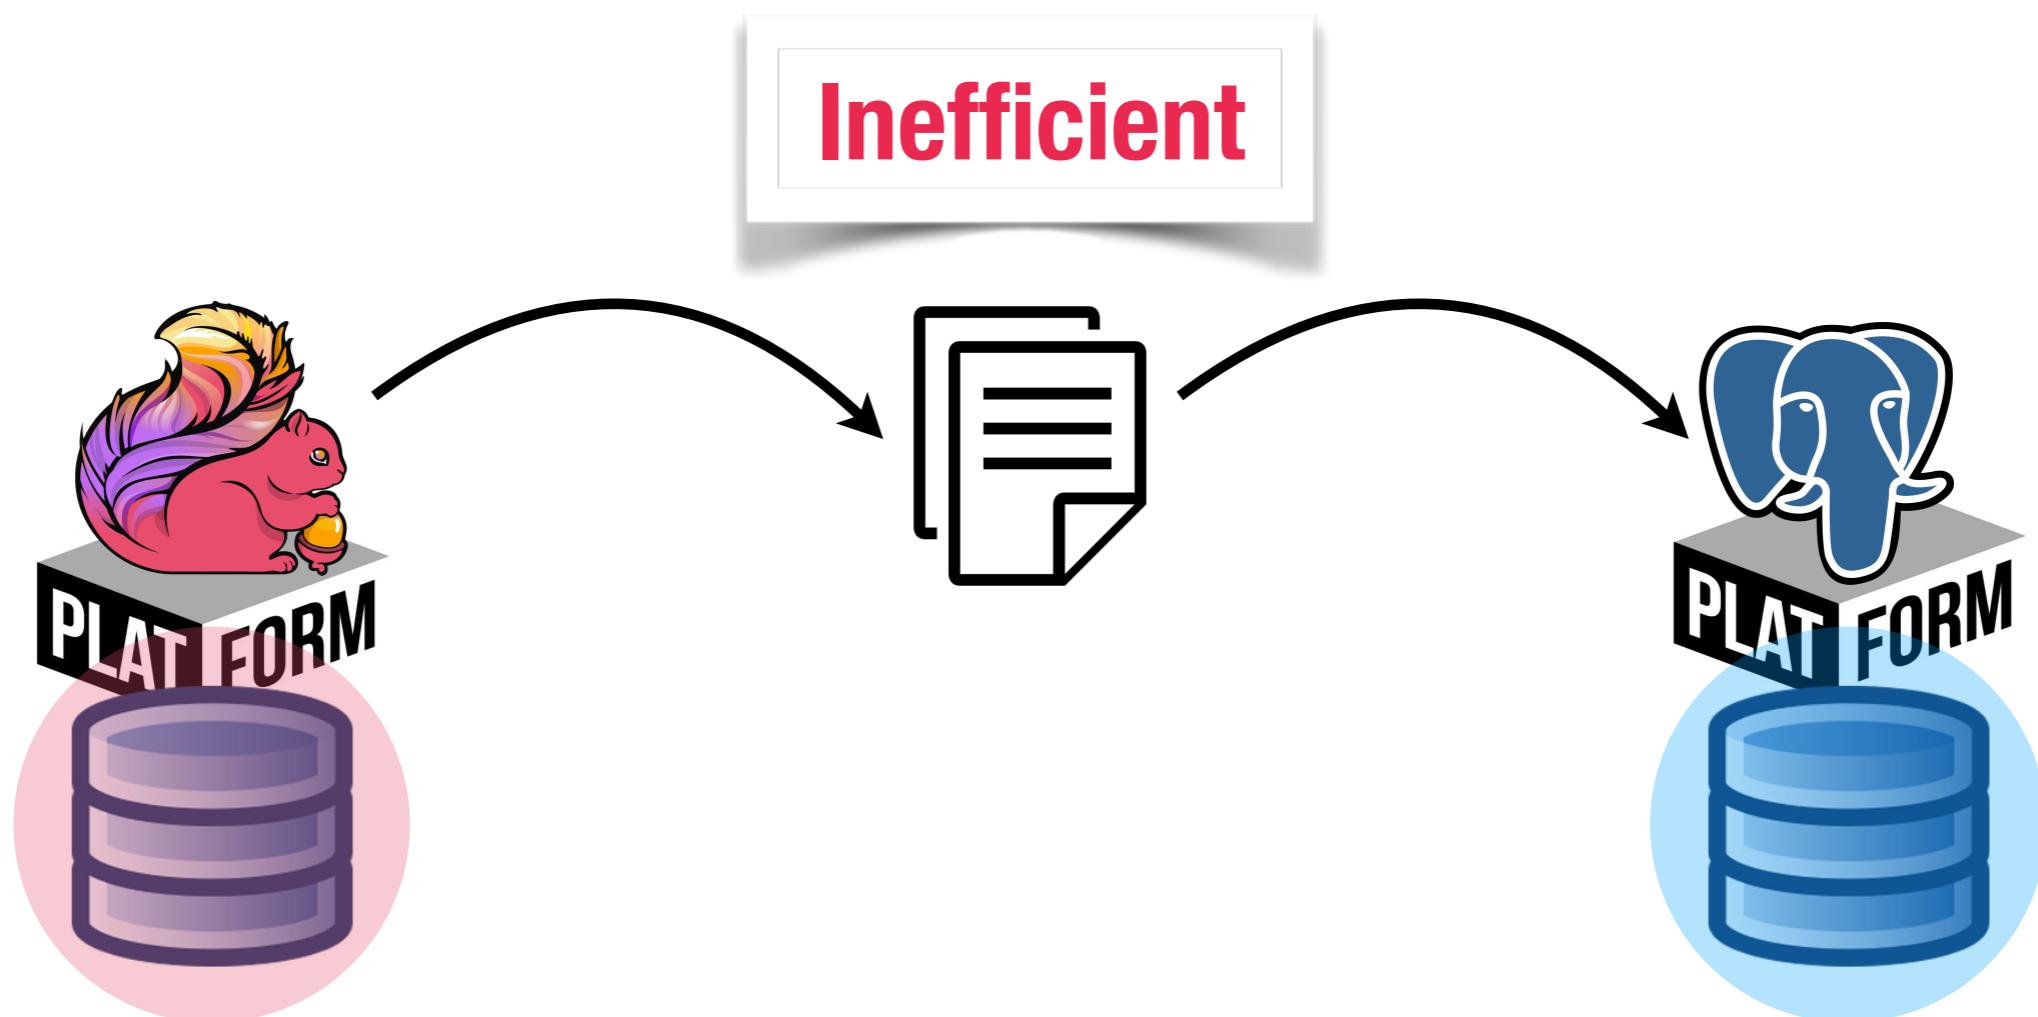
# Connecting two platforms



# Connecting two platforms

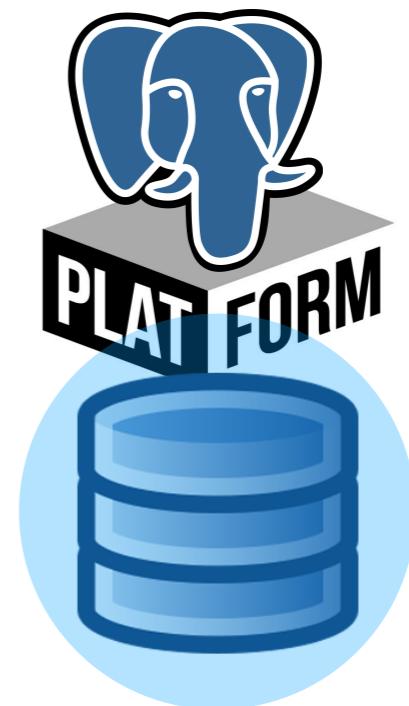
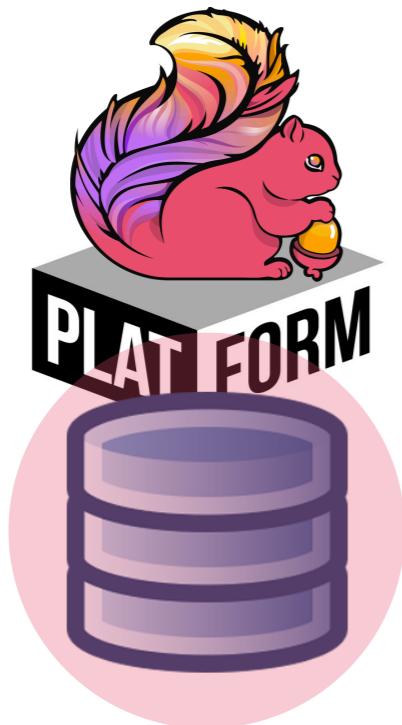


# Connecting two platforms



# PipeGen by Myria

***Binary*** data transfer via generated ***data pipes***



# PipeGen by Myria

***Binary*** data transfer via generated ***data pipes***

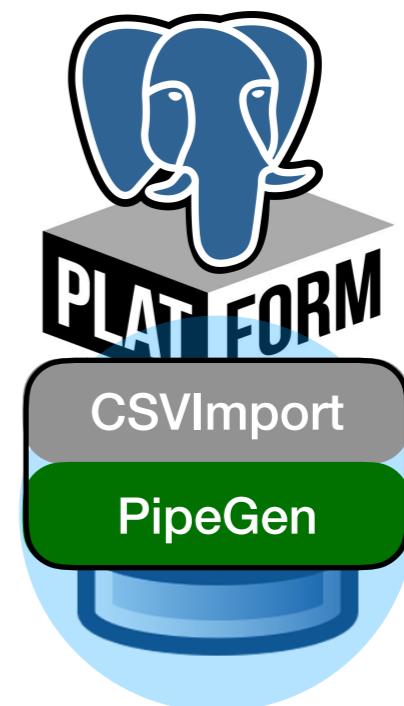
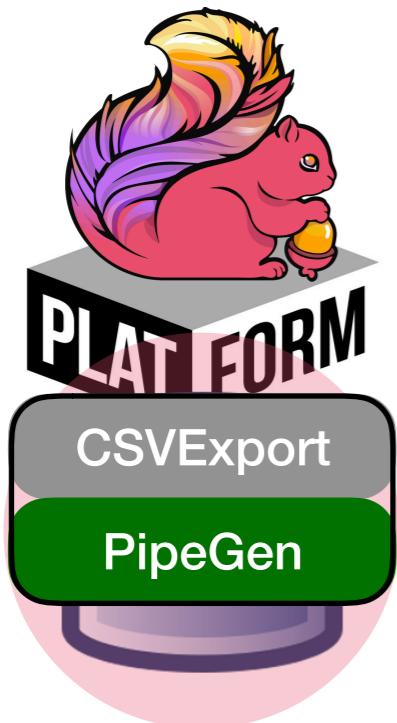
**Same Import/Export Capabilities**



# PipeGen by Myria

***Binary*** data transfer via generated ***data pipes***

**Same Import/Export Capabilities**

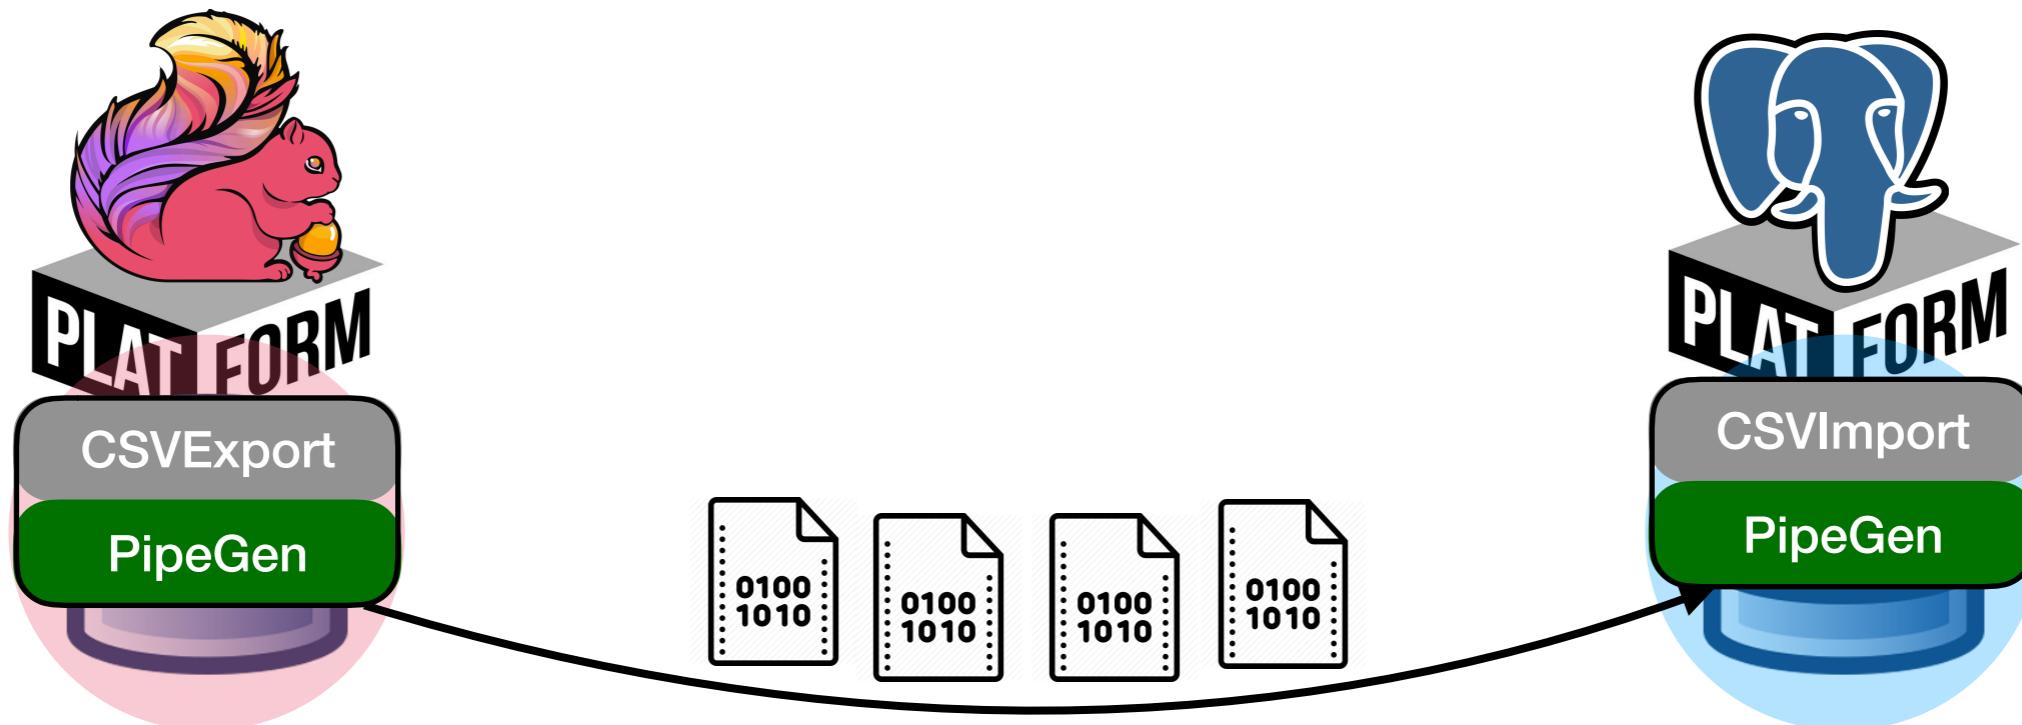


**Program Analysis**

# PipeGen by Myria

***Binary*** data transfer via generated ***data pipes***

**Same Import/Export Capabilities**

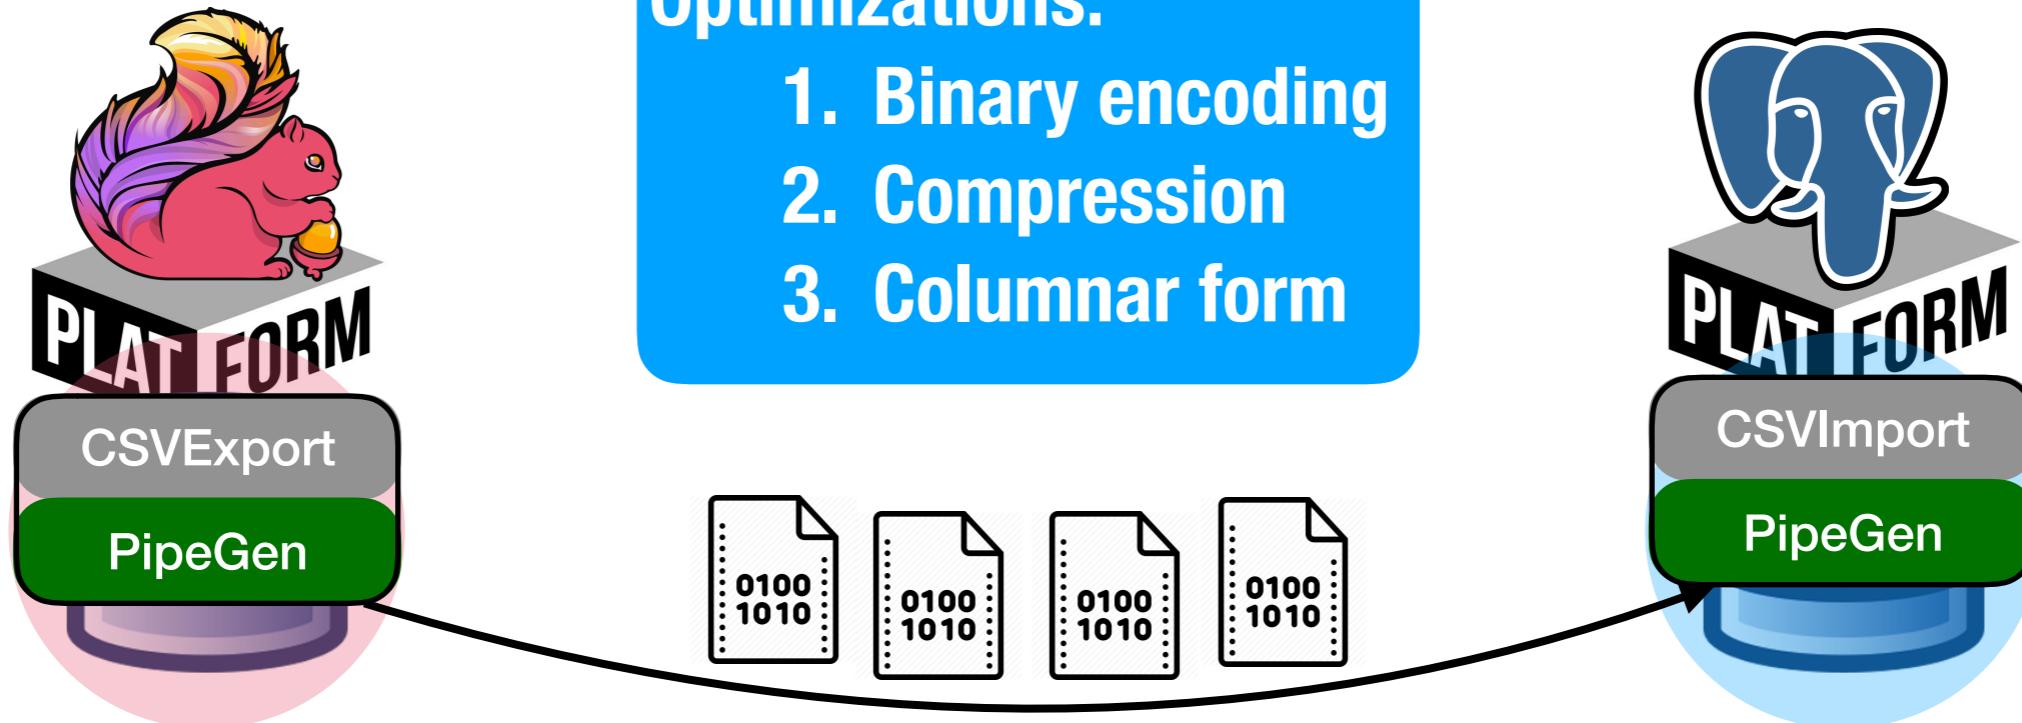


**Program Analysis**

# PipeGen by Myria

***Binary*** data transfer via generated ***data pipes***

## Same Import/Export Capabilities

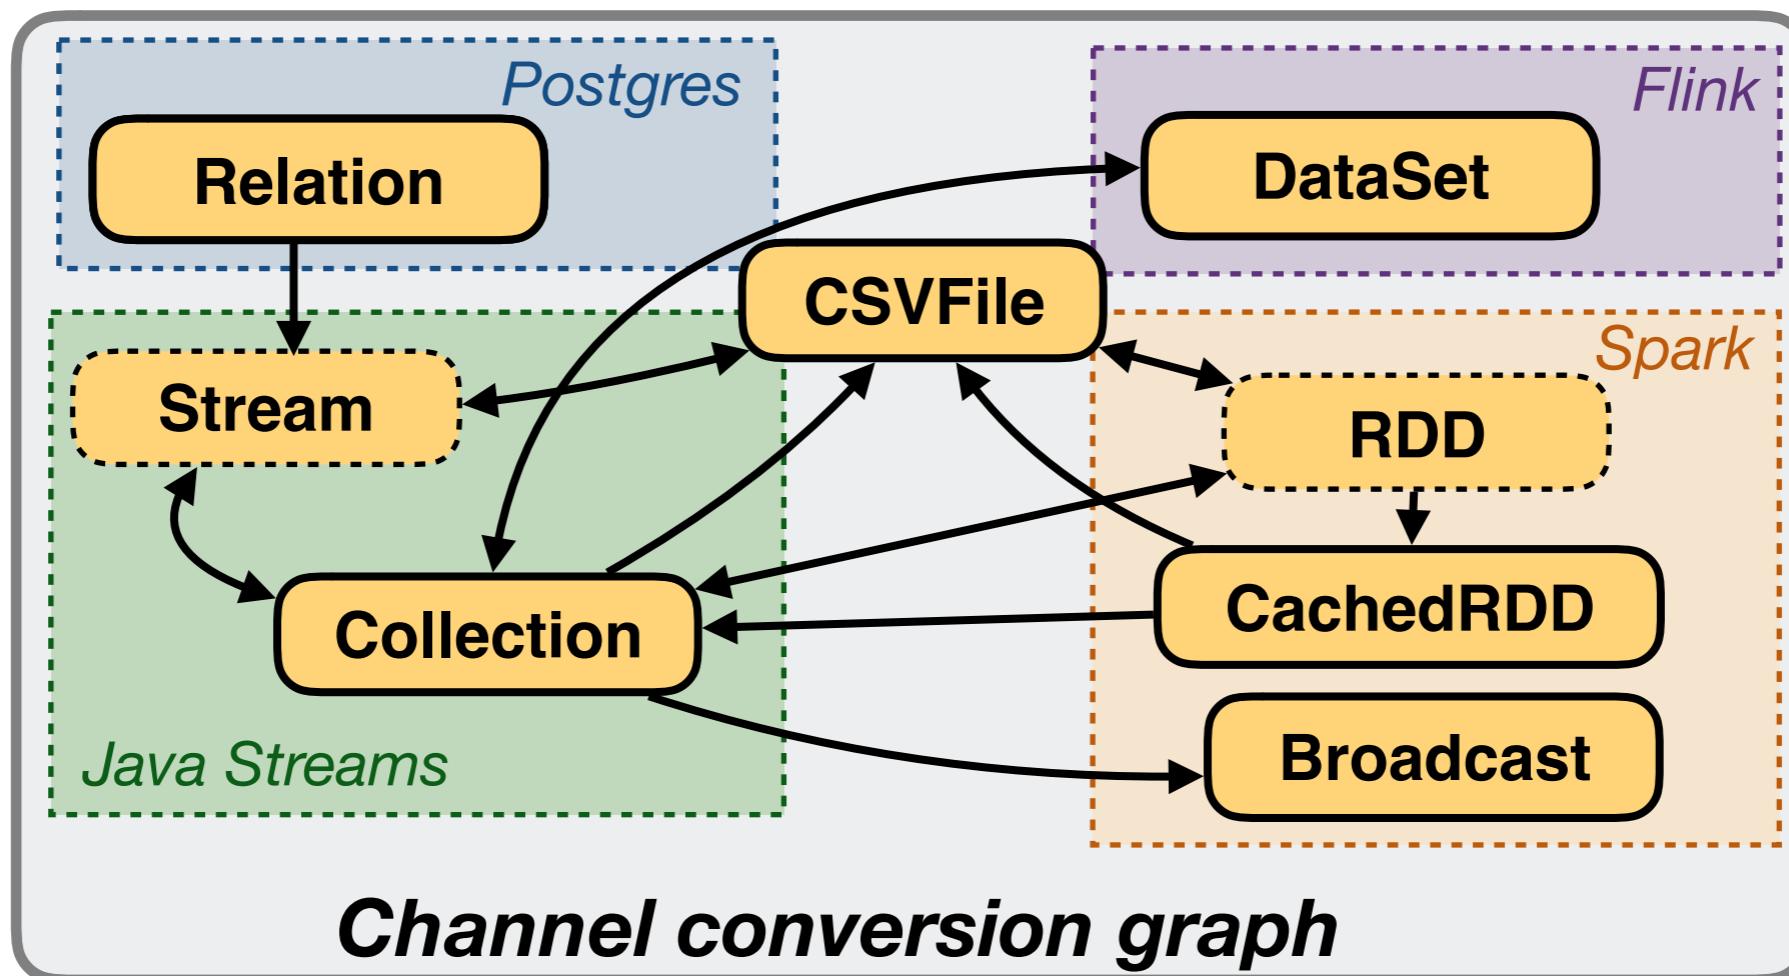


# Data Movement Features

	Connection	Transfer Means
Musketeer	Direct	HDFS Files
Myria	Direct	PipeGen

# Channels Graph by Rheem

**Graph-based** data transfer representation

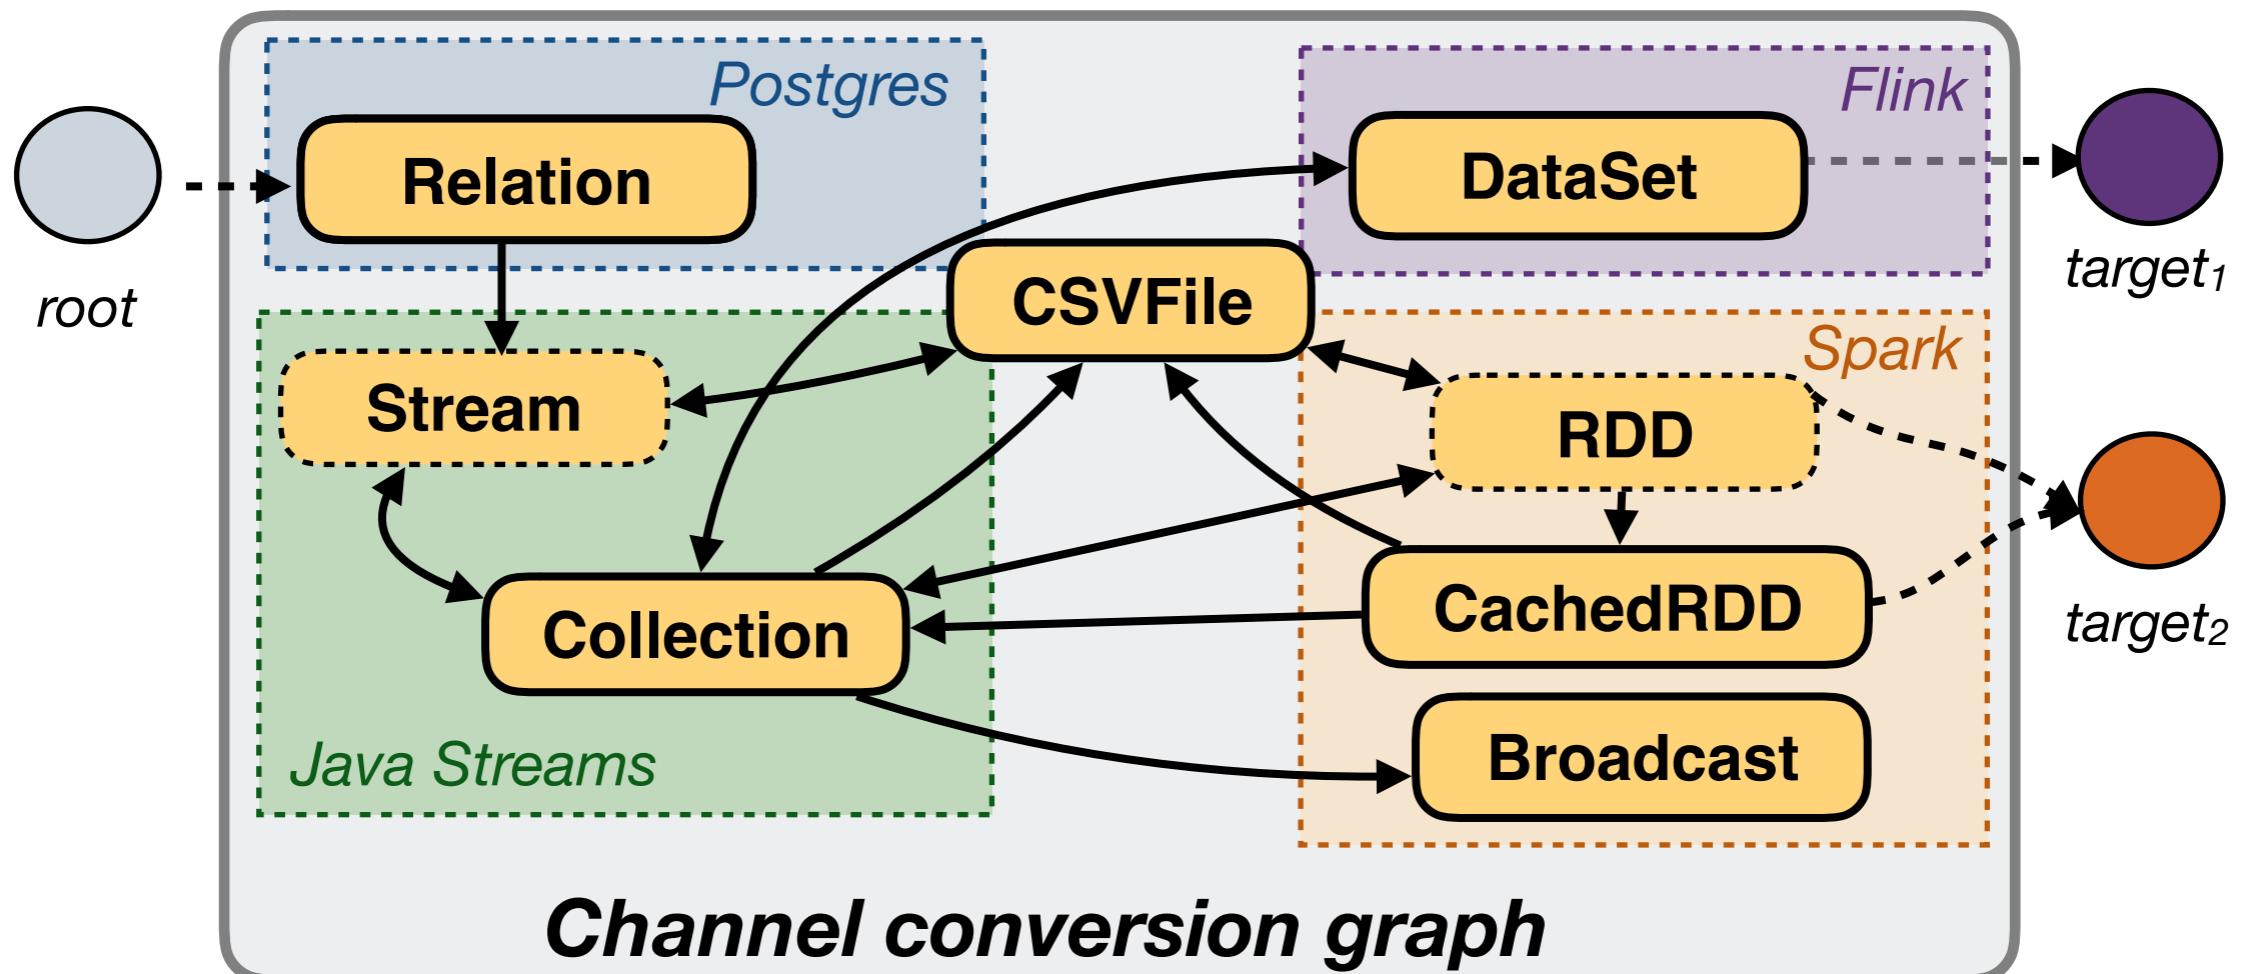


● reusable channel

● non-reusable channel

# Channels Graph by Rheem

**Graph-based** data transfer representation

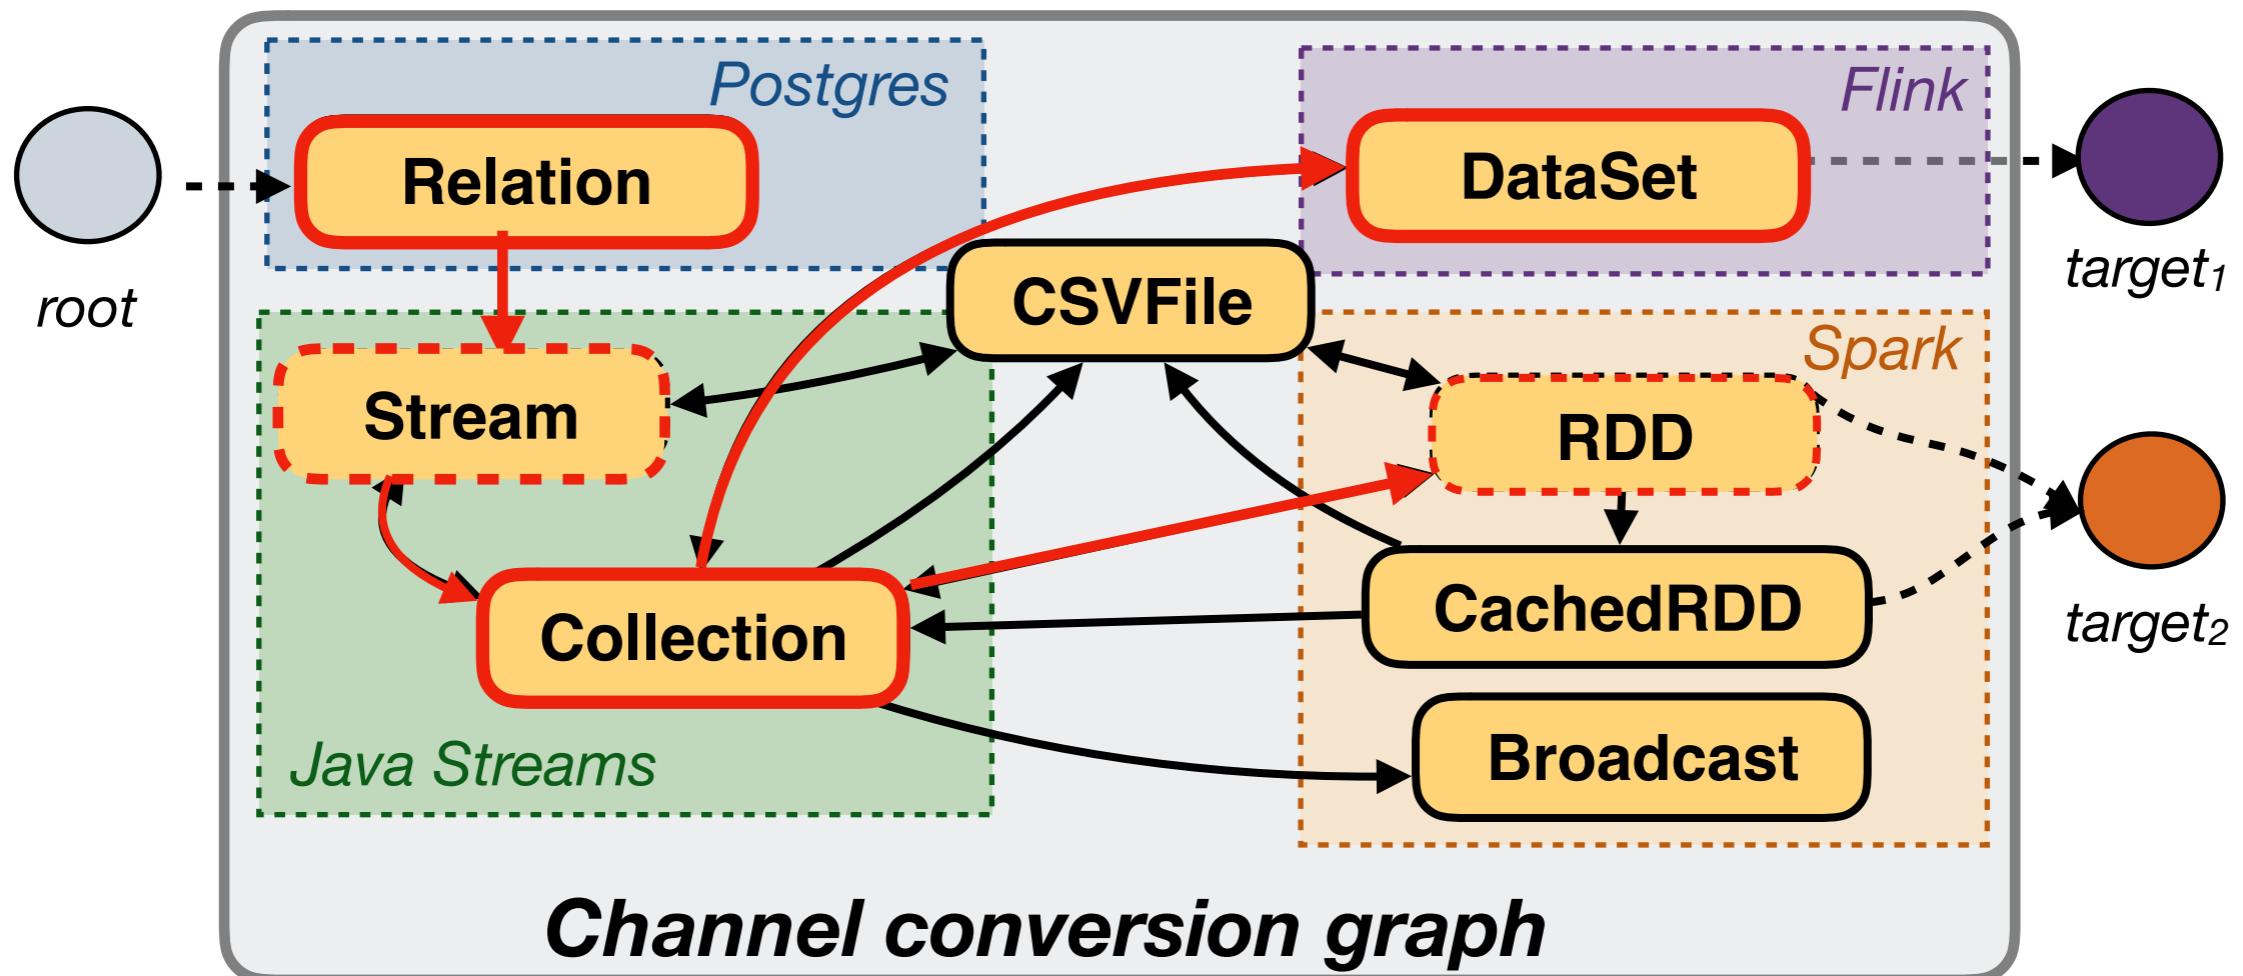


● reusable channel

● non-reusable channel

# Channels Graph by Rheem

**Graph-based** data transfer representation

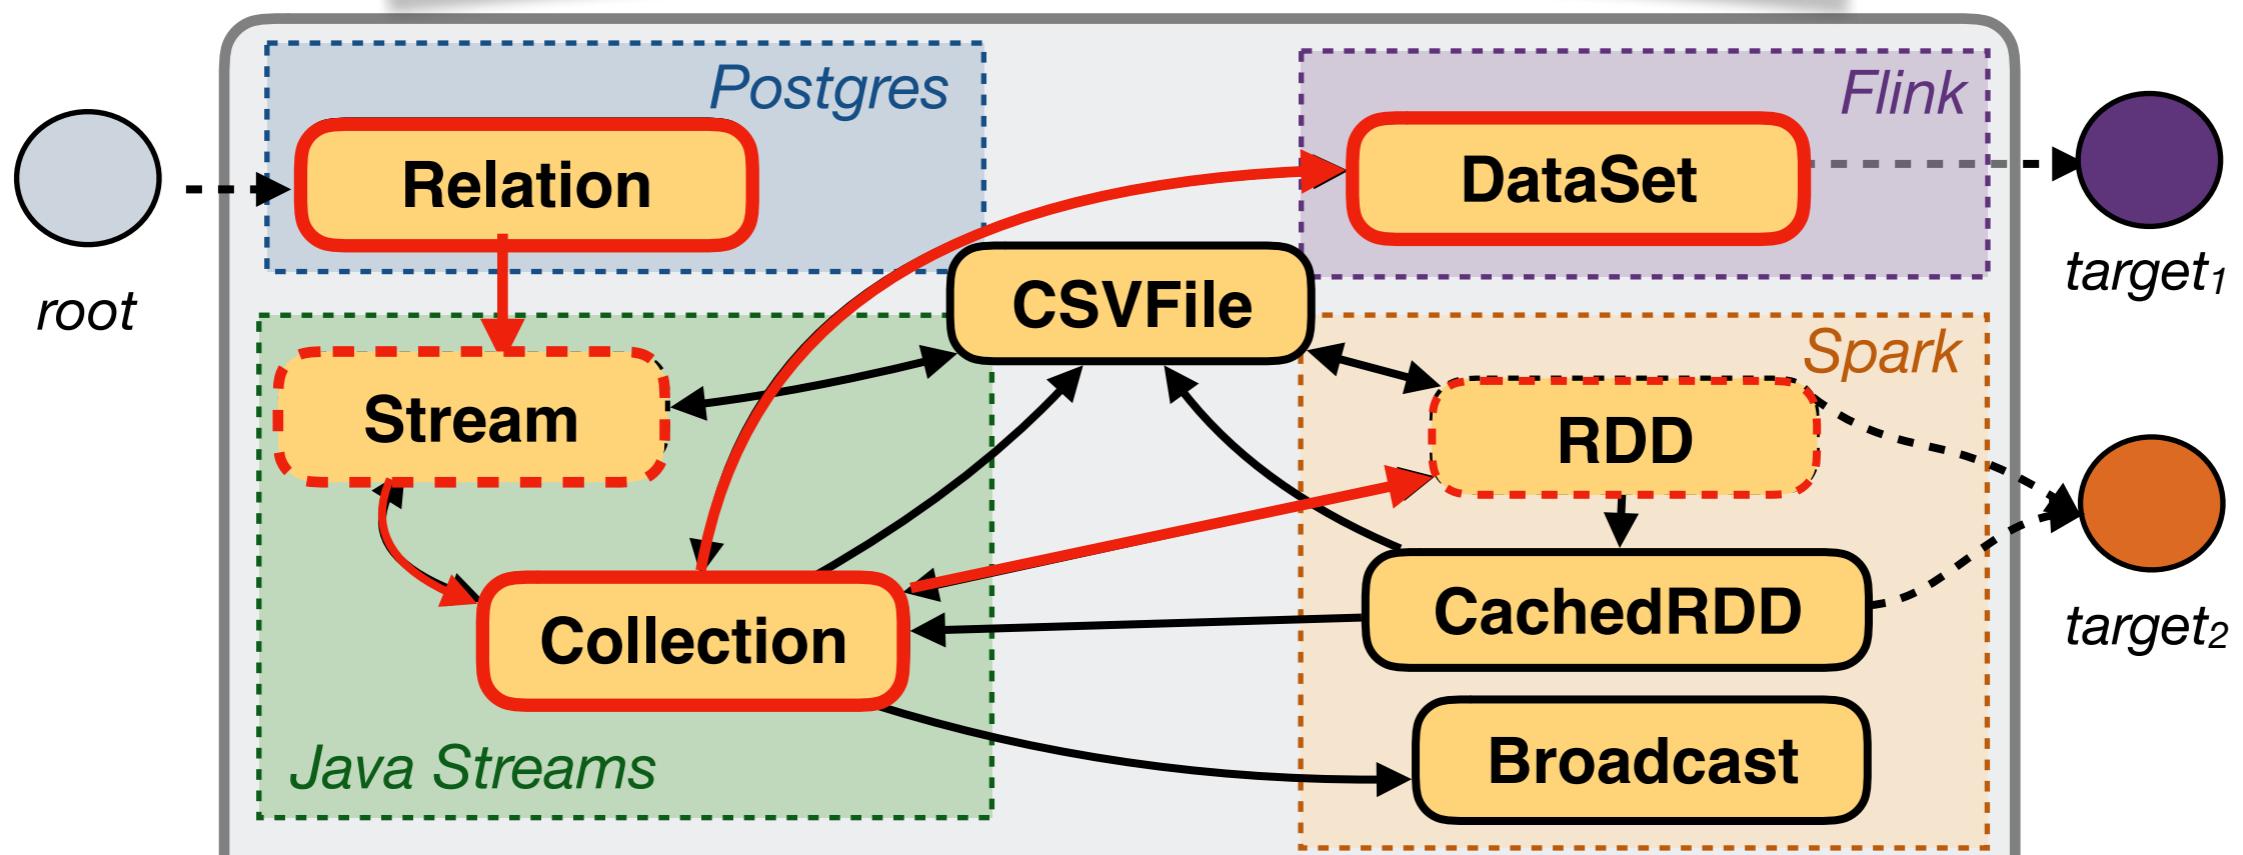


- reusable channel
- non-reusable channel

# Channels Graph by Rheem

**Graph-based** data transfer representation

## Group Steiner Tree Problem



- reusable channel
- non-reusable channel

# Data Movement Features

	Connection	Transfer Means
Musketeer	Direct	HDFS Files
Myria	Direct	PipeGen
Rheem	Graph	Any

# Data Movement Features

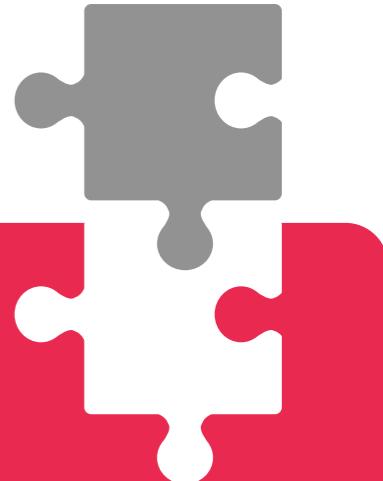
	Connection	Transfer Means
Musketeer	Direct	HDFS Files
Myria	Direct	PipeGen
Rheem	Graph	Any
Ires	Direct	Any
QoX	Graph	?
DBMS+	Direct	HDFS Files

# Challenges

1 Decoupling Applications



4 Extensibility



3

Automatic Cross-Platform  
Data Processing



# Decoupling Applications

Applications  
/Frontends

Hive

Crunch

MLlib

Mahout

Pig

## Cross-Platform System

Processing  
Platforms

Hadoop  
MR

DBMS

Storm

Spark

Flink

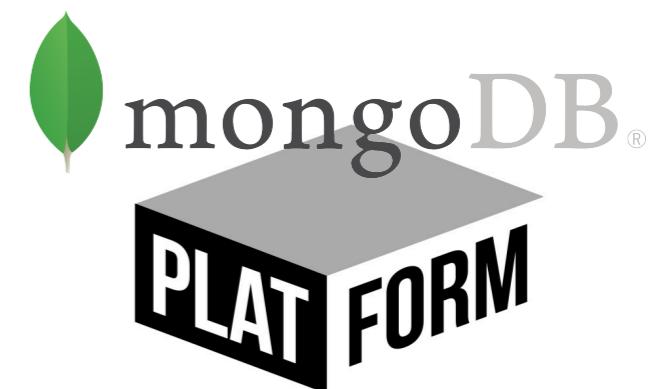
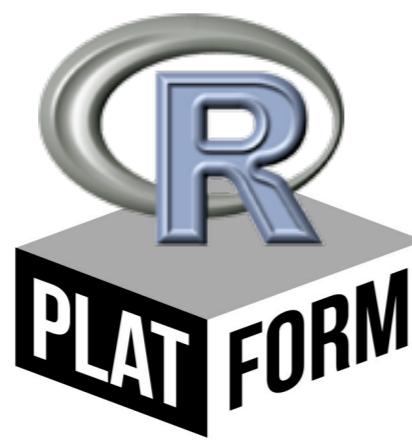
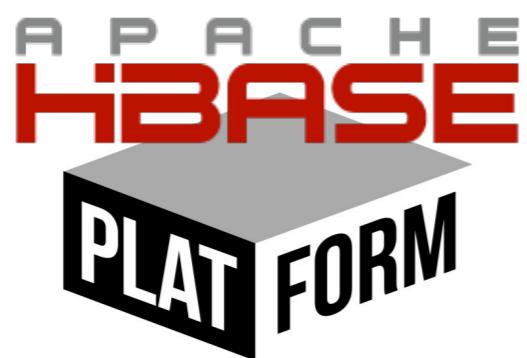
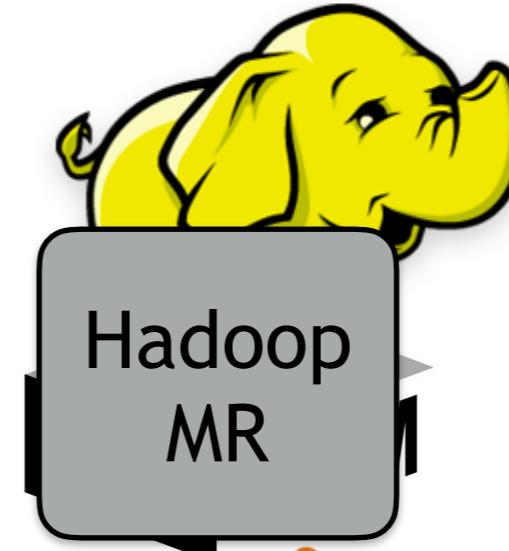
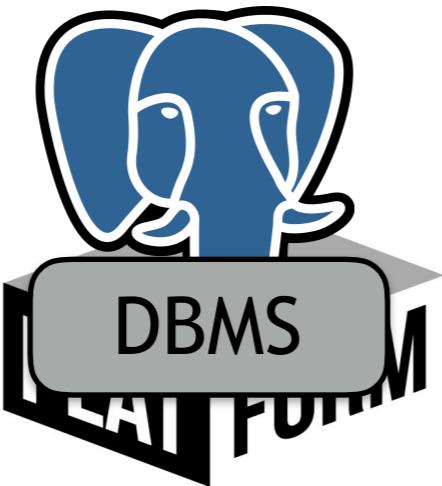
Storage  
Engines

HDFS

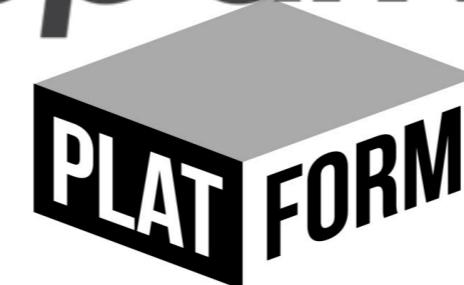
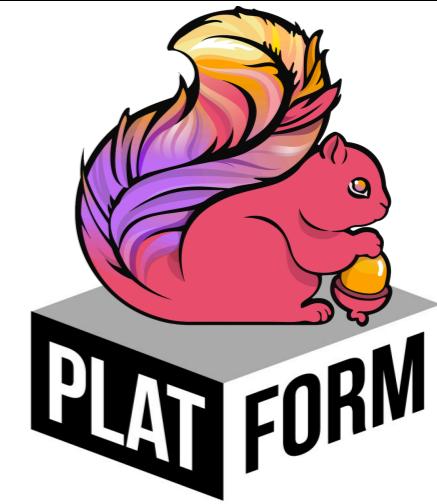
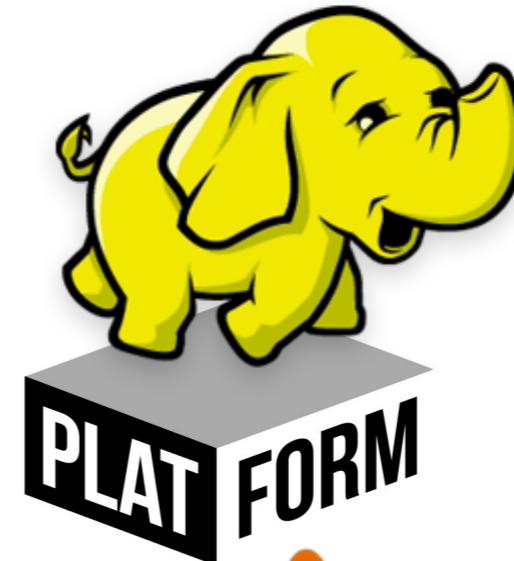
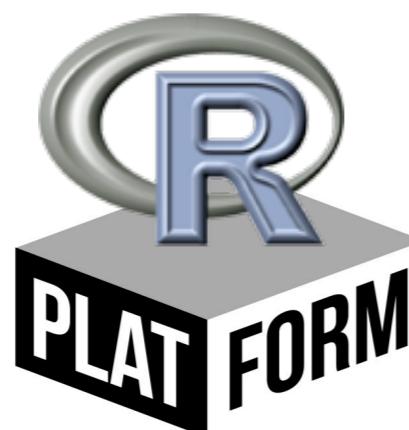
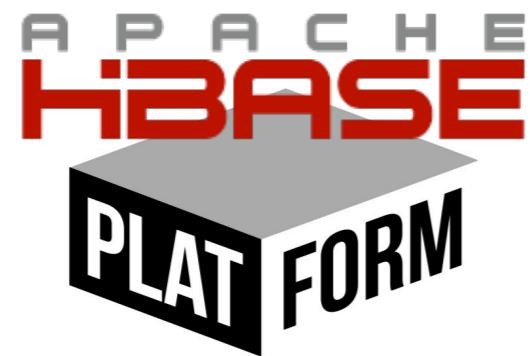
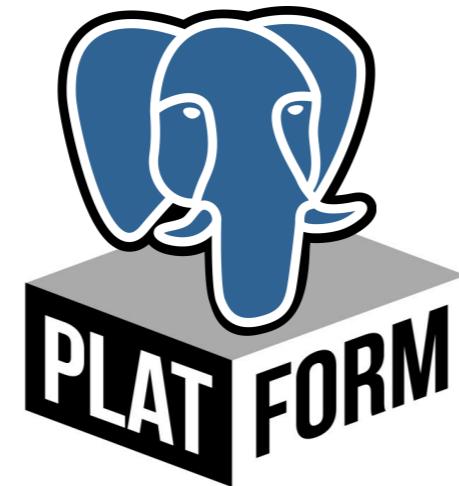
S3

Local FS

# Which Platforms to Use?



# Which Platforms to Use?



# Cross-Platform Optimization



**Cost-based**



**Rule-based**

# Cross-Platform Optimization

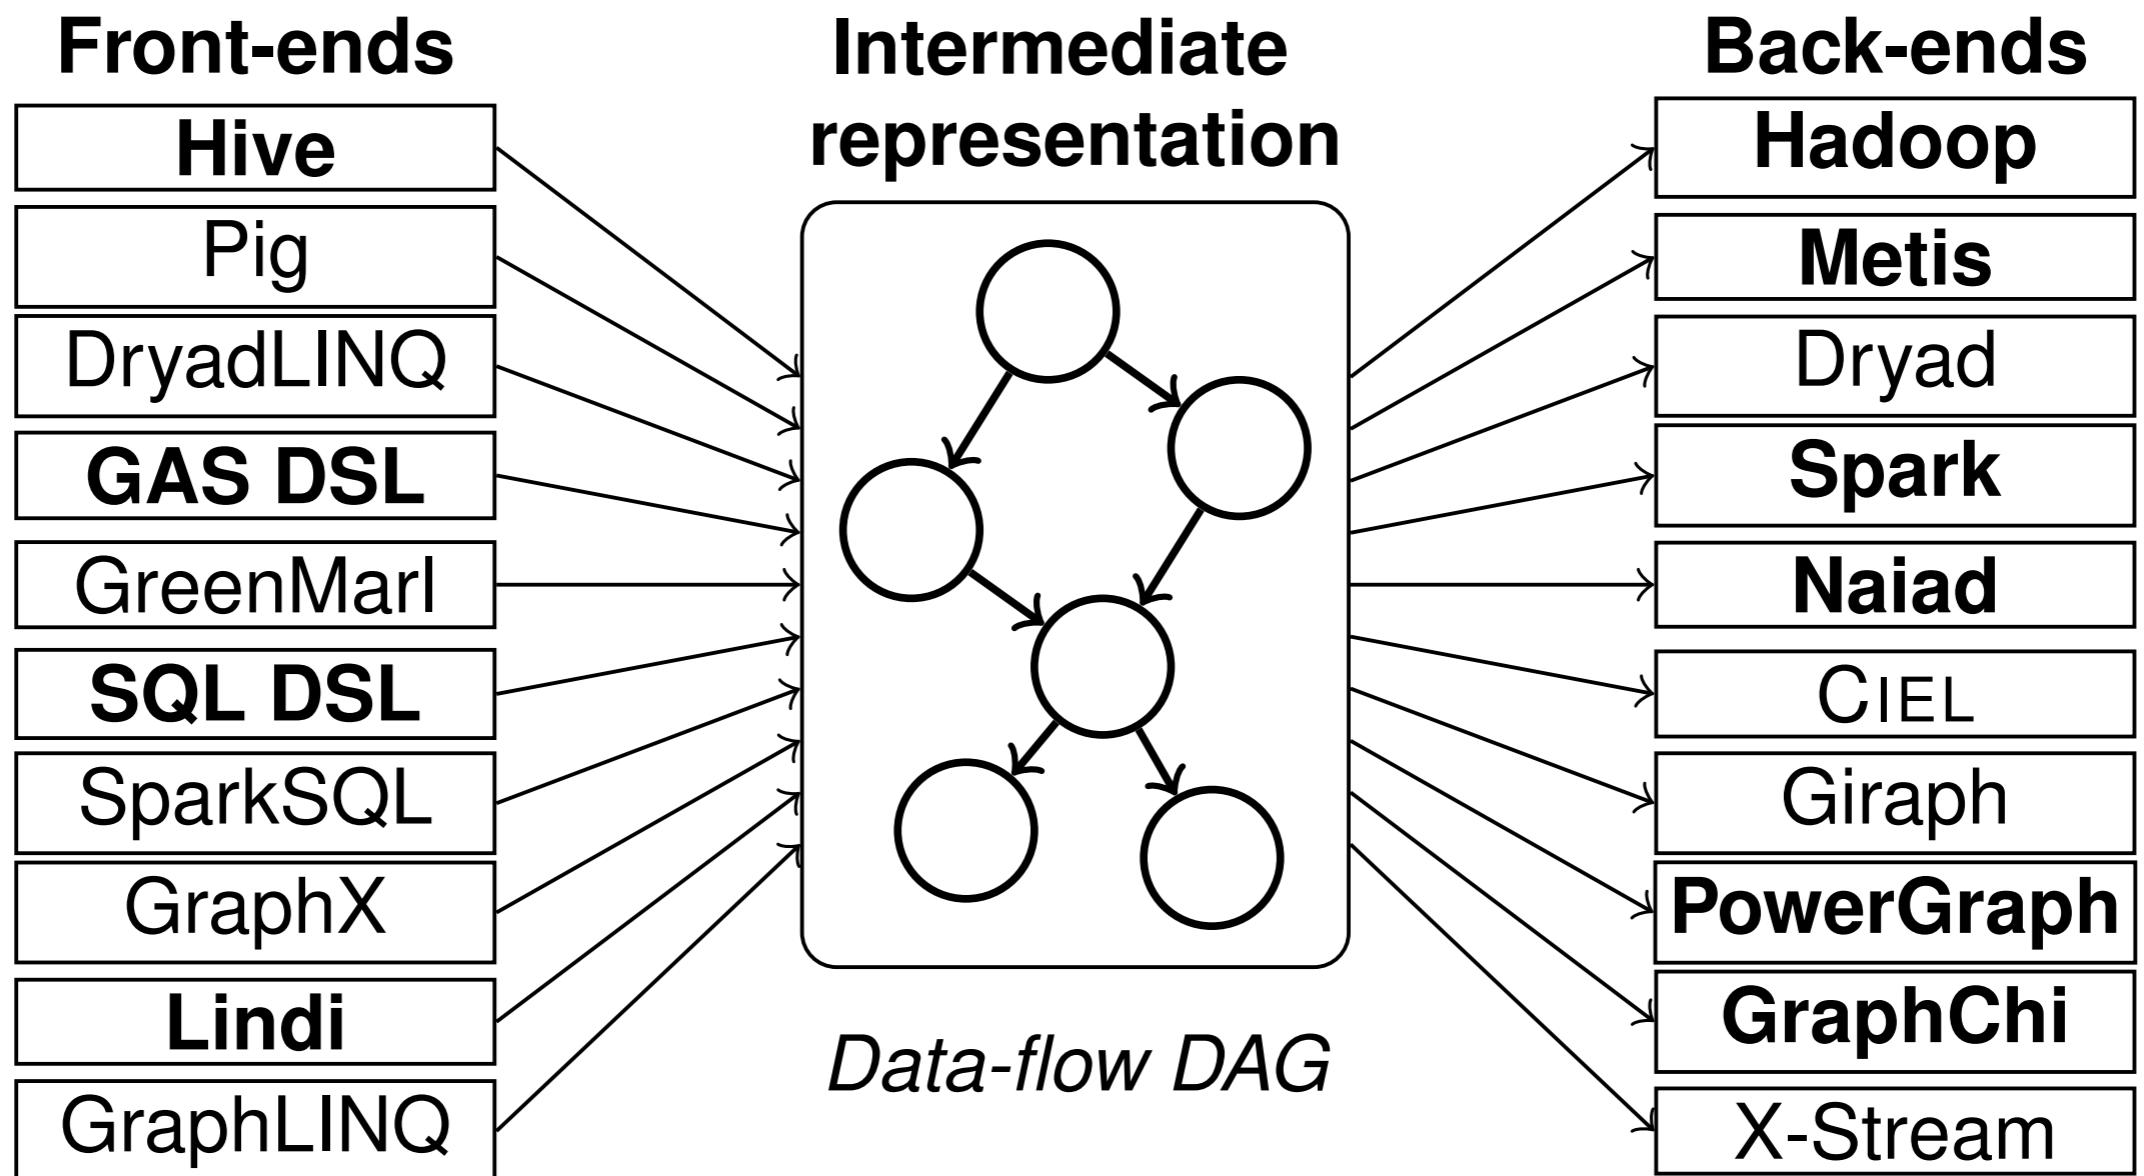


**Cost-based**

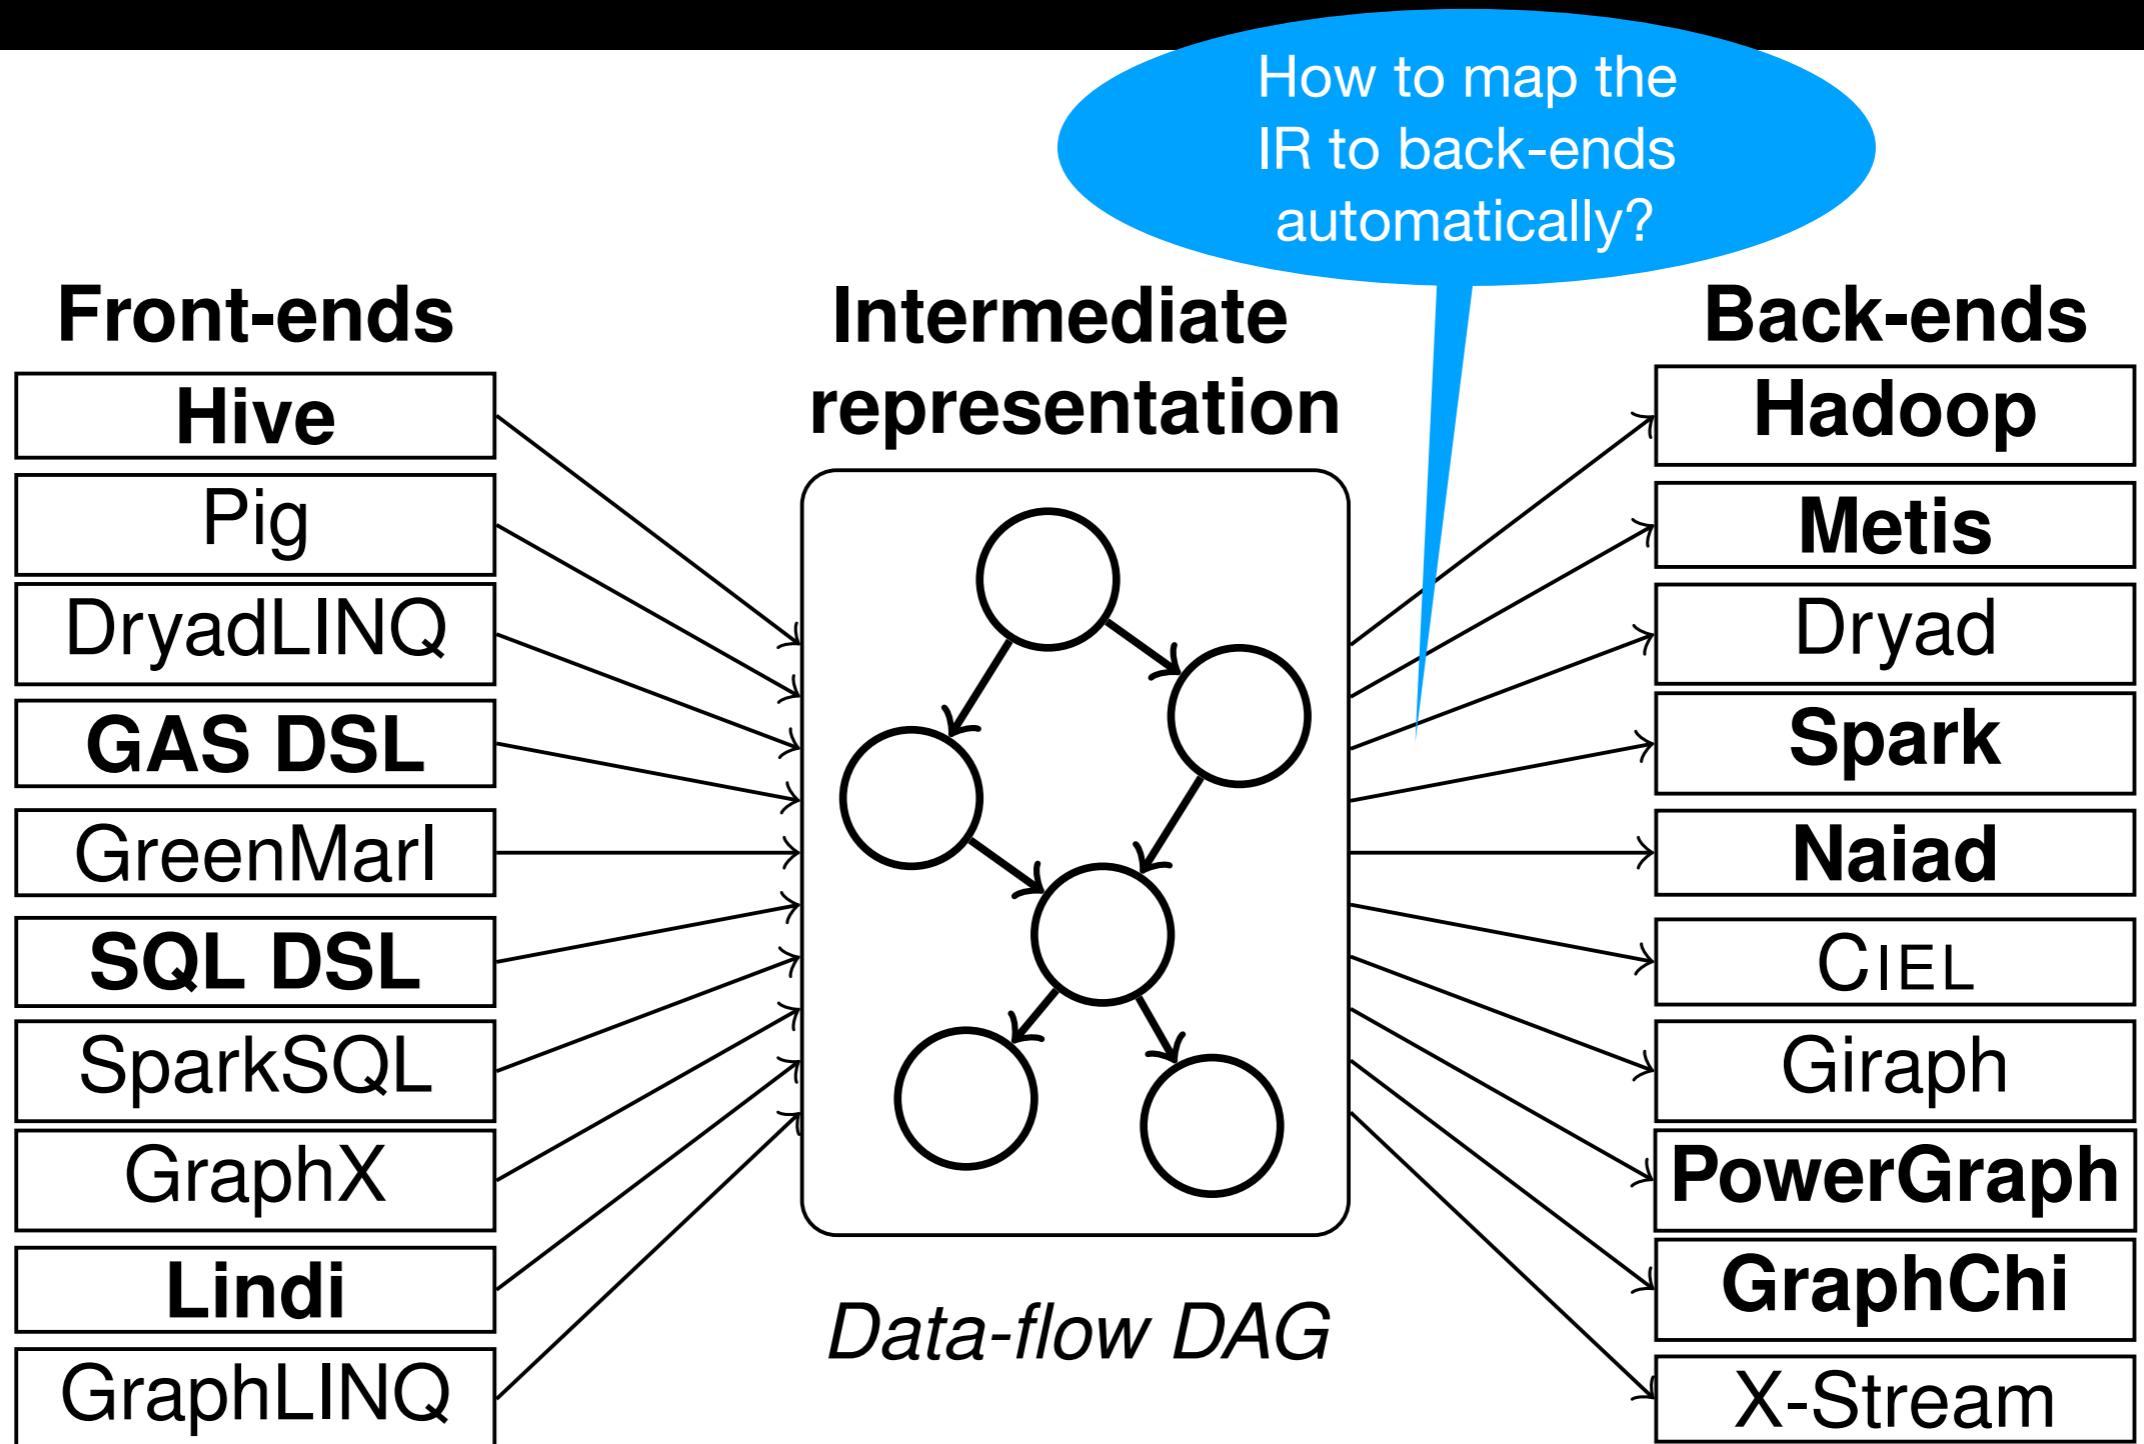


**Rule-based**

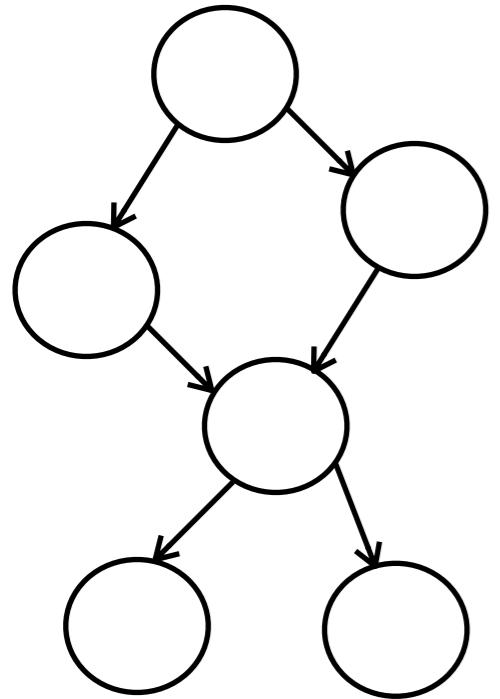
# Musketeer



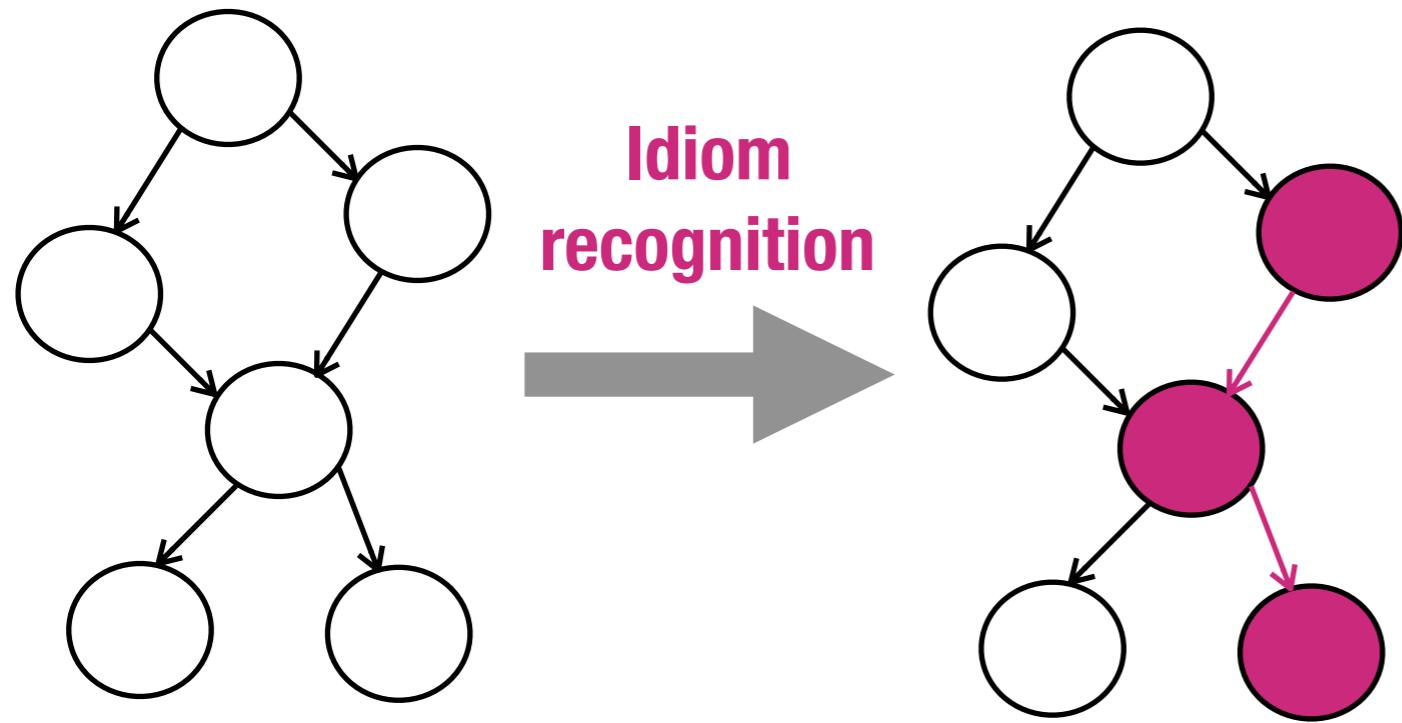
# Musketeer



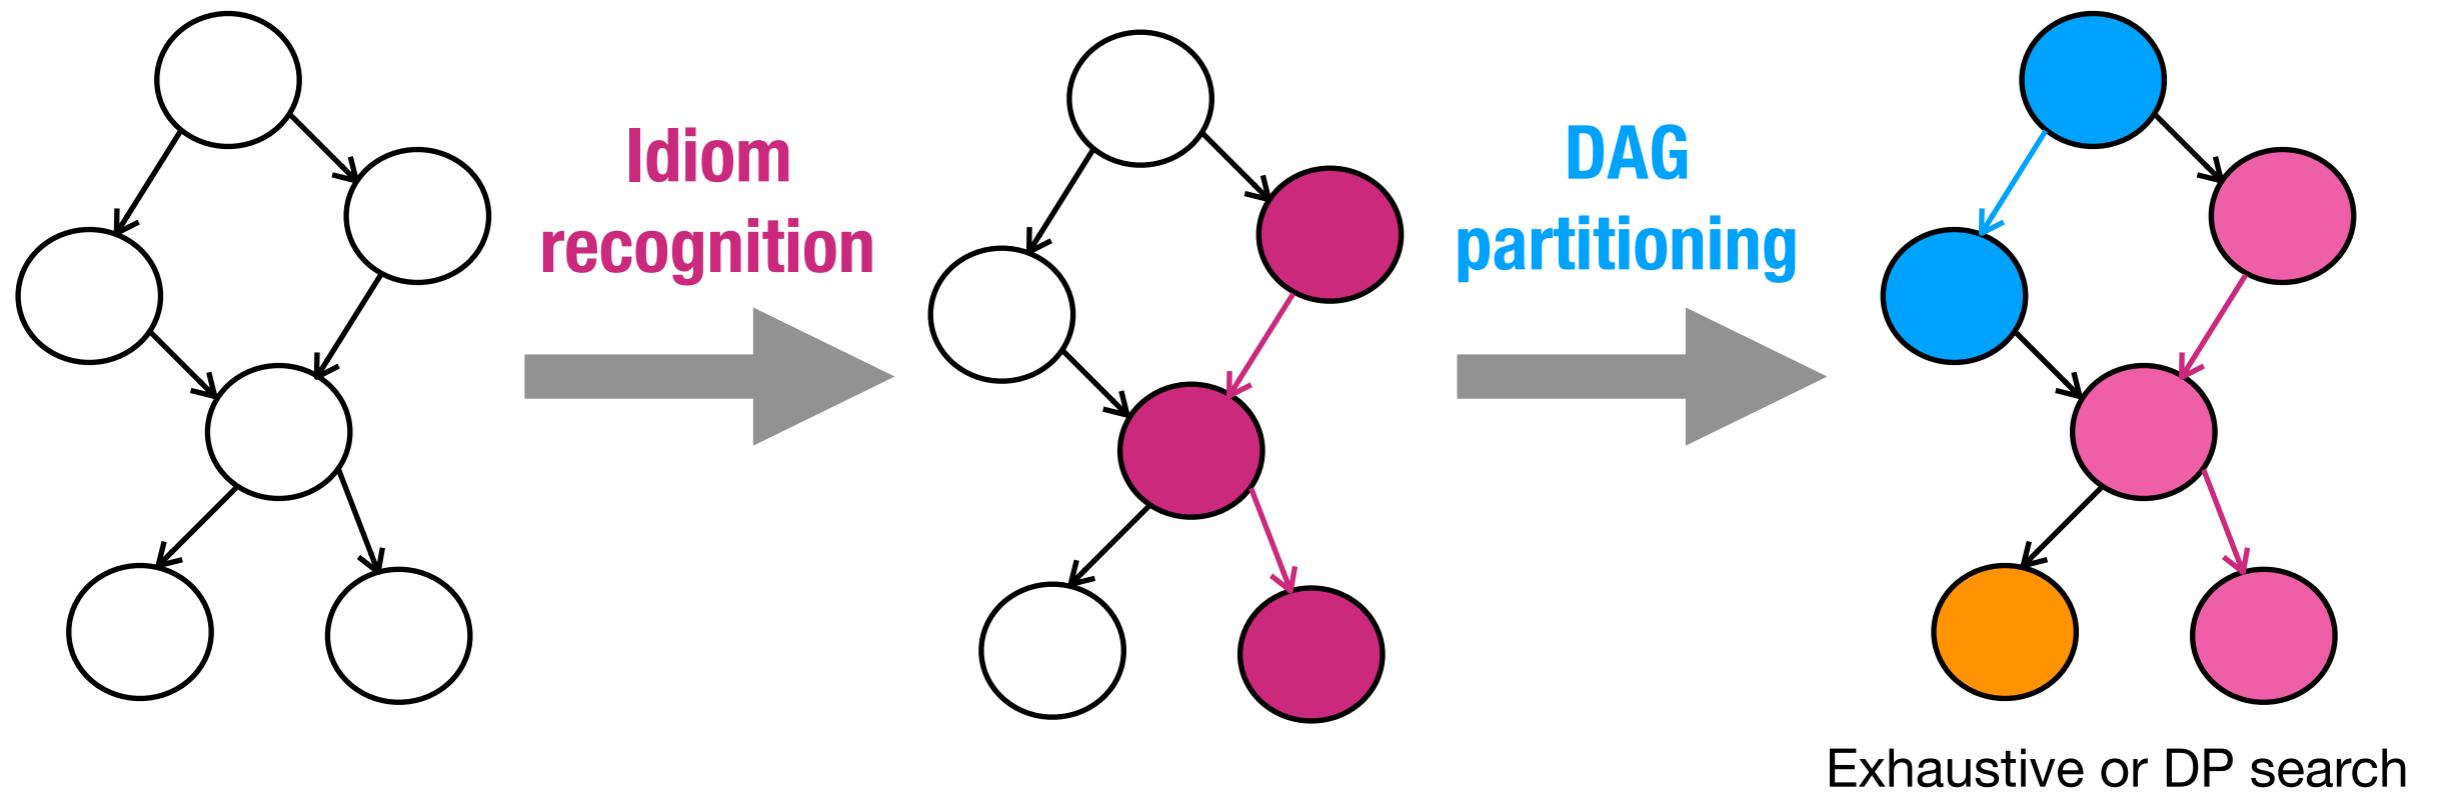
# Musketeer



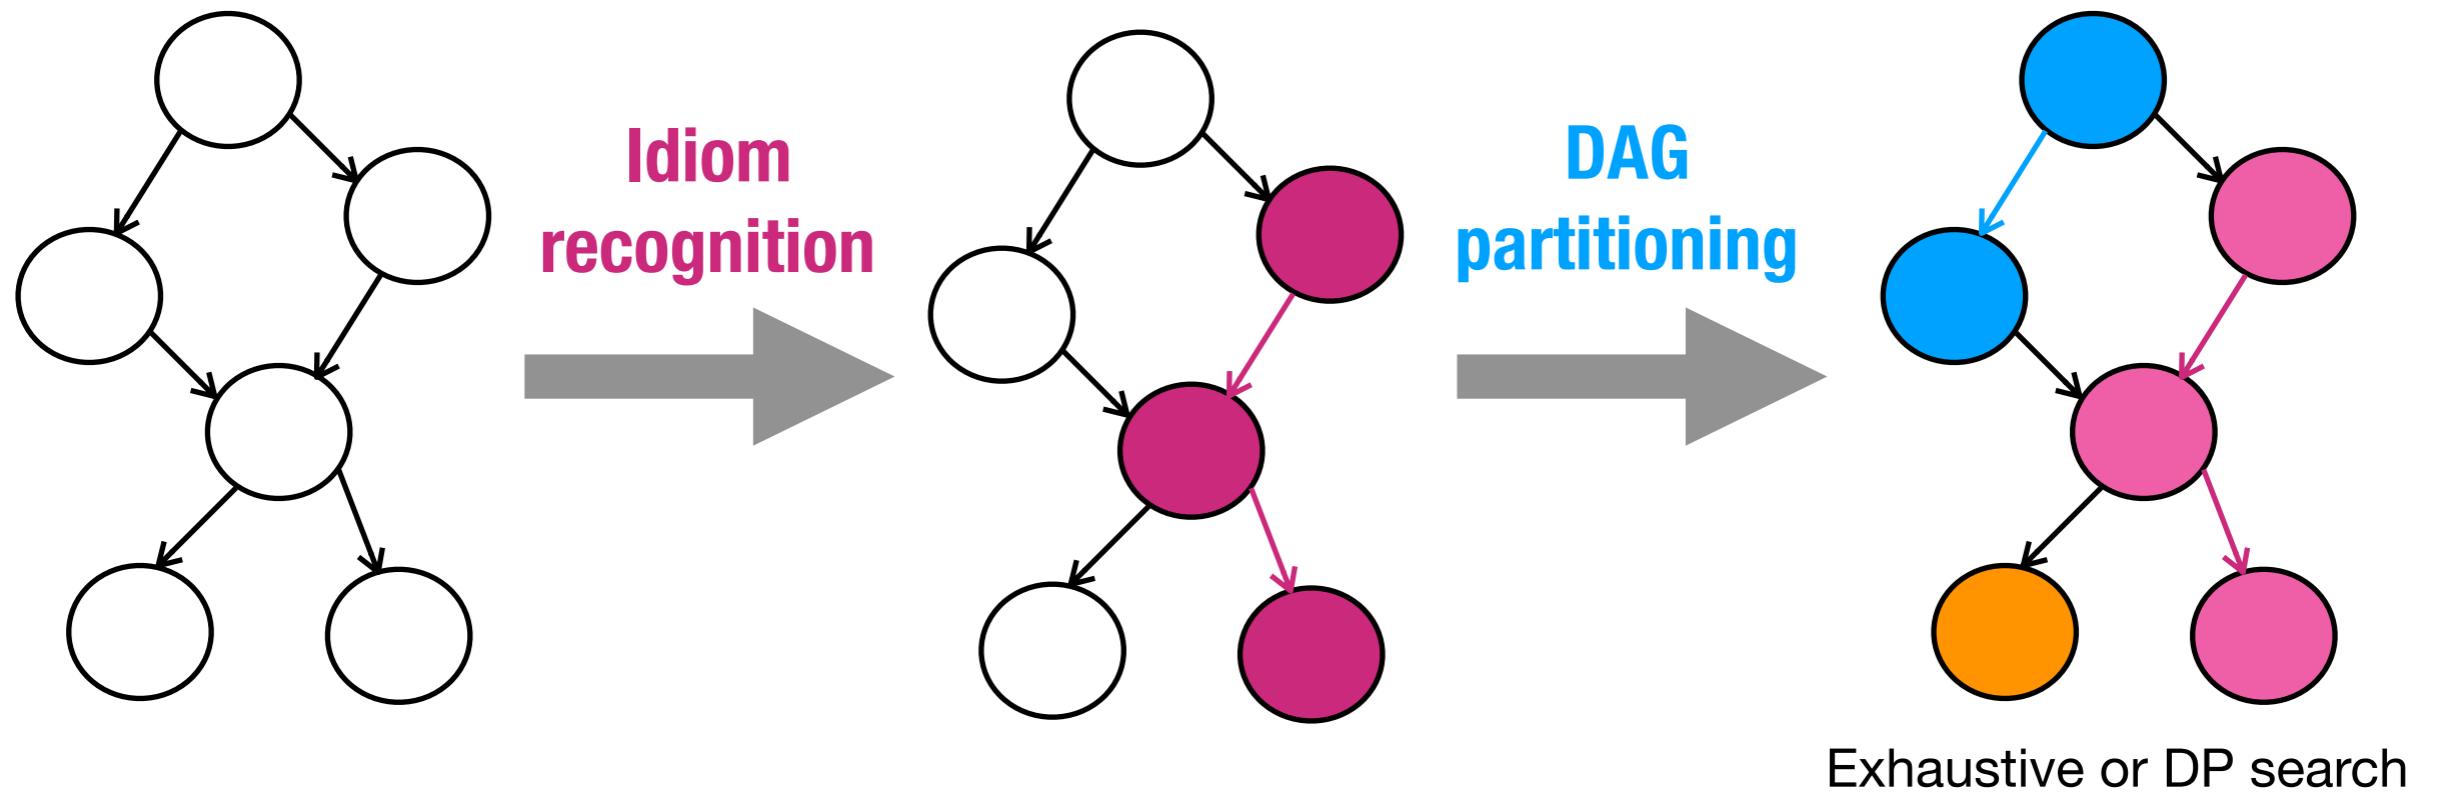
# Musketeer



# Musketeer



# Musketeer



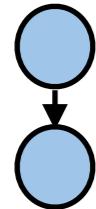
**Analytical Cost Model**

# CP Optimization Features

	Method	Enumeration	Cost Model	Data Movement Cost	User-Specified Platform
Musketeer	Cost-based	Exhaustive/DP	Analytical	No	No

# Rheem

*RHEEM plan  
(platform-agnostic)*



**Rheem's optimizer**

# Rheem

*RHEEM plan  
(platform-agnostic)*



**Plan  
Inflation**

**Rheem's optimizer**

# Rheem

Plan Inflation

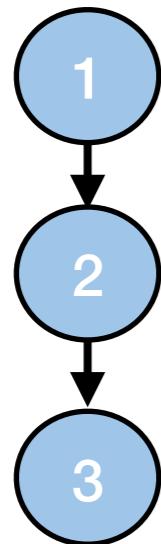
## Rheem plan



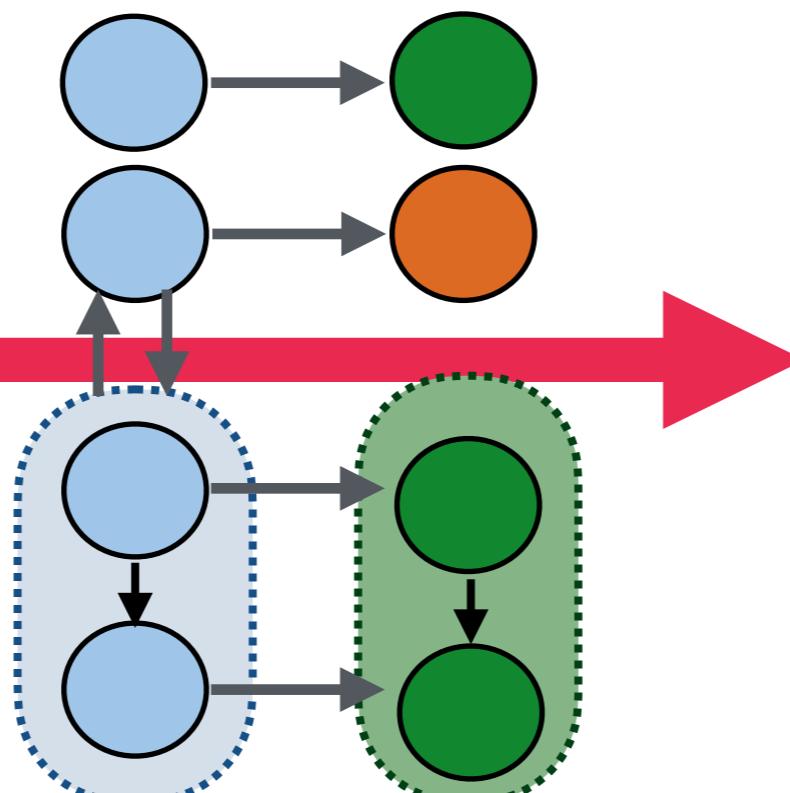
# Rheem

## Plan Inflation

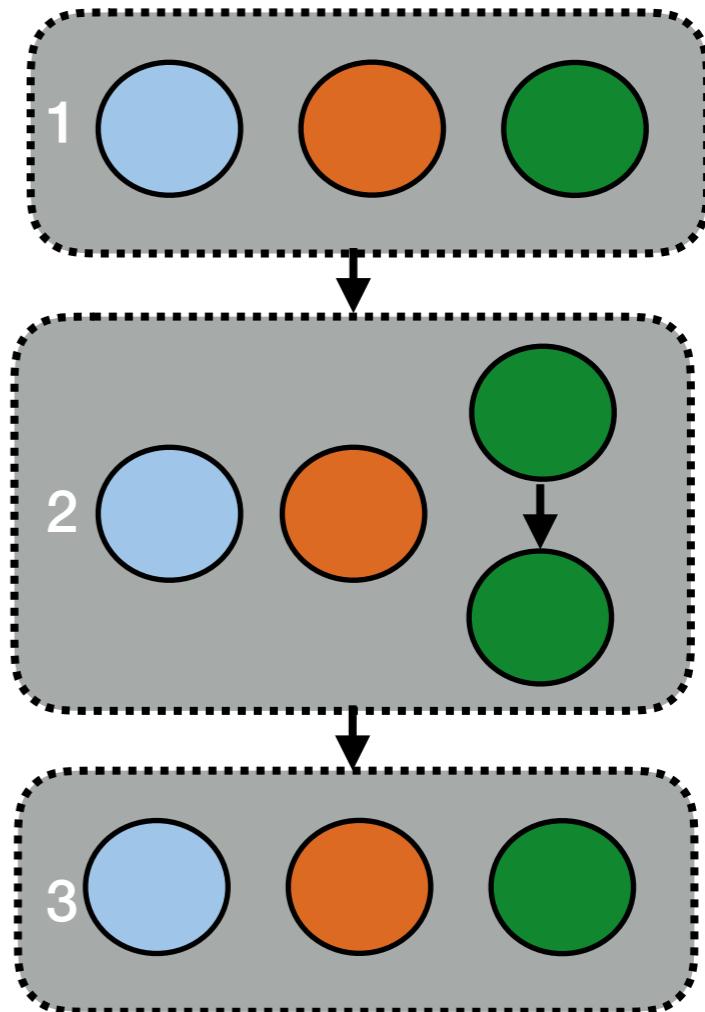
**Rheem plan**



*Mappings*



**Inflated plan**



# Rheem

*RHEEM plan  
(platform-agnostic)*

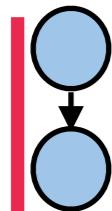


**Plan  
Inflation**

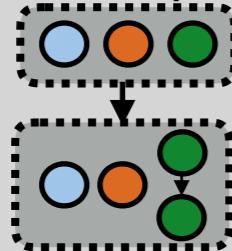
**Rheem's optimizer**

# Rheem

*RHEEM plan  
(platform-agnostic)*



*Inflated plan*



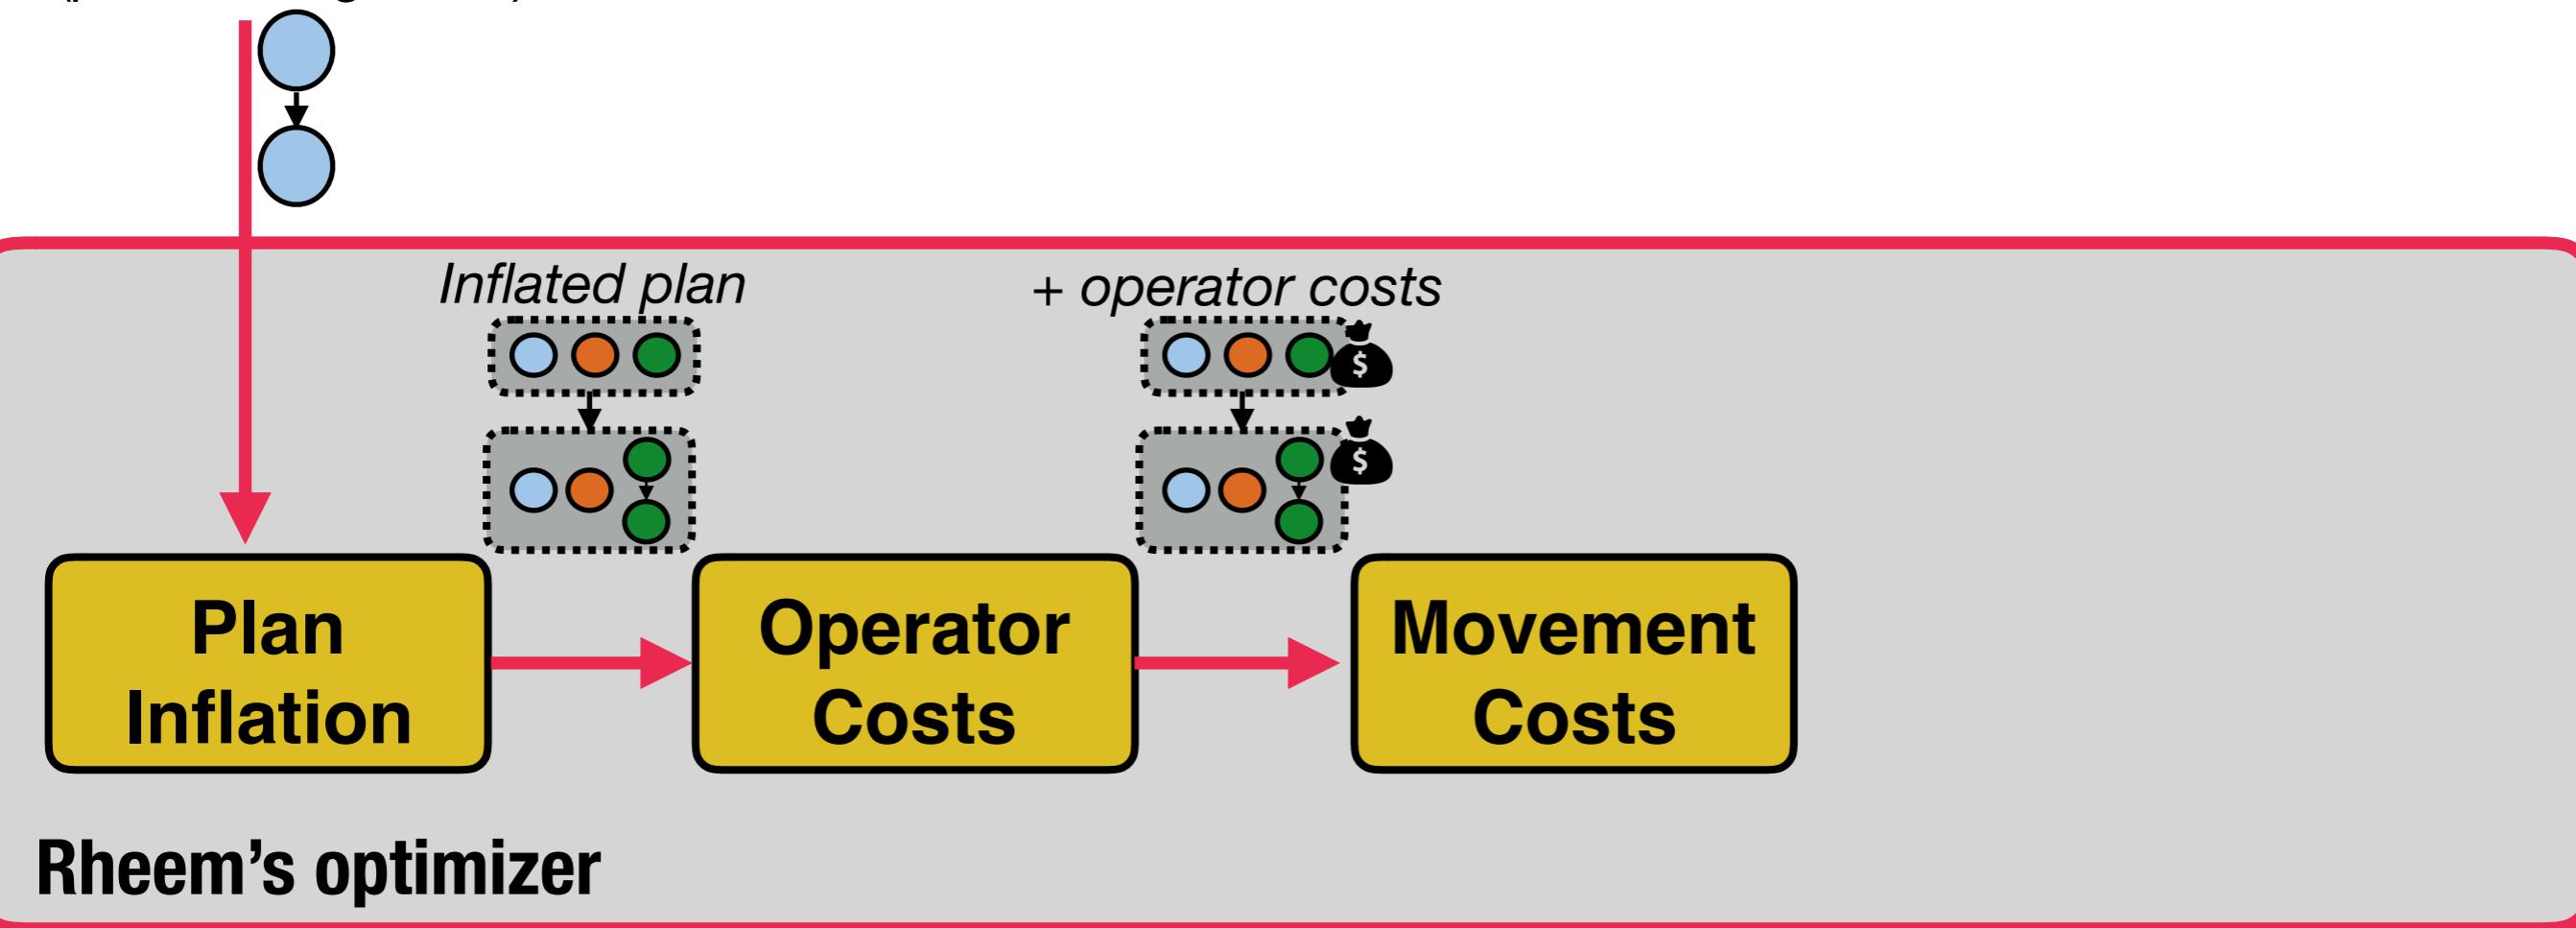
**Plan  
Inflation**

**Operator  
Costs**

**Rheem's optimizer**

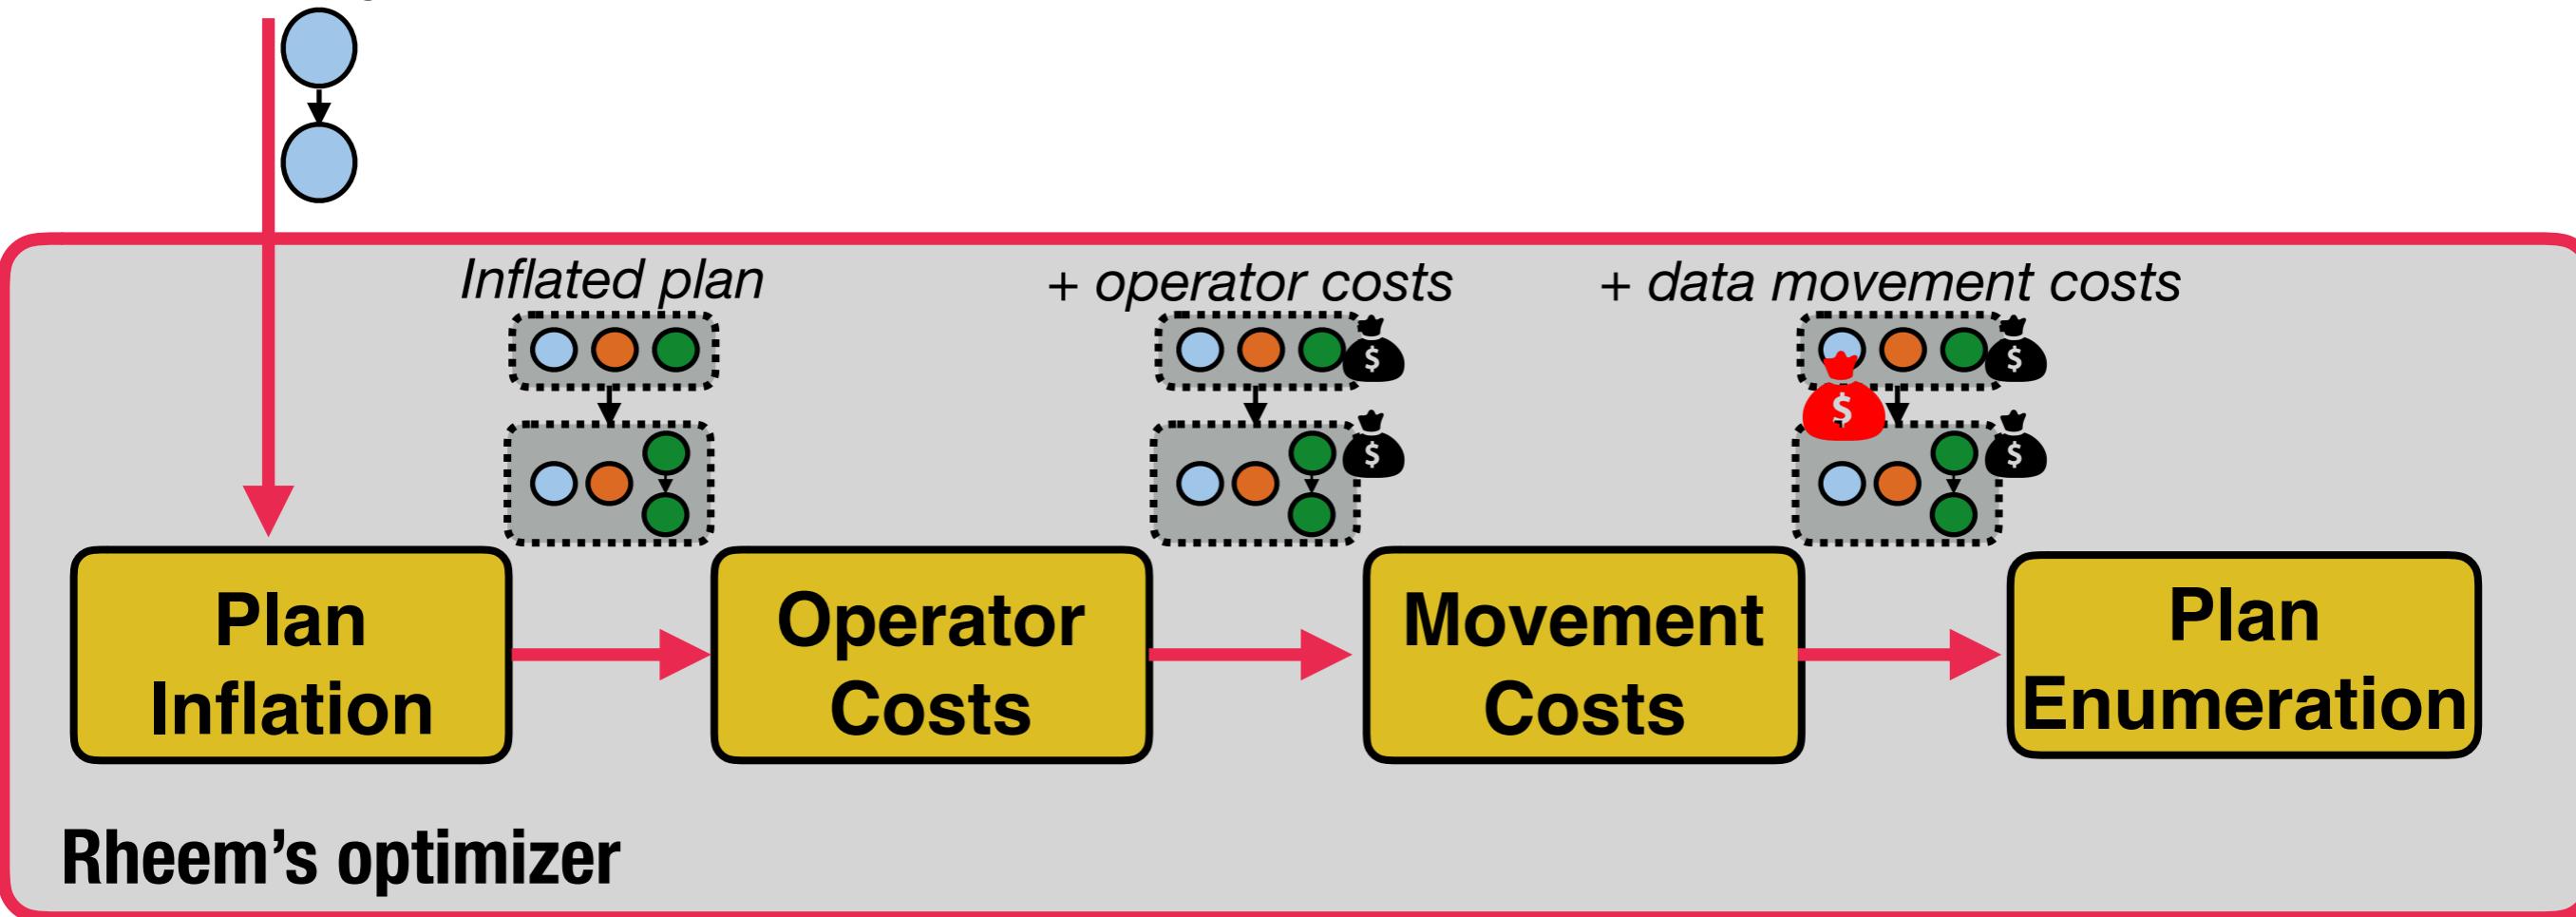
# Rheem

*RHEEM plan  
(platform-agnostic)*



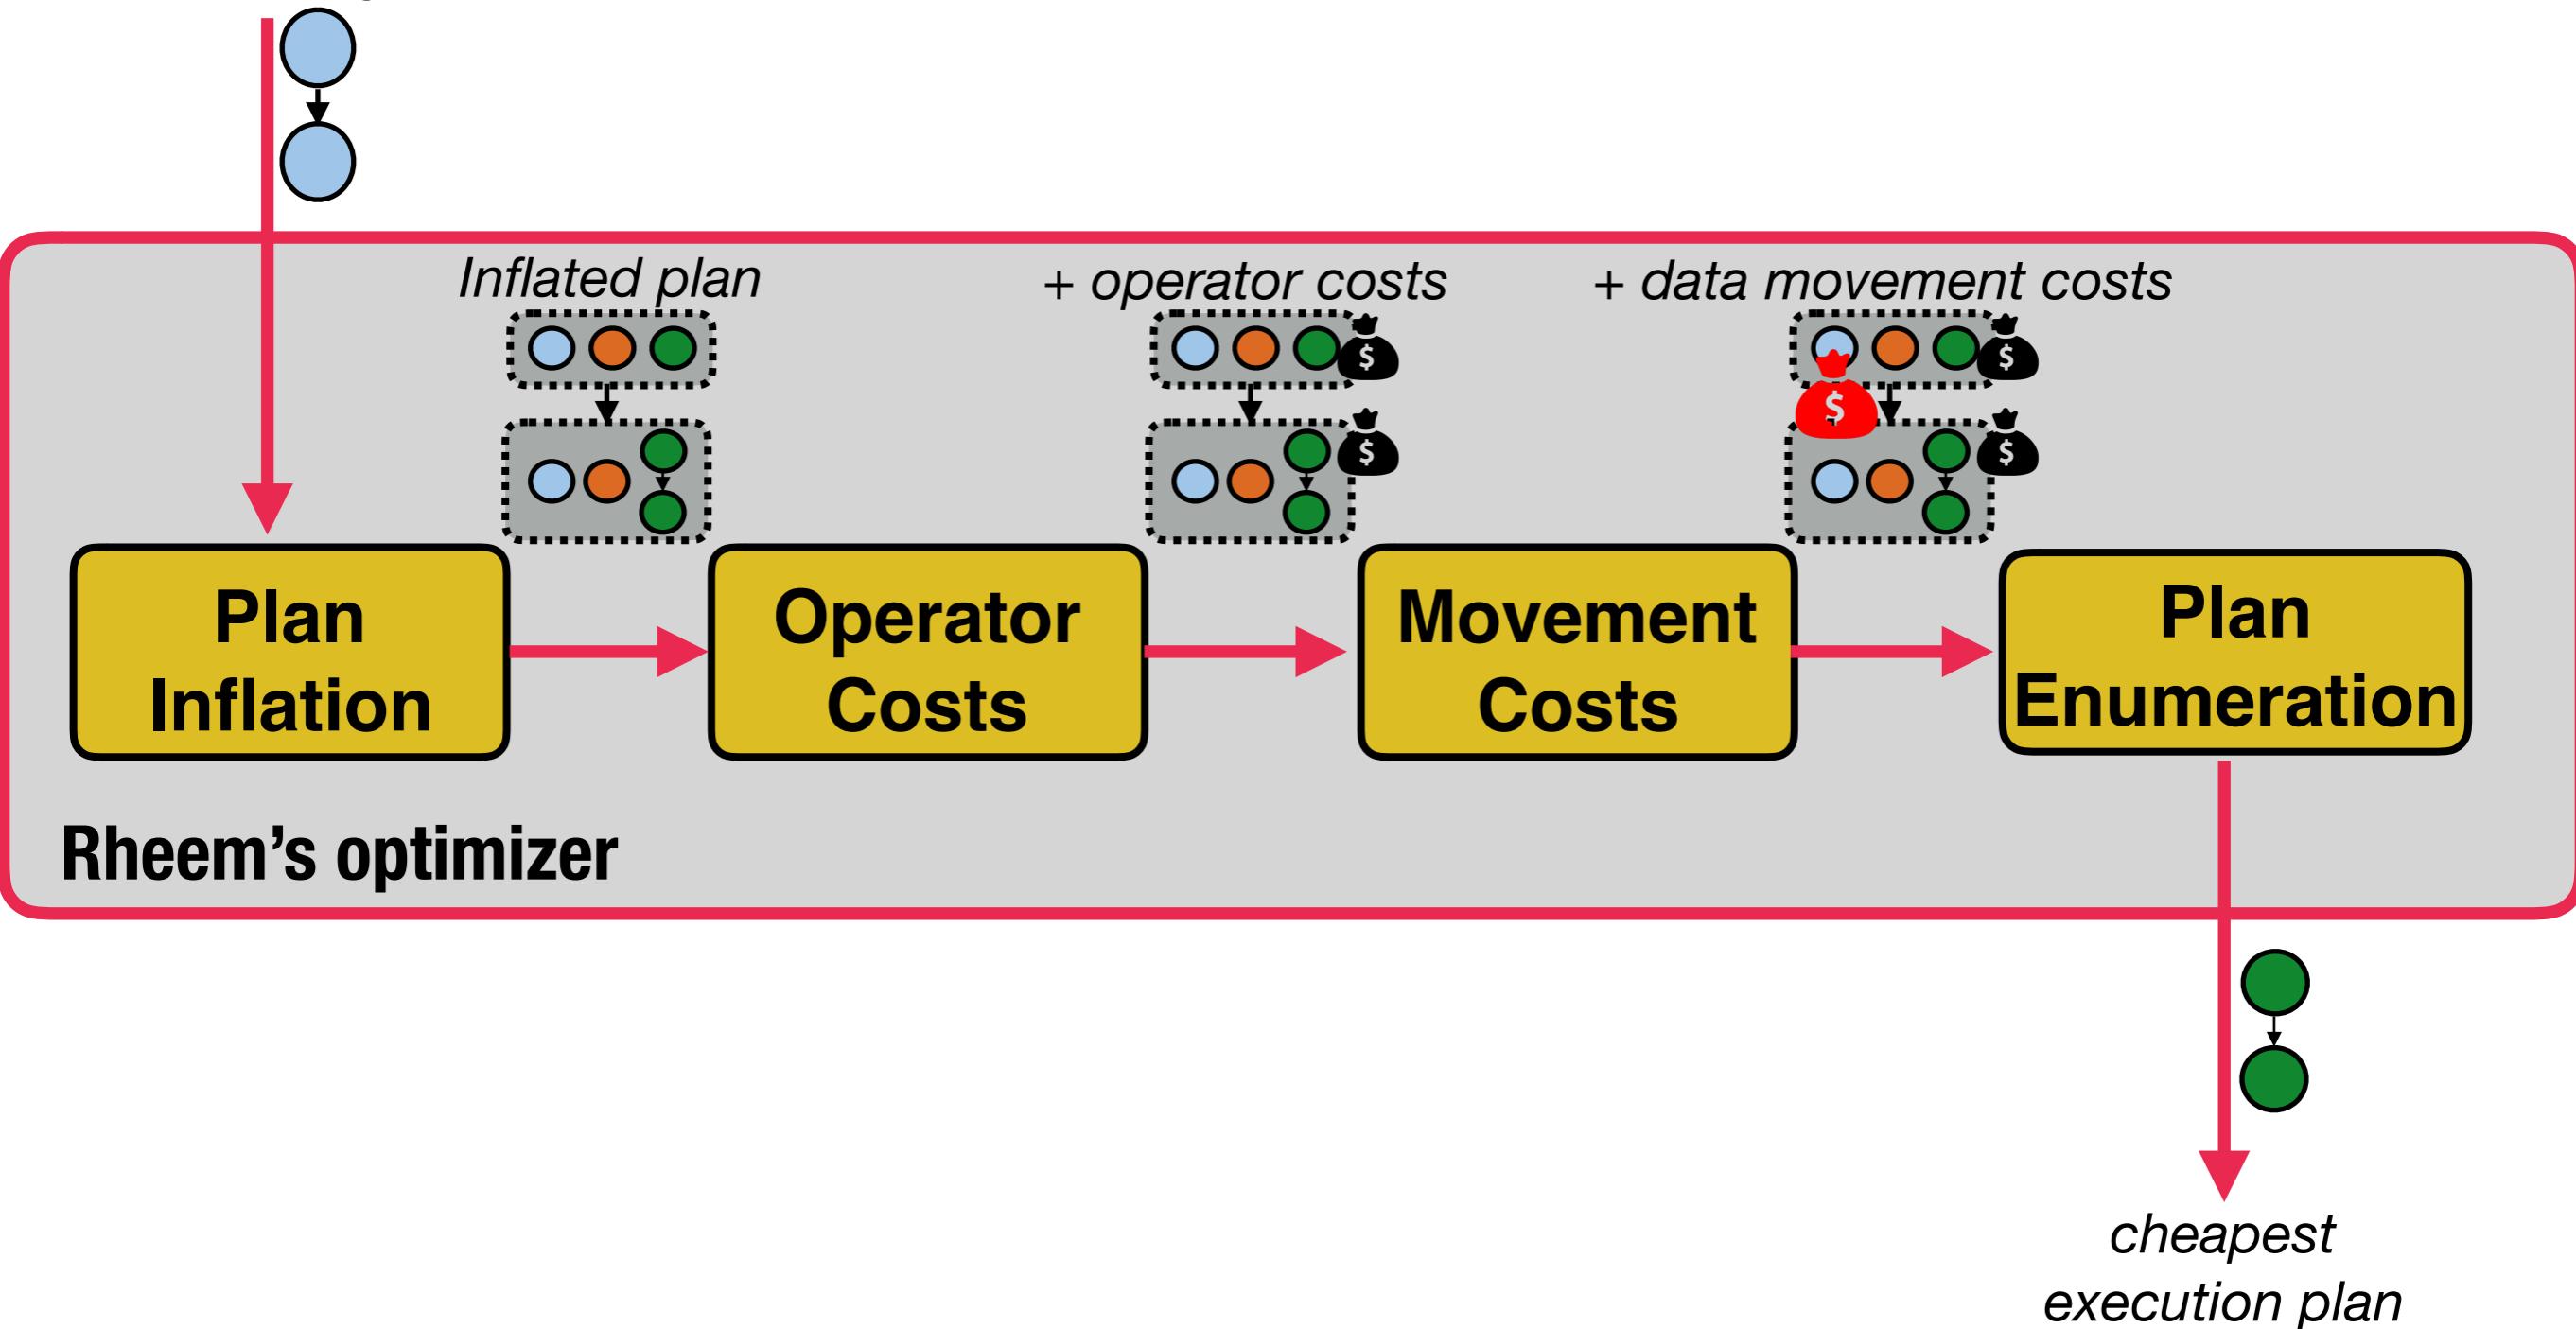
# Rheem

*RHEEM plan  
(platform-agnostic)*



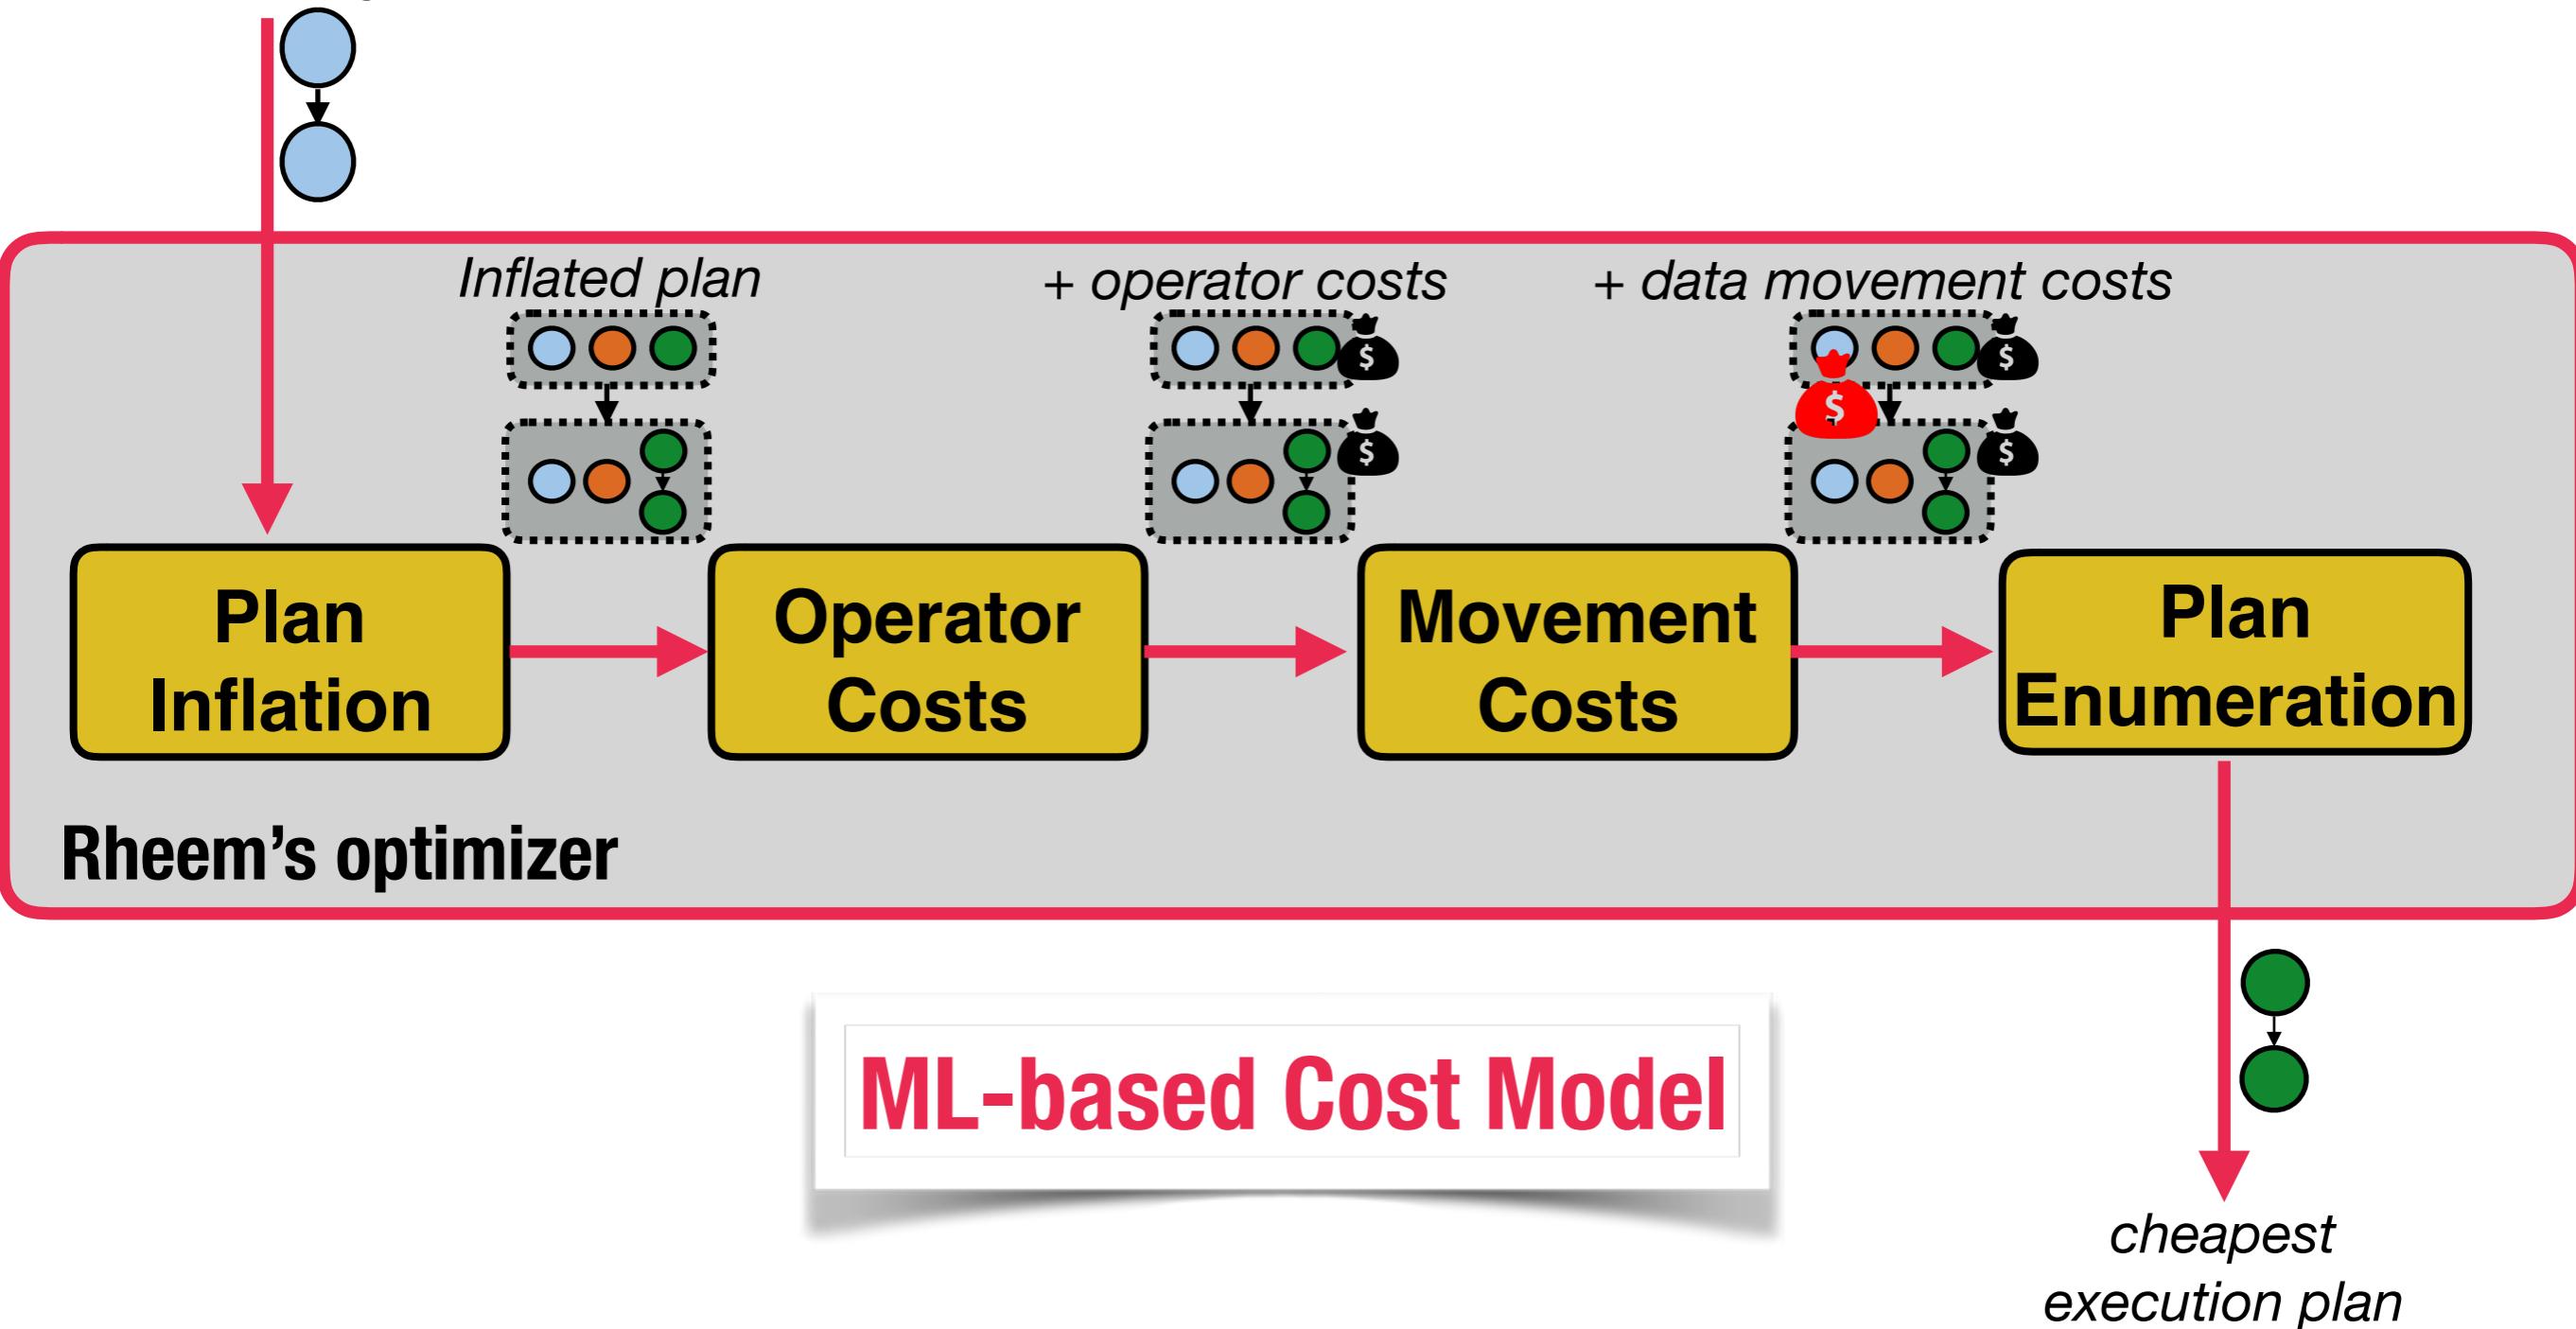
# Rheem

*RHEEM plan  
(platform-agnostic)*



# Rheem

*RHEEM plan  
(platform-agnostic)*



# CP Optimization Features

	Method	Enumeration	Cost Model	Data Movement Cost	User-Specified Platform
Musketeer	Cost-based	Exhaustive/DP	Analytical	No	No
Rheem	Cost-based	Enumeration algebra	ML-based	Yes	Operator level

# CP Optimization Features

	Method	Enumeration	Cost Model	Data Movement Cost	User-Specified Platform
Musketeer	Cost-based	Exhaustive/DP	Analytical	No	No
Rheem	Cost-based	Enumeration algebra	ML-based	Yes	Operator level
Ires	Cost-based	DP	ML-based	Yes	Operator level
QoX	Cost-based	Greedy	Analytical	Yes	No
Cyclops	Cost-based	?	ML-based	No	Task level

# Cross-Platform Optimization

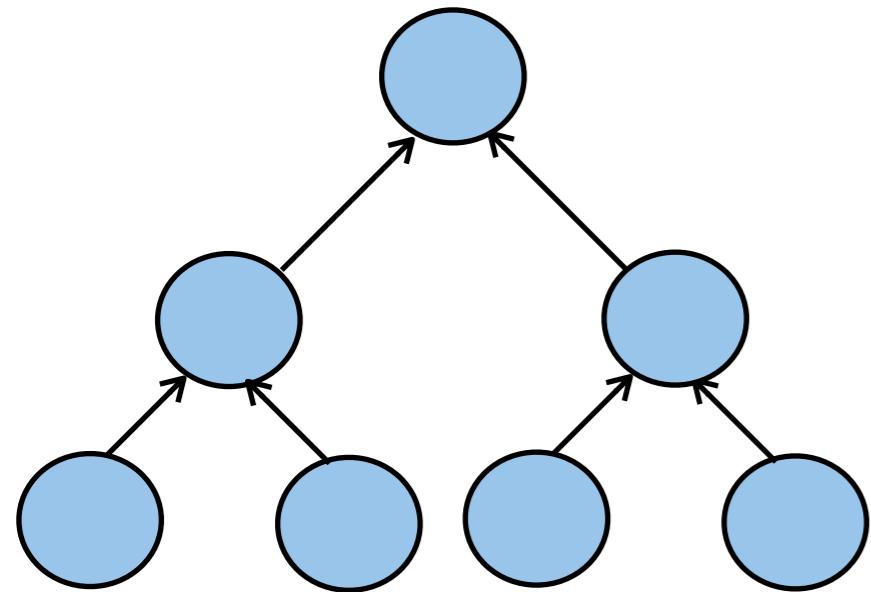


**Cost-based**



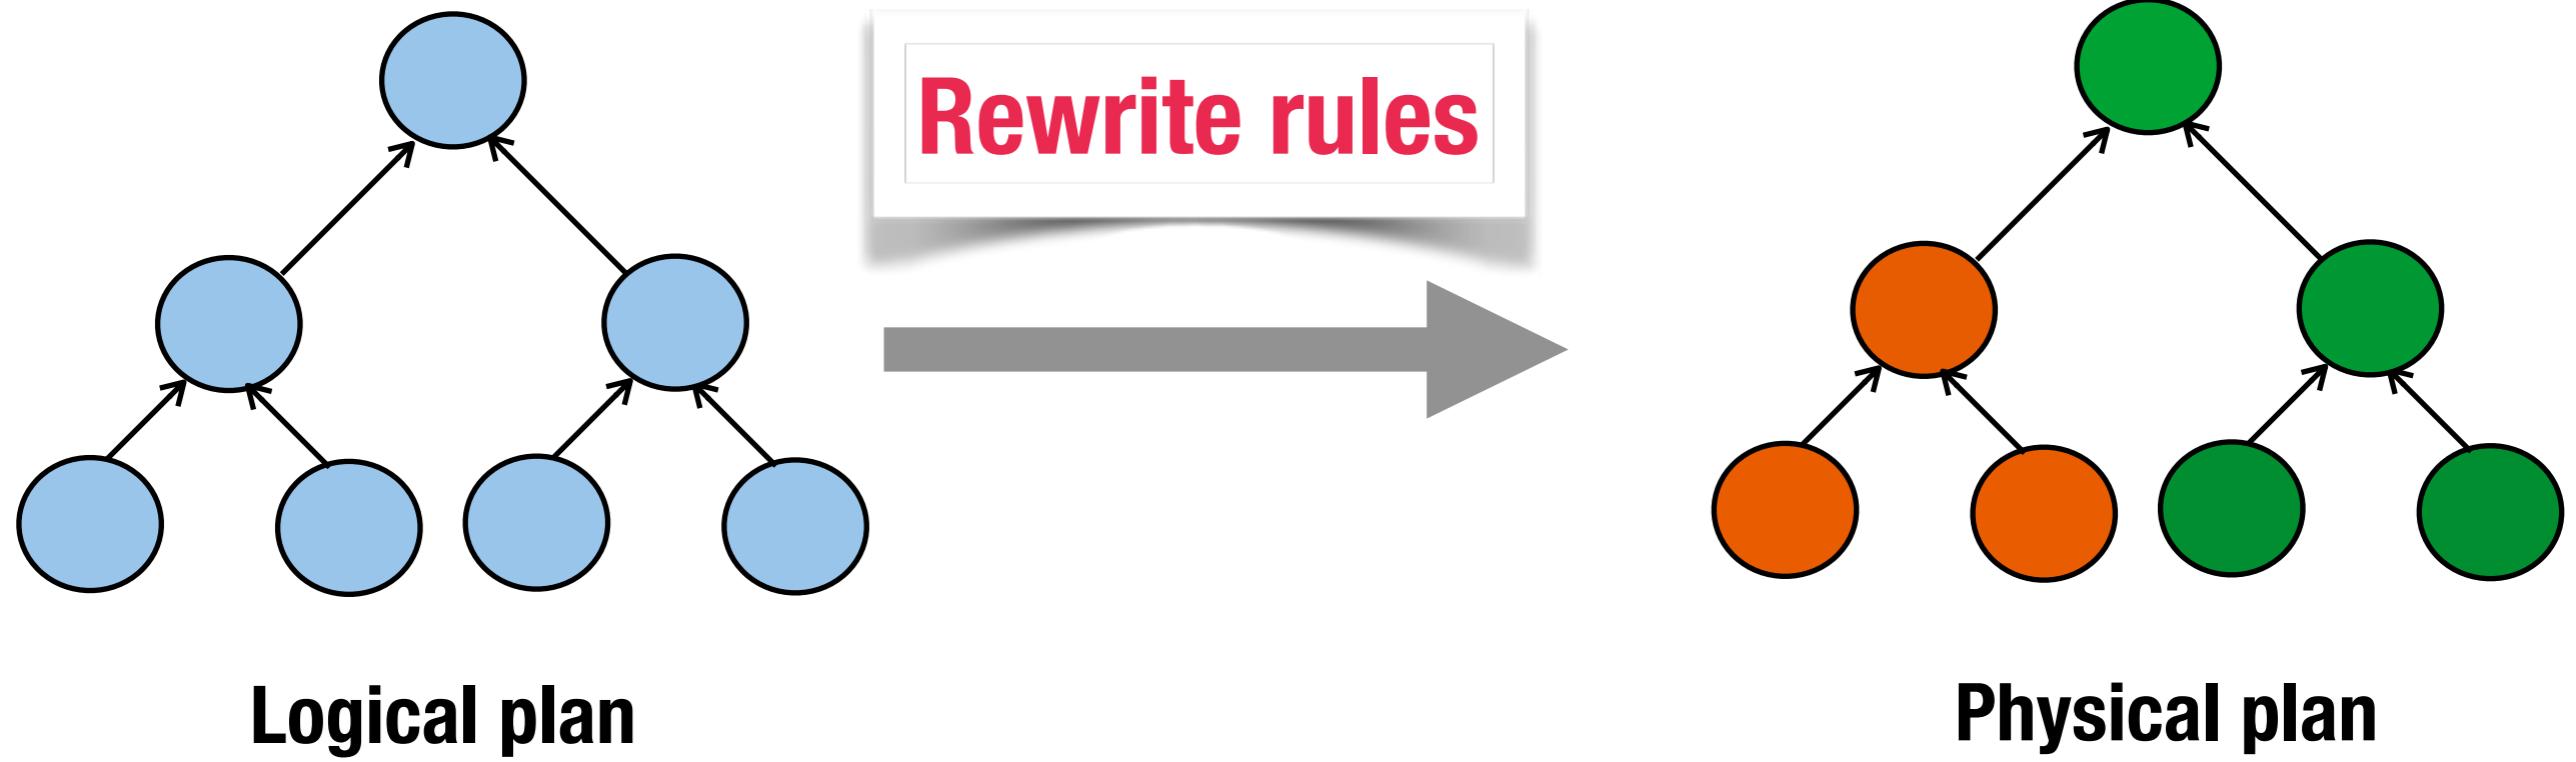
**Rule-based**

# Myria

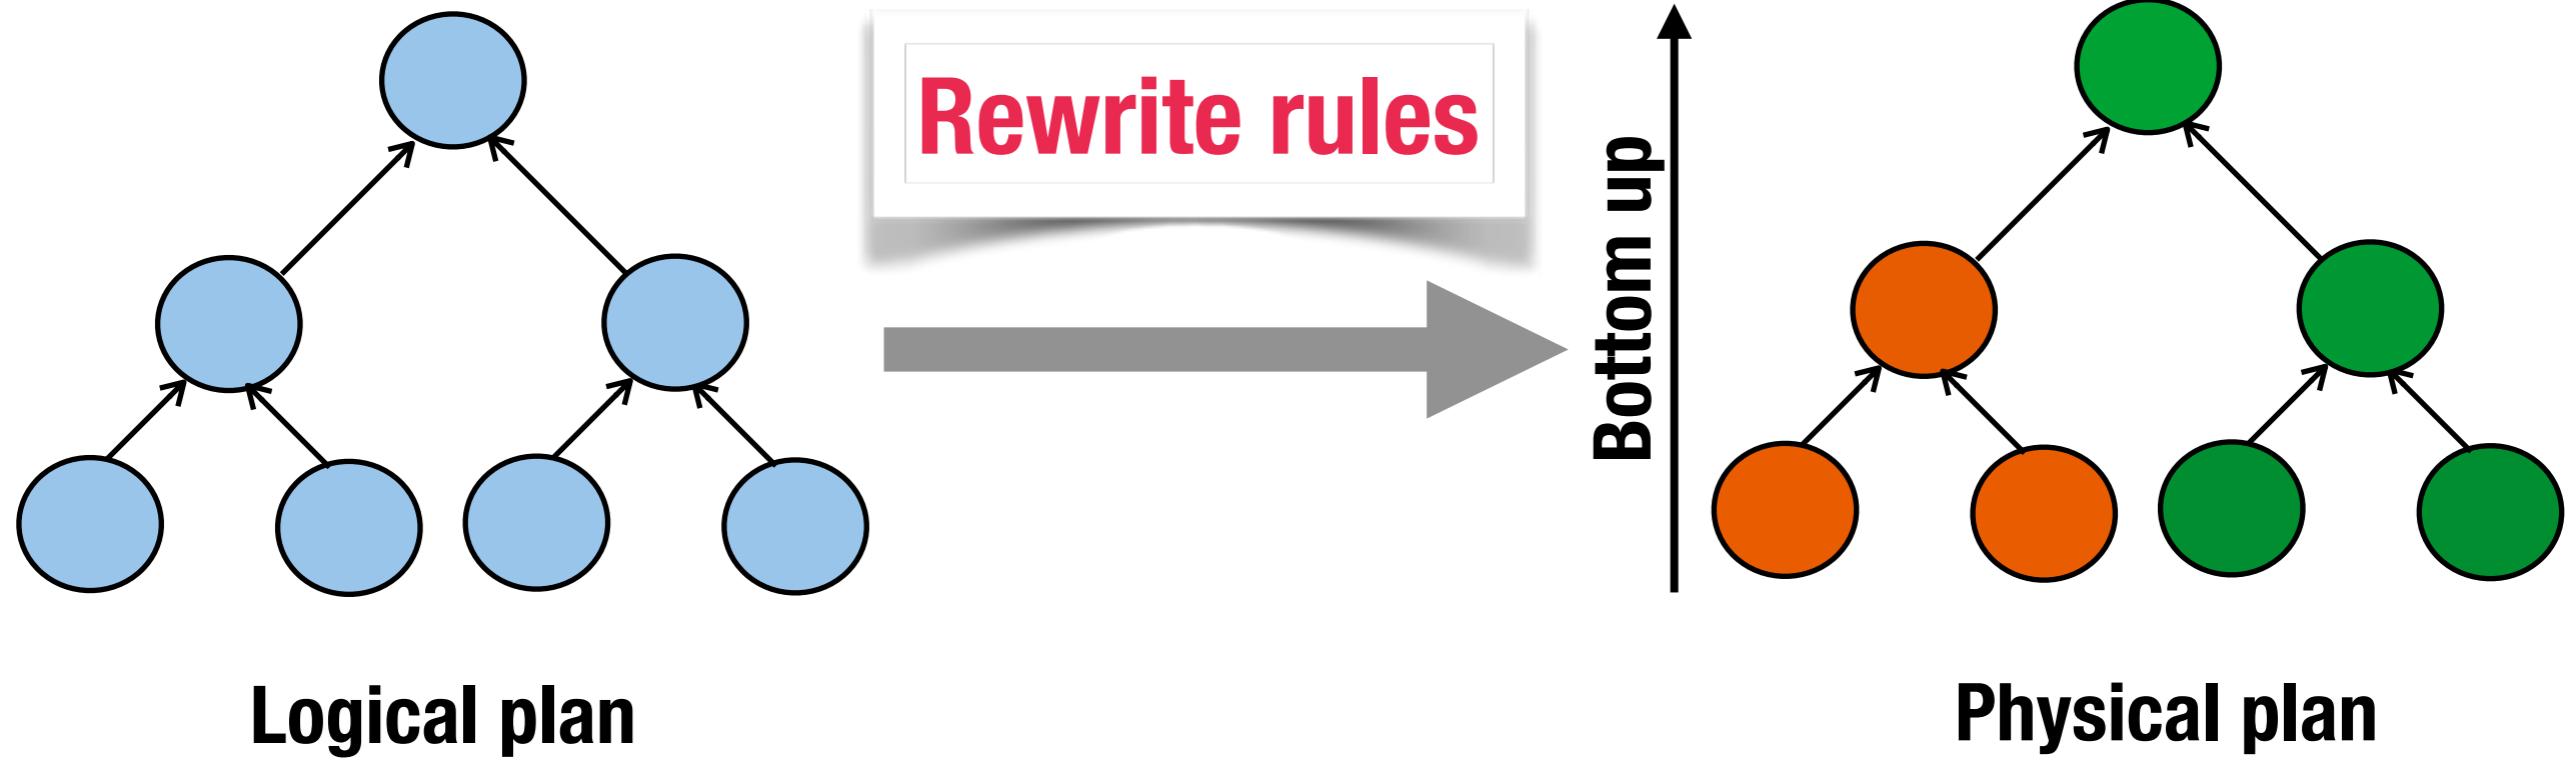


**Logical plan**

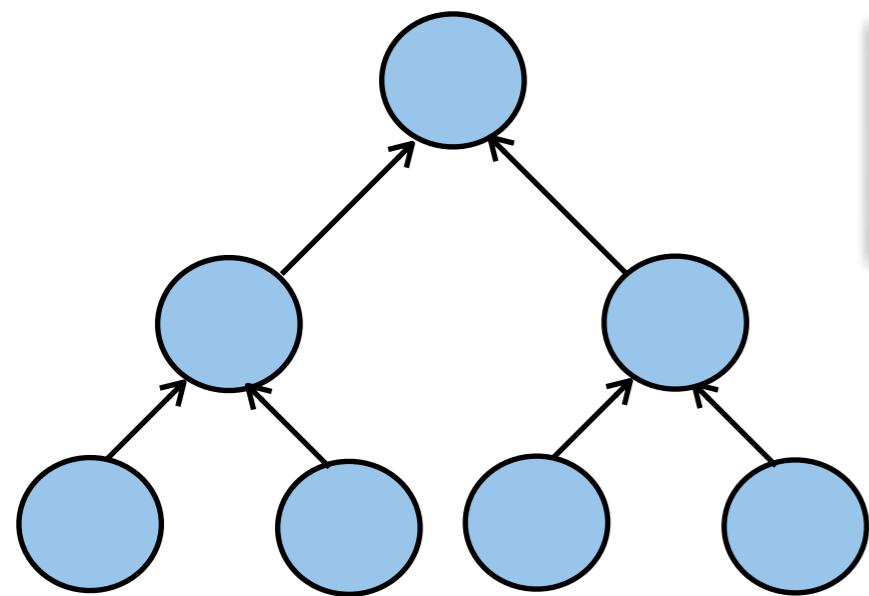
# Myria



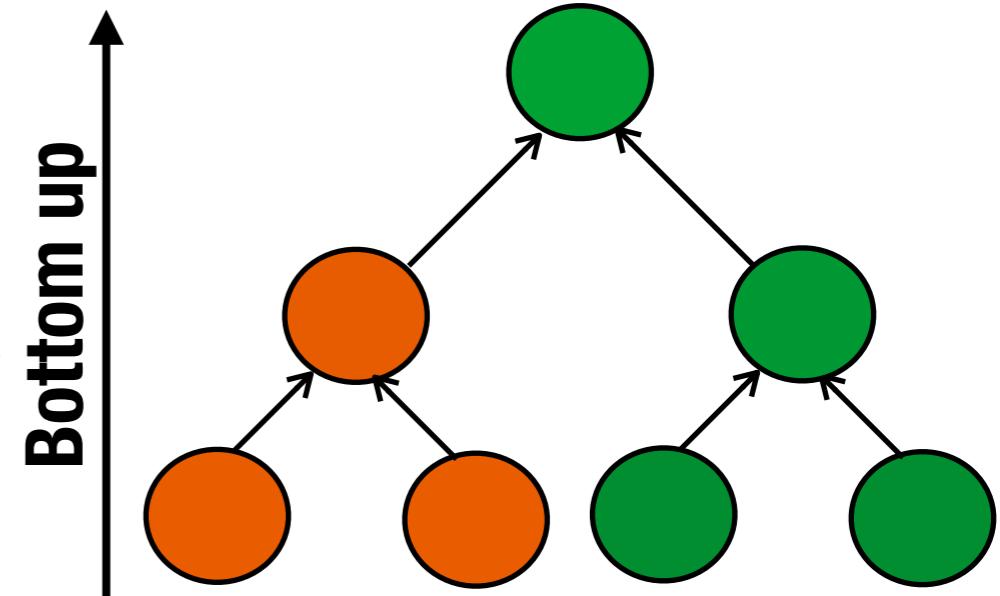
# Myria



# Myria



Logical plan



Physical plan

Leaf nodes executed  
where the data resides

# CP Optimization Features

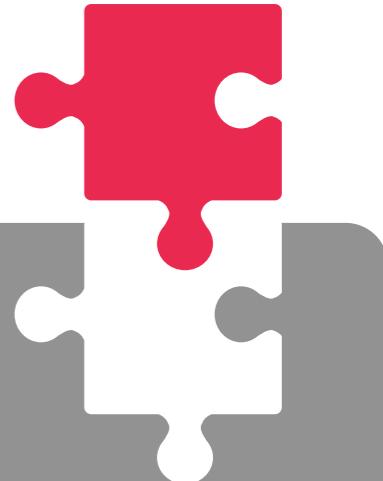
	<b>Method</b>	<b>Enumeration</b>	<b>Cost Model</b>	<b>Data Movement Cost</b>	<b>User-Specified Platform</b>
<b>Musketeer</b>	Cost-based	Exhaustive/DP	Analytical	No	No
<b>Rheem</b>	Cost-based	Enumeration algebra	ML-based	Yes	Operator level
<b>Ires</b>	Cost-based	DP	ML-based	Yes	Operator level
<b>QoX</b>	Cost-based	Greedy	Analytical	Yes	No
<b>Cyclops</b>	Cost-based	?	ML-based	No	Task level
<b>Myria</b>	Rule-based	Bottom up	NA	NA	No

# Challenges

1 Decoupling Applications

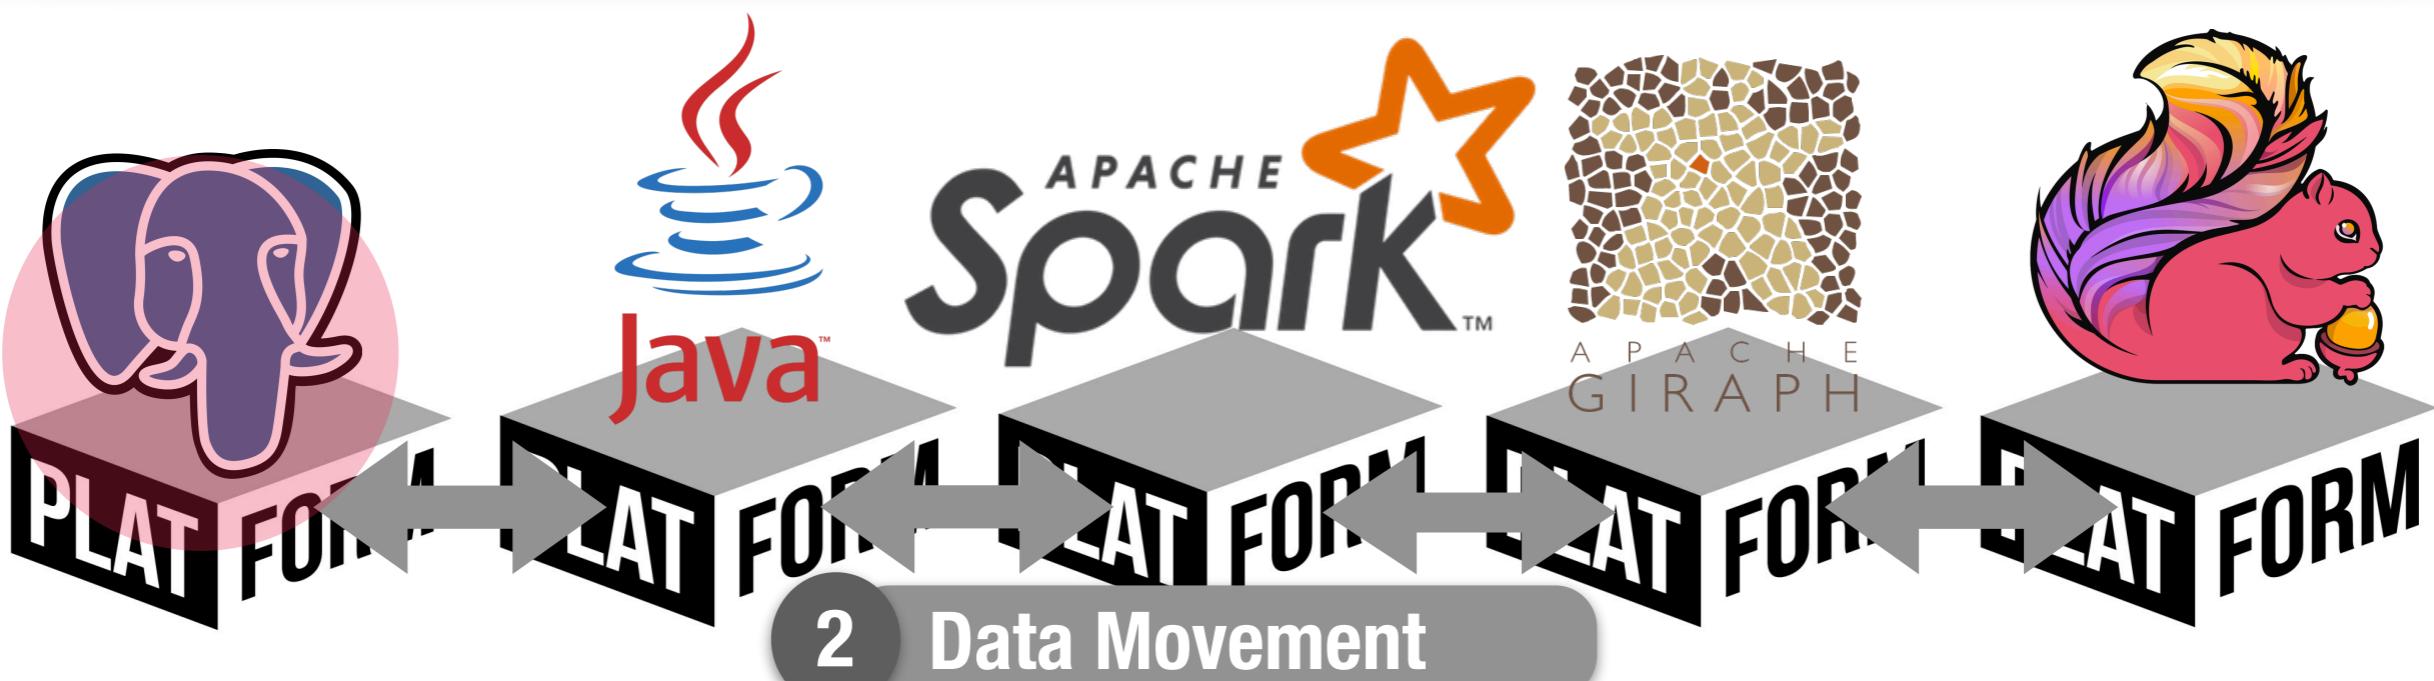


4 Extensibility



3

Automatic Cross-Platform  
Data Processing

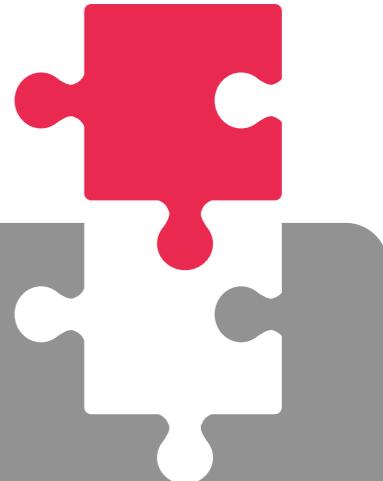


# Adding/Modifying a Platform

1 Decoupling Applications

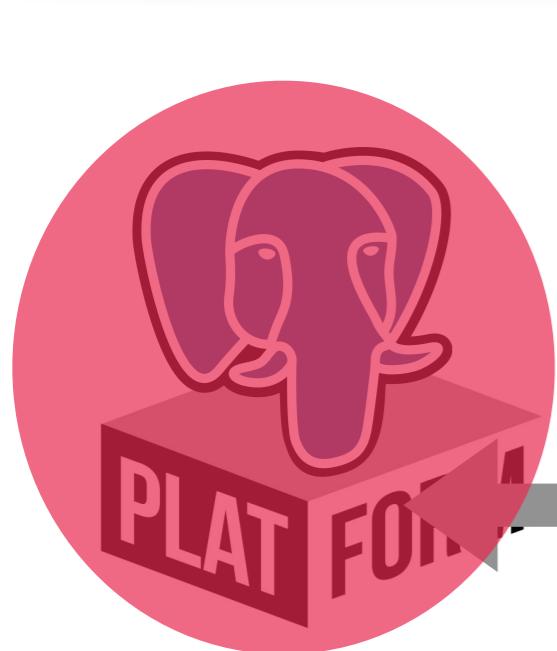


4 Extensibility



3

Automatic Cross-Platform  
Data Processing



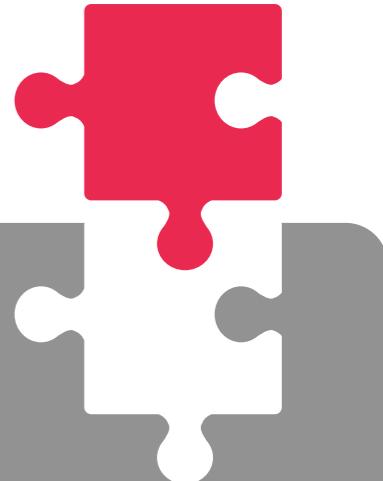
2 Data Movement

# Adding/Modifying a Platform

1 Decoupling Applications

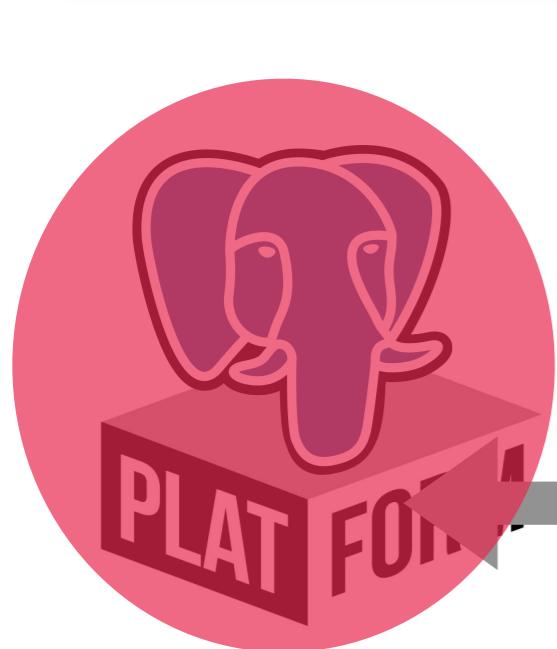


4 Extensibility



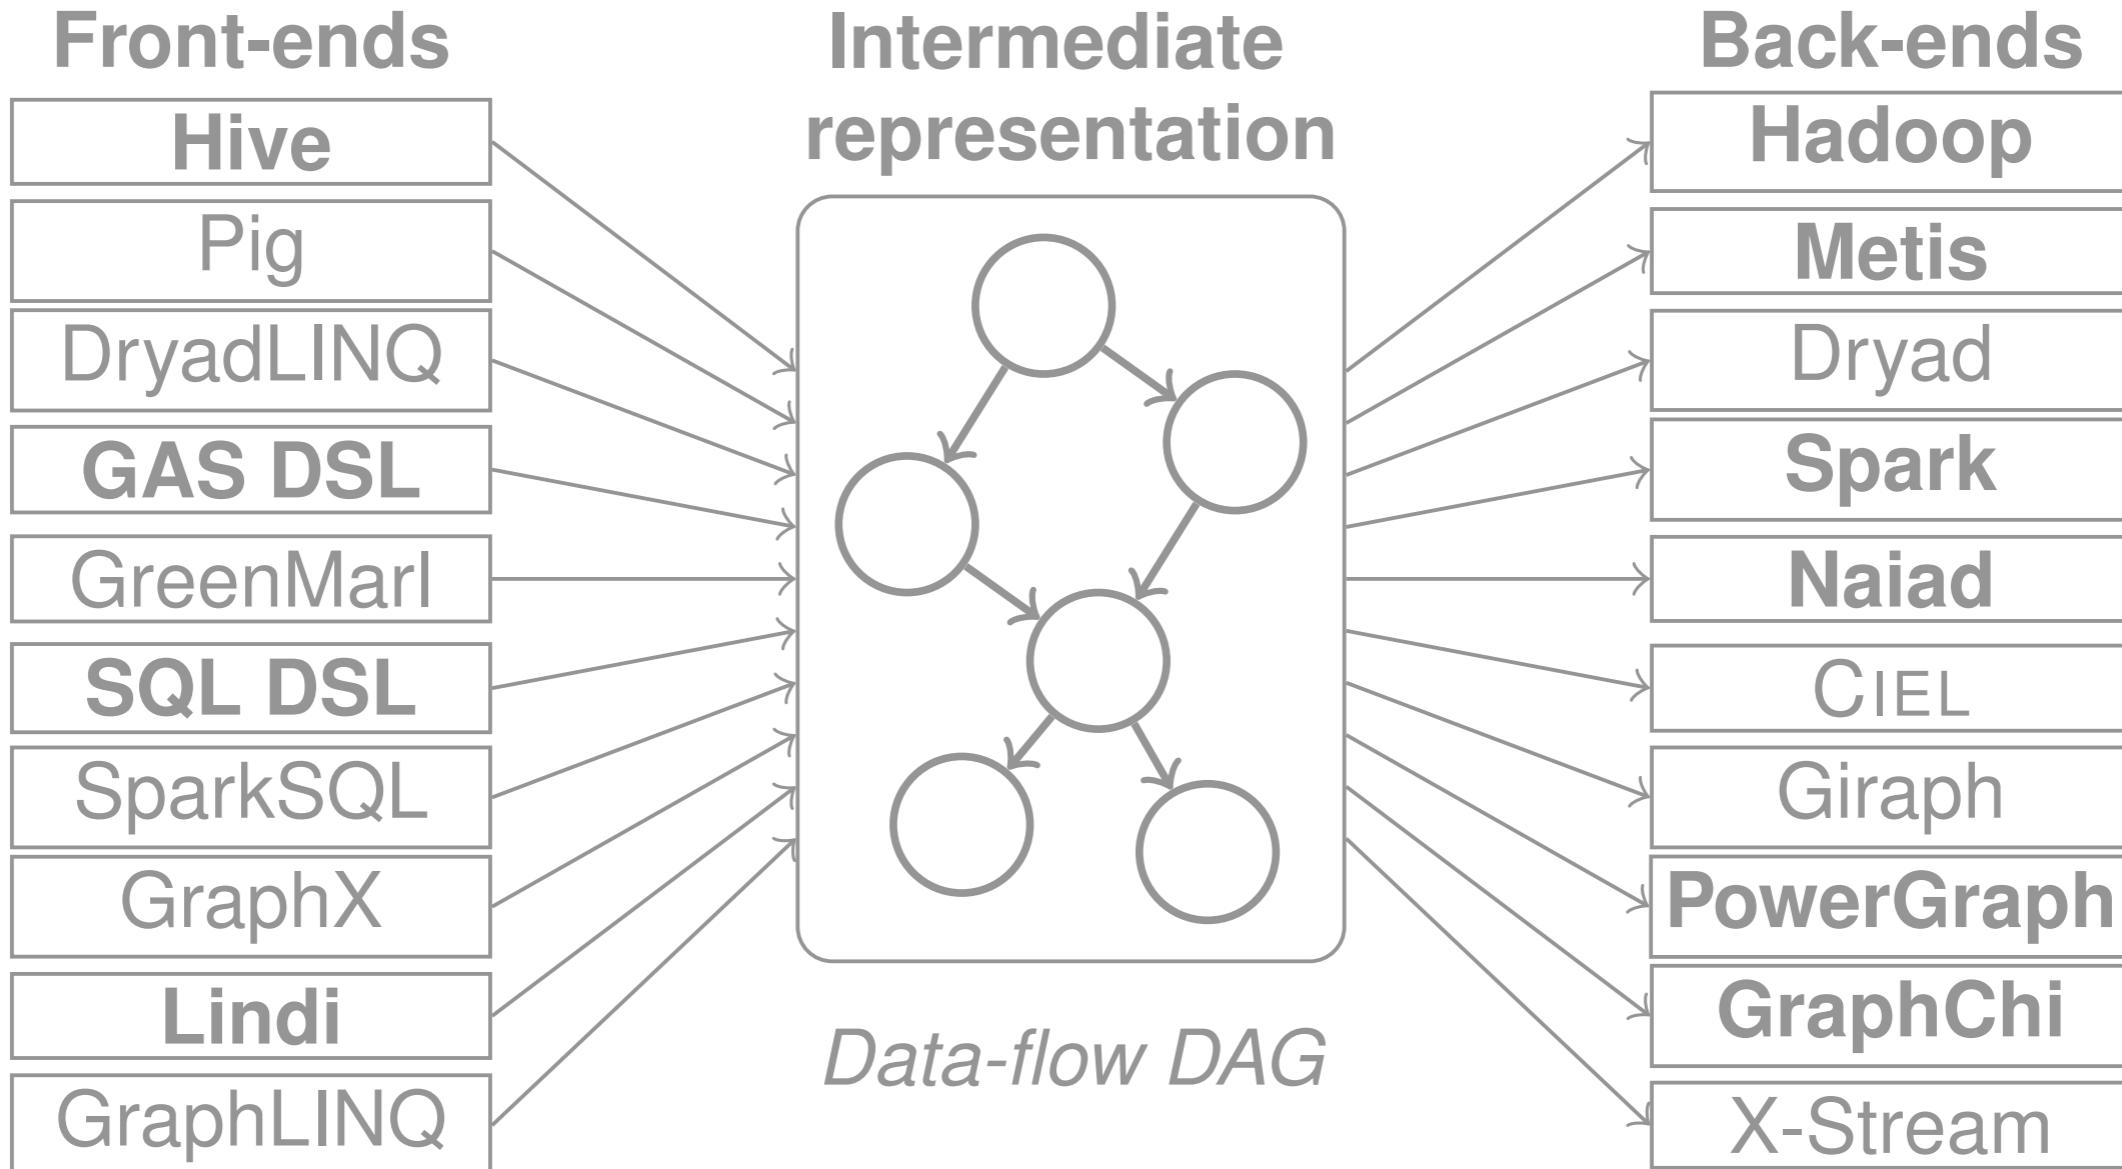
3

Automatic Cross-Platform  
Data Processing



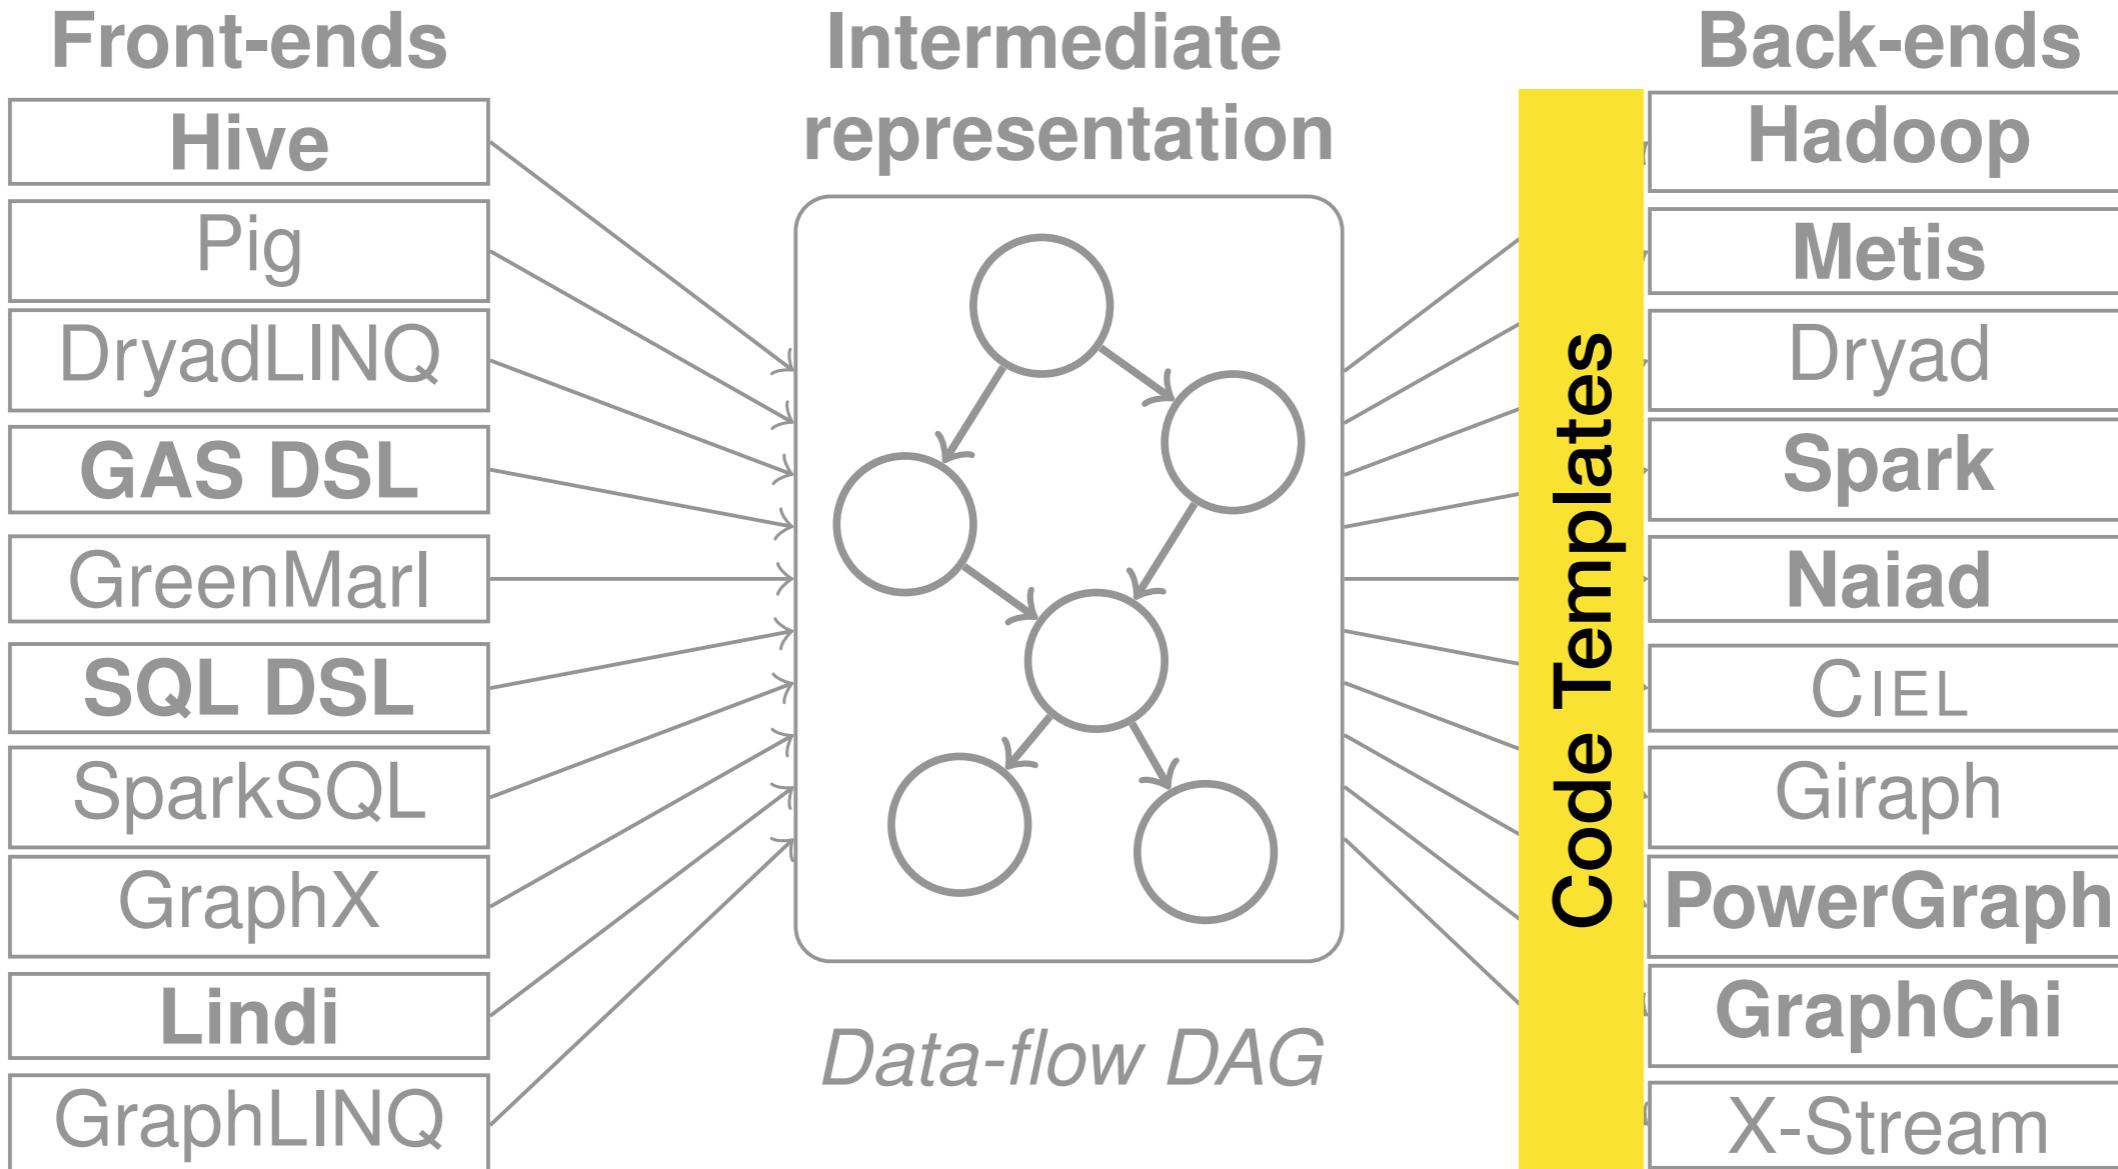
# Musketeer

***Code templates*** describing a platform



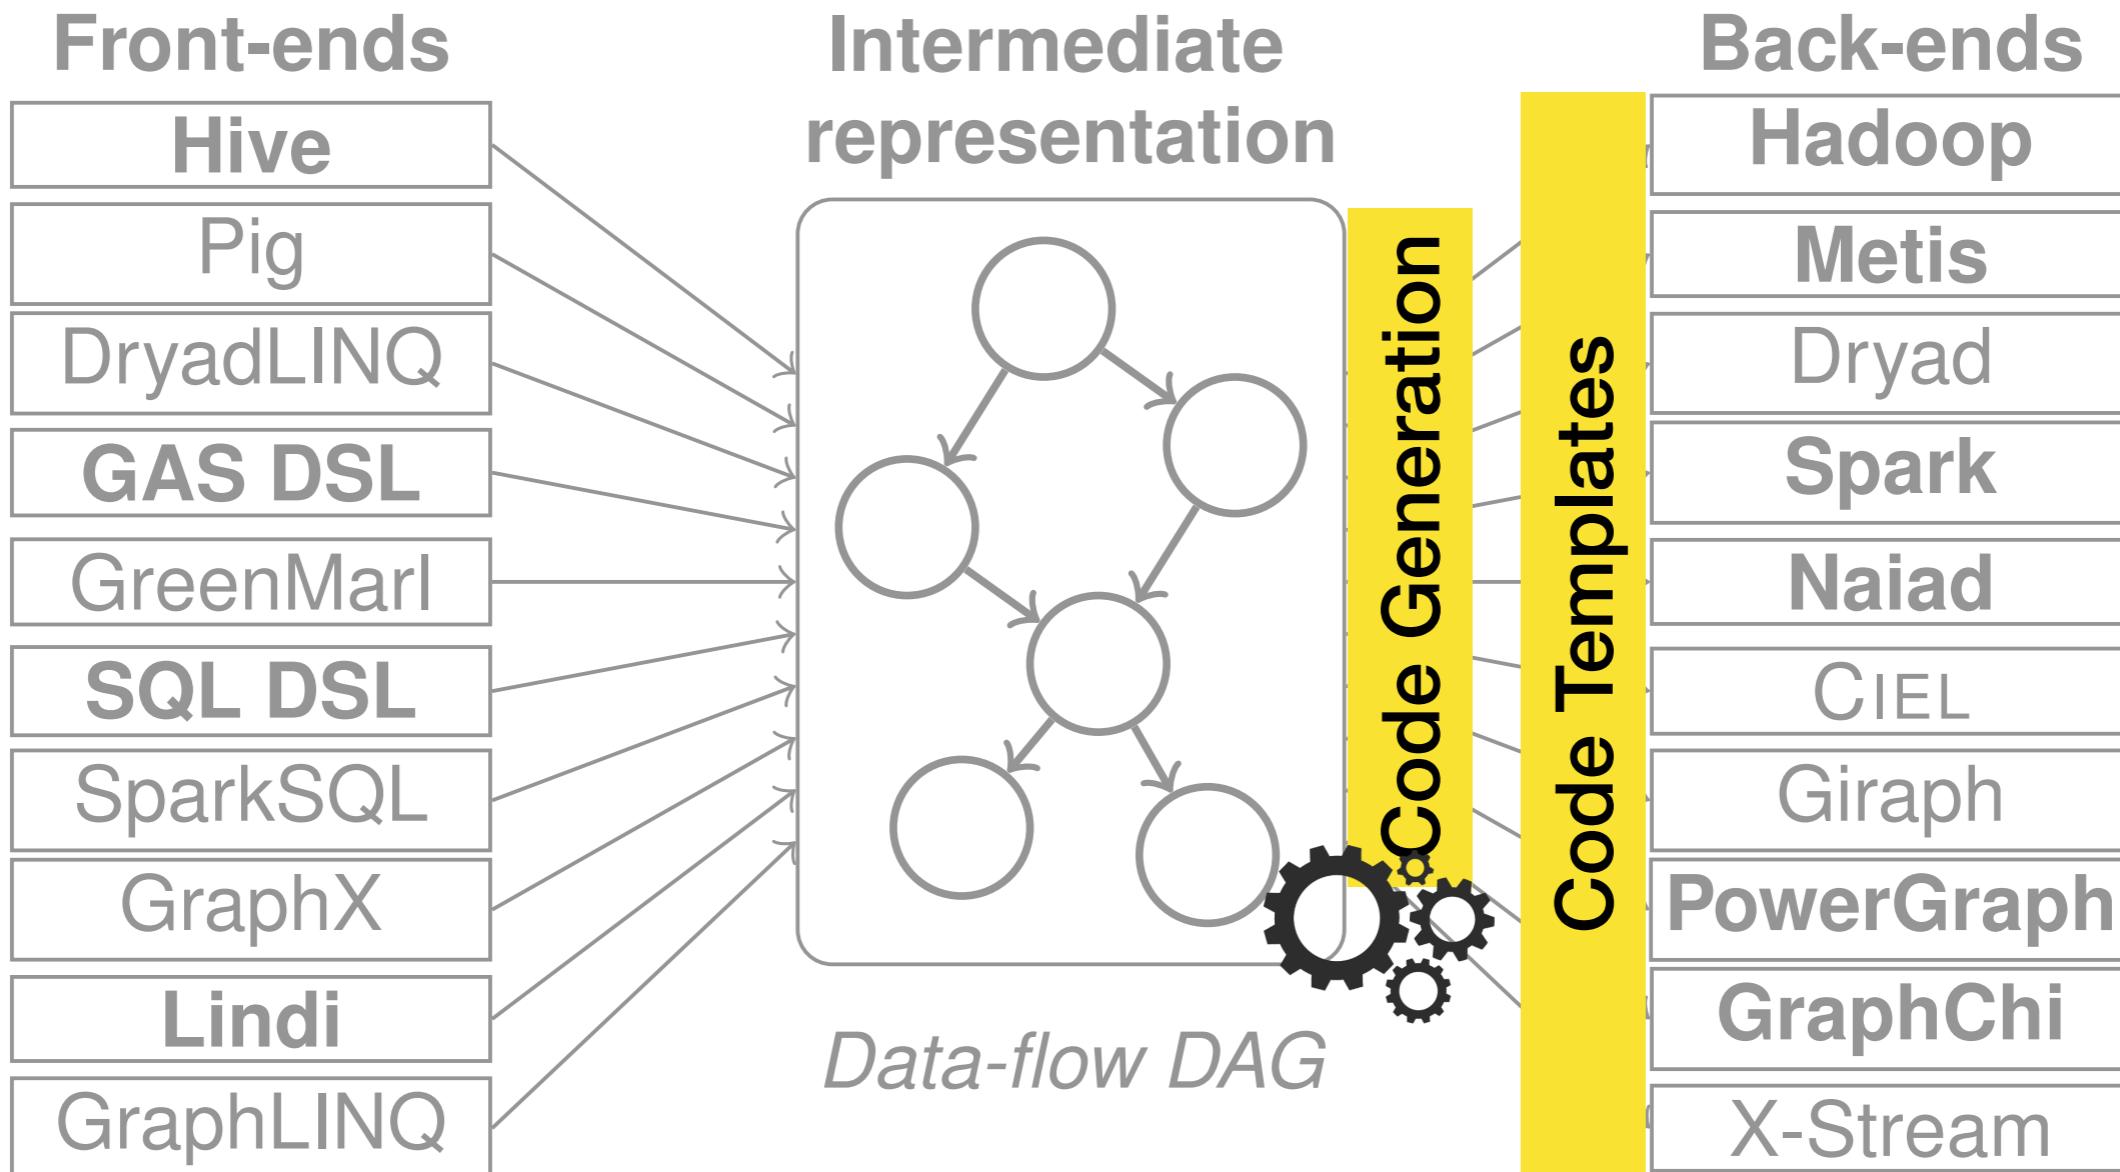
# Musketeer

***Code templates*** describing a platform



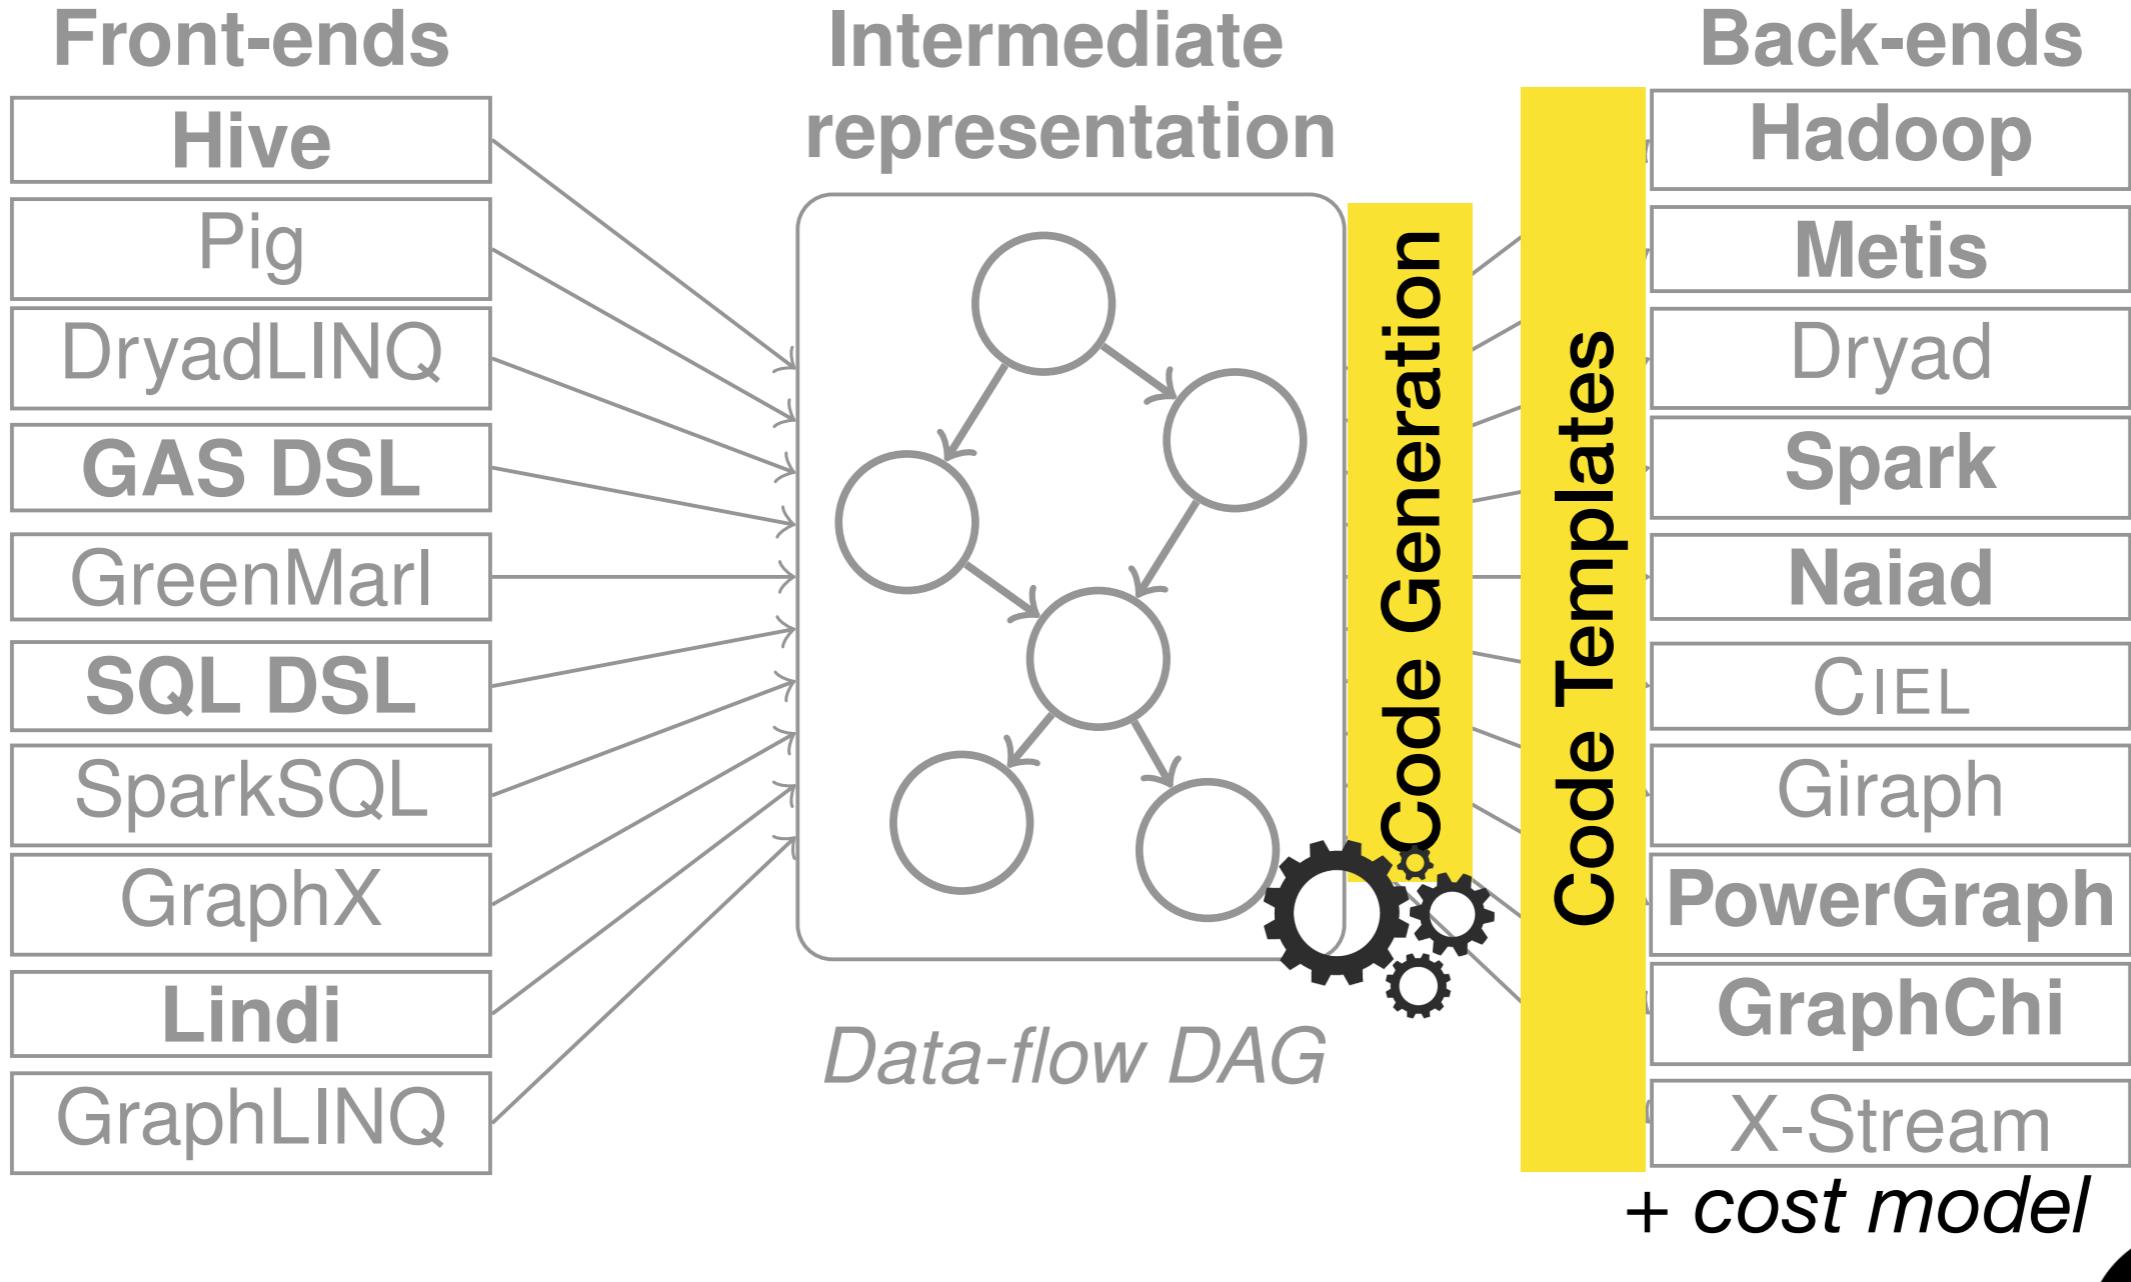
# Musketeer

***Code templates*** describing a platform



# Musketeer

***Code templates*** describing a platform



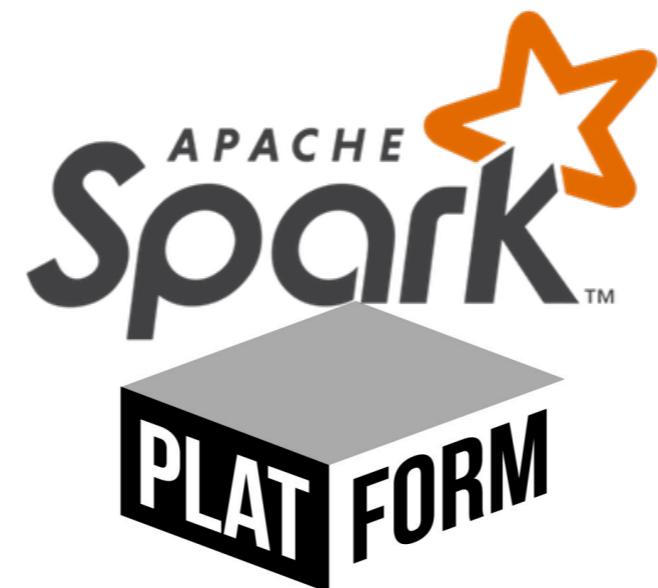
# Rheem

*Execution operators* describing a platform

Rheem Operators



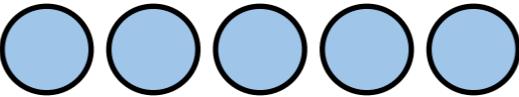
Rheem



# Rheem

*Execution operators* describing a platform

Rheem Operators



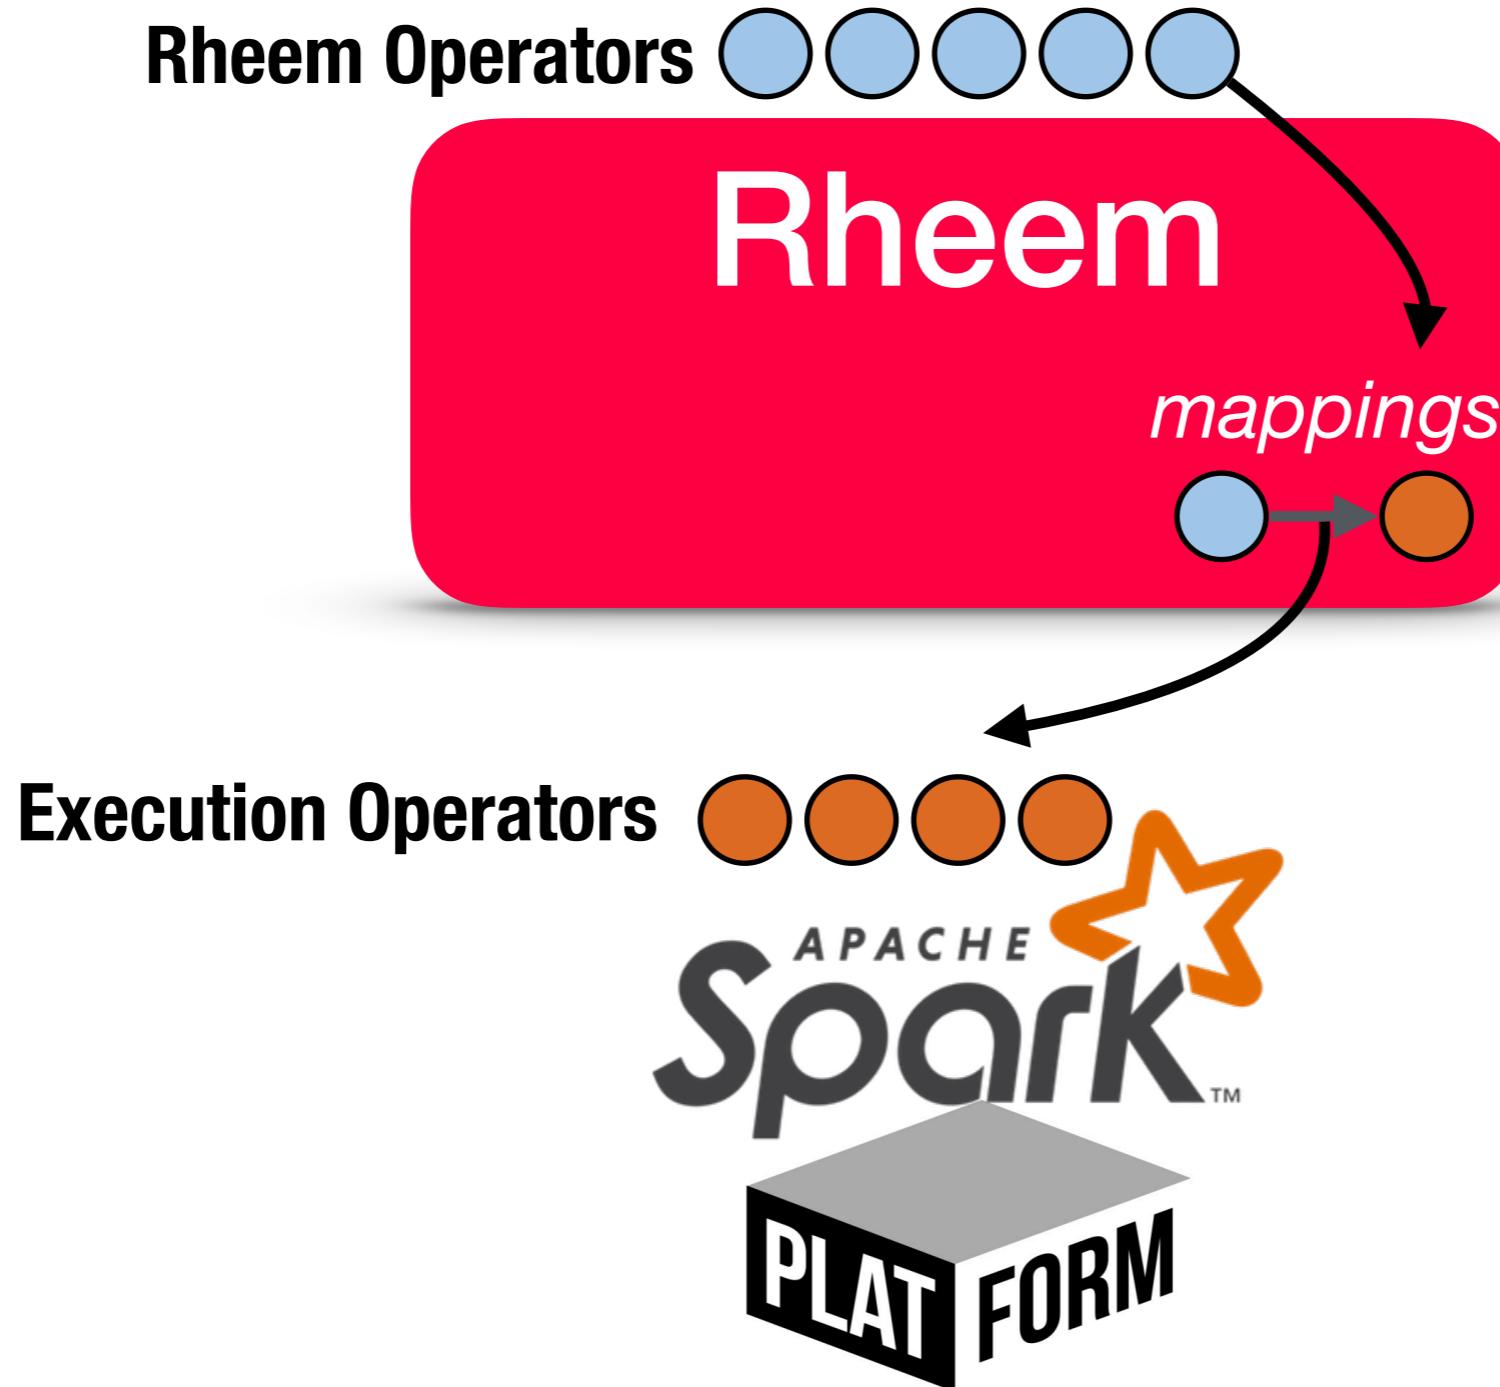
Rheem

Execution Operators



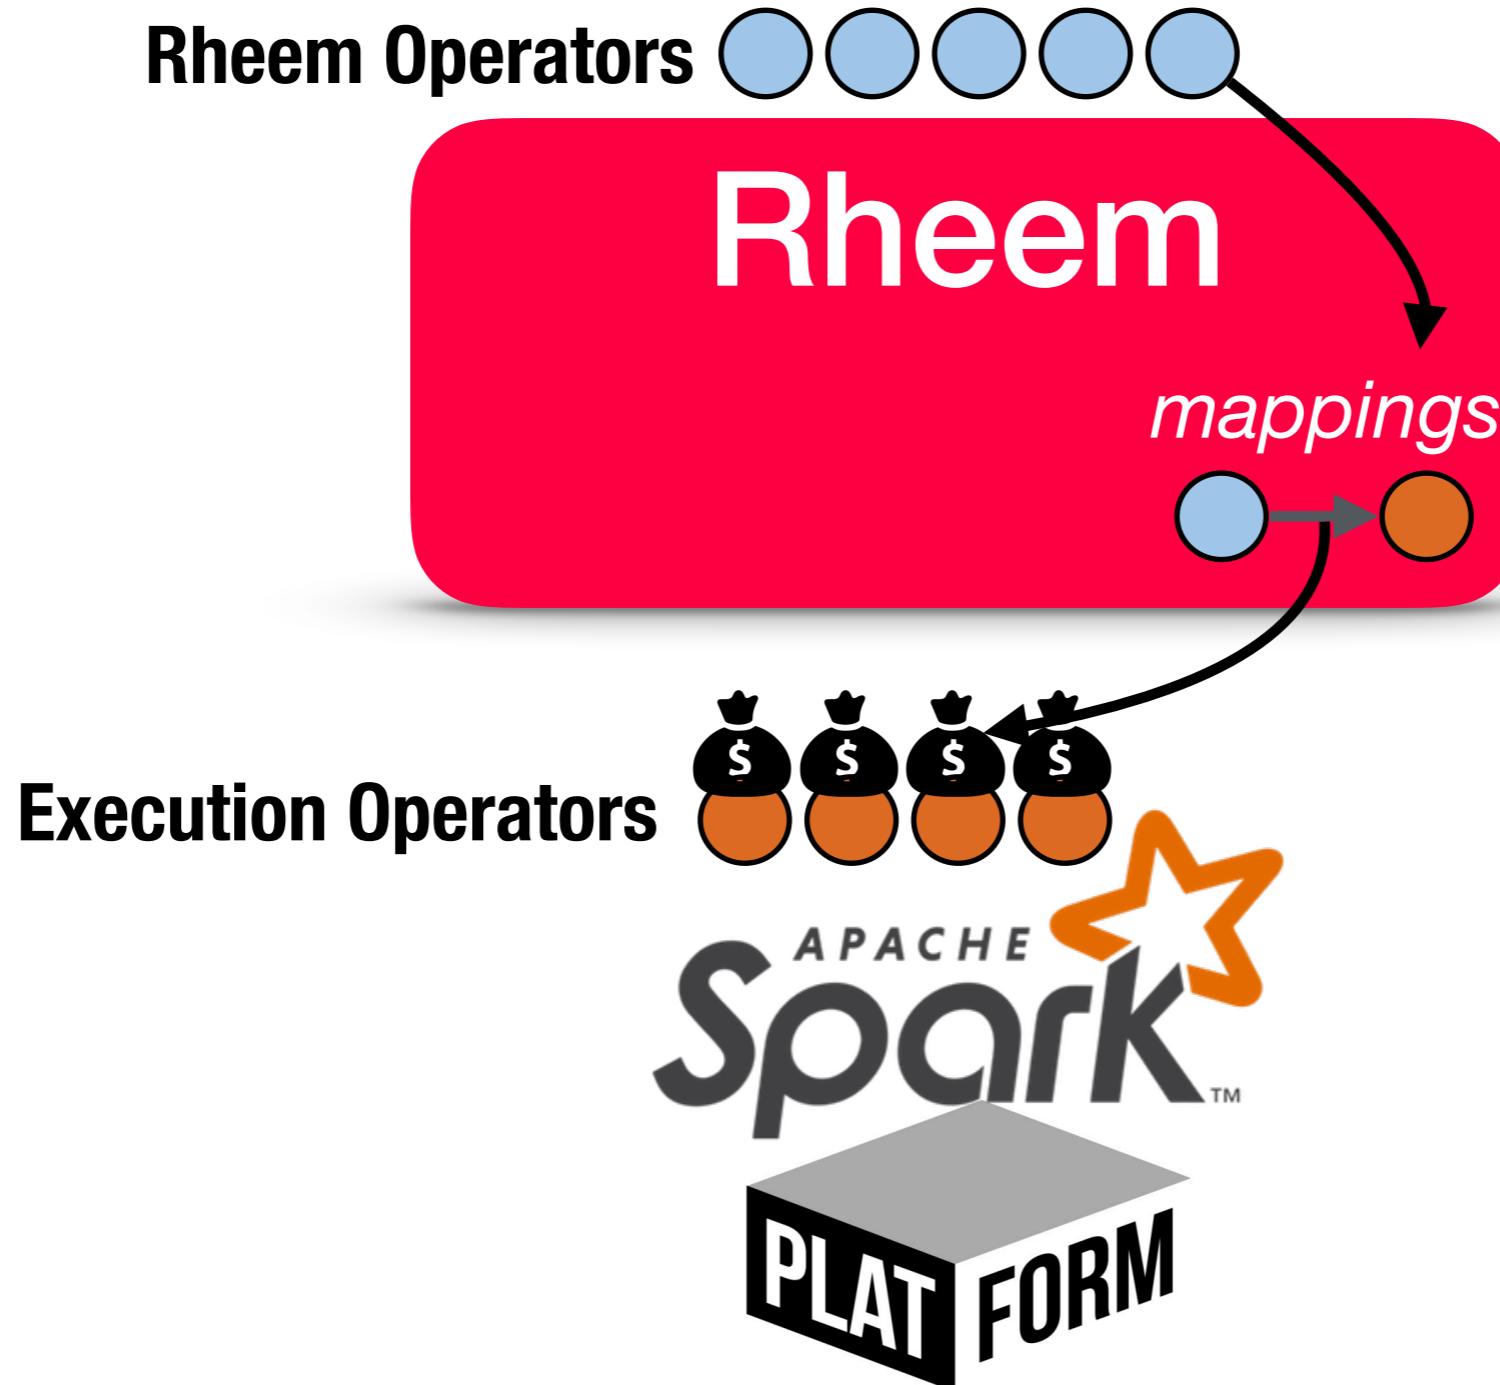
# Rheem

*Execution operators* describing a platform



# Rheem

*Execution operators* describing a platform



# Myria

**AST** describing a platform



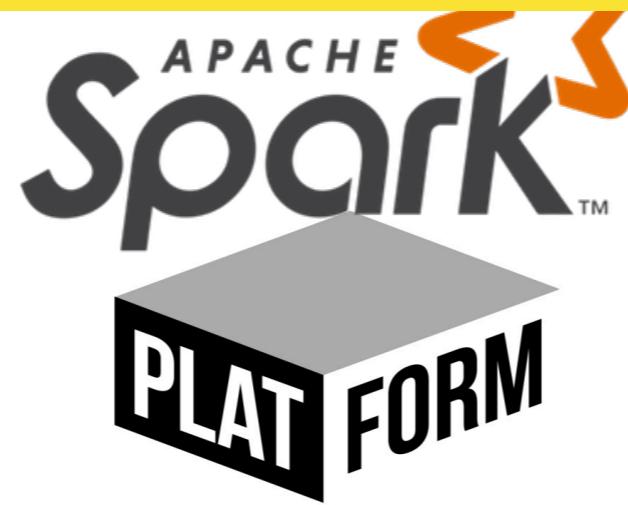
# Myria

**AST** describing a platform

MyriaL

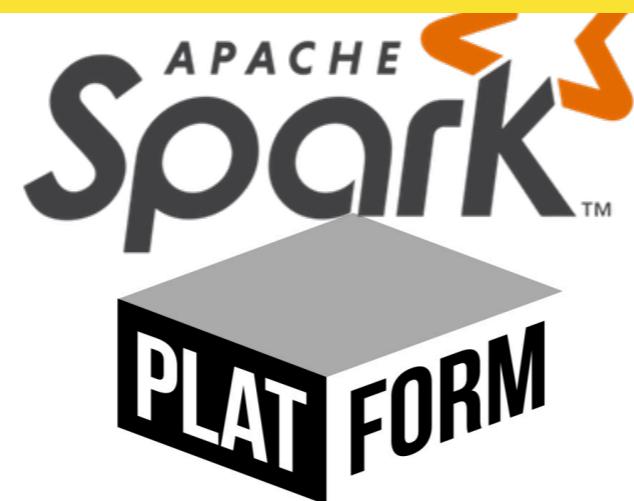
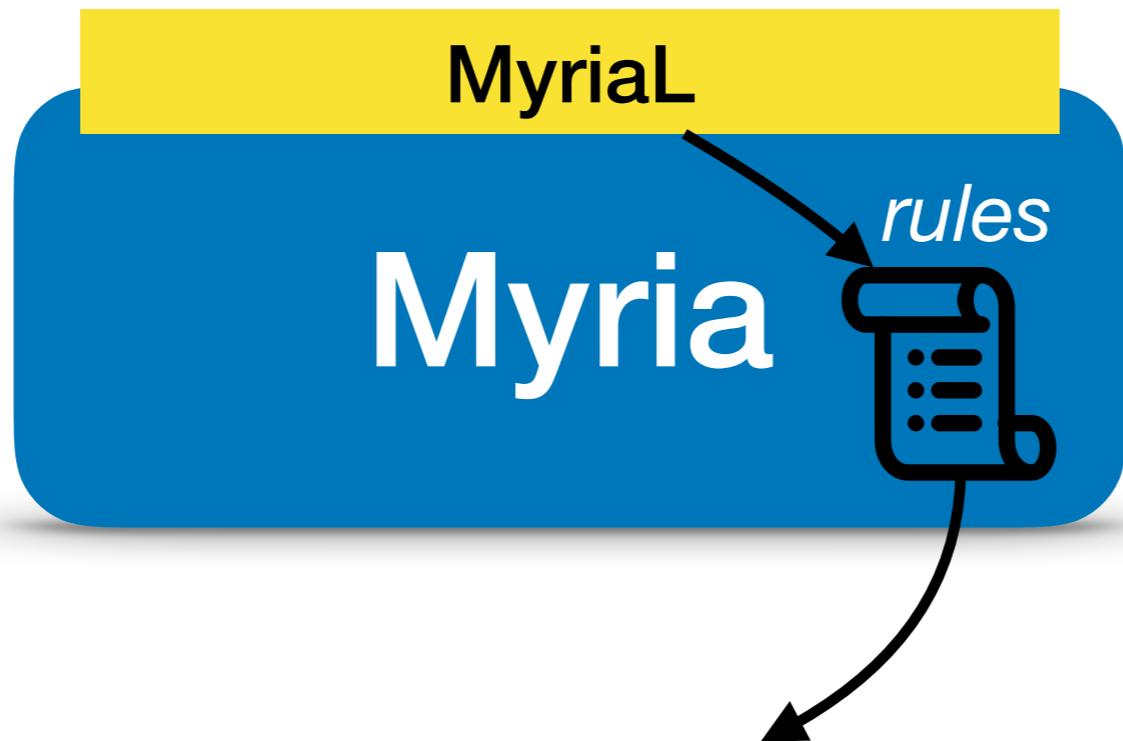
Myria

Abstract Syntax Tree (AST)



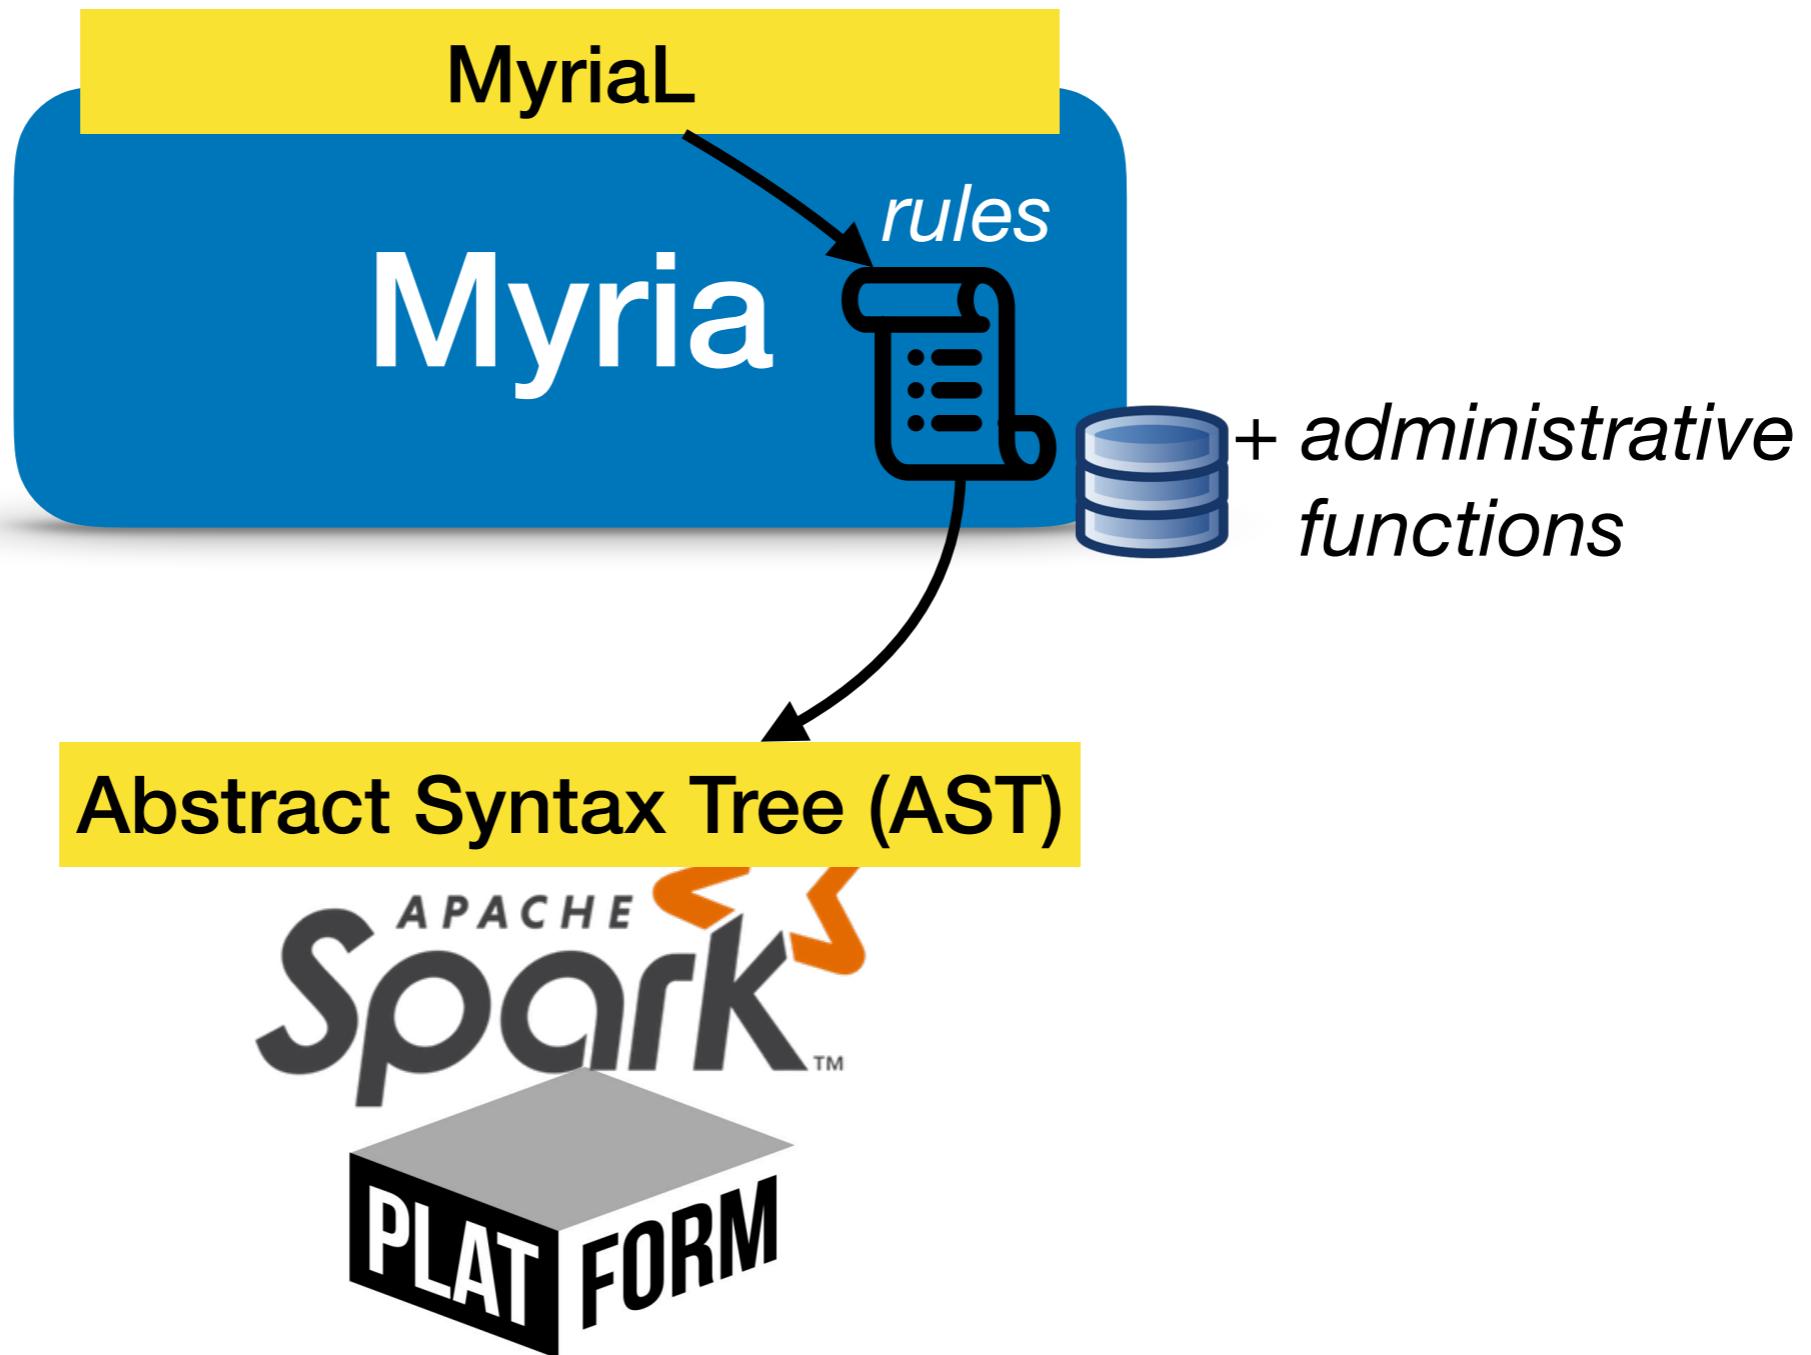
# Myria

**AST** describing a platform



# Myria

**AST** describing a platform

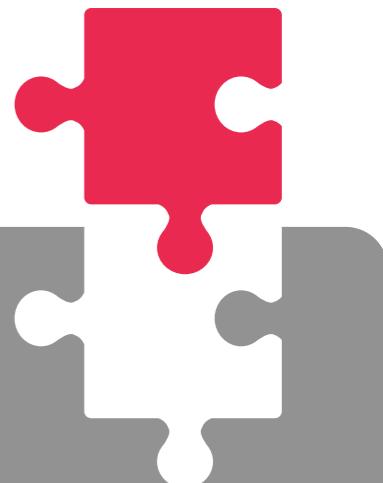


# Extending the Processing Model

1 Decoupling Applications

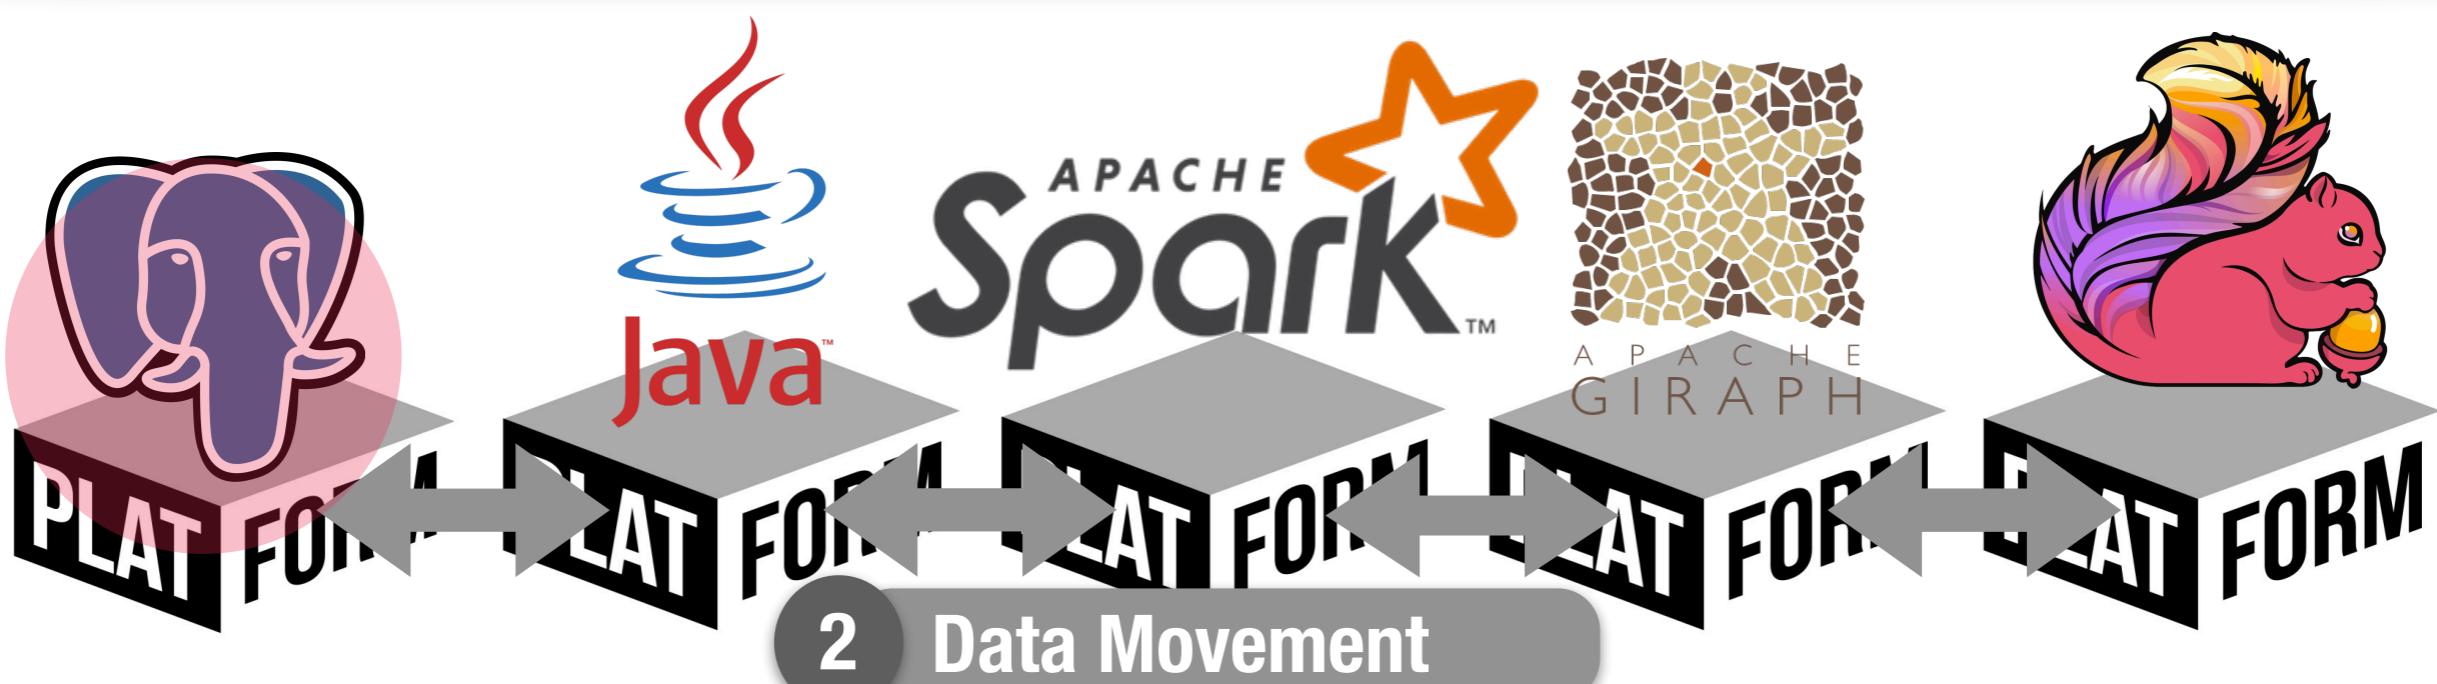


4 Extensibility



3

Automatic Cross-Platform  
Data Processing

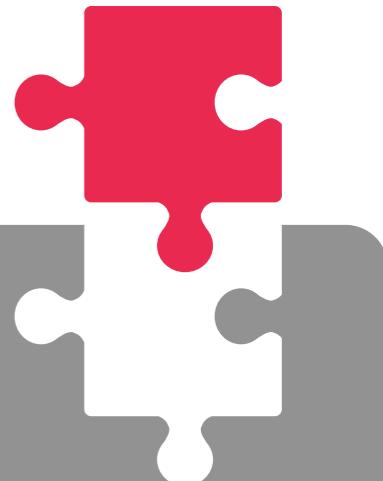


# Extending the Processing Model

1 Decoupling Applications

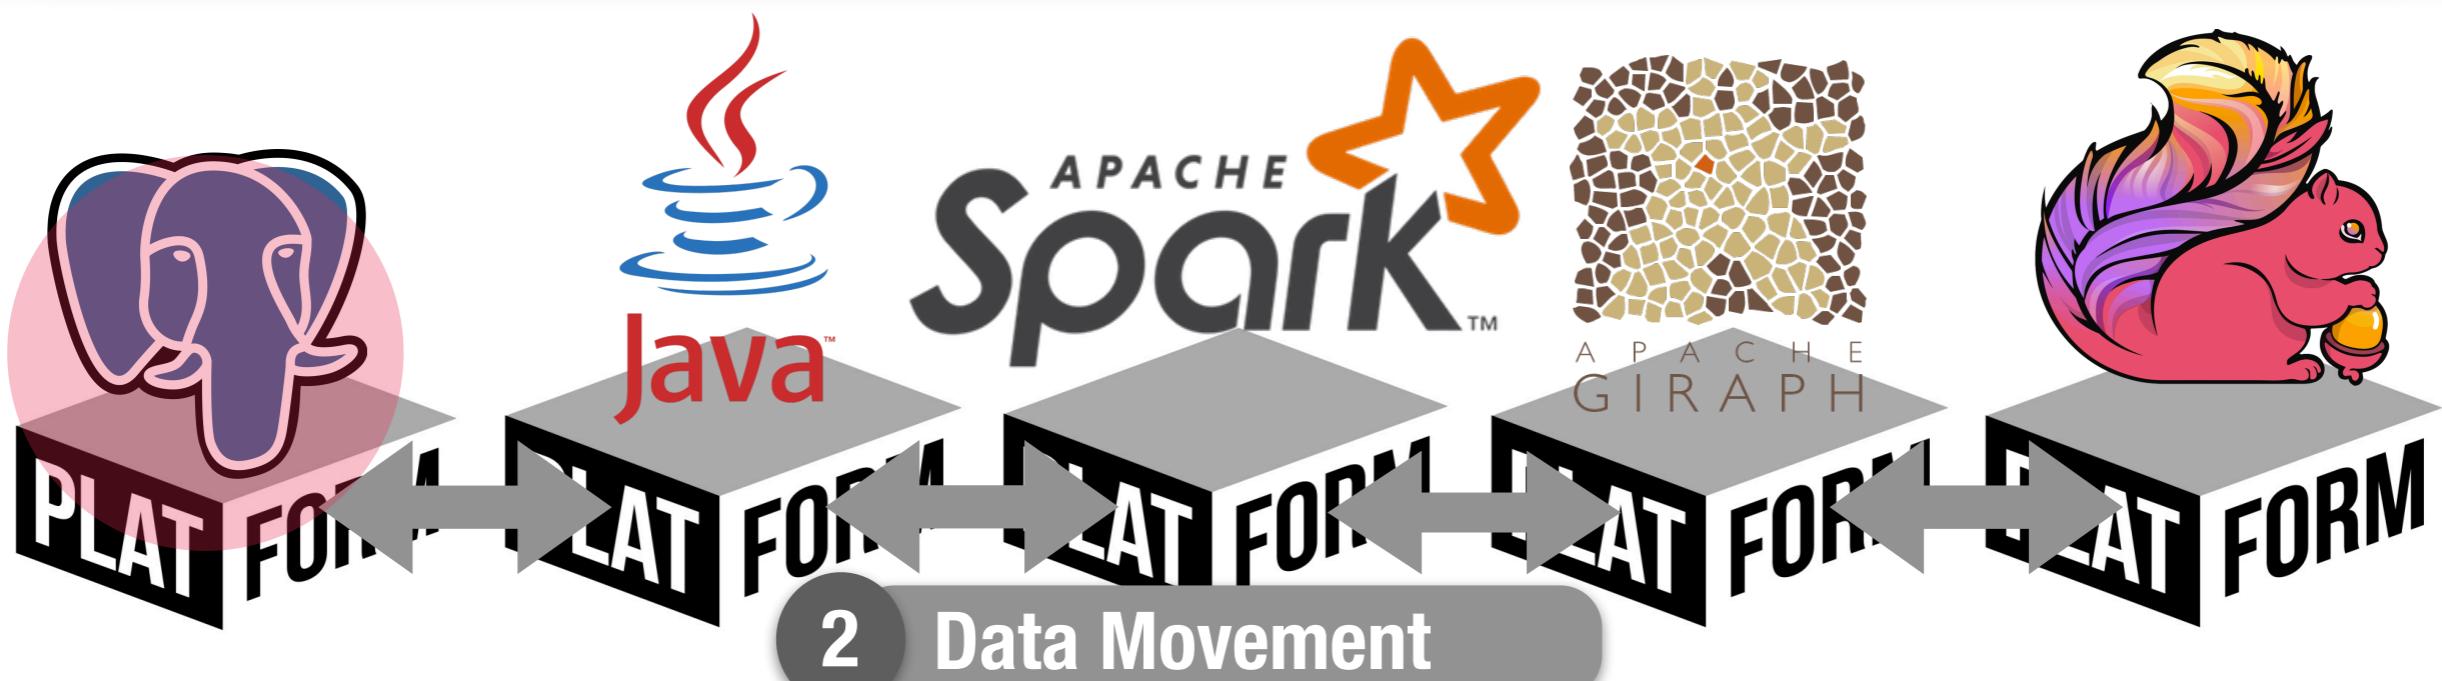


4 Extensibility

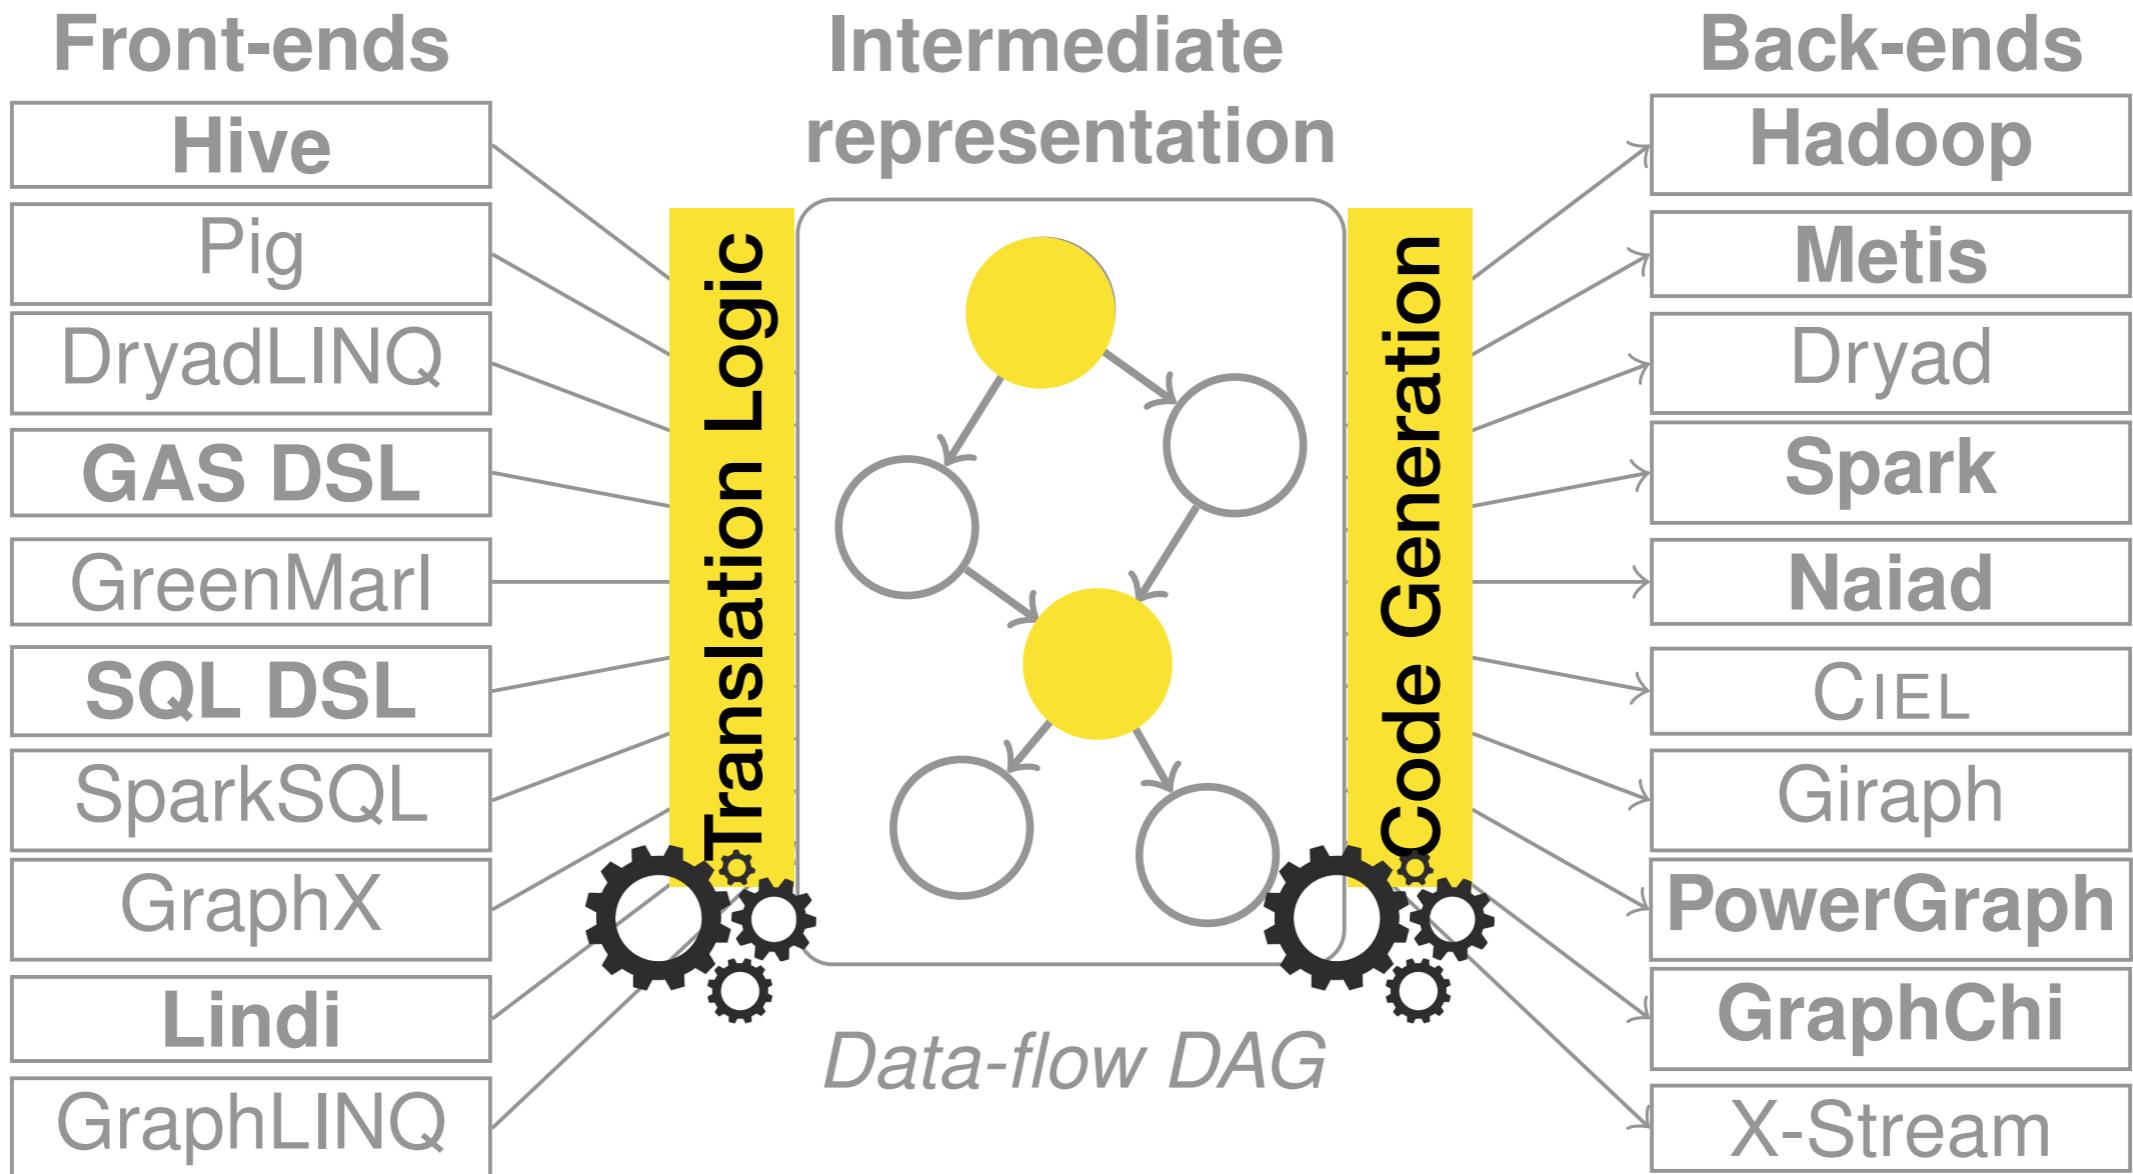


3

Automatic Cross-Platform  
Data Processing



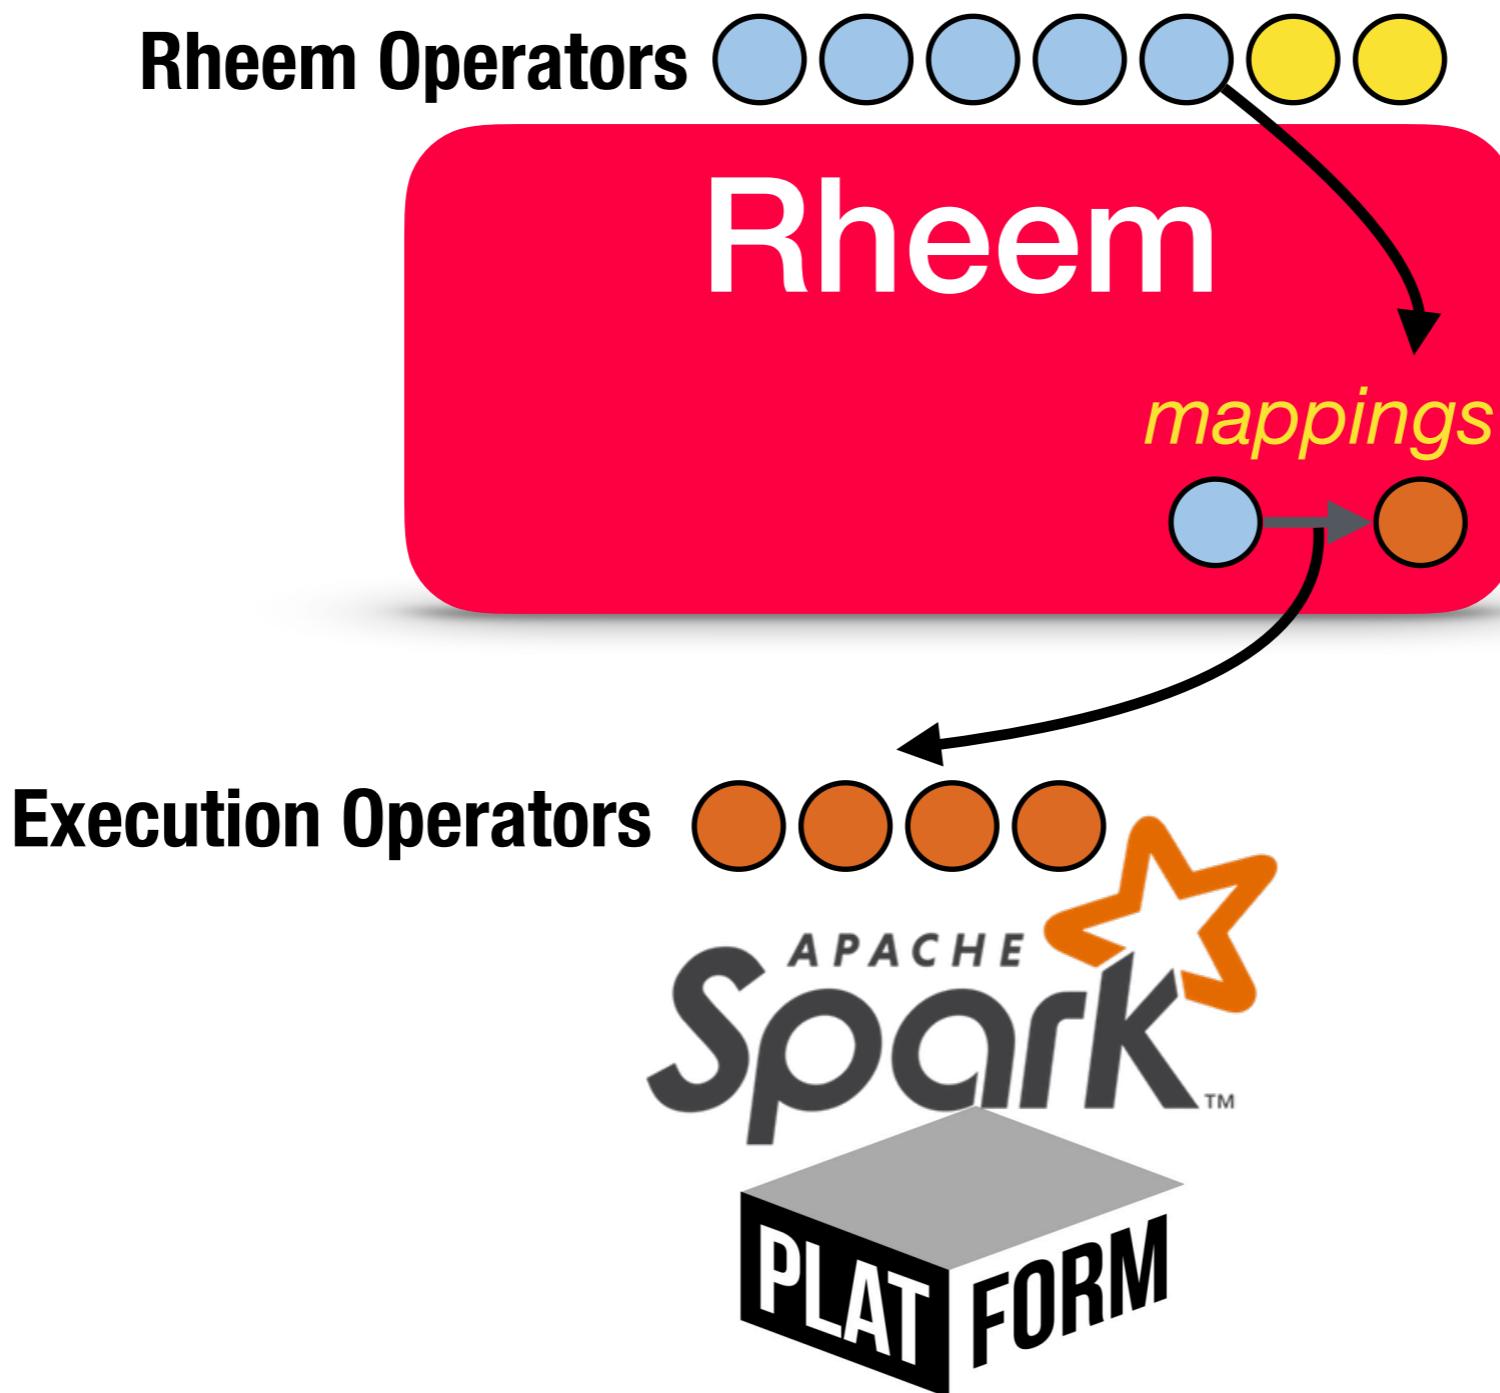
# Musketeer



# Extensibility Features

	Platform Description	IR Extensibility	Difficulty
Musketeer	Code Templates	Internal Code Changes	High
Rheem	Execution Operators		
Myria	AST		

# Rheem



# Extensibility Features

	Platform Description	IR Extensibility	Difficulty
Musketeer	Code Templates	Internal Code Changes	High
Rheem	Execution Operators	Rheem Operators	Medium
Myria	AST		

# Extensibility Features

	Platform Description	IR Extensibility	Difficulty
Musketeer	Code Templates	Internal Code Changes	High
Rheem	Execution Operators	Rheem Operators	Medium
Myria	AST	—	Medium
BigDawg	Language Translators	—	High
Cyclops	?	—	High
QoX	Internal Code Changes	—	High

# Cross-Platform Data Processing

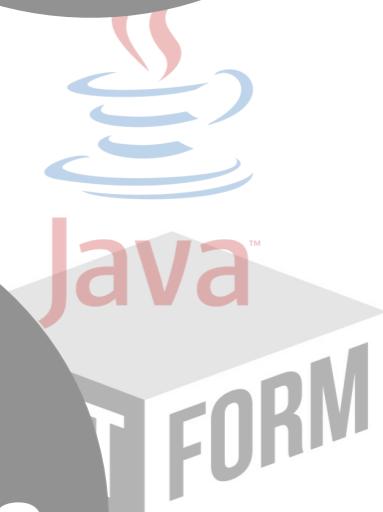
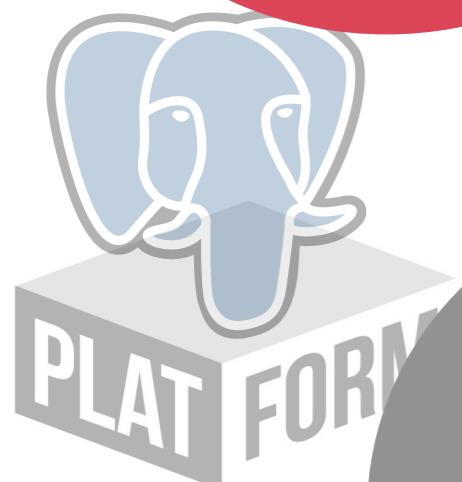
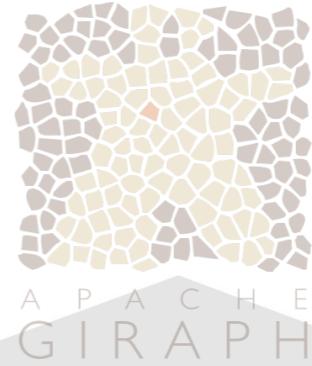
Summary

Motivation

Use Cases

Current  
Efforts

Challenges



# Recap

	Processing Model	CP Optimizer	Data Movement Cost	Extensibility
Musketeer	Dataflow DAG + Loop	Yes	No	High
Rheem	Dataflow DAG + Loop	Yes	Yes	Medium
Ires	Workflow graph	Yes	Yes	Medium
BigDawg	Query Plan Tree	No	NA	High
Myria	Extended relational	Yes	No	High
QoX	Dataflow graph	Yes	Yes	High
Cyclops	Windowed aggregation	Yes	Yes	High
Apache Beam	Dataflow DAG	No	NA	High
CloudMdsQL	Query Plan Tree	No	NA	High

# Recap

	Processing Model	CP Optimizer	Data Movement Cost	Extensibility
Musketeer	Dataflow DAG + Loop	Yes	No	High
Rheem	Dataflow DAG + Loop	Yes	Yes	Medium
<b>Fixed Processing Model</b>			Yes	Medium
BigDawg	Query Plan Tree	No	NA	High
Myria	Extended relational	Yes	No	High
QoX	Dataflow graph	Yes	Yes	High
Cyclops	Windowed aggregation	Yes	Yes	High
Apache Beam	Dataflow DAG	No	NA	High
CloudMdsQL	Query Plan Tree	No	NA	High

# Recap

	Processing Model	CP Optimizer	Data Movement Cost	Extensibility
Musketeer	Dataflow DAG + Loop	Yes	No	High
Rheem	Dataflow DAG + Loop	Yes	Yes	Medium
<b>Fixed Processing Model</b>			Yes	Medium
BigDawg	Query Plan Tree	No	NA	High
Myria	Extended	<b>Time-Consuming &amp; Hard Cost Tuning</b>		
QoX	Dataflow			High
Cyclops	Windowed aggregation	Yes	Yes	High
Apache Beam	Dataflow DAG	No	NA	High
CloudMdsQL	Query Plan Tree	No	NA	High

# Recap

	Processing Model	CP Optimizer	Data Movement Cost	Extensibility
Musketeer	Dataflow DAG + Loop	Yes	No	High
Rheem	Dataflow DAG + Loop	Yes	Yes	Medium
<b>Fixed Processing Model</b>				
BigDawg	Query Plan Tree	No	NA	High
Myria	Extended	<b>Time-Consuming &amp; Hard Cost Tuning</b>		
QoX	Dataflow	<b>Time-Consuming &amp; Hard Cost Tuning</b>		
Cyclops	Windowed aggregation	Yes	Yes	High
Apache Beam	Dataflow DAG	No	NA	High
CloudMdsQL	Query Plan Tree	No	NA	High

# Open Problems

## Flexible Processing Model

How to support arbitrary processing models?

# Open Problems

## Flexible Processing Model

How to support arbitrary processing models?

## Fully Automatic Cost Tuning

- How to match real workloads without knowing them?
- Weak control over the underlying platforms

# Open Problems

## Flexible Processing Model

How to support arbitrary processing models?

## Fully Automatic Cost Tuning

- How to match real workloads without knowing them?
- Weak control over the underlying platforms

## Platform Integration

How to integrate a platform fully automatically?

# Open Problems

## Flexible Processing Model

How to support arbitrary processing models?

## Fully Automatic Cost Tuning

- How to match real workloads without knowing them?
- Weak control over the underlying platforms

## Platform Integration

How to integrate a platform fully automatically?

## Fast Data Transfer

What is the right intermediate data representation?

# Open Problems

## Flexible Processing Model

How to support arbitrary processing models?

## Fully Automatic Cost Tuning

- How to match real workloads without knowing them?
- Weak control over the underlying platforms

## Platform Integration

How to integrate a platform fully automatically?

## Fast Data Transfer

What is the right intermediate data representation?

## Flexible High-Level Languages

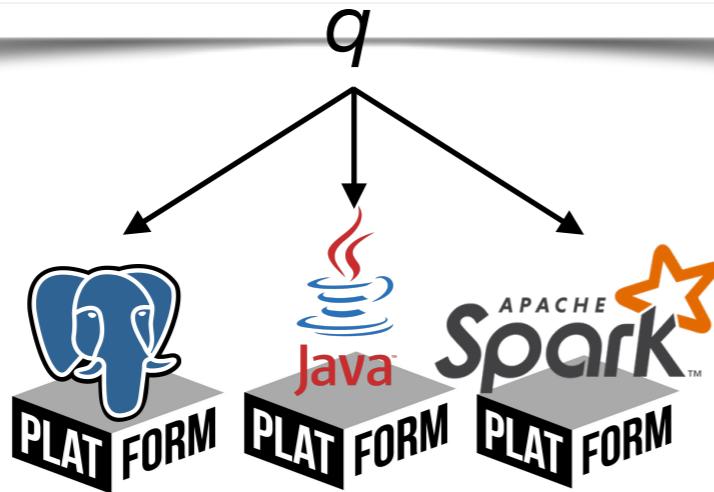
How to adapt to the needs of applications?

# Take Away

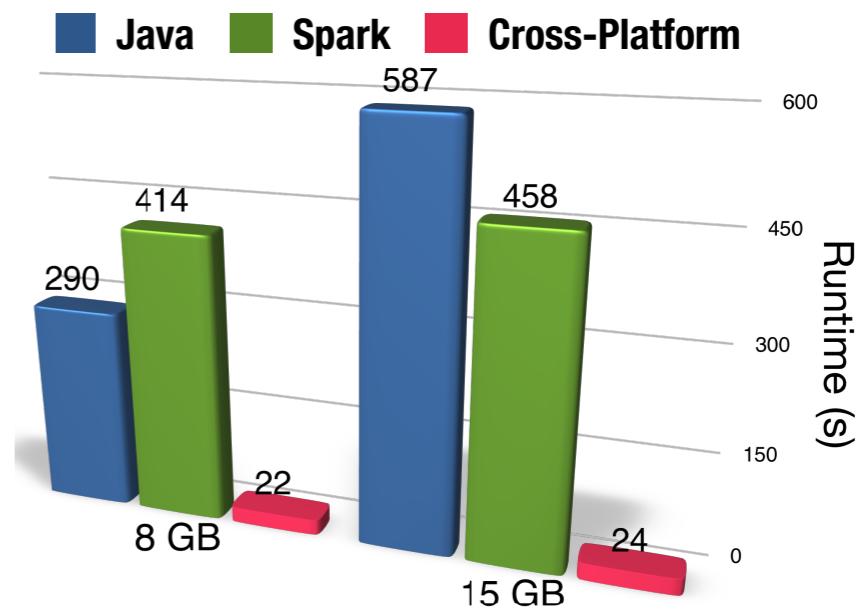
real need for *data processing independence*

	Monoplatform	Polyplatform
Monostore	Decoupled Single Platform	Opportunistic CP Mandatory CP
Polystore	Decoupled Single Platform	Polystore CP

— Queries Beyond a Single Platform —



SGD



Processor	Dataflow		MapReduce	
	Processing Model	CP Optimizer	Data Movement Cost	Extensibility
Musketeer	Dataflow DAG + Loop	Yes	No	High
Flume	Dataflow DAG +	..	..	Medium

No Perfect Current Solution

Processor	Time Consuming & Hard Cost Tuning			
	Myria	QoX	Cyclops	Apache Beam
Myria	Extended	..	..	..
QoX	Dataflow graph	Yes	Yes	Yes
Cyclops	Windowed aggregation	Yes	Yes	Yes
Apache Beam	Dataflow DAG	No	NA	High

## OPEN PROBLEMS

1. Fully Automatic Cost Tuning

2. Flexible Processing Model

3. Platform Integration

4. Fast Data Movement

5. Flexible High-Level Languages

# References

- I. Gog et al. *Musketeer: all for one, one for all in data processing systems*. In EuroSys, 2015.
- D. Agrawal et al. *Rheem: Enabling Multi-Platform Task Execution*. In SIGMOD, 2016.
- D. Agrawal et al. *Road to Freedom in Big Data Analytics*. In EDBT, 2016.
- K. Doka et al. *Mix 'n' match multi-engine analytics*. In IEEE BigData, 2016.
- J. Wang et al. *The Myria Big Data Management and Analytics System and Cloud Services*. In CIDR, 2017.
- A. Simitsis et al. *Optimizing Analytic Data Flows for Multiple Execution Engines*. In SIGMOD, 2012.
- V. Gadepally et al. The BigDawg Polystore System and Architecture. IEEE High Performance Extreme Computing 2016.
- A. J. Elmore et al. *A Demonstration of the BigDAWG Polystore System*. PVLDB, 8(12):1908–1911, 2015.
- H. Lim et al. *How to Fit when No One Size Fits*. In CIDR, 2013.
- B. Haynes et al. *PipeGen: Data Pipe Generator for Hybrid Analytics*. In SoCC, 2016.
- S. Palkar et al. *Weld: A Common Runtime for High Performance Data Analysis*. In CIDR, 2017.
- M. Boehm et al. *SystemML: Declarative Machine Learning on Spark*. PVLDB, 9(13):1425–1436, 2016.
- B. Kolev et al. *CloudMdsQL: querying heterogeneous cloud data stores with a common language*. Journal of Distributed and Parallel Databases, 2016.
- C. Chambers et al. *FlumeJava: Easy, Efficient Data-Parallel Pipelines*. SIGPlan 2010.