



# Cloud-based RDF data management

**Zoi Kaoudi**

IMIS, Athena Research Center

**Ioana Manolescu**

INRIA Saclay, France

**SIGMOD 2014 Tutorial**

# History and acknowledgments

First presented at ICDE 2013 (short version: 1.5 hours)

Long version previously presented at the Workshop on Open Data  
(WOD) 2013

Companion paper:



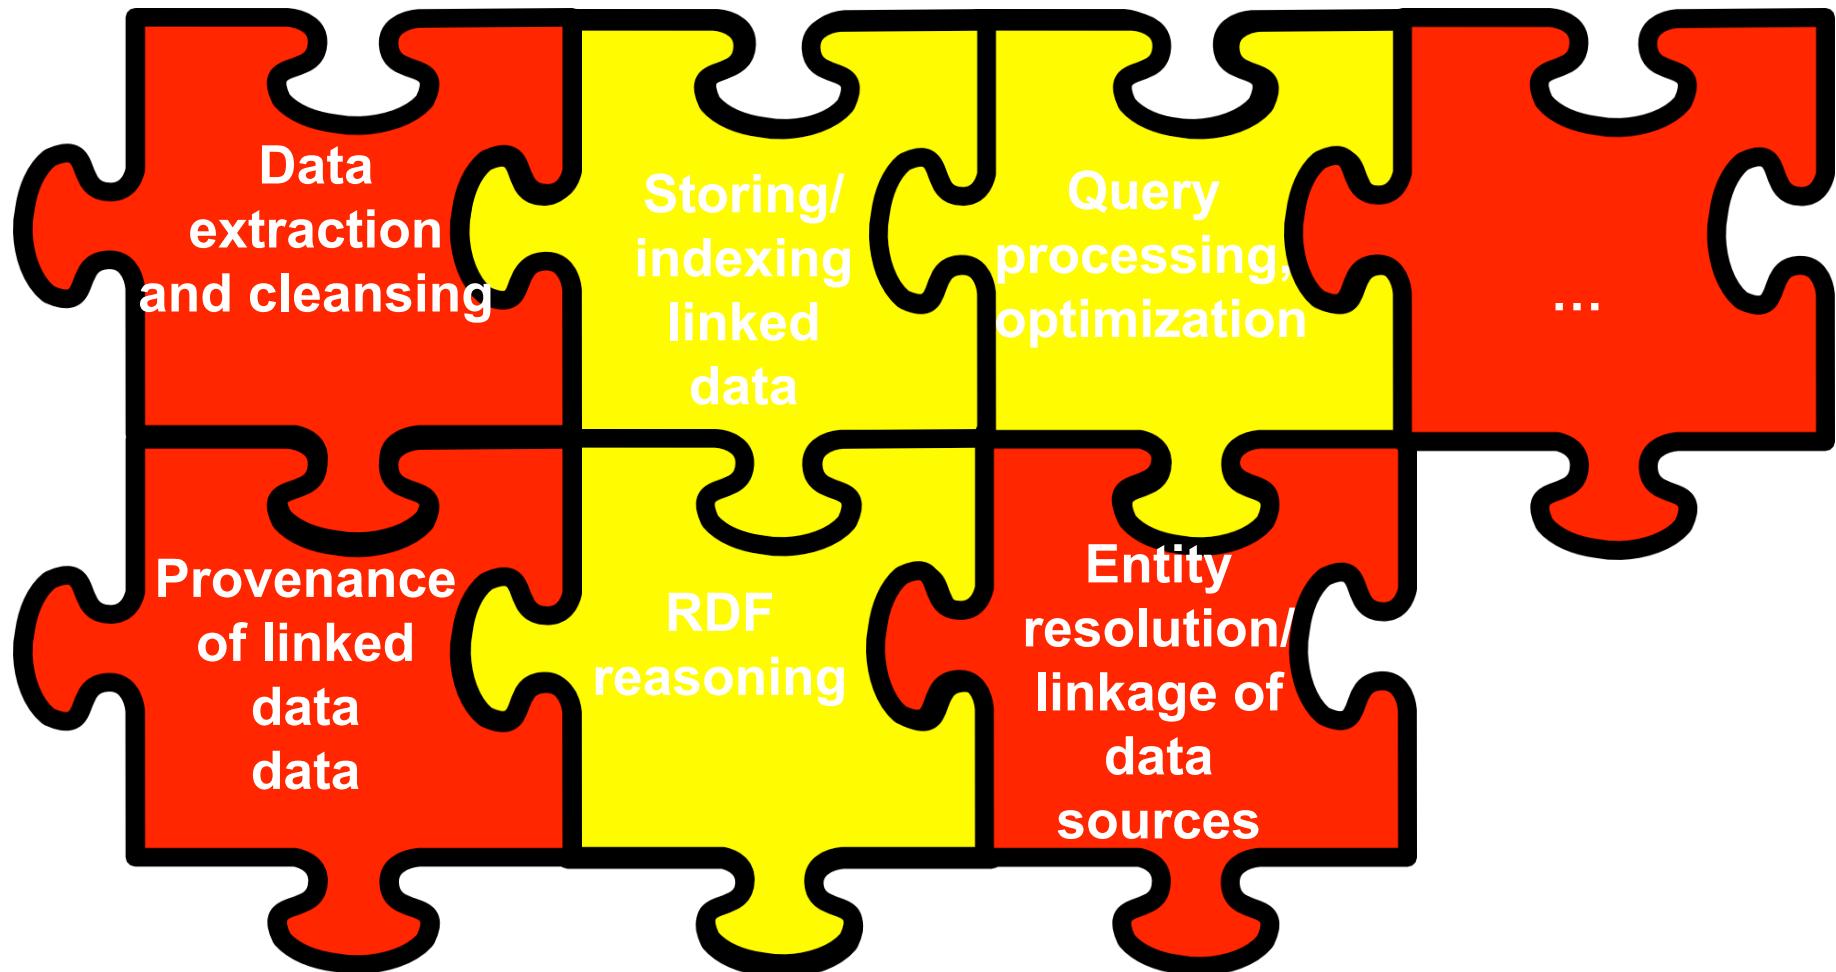
“RDF in the clouds: A Survey”, accepted for publication in the  
VLDB Journal (June 5)

# Linked [Open] Data

- Tim Berners-Lee in TED 2009 conference  
*“Linked Data is the Semantic Web done right”*
- Linked Data principles:
  - Use URIs to represent data
  - Use HTTP URIs (make data de-referenceable)
  - Give information for a piece of data using W3C standards
  - Include links to other URIs
- Linked Data Management → RDF data management

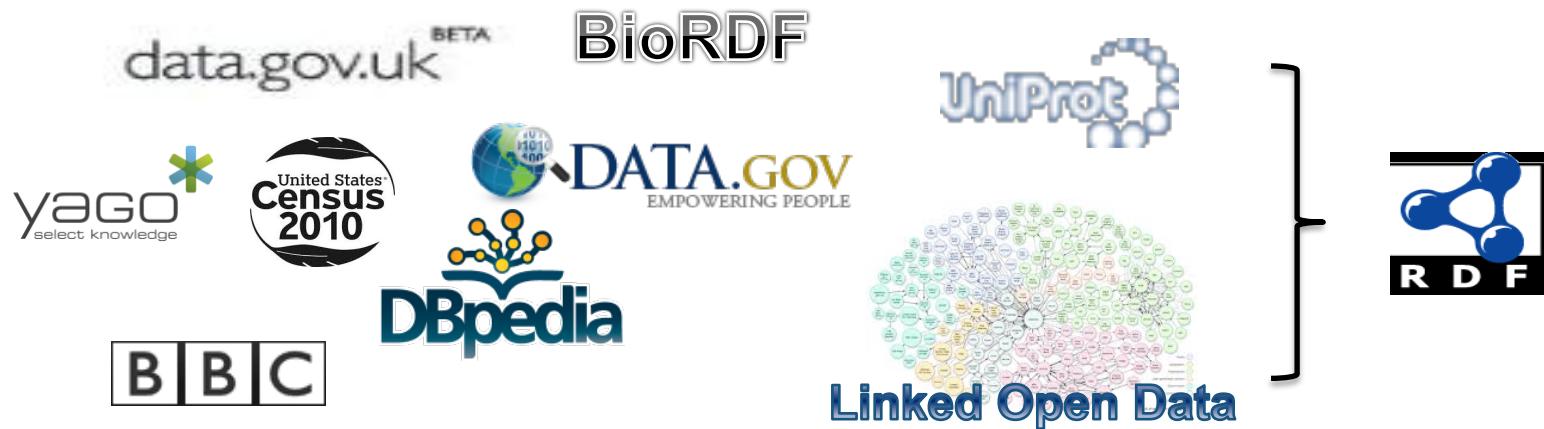


# Linked data management tasks



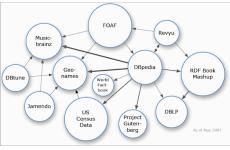
# Where is RDF found today?

- Many available open data sources in RDF

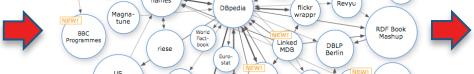


# Linked data is growing

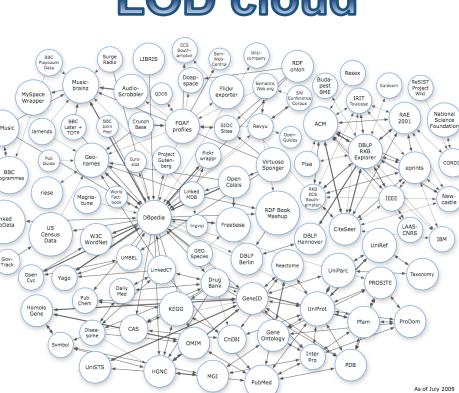
## LOD cloud



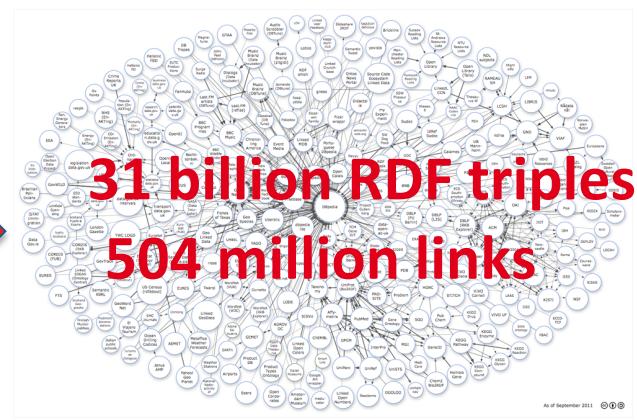
2007



2008



2009



2011



# RDF data sources in numbers

LOD cloud



BioRDF



Linked Open Data cloud

**31 billion** triples, Sept.

RDF-encoded Wikipedia

**1.89 billion** triples

RDF-encoded biological data

**2.7 billion** triples

US government data in RDF

**5 billion** triples

crawled Web data

**2 billion** triples

US population statistics

**1 billion** triples

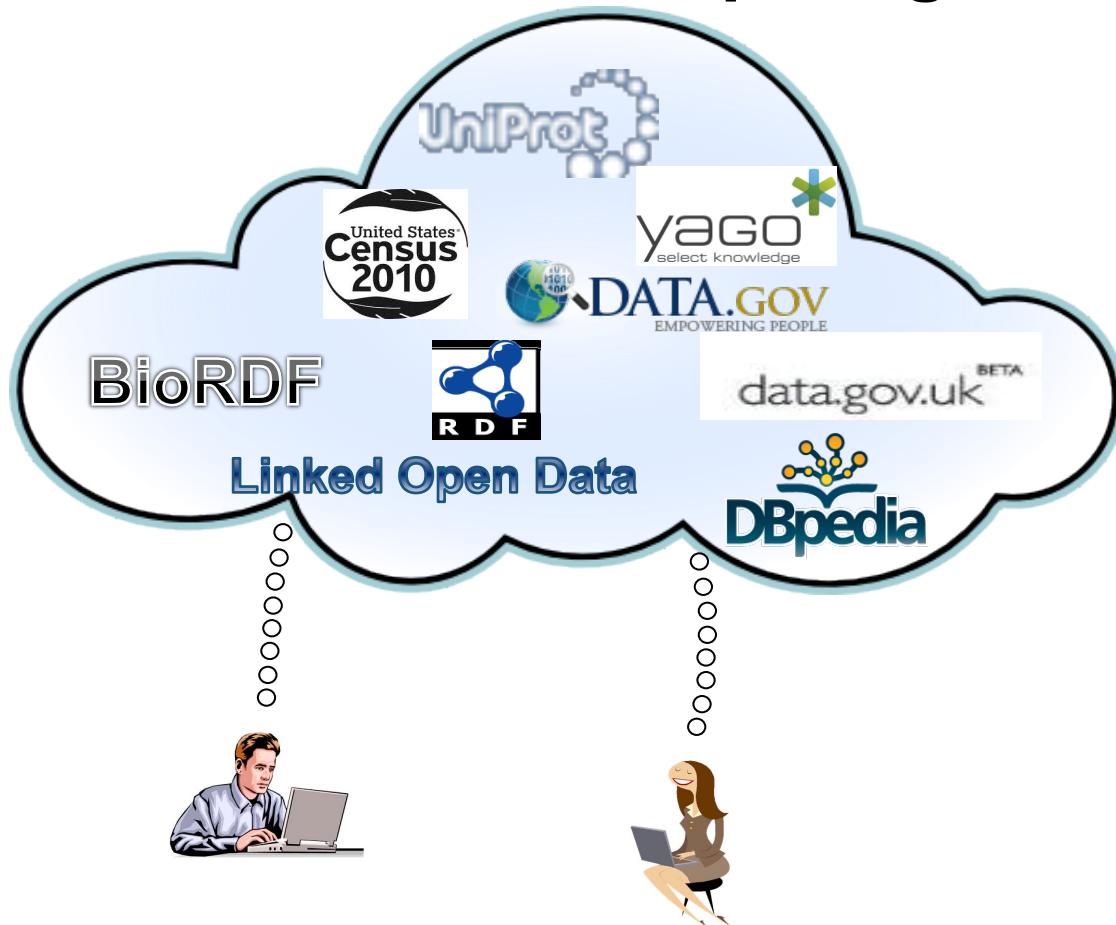
facts from Wikipedia, Wordnet,  
Geonames

**0.12 billion** triples

... and many more ... <http://www.w3.org/wiki/DataSetRDFDumps>  
<http://datahub.io/>

# Managing large volumes of RDF data

## Cloud computing



Scalability

Elasticity

Fault-tolerance

# What this seminar is about

## RDF data management in the cloud *(or: large-scale distributed repositories)*

- Storing RDF data in the cloud
- Evaluation of RDF queries in the cloud
- Inference of RDF data in the cloud



# Outline

1. Introduction to RDF
2. Recall of Cloud platforms
3. Analysis of RDF systems in the cloud
  - Storage
  - Querying
  - Inference
4. Open issues

# 1

## RDF BASICS

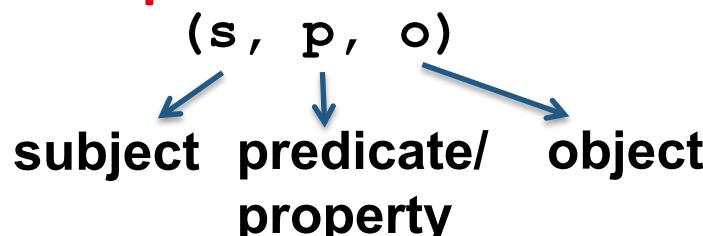
# Resource Description Framework (RDF)



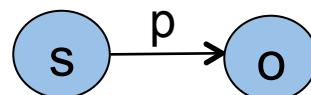
- RDF: representation of **resources** on the Web identified by Universal Resource Identifiers (**URIs**) U

[http://live.dbpedia.org/page/Pablo\\_Picasso](http://live.dbpedia.org/page/Pablo_Picasso)

- **Literals**: strings, numbers, dates, etc. L
- **Blank nodes**: existential identifiers B
- RDF data: facts = **triples**



- A *well-formed* triple belongs to  $(U+B) \times (U+B) \times (U+B+L)$
- RDF as graphs



# RDF Schema

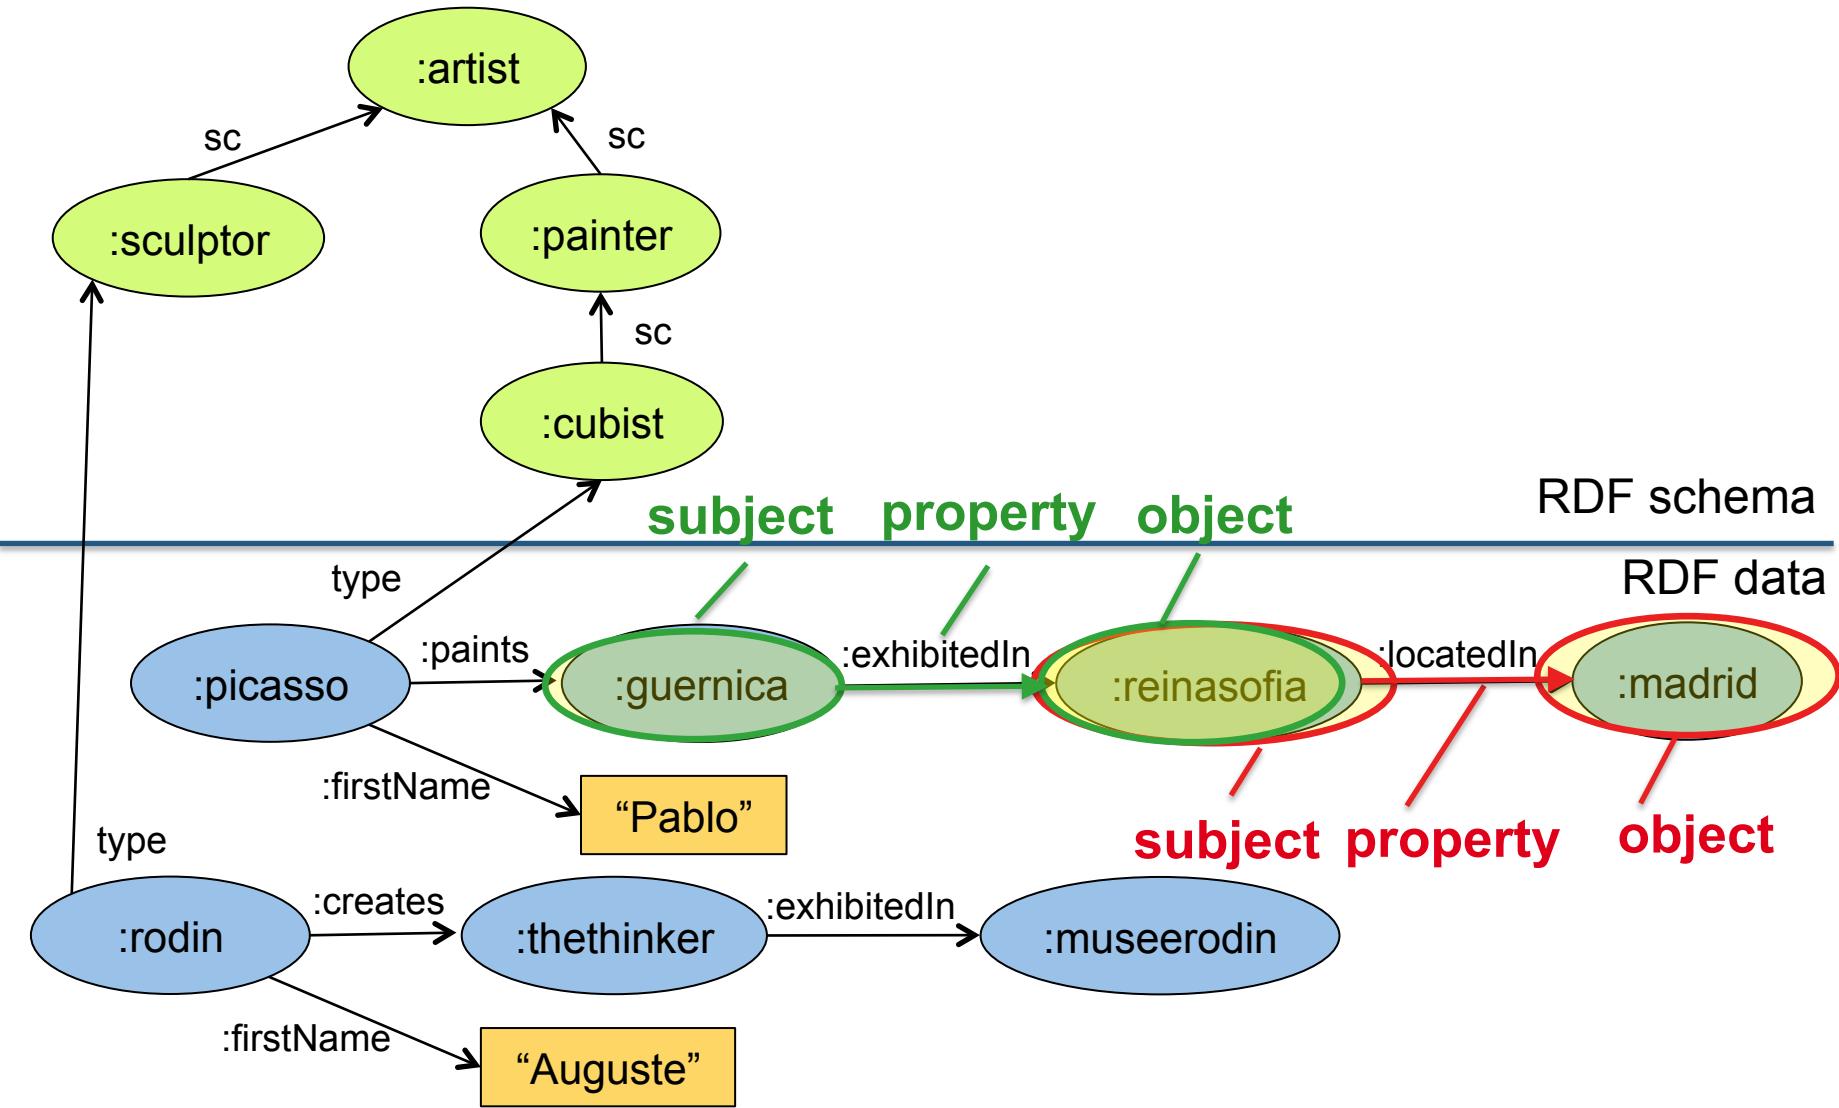


- RDF **schema** (RDFS): specifies **concepts**, **properties** and **relationships** between them
  - Concept/class: <http://live.dbpedia.org/ontology/Artist>
  - Property: <http://dbpedia.org/property/birthPlace>
  - Relationships:  
(dbpedia-owl:Artist, rdfs:subClassOf, dbpedia-owl:Person)  
(dbpprop:birthplace, rdfs:domain, dbpedia-owl:Person)  
(dbpedia:Pablo\_Picasso, rdf:type, dbpedia-owl:Artist)

We skip namespaces and use the following notation for rdf-specific terms:

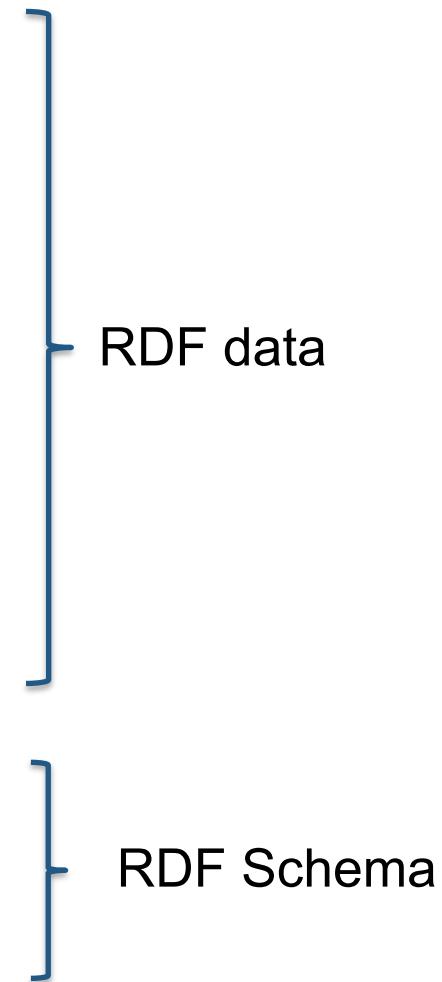
```
(:Artist, sc, :Person)
(:birthplace, domain, :Person
(:Pablo_Picasso, type, :Artist)
```

# Running example: RDF database



## Running example: RDF database (triple syntax)

```
:picasso type :cubist .  
:picasso :firstName "Pablo"  
:picasso :paints :guernica .  
:guernica :exhibitedIn :reinasofia .  
:reinasofia :locatedIn :madrid .  
:rodin type :sculptor .  
:rodin :firstname "Auguste" .  
:rodin :creates :thethinker .  
:thethinker :exhibitedIn :museerodin .  
  
:sculptor sc :artist .  
:painter sc :artist .  
:cubist sc :painter.
```



# SPARQL Query Language



- Basic building block
  - **Triple pattern**  $(s, p, o)$  from  $UBV \times UBV \times UBLV$
- SQL-based syntax

- Basic graph pattern (**BGP**) queries:

SELECT  $?x_1, ?x_2, \dots, ?x_n$  —> **distinguished variables**

WHERE {

$s_1 p_1 o_1 .$

$\dots$

$s_k p_k o_k .$

}

}

**Conjunction of triple patterns**

**U:** URIs  
**B:** blank nodes  
**L:** literals  
**V:** variables

# SPARQL Query Language



- BGP queries in conjunctive form:

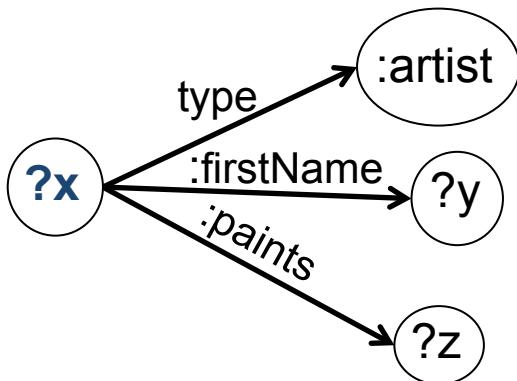
–  $\underbrace{?x_1, ?x_2, \dots, ?x_n}_{\text{distinguished variables}} : \underbrace{(s_1, p_1, o_1) \wedge \dots \wedge (s_k, p_k, o_k)}_{\text{triple patterns}}$

- BGP queries: Select-Project-Join fragment of relational algebra
  - **Select** →  $?x: (?x, :paints, ?y)$
  - **Project** →  $?x: (?x, :paints, ?y)$
  - **Join** →  $?x: (?x, :paints, ?y) \wedge (?y, :exhibited, :reinasofia)$
- More features in SPARQL 1.1: W3C recommendation, March 2013

# Typical SPARQL query shapes

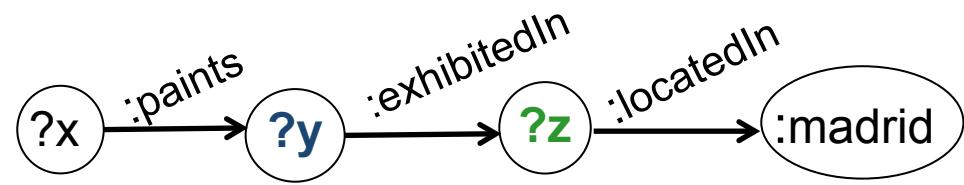
## Star query

```
SELECT ?y ?z  
WHERE {?x type :artist .  
       ?x :firstName ?y .  
       ?x :paints ?z .}
```


$$\text{?y ?z: } (\text{?x type :artist} \wedge  
(\text{?x :firstName ?y}) \wedge  
(\text{?x :paints ?z}))$$

## Path query

```
SELECT ?x ?y ?z  
WHERE {?x :paints ?y .  
       ?y :exhibitedIn ?z .  
       ?z :locatedIn :madrid.}
```


$$\text{?x ?y ?z: } (\text{?x :paints ?y}) \wedge  
(\text{?y :exhibitedIn ?z}) \wedge  
(\text{?z :locatedIn :madrid})$$

# Centralized RDF stores

- How to store RDF in a centralized setting?
- Approaches relying on RDBMs
  - Big triple tables

Triples

subject	property	object
:picasso	type	:cubist
:picasso	:firstName	“Pablo”
picasso	:paints	:guernica
:guernica	:exhibitedIn	:reinasofia
:reinasofia	:locatedIn	:madrid
:rodin	type	:sculptor
:rodin	:creates	:thethinker
...	...	...

# Centralized RDF stores

- How to store RDF in a centralized setting?
- Approaches relying on RDBMs
  - Big triple tables
  - **Property tables** [Alexaki01, Wilkison06]

Artists

subject	type	firstName	paints	creates
:picasso	:cubist	“Pablo”	:guernica	null
:rodin	:sculptor	“Auguste”	null	:thethinker

Artifacts

subject	exhibitedIn
:guernica	:reinasofia
:thethinker	:museerodin

subject	...
...	...
...	...

# Centralized RDF stores

- How to store RDF in a centralized setting?
- Approaches relying on RDBMs

- Big triple tables

- Property tables [Alexaki01, Wilkison06]

- **Vertical partitioning** [Abadi07]

type

subject	
:picasso	
:rodin	

Fits well in  
column-stores

paints

object
:guernica



Benchmarking [Theoharis05]

firstName

edIn

subject	object
:picasso	“Pablo”
:rodin	“Auguste”

subject	object
:guernica	:reinasofia
:thethinker	museerodin

subject	object
...	...

# Centralized RDF stores

- How to store RDF in a centralized setting?
- Approaches relying on RDBMs
  - Big triple tables
  - Property tables [Alexaki01, Wilkison06]
  - Vertical partitioning [Abadi07]
- Native RDF stores
  - RDF-3X [Neumann10], Hexastore [Weiss08]
  - 6 indexes (1 for each permutation of s p o)
  - compressed indices
  - aggregated indices

# Dictionary encoding

- RDF contains long strings of URIs

Size of triples

	min	max	avg
	80 bytes	2100 bytes	240 bytes

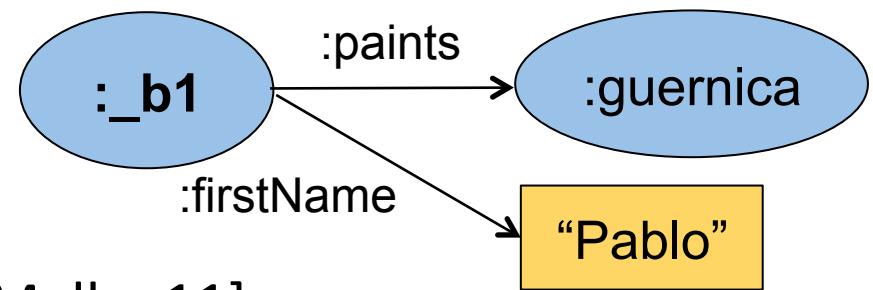


Source <http://gromgull.net/blog/category/semantic-web/billion-triple-challenge/>

Dictionary		Triples		Triples	
<b>id</b>	<b>resource</b>	<b>pro</b>	<b>subject</b>	<b>property</b>	<b>object</b>
1	:picasso	t	1	3	4
2	:rodin	:fir	1	5	10
3	type	:p	1	7	6
4	:cubist	:exh	6	12	9
5	:firstName	:loc	...	...	...
6	:guernica	type		:sculptor	
...	...	:creates		:thethinker	
...	...	...		...	

# RDF is not just a 3-attribute table

- Blank nodes
  - unknown/anonymous resources
  - existential variables



- for more “On blank nodes” [Mallea11]
- Schema queries
  - $?x: (?x, sc, :artist) \wedge (:paints, domain, ?x)$
- Data and schema can be queried together
  - $?x: (?x, :paints, ?y) \wedge (?x, type, ?z) \wedge (?z, sc, :artist)$

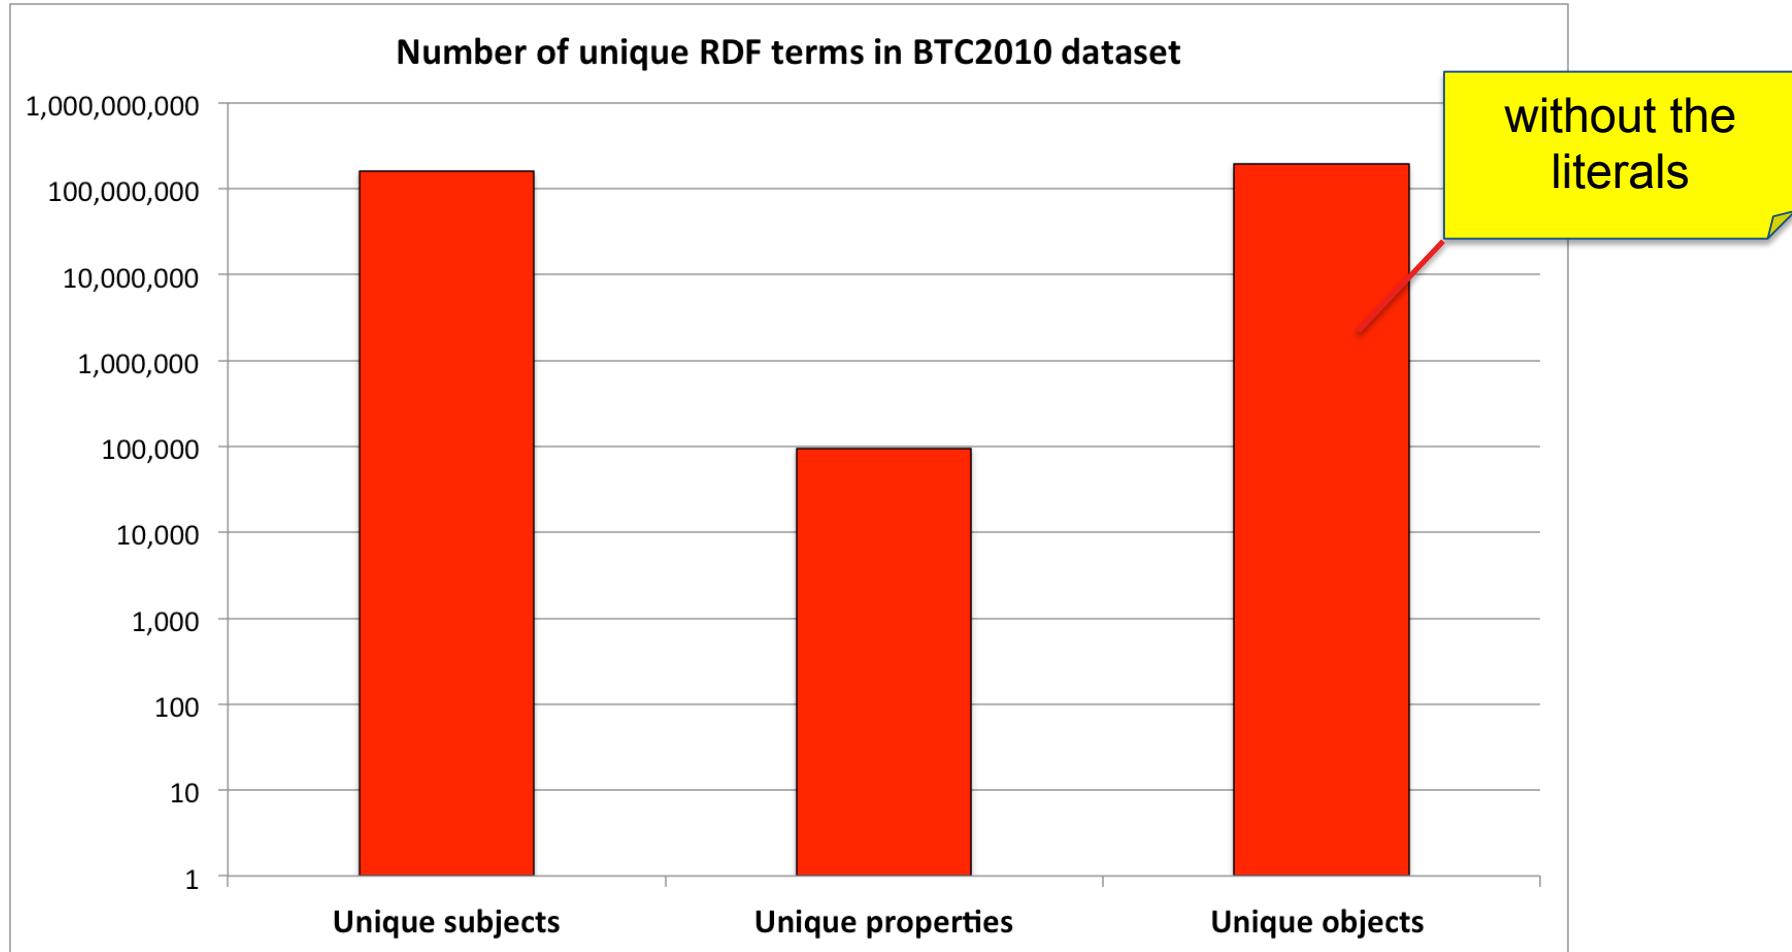
Data query

Schema query

# RDF is not just a 3-attribute table

- Properties in RDF are treated as first class citizens (they are also resources)
  - (`:paints`, `domain`, `:painter`)
  - joins between properties and subjects/objects  
`?x: (:paints, domain, :painter) ∧ (?y, domain, :painter)`
- RDFS inference
- and more ...

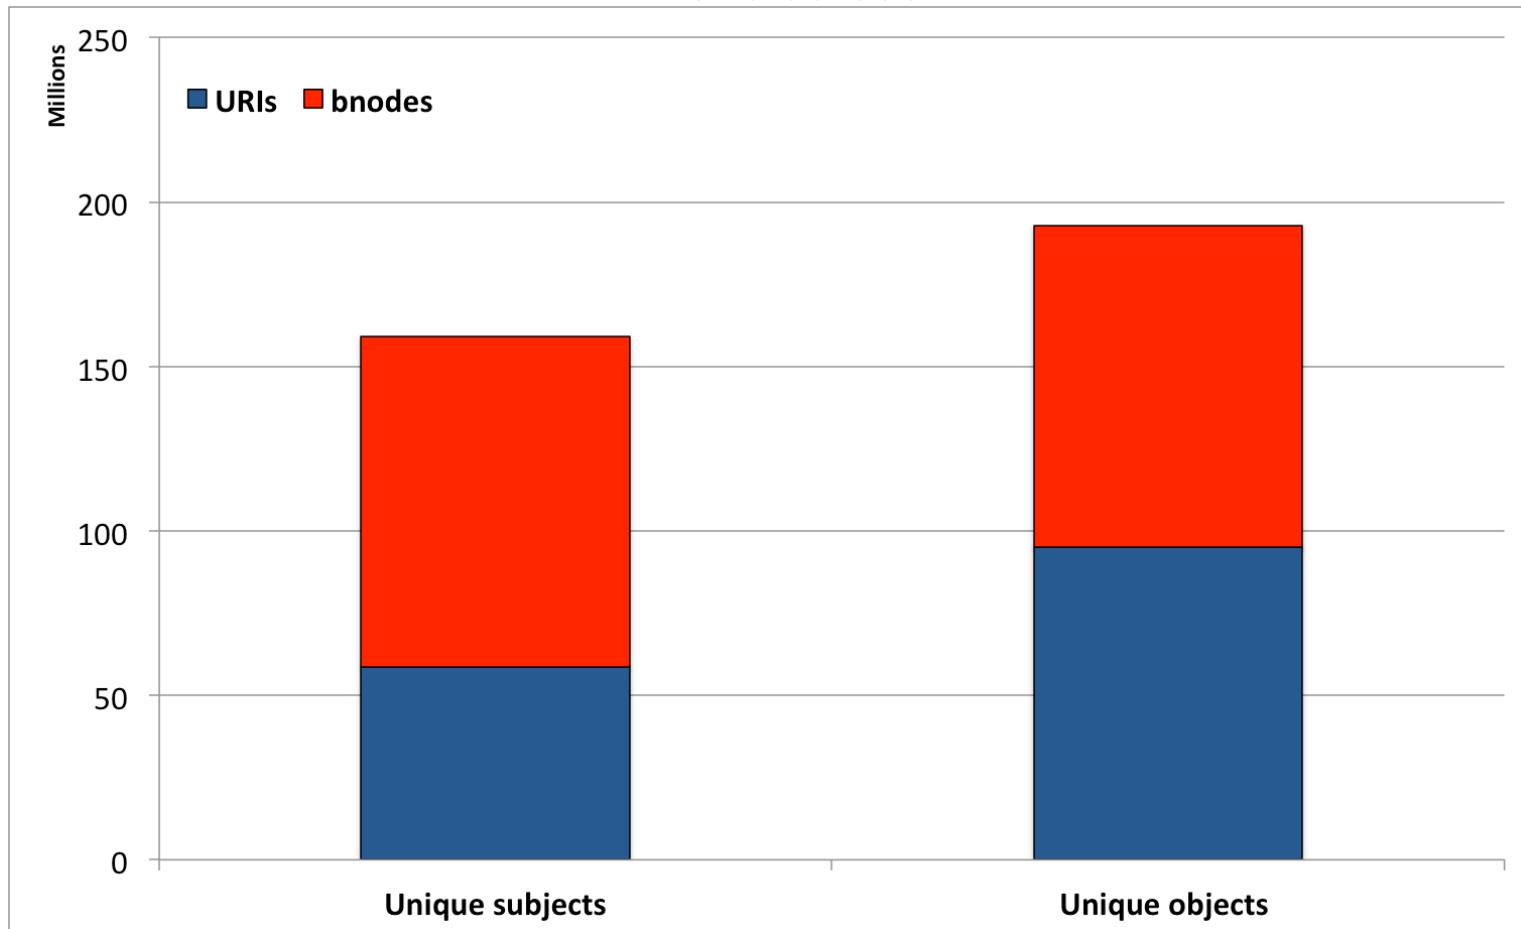
# How does real RDF look?



Source: <http://gromgull.net/blog/2010/09/btc2010-basic-stats/>

# How does real RDF look?

BTC2010 dataset

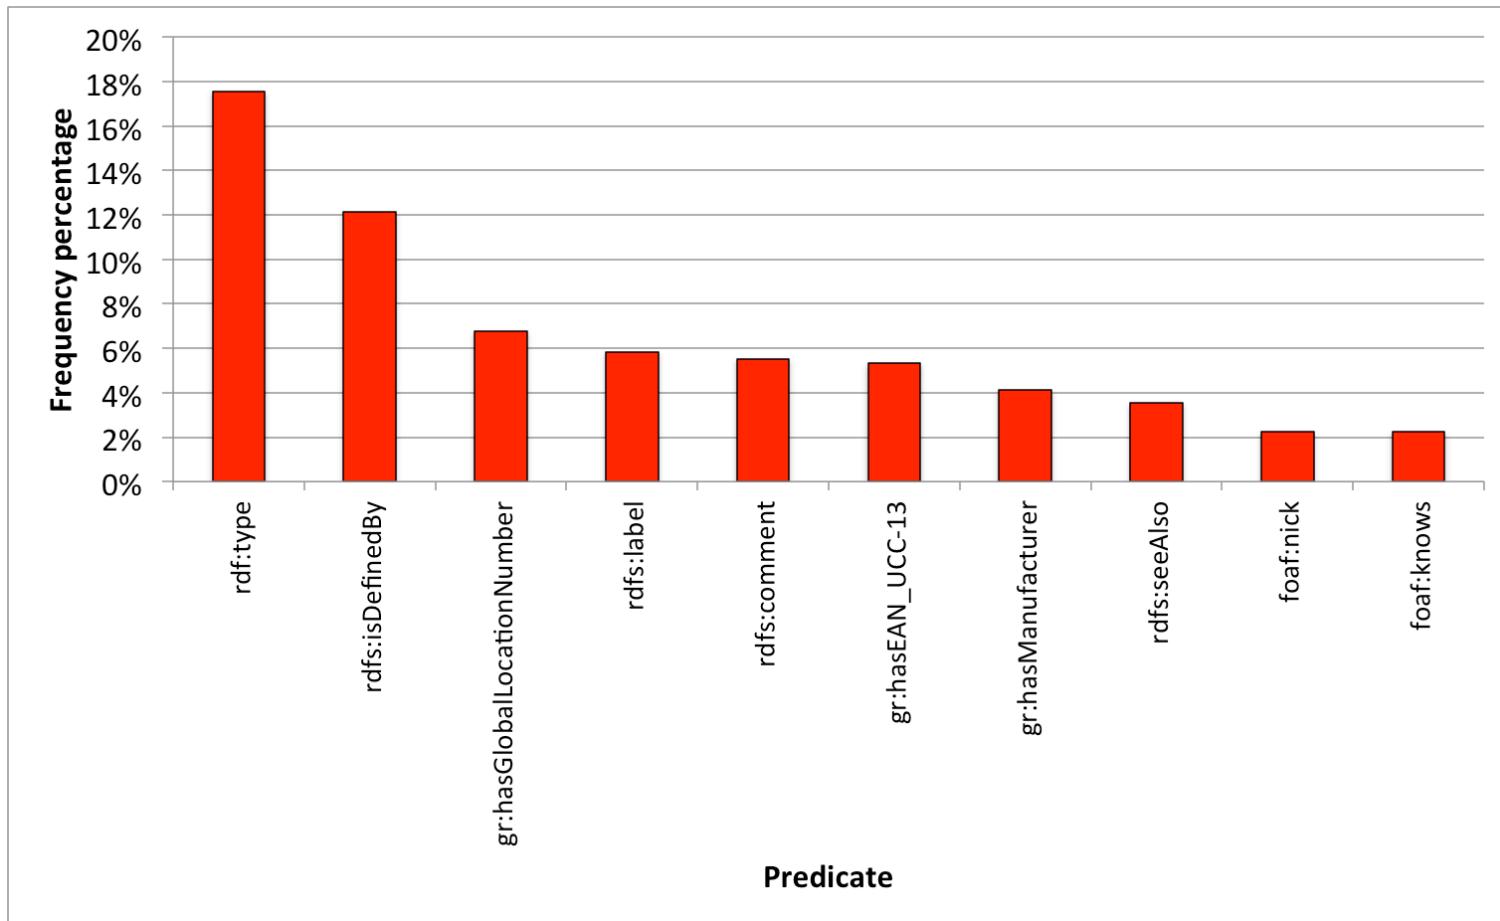


Source:

<http://gromgull.net/blog/2010/09/btc2010-basic-stats/>

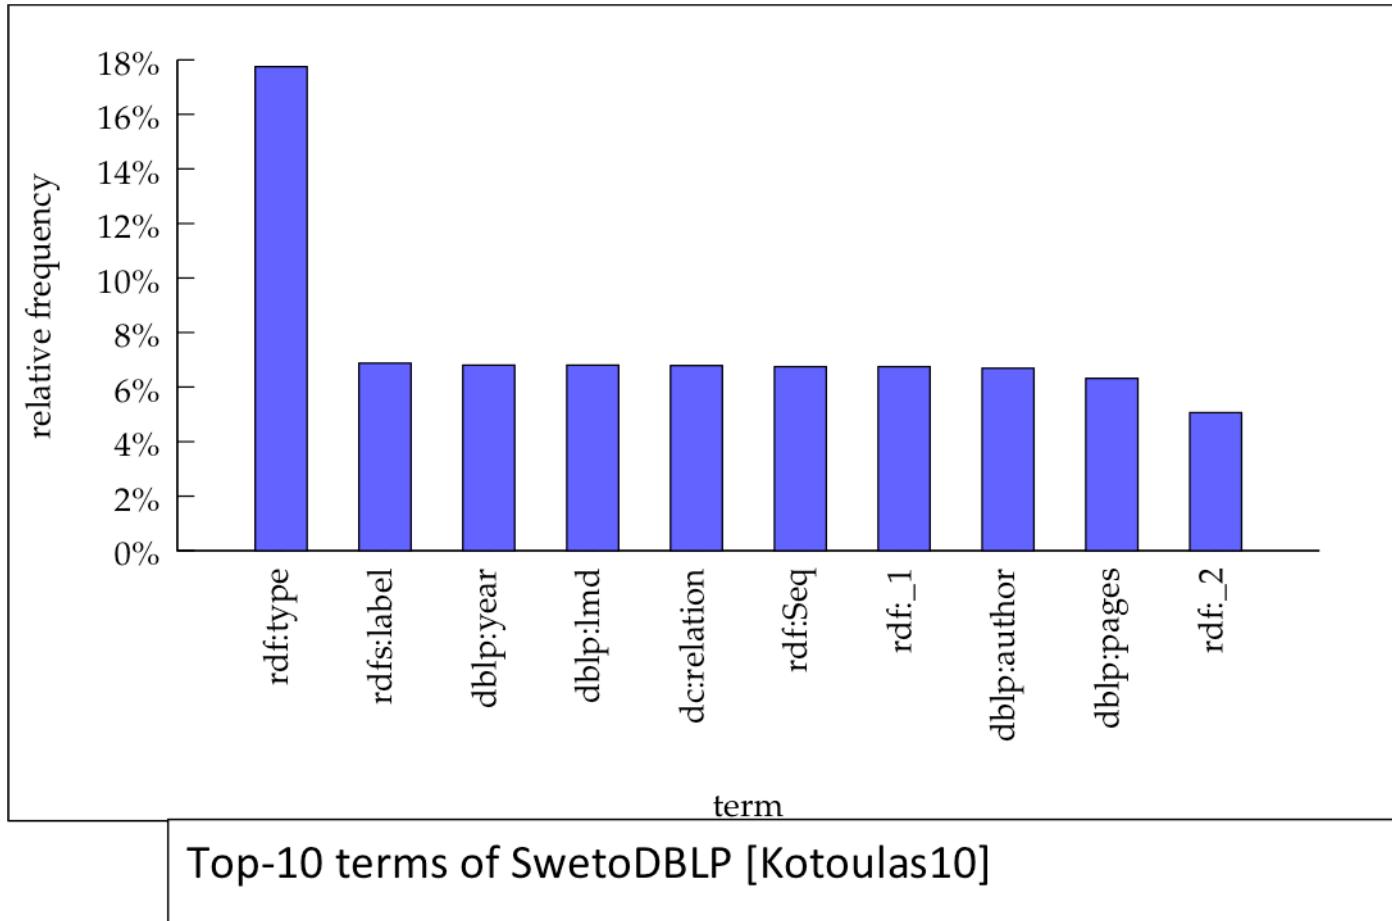
# How does real RDF look?

BTC2010 dataset



Source: <http://gromgull.net/blog/2010/09/btc2010-basic-stats/>

# How does real RDF look?



# 2

## CLOUD BASICS

# Cloud computing

*The origin of the term **cloud computing** was derived from [...] drawings of stylized clouds to denote networks [...] The cloud symbol was used to represent the **Internet** as early as 1994.* [Wikipedia: [http://en.wikipedia.org/wiki/Cloud\\_computing](http://en.wikipedia.org/wiki/Cloud_computing)]

- Use of *computing resources as a service* over the network
- Technical advantages:
  - Scalability
  - Elasticity
  - Reliability
- Non-technical advantage:
  - Outsourcing system maintenance

# Cloud computing service models

Google Apps

Software-as-a-Service (SaaS)



Applications

Windows Azure



Platform-as-a-Service (PaaS)



Platforms

amazon  
webservices™

Elastic Compute Cloud (EC2)

Infrastructure-as-a-Service (IaaS)



Hardware

**pay-per-use**



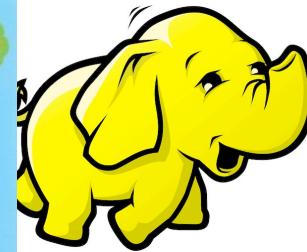
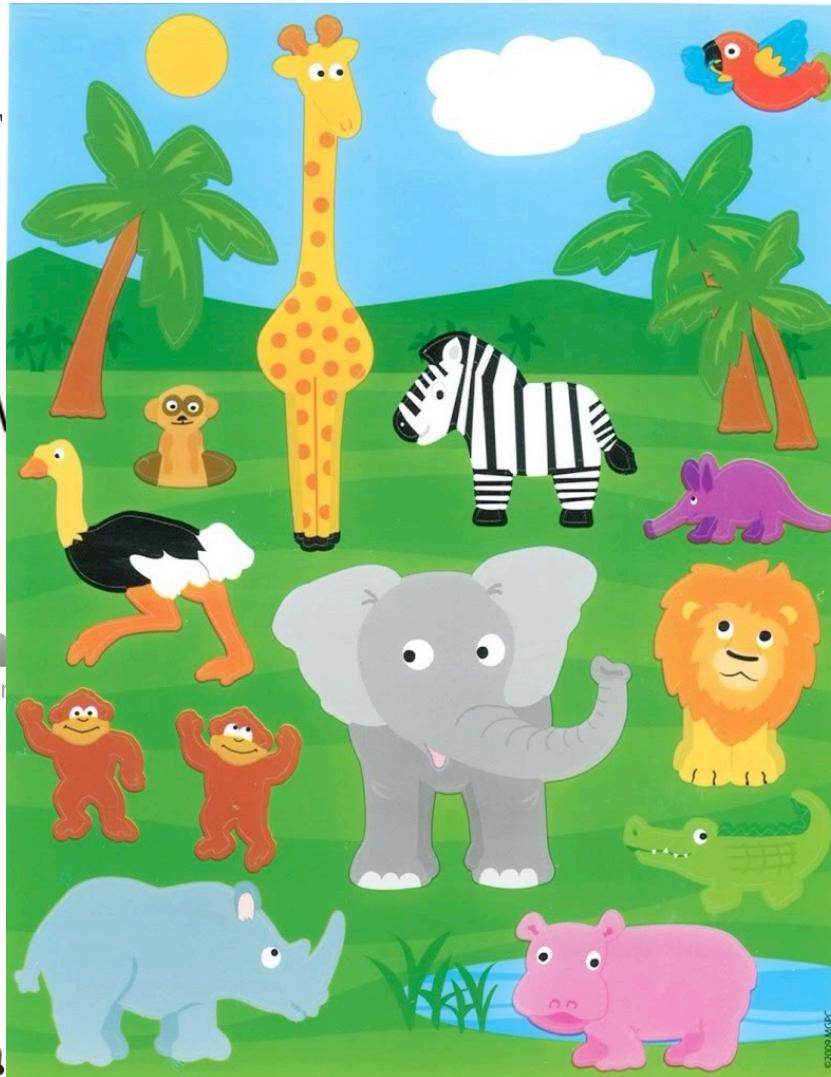
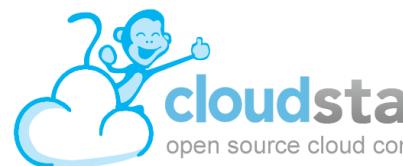
# Cloud computing and data management

- Relational databases in the cloud (DBaaS)
  - IBM DB2 on SmartCloud
  - Amazon RDS (MySQL, Oracle or SQL Server)
- NoSQL systems [Catell11]:
  - Key-value stores (memcached, DynamoDB, SimpleDB)
  - Extensible record stores (BigTable, Cassandra, HBase, Accumulo)
  - Document stores (MongoDB, CouchDB)
  - Graph databases (Pregel, Apache Giraph, Neo4J)
- Parallel data processing paradigms:
  - MapReduce (Hadoop, Amazon EMR), PACTs/Stratosphere, Spark

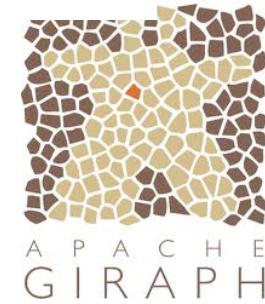
# Cloud ecosystem



Google  
App Engine



APACHE  
ASE



# Cloud infrastructures considered here

- I. Distributed key-value stores
- II. MapReduce-based parallel processing

# Distributed key-value stores

- Storing **schema-less** data in key-value pairs

- For each **key** there may be multiple key 1 → **(attribute-value)** pairs

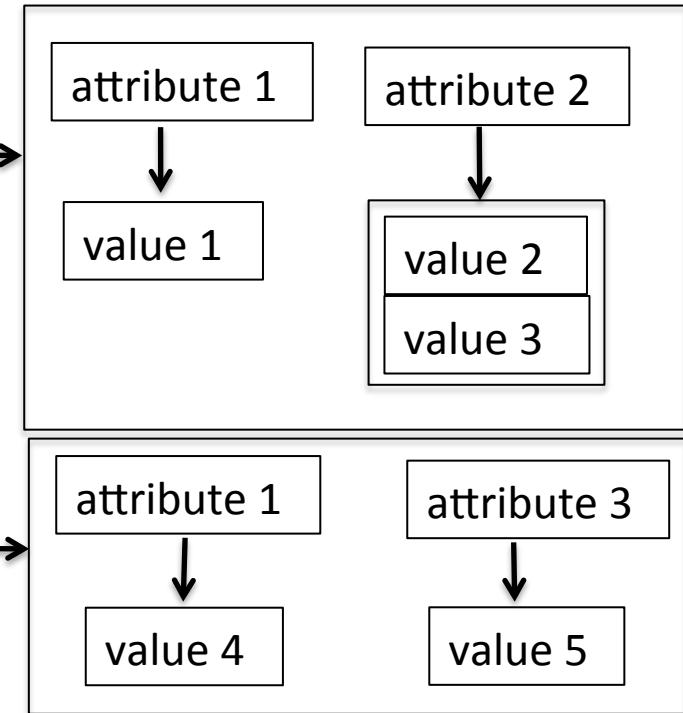
- Multi-valued attributes

- **Index:** Hash-based or lexicographically sorted on the **key** key 2 →

- Basic APIs:

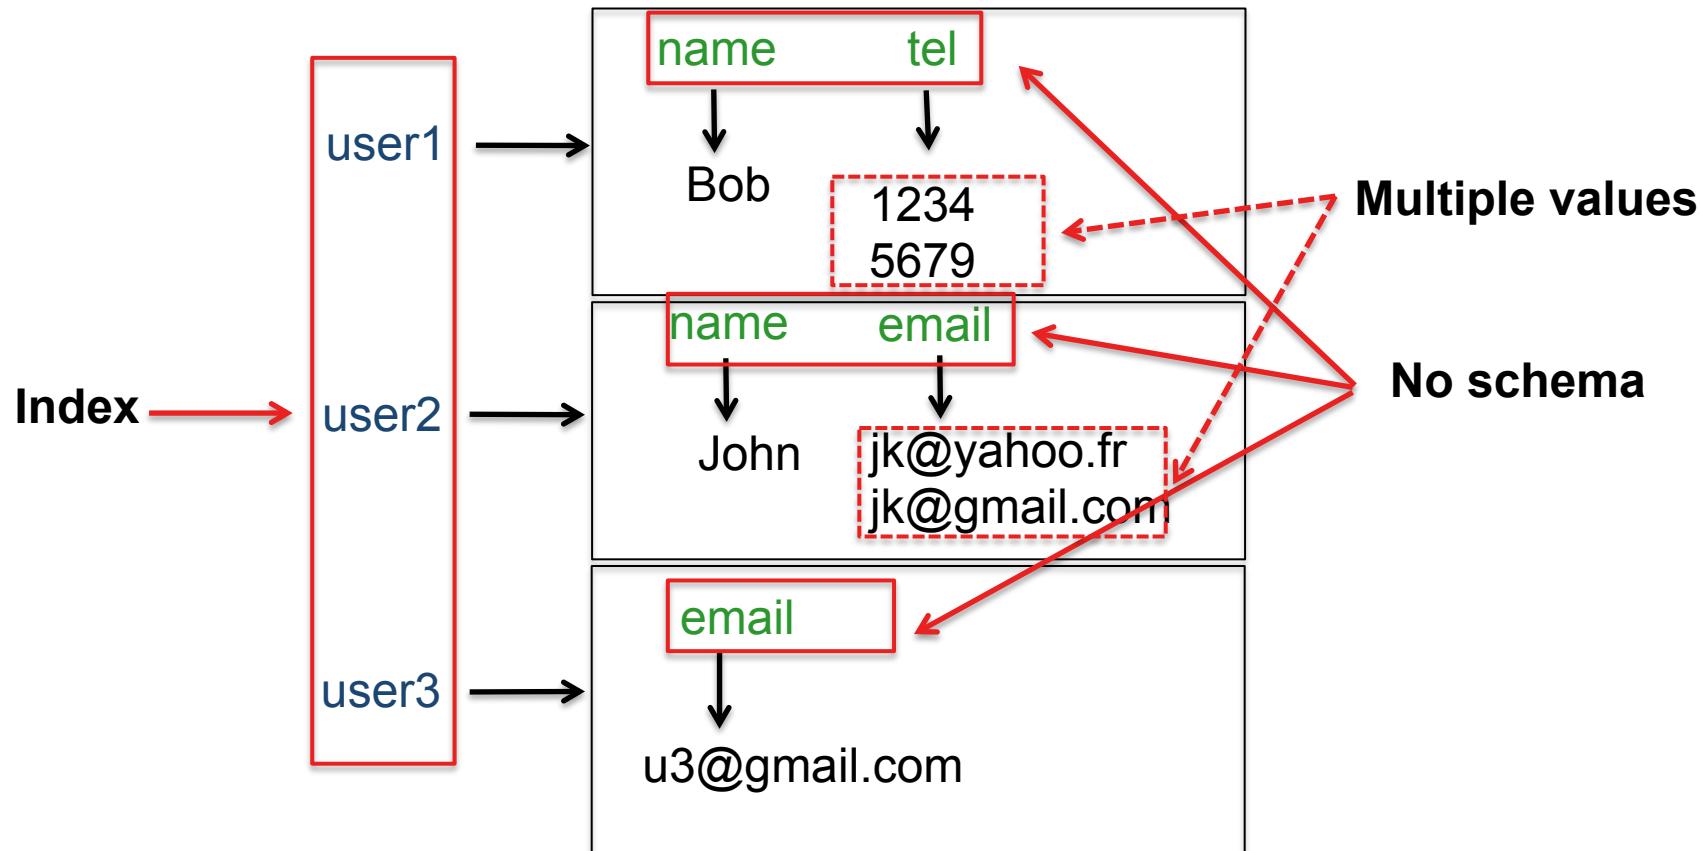
Put( $k, \langle a, v \rangle$ ), Get( $k$ ), Delete( $k$ )

- No support for operations across tables (e.g., join!)

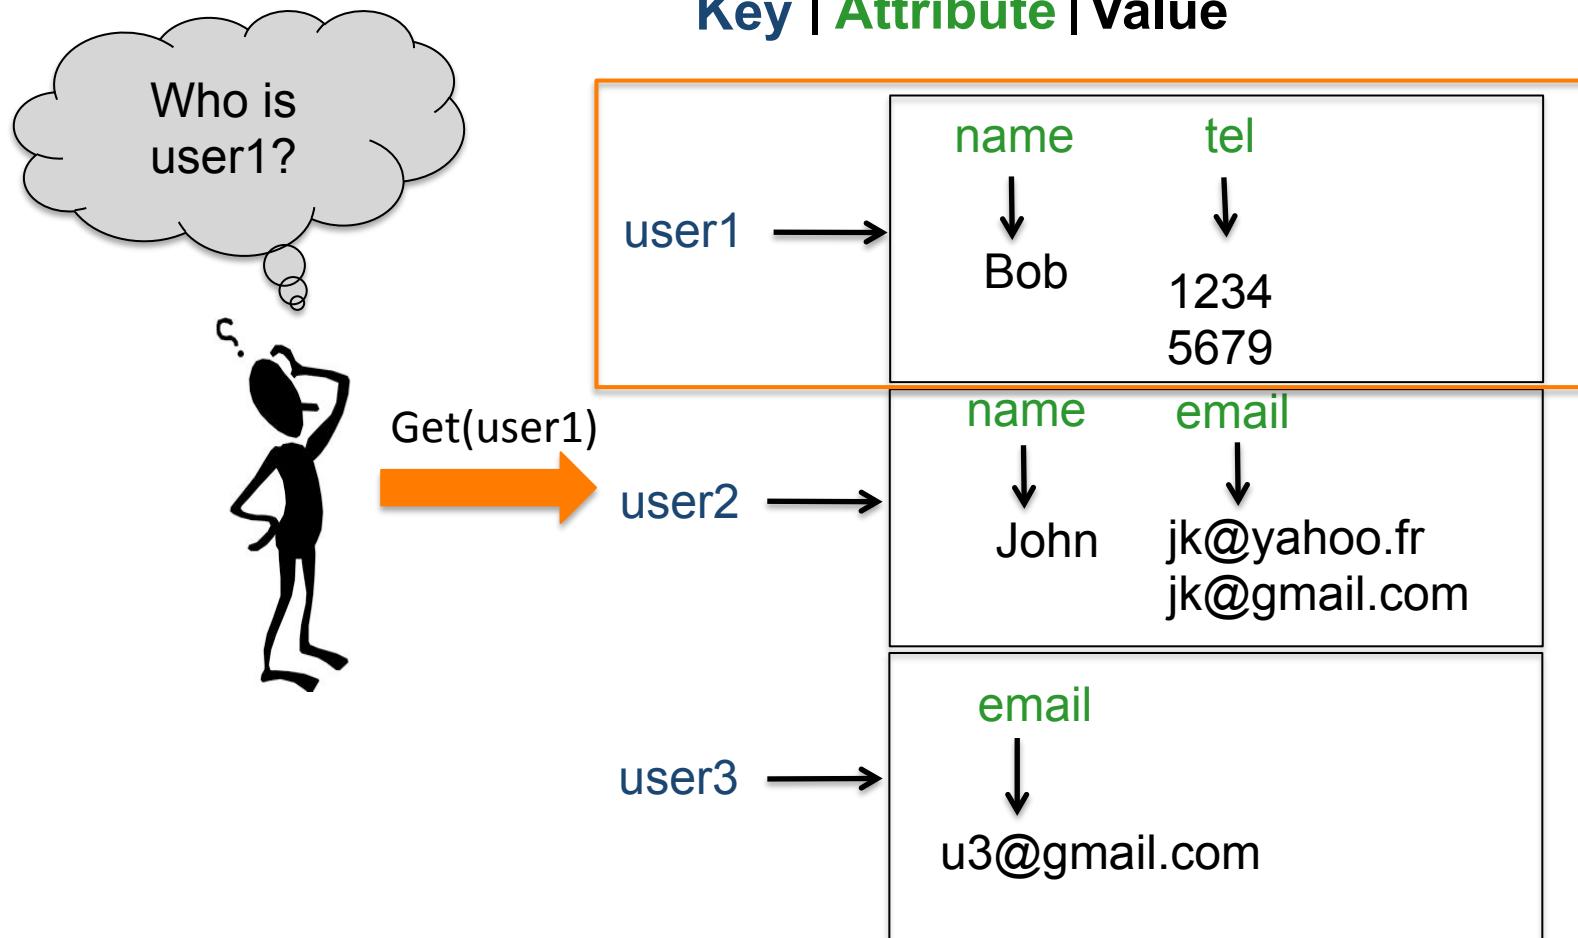


# Distributed key-value stores

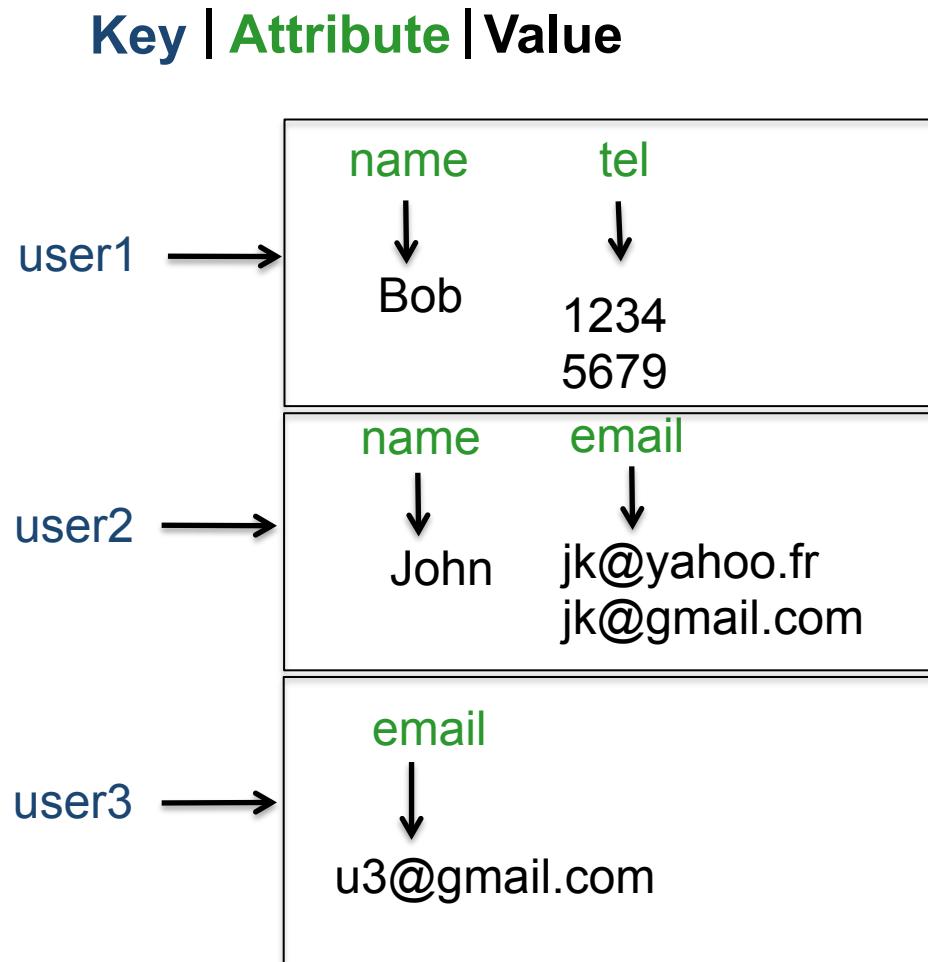
Key | Attribute | Value



# Distributed key-value stores

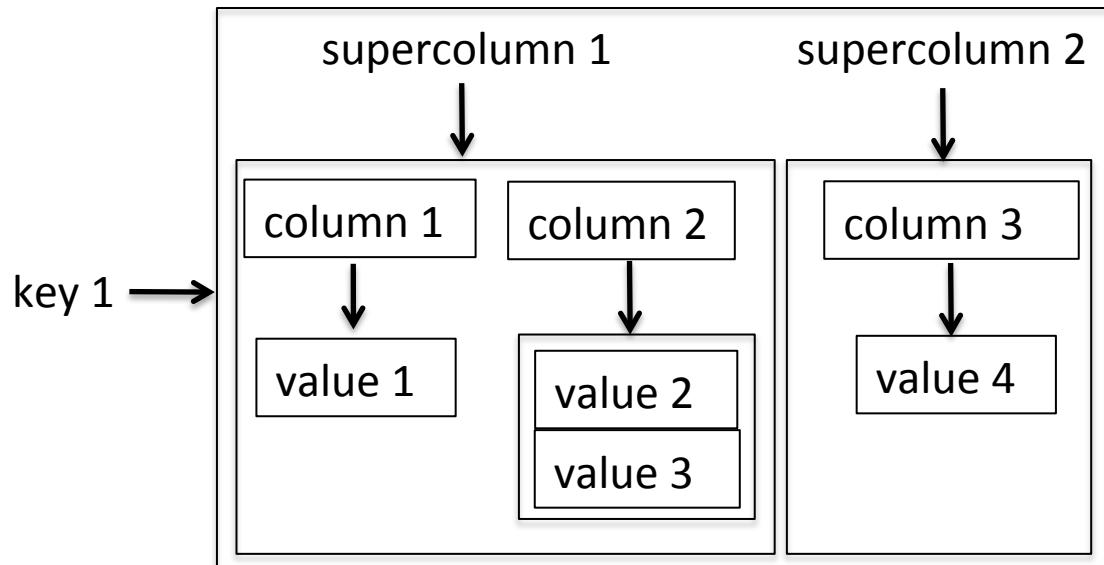


# Distributed key-value stores



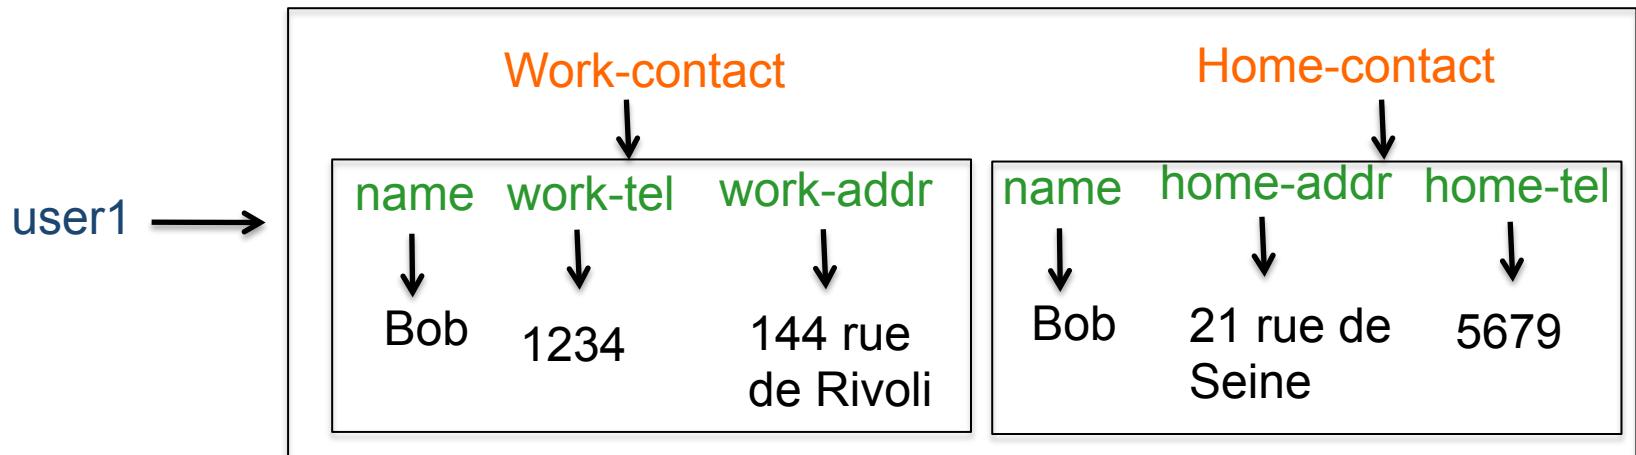
# Extended key-value stores

- Google BigTable, Apache Cassandra
- **Supercolumn**: column container
  - K|SC|C|V
- **Sorted** index on the **key**
- **Sorted** index on the **supercolumn**
- **Secondary** index
  - maps values to keys



# Extended key-value stores

Key | SuperColumn | Column | Value



# Cloud infrastructures

- I. Distributed key-value stores
- II. MapReduce-based parallel processing

# MapReduce in a nutshell

- Issued from Google [DG04]
- Massively parallel processing on commodity machines
- Master/slave architecture
- Input data is split horizontally into chunks
  - distributed file system (DFS)
- Simple API:
  - Map (*key, value*) → {*ikey, ivalue*}
  - Reduce (*ikey, {ivalue}*) → (*key', value'*)

Many “more modern” extensions.

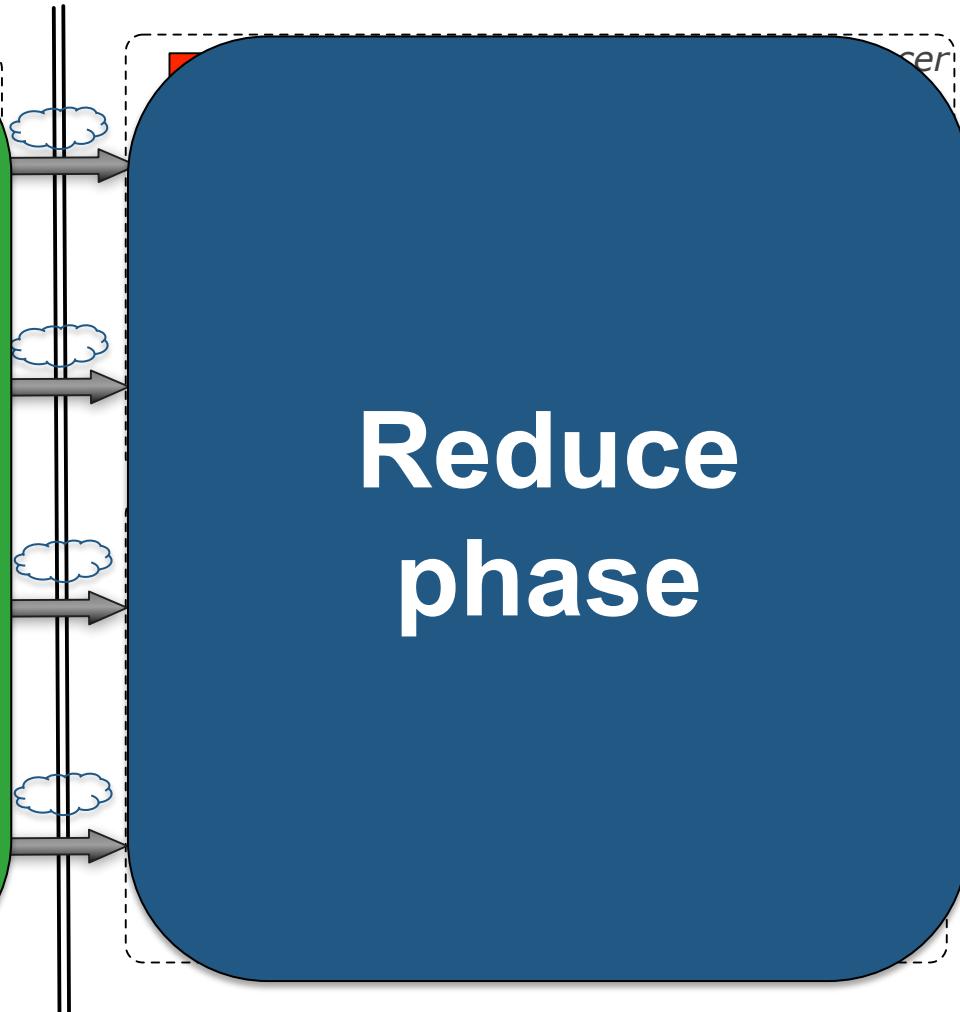
Extensively used in systems we cover.

# MapReduce in a nutshell

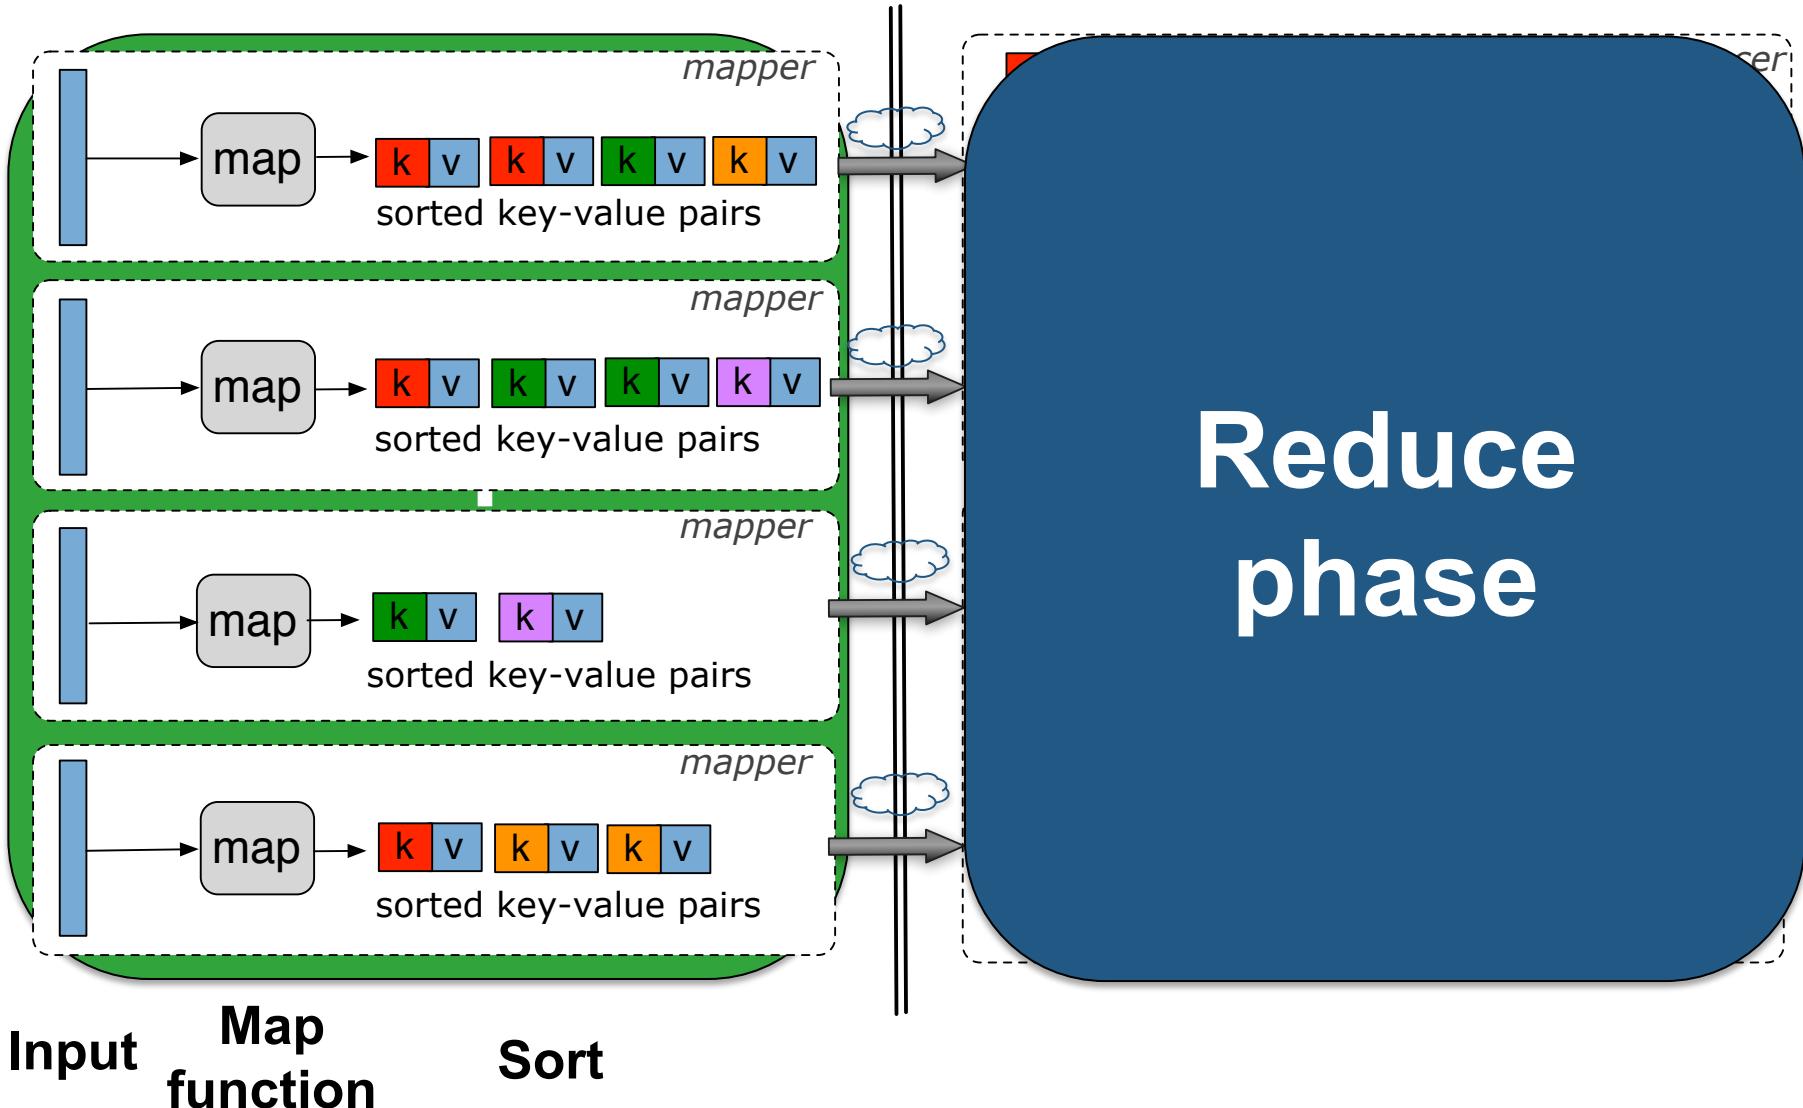
- **Map phase:** independent processes (mappers) which run **in parallel**
  - operate on the chunks of input data
  - output intermediate results
- **Shuffle phase:** Intermediate results are shuffled through the network
- **Reduce phase:** independent processes (reducers) which run **in parallel**
  - group intermediate results of the map phase
  - operate on the groups
  - output final results
- Key-value pairs with the **same key *ikey*** meet at the **same reducer**
  - e.g.,  $\text{hash}(ikey) \bmod R$
- Open source implementation: **Hadoop** and **HDFS**

# MapReduce illustration

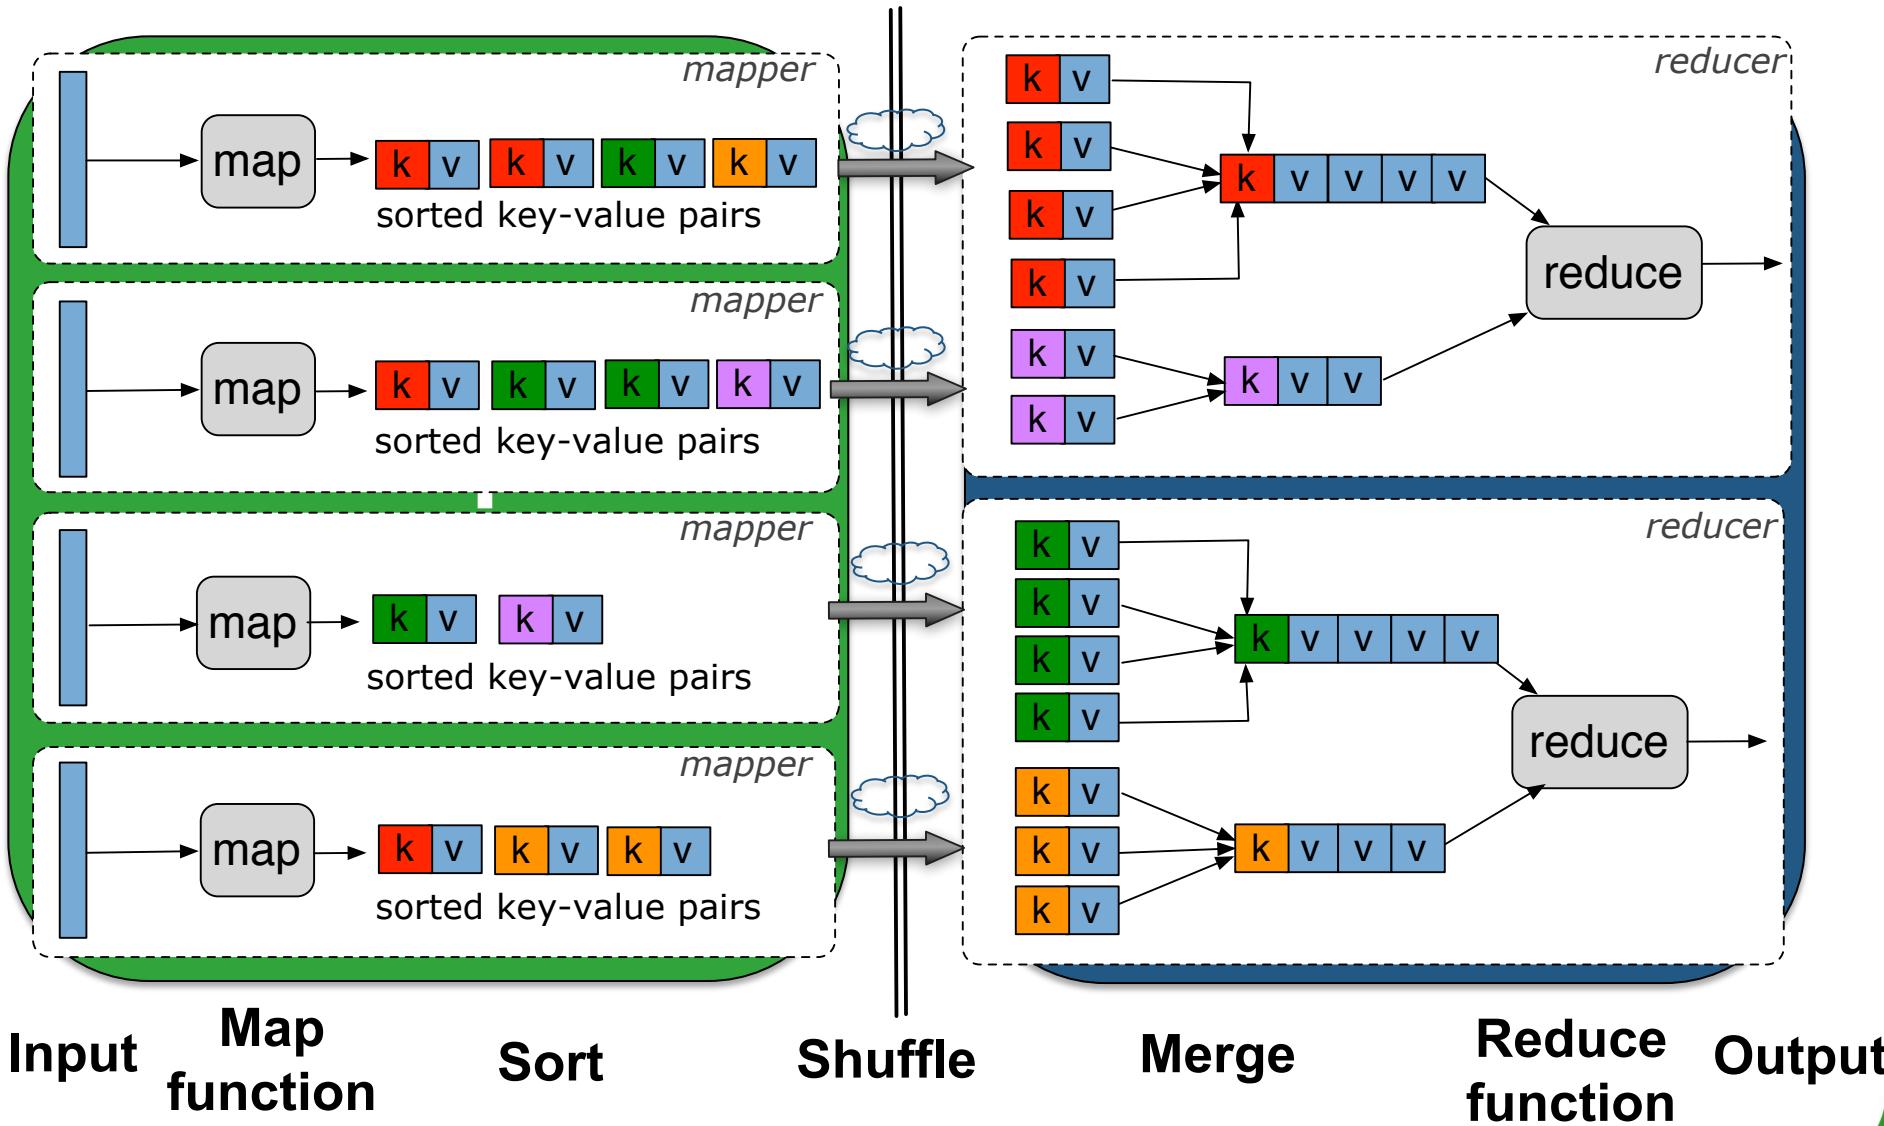
Map phase      Reduce phase



# MapReduce illustration

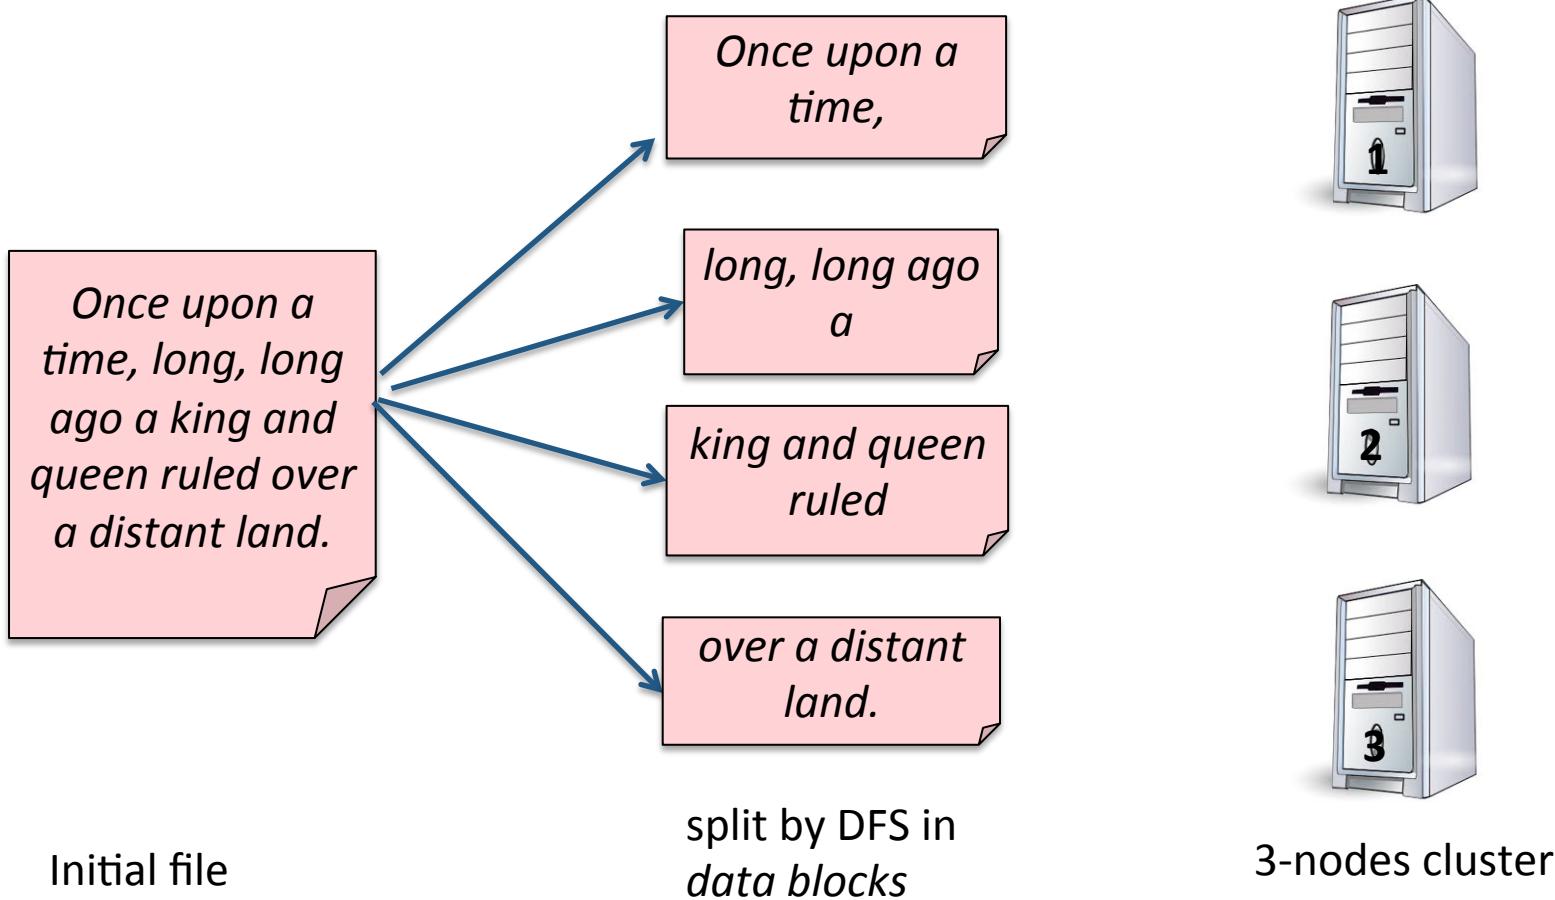


# MapReduce illustration

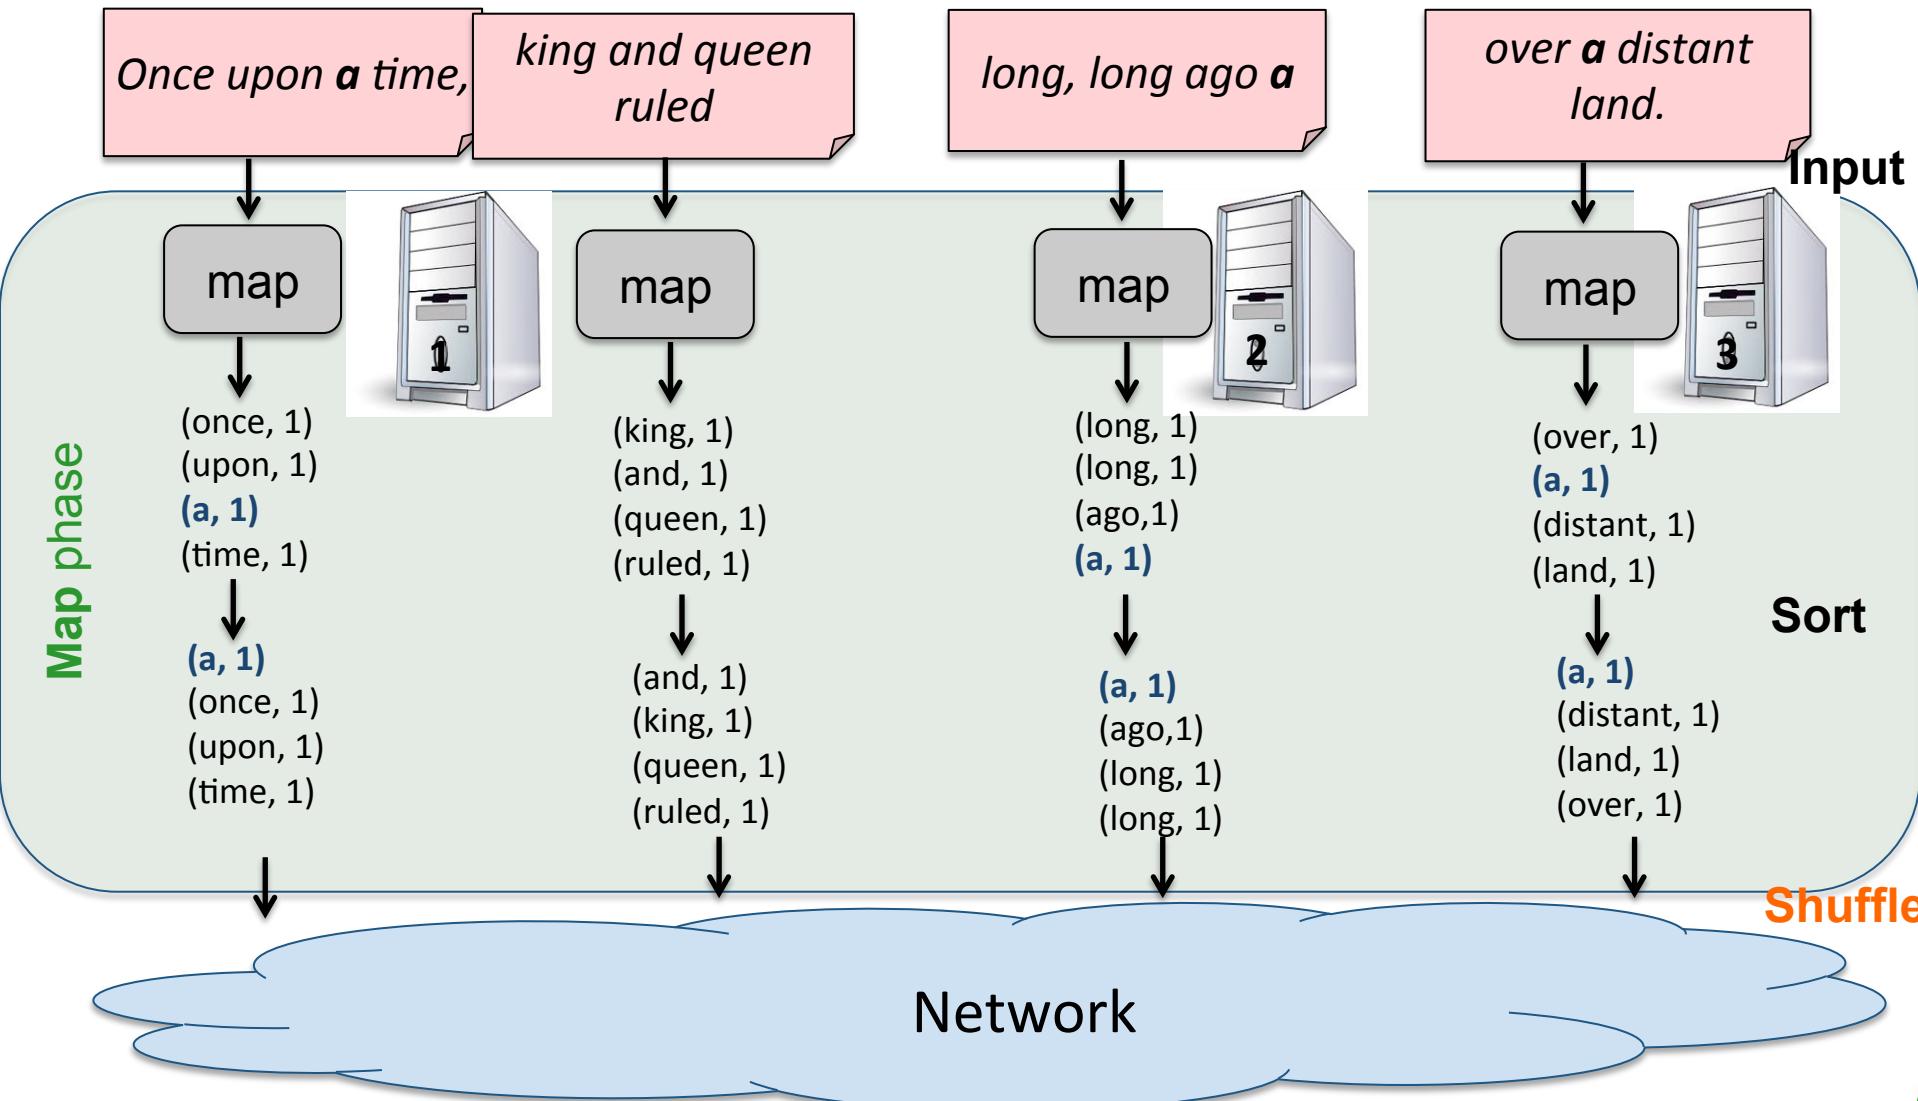


# MapReduce by words-count example

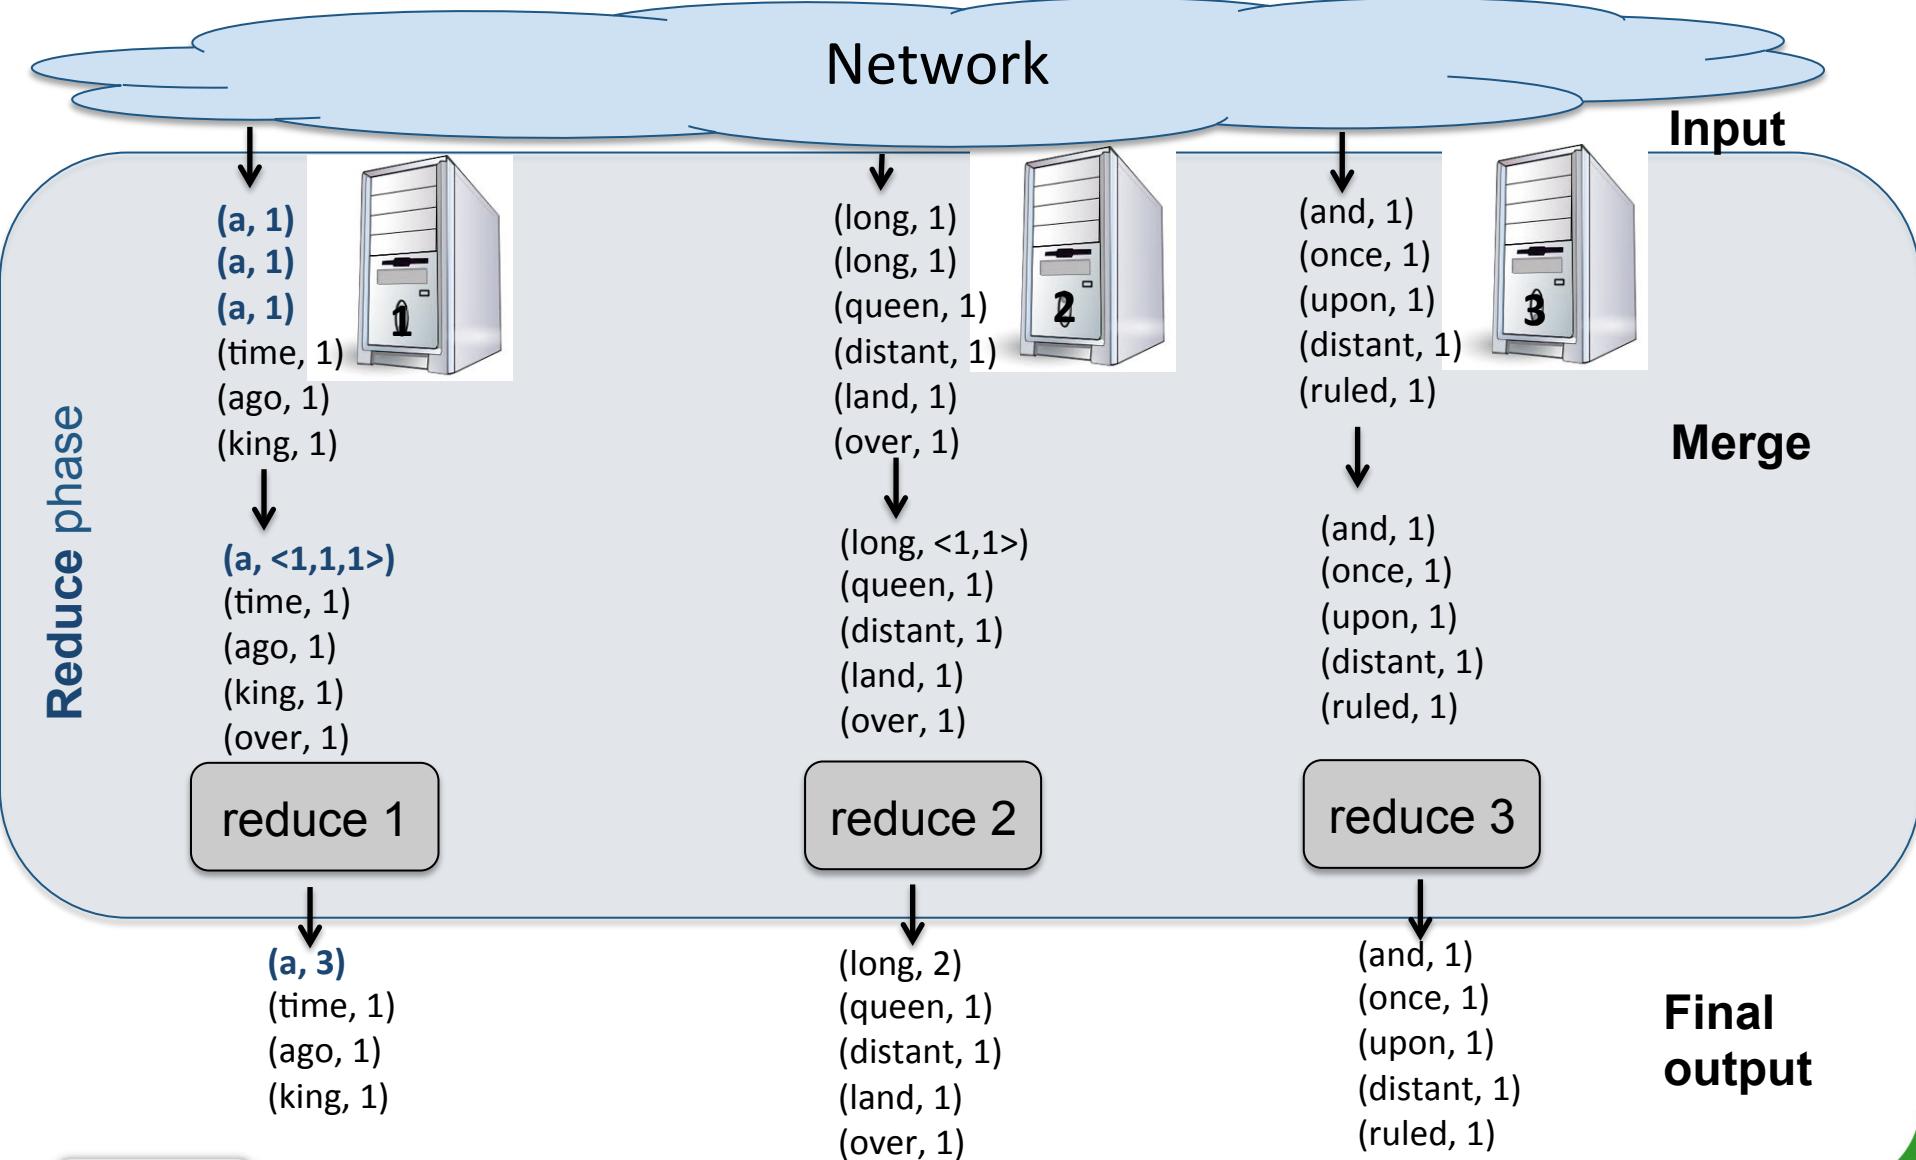
- Count the number of words in parallel



# MapReduce by words-count example (Map phase)



# MapReduce by words-count example (Reduce phase)



# MapReduce code 😊

```
1 package org.myorg;
2
3 import java.io.IOException;
4 import java.util.*;
5
6 import org.apache.hadoop.fs.Path;
7 import org.apache.hadoop.conf.*;
8 import org.apache.hadoop.io.*;
9 import org.apache.hadoop.mapreduce.*;
10 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
11 import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
12 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
13 import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
14
15 public class WordCount {
16
17     public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {
18         private final static IntWritable one = new IntWritable(1);
19         private Text word = new Text();
20
21         public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
22             String line = value.toString();
23             StringTokenizer tokenizer = new StringTokenizer(line);
24             while (tokenizer.hasMoreTokens()) {
25                 word.set(tokenizer.nextToken());
26                 context.write(word, one);
27             }
28         }
29     }
30
31     public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {
32
33         public void reduce(Text key, Iterable<IntWritable> values, Context context)
34             throws IOException, InterruptedException {
35             int sum = 0;
36             for (IntWritable val : values) {
37                 sum += val.get();
38             }
39             context.write(key, new IntWritable(sum));
40         }
41     }
42
43     public static void main(String[] args) throws Exception {
44         Configuration conf = new Configuration();
45
46         Job job = new Job(conf, "wordcount");
47
48         job.setOutputKeyClass(Text.class);
49         job.setOutputValueClass(IntWritable.class);
50
51         job.setMapperClass(Map.class);
52         job.setReducerClass(Reduce.class);
53
54         job.setInputFormatClass(TextInputFormat.class);
55         job.setOutputFormatClass(TextOutputFormat.class);
56
57         TextInputFormat.addInputPath(job, new Path(args[0]));
58         FileOutputFormat.setOutputPath(job, new Path(args[1]));
59
60         job.waitForCompletion(true);
61     }
62
63 }
```

Only 63 lines for a distributed/parallel application!

<http://wiki.apache.org/hadoop/WordCount>

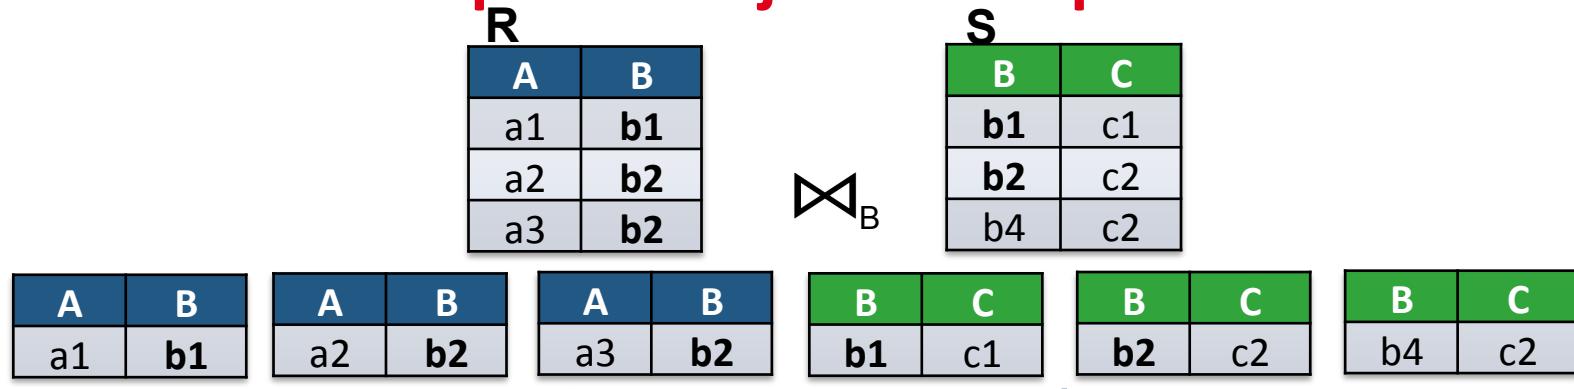
# MapReduce - pros/cons

- + High parallelization
- + Scalable
  - number of nodes
  - size of data
- + Fault-tolerant
  - replication of data
- + Easy to implement
- Operates on only one input
- No indexes
  - scan of all data
- Blocking operations
  - e.g., reduce does not begin until all tasks have finished

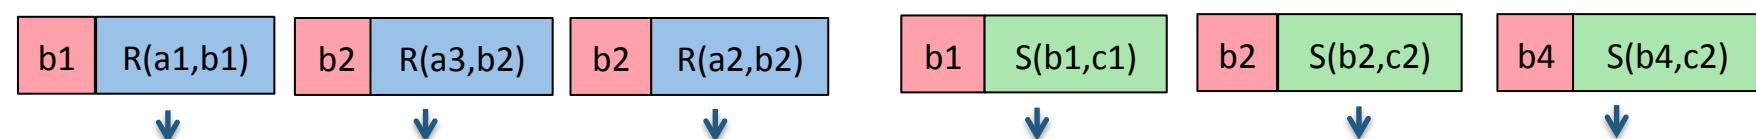
# Joins in MapReduce

- Two main strategies:
  - symmetric hash join (repartition join)
  - replicated join (broadcast join)
- Representative works:
  - [Blanas10] on different join methods in MapReduce
  - [Afrati11] on optimizing multiway joins in MapReduce
  - [Wu11] on query optimization in MapReduce

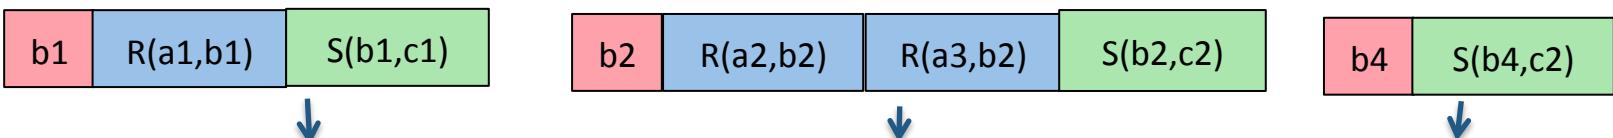
# Repartition join in MapReduce



Map



Shuffle



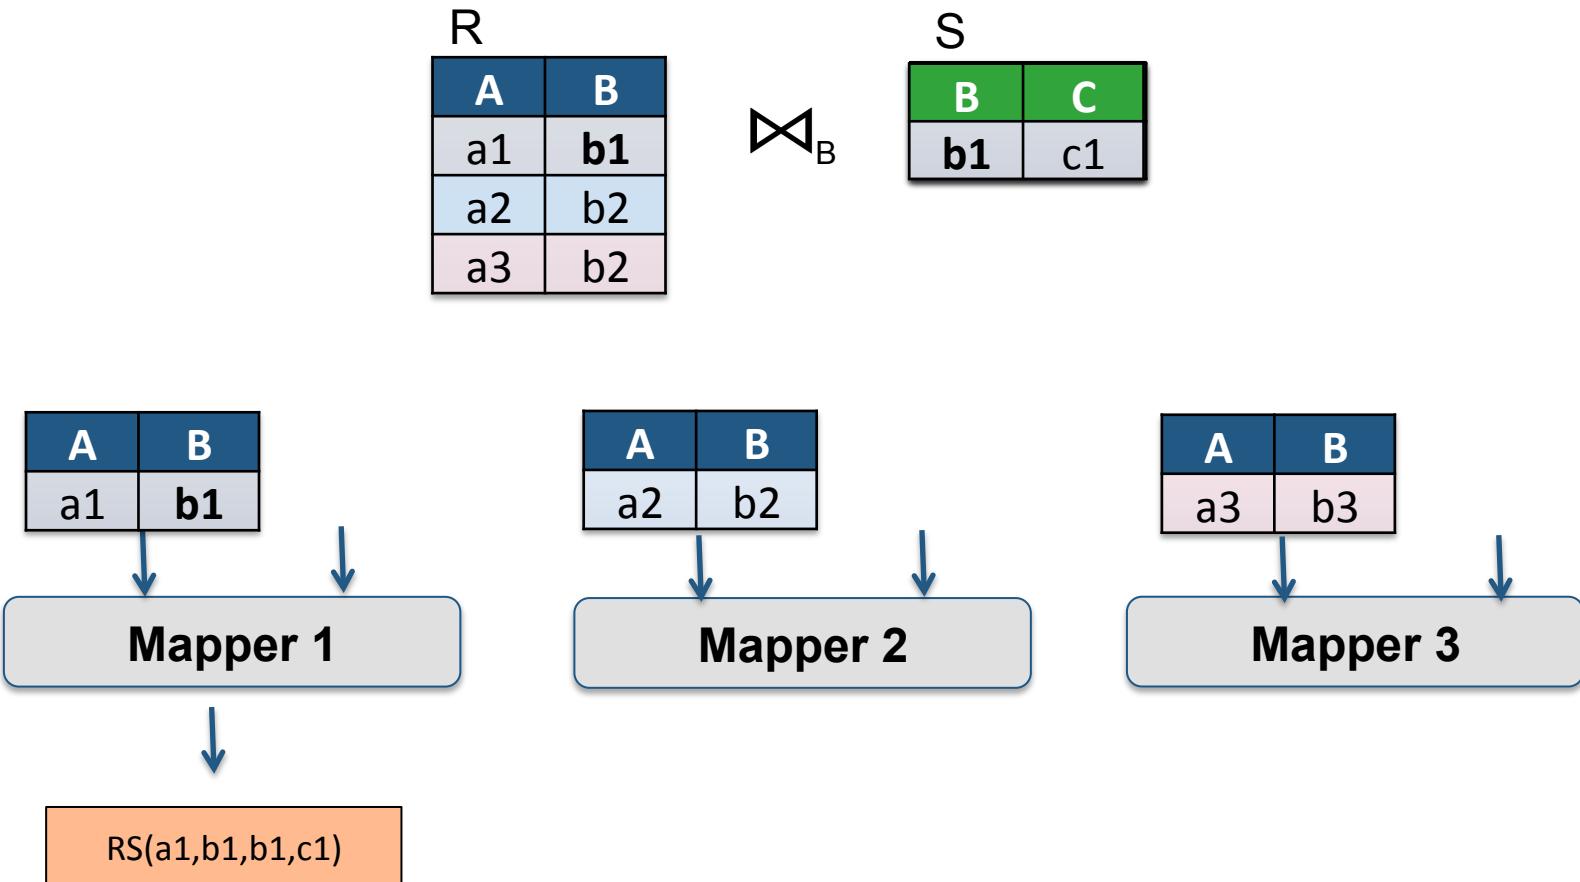
Reduce

RS(a1,b1,b1,c1)

RS(a2,b2,b2,c2)

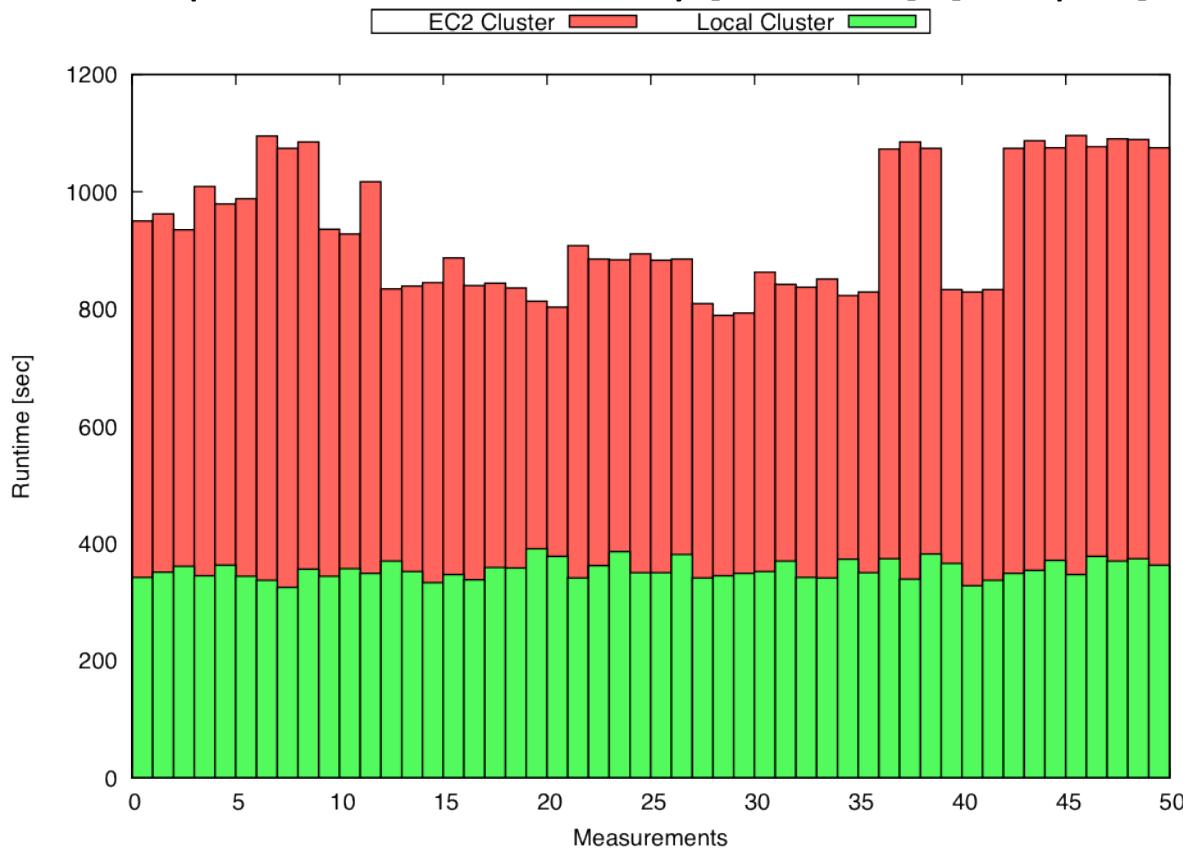
RS(a3,b2,b2,c2)

# Replicated (broadcast) join in MapReduce



# Variability in the clouds

Cloud = performance variability [Schad10], [Iosup11]



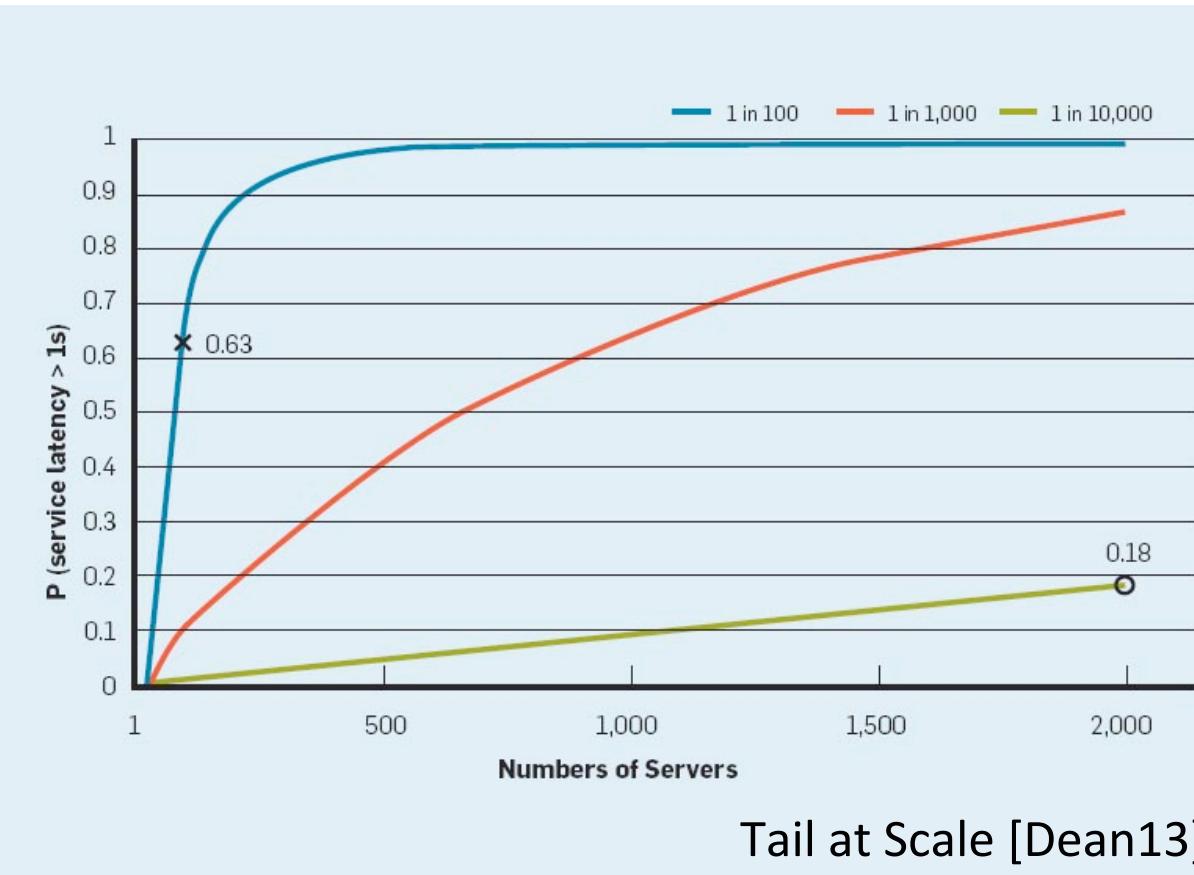
**Figure 1: Runtime for a MapReduce job.**  
from [Schad10]

# Why the variable performance?

- **Software**-based reasons:
  - Shared resources (cpu, memory, network bandwidth)
  - Daemons running in the background
  - Global resource sharing such as network switches and shared file systems
  - Maintenance activities running in the background
  - Multiple layers of queuing
- **Hardware**-based reasons:
  - Garbage collection for SSD
  - Energy management

# Variability in the clouds

Variability is amplified by scale



# 3

## ANALYSIS OF MASSIVELY DISTRIBUTED RDF DATA MANAGEMENT PLATFORMS

# RDF systems in the cloud

- Relatively **new** and **fast-growing** research area
- Numerous systems
  - **AMADA** [Bugiotti11, Arandar12]
  - CumulusRDF [Ladwig11]
  - EAGRE [Zhang13]
  - **Graph partitioning** [Huang11]
  - H2RDF [Papailiou12]
  - **HadoopRDF** [Husain11]
  - MAPSIN [Schätzle12]
  - PigSPARQL [Schätzle11]
  - RAPID+ [Ravindra11]
  - **Rya** [Punnoose12]
  - Stratustore [Ladwig11]
  - **SHARD** [Rohloff10]
  - **TriAD** [Gurajada14]
  - **Trinity.RDF** [Zeng13]
  - **WebPie** [Urbani09]
  - QueryPie [Urbani11]
  - ...

# Analysis dimensions

**1. Data storage**

**2. Query processing**

**3. RDFS entailment**

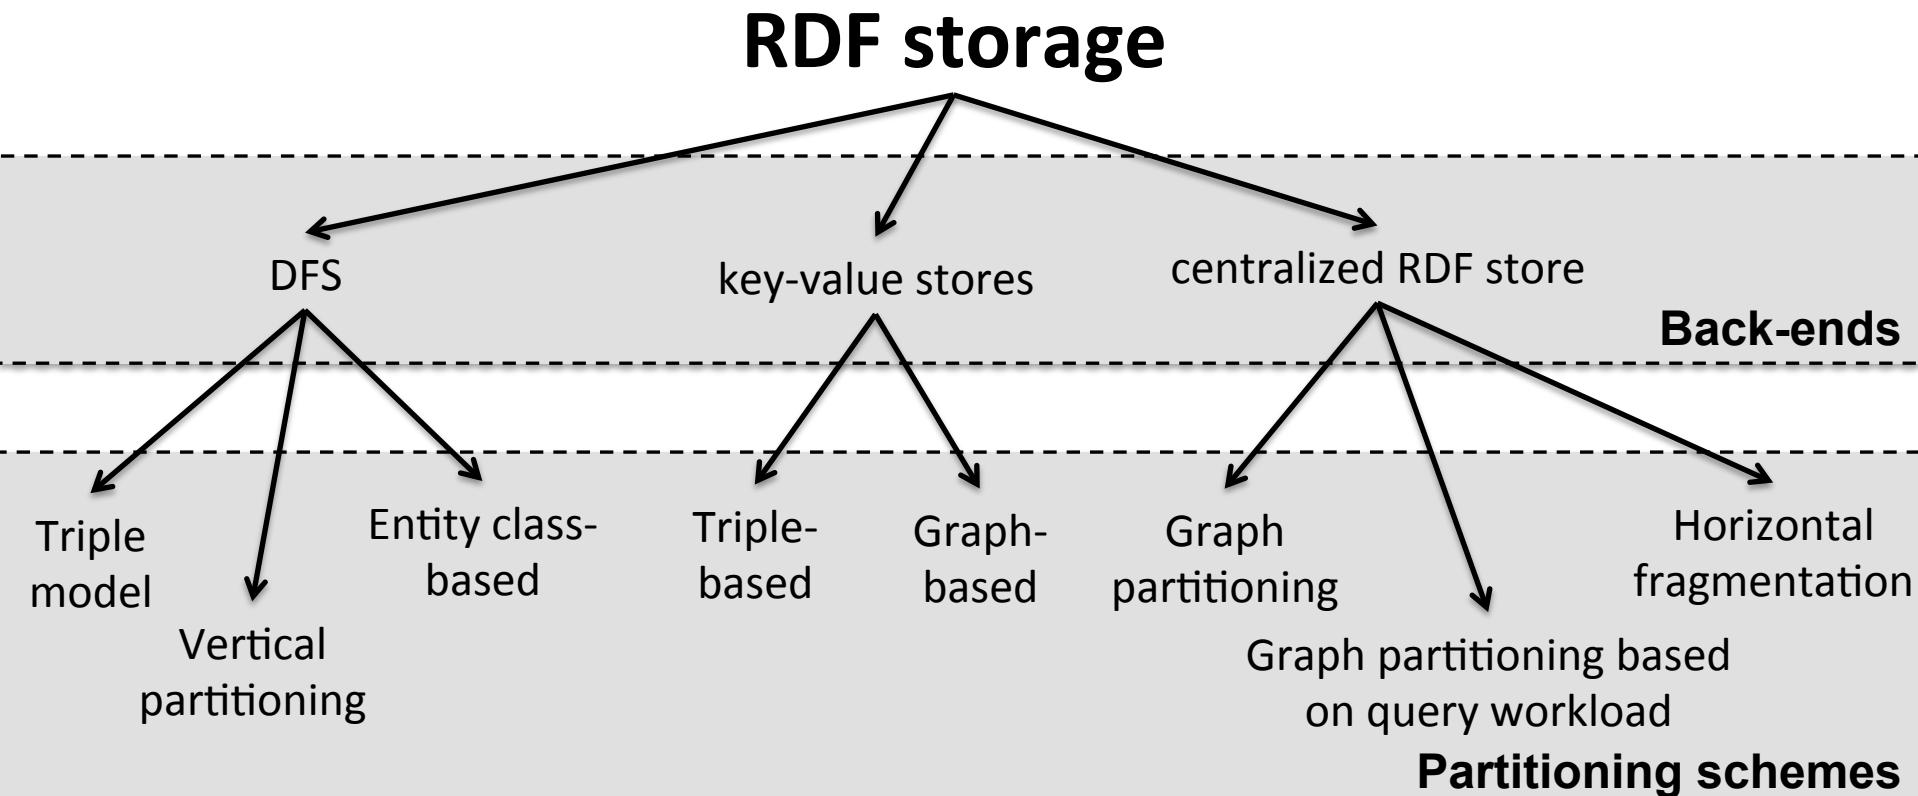
# Analysis dimensions

**1. Data storage**

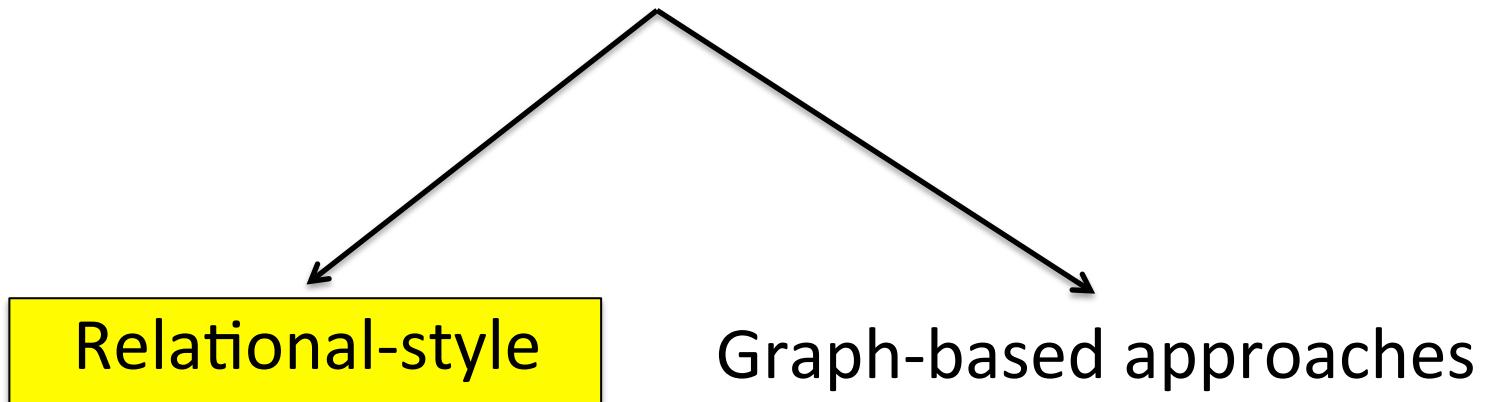
**2. Query processing**

**3. RDFS entailment**

# Data storage and partitioning schemes

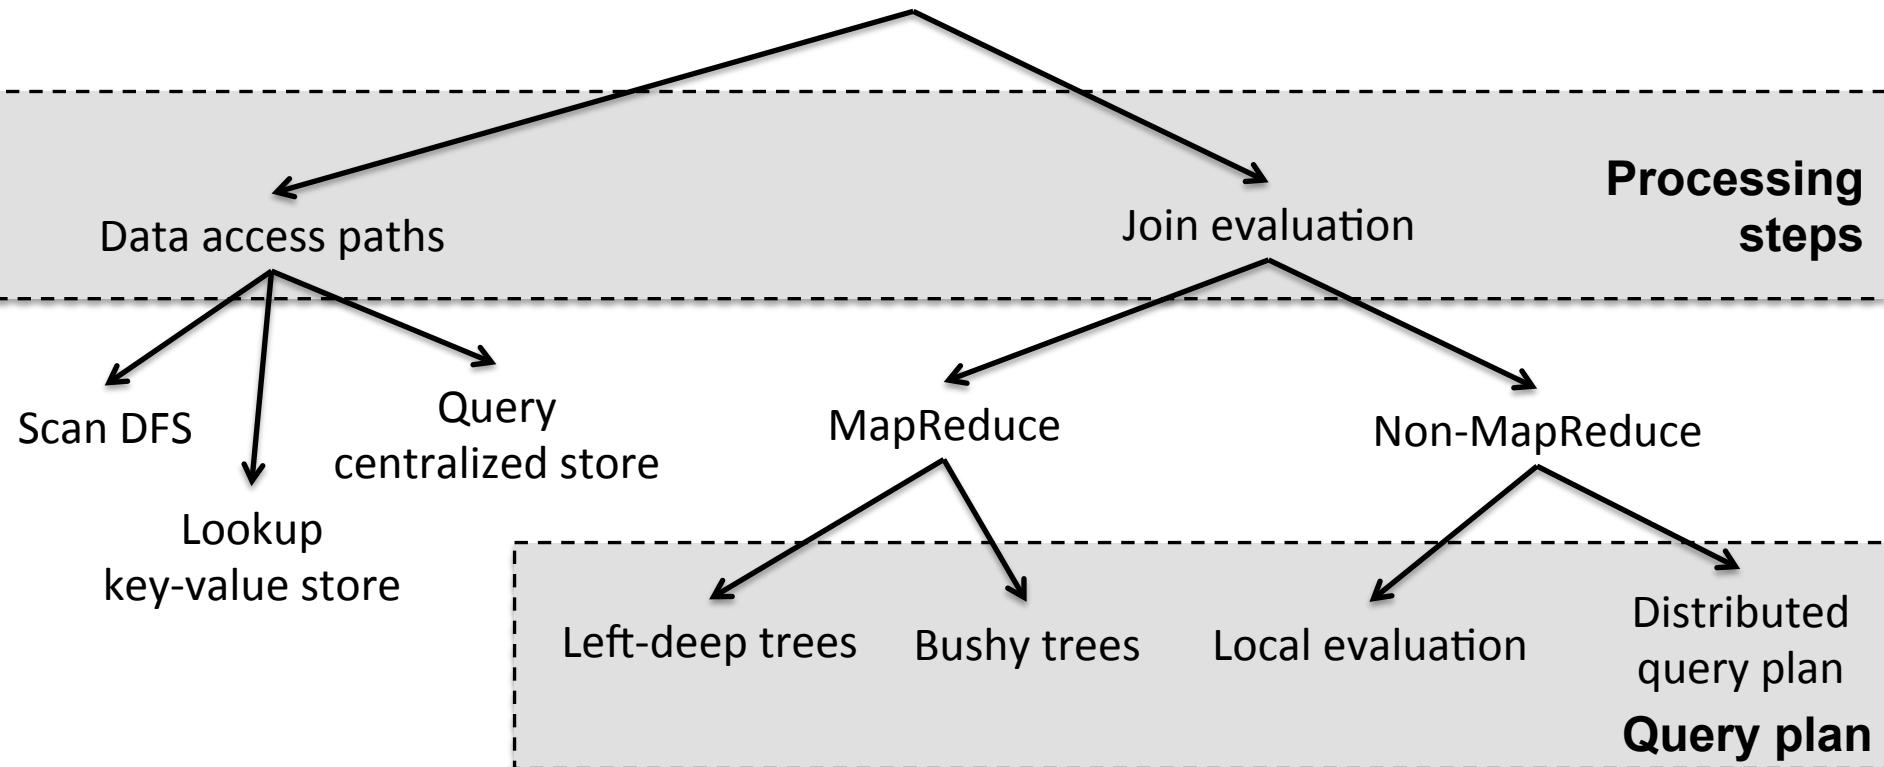


# Query processing



# Query processing

## Relational-style query processing



# Categorization based on storage and query processing

- I. Based on MapReduce and DFS
- II. Based on **key-value** stores
- III. Based on **graph platforms**
- IV. Based on multiple **centralized** (RDF) stores
- V. **Hybrid** systems

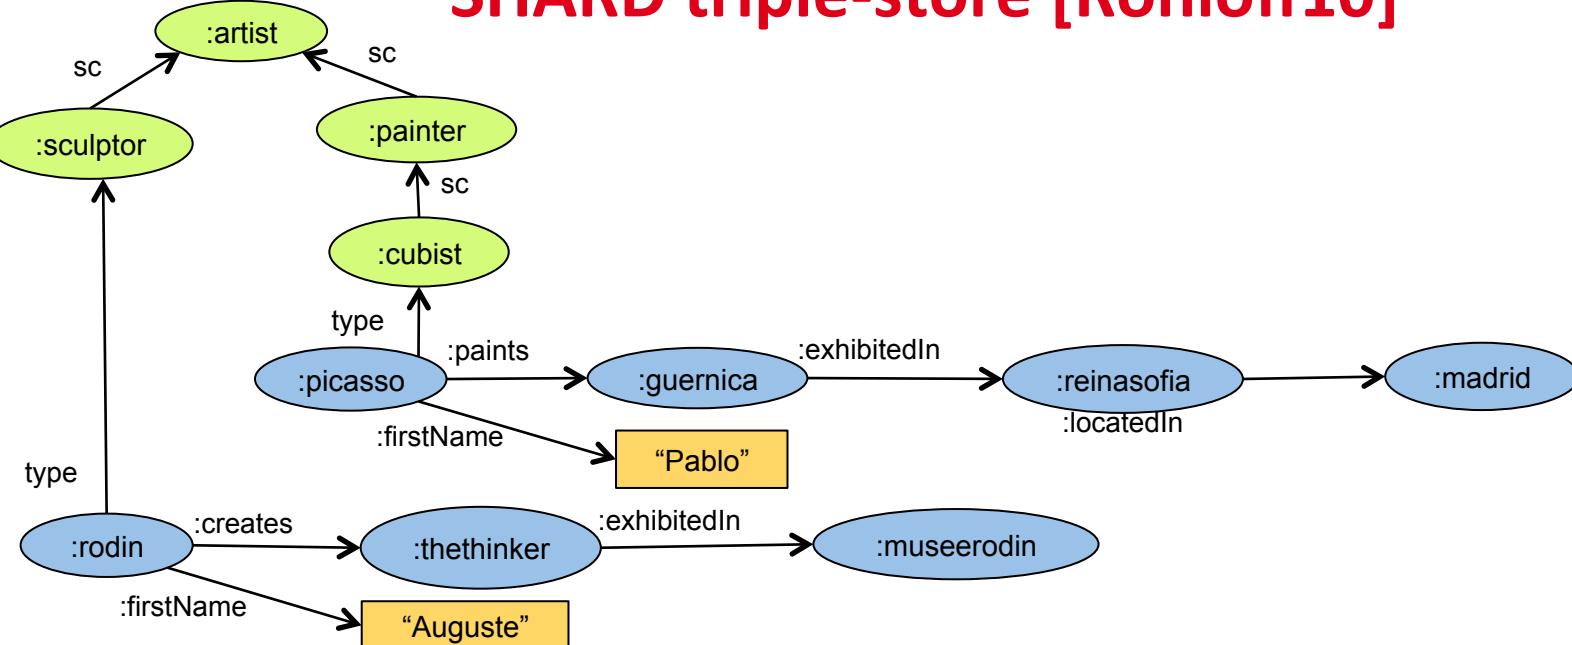
# 3.1

## MapReduce/DFS-based approaches

# I. MapReduce/DFS-based approaches

- Store RDF triples in **DFS** (usually HDFS)
- Translate the query to **MapReduce** jobs
- Run **MapReduce** to answer the queries (usually Hadoop)
- Representative works:
  - SHARD [Rohloff10]
  - HadoopRDF [Husain11]
  - RAPID+ [Ravindra11]
  - PigSPARQL [Schätzle11]
  - TriAD [Gujarada14]

# SHARD triple-store [Rohloff10]



**Storage:** each line holds all triples about a given subject

File1.rdf

```
picasso type :cubist :paints :guernica :firstName "Pablo"
guernica :exhibitedIn :reinasofia
reinasofia :locatedIn :madrid
rodin type :sculptor :firstName
"Auguste" :creates :thethinker
thethinker :exhibitedIn :museerodin
```

File2.rdf

```
:sculptor sc :artist
:cubist sc :painter
:painter sc :artist
```

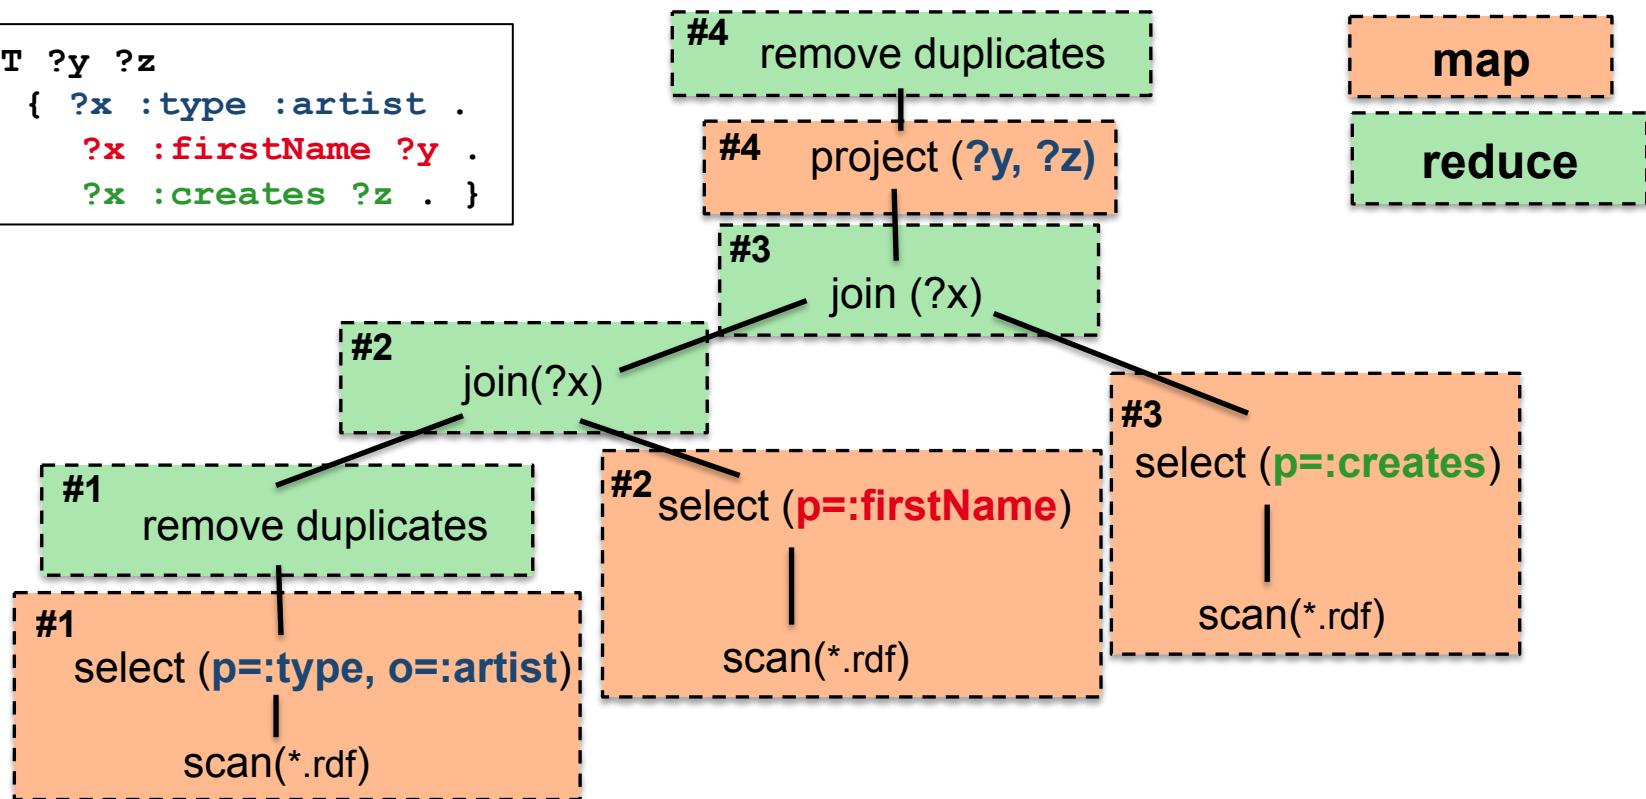
HDFS

# SHARD triple-store [Rohloff10]

Query processing: Iterative MapReduce jobs

- 1 MapReduce job for each triple pattern + 1 for the projection at the end
- Map phase: **match** triple pattern by scanning all triples
- Reduce phase: **join** triple pattern with intermediate results from previous ones

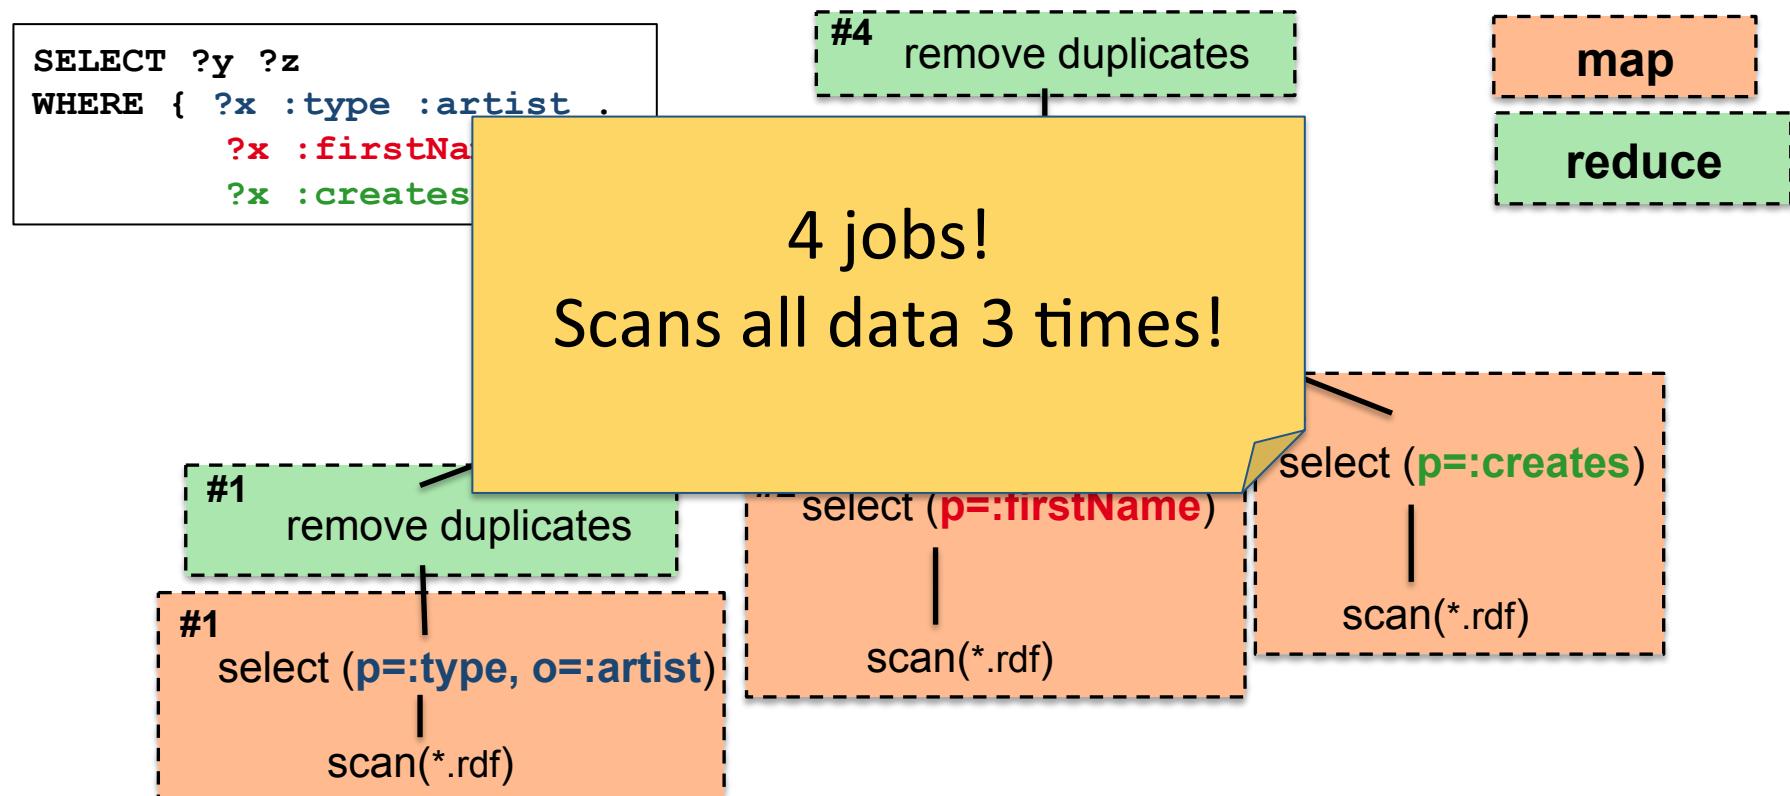
```
SELECT ?y ?z
WHERE { ?x :type :artist .
         ?x :firstName ?y .
         ?x :creates ?z . }
```



# SHARD triple-store [Rohloff10]

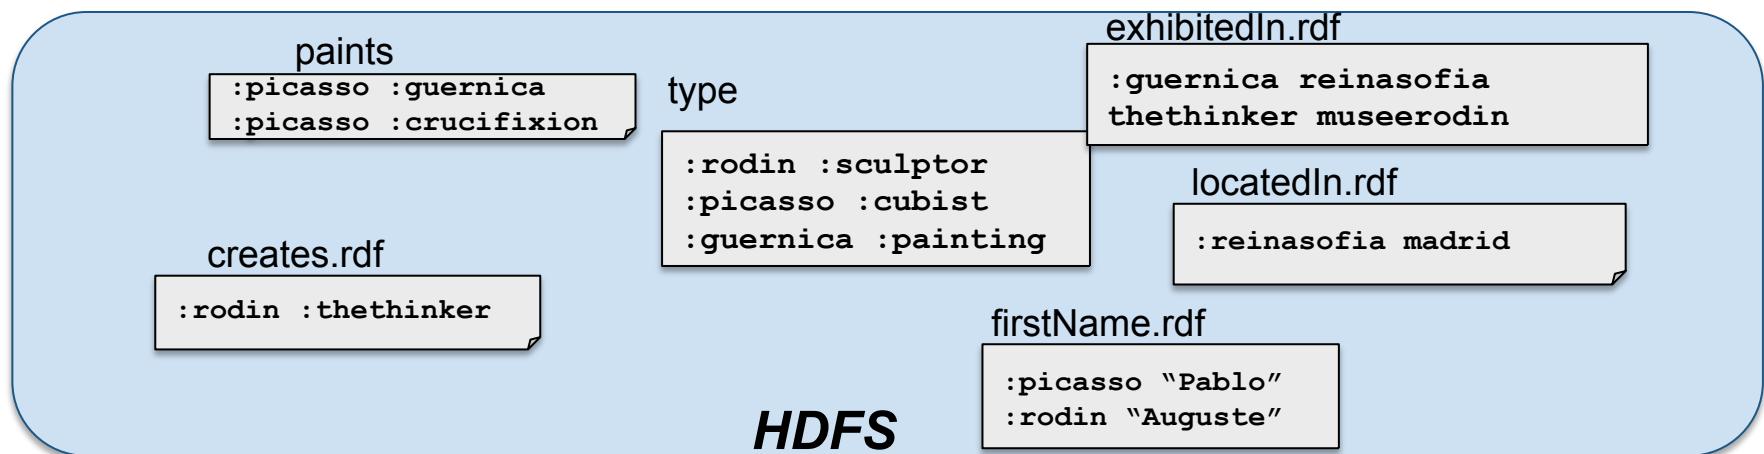
Query processing: Iterative MapReduce jobs

- 1 MapReduce job for each triple pattern + 1 for the projection at the end
- Map phase: **match** triple pattern by scanning all triples
- Reduce phase: **join** triple pattern with intermediate results from previous ones



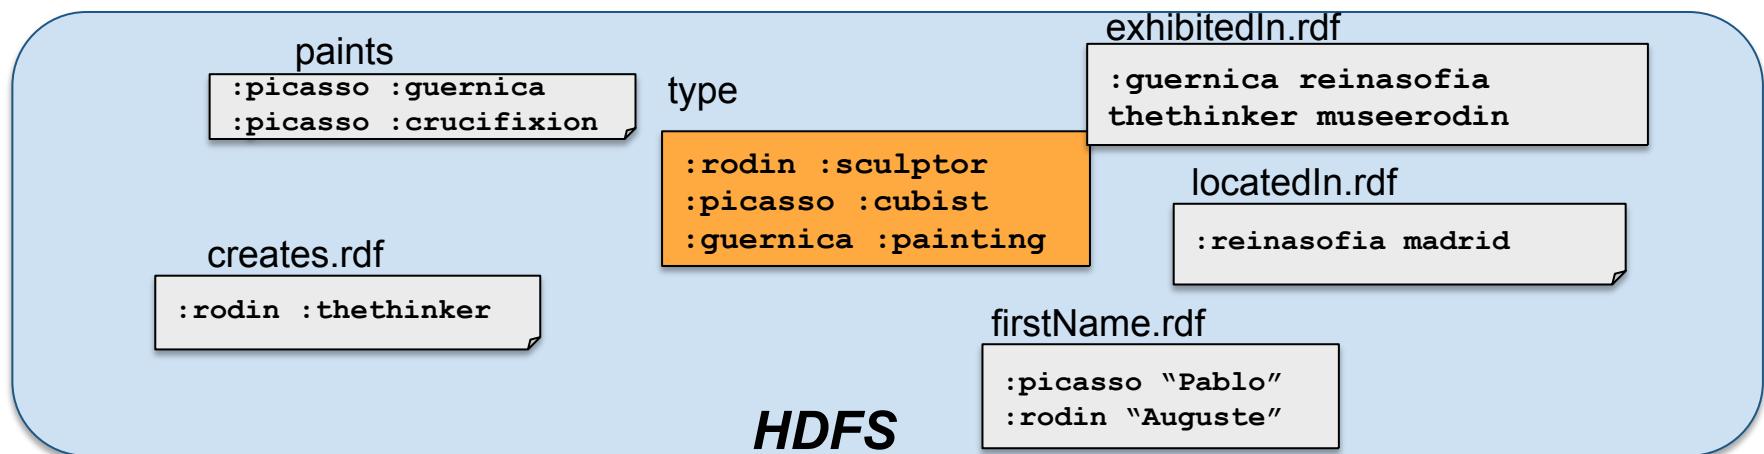
# HadoopRDF - storage [Husain11]

- RDF triples grouped by the **property names** in files (vertical partitioning of centralized RDF stores)



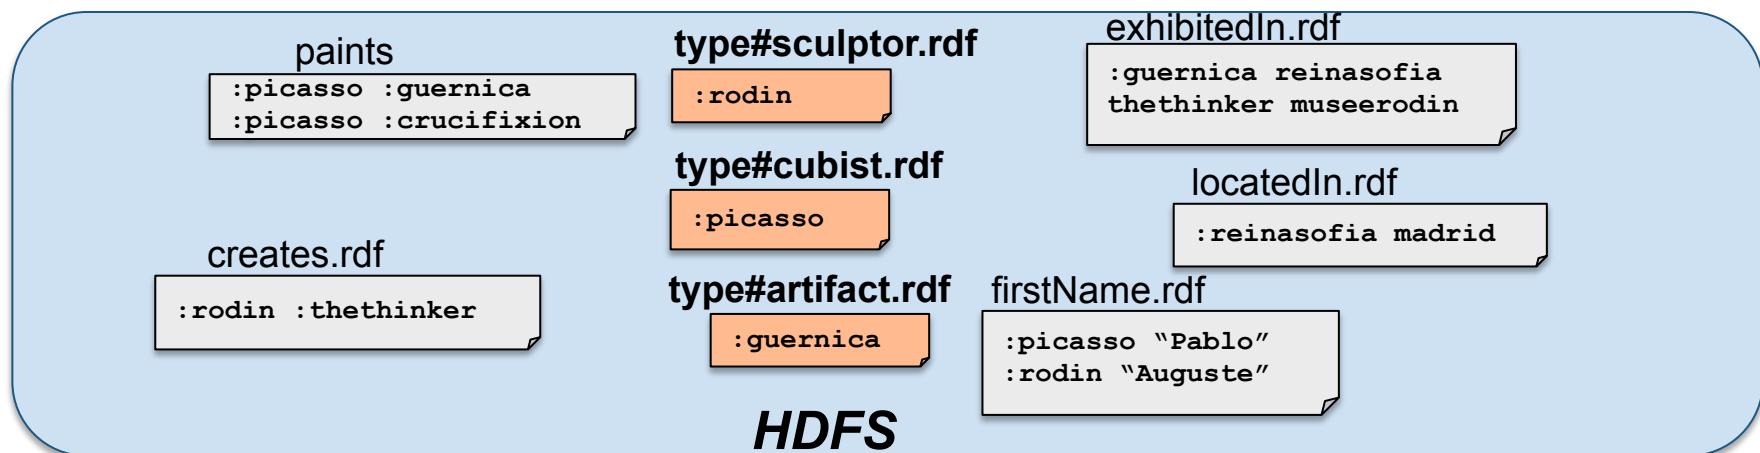
# HadoopRDF - storage

- RDF triples grouped by the **property names** in files (vertical partitioning of centralized RDF stores)
- The file for the property **type** partitioned based on the object value



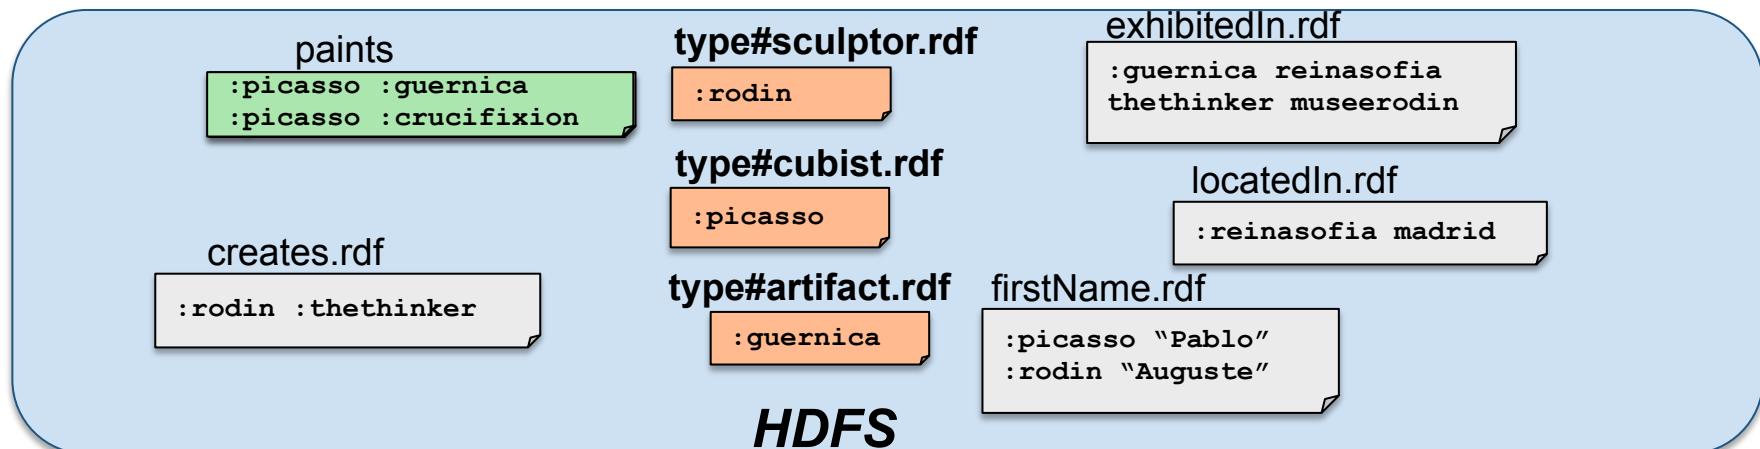
# HadoopRDF - storage

- RDF triples grouped by the **property names** in files (vertical partitioning of centralized RDF stores)
- The file for the property **type** partitioned based on the object value



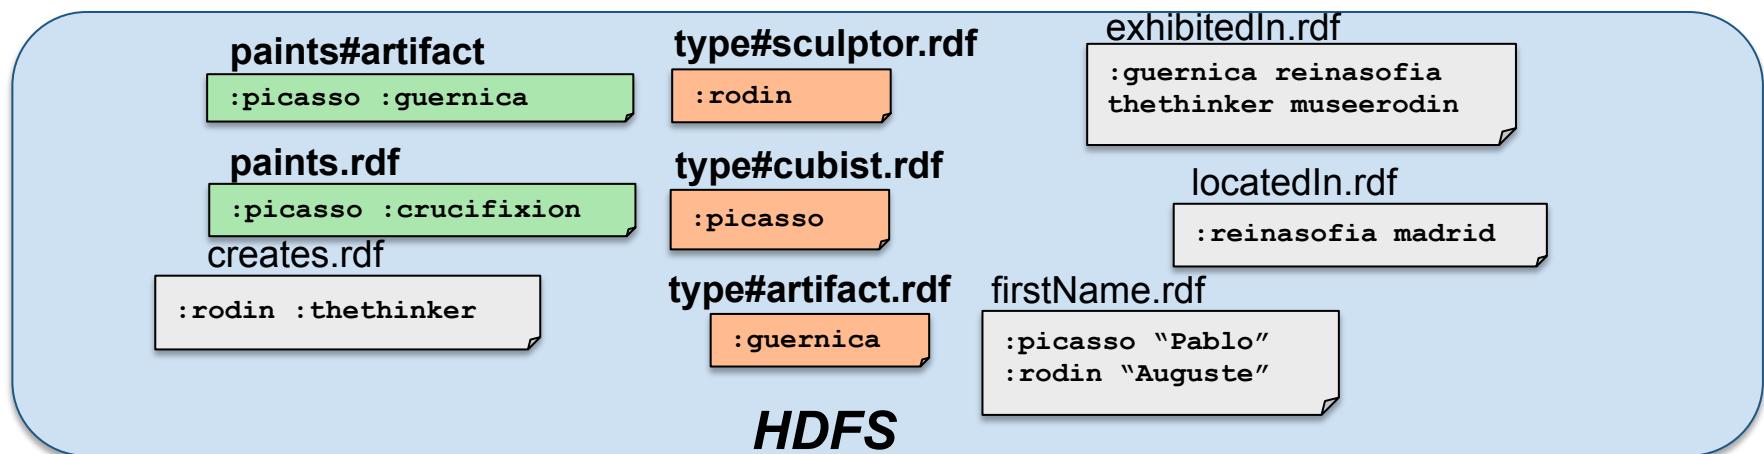
# HadoopRDF storage

- RDF triples grouped by the **property names** in files (vertical partitioning of centralized RDF stores)
- The file for the property **type** partitioned based on the object value
- All other property files partitioned by their **object type** (if there is such information)



# HadoopRDF storage

- RDF triples grouped by the **property names** in files (vertical partitioning of centralized RDF stores)
- The file for the property **type** partitioned based on the object value
- All other property files partitioned by their **object type** (if there is such information)



# HadoopRDF storage features

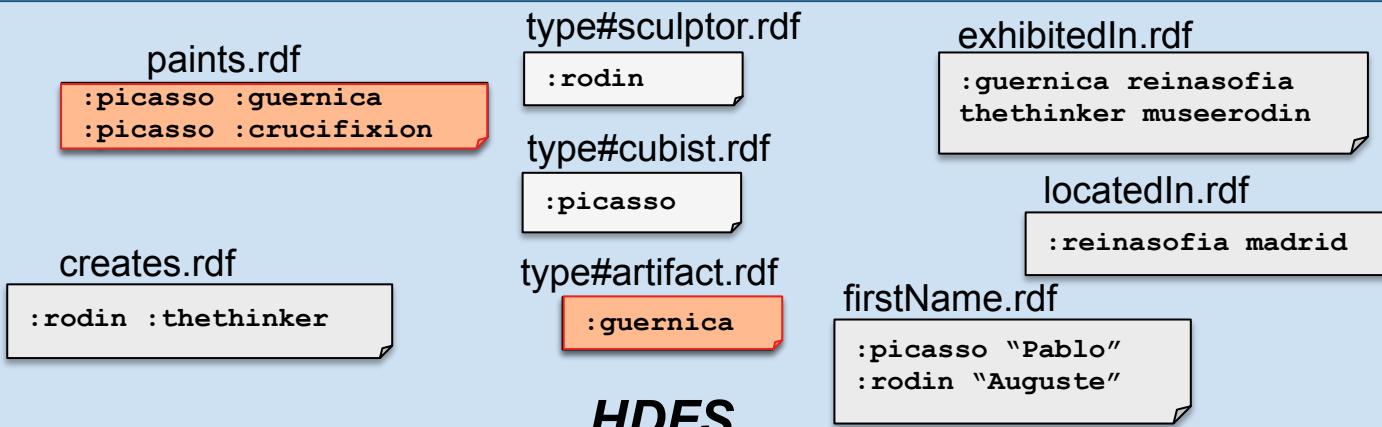
- Space gains
- Query optimization (minimization?)

Step	Files	Size (GB)	Space Gain
N-Triples	20020	24	-
PS	17	7.1	70.42%
POS	41	6.6	7.04%

```
SELECT ?x
WHERE { ?x :paints ?y .
         ?y type :artifact . }
```

Without object value partitioning, we would have to scan files:

{paints.rdf, type#artifact.rdf}



# HadoopRDF storage features

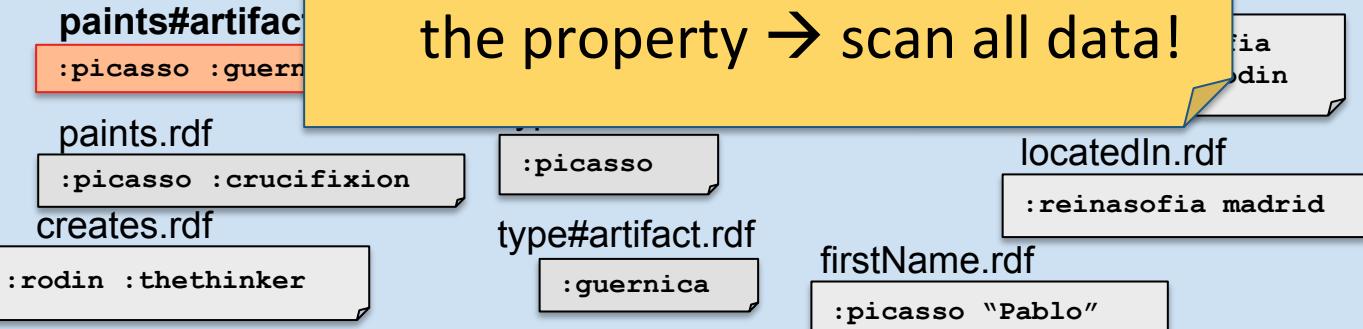
- Space gains
- Query optimization

```
SELECT ?x
WHERE { ?x :paints ?y .
         ?y type :artifact . }
```

Step	Files	Size (GB)	Space Gain
N-Triples	20020	24	-
PS	17	7.1	70.42%
POS	41	6.6	7.04%

With object value partitioning, we only need to scan: {paints#artifact.rdf}

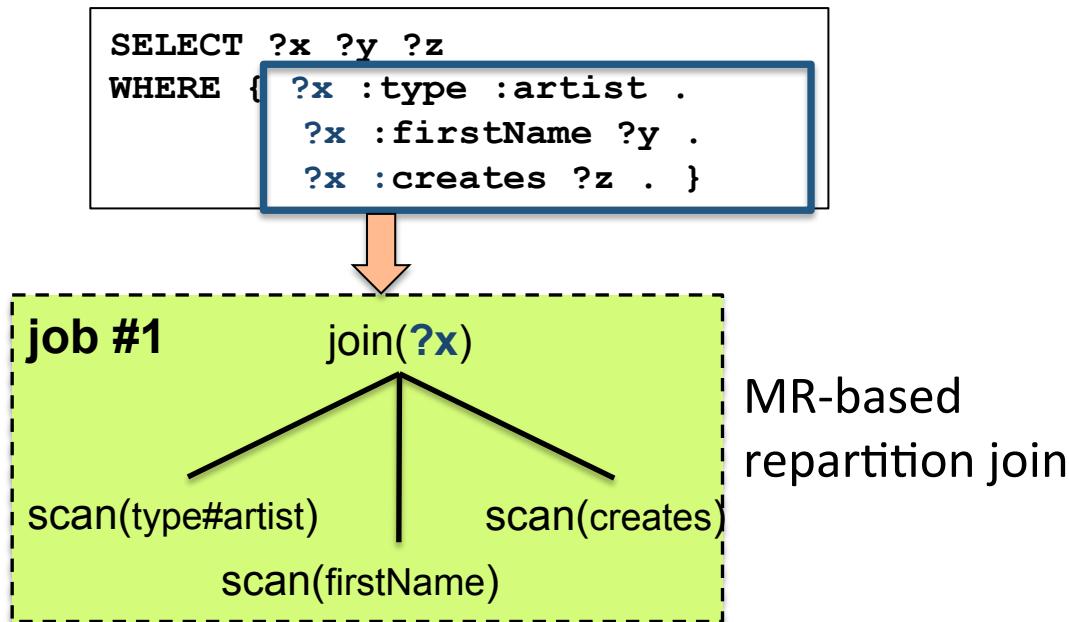
Issue:  
For queries with a variable in  
the property → scan all data!



**HDFS**

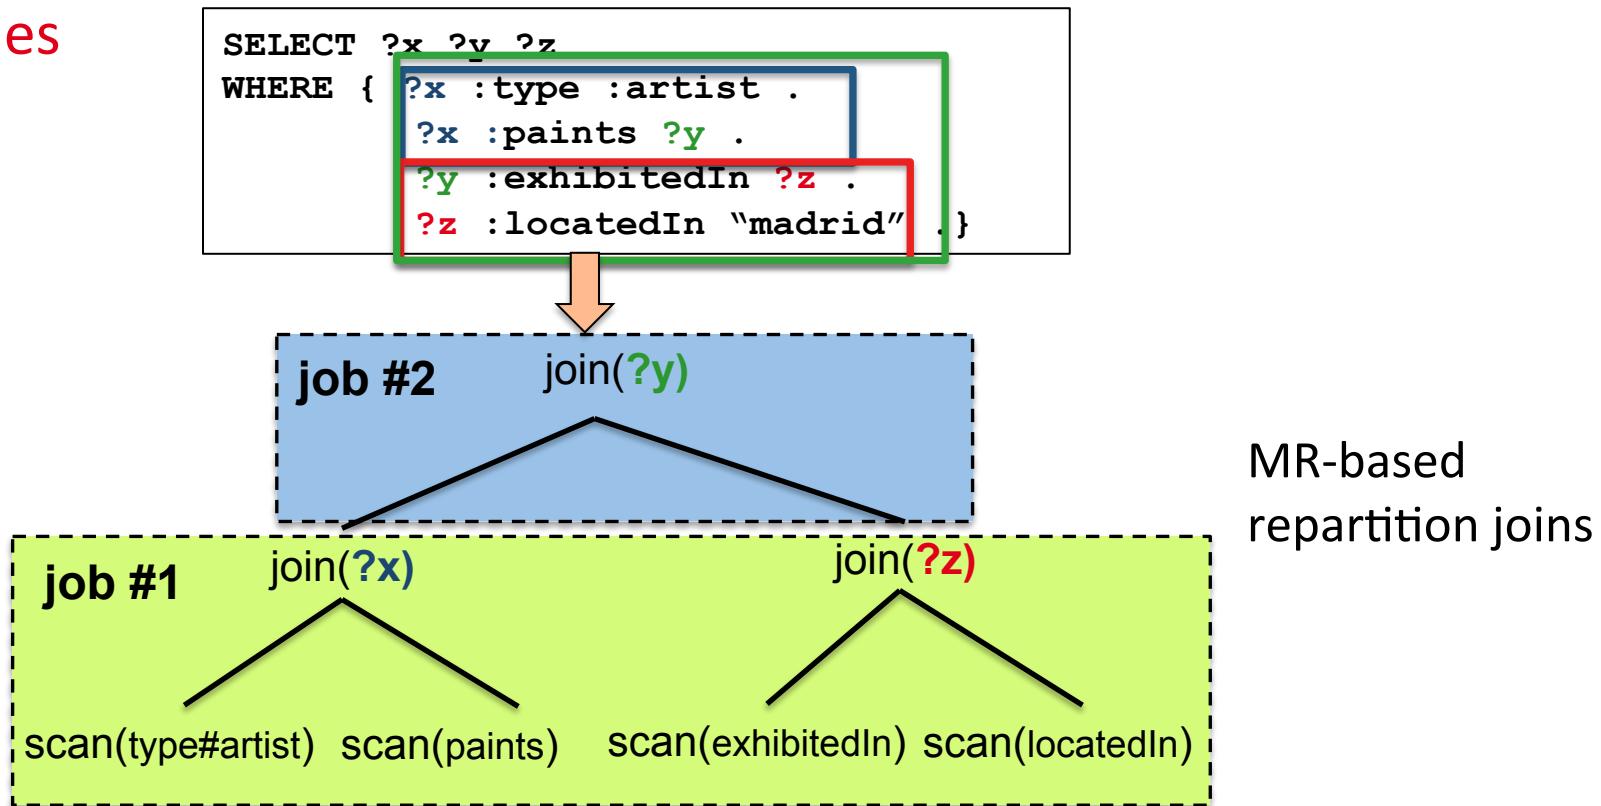
# HadoopRDF - querying

- File selection based on the property (and object) of the triple
- Joining two or more triple patterns on the **same variable**



# HadoopRDF querying

- File selection based on the property (and object) of the triple
- Joining two or more triple patterns on the **same variable**
- Several joins may be executed **in parallel** if they are on **different variables**



# HadoopRDF query processing

*How to decompose a query in elementary operations ?*

*What query plan to build ?*

- Observations
  - Significant overhead when running a job
  - Traditional selectivity-based optimization techniques may result to plans with more jobs → bigger response times

Dataset	2 Job Plan	3 Job Plan	Difference
LUBM_10000	4920	9180	4260
LUBM_20000	31020	36540	5520
LUBM_30000	80460	93947	13487

- Heuristic: Query **decomposition** to **minimize** the number of Hadoop **jobs**

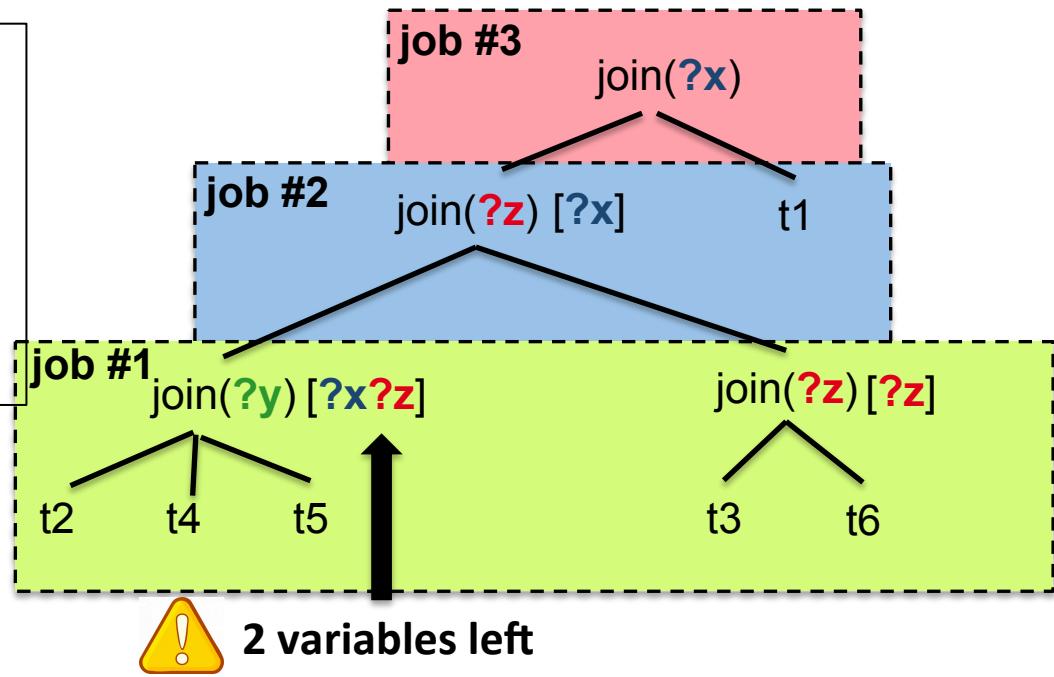


# HadoopRDF query planning

- Greedy algorithm
- As many joins as possible in one job
- Each triple pattern is used only once in each join
- Leave the least number of remaining variables to be joined

```
SELECT ?x ?y ?z
WHERE { ?x type :artist .          (t1)
        ?y type :artifact .        (t2)
        ?z type :museum .         (t3)
        ?x :paints ?y .          (t4)
        ?y :exhibitedIn ?z .    (t5)
        ?z :locatedIn "madrid" (t6)
}
```

“Complete variable elimination”  
of  $\text{?y}$

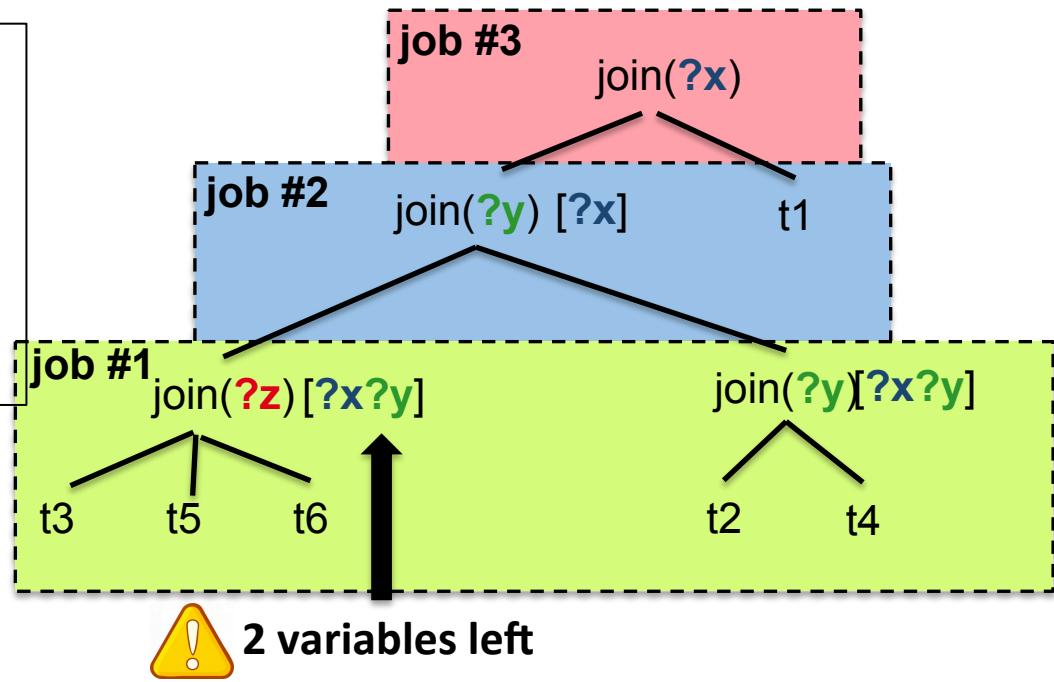


# HadoopRDF query planning

- Greedy algorithm
- As many joins as possible in one job
- Each triple pattern is used only once in each join
- Leave the least number of remaining variables to be joined

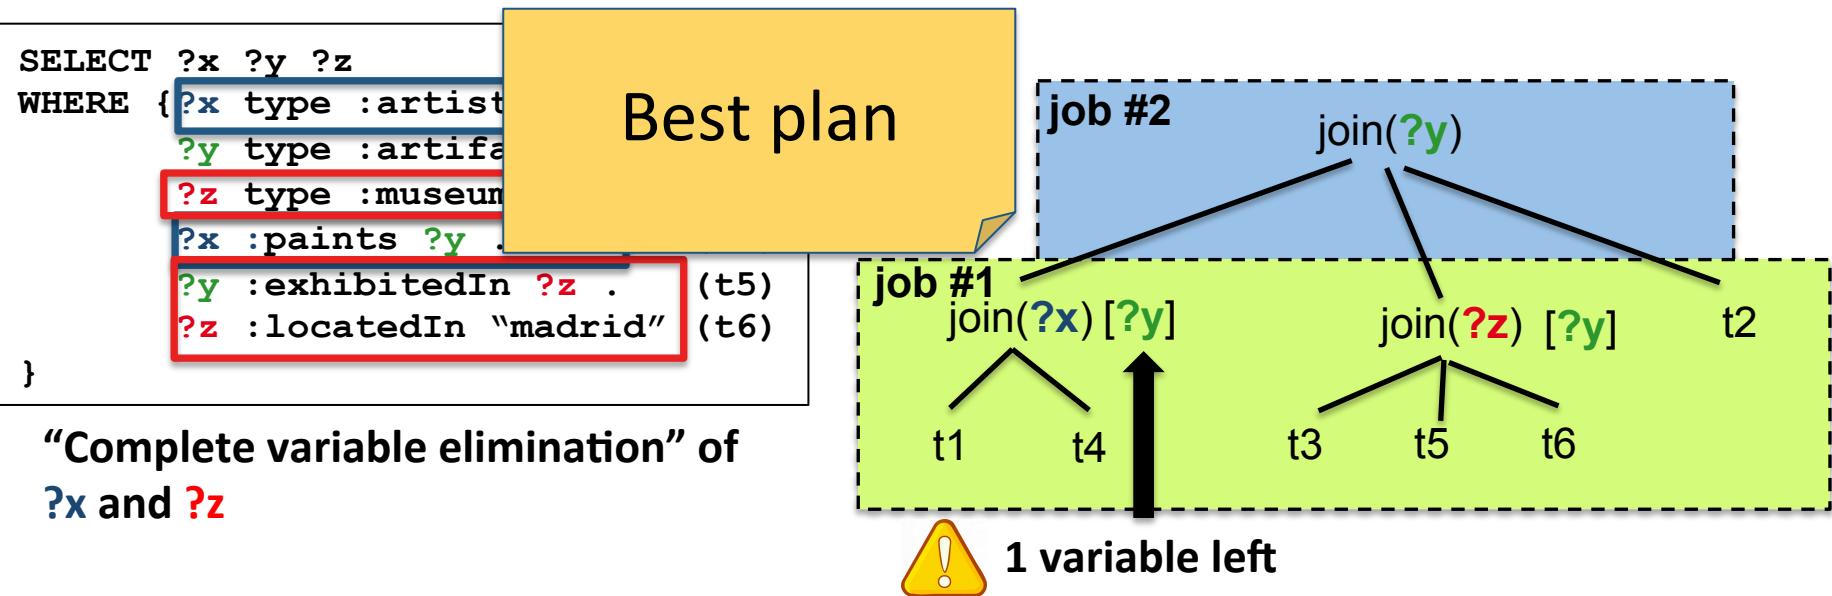
```
SELECT ?x ?y ?z
WHERE { ?x type :artist .          (t1)
        ?y type :artifact .        (t2)
        ?z type :museum .         (t3)
        ?x :paints ?y .          (t4)
        ?y :exhibitedIn ?z .     (t5)
        ?z :locatedIn "madrid" . (t6)
}
```

“Complete variable elimination”  
of  $\text{?z}$



# HadoopRDF query planning

- Greedy algorithm
- As many joins as possible in one job
- Each triple pattern is used only once in each join
- Leave the least number of remaining variables to be joined



## TriAD [Gurajada14]

Based not on MapReduce but on custom shared-nothing parallel processing framework using *asynchronous message passing*

- Advantage: asynchronous parallelism (join in stage N does not need to wait for join in stage N-1)

Storage: METIS partitioning (disjoint subgraphs)

Query processing:

- Custom distributed nested loops-, hash- and merge-joins
  - Layers « like MapReduce » but asynchronous execution
- Also:
  - Intra-node multi-threading
  - Optimizations to reduce intermediary results, ...

# Summary: MapReduce/DFS-based approaches

Pros	Cons
Scalability	No indexes – scan all data
Fault-tolerance	Shuffling of data
Fast data uploading	Semi-support for joins
Fast and easy to implement	Big initialization overhead
Good for “analytical” queries	High query response times

# 3.2

## Approaches based on key-value stores

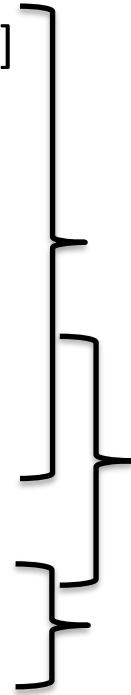
## II. Approaches based on key-value stores

- Store RDF triples in a key-value store
  - Indexes on RDF triples
- Lookups in the key-value store for each query triple pattern
- How to perform the joins?
  - Centralized join
  - Use MapReduce
  - Graph exploration

## II. Approaches based on key-value stores

- Representative works

- CumulusRDF [Ladwig11]
- Stratostore [Stein10]
- Rya [Punnoose12]
- AMADA [Bugiotti12]
- H2RDF [Papailiou12]
- H2RDF+[Papailiou13]
- MAPSIN [Schätzle12]
- Trinity.RDF [Zeng13]



**Local  
query processing**

**MapReduce-based  
query processing**

**Graph exploration**

# Storing RDF in key-value stores

- Rya uses Apache Accumulo
  - SPO|-|- POS|-|- OSP|-|-
- H2RDF uses Apache HBase
  - SP|O|- PO|S|- OS|P|-
- AMADA uses Amazon DynamoDB
  - S|P|O P|O|S O|S|P
- MAPSIN uses Apache HBase
  - S|P|O|- O|P|S|-
- Stratustore uses Amazon SimpleDB
  - S|P|O
- CumulusRDF uses Apache Cassandra
  - hierarchical layout: S|P|O|- P|O|S|- O|S|P|-
  - flat layout: S|PO|- PO|S|- PO|“P”| P O|SP|-

**Sorted Index**

**Sorted Index**

**Hash Index**

**Sorted Index**

**Hash Index**

**Hash Index    Sorted Index**

# Storing RDF in key-value stores

- Rya uses Apache Accumulo
  - SPO|-|- POS|-|- OSP|-|-
- H2RDF uses Apache HBase
  - SP|O|- PO|S|- OS|P|-
- AMADA uses Am
  - S|P|O P|O|S|P|
- MAPSIN uses Apa
  - S|P|O|- O|S|P|
- Stratustore uses Amazon SimpleDB
  - S|P|O
- CumulusRDF uses Apache Cassandra
  - hierarchical layout: S|{P}O|- P|{O}S|- O|{S}P|-
  - flat layout: S|PO|- PO|S|- PO|“P”| P O|SP|-

3 indexes:  
SPO, POS, OSP

# Rya storage - example

SPO		
Key	(attribute, value)	
:picasso, :firstName, "Pablo"	-	
:picasso, :p	Key	(attribute, value)
:picasso, ty	:exhibitedIn, :reinasofia, :guernica	-
:guernica, :	:firstName, "Pablo", :picasso,	-
...		
:paints, :gu	Key	(attribute, value)
type, :cubis	:cubist, :picasso, type,	-
...	:guernica, :picasso, :paints	-
"Pablo", :picasso, :firstName	-	
:reinasofia, :guernica, :exhibitedIn	-	
...	-	

POS

OSP

# Aggressive indexing in key-value stores

- H2RDF+ (Apache HBase) **6** indices + **12** aggregated ones:

SPO| -|-  
POS| -|-  
OSP| -|-  
PSO| -|-  
OPS| -|-  
SOP| -|-

SP | cnt(O)| -  
PO | cnt(S)| -  
OS | cnt(P)| -

S | cnt(P)| avg(O)  
P | cnt(O)| avg(S)  
O | cnt(S)| avg(P)  
(S) | avg(O)  
(P) | avg(S)  
(O) | avg(P)

aggressive compression



merge joins

selectivity estimation

# Strong point of key-value stores: selective data access

## Different access paths for triple pattern matching

Triple pattern	Rya (Accumulo)	H2RDF (HBase)	AMADA (DynamoDB)	MAPSIN (HBase)
(s, p, o)	any lookup	any lookup + select	any lookup + select	any lookup + sselect
(s, p, ?o)	SPO range scan	SP O lookup	S P O lookup + select	S P O lookup + sselect
(s, ?p, o)	OSP range scan	OS P lookup	O S P lookup + select	O S P lookup + sselect
(s, ?p, ?o)	SPO range scan	SP O range scan	S P O lookup	S P O lookup
(?s, p, o)	POS range scan	PO S lookup	P O S lookup + select	O S P lookup + sselect
(?s, p, ?o)	POS range scan	PO S range scan	P O S lookup	any scan + sselect
(?s, ?p, o)	OSP range scan	OS P range scan	O S P lookup	O S P lookup
(?s, ?p, ?o)	any scan	any scan	any scan	any scan

select: filter on client side

sselect: filter on server side

# Weak point of key-value stores: joins

Key-value stores do not support **joins** →

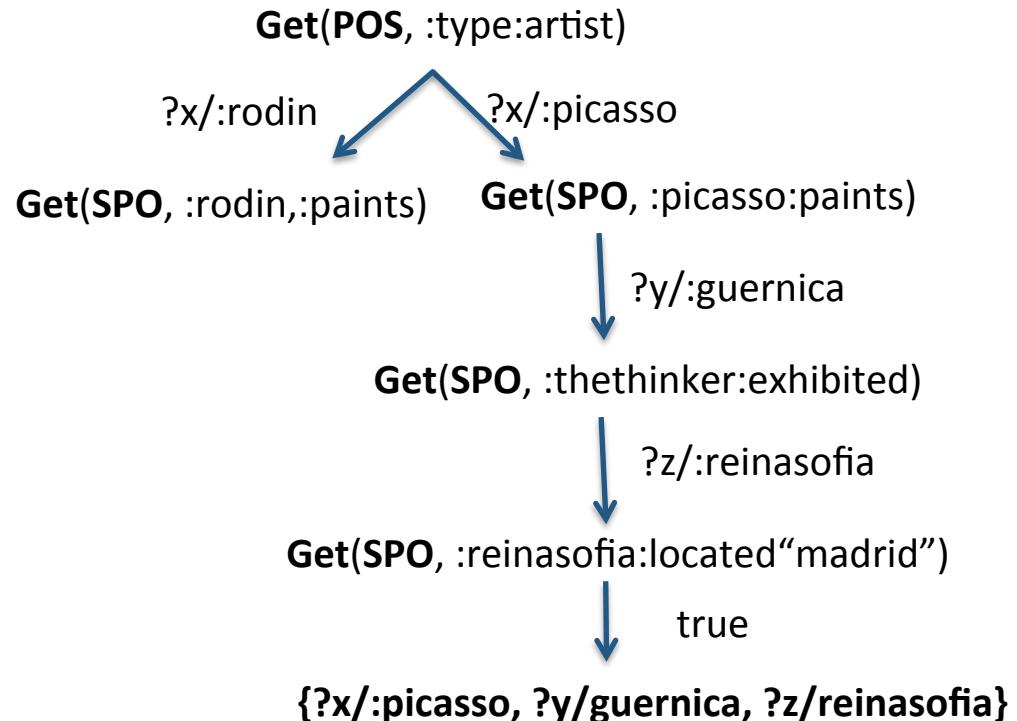
Look-ups in the store, joins out of the store

- **Rya**: index nested loops by query rewriting and lookups
- **H2RDF+**: 2 strategies
  - selective joins → centralized – sort-merge joins
  - non-selective → MapReduce (merge-join)
- **AMADA**: memory hash join
- **MAPSIN**: MapReduce (map phase only)
- **Stratustore**: SimpleDB SELECT query for each star-join and rest of the joins locally
- **CumulusRDF**: Only single triple pattern queries are supported

# Rya joining strategy

- Index nested loops
- A lookup for the 1st triple pattern
- For each triple pattern  $t_i$ , use the  $n$  results from triple pattern  $t_{i-1}$  to rewrite  $t_i$  and then perform  $n$  lookups for the rewritten  $t_i$

```
SELECT ?x ?y ?z
WHERE { ?x :type :artist .
         ?x :paints ?y .
         ?y :exhibitedIn ?z
         ?z :locatedIn "madrid" . }
```

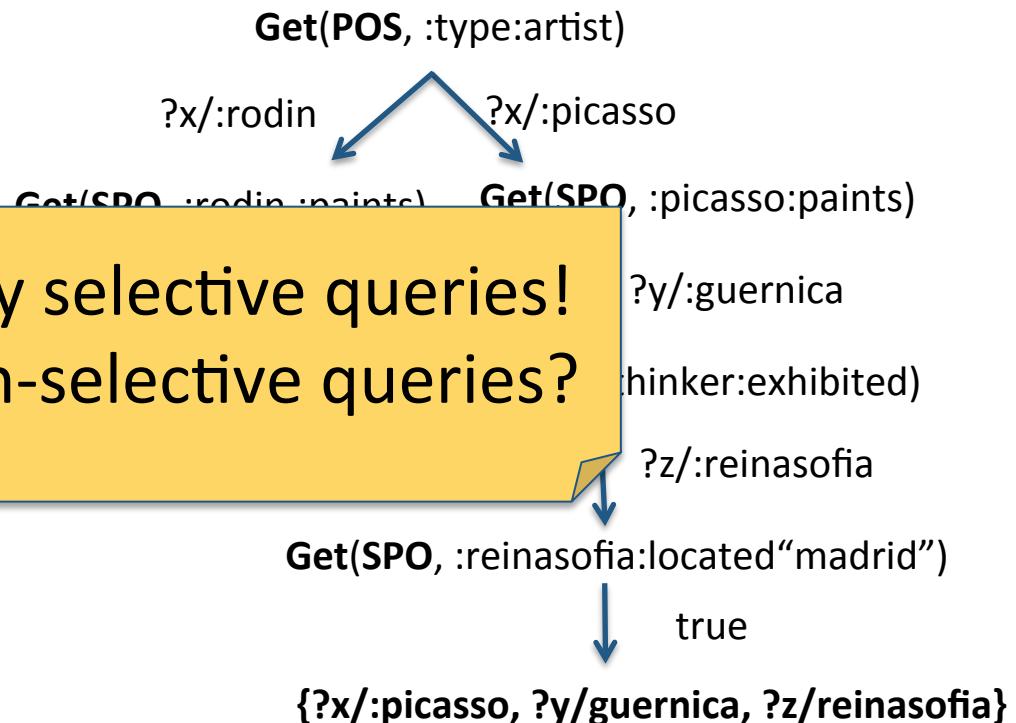


# Rya joining strategy

- Index nested loops
- A lookup for the 1st triple pattern
- For each triple pattern  $t_i$ , use the  $n$  results from triple pattern  $t_{i-1}$  to rewrite  $t_i$  and then perform  $n$  lookups for the rewritten  $t_i$

```
SELECT ?x ?y ?z  
WHERE { ?x :type :artist .  
        ?x :paints ?y .  
        ?y :exhibitedIn ?z  
        ?z :1 }
```

Efficient for very selective queries!  
What about non-selective queries?

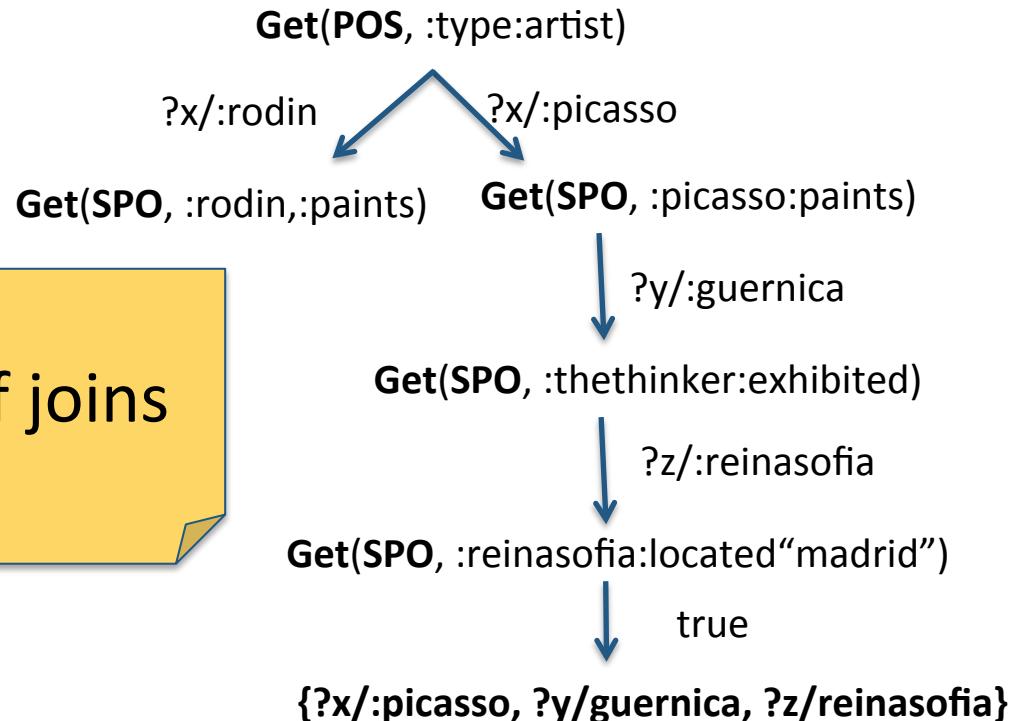


# Rya: optimizations

- Index nested loops
- A lookup for the 1st triple pattern
- For each triple pattern  $t_i$ , use the  $n$  results from triple pattern  $t_{i-1}$  to rewrite  $t_i$  and then perform  $n$  lookups for the rewritten  $t_i$

```
SELECT ?x ?y ?z
WHERE { ?x :type :artist .
         ?x :paints ?y .
         ?y :exhibitedIn ?z
         ?z :locatedIn "madrid" . }
```

## 1. Parallelization of joins

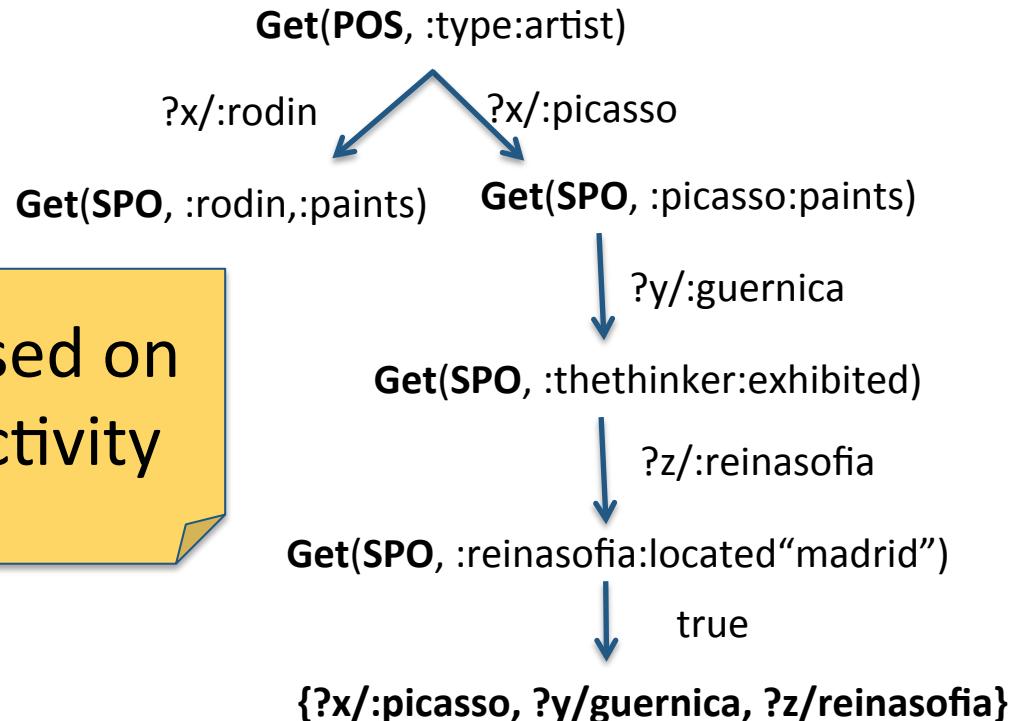


# Rya: optimizations

- Index nested loops
- A lookup for the 1st triple pattern
- For each triple pattern  $t_i$ , use the  $n$  results from triple pattern  $t_{i-1}$  to rewrite  $t_i$  and then perform  $n$  lookups for the rewritten  $t_i$

```
SELECT ?x ?y ?z
WHERE { ?x :type :artist .
         ?x :paints ?y .
         ?y :exhibitedIn ?z
         ?z :locatedIn "madrid". }
```

2. Join ordering based on triple pattern selectivity



# Summary: Approaches based on key-value stores

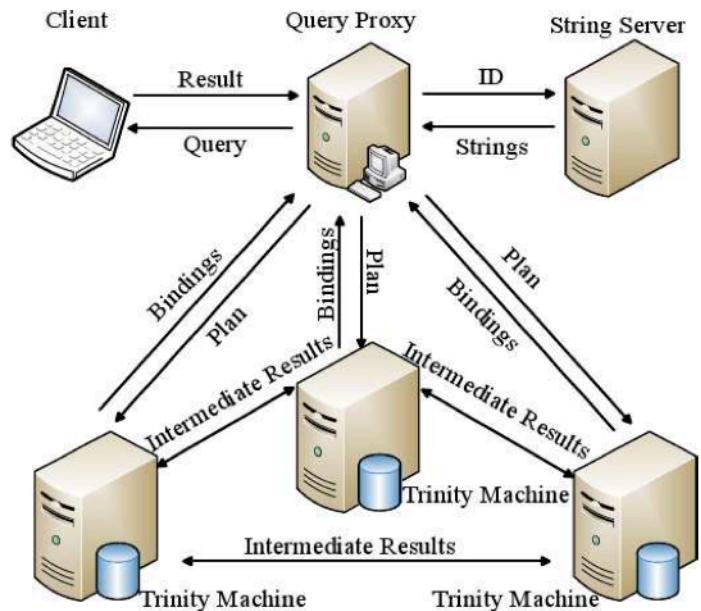
Pros	Cons
Scalability	Low schema expressivity
Fault-tolerance	Joins implemented by the user
Very fast lookups	Joins on the client-side
Efficient for selective queries	

# 3.3

## Approaches based on graph platforms

# Trinity.RDF [Zeng13]

- Built on top of the Trinity in-memory key-value store
- Takes advantage of the **graph-based model** of RDF for **query processing**
- **Graph exploration** instead of joins
  - greedy pruning of candidate matches
  - only if graph exploration can be implemented more efficiently than join
  - similar to index-nested-loops join



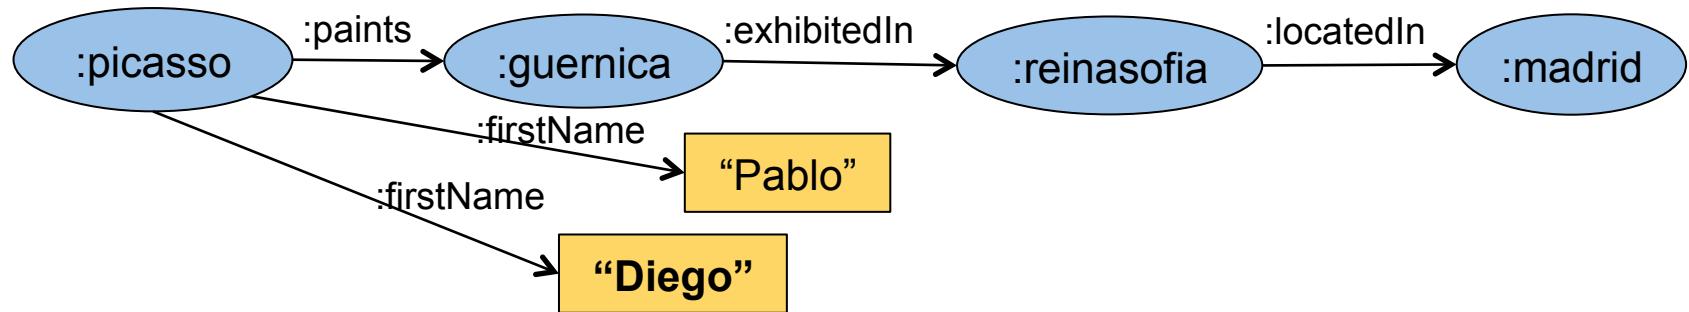
# Trinity.RDF: RDF graph storage alternatives

- Each node is assigned to a machine through hashing
- Two modeling choices
  - 1. (**key**, **value**) : (**node\_id**,  $\langle \text{in-adjacency-list}, \text{out-adjacency-list} \rangle$ )  
(property, subject)      (property, object)

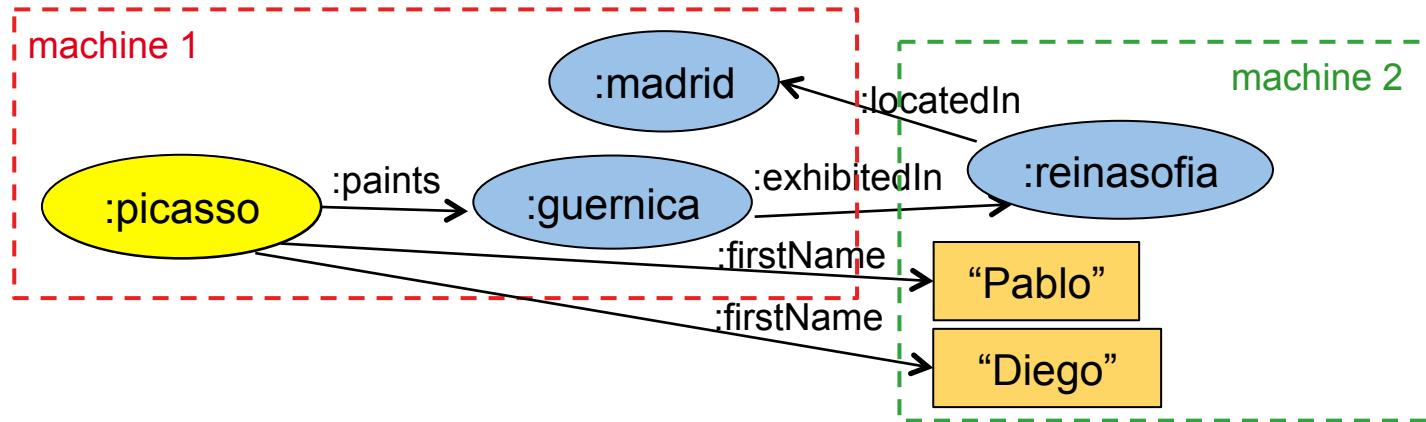


$(n_1, \langle \{(l_1, n_0)\} \{(l_2, n_2)\} \rangle)$   
 $(n_0, \langle \{-\} \{(l_1, n_1)\} \rangle)$   
 $(n_2, \langle \{(l_2, n_2)\} \{-\} \rangle)$

# Trinity.RDF: RDF graph storage alternatives (1)



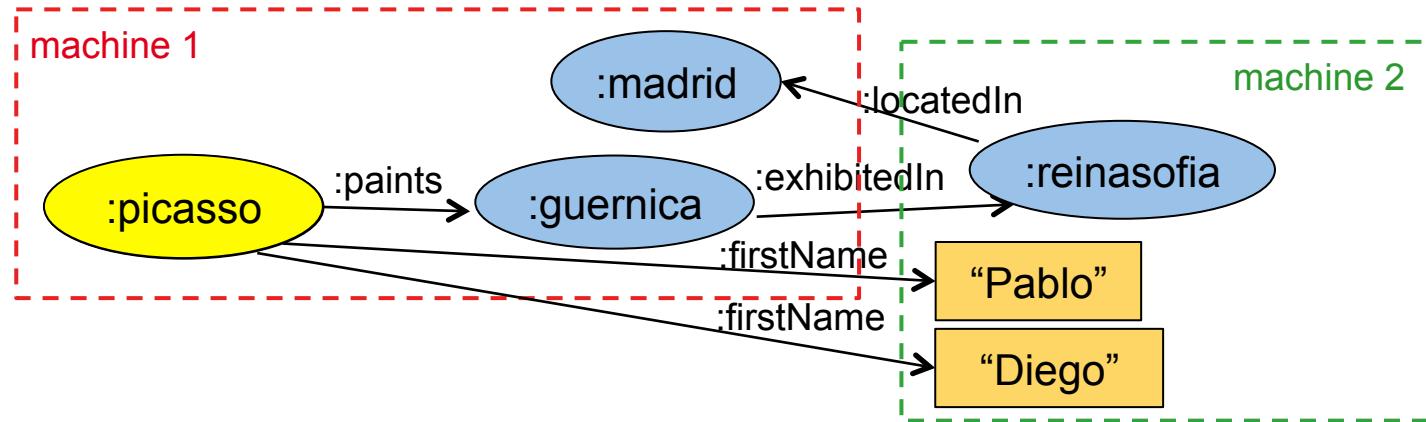
# Trinity.RDF: RDF graph storage alternatives (1)



Key	Value
:picasso	In: - Out: (:paints, :guernica) Out: (:firstName, "Pablo") Out: (:firstName, "Diego")
:guernica	In: (:paints, :picasso) Out: (:exhibitedIn, :reinasofia)
:madrid	Out: (:locatedIn, :reinasofia)

Key	Value
:reinasofia	In: (:exhibitedIn, :guernica) Out: (:locatedIn, :madrid)
"Pablo"	In: (:firstName, :picasso)
"Diego"	In: (:firstName, :picasso)

# Trinity.RDF: RDF graph storage alternatives (1)



machine 1

machine 2

Key	Value
:picasso	In: - Out: (:paints, :guernica) <b>Out: (:firstName, "Pablo")</b> <b>Out: (:firstName, "Diego")</b>
:guernica	In: (:paints, :picasso) <b>Out: (:exhibitedIn, :reinasofia)</b>
:madrid	<b>Out: (:locatedIn, :reinasofia)</b>

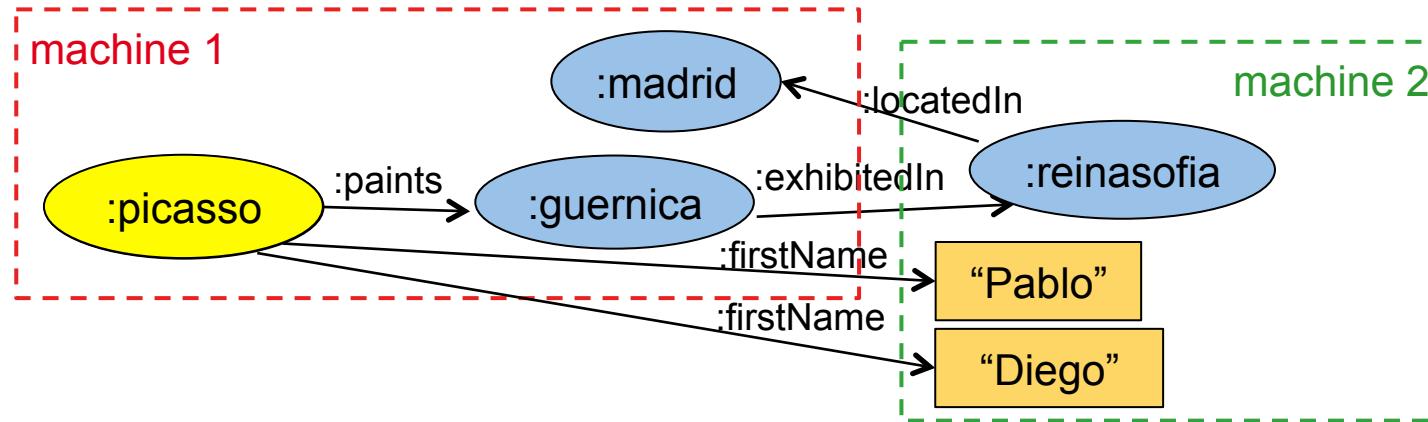
Key	Value
:reinasofia	In: (:exhibitedIn, :guernica) Out: (:locatedIn, :madrid)
"Pablo"	In: (:firstName, :picasso)
"Diego"	<b>In: (:firstName, :picasso)</b>

2 messages should be sent to continue the node exploration from *:picasso* through *:firstName*

## Trinity.RDF: RDF graph storage alternatives (2)

- Each node is assigned to a machine through hashing
- Two modeling choices
  - 1. (**key**, **value**) : (**node\_id**, <in-adjacency-list, out-adjacency-list>)
  - 2. (**key**, **value**): (**node\_id**, <in1, in2..., out1, out2...>)
    - (**in<sub>i</sub>**, **in-adjacency-list<sub>i</sub>**): on the same machine i
    - (**out<sub>i</sub>**, **out-adjacency-list<sub>i</sub>**): on the same machine i

## Trinity.RDF: RDF graph storage alternatives (2)

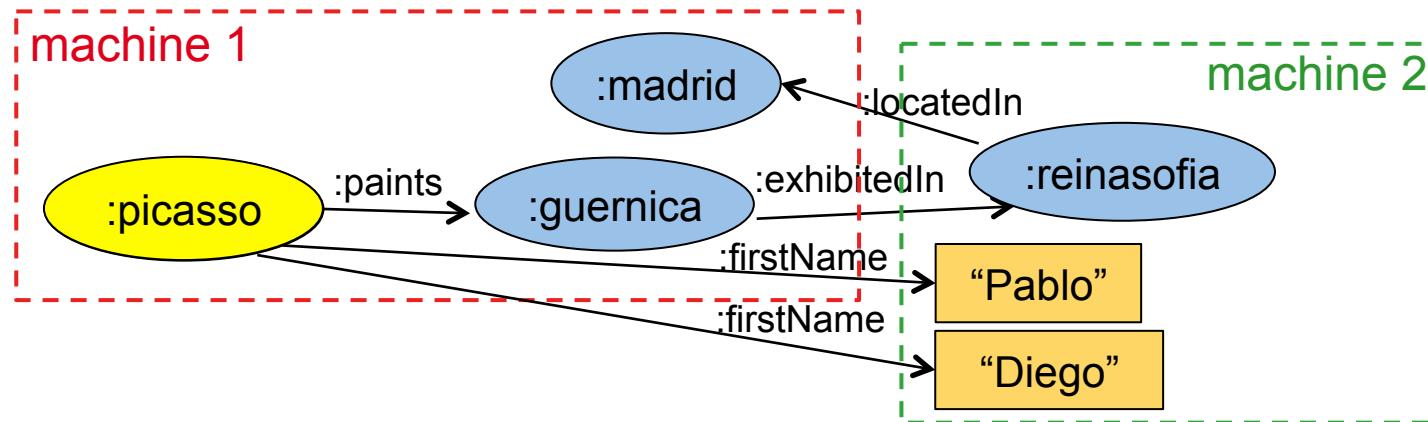


Key	Value
:picasso	out1 out2
:guernica	in1 out3
:madrid	out4
out1	(:paints, :guernica)
in1	(:paints, :picasso)
in2	(:exhibitedIn, :guernica)
out5	(:locatedIn, :madrid)
in3	(:firstName, :picasso)

Key	Value
:reinasofia	in2 out5
"Pablo"	in3
out2	(:firstName, "Pablo") (:firstName, "Diego")
out3	(:exhibitedIn, :reinasofia)
out4	(:locatedIn, :reinasofia)

only 1 message needs to be sent

## Trinity.RDF: RDF graph storage alternatives (2)



Key	Value	Key	Value
:picasso	out1 out2		
:guernica	in1 out3		
:madrid	out4		
out1	(:paints, :guernica)		(:locatedIn, "Pablo")
in1	(:paints, :picasso)		(:firstName, "Diego")
in2	(:exhibitedIn, :guernica)		(:exhibitedIn, :reinasofia)
out5	(:locatedIn, :madrid)		(:locatedIn, :reinasofia)
in3	(:firstName, :picasso)		

No more than #machines messages are sent

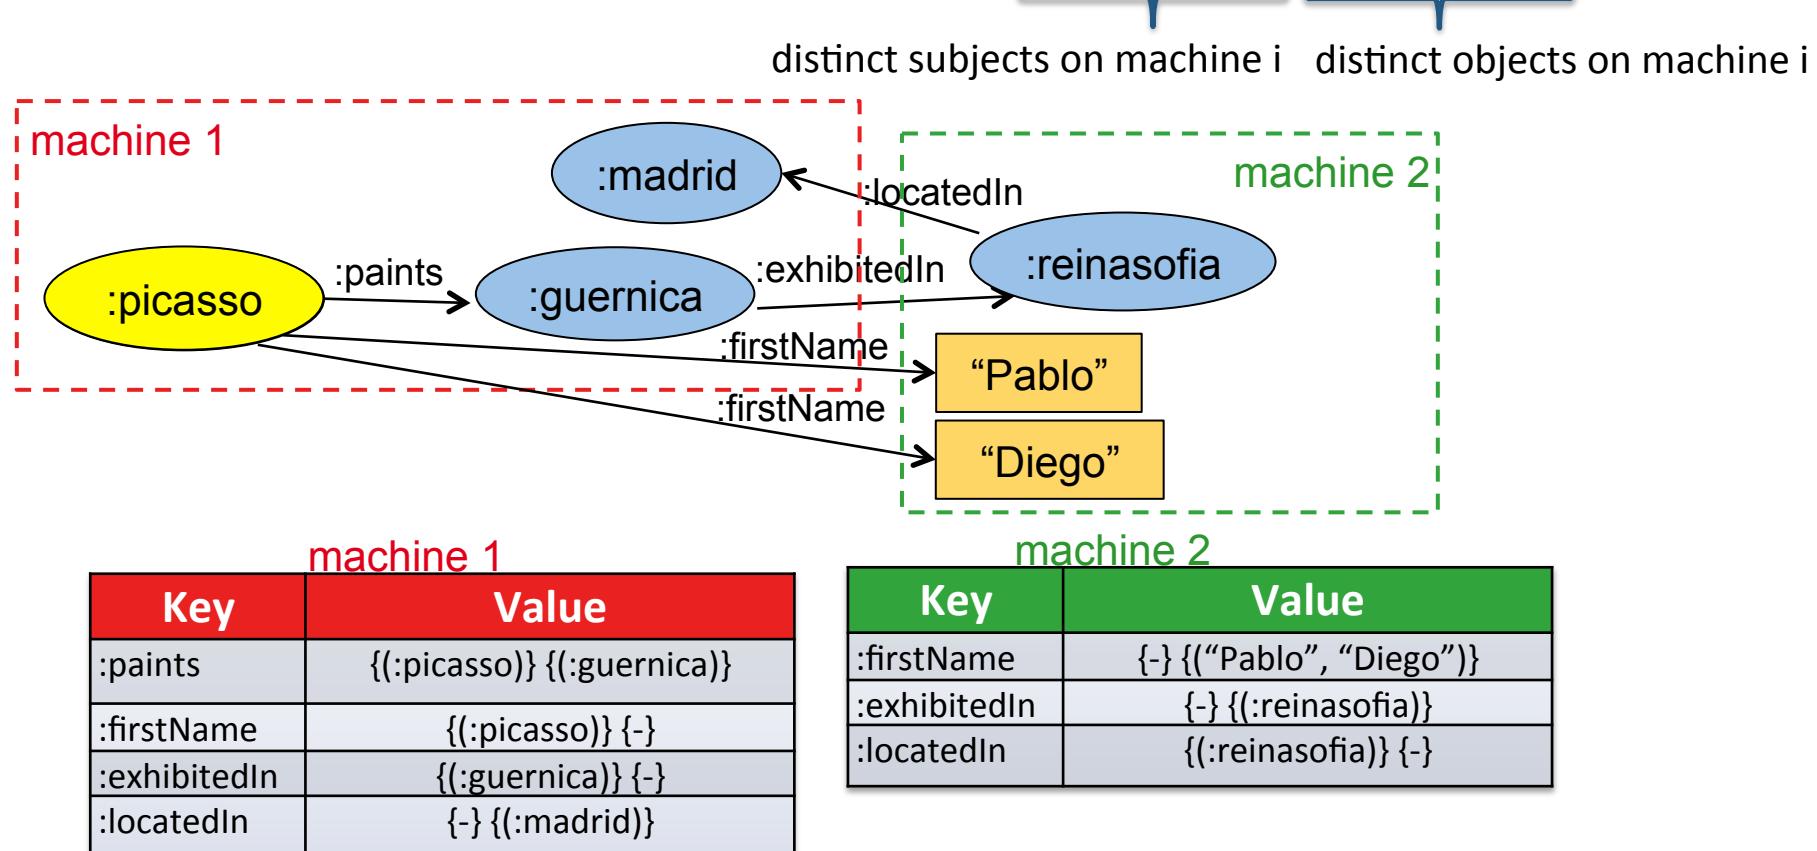
only 1 message needs to be sent

# Trinity.RDF: choosing between the storage alternatives

- Each node is assigned to a machine through hashing
- Two modeling choices
  - 1. (**key**, **value**) : (**node\_id**, <in-adjacency-list, out-adjacency-list>)
  - 2. (**key**, **value**): (**node\_id**, <in1, in2..., out1, out2...>)
    - (**in<sub>i</sub>**, in-adjacency-list<sub>i</sub>): on the same machine i
    - (**out<sub>i</sub>**, out-adjacency-list<sub>i</sub>): on the same machine i
- Deciding whether to use model (1) or (2):
  - **threshold** on the number of neighbors of a node
  - **Local index** on the predicate (in each machine)
  - Aggregate index recording whether a node has a given predicate

# Trinity.RDF: local predicate index

- To support queries of the form ( $?x, p, ?y$ ):
- Global predicate indexing (predicate,  $\langle \text{subject-list}_i, \text{object-list}_i \rangle$ )

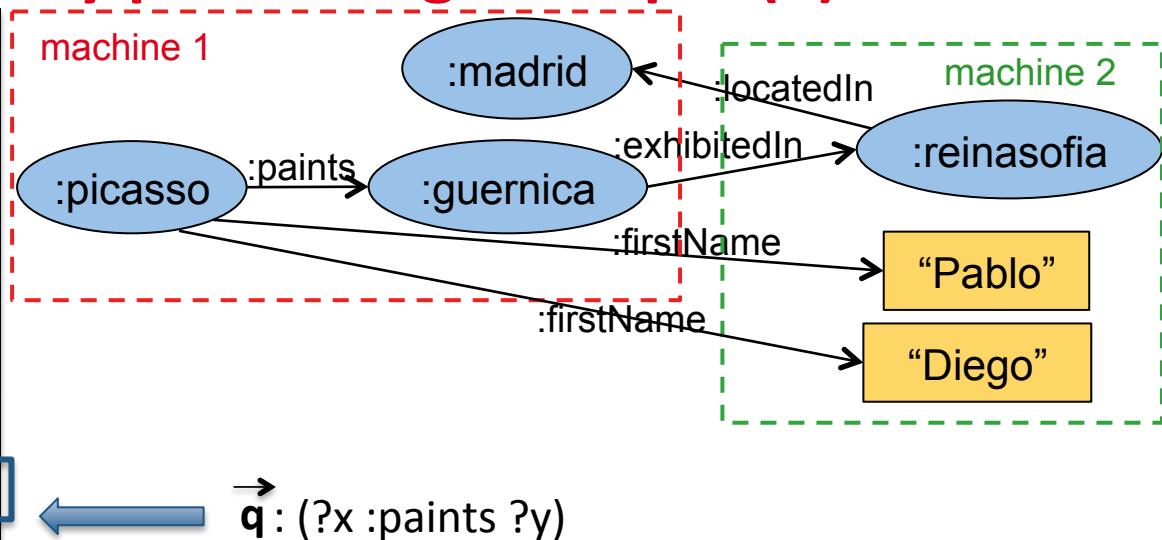


# Trinity.RDF: triple pattern matching

- Single triple pattern ( $s, p, o$ )
  - $s$  or  $o$  is constant → start exploring the graph from this node
  - $s$  and  $o$  are variables →
    - Use POS to find the nodes that are connected by the edge  $p$
    - For each node → graph exploration and filtering out edges that aren't  $p$

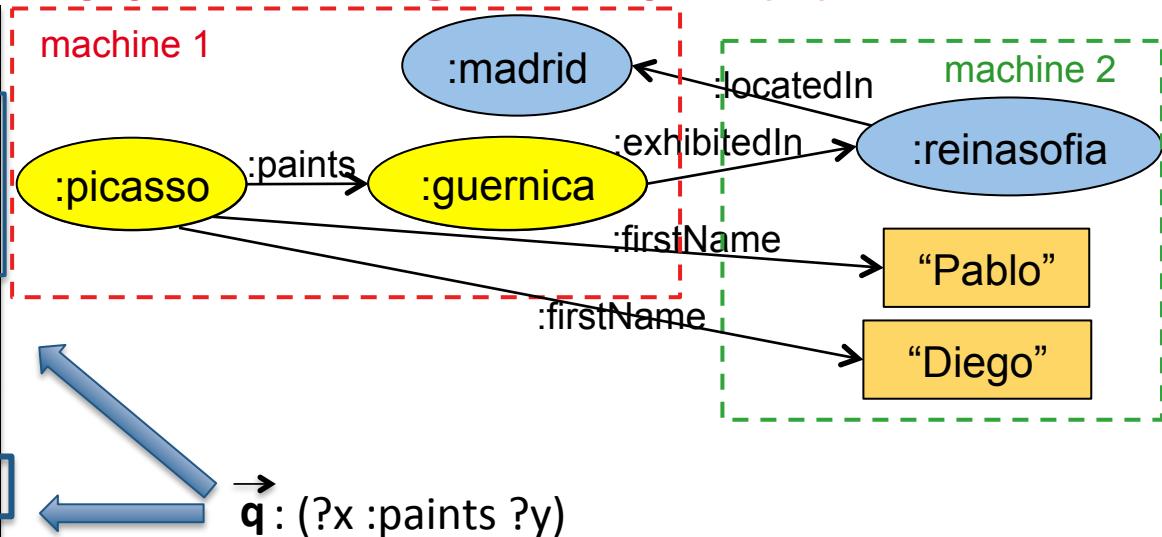
# Trinity.RDF: query processing example (1)

Key	Value
:picasso	In: - Out: (:paints, :guernica) Out: (:firstName, "Pablo") Out: (:firstName, "Diego")
:guernica	In: (:paints, :picasso) Out: (:exhibitedIn, :reinasofia)
:madrid	Out: (:locatedIn, :reinasofia)
:paints	{(:picasso)} {(:guernica)}
:firstName	{(:picasso)} {-}
:exhibitedIn	{(:guernica)} {-}
:locatedIn	{-} {(:madrid)}
Key	Value
:reinasofia	In: (:exhibitedIn, :guernica) Out: (:locatedIn, :madrid)
"Pablo"	In: (:firstName, :picasso)
"Diego"	In: (:firstName, :picasso)
:firstName	{-} {("Pablo", "Diego")}
:exhibitedIn	{-} {(:reinasofia)}
:locatedIn	{(:reinasofia)} {-}



# Trinity.RDF query processing example (1)

Key	Value
:picasso	In: - Out: (:paints, :guernica) Out: (:firstName, "Pablo") Out: (:firstName, "Diego")
:guernica	In: (:paints, :picasso) Out: (:exhibitedIn, :reinasofia)
:madrid	Out: (:locatedIn, :reinasofia)
:paints	{(:picasso)} {(:guernica)}
:firstName	{(:picasso)} {-}
:exhibitedIn	{(:guernica)} {-}
:locatedIn	{-} {(:madrid)}
Key	Value
:reinasofia	In: (:exhibitedIn, :guernica) Out: (:locatedIn, :madrid)
"Pablo"	In: (:firstName, :picasso)
"Diego"	In: (:firstName, :picasso)
:firstName	{-} {("Pablo", "Diego")}
:exhibitedIn	{-} {(:reinasofia)}
:locatedIn	{(:reinasofia)} {-}



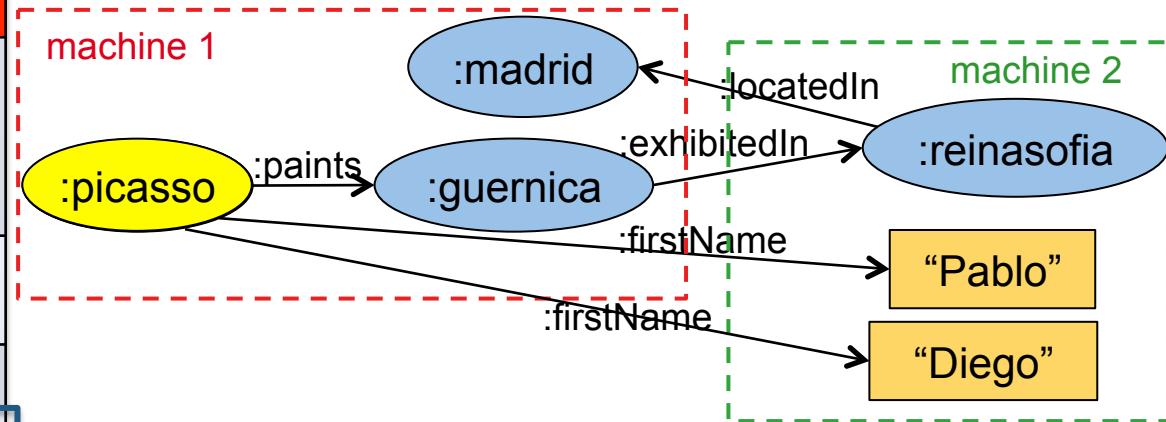
# Trinity.RDF query processing

- Single triple pattern ( $s, p, o$ )
  - $s$  or  $o$  is constant → start exploring the graph from this node
  - $s$  and  $o$  is variable →
    - Use POS to find the nodes that are connected by the edge  $p$
    - For each node → graph exploration and filtering out edges that aren't  $p$
- Multiple triple patterns:
  - process sequentially by graph exploration
  - prune intermediate results

## Trinity.RDF query processing example (2)

Key	Value
:picasso	In: - Out: (:paints, :guernica) Out: (:firstName, "Pablo") Out: (:firstName, "Diego")
:guernica	In: (:paints, :picasso) Out: (:exhibitedIn, :reinasofia)
:madrid	Out: (:locatedIn, :reinasofia)
:paints	{(:picasso)} {(:guernica)}
:firstName	{(:picasso)} {-}
:exhibitedIn	{(:guernica)} {-}
:locatedIn	{-} {(:madrid)}

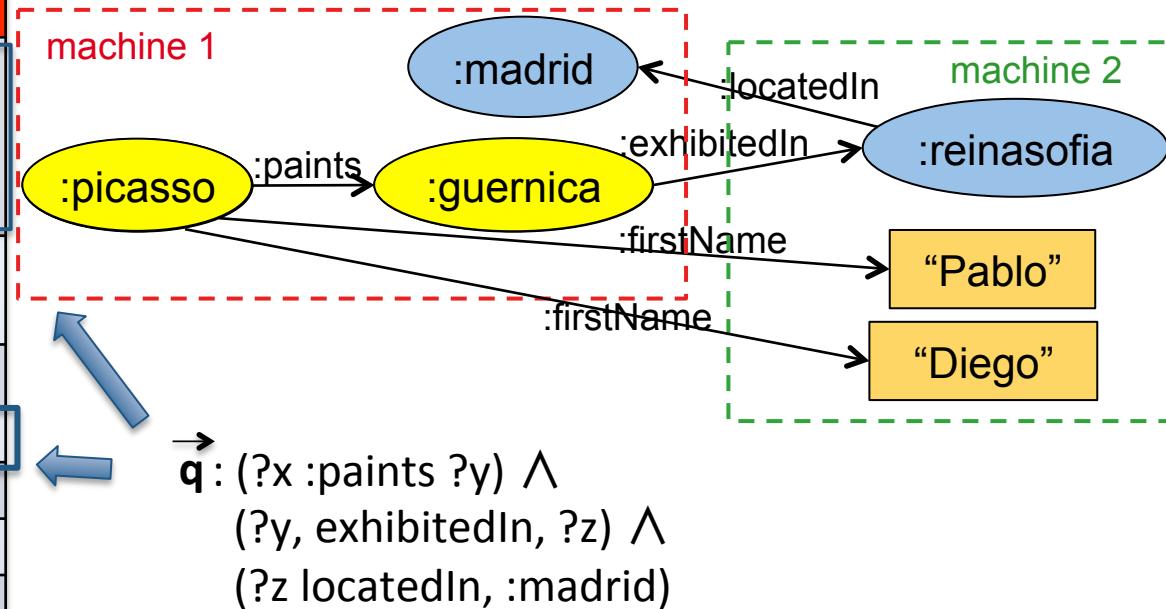
Key	Value
:reinasofia	In: (:exhibitedIn, :guernica) Out: (:locatedIn, :madrid)
"Pablo"	In: (:firstName, :picasso)
"Diego"	In: (:firstName, :picasso)
:firstName	{-} {("Pablo", "Diego")}
:exhibitedIn	{-} {(:reinasofia)}
:locatedIn	{(:reinasofia)} {-}



→  $q: (?x :paints ?y) \wedge$   
 $(?y, \text{exhibitedIn}, ?z) \wedge$   
 $(?z \text{ locatedIn}, :madrid)$

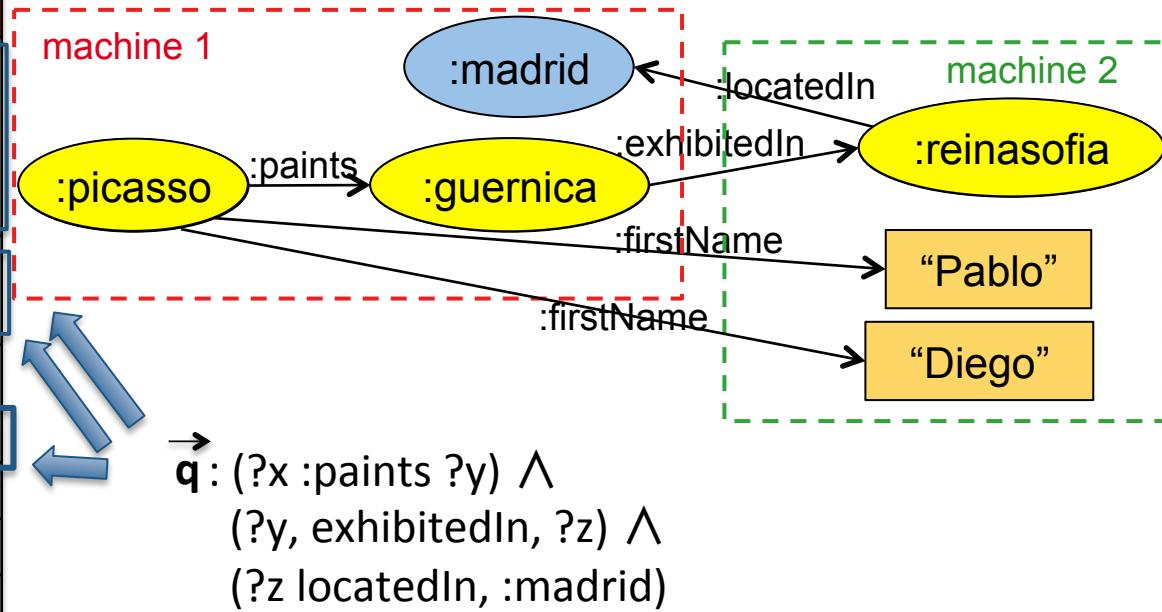
## Trinity.RDF query processing example (2)

Key	Value
:picasso	In: - Out: (:paints, :guernica) Out: (:firstName, "Pablo") Out: (:firstName, "Diego")
:guernica	In: (:paints, :picasso) Out: (:exhibitedIn, :reinasofia)
:madrid	Out: (:locatedIn, :reinasofia)
:paints	{(:picasso)} {(:guernica)}
:firstName	{(:picasso)} {-}
:exhibitedIn	{(:guernica)} {-}
:locatedIn	{-} {(:madrid)}
Key	Value
:reinasofia	In: (:exhibitedIn, :guernica) Out: (:locatedIn, :madrid)
"Pablo"	In: (:firstName, :picasso)
"Diego"	In: (:firstName, :picasso)
:firstName	{-} {("Pablo", "Diego")}
:exhibitedIn	{-} {(:reinasofia)}
:locatedIn	{(:reinasofia)} {-}



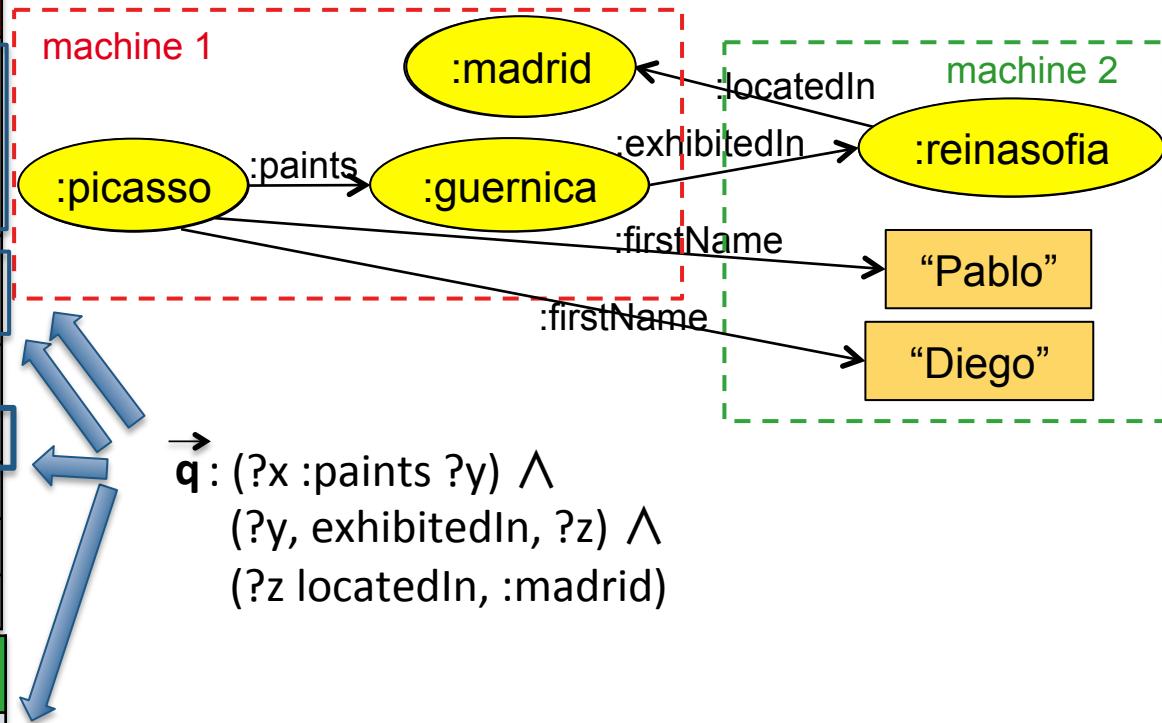
## Trinity.RDF query processing example (2)

Key	Value
:picasso	In: - Out: (:paints, :guernica) Out: (:firstName, "Pablo") Out: (:firstName, "Diego")
:guernica	In: (:paints, :picasso) Out: (:exhibitedIn, :reinasofia)
:madrid	Out: (:locatedIn, :reinasofia)
:paints	{(:picasso)} {(:guernica)}
:firstName	{(:picasso)} {-}
:exhibitedIn	{(:guernica)} {-}
:locatedIn	{-} {(:madrid)}
Key	Value
:reinasofia	In: (:exhibitedIn, :guernica) Out: (:locatedIn, :madrid)
"Pablo"	In: (:firstName, :picasso)
"Diego"	In: (:firstName, :picasso)
:firstName	{-} {("Pablo", "Diego")}
:exhibitedIn	{-} {(:reinasofia)}
:locatedIn	{(:reinasofia)} {-}



## Trinity.RDF query processing example (2)

Key	Value
:picasso	In: - Out: (:paints, :guernica) Out: (:firstName, "Pablo") Out: (:firstName, "Diego")
:guernica	In: (:paints, :picasso) Out: (:exhibitedIn, :reinasofia)
:madrid	Out: (:locatedIn, :reinasofia)
:paints	{(:picasso)} {(:guernica)}
:firstName	{(:picasso)} {-}
:exhibitedIn	{(:guernica)} {-}
:locatedIn	{-} {(:madrid)}
Key	Value
:reinasofia	In: (:exhibitedIn, :guernica) Out: (:locatedIn, :madrid)
"Pablo"	In: (:firstName, :picasso)
"Diego"	In: (:firstName, :picasso)
:firstName	{-} {("Pablo", "Diego")}
:exhibitedIn	{-} {(:reinasofia)}
:locatedIn	{(:reinasofia)} {-}



# Trinity.RDF query processing

- Single triple pattern ( $s, p, o$ )
  - $s$  or  $o$  is constant → start exploring the graph from this node
  - $s$  and  $o$  is variable →
    - Use POS to find the nodes that are connected by the edge  $p$
    - For each node → graph exploration and filtering out edges that aren't  $p$
- Multiple triple patterns:
  - process sequentially by graph exploration
  - prunes intermediate results
- Final join in a proxy for removing invalid results (small join)
- Ordering matters!
- In-memory system; hardware strongly impacts performance

## Summary: Approaches based on graph stores

Pros	Cons
Replace joins by navigation	Replace joins by navigation Hard to enforce data locality

# 3.4

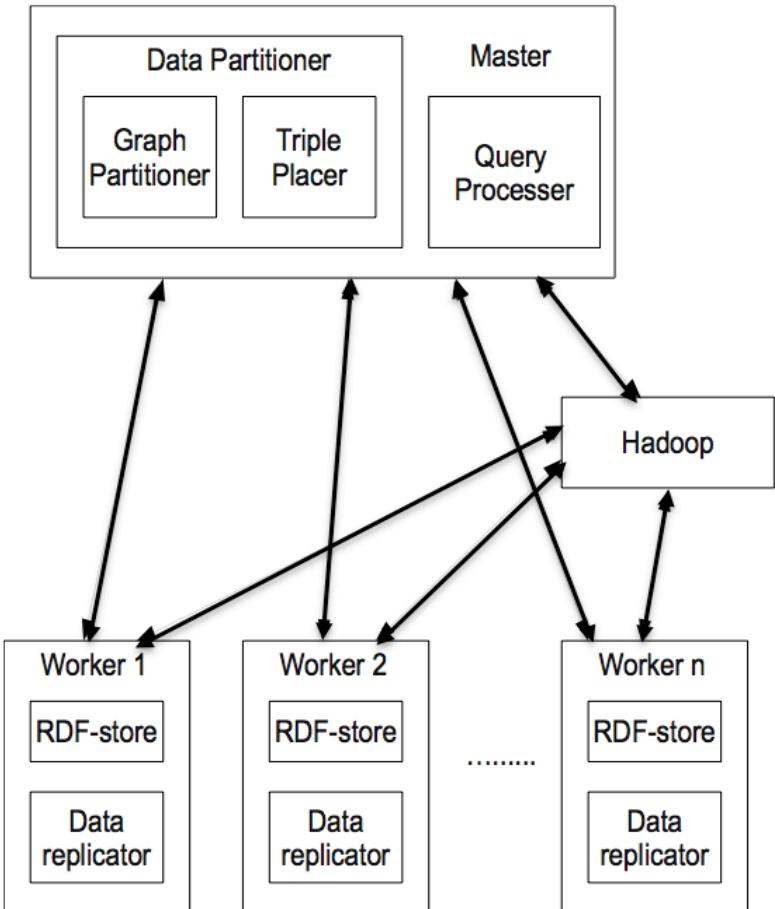
Multiple centralized (RDF) stores

## IV. Multiple centralized (RDF) stores

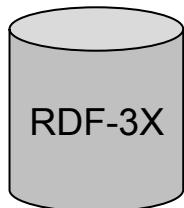
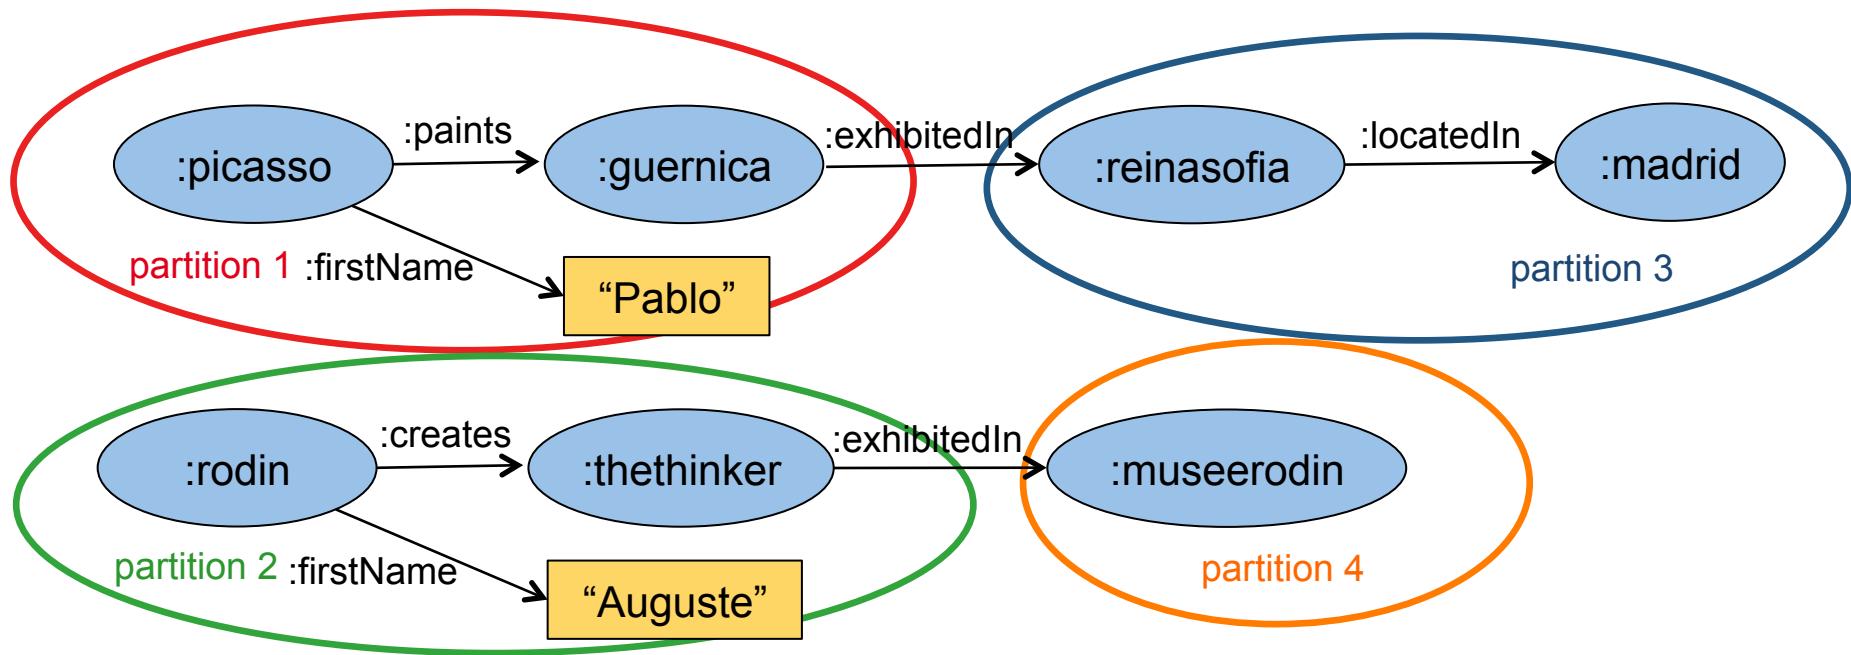
- Data partitioning using known **graph** tools
  - [Huang11] uses **METIS**: a graph partitioning tool
  - [Hose13] extends [Huang11] for known query workloads
- Data partitioning using **hashing**
  - SHAPE [Lee13] uses a semantic hashing method
- **Horizontal** fragmentation
  - Partout [Galarraga12]

# Graph partitioning [Huang11]

- Data **partitioning**
  - METIS graph partitioner [METIS]
  - **vertex** partitioning
  - $\# \text{partitions} = \# \text{machines}$
  - **minimum edge cut**
  - **type** triples are excluded from the partitioning process
- Data **placement**
  - partition triples by the subject
  - **replication** across the partitions
  - directed/undirected n-hop guarantee
- Data **storage**
  - each machine stores triples into **RDF-3X**



# Graph partitioning [Huang11] - storage



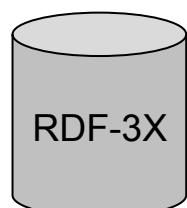
node 1



node 2

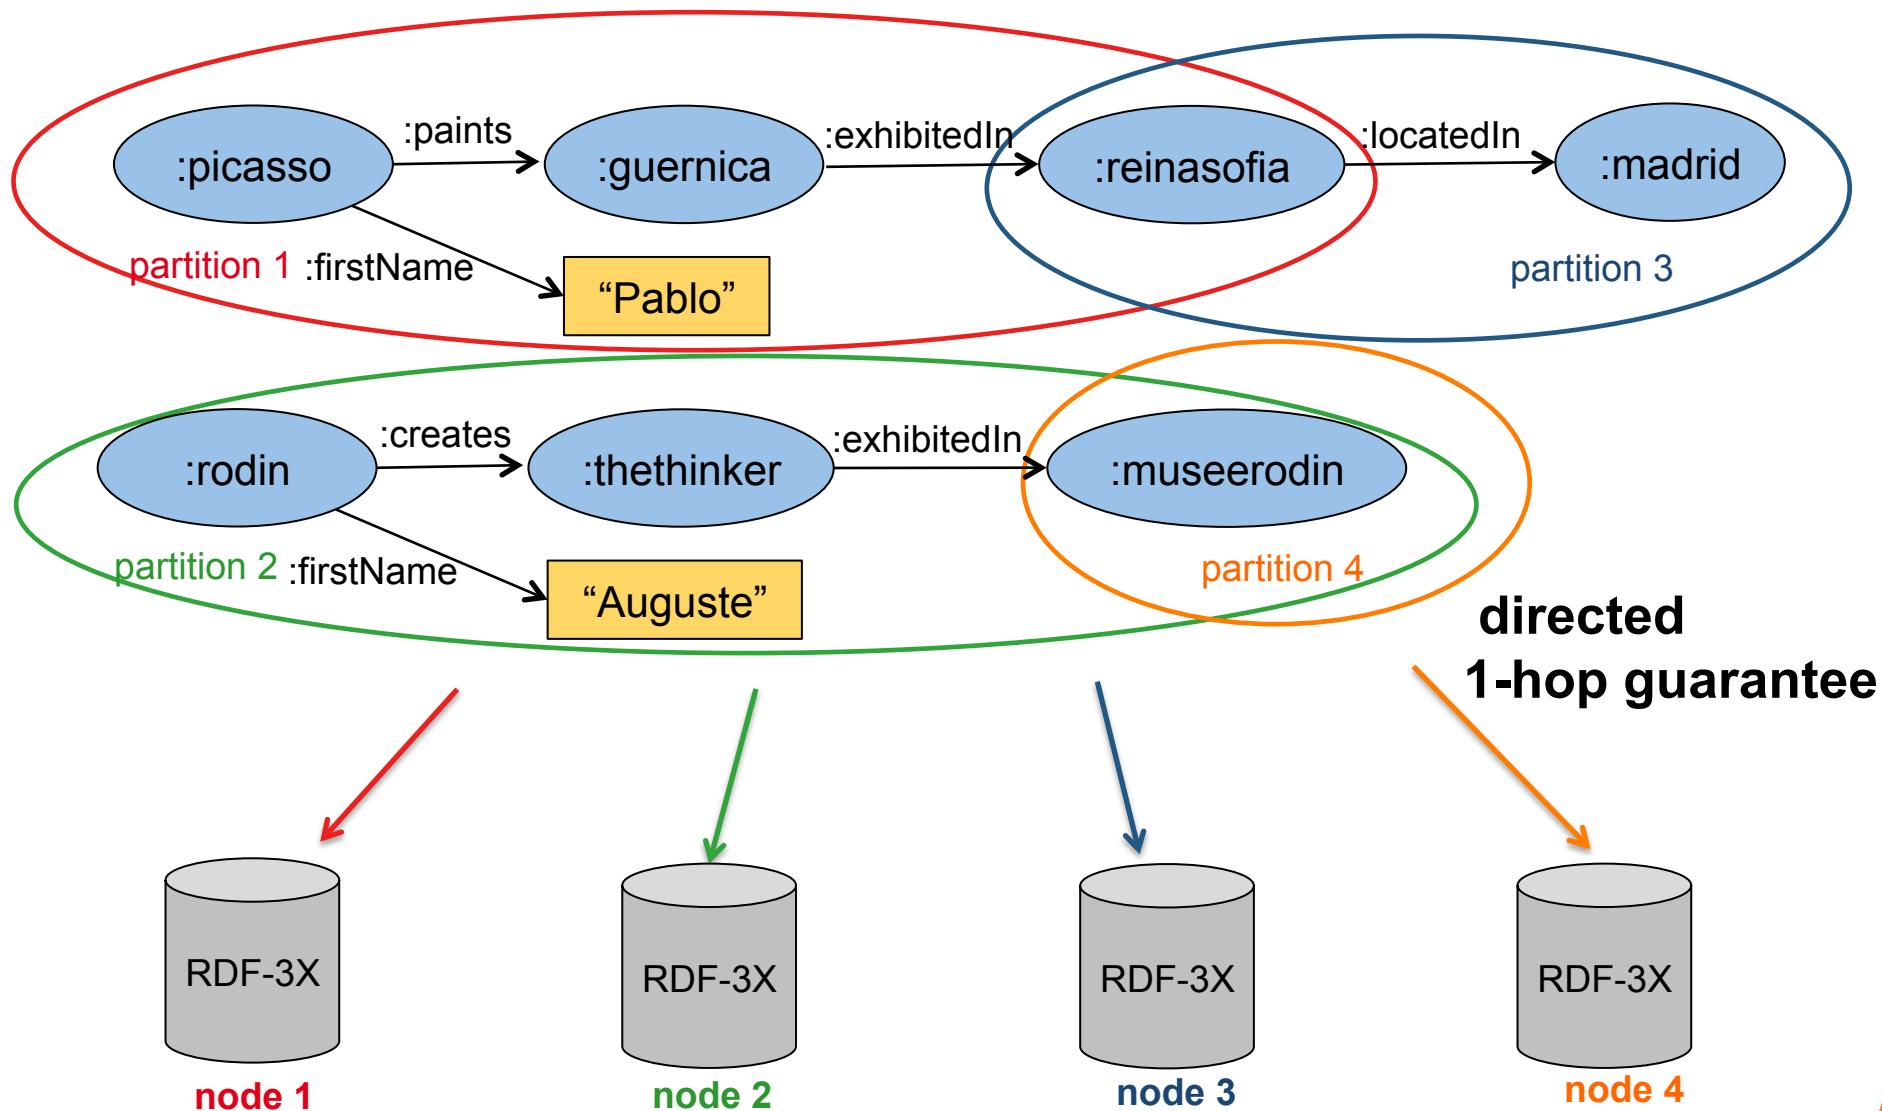


node 3

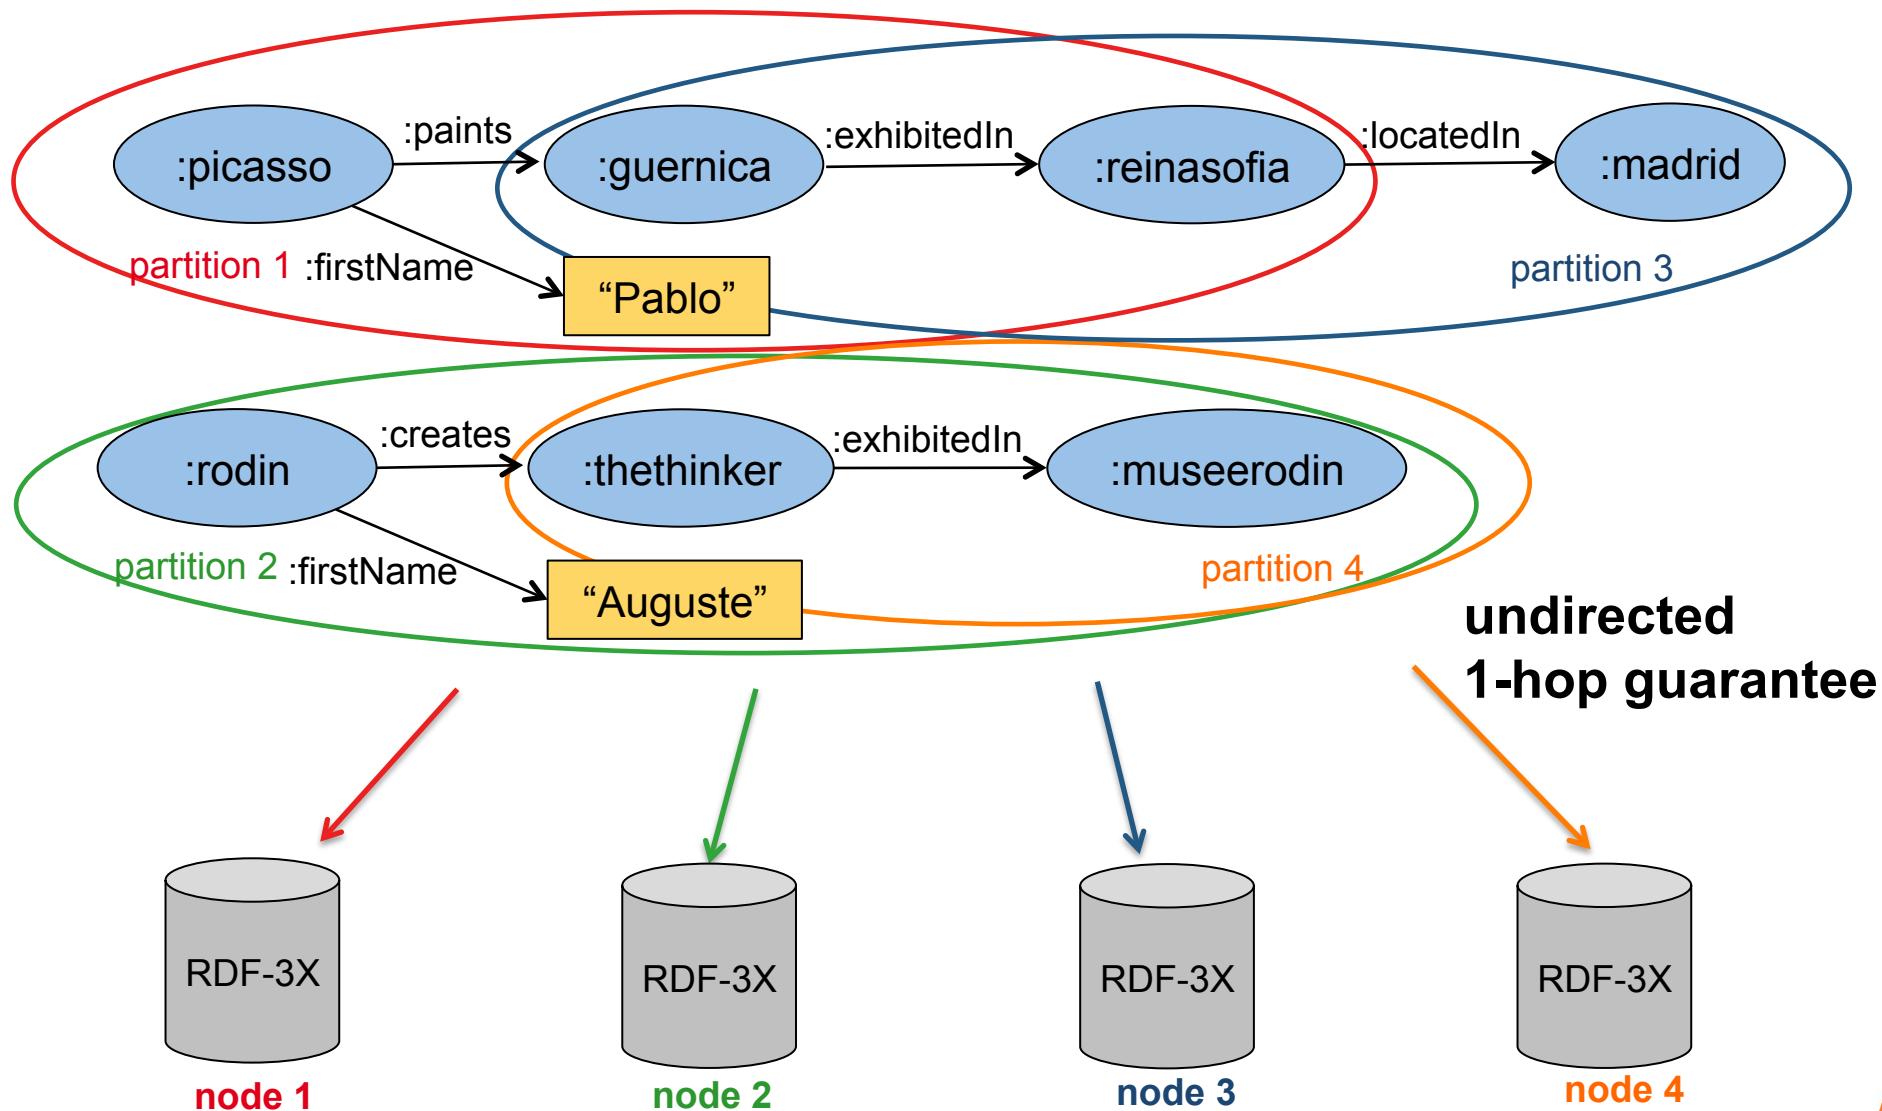


node 4

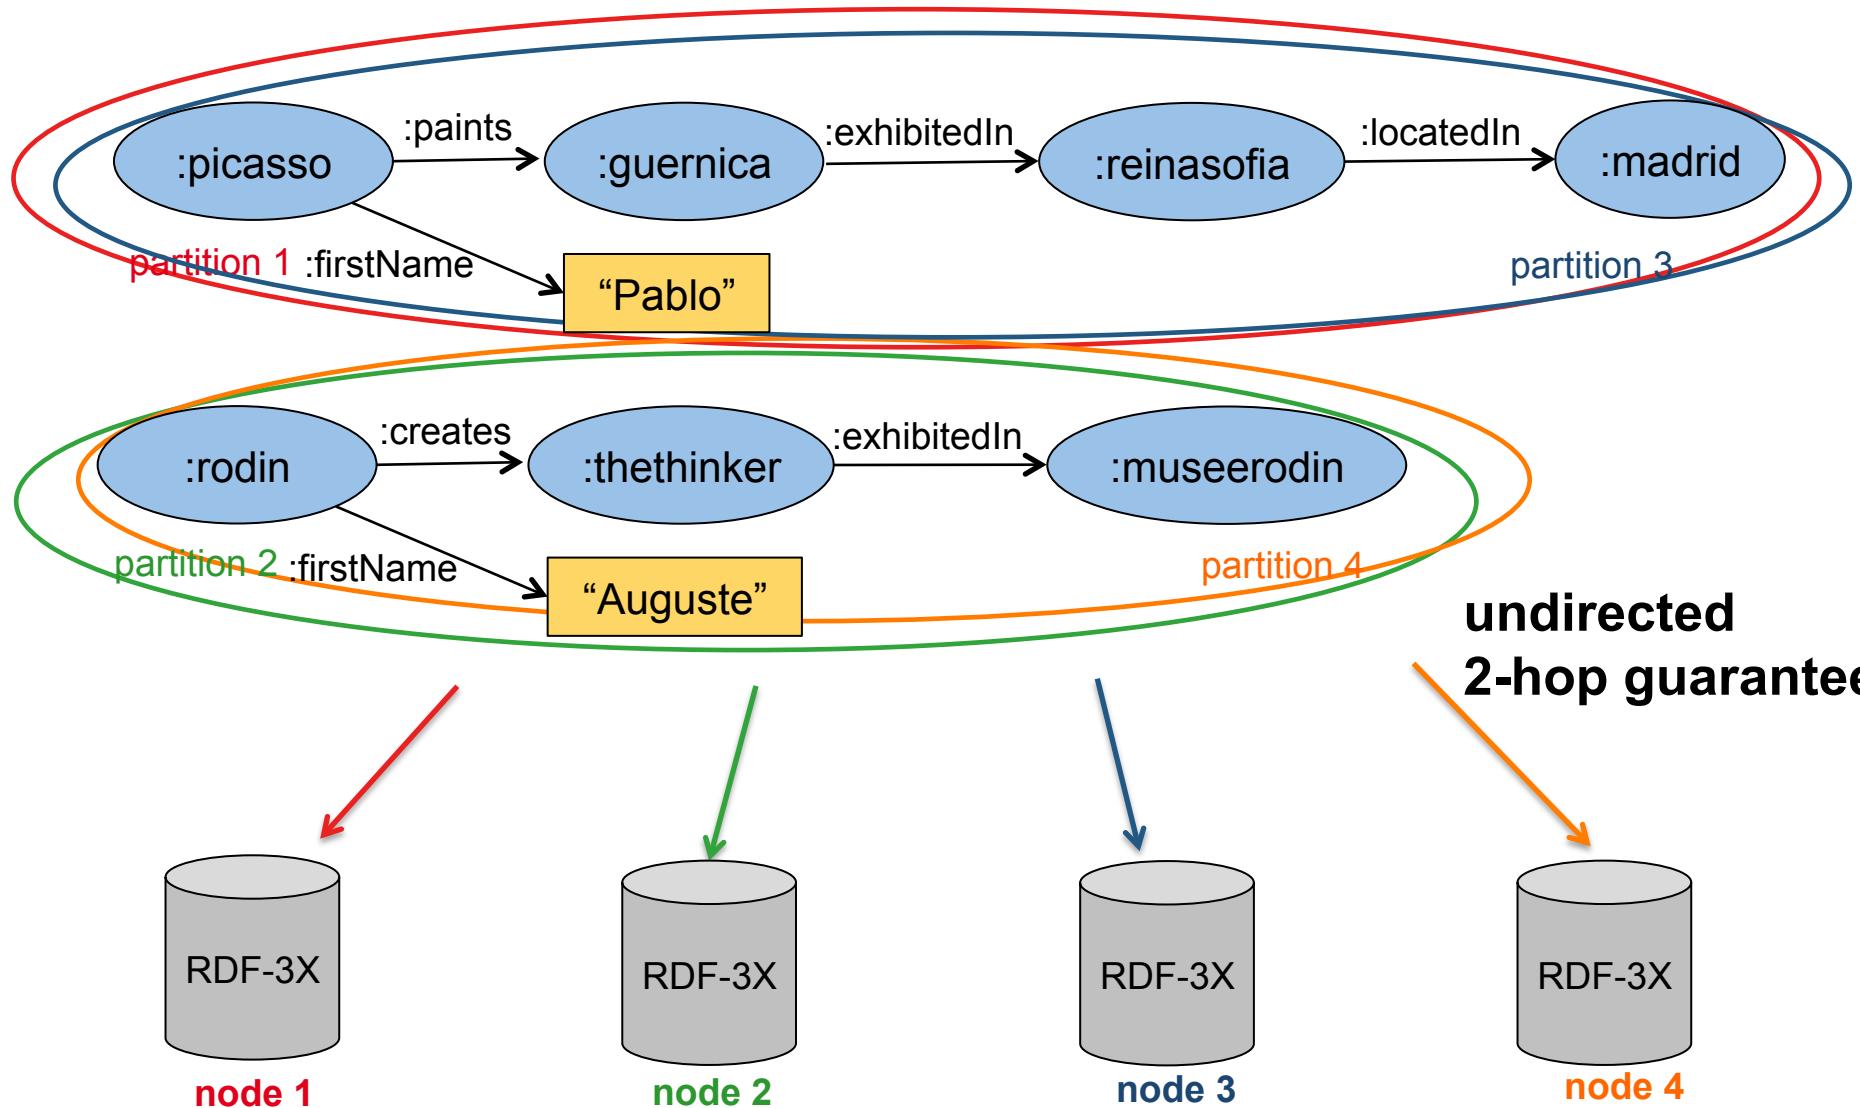
# Graph partitioning [Huang11] - storage



# Graph partitioning [Huang11] - storage

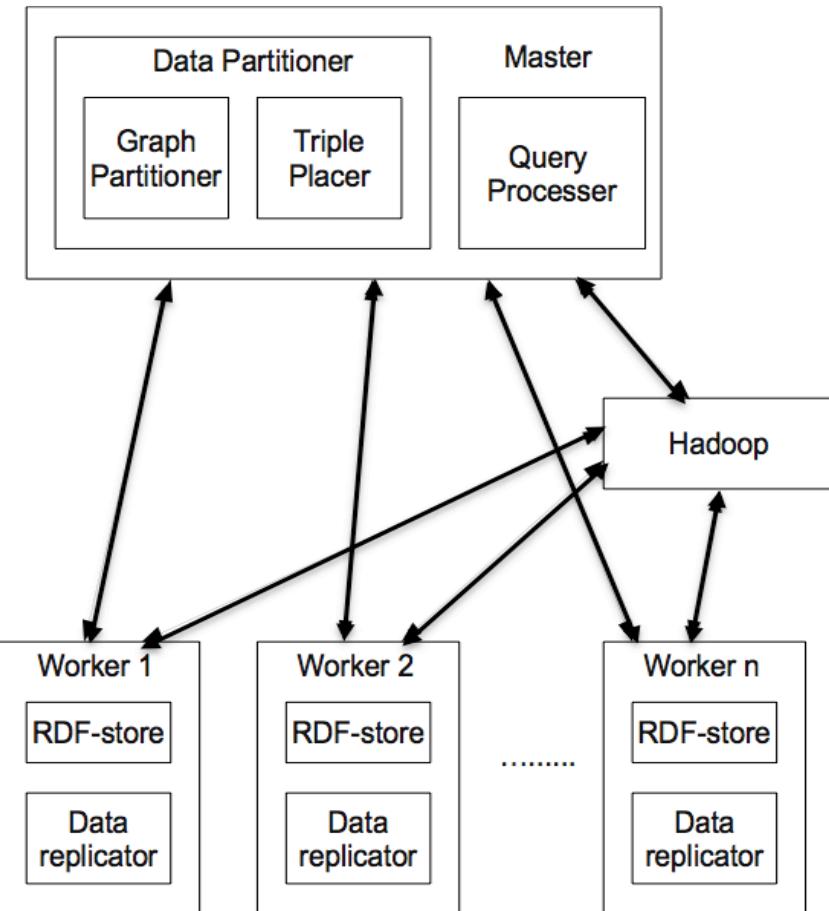


# Graph partitioning [Huang11] - storage



# Query processing [Huang11]

- Query decomposition for parallelization
- Check if q is PWOC (parallelizable without communication)
  - Yes: answer only from RDF-3X and union results
  - No: answer subqueries from RDF-3X and join them in Hadoop
- # jobs grows with # subqueries
- Heuristic optimization
  - decomposition so as to have minimum number of subqueries → minimal edge partitioning of the query graph into sub-graphs of bounded diameter

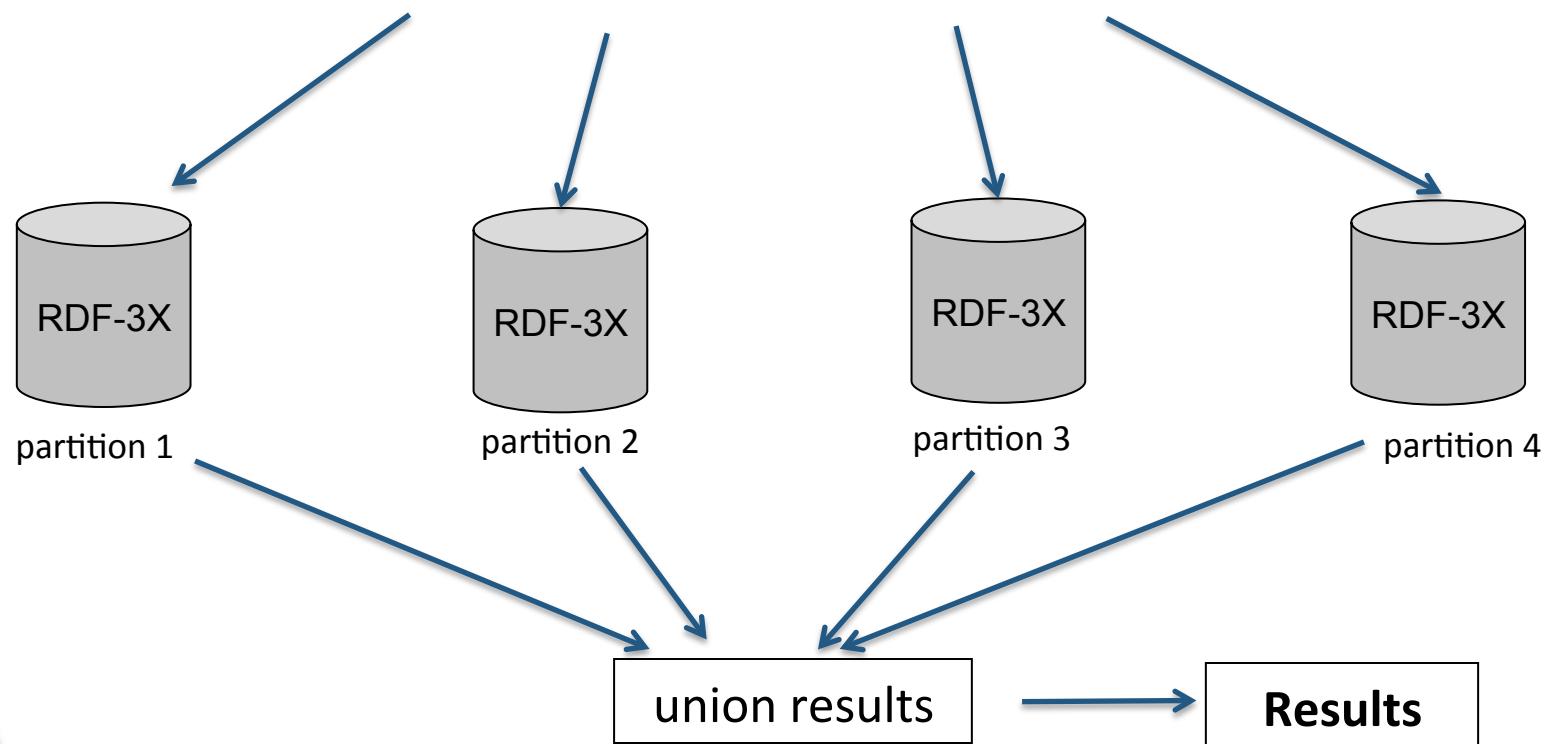


# Query processing [Huang11]

Replication with  
1-hop guarantee

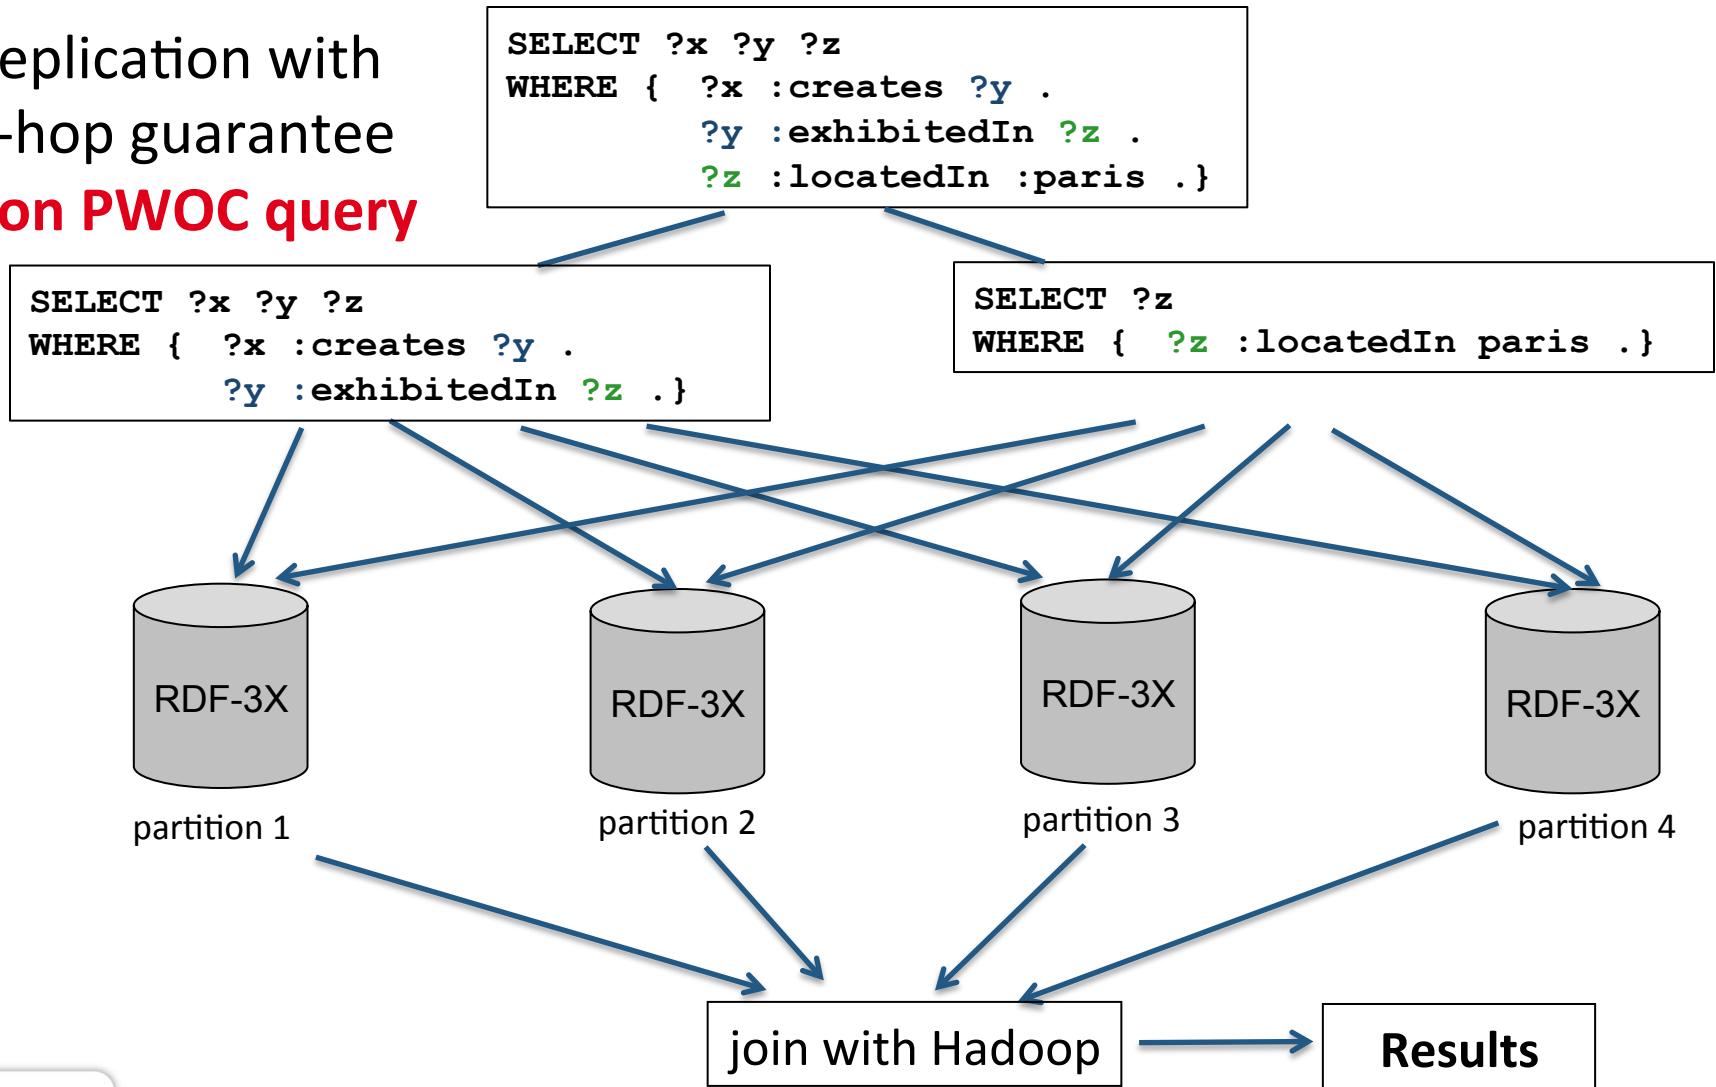
PWOC query

```
SELECT ?x ?y ?z
WHERE {
    ?x type :artist .
    ?x :firstName ?y .
    ?x :creates ?z . }
```



# Query processing [Huang11]

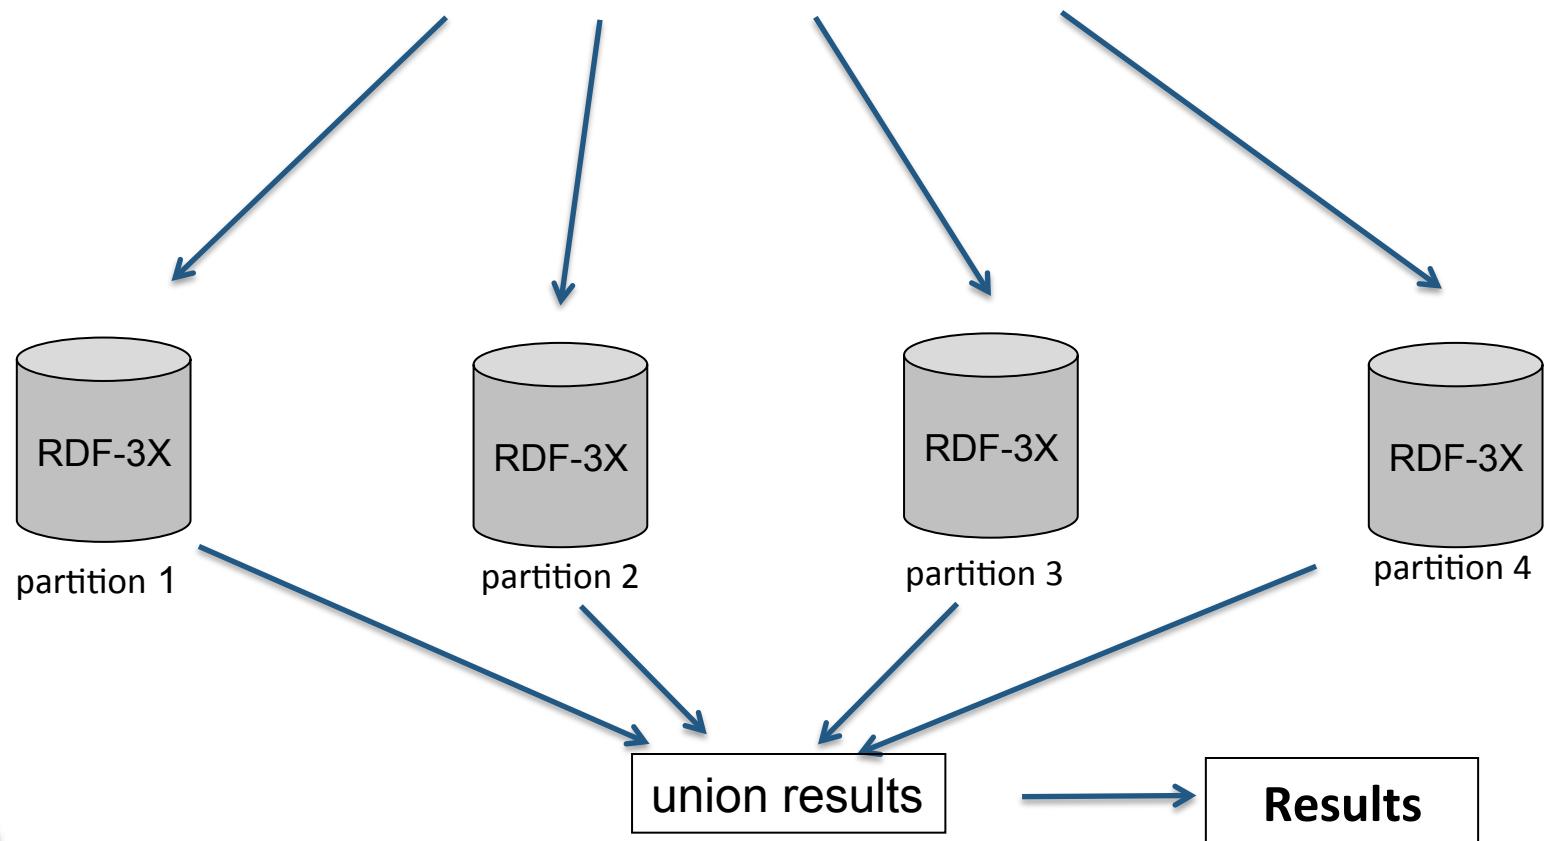
replication with  
1-hop guarantee  
**non PWOC query**



# Query processing [Huang11]

replication with  
2-hop guarantee  
**PWOC query**

```
SELECT ?x ?y ?z
WHERE { ?x :creates ?y .
        ?y :exhibitedIn ?z .
        ?z :locatedIn :paris . }
```



## Concluding remarks [Huang11]

- + Data fragmentation inspired by RDF structure
- + Reduces shuffling
- Not scalable for big datasets (Metis is centralized)
- Sub-queries broadcast to all machines → significant total effort

# SHAPE (Semantic HAsh Partitioning) [Lee13]

Introduces **semantic hashing** partitioning technique

## 1. Identify **triple groups**

- Subject triple group (pairwise disjoint)
  - All triples have the same subject (group anchor)
  - Favors **subject-star** queries
- Object triple groups (pairwise disjoint)
  - All triples have the same object (group anchor)
  - Favors **object-star** queries
- Subject-object triple groups (each triple is in 2 groups)
  - All triples have the same subject and the same object (group anchor)
  - Favors **subject- and object-star** queries

# SHAPE (Semantic HAsh Partitioning)

Introduces semantic hashing partitioning technique

## 2. Baseline hash partition

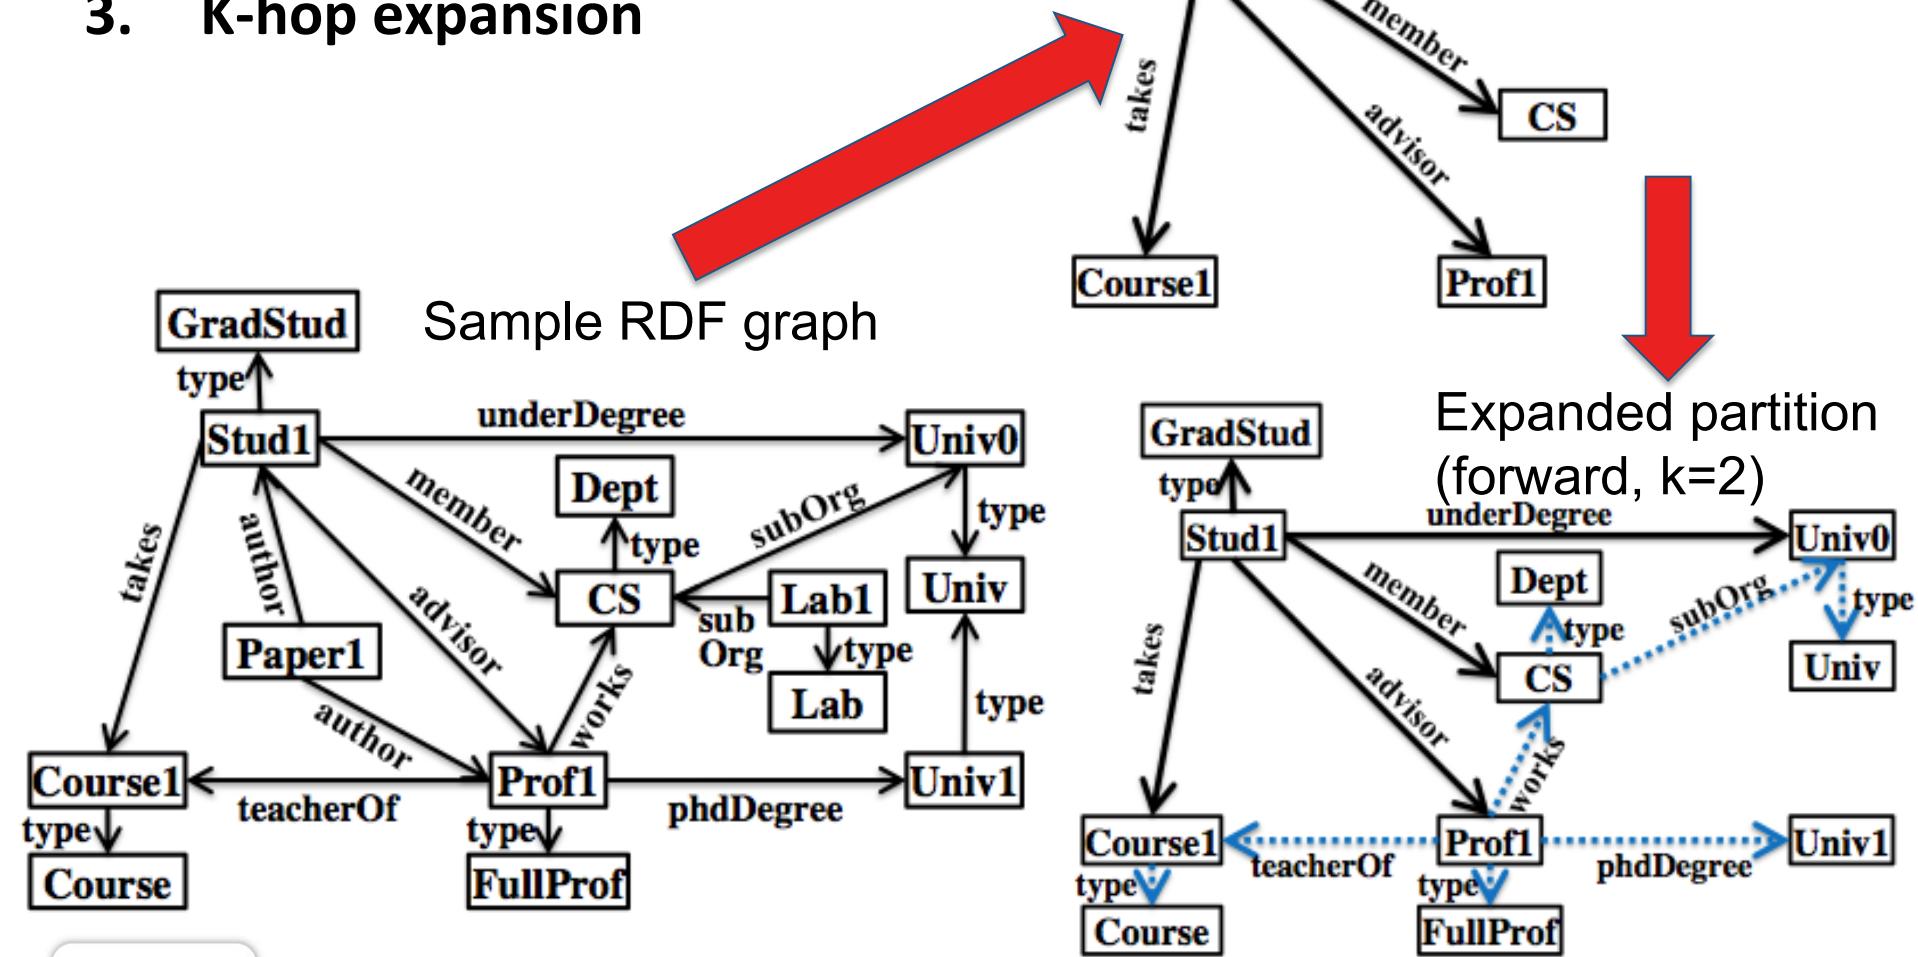
- Apply a hash function on each triple group anchor → baseline hash partitions for s (or o, or (s,o))

## 3. K-hop expansion

- Add to each partition all the triples whose *shortest distance* to *any anchor* of the partition is *at most k* (e.g., k=2)
  - Distance can be measured forward ( $s \rightarrow o$ ), backward ( $o \rightarrow s$ ) or in any direction (minimum length between  $s \rightarrow o$  and  $o \rightarrow s$ )
  - The goal is to increase the chances of PWOC query evaluation

# SHAPE (Semantic HAsh Partitioning)

2. Baseline hash partition
3. K-hop expansion



# SHAPE (Semantic HAsh Partitioning)

1. Identify **triple groups** (based on s; o; or (s, o))
2. **Baseline** hash partition
  - Hash triple groups on s (*or o, or (s,o)*)
3. **K-hop expansion**
  - Add to each partition all the triples whose *shortest distance to any anchor* of the partition is *at most k*
3. **Selective k-hop expansion**
  - Avoid some expansions which are « unlikely to help »
5. **URI-based optimization**
  - Idea: common URI prefix predicts correlated access
  - Hash on URI prefix of a group anchor, not on full URI

Each hash partition group = 1 RDF-3X server; joins on Hadoop

## Summary: Multiple centralized (RDF) stores

Pros	Cons
Reduced shuffle cost RDF-specific local operations	Fault tolerance must be implemented on top of the centralized servers  Scaling up the initial graph partitioning is non-trivial

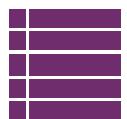
# 3.5

## Hybrid systems

## V. Hybrid RDF stores

- AMADA [Bugiotti12, Aranda12]
  - Generic cloud-based architecture
  - Built on top of **Amazon** Web Services as a **SaaS**
  - **Route the queries** to a small **subset** of the data that are likely to have matches
  - Trade-offs between **efficiency** and **monetary costs**

# AMADA architecture



## Key-value store

Ideal for indexing and querying **small, structured data**



## Storage for raw data

Ideal for large files



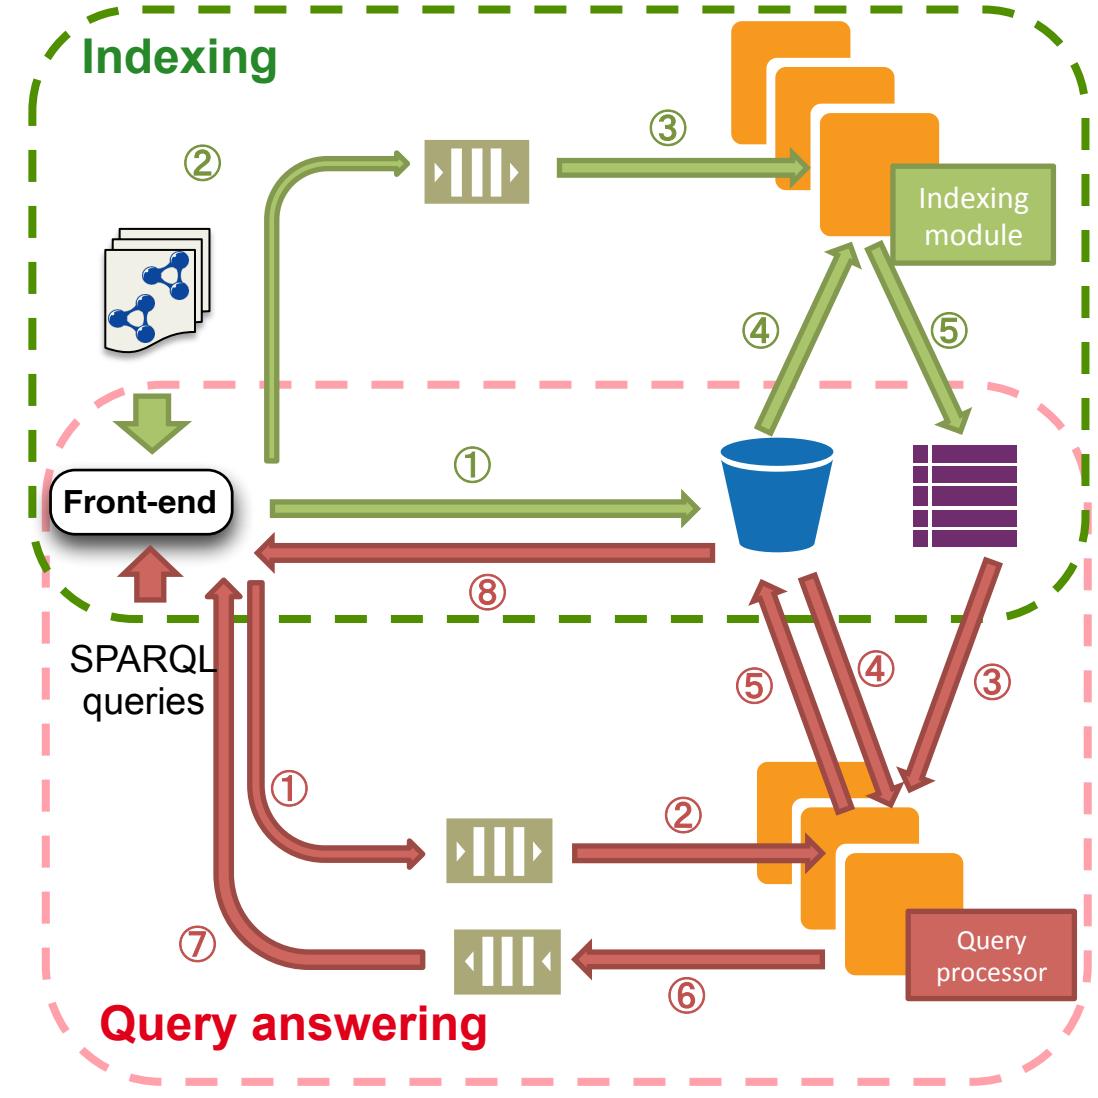
## Virtual machines

**Resizable computing capacity** in the cloud



## Messaging queues

**Queues** for communication between distributed components



# AMADA indexing and storage

**S table**

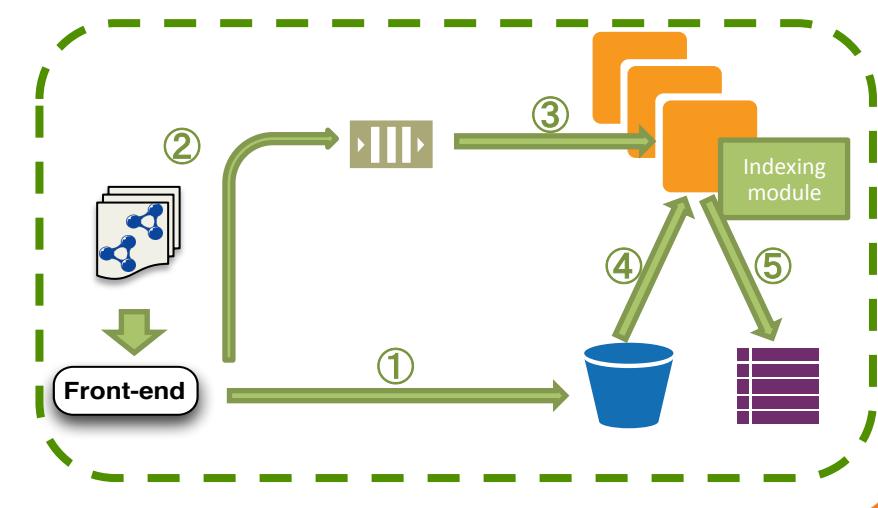
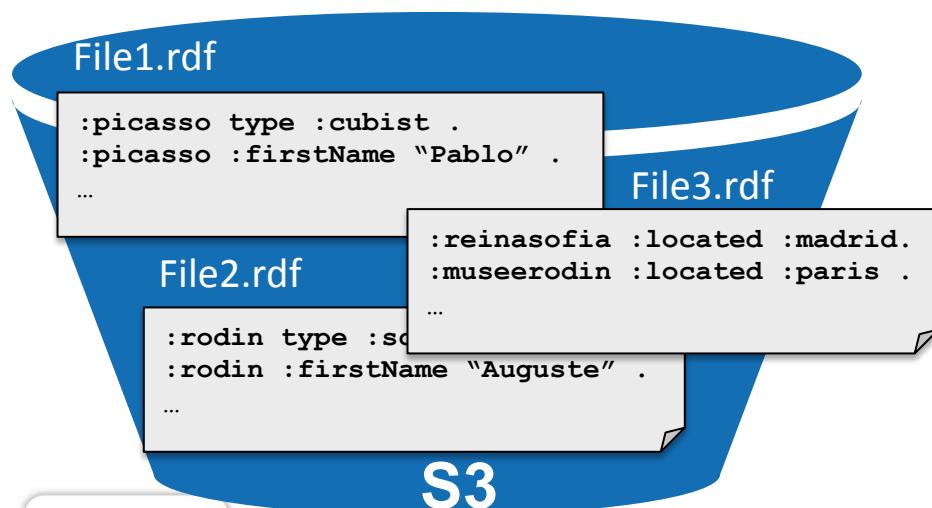
key	(attribute, value)
:picasso	(file1, -)
:guernica	(file1, -)
:reinasofia	(file1, -)
:rodin	(file2, -)
:thethinker	(file2, -)
:museerodin	(file2, -)
...	...

**P table**

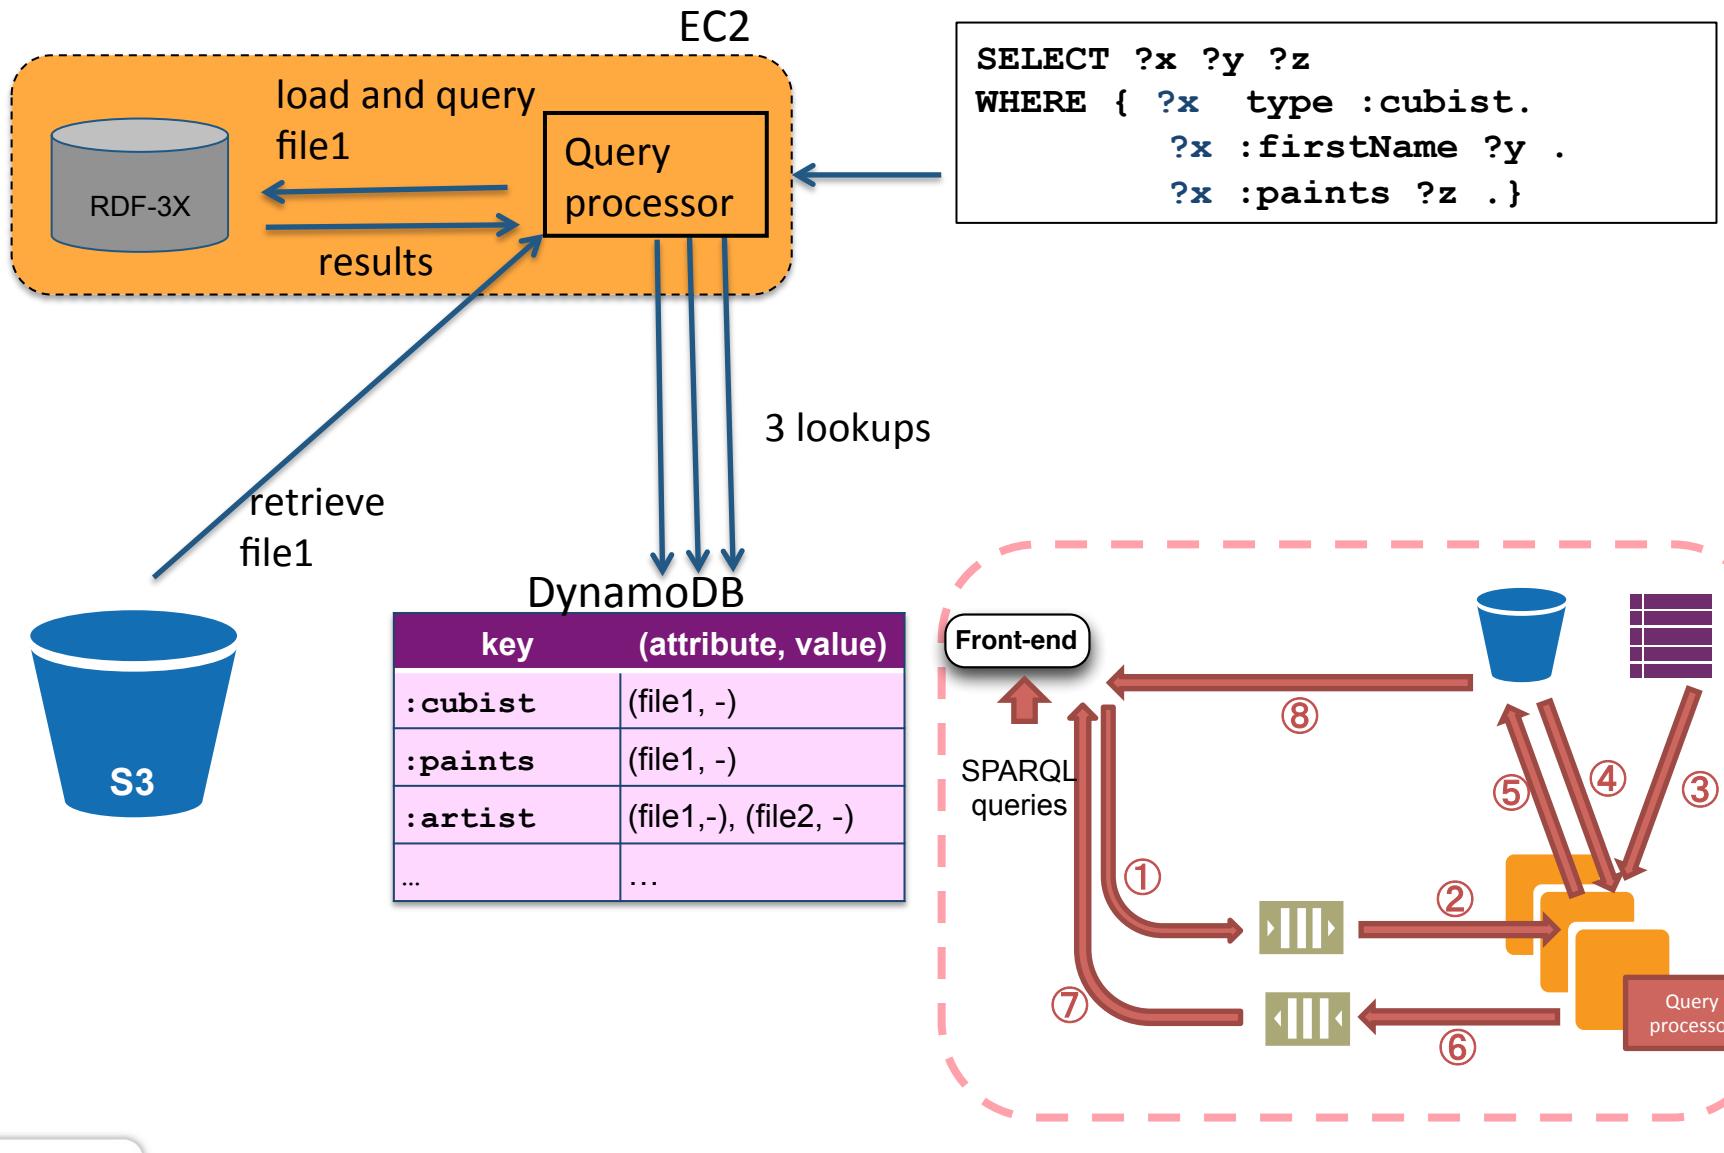
key	(attribute, value)
type	(file1,-), (file2, -)
:firstName	(file1,-), (file2, -)
:paints	(file1,-)
:exhibitedIn	(file1,-), (file2, -)
:locatedIn	(file1,-)
:creates	(file2,-)
...	...

**O table**

key	(attribute, value)
:cubist	(file1,-)
:guernica	(file1,-)
"Pablo"	(file1,-)
:sculptor	(file2,-)
:thethinker	(file2,-)
:museerodin	(file2,-)
...	...



# AMADA querying



# AMADA indexing strategies

- Source indexing strategies
  - Identify files (F) that contribute an answer to a query
  - Term-based: T|F|- (T: URIs, literals etc. without caring if it's S, P, or O)
  - Attribute-based: S|F|- P|F|- O|F|-
  - Attribute-subset: S|F|- P|F|- O|F|- SP|F|- PS|F|- SO|F|- SPO|F|-
- Distributed data indexing (as in Part II)
  - RDF data itself is indexed in the key-value store
  - S|P|O P|O|S O|S|P
- Different trade-offs

T table

# Term-based example

key	(attribute, value)
:picasso	(file1, -)
:guernica	(file1, -), (file4, -)
:reinasofia	(file1, -)
:rodin	(file2, -)
:museerodin	(file2, -)
:firstName	(file1, -), (file2, -), (file5, -)
:paints	(file1, -)
...	...

```
SELECT ?y  
WHERE {?x :firstName ?y .  
?x :paints :guernica .}
```

Lookups

$$\text{Get}(T, \text{:firstName}) \cup [\text{Get}(T, \text{:paints}) \cap \text{Get}(T, \text{:guernica})]$$

$\{ \text{file1}, \text{file2}, \text{file5} \}$        $\{ \text{file1} \}$        $\{ \text{file1}, \text{file4} \}$

$$\{ \text{file1} \}$$

$\{ \text{file1}, \text{file2}, \text{file5} \}$

{file1, file2, file5}

S3

file1.rdf

file2.rdf

file3.rdf

file4.rdf

file5.rdf

:picasso type :cubist .  
:picasso :firstName "Pablo" .  
...



:rodin type :sculptor .  
:rodin :firstName "Auguste" .  
...



:reinasofia :located :madrid.  
:museerodin :located :paris .  
...

:guernica :exhibited :reinasofia.  
:thethinker :exhibited :museerodin.  
...

:cubist sc :painter.  
:firstName domain :artist .  
...



# Attribute-based example

**S table**

key	(attribute, value)
:picasso	(file1, -)
:guernica	(file4, -)
:reinasofia	(file1, -)
:rodin	(file2, -)
:thethinker	(file2, -)
:firstName	(file4, -)
...	...

```
SELECT ?y
WHERE {?x :firstName ?y .
       ?x :paints :guernica .}
```

file1.rdf

```
:picasso type :cubist .
:picasso :firstName "Pablo" .
...
```

file2.rdf

```
:rodin type :sculptor .
:rodin :firstName "Auguste" .
...
```

file3.rdf

```
:reinasofia :located :madrid.
:museerodin :located :paris .
...
```

file4.rdf

```
:guernica :exhibited :reinasofia.
:thethinker :exhibited :museerodin.
...
```

S3

file5.rdf

```
:cubist sc :painter.
:firstName domain :artist .
...
```

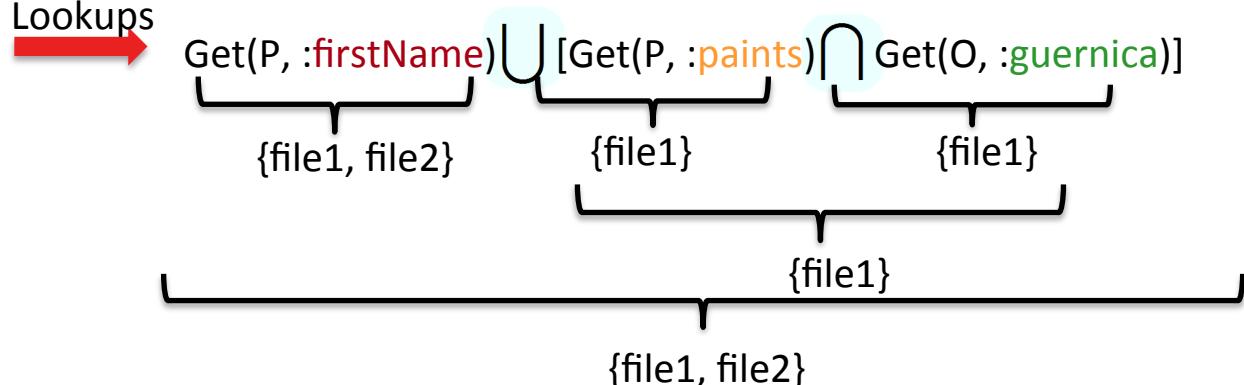
**P table**

key	(attribute, value)
type	(file1,-), (file3, -), (file4, -)
:firstName	(file1,-), (file2, -)
:paints	(file1,-)
:exhibitedIn	(file1,-), (file2, -)
:locatedIn	(file1,-)
...	...

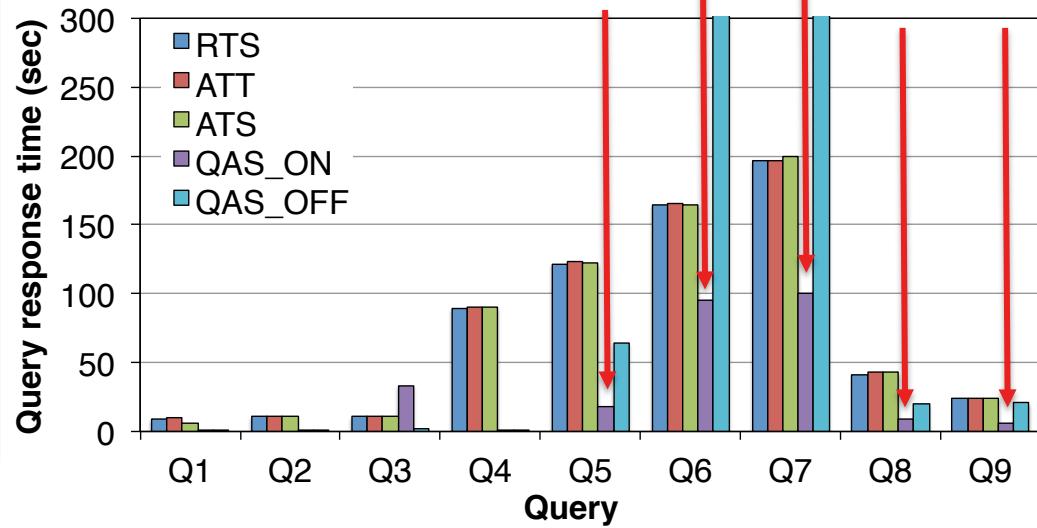
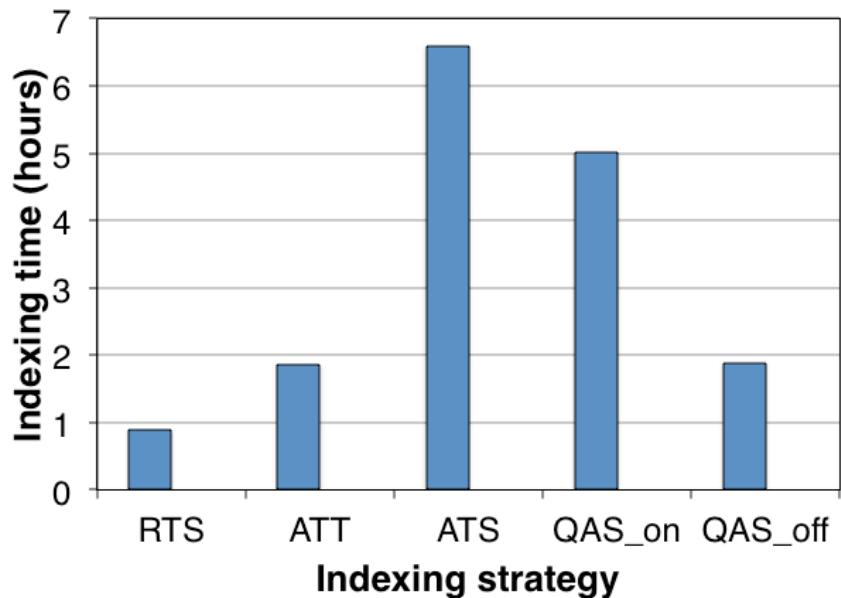
**O table**

key	(attribute, value)
:cubist	(file1,-)
:guernica	(file1,-)
"Pablo"	(file1,-)
:sculptor	(file2,-)
:thethinker	(file2,-)
...	...

Lookups



# AMADA results



**RTS:** term-based

**ATT:** attribute-based

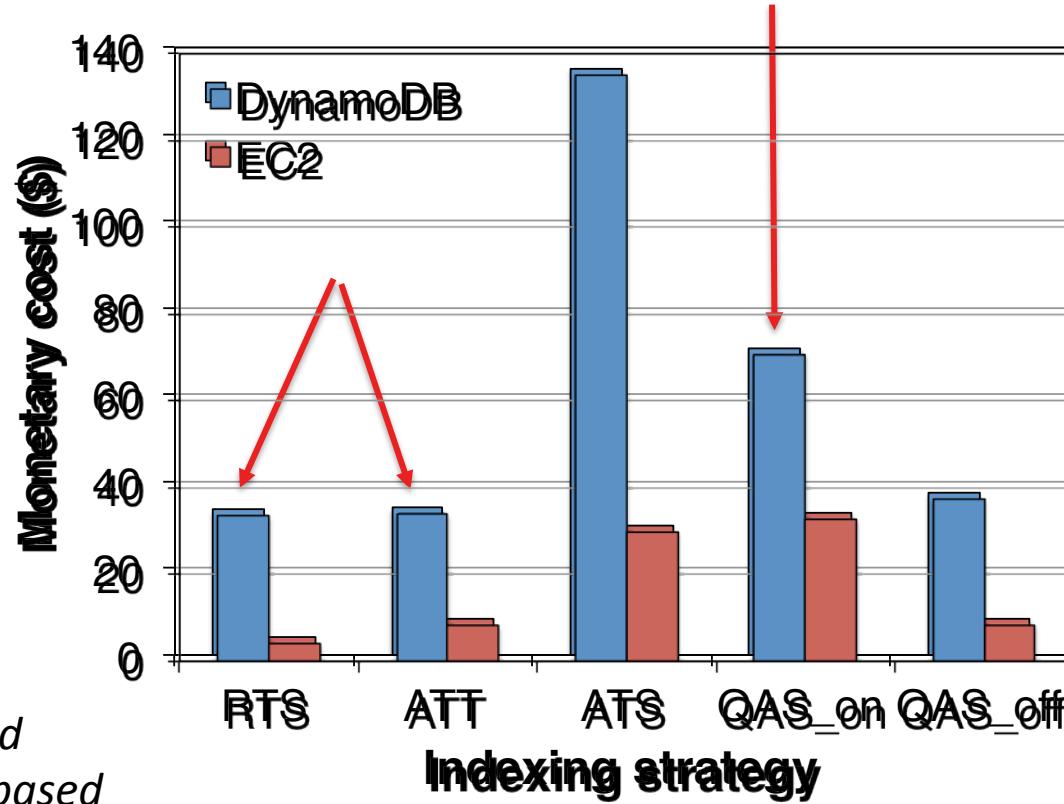
**ATS:** attribute-subset

**QAS:** answering from the index

on: w compression

off: w/o compression

# What about the money spent?



RTS: *term-based*

ATT: *attribute-based*

ATS: *attribute-subset*

QAS: *answering from the index*

on: *w compression*

off: *w/o compression*

## Summary: Hybrid RDF stores

Pros	Cons
Scalability	Depends on data partitioning
Fault-tolerance	Separation between storage and processing services
Elastic allocation of resources	Low latency for non-selective queries
Inter-query parallelism	No control on the services
Benefit from higher-level services (PaaS)	

# Classification of systems

Query processing	Centralized RDF store	Data storage		
		key-value store	DFS	
		Trinity.RDF		
	Other parallel	Partout WARP		
	MapReduce	MAPSIN	SHARD HadoopRDF RAPID+ PigSPARQL EAGRE	
Locally	Graph-partitioning	H2RDF+	Rya Statustore CumulusRDF	
	AMADA			

# Analysis dimensions

## 1. Data storage

- Key-value stores
- DFS (distributed file system)
- Single-site RDF stores or data storage services supplied by cloud providers

## 2. Query processing

- MapReduce-based
- Local processing (joins are usually performed at a single site)
- Distributed graph-based exploration

## 3. RDFS entailment

# Analysis dimensions

## 1. Data storage

- Key-value stores
- DFS (distributed file system)
- Single-site RDF stores or data storage services supplied by cloud providers

## 2. Query processing

- MapReduce-based
- Local processing (joins are usually performed at a single site)
- Distributed graph-based exploration

## 3. RDFS entailment

# Entailment in RDF

- RDF schema statements and set of entailment rules lead to implicit (entailed) RDF data

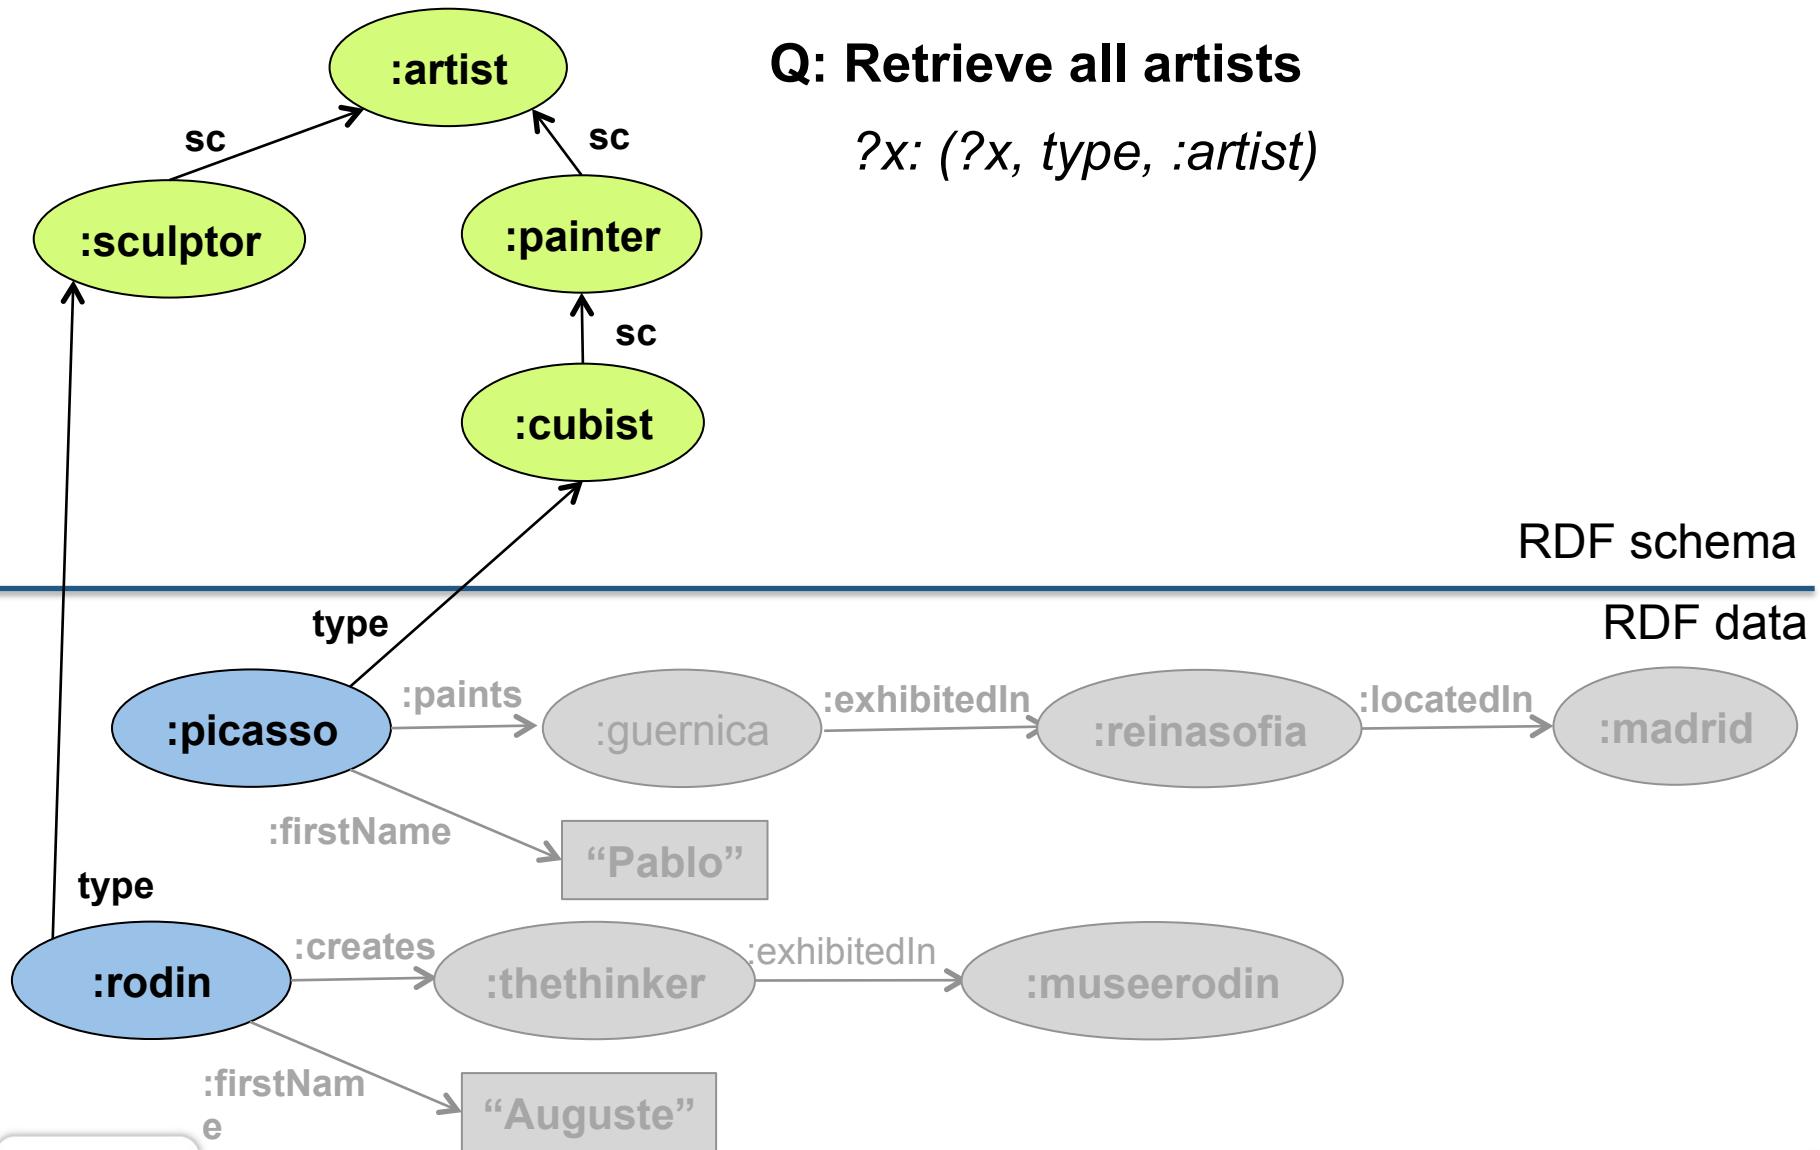
Sample rules

$$(X, \text{type}, A) \wedge (A, \text{sc}, B) \rightarrow (X, \text{type}, B)$$
$$(X, P, O) \wedge (P, \text{domain}, A) \rightarrow (X, \text{type}, A)$$

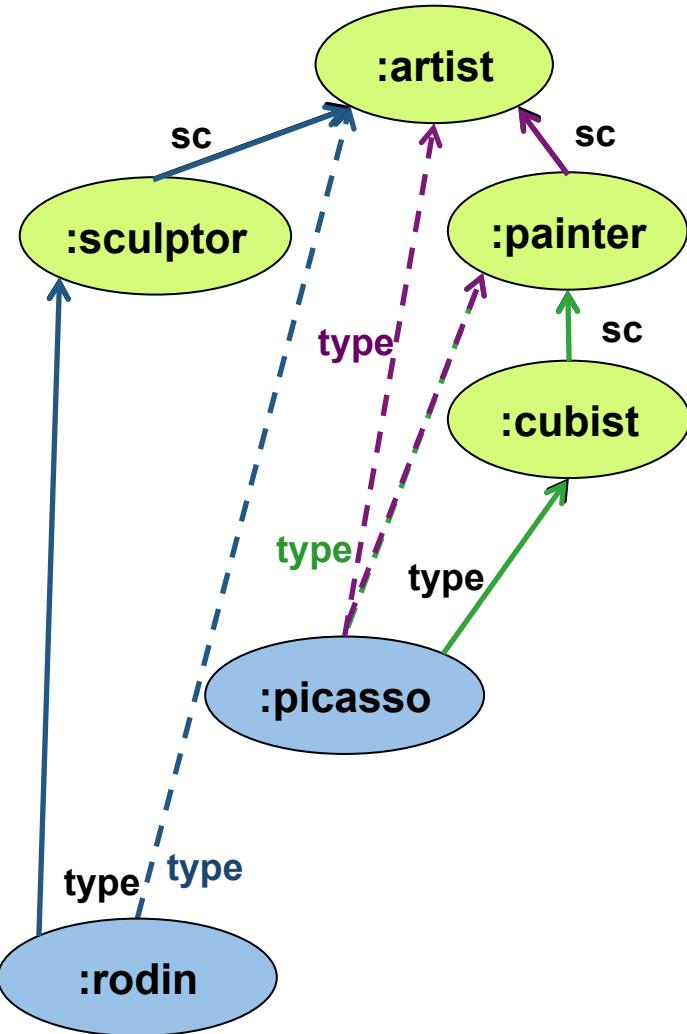
$$(:\text{picasso}, :p\text{aints}, :g\text{uernica}) \wedge (:p\text{aints}, \text{domain}, :\text{painter}) \rightarrow$$
$$(:\text{picasso}, \text{type}, :\text{painter})$$

- RDFS entailment : *Given an RDF(S) database DB and a triple t, is t entailed from DB (or does t logically follows from DB)?*

# Entailed triples and RDF querying



# Entailed triples and RDF querying



Q: Retrieve all artists ?x: ( $?x, \text{ type }, :artist$ )

A:	<table border="1"><tr><td><math>?x</math></td></tr><tr><td>-</td></tr></table>	$?x$	-
$?x$			
-			

no entailment

Rule:

$$(X, \text{ type }, A) \wedge (A, \text{ sc }, B) \rightarrow (X, \text{ type }, B)$$

Implicit / entailed triples

( $:rodin, \text{ type }, :artist$ )

( $:picasso, \text{ type }, :painter$ )

( $:picasso, \text{ type }, :artist$ )

A:	<table border="1"><tr><td><math>?x</math></td></tr><tr><td><math>:picasso</math></td></tr><tr><td><math>:rodin</math></td></tr></table>	$?x$	$:picasso$	$:rodin$
$?x$				
$:picasso$				
$:rodin$				

with entailment

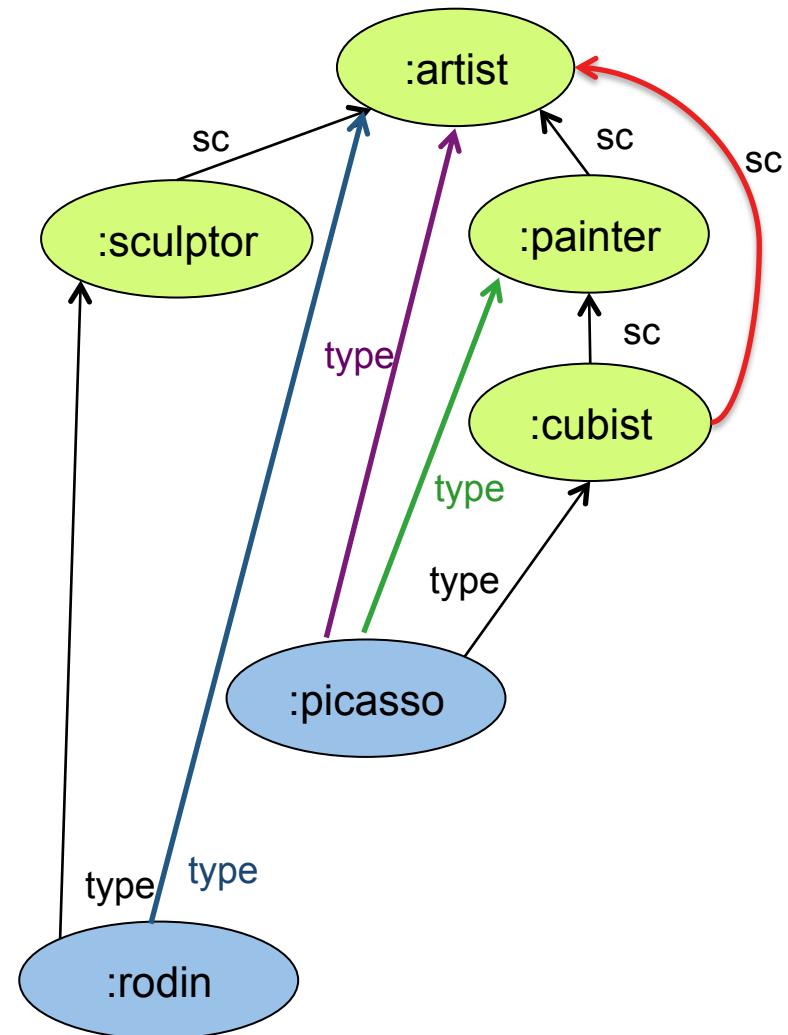
# RDFS entailment techniques

- RDFS closure
  - compute all implicit triples and store/materialize them in the database
- Query reformulation
  - transform/rewrite the query so that you get all the answers (even the ones that come from implicit triples)
- Hybrid approaches
  - compute the schema closure, and then reformulate the query
  - magic sets: based on the queries, pre-compute and store only the implicit triples that are required for the answer

# RDFS entailment – RDFS closure

Before querying

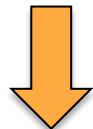
subject	property	object
:picasso	type	:cubist
:rodin	type	:sculptor
:sculptor	sc	:artist
:cubist	sc	:painter
:painter	sc	:artist
:cubist	sc	:painter
:rodin	type	:artist
:picasso	type	:painter
:picasso	type	:artist



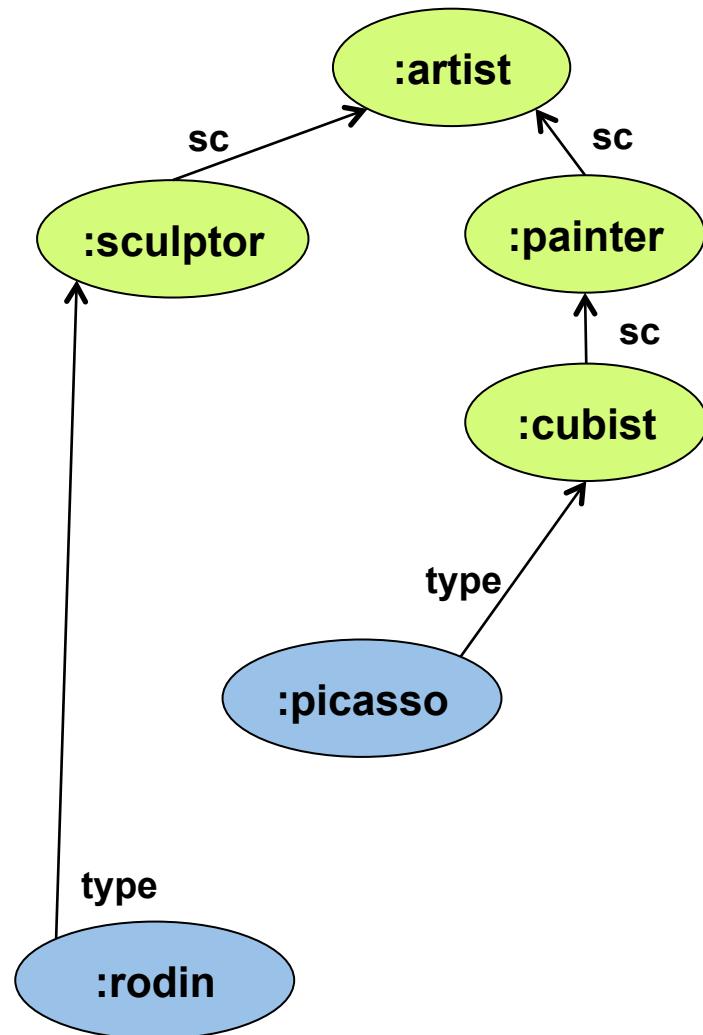
# RDFS entailment – query reformulation

At query time

?x: (?x, type, :artist)



?x: (?x, type, :artist)  $\vee$   
(?x, type, :sculptor)  $\vee$   
(?y, type, :painter)  $\vee$   
(?x, type, :cubist)



# RDFS entailment techniques

- RDFS closure
  - + Simple query evaluation
  - Storage overhead
  - Maintenance costs upon updates
- Query reformulation
  - + No storage overhead
  - + Robust to data/schema changes
  - Evaluation inefficient for complex queries
- Hybrid approaches combine the two
- For a comparison on these techniques:
  - centralized [Goasdoué13] and distributed [Kaoudi13]

# Categorization based on RDFS entailment

- I. Computing RDFS closure
- II. Approaches based on query reformulation
- III. Hybrid approaches

# I. Computing RDFS closure in the cloud

- RDFS closure computation in MapReduce Hadoop: **WebPie** [Urbani09]
  - RDF data is stored in HDFS
  - Compute all entailed triples using MapReduce
  - Optimization techniques
- Full RDFS closure computation in parallel (MPI-based) [Weaver09]
  - embarrassingly parallel algorithm
  - RDFS triples are kept in memory in each process

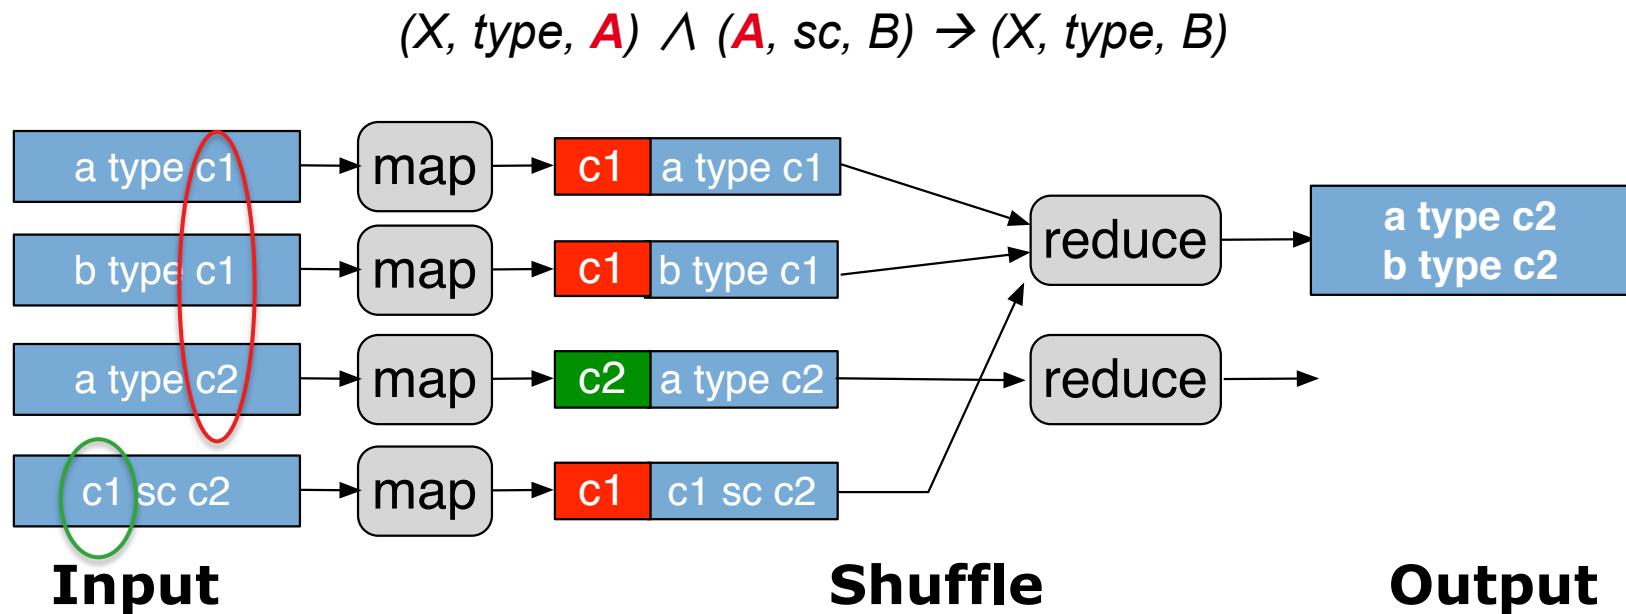
# How do we execute the RDFS rules in MapReduce?

Rule Name	If E contains:	then add:
rdfs1	uuu aaa lll. where lll is a plain literal (with or without a language tag).	_:nnn rdf:type rdfs:Literal . where _:nnn identifies a blank node allocated to lll by rule rule lg.
rdfs2	aaa rdfs:domain XXX . uuu aaa yyy .	uuu rdf:type XXX .
rdfs3	aaa rdfs:range XXX . uuu aaa vvv .	vvv rdf:type XXX .
rdfs4a	uuu aaa xxx .	
rdfs4b	uuu aaa vvv .	
rdfs5	uuu rdfs:subPropertyOf vvv vvv rdfs:subPropertyOf xxx .	
rdfs6	uuu rdf:type rdf:Property .	
rdfs7	aaa rdfs:subPropertyOf bbb . uuu aaa yyy .	
rdfs8	uuu rdf:type rdfs:Class .	uuu rdfs:subClassOf rdfs:Resource .
rdfs9	uuu rdfs:subClassOf XXX . vvv rdf:type uuu	vvv rdf:type XXX .
rdfs10	uuu rdf:type rdfs:Class .	uuu rdfs:subClassOf uuu .
rdfs11	uuu rdfs:subClassOf vvv vvv rdfs:subClassOf XXX .	uuu rdfs:subClassOf XXX .
rdfs12	uuu rdf:type rdfs:ContainerMembershipProperty .	uuu rdfs:subPropertyOf rdfs:member .
rdfs13	uuu rdf:type rdfs:Datatype .	uuu rdfs:subClassOf rdfs:Literal .

There is a common constant for rules with 2 atoms → Join

# RDFS closure in MapReduce [Urbani09]

- Implement as a join in MapReduce each rule that has two atoms in its body



# How do we execute the RDFS rules in MapReduce?

Rule Name	If E contains:	then add:
rdfs1	uuu aaa lll. where lll is a plain literal (with or without a language tag).	_:nnn rdf:type rdfs:Literal . where _:nnn identifies a blank node allocated to lll by rule rule lg.
rdfs2	aaa rdfs:domain XXX . uuu aaa yyy .	uuu rdf:type XXX .
rdfs3	aaa rdfs:range XXX . uuu aaa vvv .	vvv rdf:type XXX .
rdfs4a	uuu aaa xxx .	
rdfs4b	uuu aaa vvv .	
rdfs5	uuu rdfs:subPropertyOf vvv vvv rdfs:subPropertyOf XXX .	
rdfs6	uuu rdf:type rdf:Property	
rdfs7	aaa rdfs:subPropertyOf bbb uuu aaa yyy .	
rdfs8	uuu rdf:type rdfs:Class .	
rdfs9	uuu rdfs:subClassOf XXX . vvv rdf:type uuu .	vvv rdf:type XXX .
rdfs10	uuu rdf:type rdfs:Class .	uuu rdfs:subClassOf uuu .
rdfs11	uuu rdfs:subClassOf vvv . vvv rdfs:subClassOf XXX .	uuu rdfs:subClassOf XXX .
rdfs12	uuu rdf:type rdfs:ContainerMembershipProperty .	uuu rdfs:subPropertyOf rdfs:member .
rdfs13	uuu rdf:type rdfs:Datatype .	uuu rdfs:subClassOf rdfs:Literal .

Entailed triples are used as input in the rules → iterate jobs until no new data is produced (fixpoint)

# RDFS closure in MapReduce [Urbani09]

- Optimizations

# RDFS rules

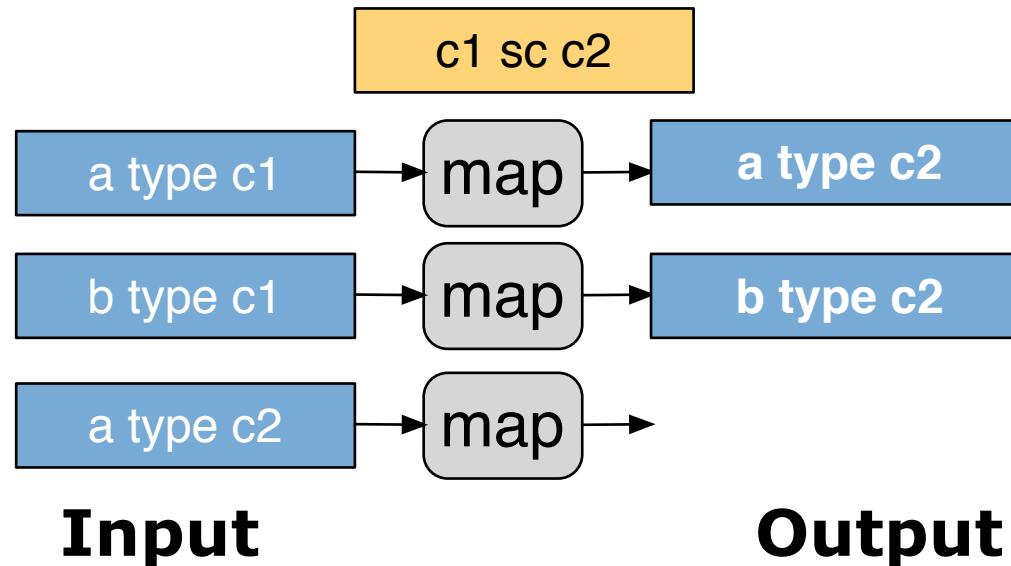
Rule Name	If E contains:	then add:
rdfs1	uuu aaa lll. where lll is a plain literal (with or without a language tag).	_:nnn rdf:type rdfs:Literal . where _:nnn identifies a blank node allocated to lll by rule rule lg.
rdfs2	aaa rdfs:domain XXX . uuu aaa yyy .	uuu rdf:type XXX .
rdfs3	aaa rdfs:range XXX . uuu aaa vvv .	vvv rdf:type XXX .
rdfs4a	uuu aaa xxx .	uuu rdf:type rdfs:Resource .
rdfs4b	uuu aaa vvv .	vvv rdf:type rdfs:Resource
rdfs5	uuu rdfs:subPropertyOf vvv . vvv rdfs:subPropertyOf xxx .	
rdfs6	uuu rdf:type rdf:Property .	
rdfs7	aaa rdfs:subPropertyOf bbb . uuu aaa yyy .	
rdfs8	uuu rdf:type rdfs:Class .	uuu rdfs:subClassOf rdfs:Resource .
rdfs9	uuu rdfs:subClassOf XXX . vvv rdf:type uuu .	vvv rdf:type XXX .
rdfs10	uuu rdf:type rdfs:Class .	uuu rdfs:subClassOf uuu .
rdfs11	uuu rdfs:subClassOf vvv . vvv rdfs:subClassOf XXX .	uuu rdfs:subClassOf XXX .
rdfs12	uuu rdf:type rdfs:ContainerMembershipProperty .	uuu rdfs:subPropertyOf rdfs:member .
rdfs13	uuu rdf:type rdfs:Datatype .	uuu rdfs:subClassOf rdfs:Literal .

At least 1 of the atoms is  
a schema triple!

# RDFS closure in MapReduce [Urbani09]

- Optimizations
  - Keep the RDF schema triples in memory of each node

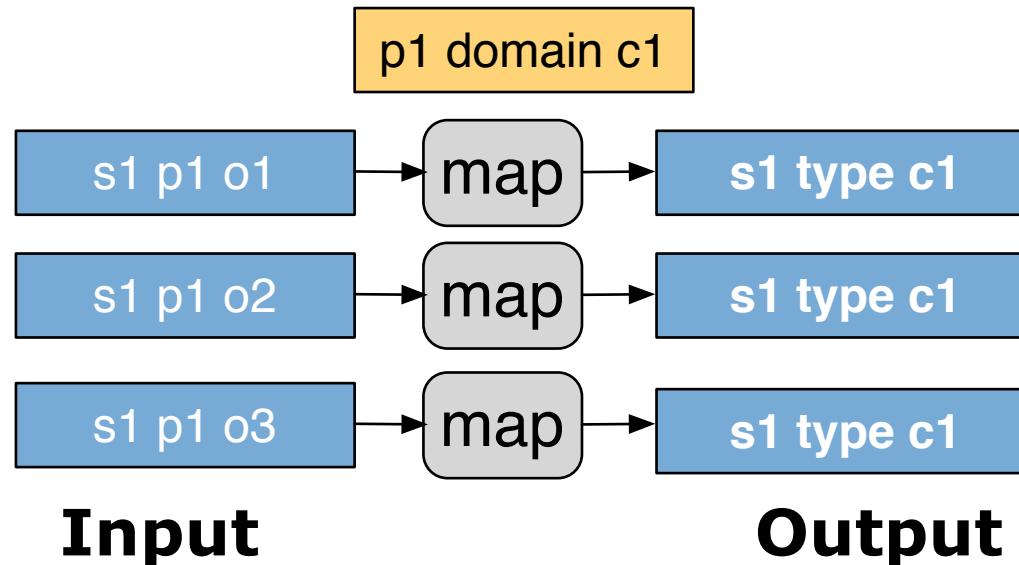
$$(X, \text{type}, A) \wedge (A, \text{sc}, B) \rightarrow (X, \text{type}, B)$$



# RDFS closure in MapReduce [Urbani09]

- Optimizations
  - Keep the RDF schema triples in memory of each node

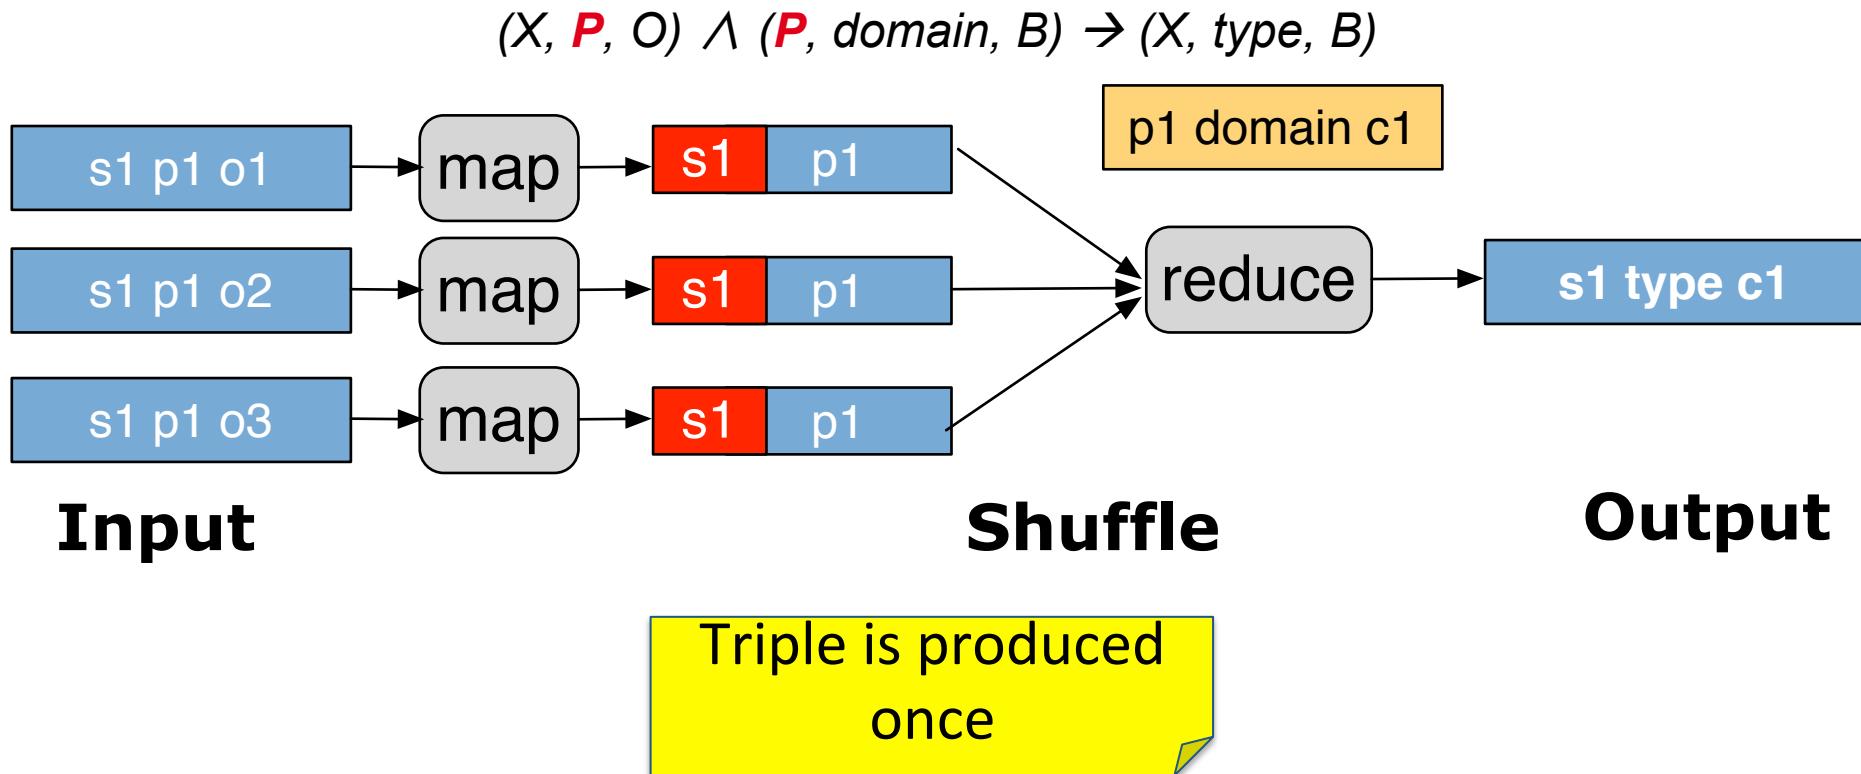
$$(X, \textcolor{red}{P}, O) \wedge (\textcolor{red}{P}, \textit{domain}, B) \rightarrow (X, \textit{type}, B)$$



Same triple is  
produced 3 times!

# RDFS closure in MapReduce [Urbani09]

- Optimizations
  - Keep the RDF schema triples in memory of each node
  - Data grouping to avoid duplicates



# RDFS closure in MapReduce [Urbani09]

- Optimizations
  - Keep the RDF schema triples in memory of each node
  - Data grouping to avoid duplicates
  - Ordering of rules to limit iterations

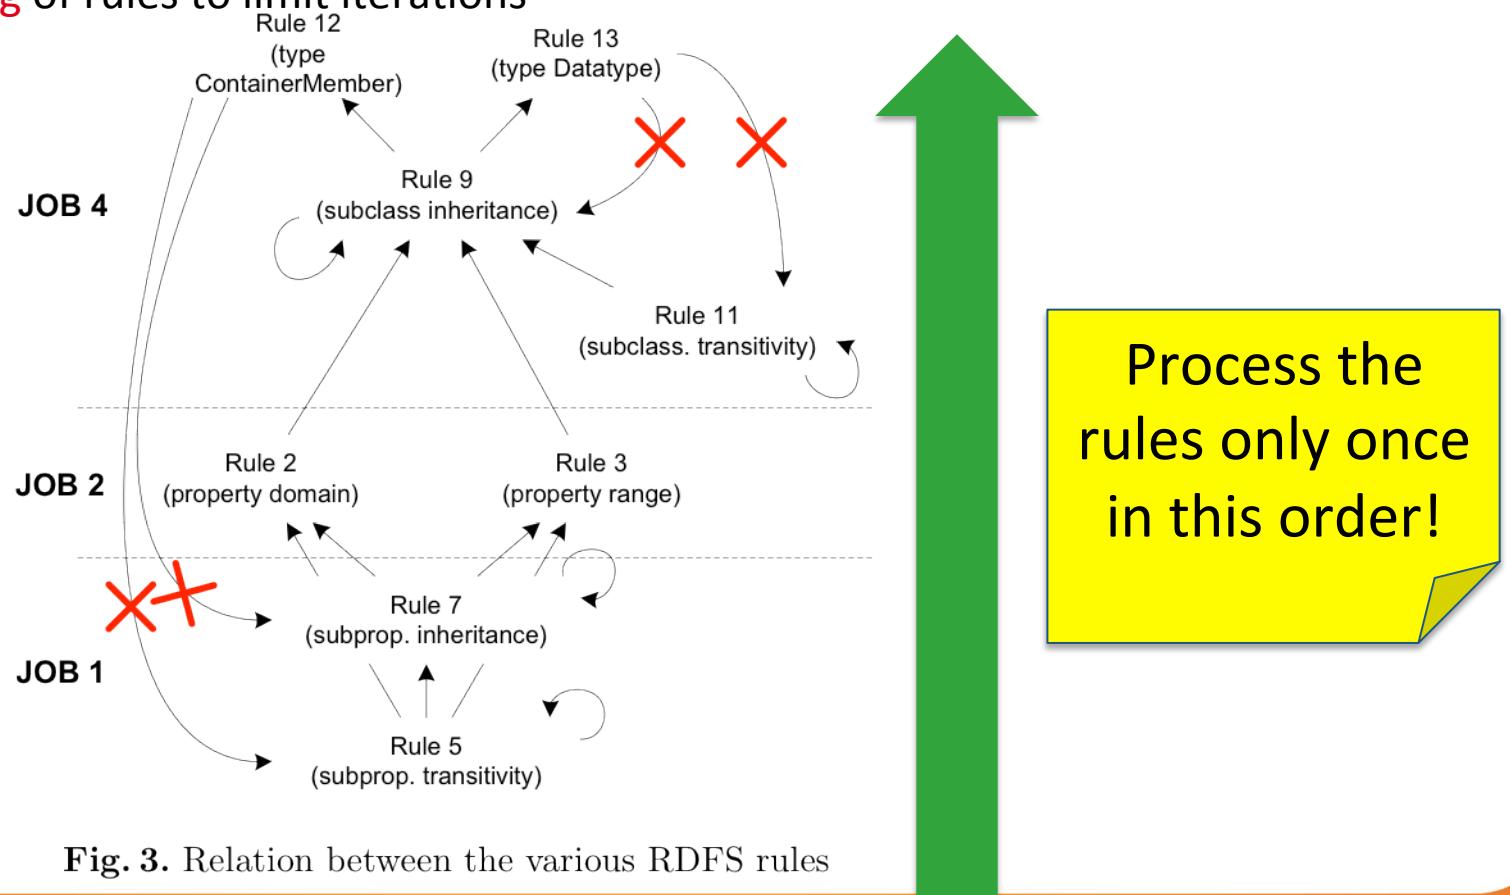


Fig. 3. Relation between the various RDFS rules

## II. Query reformulation in the cloud

- Query rewriting of BGP queries may lead to complex queries

$$(\exists x, \text{type}, \text{:artist}) \wedge (\exists x, \text{:creates}, ?y) \wedge (\exists y, \text{:exhibitedIn}, ?z)$$


$$\begin{aligned} &(\exists x, \text{type}, \text{:artist}) \vee \\ &(\exists x, \text{type}, \text{:sculptor}) \vee \\ &(\exists y, \text{type}, \text{:painter}) \vee \\ &(\exists x, \text{type}, \text{:cubist}) \end{aligned}$$

$$(\exists x, \text{:creates}, ?y) \vee (\exists x, \text{:paints}, ?y) \vee$$

- Conjunction of unions of atomic queries

- HadoopRDF [Husain11]

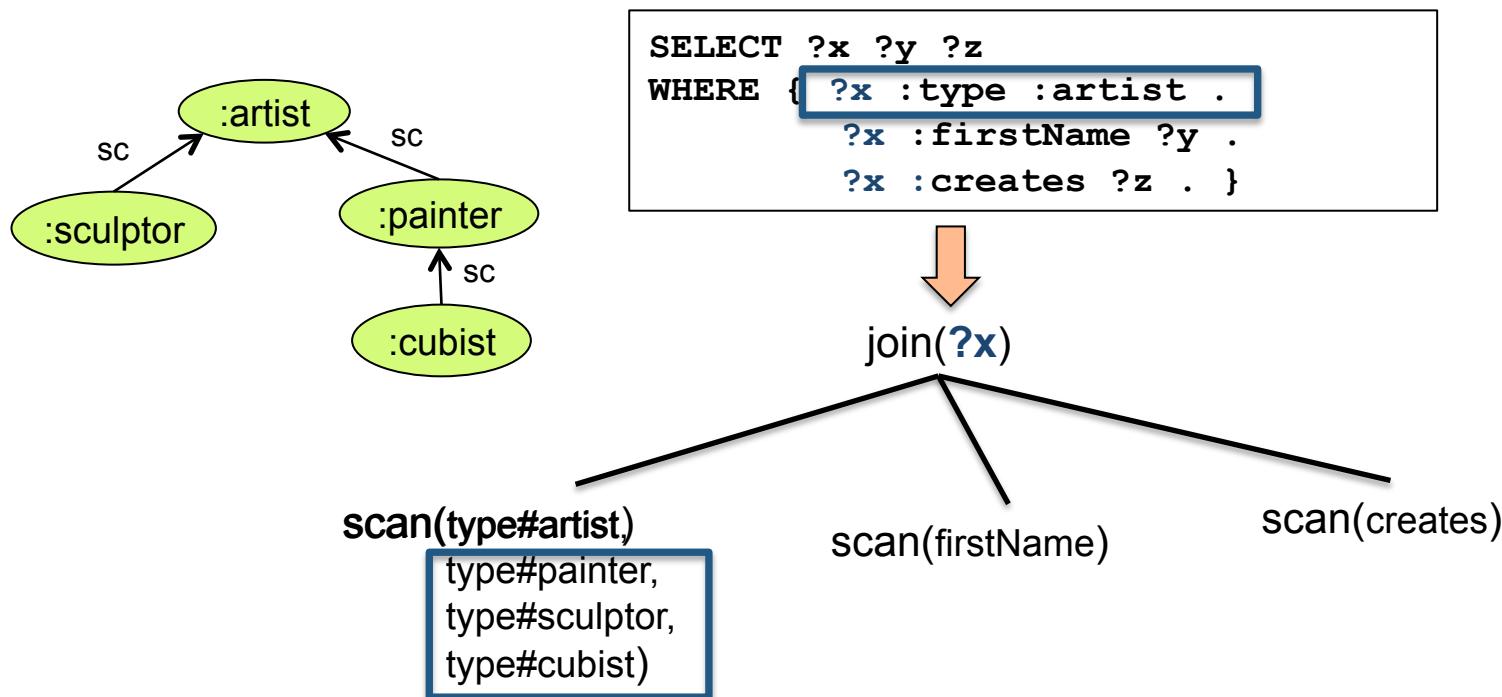
- Or union of conjunctive queries?

$$\begin{aligned} &[(\exists x, \text{type}, \text{:artist}) \wedge (\exists x, \text{:creates}, ?y) \wedge (\exists y, \text{:exhibitedIn}, ?z)] \vee \\ &[(\exists x, \text{type}, \text{:artist}) \wedge (\exists x, \text{:paints}, ?y) \wedge (\exists y, \text{:exhibitedIn}, ?z)] \vee \\ &[(\exists x, \text{type}, \text{:sculptor}) \wedge (\exists x, \text{:creates}, ?y) \wedge (\exists y, \text{:exhibitedIn}, ?z)] \vee \\ &[(\exists x, \text{type}, \text{:sculptor}) \wedge (\exists x, \text{:paints}, ?y) \wedge (\exists y, \text{:exhibitedIn}, ?z)] \vee \\ &\dots \end{aligned}$$

- Transitive closure computation and recursive query processing in MapReduce [Afrati12, Bu12] may apply to RDFS reasoning

# Query reformulation in HadoopRDF [Husain11]

- Query reformulation during the file selection of the query processing



## III. Hybrid approaches

- Precompute the closure of the schema
  - Schema is usually much smaller than the data
  - Schema does not change often
  - It is always used in the RDFS rules
- Reformulate the query (faster reformulations)
- Representative works
  - QueryPie [Urbani11]
  - Rya [Punnoose12]

# Analysis dimensions

## 1. Data storage

- Key-value stores
- DFS (distributed file system)
- Single-site RDF stores or data storage services supplied by cloud providers

## 2. Query processing

- MapReduce-based
- Local processing (joins are usually performed at a single site)
- Distributed graph-based exploration

## 3. RDFS entailment

- Materialization of RDFS closure
- Query reformulation
- Hybrid (only the schema closure is precomputed)

# 4

## SUMMARY AND OPEN ISSUES

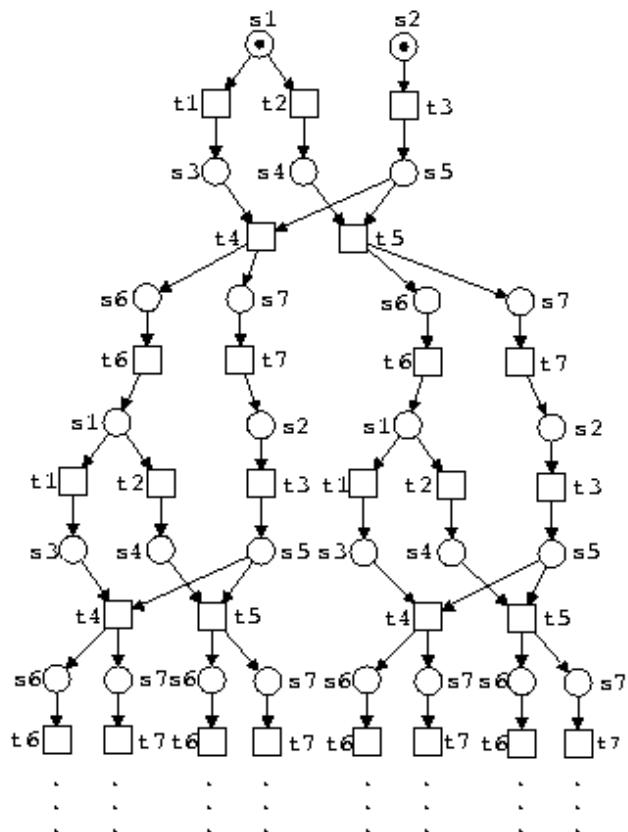
# Summary

Main problems to be solved in a (distributed) RDF store

- Storage
- Query evaluation
- Reasoning (for query answering)

Plenty of techniques for centralized and distributed settings

- Oftentimes incomparable  
(partial order)
- Zoo!



## Do we have a winner?

- No, at least not yet ☺
- Performance depends on data, workload...
  - Star queries or path queries?
  - Selective or analytics-style queries?
- Benchmarks on how **real** SPARQL queries look:
  - [Duan11]
  - [Arias11]
  - [Picalausa11]
  - Linked Data benchmark council ([ldbc.eu](http://ldbc.eu))

# Where do we go from here?

- Room for optimization
  - Improving query response time in the first place
  - Applying traditional optimization techniques (join ordering, statistics, etc.)?
- RDF data partitioning
  - Are the existing graph partitioning algorithms satisfactory?
- RDFS entailment
  - Query reformulation
- More SPARQL features
  - Blank nodes, optional, negation, property paths, etc.

# Where do we go from here?

- RDF **views/indexes** in the cloud
- SPARQL multi-query optimization in the cloud
- RDF **updates** in the cloud
- RDF data **analytics**
- How does cloud **variability** affect RDF data management?

# Thank you /Questions?

Companion paper:



New!

“RDF in the clouds: A Survey”, VLDB Journal

<http://web.imis.athena-innovation.gr/~zoi/>

<http://pages.saclay.inria.fr/ioana.manolescu>

# References

- [Abadi07] D. J. Abadi, A. Marcus, S. Madden, K. J. Hollenbach, “Scalable Semantic Web Data Management Using Vertical Partitioning”, in VLDB 2007.
- [Afrati10] F. N. Afrati and J. D. Ullman, “Optimizing Joins in a Map-Reduce Environment,” in EDBT 2010.
- [Afrati11a] F. N. Afrati and J. D. Ullman, “Optimizing Multiway Joins in a Map-Reduce Environment,” IEEE Trans. Knowl. Data Eng., 2011.
- [Afrati11b] F. N. Afrati, V. R. Borkar, M. J. Carey, N. Polyzotis, and J. D. Ullman, “Map-Reduce Extensions and Recursive Queries,” in EDBT 2011.
- [Afrati12] F. N. Afrati, J. D. Ullman, “Transitive closure and recursive Datalog implemented on clusters” in EDBT 2012.
- [Alexaki01] S. Alexaki, V. Christophides, G. Karvounarakis, D. Plexousakis. “On Storing Voluminous RDF Descriptions: The Case of Web Portal Catalogs”, in WebDB 2001.
- [Aranda12] A. Aranda-Andujar, F. Bugiotti, J. Camacho-Rodriguez, D. Colazzo, F. Goasdoué, Z. Kaoudi, and I. Manolescu, “Amada: Web Data Repositories in the Amazon Cloud (demo)”, in CIKM 2012.
- [Arias11] M. Arias, J. D. Fernandez, M. A. Martínez-Prieto, “An Empirical Study of Real-World SPARQL Queries”, in USEWOD 2011.
- [Battre06] D. Battre, A. Hoing, F. Heine, O. Kao, “On Triple Dissemination, Forward-Chaining, and Load Balancing in DHT based RDF stores”, in DBISP2P 2006.
- [Blanas10] S. Blanas, J. M. Patel, V. Ercegovac, J. Rao, E. J. Shekita, and Y. Tian, “A Comparison of Join Algorithms for Log Processing in MapReduce,” in SIGMOD 2010.

# References

- [Bu12] Y. Bu, B. Howe, M. Balazinska, and M. D. Ernst, “The HaLoop Approach to Large-Scale Iterative Data Analysis,” VLDB J, 2012.
- [Bugiotti12] F. Bugiotti, F. Goasdoué, Z. Kaoudi, and I. Manolescu, “RDF Data Management in the Amazon Cloud,” in ICDT/EDBT Workshops 2012.
- [Cattell11] R. Cattell, “Scalable SQL and NoSQL data stores,” SIGMOD Record, May 2011.
- [Dean13] J. Dean, L. A. Barroso, “Tail at Scale”, communication of the ACM, February 2013.
- [Duan11] S. Duan, A. Kementsietsidis, K. Srinivas, O. Udrea, “Apples and oranges: a comparison of RDF benchmarks and real RDF datasets”, in SIGMOD 2011.
- [Galarraga12] L. Galarraga, L. Hose, R. Schenkel, “Partout:A Distributed Engine for Efficient RDF Processing”, Technical Report, 2012.
- [Goasdoué13] F. Goasdoué, I. Manolescu, and A. Roatis, “Efficient Query Answering against Dynamic RDF Databases”, in EDBT 2013.
- [Gurajada14] S. Gurajada, S. Seufert, I. Miliaraki, M. Theobald TriAD: a distributed shared-nothing RDF engine based on asynchronous message passing
- [Hayes04] P. Hayes, “RDF Semantics”, W3C Recommendation, February 2004, <http://www.w3.org/TR/rdf-mt/>.
- [Hose13] K. Hose, and R. Schenkel, “WARP: Workload-Aware Replication and Partitioning for RDF”, in DESWEB 2013.
- [Huang11] J. Huang, D. J. Abadi, and K. Ren, “Scalable SPARQL Querying of Large RDF Graphs,” PVLDB 2011.

# References

- [Husain11] M. Husain, J. McGlothlin, M. M. Masud, L. Khan, and B. M. Thuraisingham, “Heuristics-Based Query Processing for Large RDF Graphs Using Cloud Computing,” IEEE Trans. on Knowl. and Data Eng., 2011.
- [Iosup11] A. Iosup, N. Yigitbasi, D.H. J. Epema, “On the Performance Variability of Production Cloud Services”, in CCGRID 2011.
- [Kaoudi13] Z. Kaoudi, M. Koubarakis, “Distributed RDFS Reasoning over Structured Overlay Networks”, Journal on Data Semantics, 2013.
- [Kotoulas10] S. Kotoulas, E. Oren, F. Harmelen, “Mind the data skew: distributed inferencing by speeddating in elastic regions”, in WWW 2010.
- [Lee13] K. Lee and Ling Liu, “Scaling Queries over Big RDF Graphs with Semantic Hash Partitioning”, in PVLDB 2013.
- [Mallea11] A. Mallea, M. Arenas, A. Hogan, A. Polleres, “On Blank Nodes”, in ISWC 2011.
- [METIS] METIS - Serial Graph Partitioning and Fill-reducing Matrix Ordering:  
<http://glaros.dtc.umn.edu/gkhome/metis/metis/overview>
- [Neumann10] T. Neumann and G. Weikum, “The RDF-3X Engine for Scalable Management of RDF Data,” VLDBJ 2010.
- [Papailiou12] N. Papailiou, I. Konstantinou, D. Tsoumakos, N. Koziris, “H2RDF: adaptive query processing on RDF data in the cloud”, in WWW demo 2012.
- [Papailiou12] N. Papailiou, I. Konstantinou, D. Tsoumakos, P. Karras, N. Koziris, “H2RDF+: High-performance distributed joins over large-scale RDF graphs”, in BigData Conference 2013.

# References

- [Picalausa11] F. Picalausa, S. Vansumeren, “What are real SPARQL queries like?”, in SWIM 2011.
- [Punnoose12] R. Punnoose, A. Crainiceanu, and D. Rapp, “Rya: A Scalable RDF Triple Store for the Clouds,” in 1st International Workshop on Cloud Intelligence (in conjunction with VLDB), 2012.
- [Ravindra11] P. Ravindra, H. Kim, K. Anyanwu, “An Intermediate Algebra for Optimizing RDF Graph Pattern Matching on MapReduce”, in ESWC 2011.
- [Rohloff10] K. Rohloff and R. E. Schantz, “High-Performance, Massively Scalable Distributed Systems using the MapReduce Software Framework: the SHARD Triple-Store,” in Programming Support Innovations for Emerging Distributed Applications, 2010.
- [Schad10] J. Schad, J. Dittrich, J. Quiané-Ruiz, “Runtime Measurements in the Cloud: Observing, Analyzing, and Reducing Variance”, PVLDB 2010.
- [Schätzle12] A. Schätzle, M. Przyjaciel-Zablocki, C. Dorner, T. Hornung, G. Lausen, “Cascading Map-Side Joins over HBase for Scalable Join Processing”, in SSWS+HPCSW 2012.
- [Schätzle11] A. Schätzle, M. Przyjaciel-Zablocki, G. Lausen, ”PigSPARQL: mapping SPARQL to Pig Latin”, in SWIM 2011.
- [Stein10] R. Stein and V. Zacharias, “RDF On Cloud Number Nine,” in 4<sup>th</sup> Workshop on New Forms of Reasoning for the Semantic Web: Scalable and Dynamic, May 2010.
- [Theoharis05] Y. Theoharis, V. Christophides, and G. Karvounarakis. “Benchmarking Database Representations of RDF/S Stores”, in ISWC 2005.
- [Urbani09] J. Urbani, S. Kotoulas, E. Oren, and F. van Harmelen, “Scalable Distributed Reasoning using MapReduce,” in ISWC 2009.

# References

- [Urbani11] J. Urbani, F. van Harmelen, S. Schlobach, and H. Bal, “QueryPIE: Backward Reasoning for OWL Horst over Very Large Knowledge Bases,” in ISWC 2011.
- [Weaver09] J. Weaver and J. A. Hendler, “Parallel Materialization of the Finite RDFS Closure for Hundreds of Millions of Triples”, in ISWC 2009.
- [Weiss08] C. Weiss, P. Karras, A. Bernstein, “Hexastore: sextuple indexing for semantic web data management”, PVLDB 2008.
- [Wilkinson06] K. Wilkinson, “Jena Property Table Implementation”, in SSWS 2006.
- [Wu11] S. Wu, F. Li, S. Mehrotra, B. C. Ooi, “Query Optimization for Massively Parallel Data Processing”, in SOCC 2011.
- [Zeng13] K. Zeng, J. Yang, H. Wang, B. Shao, and Z. Wang, “A Distributed Graph Engine for Web Scale RDF Data”, in PVLDB 2013.
- [Zhang13] X. Zhang, L. Chein, Y. Tong, and M. Wang, “EAGRE: Towards Scalable I/O Efficient SPARQL Query Evaluation on the Cloud”, in ICDE 2013.