



# Triples in the clouds

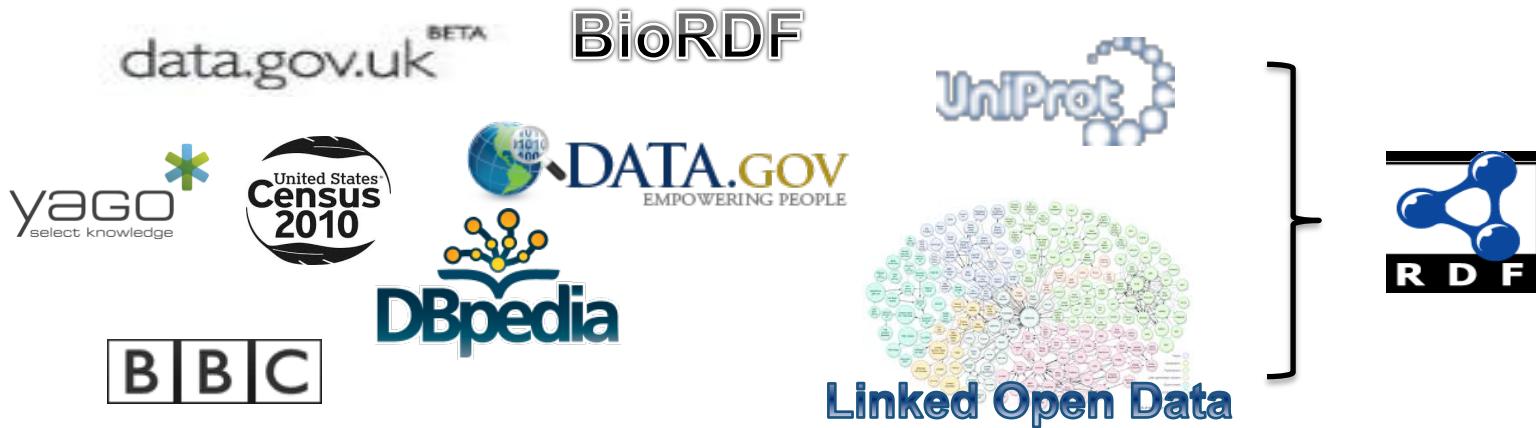
Zoi Kaoudi and Ioana Manolescu  
Inria Saclay, France

**ICDE seminar 2013**

[http://pages.saclay.inria.fr/zoi.kaoudi/rdfcloud\\_icde13.pdf](http://pages.saclay.inria.fr/zoi.kaoudi/rdfcloud_icde13.pdf)

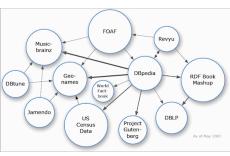
# Where is RDF found today?

- Semantic Web applications
  - Commercial database systems (Oracle, IBM, etc.)
  - Open Data applications
  - Linked Data applications
  - Many available data sources in RDF
- } **Linked Open Data**

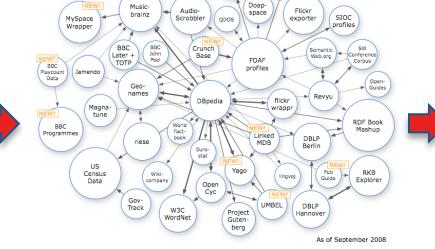


# RDF is growing

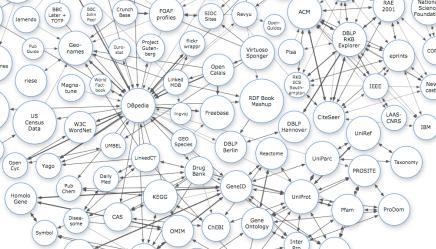
## LOD cloud



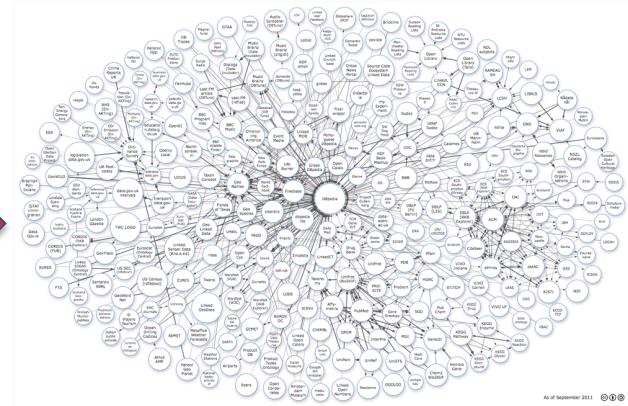
2007



2008



2009



2011



# RDF data sources in numbers



RDF-encoded Wikipedia **1.89 billion** triples

RDF-encoded biological data **2.7 billion** triples

US government data in RDF **5 billion** triples

crawled Web data **2 billion** triples

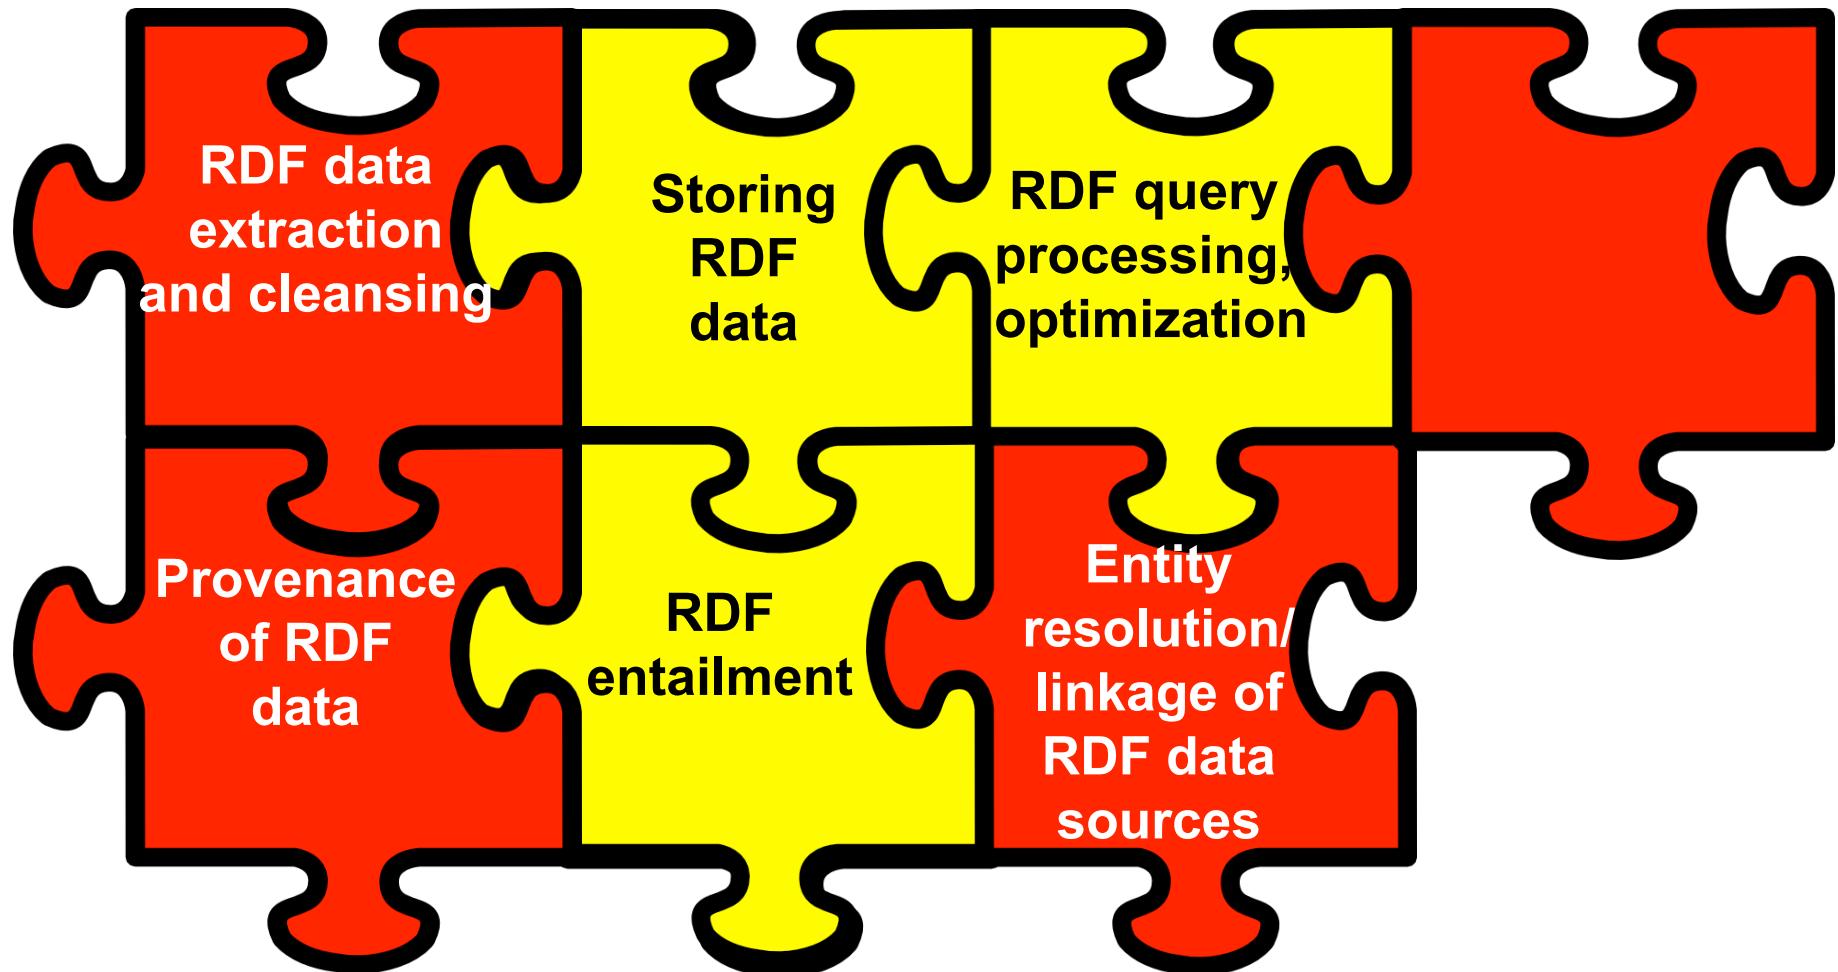
US population statistics **1 billion** triples

facts from Wikipedia, Wordnet,  
Geonames **0.12 billion** triples

**Linked Open Data cloud** **31 billion** triples, Sept. 2011

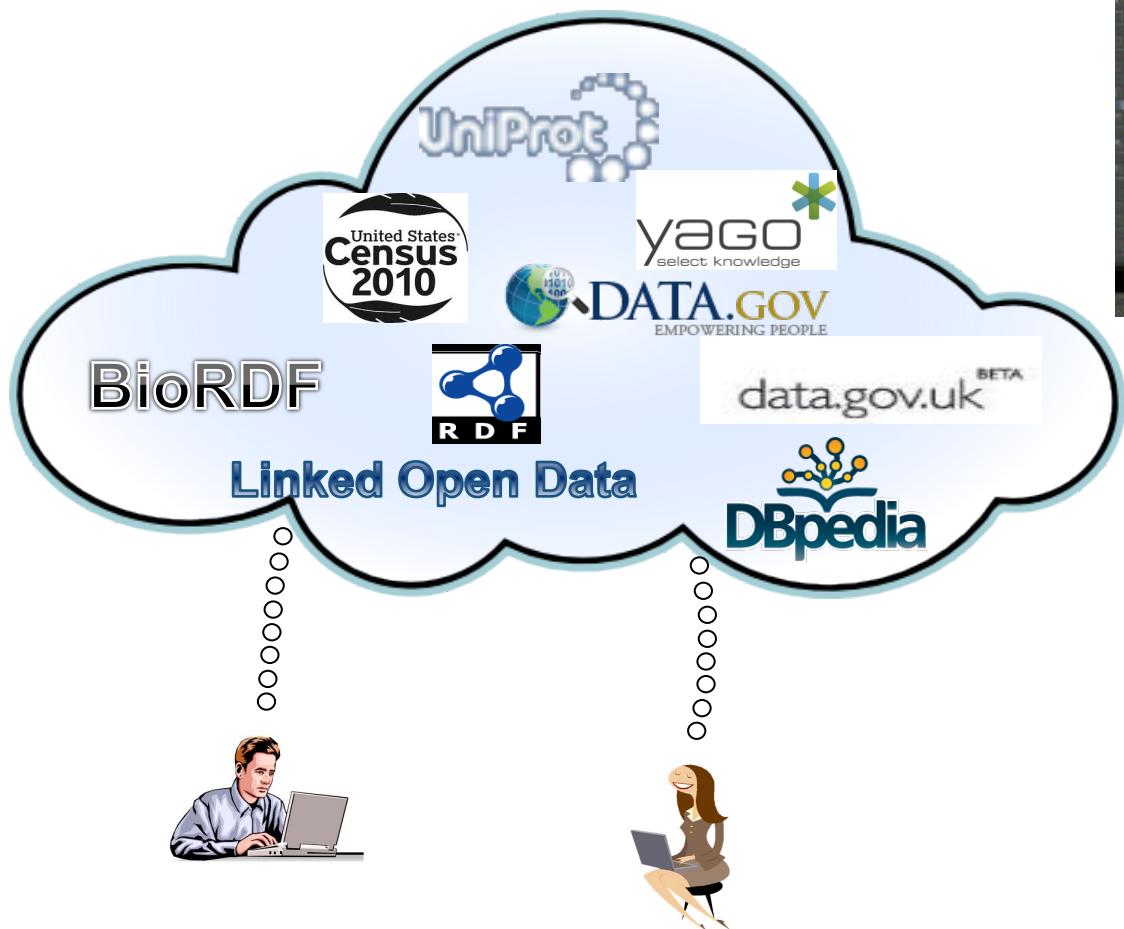
... and many more ... <http://www.w3.org/wiki/DataSetRDFDumps>  
<http://datahub.io/>

# RDF data management tasks



# Managing large volumes of RDF data

## Cloud computing



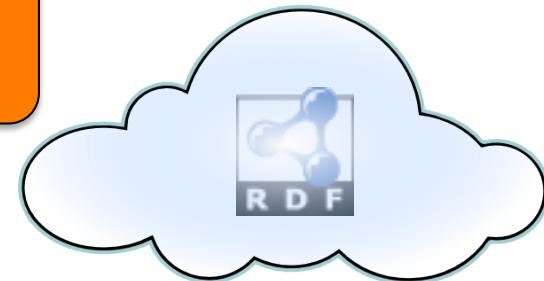
**Scalability**

**Elasticity**

**Fault-tolerance**

# What this seminar is about

## Triples in the clouds



- Storage RDF data in the cloud
- Evaluation of RDF queries in the cloud
- Inference of RDF data in the cloud

# Outline

1. RDF basics
2. Cloud basics
3. Analysis of RDF systems in the cloud
  - Storage
  - Querying
  - Inference
4. Open issues

# 1

## RDF BASICS

# Resource Description Framework (RDF)



- RDF: representation of **resources** on the Web identified by Universal Resource Identifiers (**URIs**) [http://live.dbpedia.org/page/Pablo\\_Picasso](http://live.dbpedia.org/page/Pablo_Picasso)
- RDF data: facts = **triples**  $(s, p, o)$   
A diagram illustrating an RDF triple. The triple is represented as  $(s, p, o)$ . Below this, the words "subject", "property", and "object" are written in a horizontal line. Three blue arrows point from the letters "s", "p", and "o" respectively to the words "subject", "property", and "object".
- **Literals**: strings, numbers, dates, etc.
  - Only in the object position
- RDF **schema** (RDFS): specifies **concepts**, **properties** and **relationships** between them <http://live.dbpedia.org/ontology/Artist>

`(dbpedia:Pablo_Picasso, rdf:type, dbpedia-owl:Artist)`

`(dbpedia-owl:Artist, rdfs:subClassOf, dbpedia-owl:Person)`

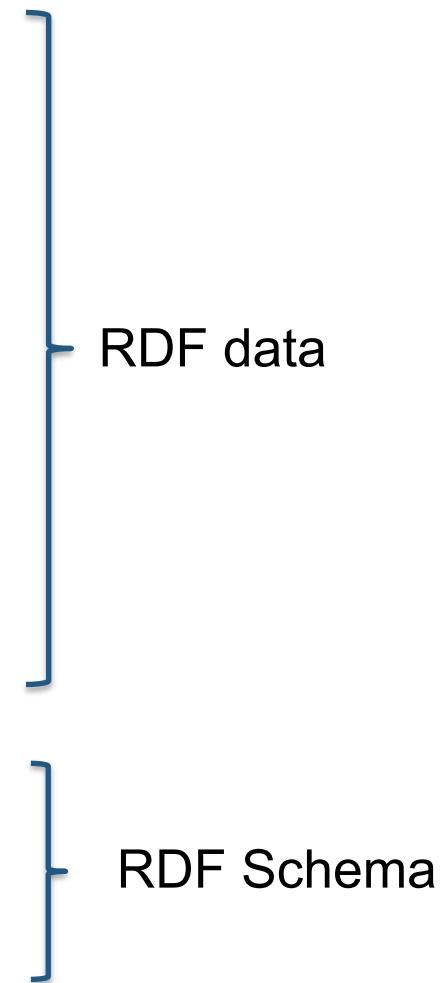
Notation:

`( :Pablo_Picasso, type, :Artist)`

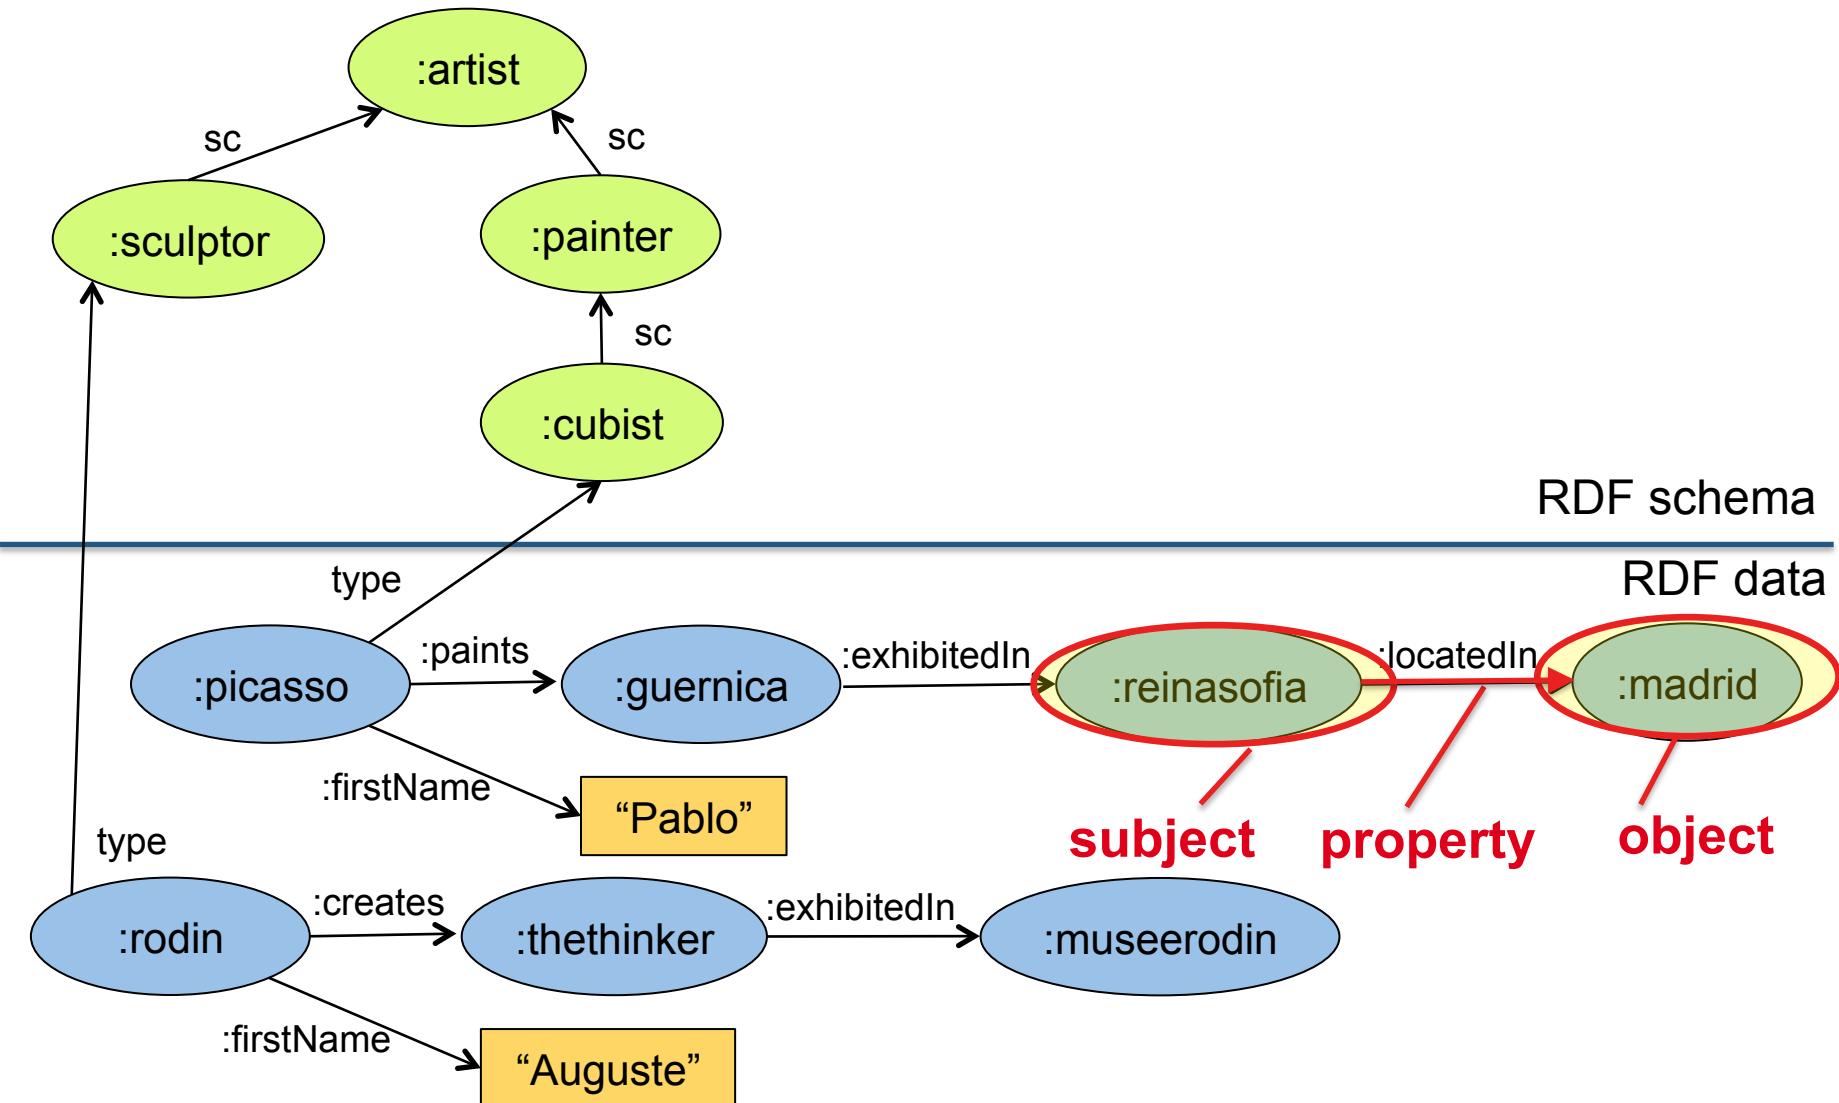
`( :Artist, sc, :Person)`

# Running example: RDF database (triple syntax)

```
:picasso type :cubist .  
:picasso :firstName "Pablo"  
:picasso :paints :guernica .  
:guernica :exhibitedIn :reinasofia .  
:reinasofia :locatedIn :madrid .  
:rodin type :sculptor .  
:rodin :firstname "Auguste" .  
:rodin :creates :thethinker .  
:thethinker :exhibitedIn :museerodin .  
  
:sculptor sc :artist .  
:painter sc :artist .  
:cubist sc :painter.
```



# Running example: RDF database



# SPARQL Query Language

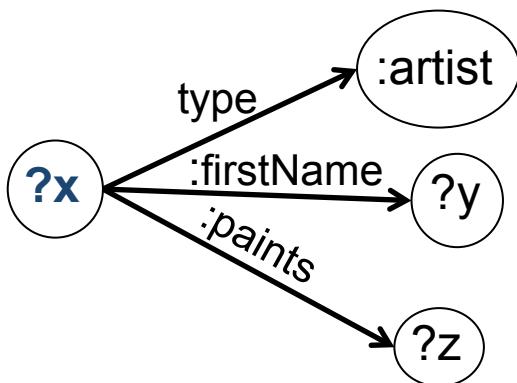


- Standard language to query RDF data
- SQL-based syntax
- Basic building block: **triple pattern**  $(s, p, o)$   
A diagram showing a triple pattern  $(s, p, o)$ . Below the pattern, three arrows point downwards from the variables  $s$ ,  $p$ , and  $o$  to the text "constant or variable".
- Basic graph pattern (BGP) queries
  - conjunctive triple pattern queries
  - $\underbrace{?x_1, ?x_2, \dots, ?x_n}_{\text{head variables}} : \underbrace{(s_1, p_1, o_1) \wedge \dots \wedge (s_k, p_k, o_k)}_{\text{triple patterns}}$
- SPARQL 1.1: W3C recommendation, March 2013

# SPARQL examples

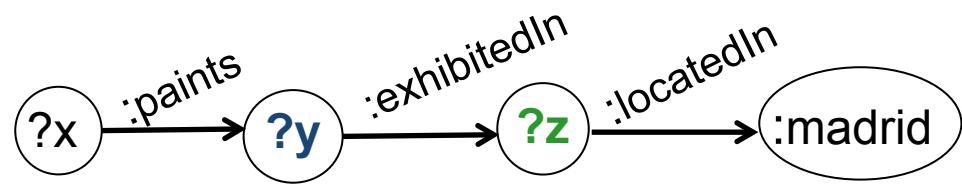
## Star query

```
SELECT ?y ?z  
WHERE {?x type :artist .  
       ?x :firstName ?y .  
       ?x :paints ?z .}
```


$$\text{?y ?z: } (\text{?x type :artist} \wedge  
(\text{?x :firstName ?y}) \wedge  
(\text{?x :paints ?z}))$$

## Path query

```
SELECT ?x ?y ?z  
WHERE {?x :paints ?y .  
       ?y :exhibitedIn ?z .  
       ?z :locatedIn :madrid.}
```


$$\text{?x ?y ?z: } (\text{?x :paints ?y}) \wedge  
(\text{?y :exhibitedIn ?z}) \wedge  
(\text{?z :locatedIn :madrid})$$

# Centralized RDF stores

- How do we usually store RDF in centralized environment?
- Approaches relying on RDBMs
  - Big **triple tables**

Triples

subject	property	object
:picasso	type	:cubist
:picasso	:firstName	“Pablo”
picasso	:paints	:guernica
:guernica	:exhibitedIn	:reinasofia
:reinasofia	:locatedIn	:madrid
:rodin	type	:sculptor
:rodin	:creates	:thethinker
...	...	...

# Centralized RDF stores

- How do we usually store RDF in centralized environment?
- Approaches relying on RDBMs
  - Big triple tables
  - **Property tables** [Alexaki01, Wilkison06]

Artists

subject	type	firstName	paints	creates
:picasso	:cubist	“Pablo”	:guernica	null
:rodin	:sculptor	“Auguste”	null	:thethinker

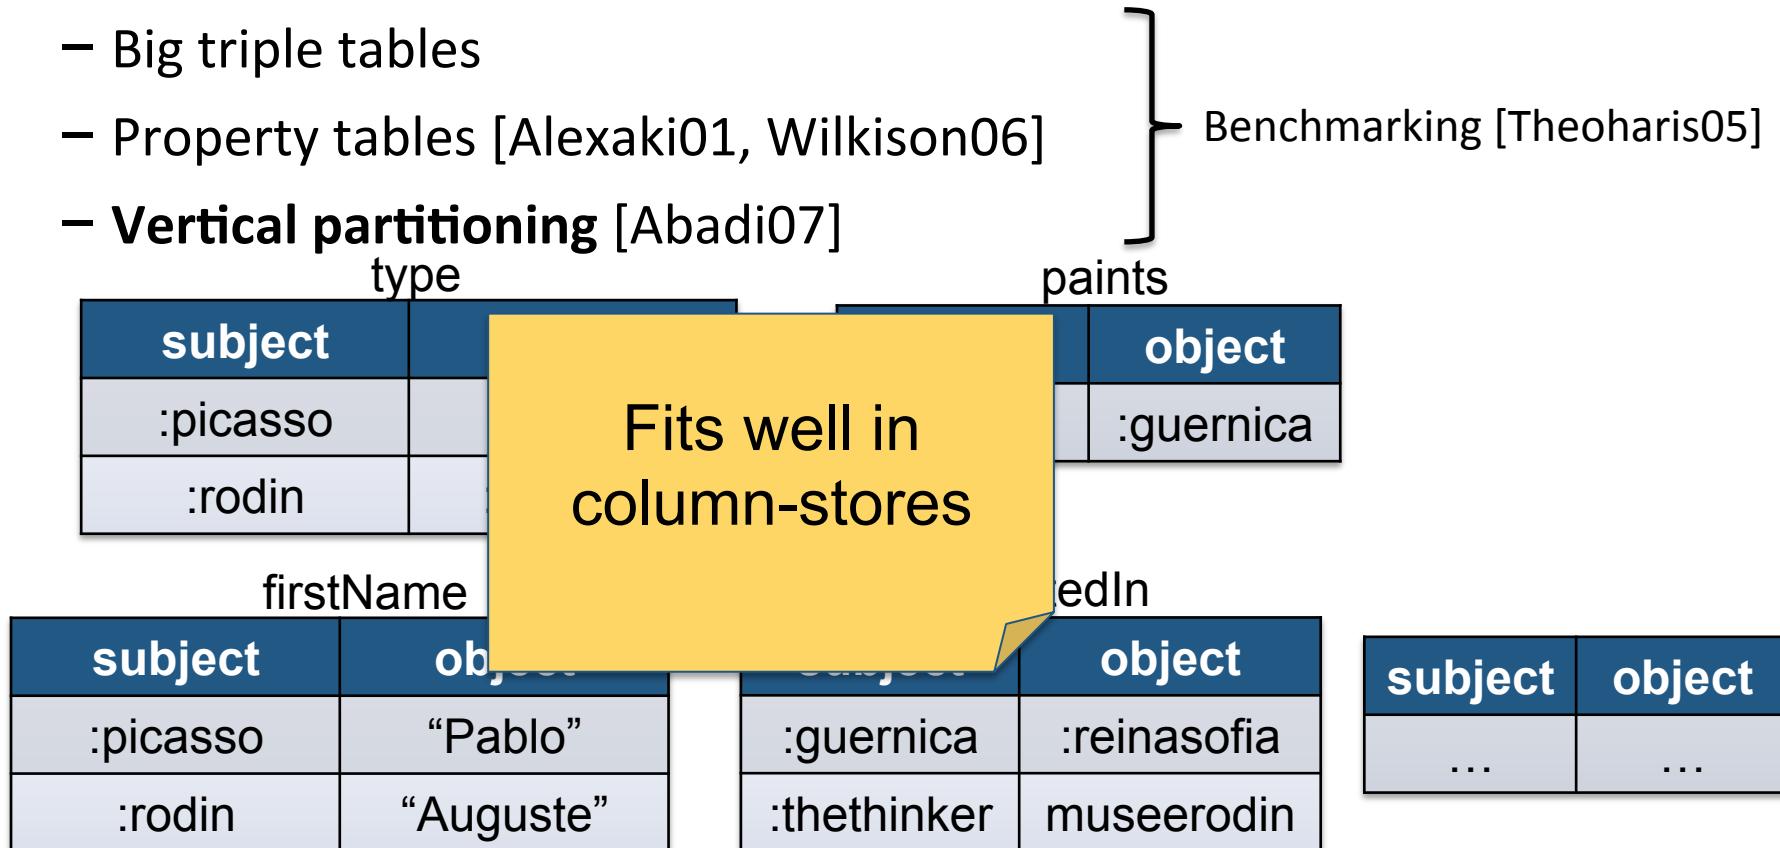
Artifacts

subject	exhibitedIn
:guernica	:reinasofia
:thethinker	:museerodin

subject	...
...	...
...	...

# Centralized RDF stores

- How do we usually store RDF in centralized environment?
- Approaches relying on RDBMs
  - Big triple tables
  - Property tables [Alexaki01, Wilkison06]
  - **Vertical partitioning** [Abadi07]



# Centralized RDF stores

- How do we usually store RDF in centralized environment?
- Approaches relying on RDBMs
  - Big triple tables
  - Property tables [Alexaki01, Wilkison06]
  - Vertical partitioning [Abadi07]
- Native RDF stores
  - RDF-3X [Neumann10], Hexastore [Weiss08]
  - 6 indexes (1 for each permutation of s p o)
  - compressed indices
  - aggregated indices

# Dictionary encoding

- RDF contains long strings of URIs

Size of triples

	min	max	avg
	80 bytes	2100 bytes	240 bytes

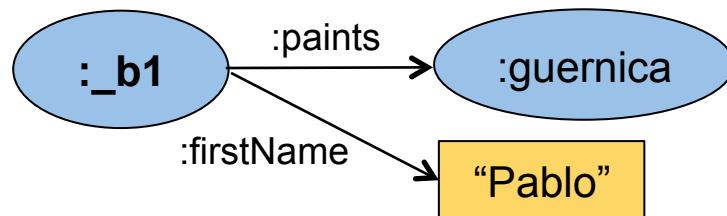


Source <http://gromgull.net/blog/category/semantic-web/billion-triple-challenge/>

Dictionary		Triples		Triples	
<b>id</b>	<b>resource</b>	<b>pro</b>	<b>subject</b>	<b>property</b>	<b>object</b>
1	:picasso	t	1	3	4
2	:rodin	:fir	1	5	10
3	type	:p	1	7	6
4	:cubist	:exh	6	12	9
5	:firstName	:loc	...	...	...
6	:guernica	type		:sculptor	
...	...	:creates		:thethinker	
...	...	...		...	

# RDF is not just a 3-attribute table

- Blank nodes
  - unknown/anonymous resources
  - existential variables



- for more on blank nodes: [Mallea11]
- Schema and data can be queried together
- Properties can also be described
  - (`:paints, domain, :painter`)
  - joins between properties and subjects/objects
- Inference
- and more ... [Hayes04]

# Inference in RDF

- RDF schema statements and set of entailment rules lead to implicit (entailed) RDF data

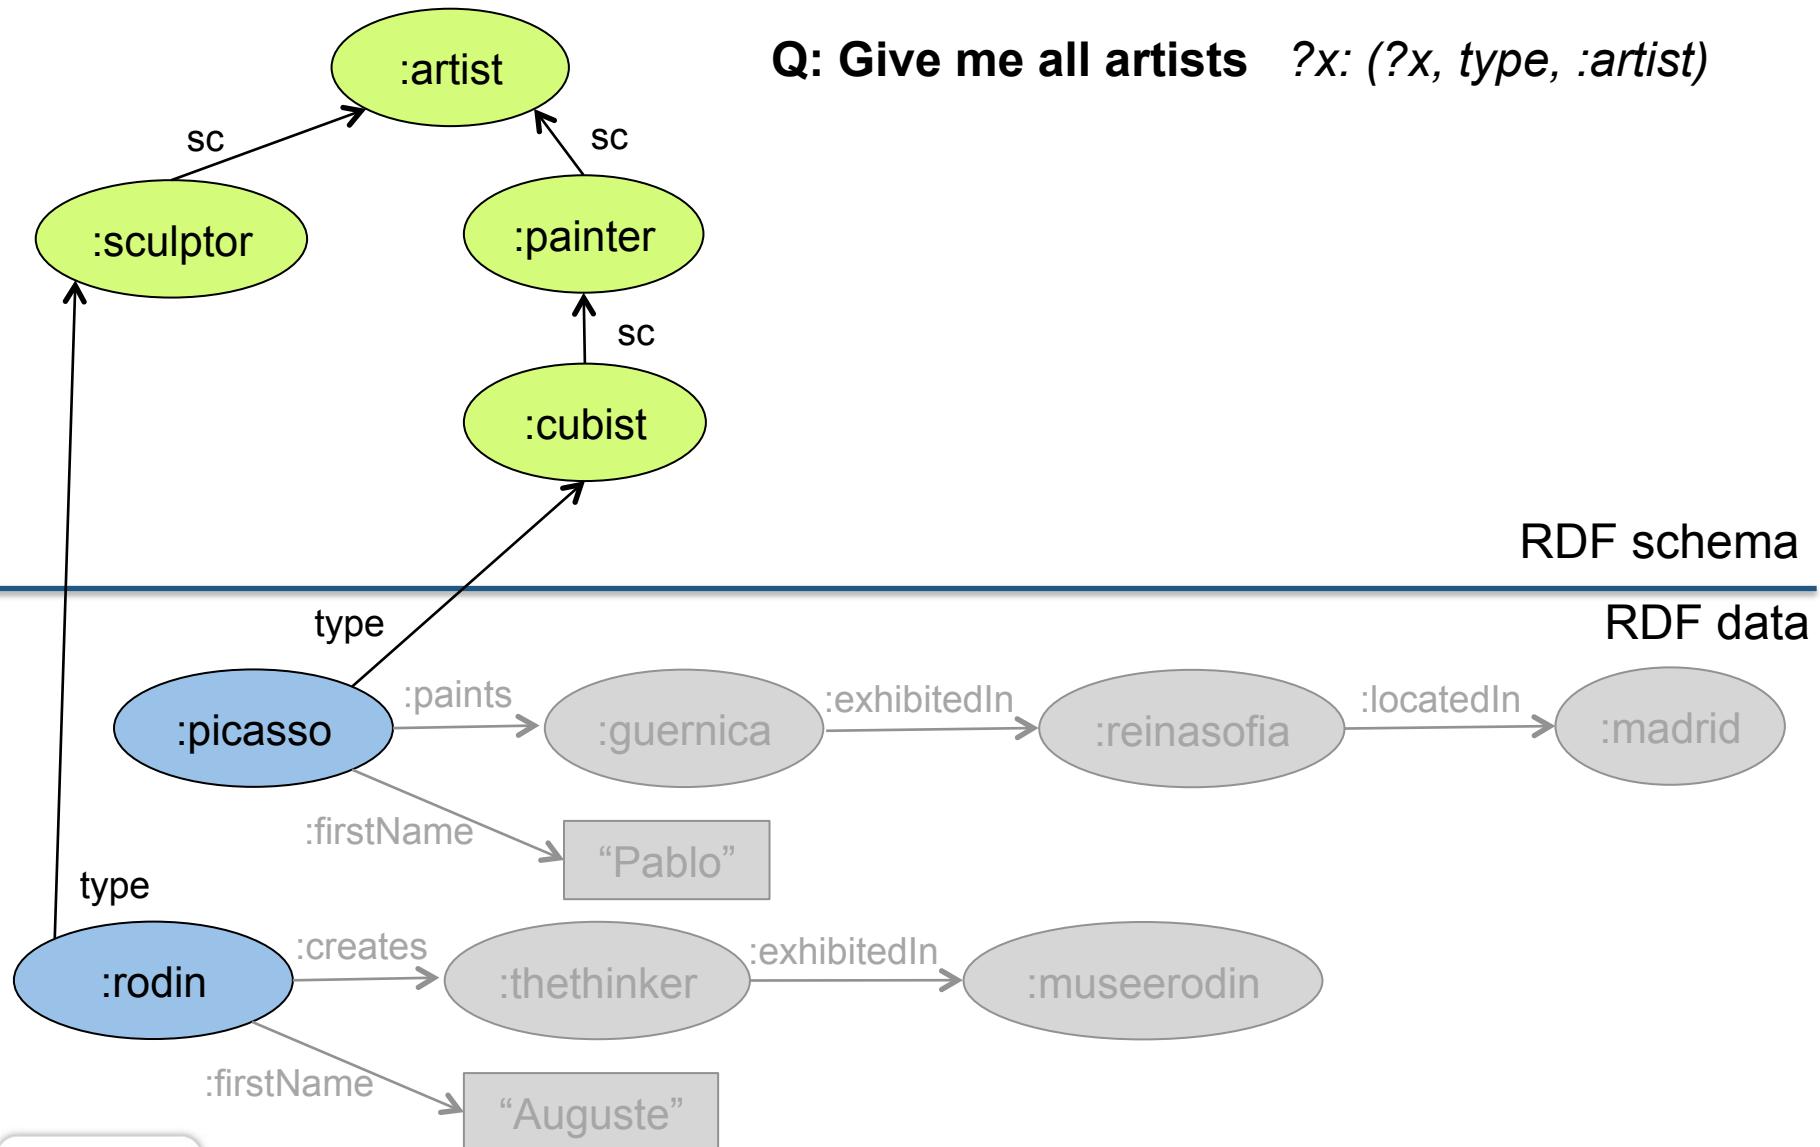
Sample rules

$$\begin{aligned}(X, \text{type}, A) \wedge (A, \text{sc}, B) &\rightarrow (X, \text{type}, B) \\ (X, P, O) \wedge (P, \text{domain}, A) &\rightarrow (X, \text{type}, A)\end{aligned}$$

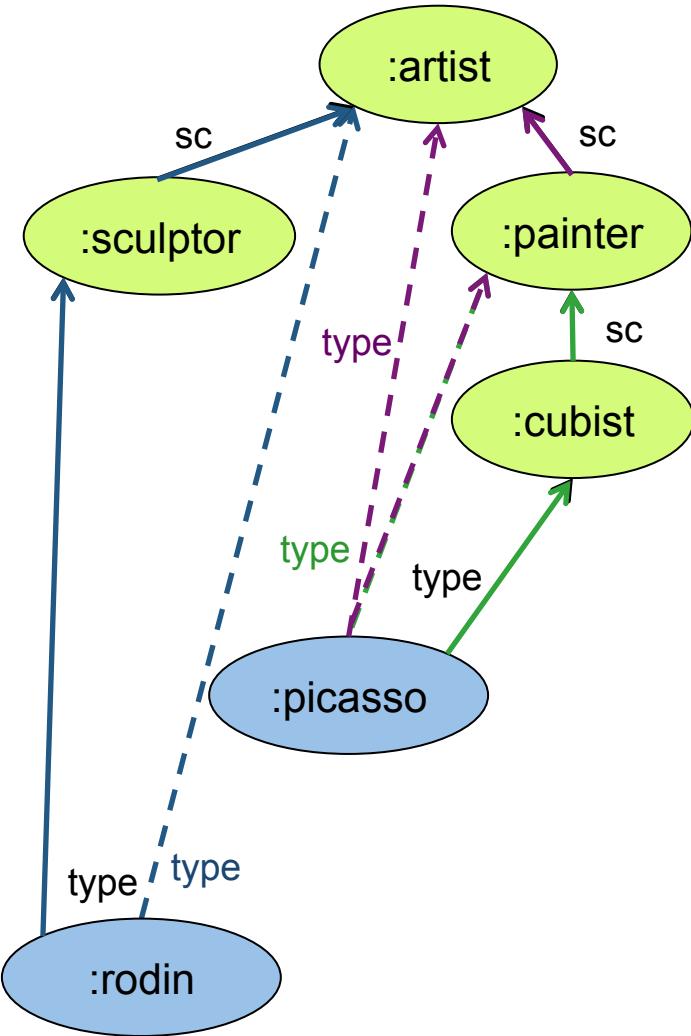
$(\text{:picasso}, \text{paints}, \text{:guernica}) \wedge (\text{:paints}, \text{domain}, \text{:painter}) \rightarrow (\text{:picasso}, \text{type}, \text{:painter})$

- RDFS entailment : Given an RDF(S) database DB and a triple  $t$ , is  $t$  entailed from DB (or does  $t$  logically follows from DB)?

# Entailed triples and RDF querying



# Entailed triples and RDF querying



**Q: Give me all artists**  $?x: (?x, \text{type}, :artist)$

**A:**

$?x$
-

no entailment

Rule:

$(X, \text{type}, A) \wedge (A, \text{sc}, B) \rightarrow (X, \text{type}, B)$

**Implicit / entailed triples**

$(:rodin, \text{type}, :artist)$

$(:picasso, \text{type}, :painter)$

$(:picasso, \text{type}, :artist)$

**A:**

$?x$
$:picasso$
$:rodin$

with entailment

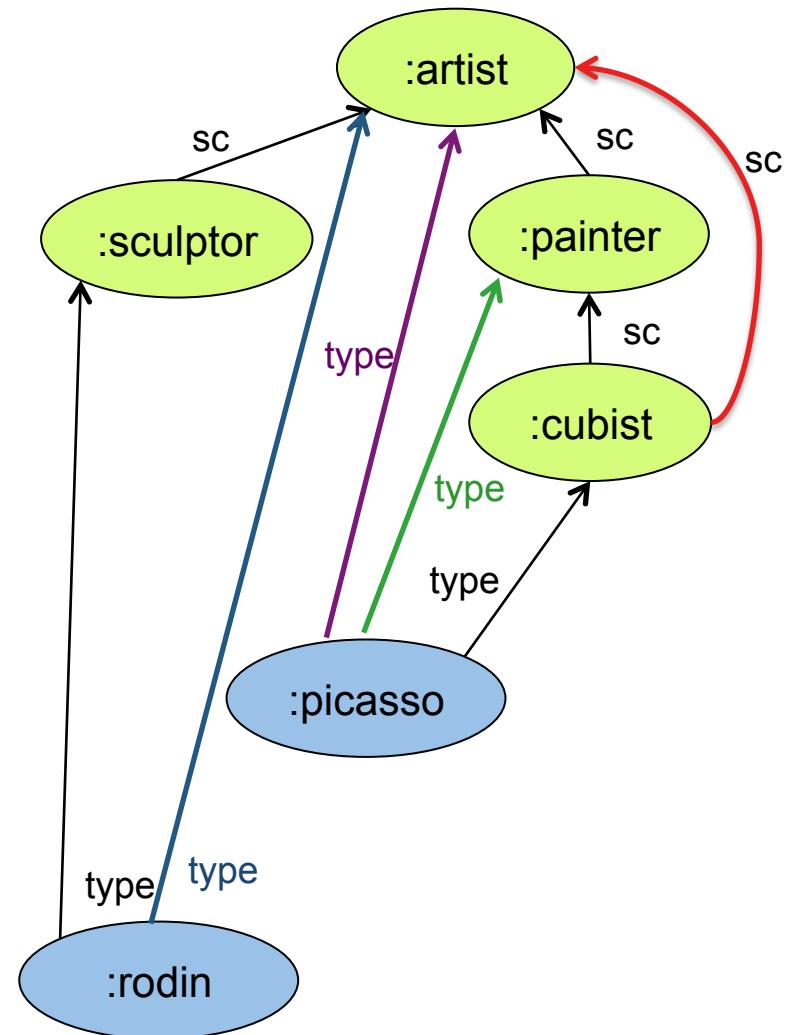
# RDFS entailment techniques

- RDFS closure
  - compute **all implicit triples** and store them in the database
- Query reformulation
  - transform the query so that you get all the answers (even the ones that come from implicit triples)
- Hybrid approaches
  - compute the **schema closure**, and then **reformulate the query**
  - **magic sets**: based on the queries, precompute and store only the implicit triples that are required for the answer
- For a comparison on these techniques:
  - centralized [Goasdoué13]
  - distributed [Kaoudi13]

# RDFS entailment – RDFS closure

- Before querying

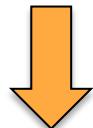
subject	property	object
:picasso	type	:cubist
:rodin	type	:sculptor
:sculptor	sc	:artist
:cubist	sc	:painter
:painter	sc	:artist
<b>:cubist</b>	<b>sc</b>	<b>:painter</b>
:rodin	type	:artist
<b>:picasso</b>	<b>type</b>	<b>:painter</b>
<b>:picasso</b>	<b>type</b>	<b>:artist</b>



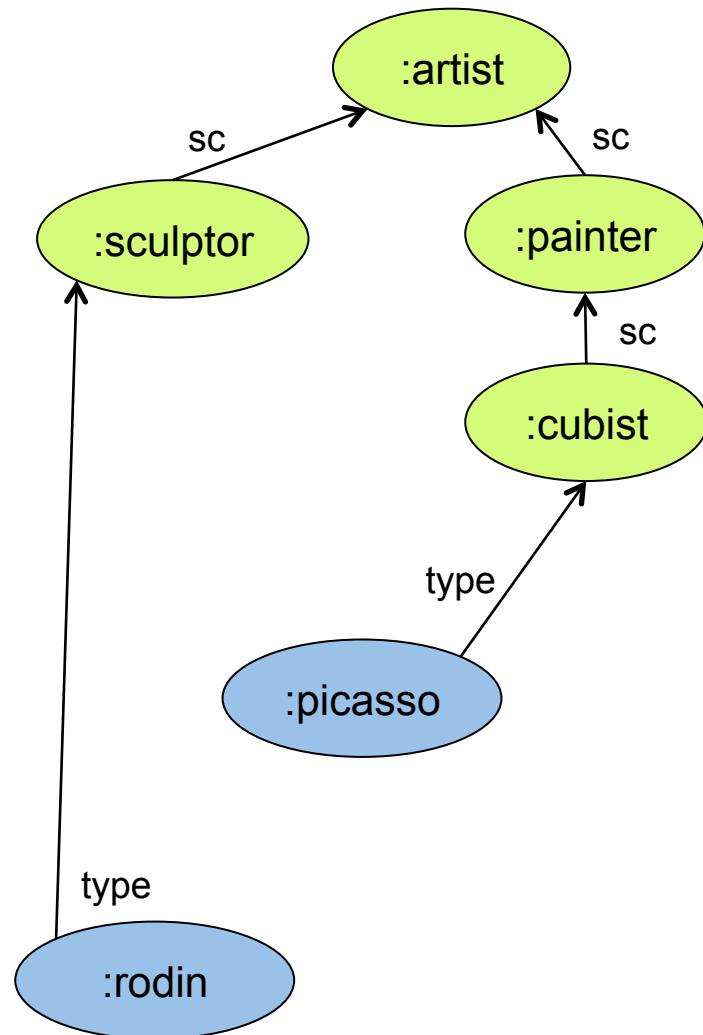
# RDFS entailment – query reformulation

- At query time

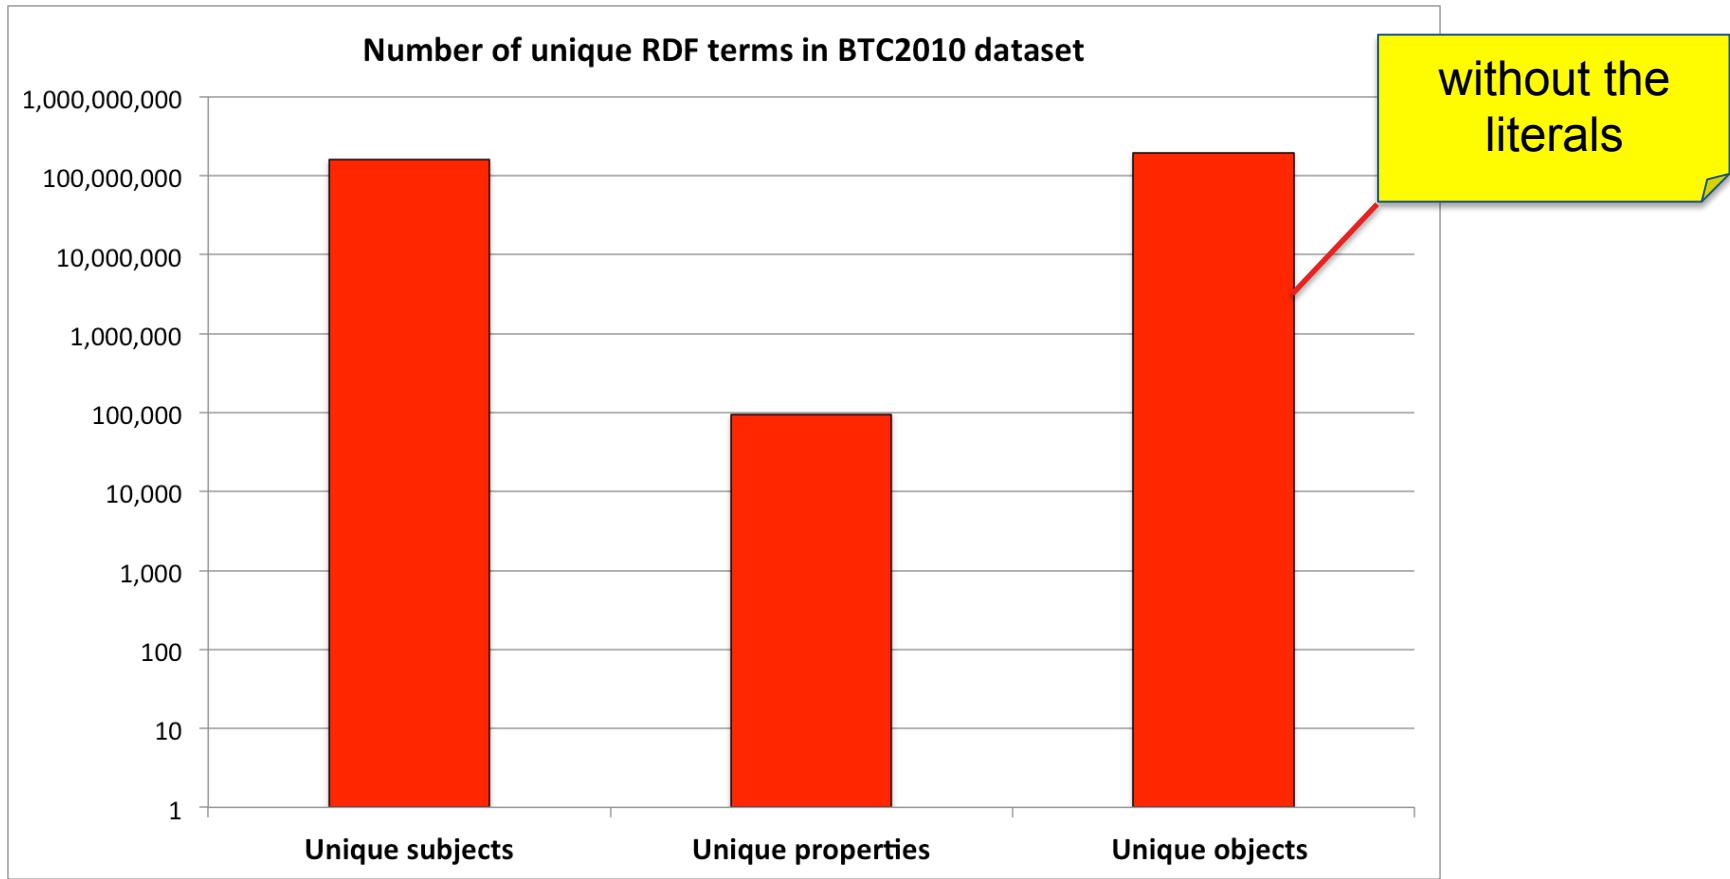
?x: (?x, type, :artist)



?x: (?x, type, :artist)  $\vee$   
(?x, type, :sculptor)  $\vee$   
(?y, type, :painter)  $\vee$   
(?x, type, :cubist)



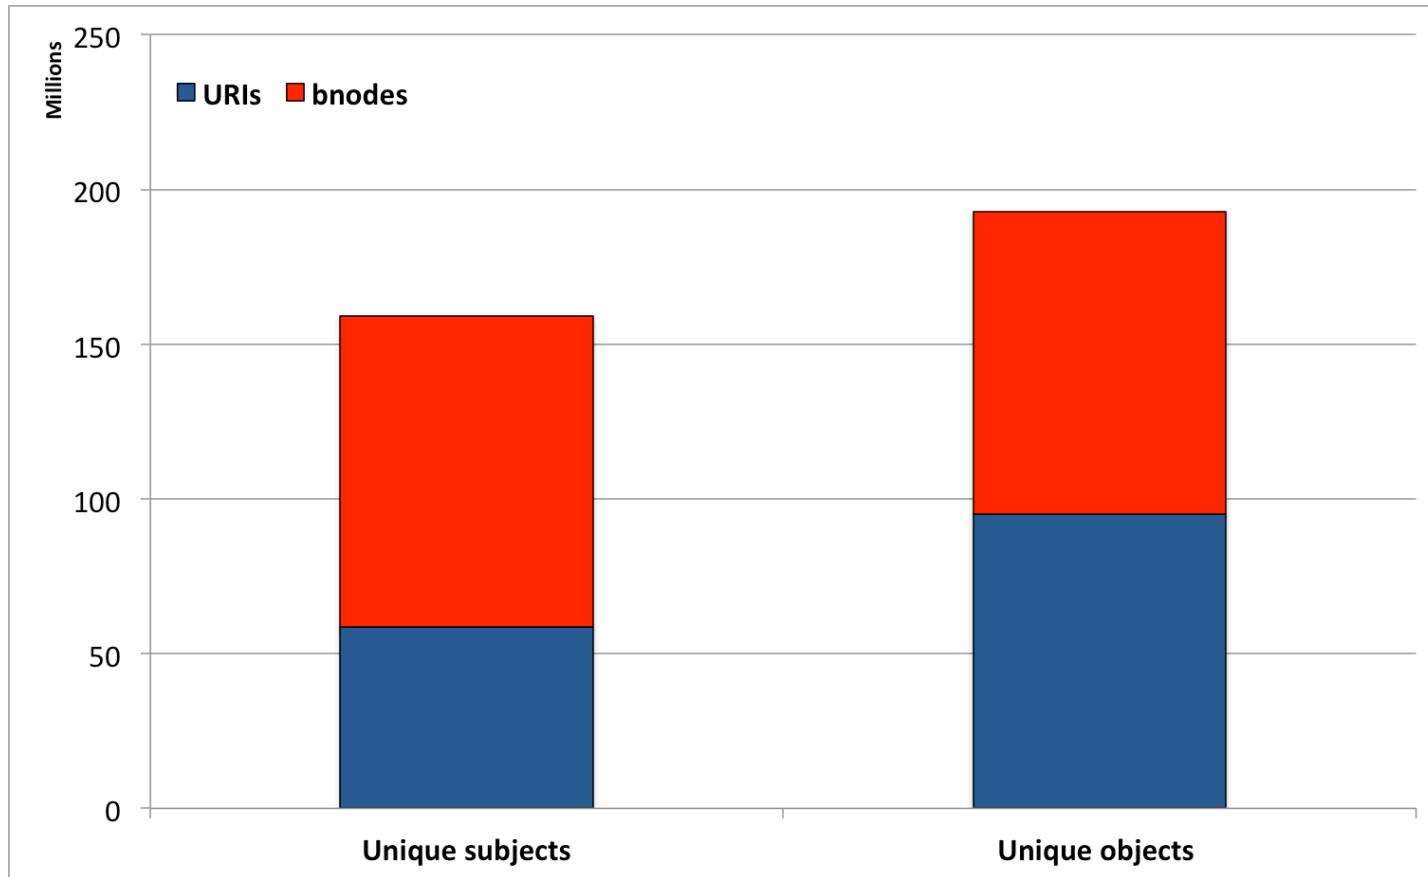
# How does real RDF look?



Data taken from <http://gromgull.net/blog/2010/09/btc2010-basic-stats/>

# How does real RDF look?

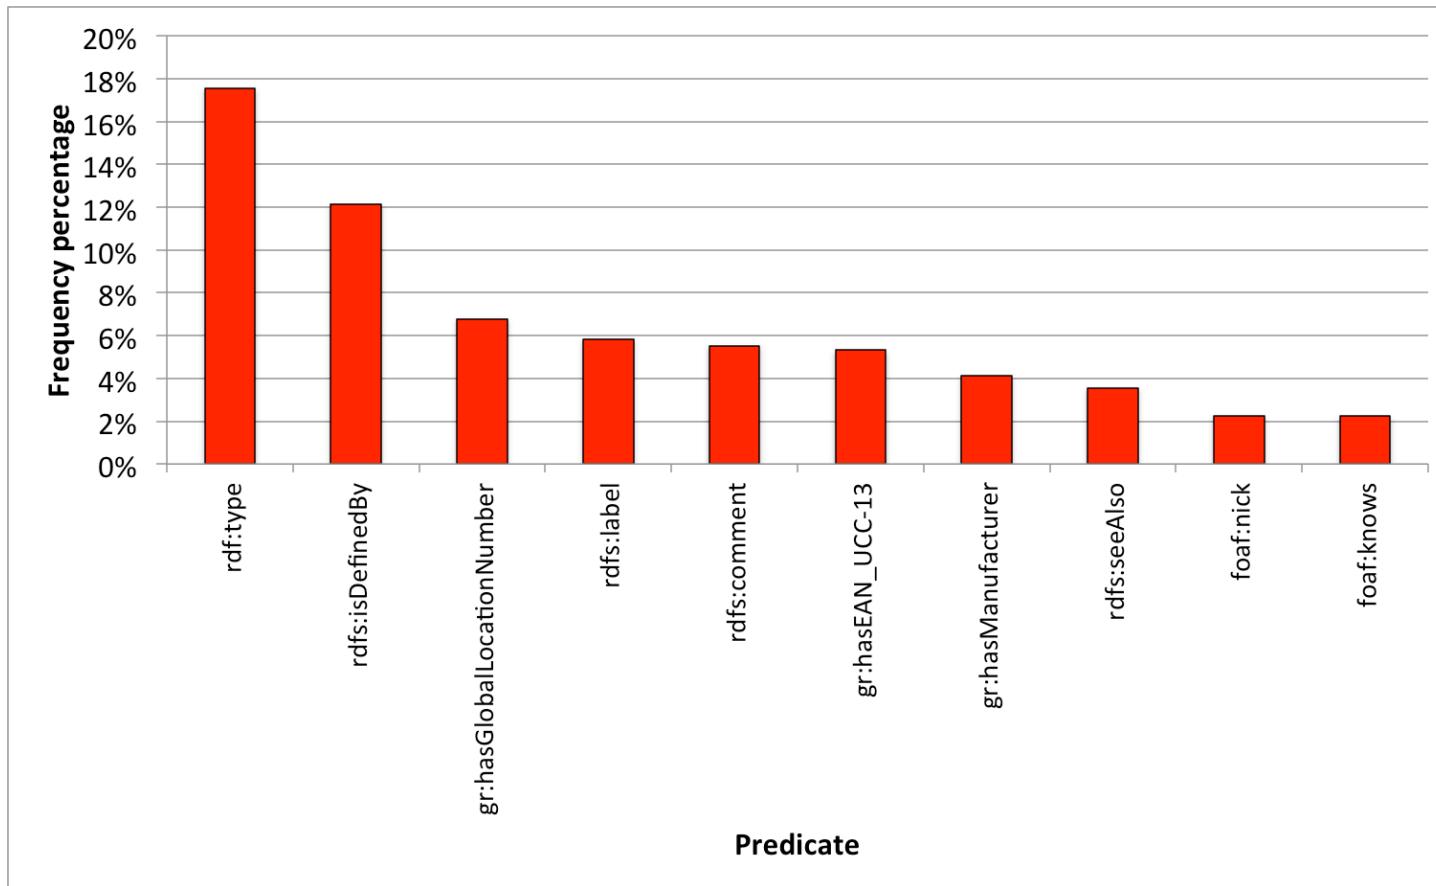
BTC2010 dataset



Data taken from <http://gromgull.net/blog/2010/09/btc2010-basic-stats/>

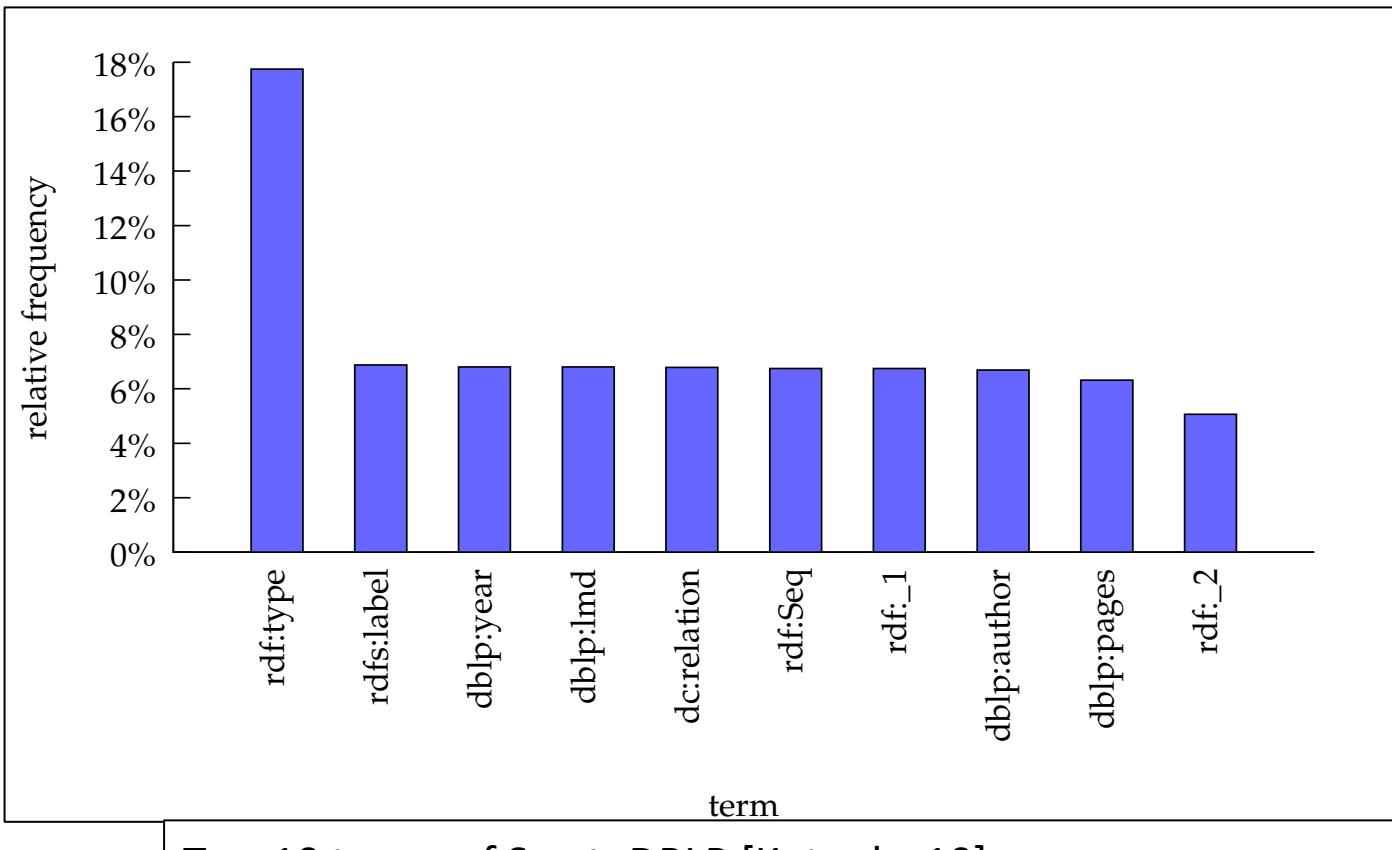
# How does real RDF look?

BTC2010 dataset



Data taken from <http://gromgull.net/blog/2010/09/btc2010-basic-stats/>

# How does real RDF look?



# 2

## CLOUD BASICS

# Cloud computing

*The origin of the term **cloud computing** was derived from [...] drawings of stylized clouds to denote networks [...]*

*The cloud symbol was used to represent the **Internet** as early as **1994**.*

[Wikipedia: [http://en.wikipedia.org/wiki/Cloud\\_computing](http://en.wikipedia.org/wiki/Cloud_computing)]

- Use of computing resources as a service over the network
- Scalability
- Elasticity
- Reliability

# Cloud computing service models

Google Apps

Software-as-a-Service (SaaS)



Applications

Windows Azure



Platform-as-a-Service (PaaS)



Platforms

amazon  
webservices™

Elastic Compute Cloud (EC2)

Infrastructure-as-a-Service (IaaS)



Hardware

## pay-per-use

# Cloud computing and data management

- Relational databases in the cloud (DBaaS)
  - IBM DB2 on SmartCloud
  - Amazon RDS (MySQL, Oracle or SQL Server)
- NoSQL systems [Catell11]:
  - Key-value stores (memcached, DynamoDB, SimpleDB)
  - Extensible record stores (BigTable, Cassandra, HBase, Accumulo)
  - Document stores (MongoDB, CouchDB)
  - Graph databases (Pregel, Apache Giraph, Neo4J)
- Parallel data processing paradigms:
  - MapReduce (Hadoop, Amazon EMR), PACTs/Stratosphere

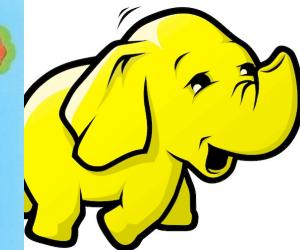
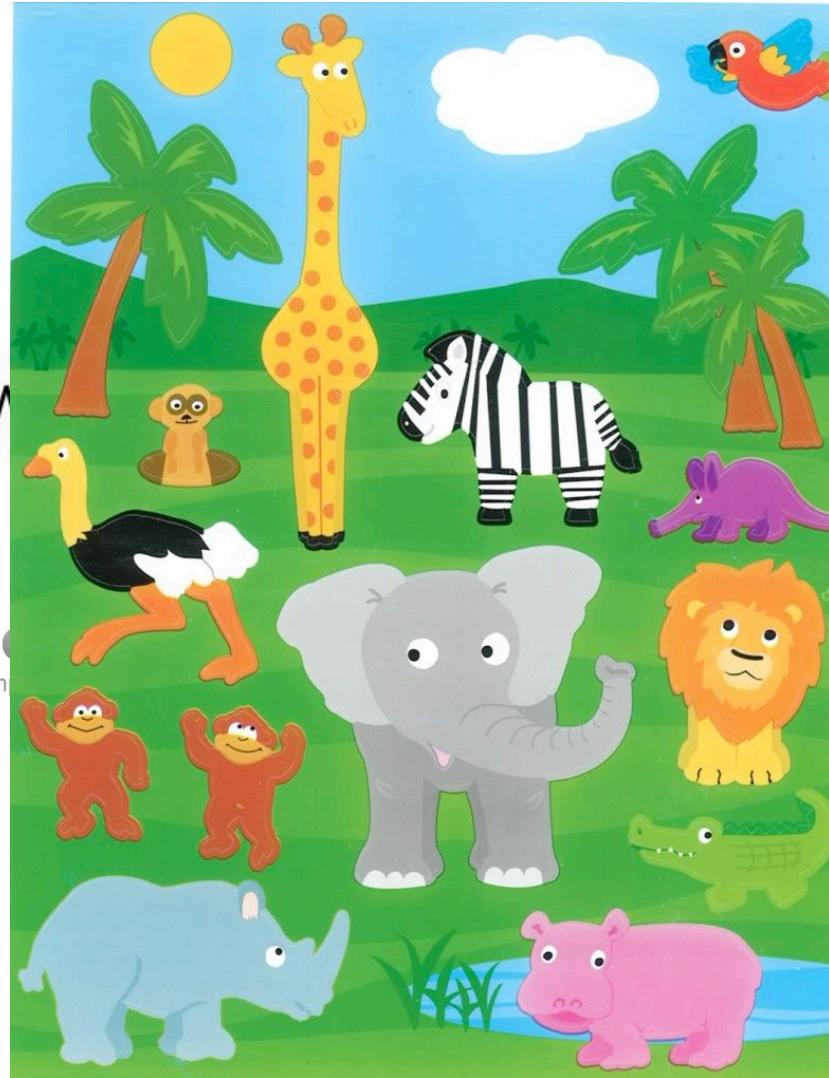
# Cloud ecosystem



Google  
App Engine



cloudstack  
open source cloud comp



CHE  
ASE

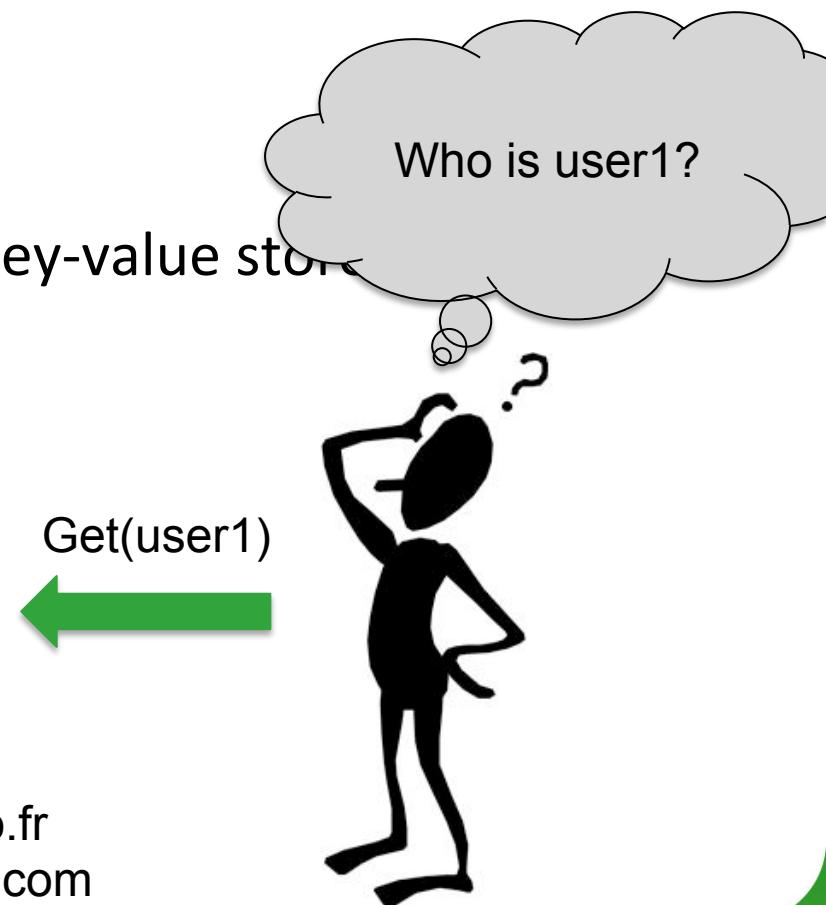


APACHE  
GIRAPH

# Distributed key-value stores

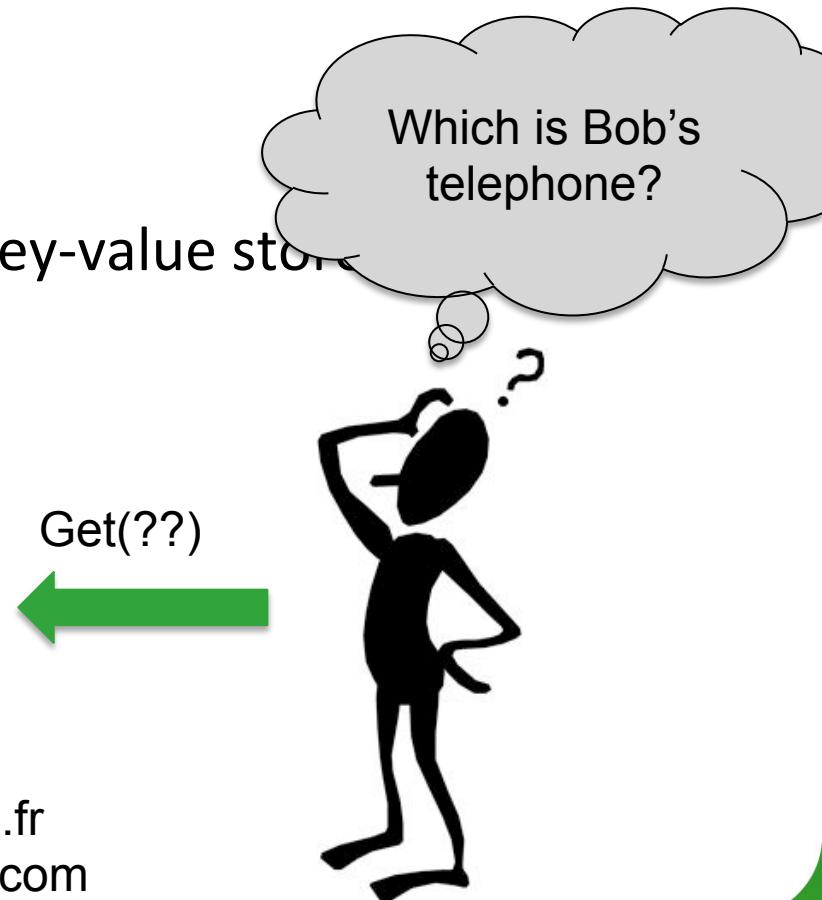
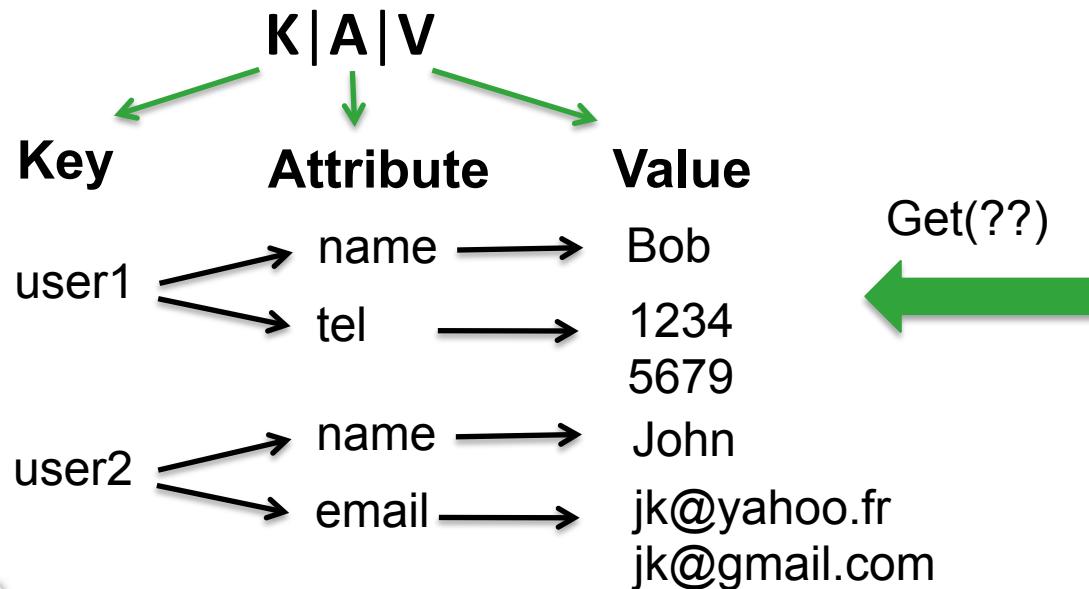
- **Index:** Hash-based or lexicographically sorted on the **key**
- Basic APIs: **Put(k, v)**, **Get(k)**, **Delete(k)**
- No schema, no joins
- Multi-value attributes
- Notation we will use for a table in a key-value store

K   A   V		
Key	Attribute	Value
user1	name	Bob
	tel	1234 5679
user2	name	John
	email	jk@yahoo.fr jk@gmail.com



# Distributed key-value stores

- **Index:** Hash-based or lexicographically sorted on the **key**
- Basic APIs: **Put(k, v)**, **Get(k)**, **Delete(k)**
- No schema, no joins
- Multi-value attributes
- Notation we will use for a table in a key-value store



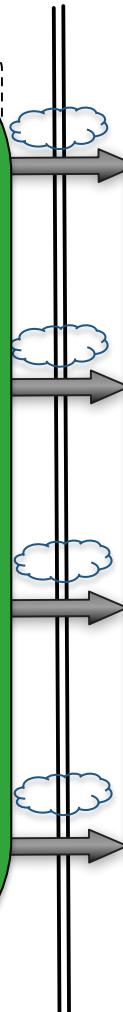
# More “advanced” distributed key-value stores

- Google BigTable, Apache Cassandra
- Hash or sorted index on the key
- Sorted index on the attribute name
- Supercolumns
  - $K | \{SC\}A | V$
  - Hash or sorted index
- Secondary index
  - maps values to keys

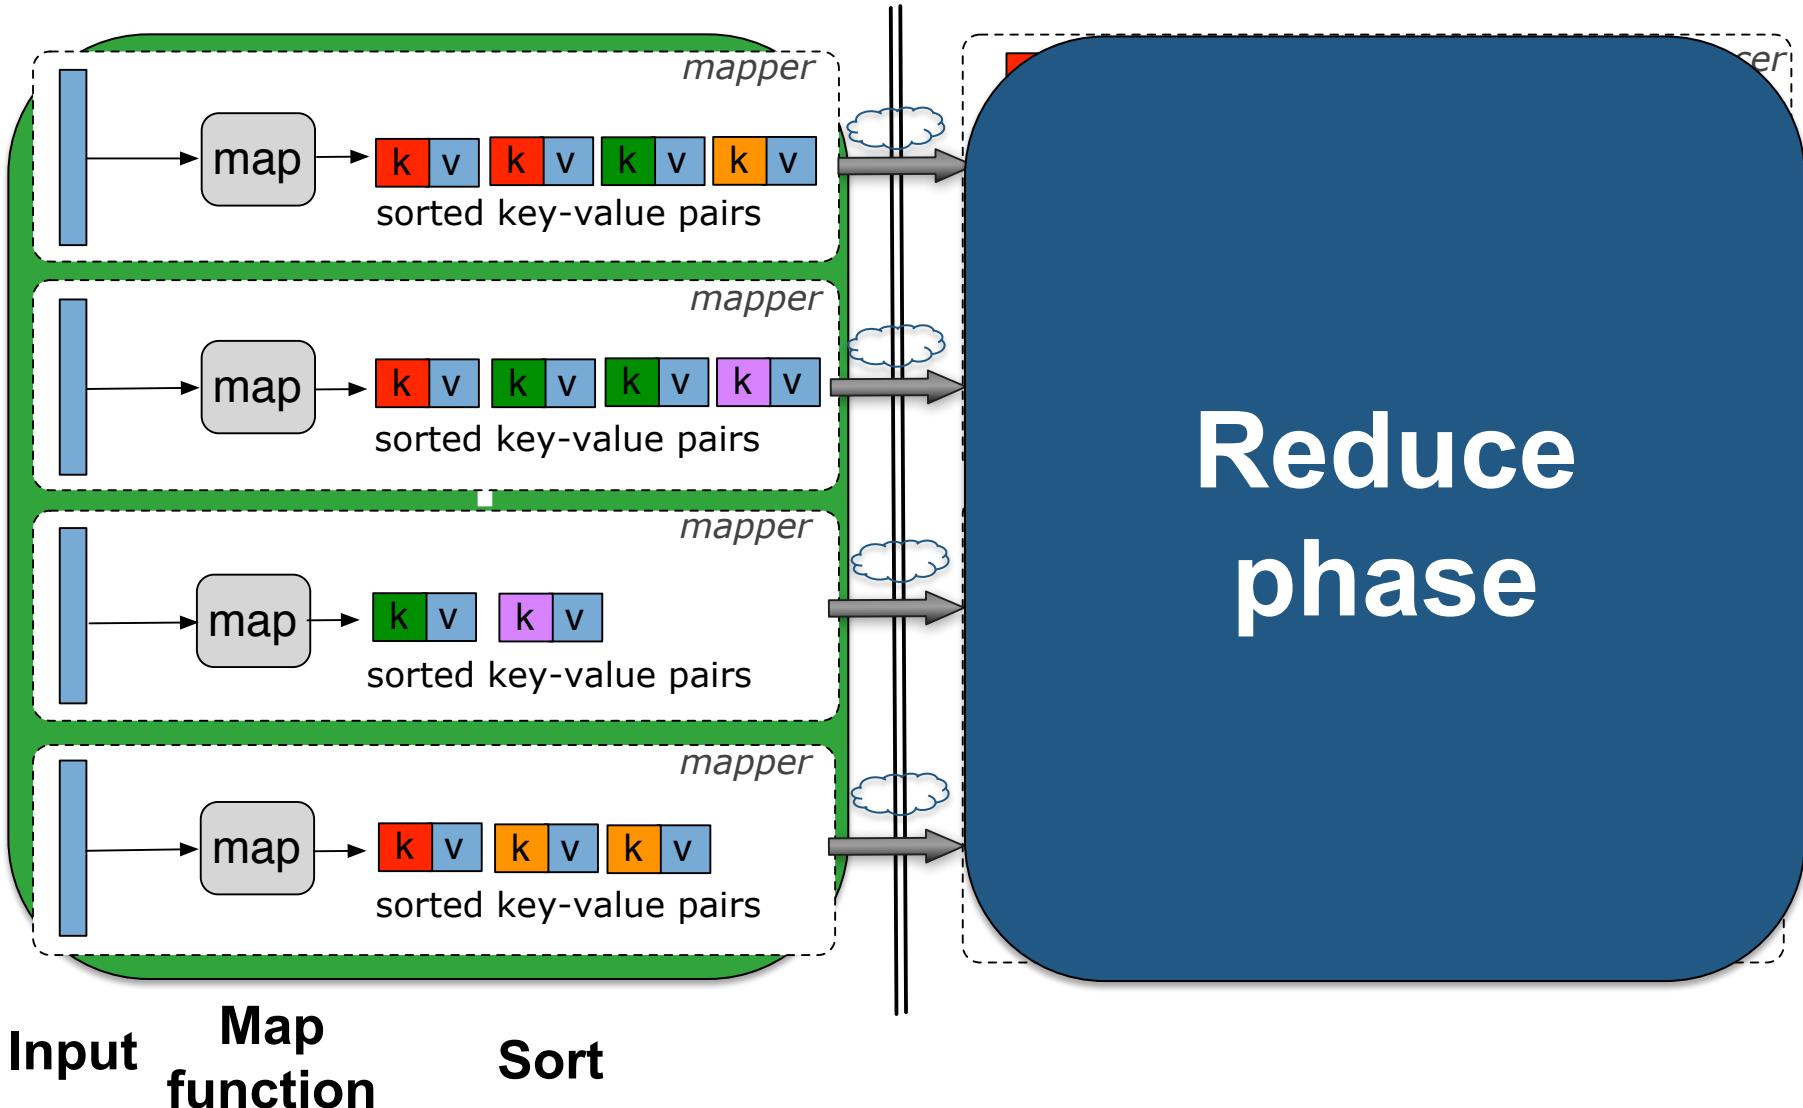
# MapReduce

Map  
phase

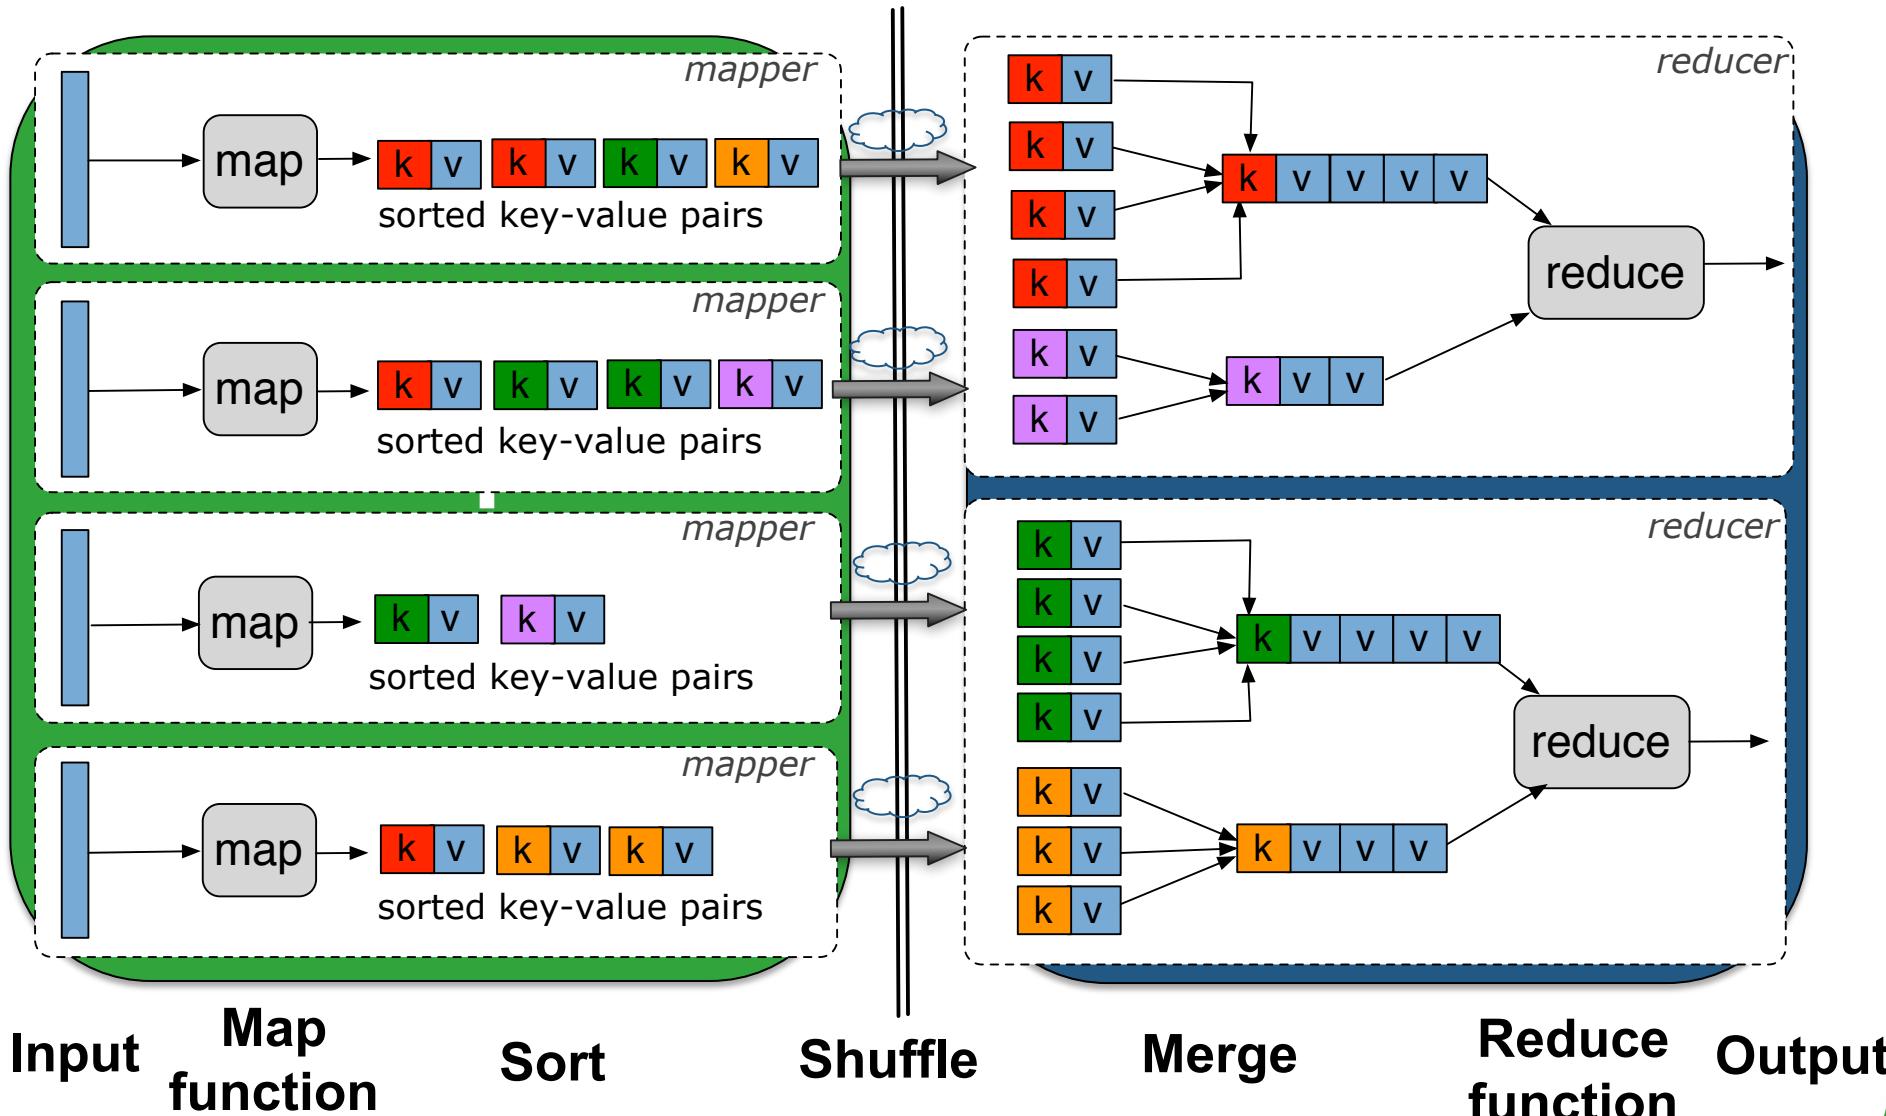
Reduce  
phase



# MapReduce



# MapReduce



# MapReduce

- High parallelization
- Scalable
- Fault-tolerant
- Operates on 1 input
- No indexes
- Blocking operations

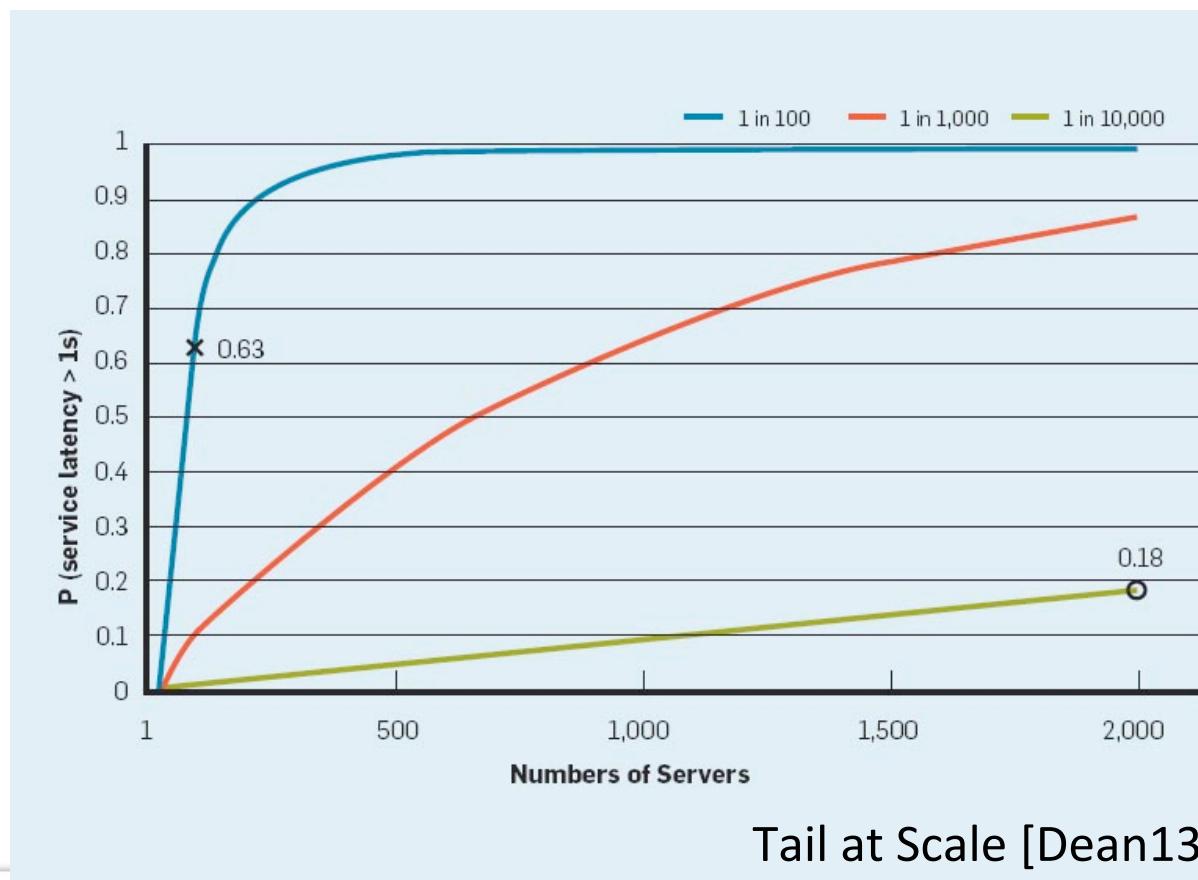
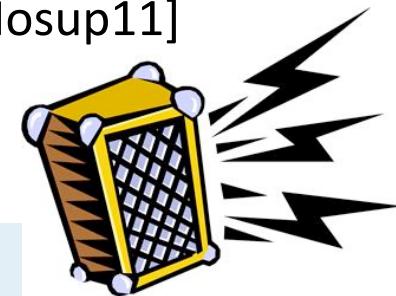
# Joins in MapReduce

- Two main strategies:
  - symmetric hash join
  - replicated join
- Representative works:
  - [Blanas10] on different join methods in MapReduce
  - [Afrati11] on optimizing multiway joins in MapReduce
  - [Wu11] on query optimization in MapReduce

# Variability in the clouds

Cloud = performance variability [Schad10], [Iosup11]

**Variability is amplified by scale**



# 3

## ANALYSIS OF RDF SYSTEMS IN THE CLOUD

# RDF stores in the cloud

- **AMADA** [Bugiotti11, Arandar12]
- CumulusRDF [Ladwig11]
- **Graph partitioning** [Huang11]
- H2RDF [Papailiou12]
- **HadoopRDF** [Husain11]
- MAPSIN [Schätzle12]
- PigSPARQL [Schätzle11]
- RAPID+ [Ravindra11]
- **Rya** [Punnoose12]
- Stratustore [Ladwig11]
- **SHARD** [Rohloff10]
- Trinity.RDF [Zeng13]
- **WebPie** [Urbani09]
- QueryPie [Urbani11]
- ...

# Analysis dimensions

**Data storage**

**Query processing**

**RDFS entailment**

# Analysis dimensions

Data storage

Query processing

RDFS entailment

# Categorization based on storage and query processing

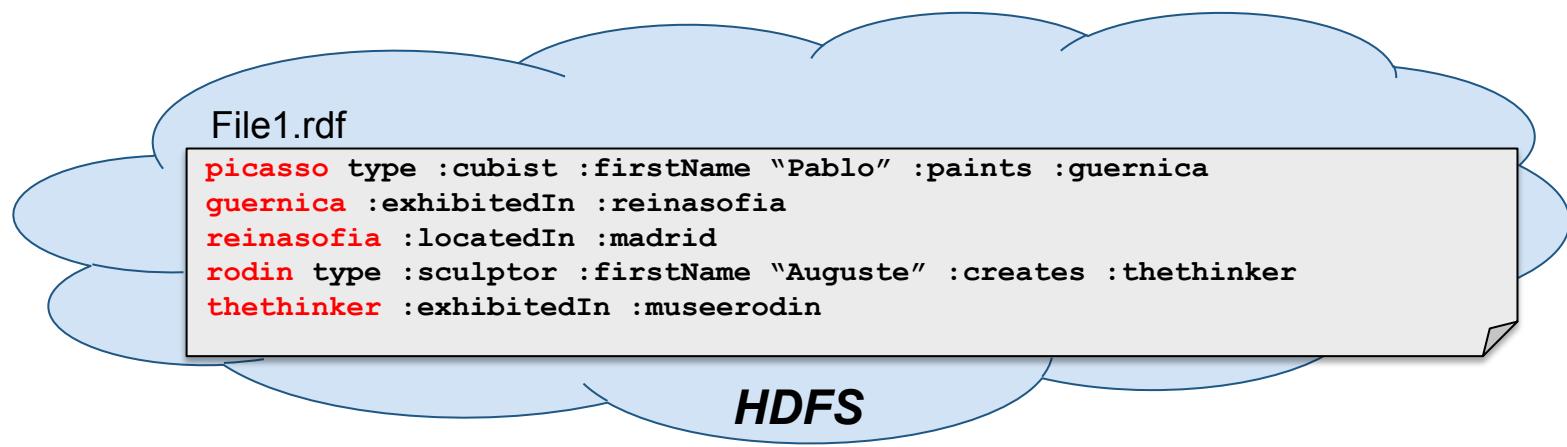
1. MapReduce-based approaches
2. Approaches based on key-value stores
3. Graph-based approaches
4. Systems based on commercial cloud services

# 1. MapReduce-based approaches

- Store RDF triples in **DFS** (usually HDFS)
- Use **MapReduce** to answer the queries (usually Hadoop)
- Representative works:
  - SHARD [Rohloff10]
  - HadoopRDF [Husain11]
  - RAPID+ [Ravindra11]
  - PigSPARQL [Schätzle11]

# SHARD triple-store [Rohloff10]

Storing: each line holds all triples about a given subject

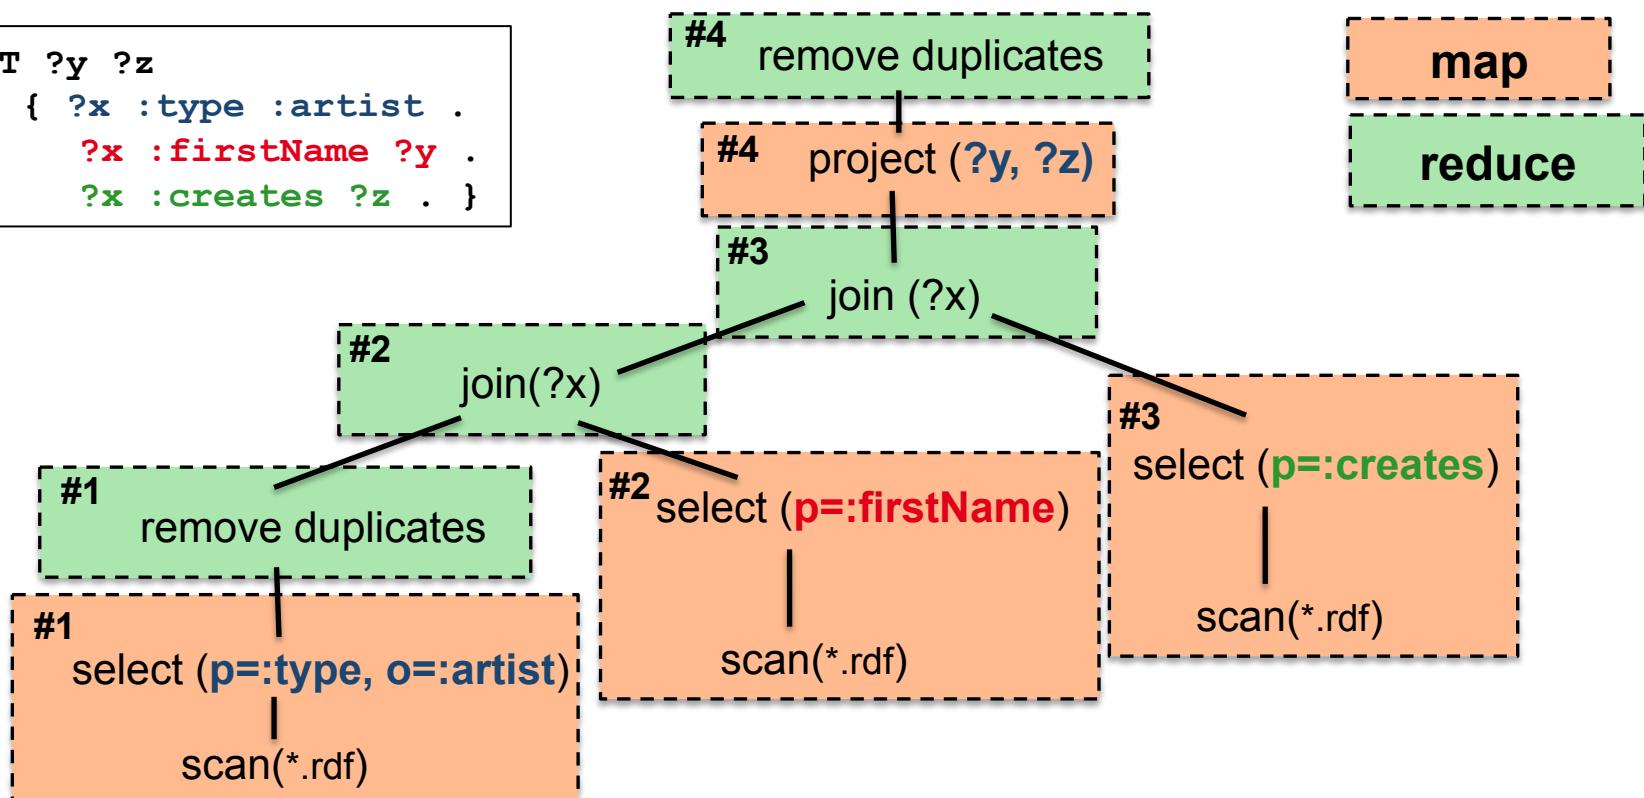


# SHARD triple-store [Rohloff10]

Query processing: Iterative MapReduce jobs

- 1 MapReduce job for each triple pattern + 1 for the projection at the end
- Map phase: **match** triple pattern by scanning all triples
- Reduce phase: **join** triple pattern with intermediate results from previous ones

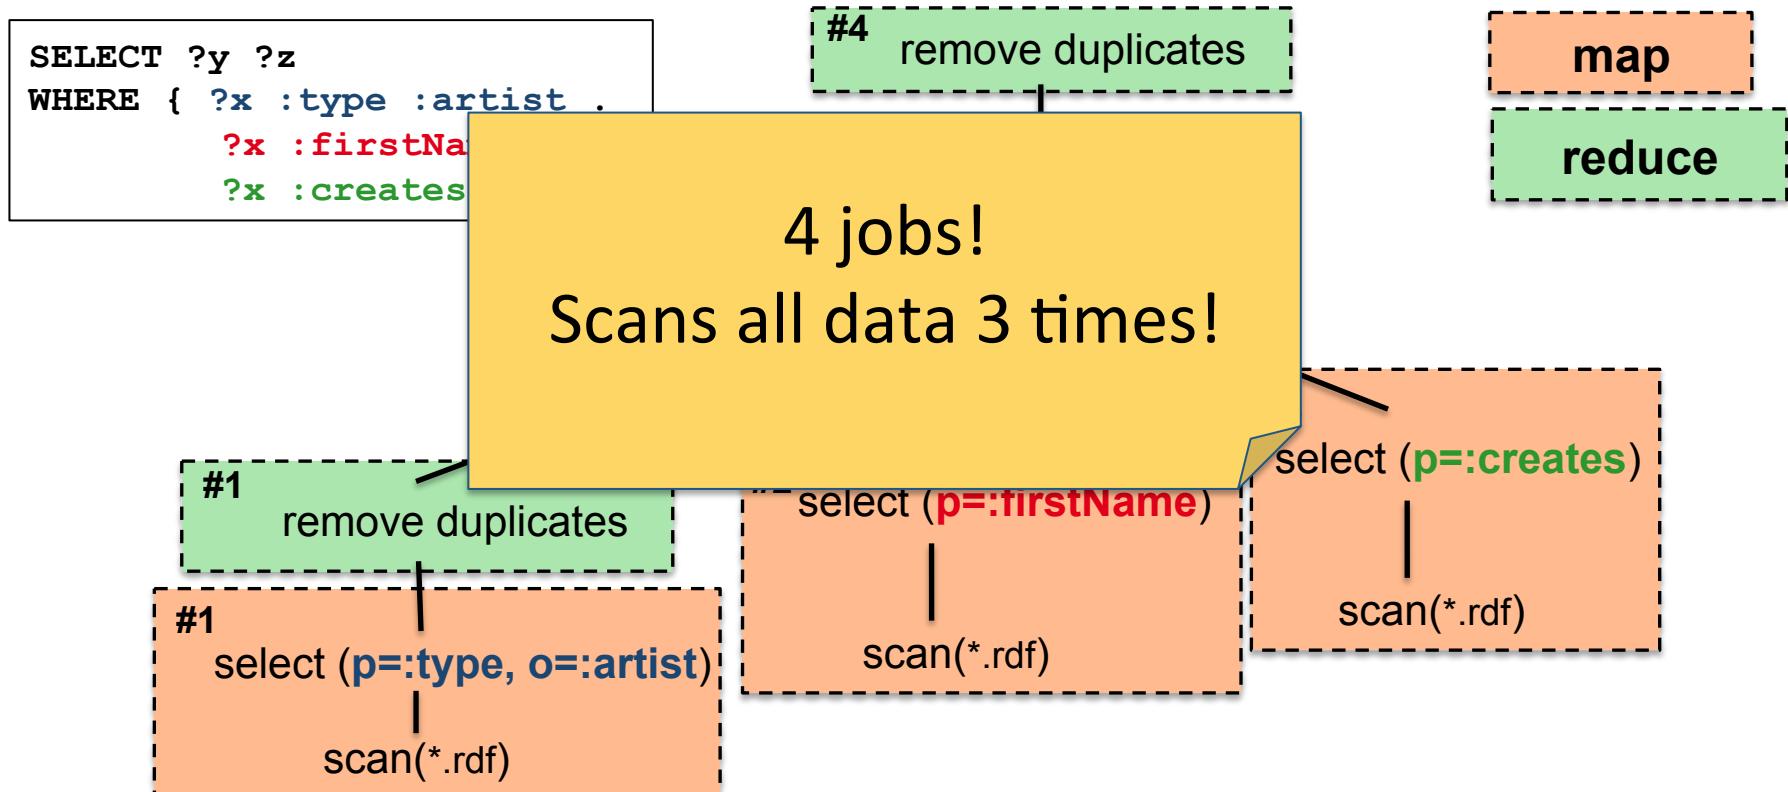
```
SELECT ?y ?z
WHERE { ?x :type :artist .
         ?x :firstName ?y .
         ?x :creates ?z . }
```



# SHARD triple-store [Rohloff10]

Query processing: Iterative MapReduce jobs

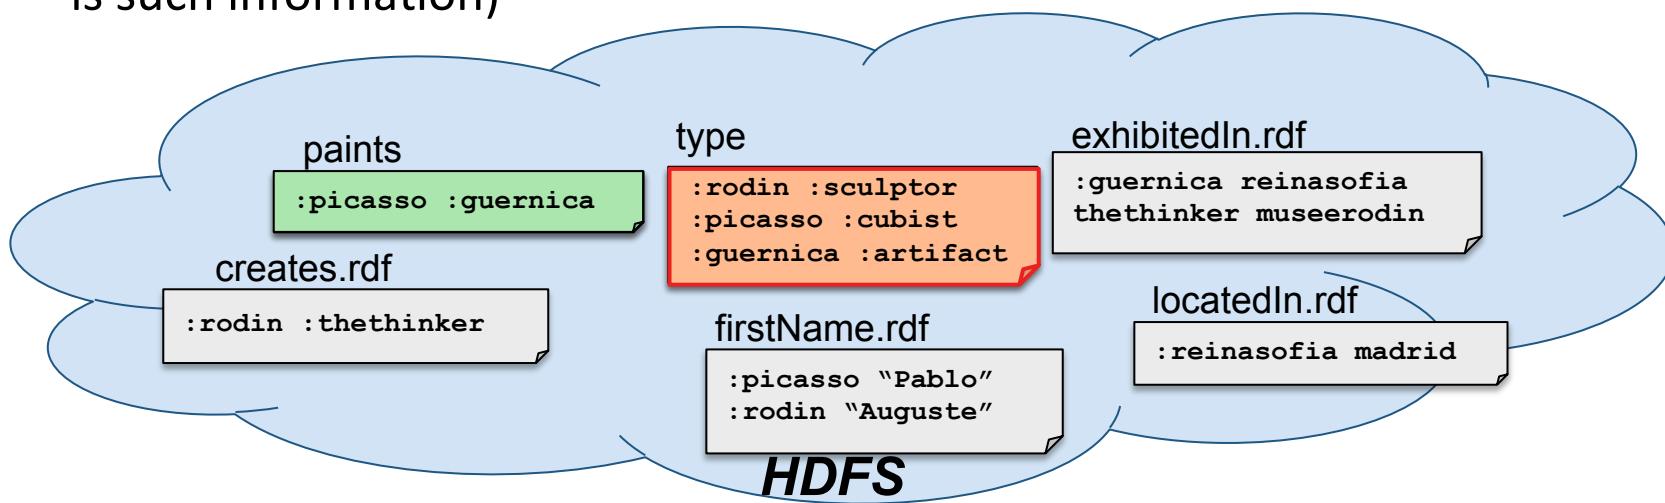
- 1 MapReduce job for each triple pattern + 1 for the projection at the end
- Map phase: **match** triple pattern by scanning all triples
- Reduce phase: **join** triple pattern with intermediate results from previous ones



# HadoopRDF [Husain11]

## Storing:

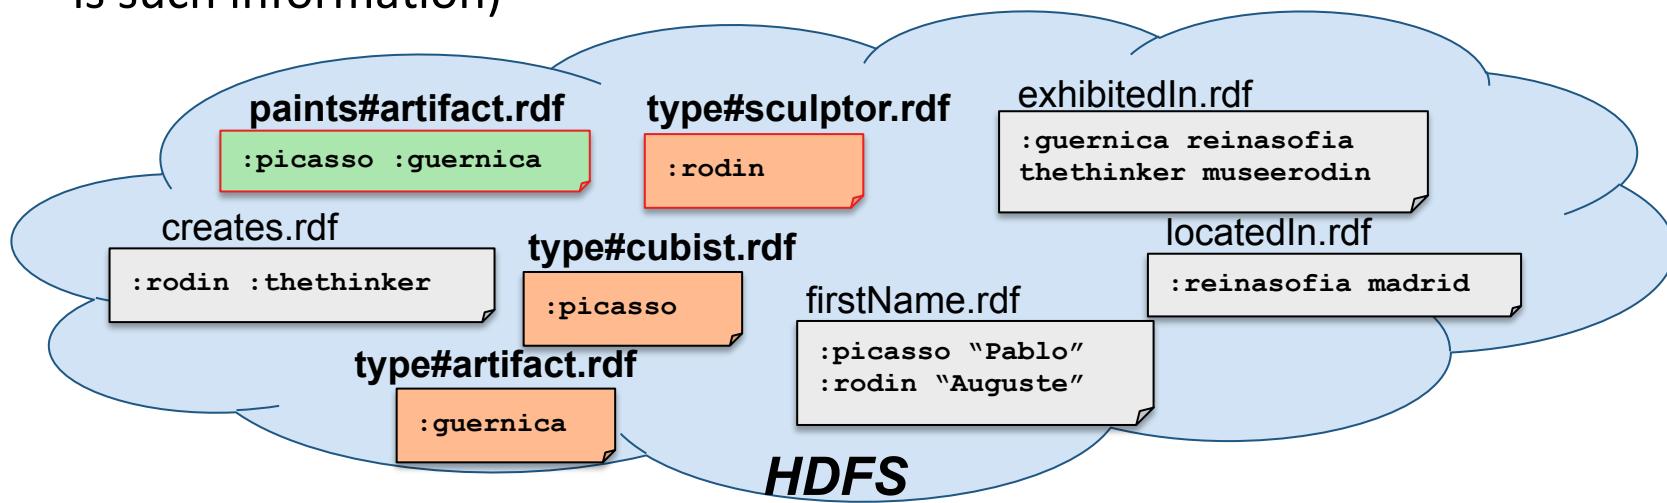
- RDF triples are grouped by the **property names** in files (like in vertical partitioning of centralized RDF stores)
- The file for the property **type** is split based on the object
- Triples with the same property are then grouped by their **object type** (if there is such information)



# HadoopRDF [Husain11]

## Storing:

- RDF triples are grouped by the **property names** in files (like in vertical partitioning of centralized RDF stores)
- The file for the property **type** is split based on the object
- Triples with the same property are then grouped by their **object type** (if there is such information)

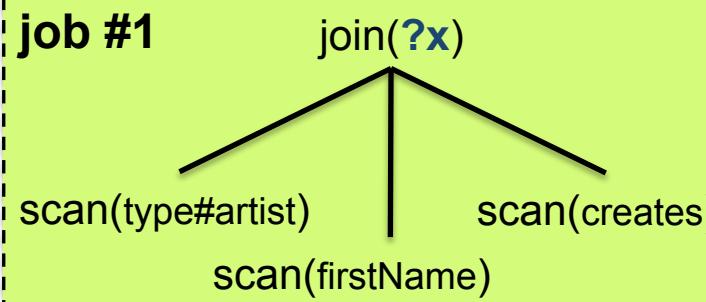


# HadoopRDF [Husain11]

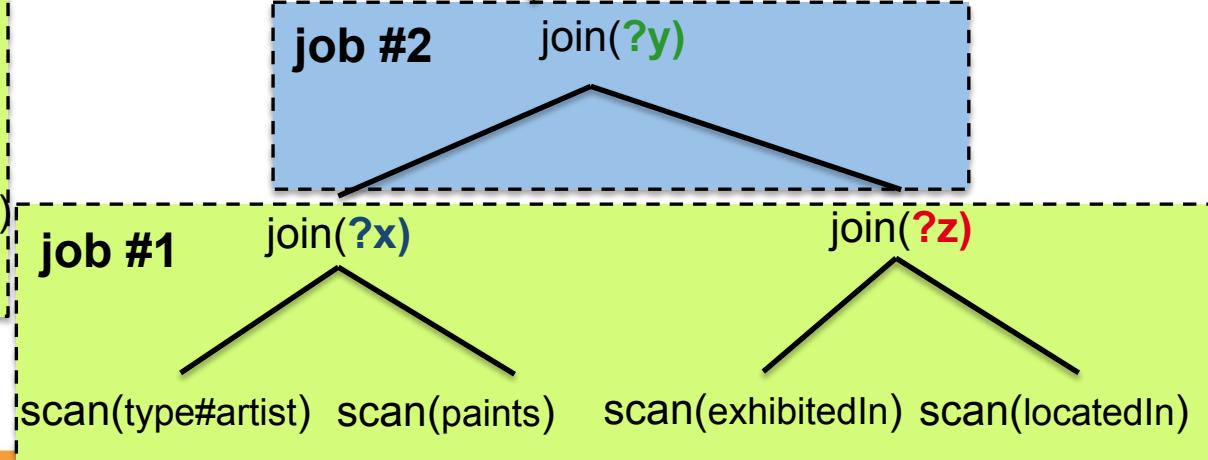
## Query processing:

- File selection based on the property (and object) of the triple
- Joining two or more triple patterns on the same single variable
- Several joins may be executed **in parallel** if they are on different variables
- Heuristic: Query **decomposition** to minimize the number of Hadoop jobs

```
SELECT ?x ?y ?z  
WHERE { ?x :type :artist .  
       ?x :firstName ?y .  
       ?x :creates ?z . }
```



```
SELECT ?x ?y ?z  
WHERE { ?x :type :artist .  
       ?x :paints ?y .  
       ?y :exhibitedIn ?z  
       ?z :locatedIn "madrid" . }
```



# 1. MapReduce-based approaches

Pros	Cons
Scalability	No indexes – scan all data
Fault-tolerance	Shuffling of data
Fast data uploading	Semi-support for joins
Fast and easy to implement	Big initialization overhead High query response times

## 2. Approaches based on key-value stores

- Store RDF triples in a key-value store
    - Indexes on RDF triples
  - Lookups to the key-value store for matching each triple pattern
  - How do I perform the joins?
    - Collect the results somewhere locally
    - Use MapReduce
  - Representative works
    - CumulusRDF [Ladwig11]
    - Stratostore [Stein10]
    - Rya [Punnoose12]
    - AMADA [Bugiotti12]
    - H2RDF [Papailiou12]
    - MAPSIN [Schätzle12]
- 
- The diagram illustrates the classification of the representative works listed in the slide. A vertical bracket on the right side groups the first five systems (CumulusRDF, Stratostore, Rya, AMADA, H2RDF) under the heading "Local query processing". Another vertical bracket further down groups the last two systems (MAPSIN) under the heading "MapReduce-based query processing".

# Storing RDF in key-value stores

- Rya uses Apache Accumulo
  - SPO|-|- POS|-|- OSP|-|-
- H2RDF uses Apache HBase
  - SP|O|- PO|S|- OS|P|-
- AMADA uses Amazon DynamoDB
  - S|P|O P|O|S O|S|P
- MAPSIN uses Apache HBase
  - S|P|O|- O|P|S|-
- Stratustore uses Amazon SimpleDB
  - S|P|O
- CumulusRDF uses Apache Cassandra
  - hierarchical layout: S|{P}O|- P|{O}S|- O|{S}P|-
  - flat layout: S|PO|- PO|S|- PO|'P'| P - O|SP|-

**Sorted Index**

**Sorted Index**

**Hash Index**

**Sorted Index**

**Hash Index**

**Hash Index    Sorted Index**

# Storing RDF in key-value stores

- Rya uses Apache Accumulo
    - SPO|-|- POS|-|- OSP|-|-
  - H2RDF uses Apache HBase
    - SP|O|- PO|S|- OS|P|-
  - AMADA uses Am
    - S|P|O P|O|S|
  - MAPSIN uses Apa
    - S|P|O|- O|S|P|-
  - Stratustore uses Amazon SimpleDB
    - S|P|O
  - CumulusRDF uses Apache Cassandra
    - hierarchical layout: S|{P}O|- P|{O}S|- O|{S}P|-
    - flat layout: S|PO|- PO|S|- PO|'P'| P - O|SP|-
- 3 indexes:  
SPO, POS, OSP**
- Sorted Index**
- Sorted Index**
- Index**
- Hash Index**
- Hash Index Sorted Index**

# Rya storage - example

SPO		
Key	(attribute, value)	
:picasso, :firstName, "Pablo"	-	
:picasso, :p	Key	(attribute, value)
:picasso, ty	:exhibitedIn, :reinasofia, :guernica	-
:guernica, :	:firstName, "Pablo", :picasso,	-
...		
:paints, :gu	Key	(attribute, value)
type, :cubis	:cubist, :picasso, type,	-
...	:guernica, :picasso, :paints	-
"Pablo", :picasso, :firstName	-	
:reinasofia, :guernica, :exhibitedIn	-	
...	-	

POS

OSP

# Strong point of key-value stores: RDF lookups

## Different access patterns for triple pattern matching

Triple pattern	Rya (Accumulo)	H2RDF (HBase)	AMADA (DynamoDB)	MAPSIN (HBase)
(s, p, o)	any lookup	any lookup + select	any lookup + select	any lookup + sselect
(s, p, ?o)	SPO range scan	SP O lookup	S P O lookup + select	S P O lookup + sselect
(s, ?p, o)	OSP range scan	OS P lookup	O S P lookup + select	O S P lookup + sselect
(s, ?p, ?o)	SPO range scan	SP O range scan	S P O lookup	S P O lookup
(?s, p, o)	POS range scan	PO S lookup	P O S lookup + select	O S P lookup + sselect
(?s, p, ?o)	POS range scan	PO S range scan	P O S lookup	any scan + sselect
(?s, ?p, o)	OSP range scan	OS P range scan	O S P lookup	O S P lookup
(?s, ?p, ?o)	any scan	any scan	any scan	any scan

select: filter on client side

sselect: filter on server side

# Weak point of key-value stores: joins

Key-value stores do not support **joins**

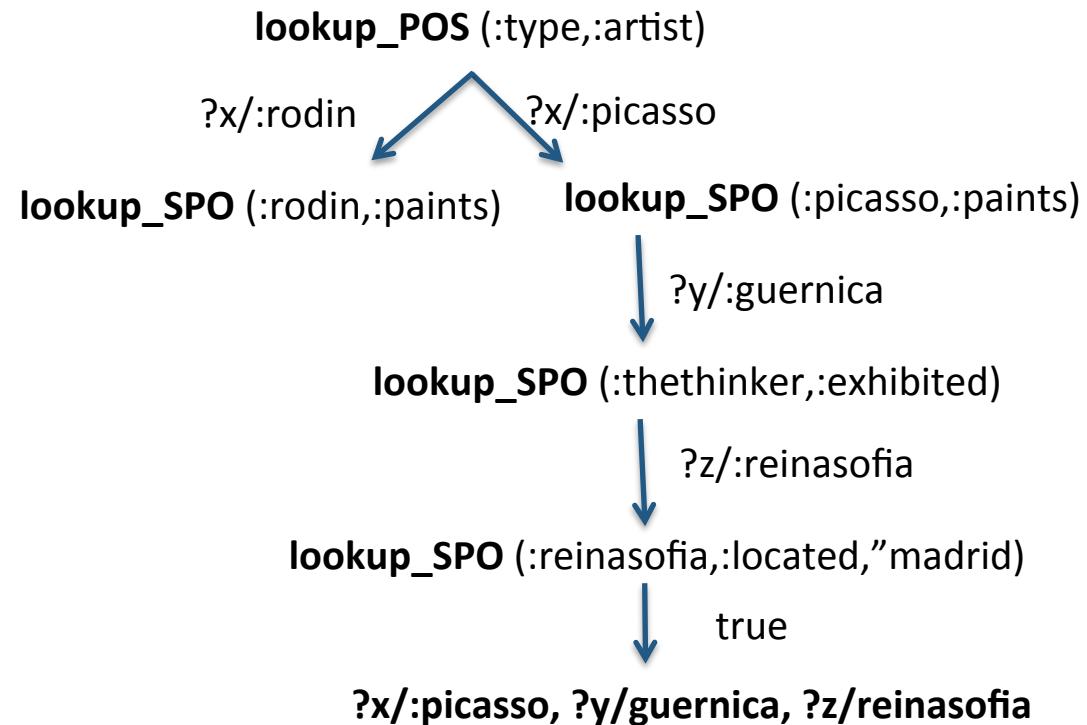
Make all lookups and perform joins out of the store

- **Rya**: index nested loops by query rewriting and lookups
- **H2RDF**: 2 strategies
  - selective joins → centralized – index nested loops
  - non-selective → MapReduce
- **AMADA**: memory hash join
- **MAPSIN**: MapReduce (map phase only)
- **Stratustore**: SimpleDB SELECT query for each star-join and rest of the joins locally
- **CumulusRDF**: Only single triple pattern queries are supported

# Rya joining strategy - example

- Index nested loops
- A lookup for the 1st triple pattern
- For each triple pattern  $t_i$ , use the  $n$  results from triple pattern  $t_{i-1}$  to rewrite  $t_i$  and then perform  $n$  lookups for the rewritten  $t_i$

```
SELECT ?x ?y ?z
WHERE { ?x :type :artist .
         ?x :paints ?y .
         ?y :exhibitedIn ?z
         ?z :locatedIn "madrid". }
```

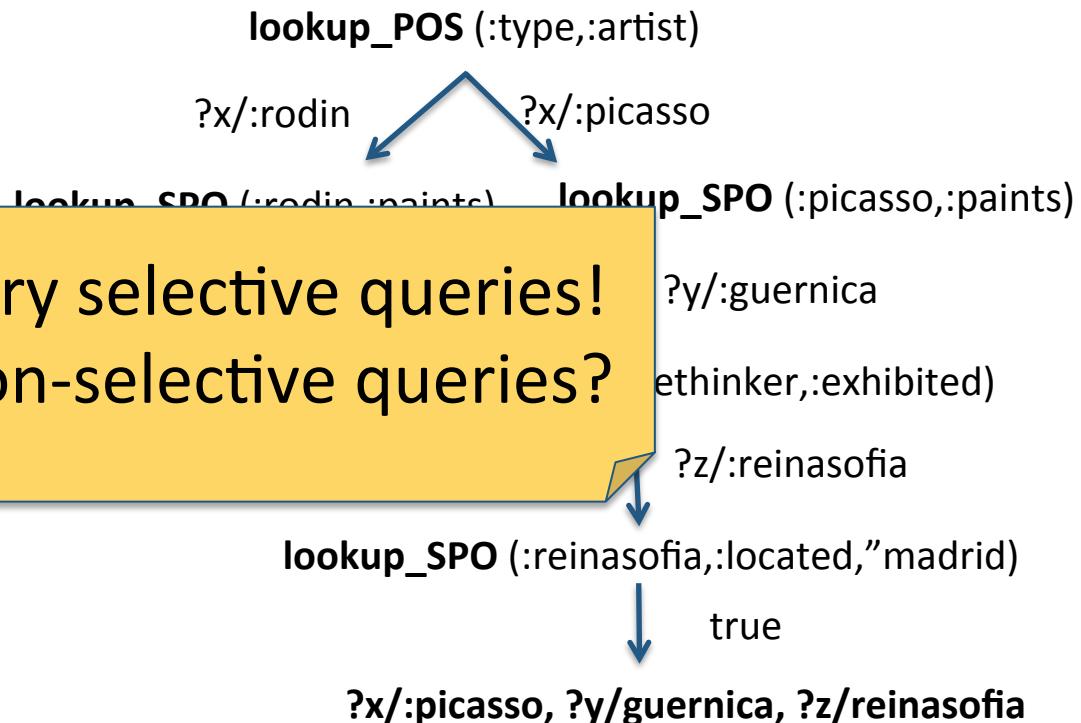


# Rya joining strategy - example

- Index nested loops
- A lookup for the 1st triple pattern
- For each triple pattern  $t_i$ , use the  $n$  results from triple pattern  $t_{i-1}$  to rewrite  $t_i$  and then perform  $n$  lookups for the rewritten  $t_i$

```
SELECT ?x ?y ?z  
WHERE { ?x :type :artist .  
        ?x :paints ?y .  
        ?y :exhibitedIn ?z  
        ?z :1 }
```

Efficient for very selective queries!  
What about non-selective queries?



## 2. Approaches based on key-value stores

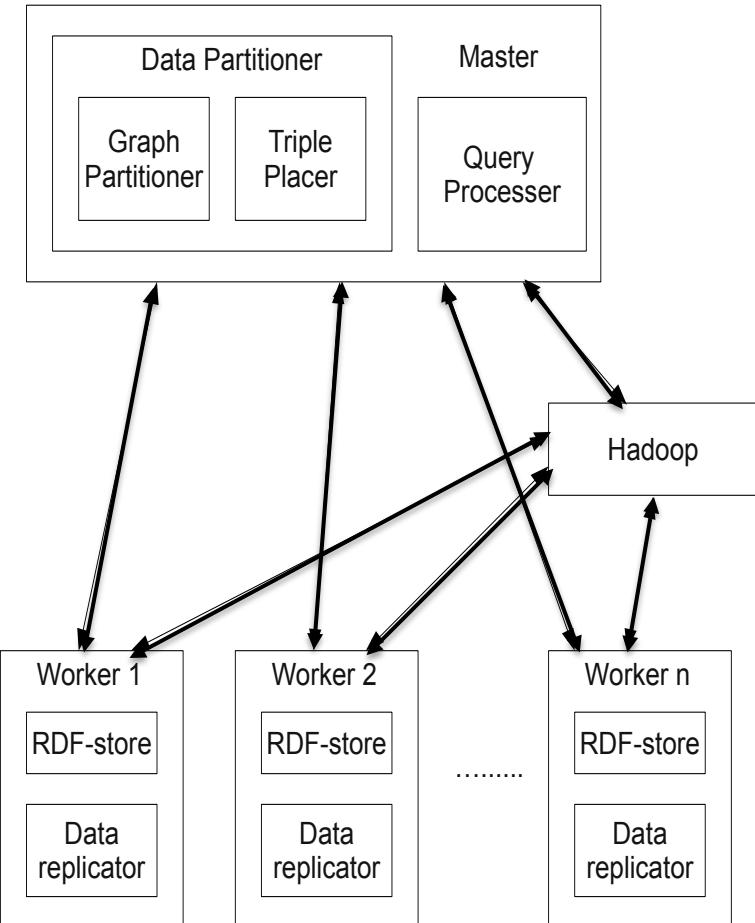
Pros	Cons
Scalability	Low schema expressivity
Fault-tolerance	Joins should be implemented by the user
Very fast lookups	
Efficient for selective queries	Joins are performed on the client-side

### 3. Graph-based approaches

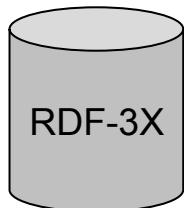
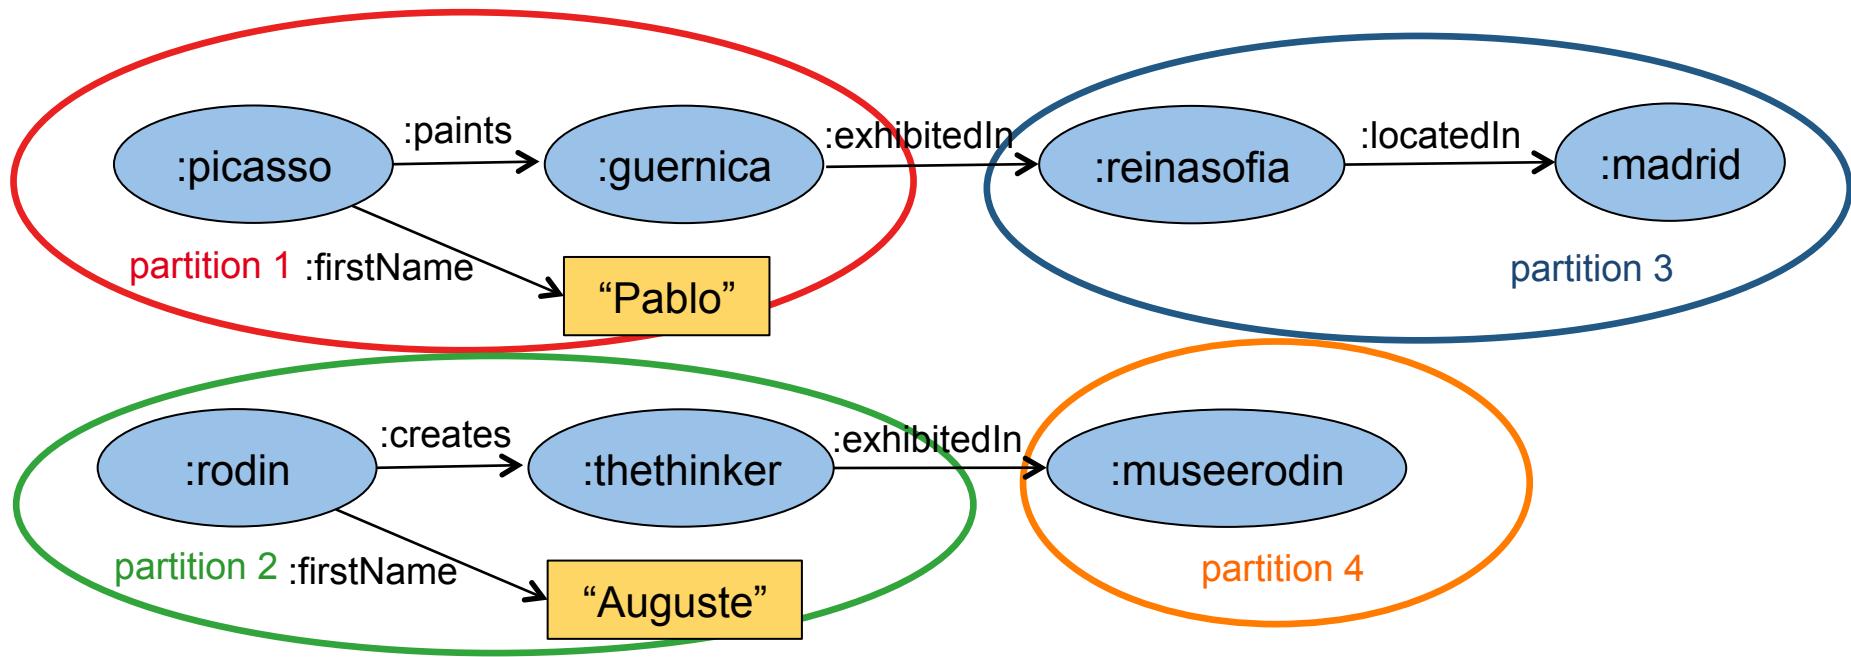
- RDF data partitioning using known **graph** tools
  - [Huang11] uses **METIS**: a graph partitioning tool
- Trinity.RDF [Zeng13] by Microsoft
  - Uses **Trinity**: a distributed graph system over a **memory cloud**

# Graph partitioning [Huang11]

- Data partitioning
  - METIS graph partitioner [METIS]
  - vertex partitioning
  - #partitions = #machines
  - minimum edge cut
  - **type** triples are excluded from the partitioning process
- Data placement
  - place the triple in the partition in which the subject belongs to
  - replication across the partitions
  - directed/undirected n-hop guarantee
- Data storage
  - each machine stores triples into **RDF-3X**



# Graph partitioning [Huang11] - storage



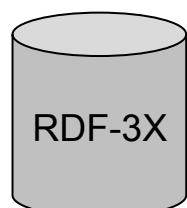
**node 1**



**node 2**

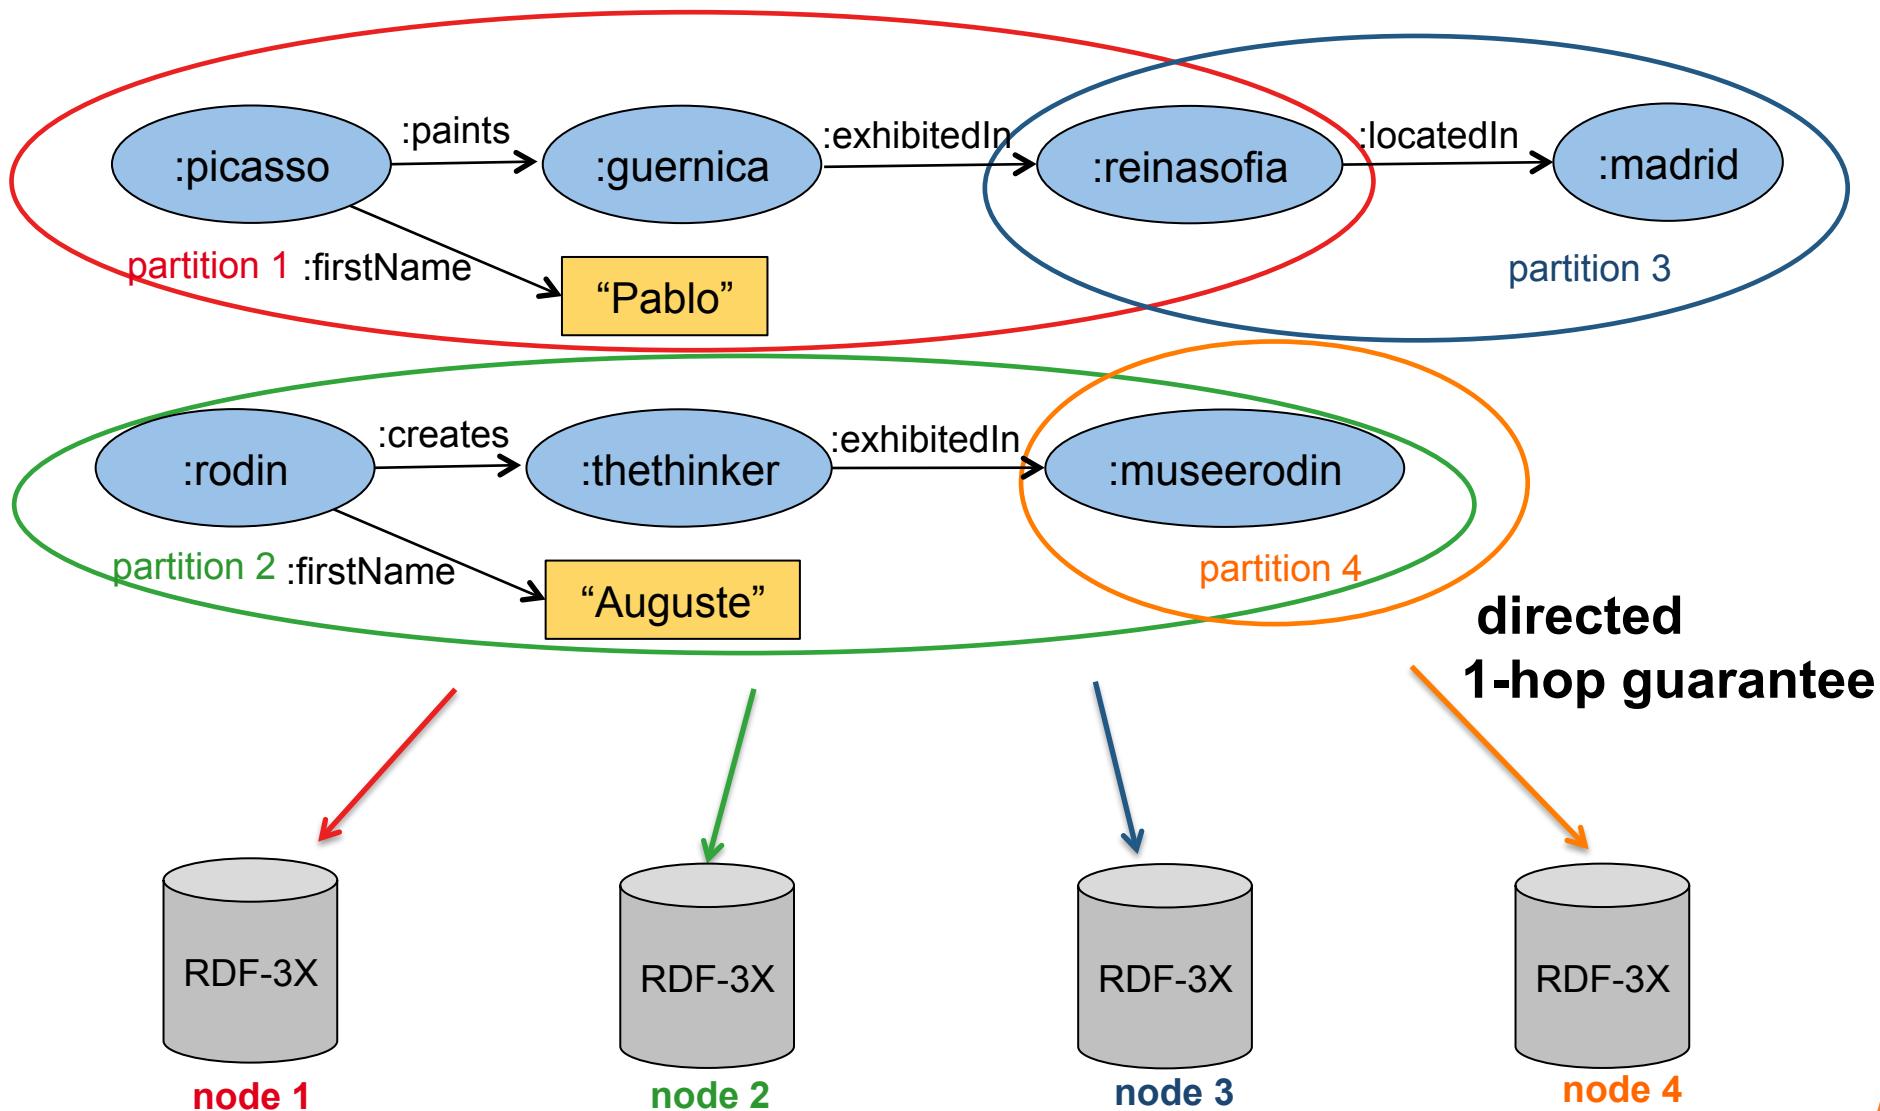


**node 3**

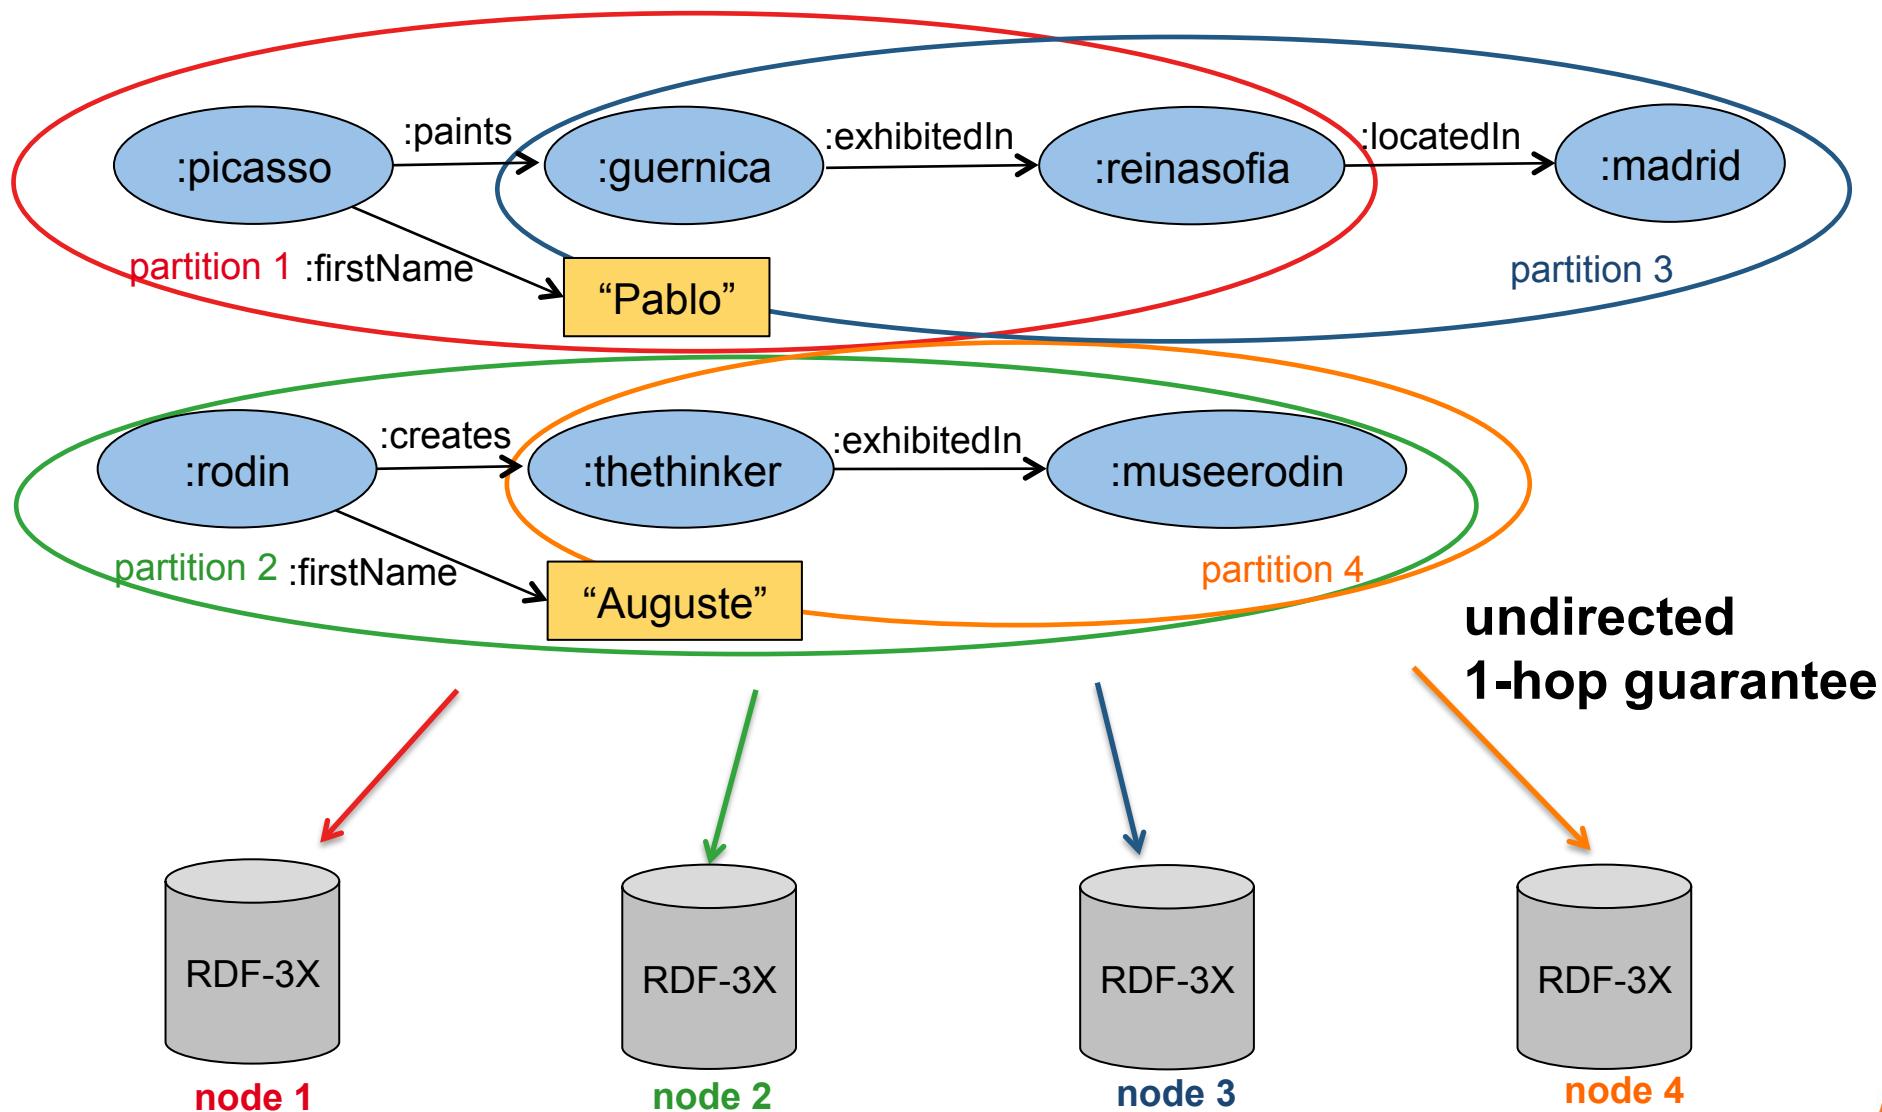


**node 4**

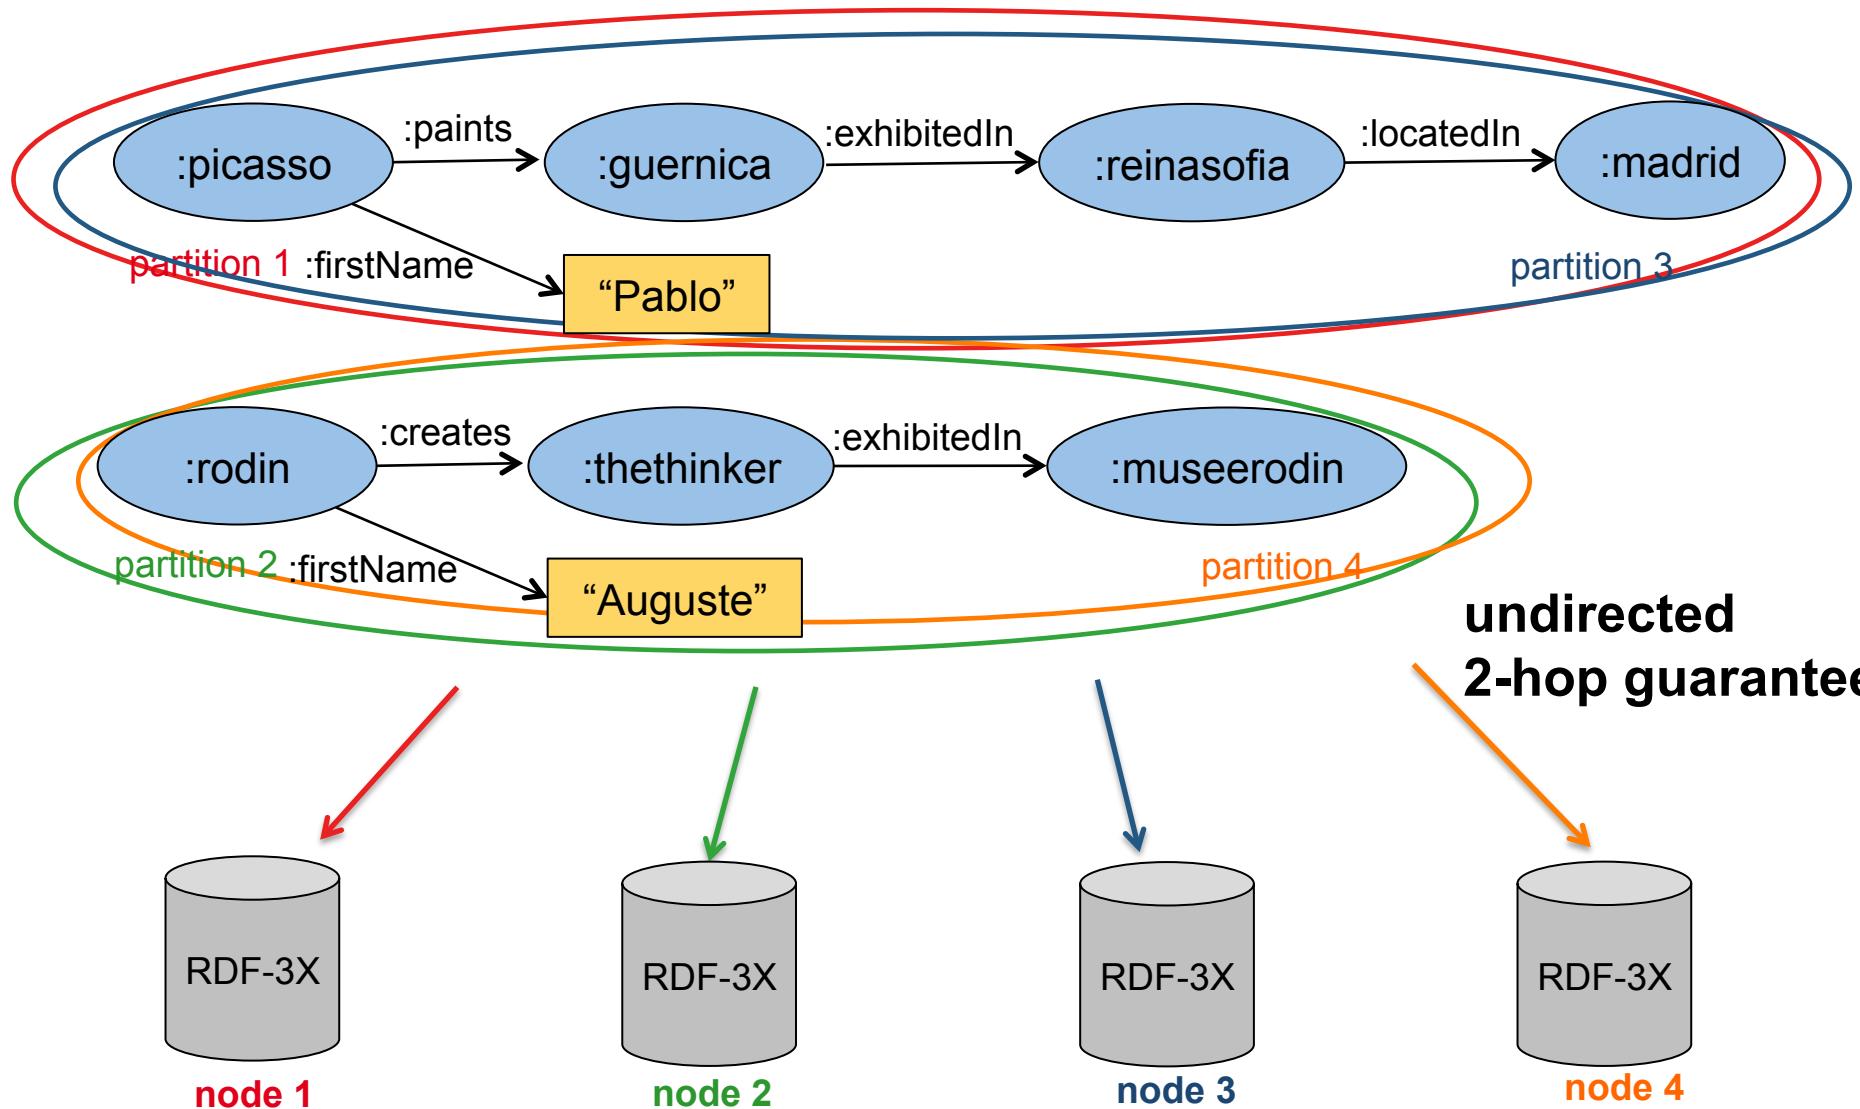
# Graph partitioning [Huang11] - storage



# Graph partitioning [Huang11] - storage

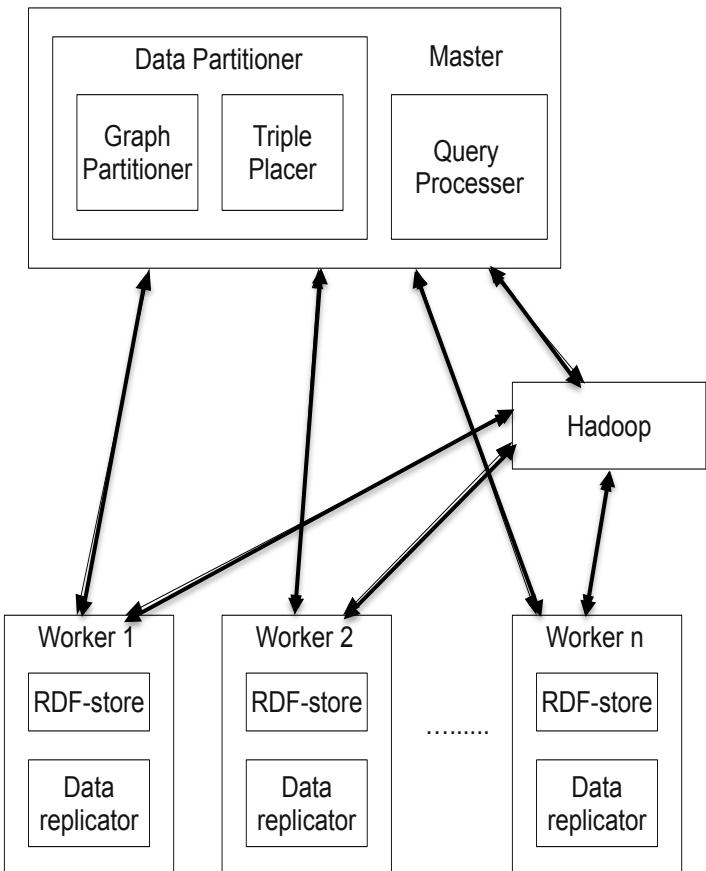


# Graph partitioning [Huang11] - storage



# Graph partitioning [Huang11] - querying

- Query **decomposition** for parallelization
- Check if q is PWOC (**parallelizable without communication**)
  - Yes: answer only from RDF-3X and union results
  - No: answer subqueries from RDF-3X and join them with Hadoop
- # jobs = # subqueries
- **Heuristic optimization**
  - decomposition so as to have **minimum number of subqueries** → minimal edge partitioning of a graph into subgraphs of bounded diameter

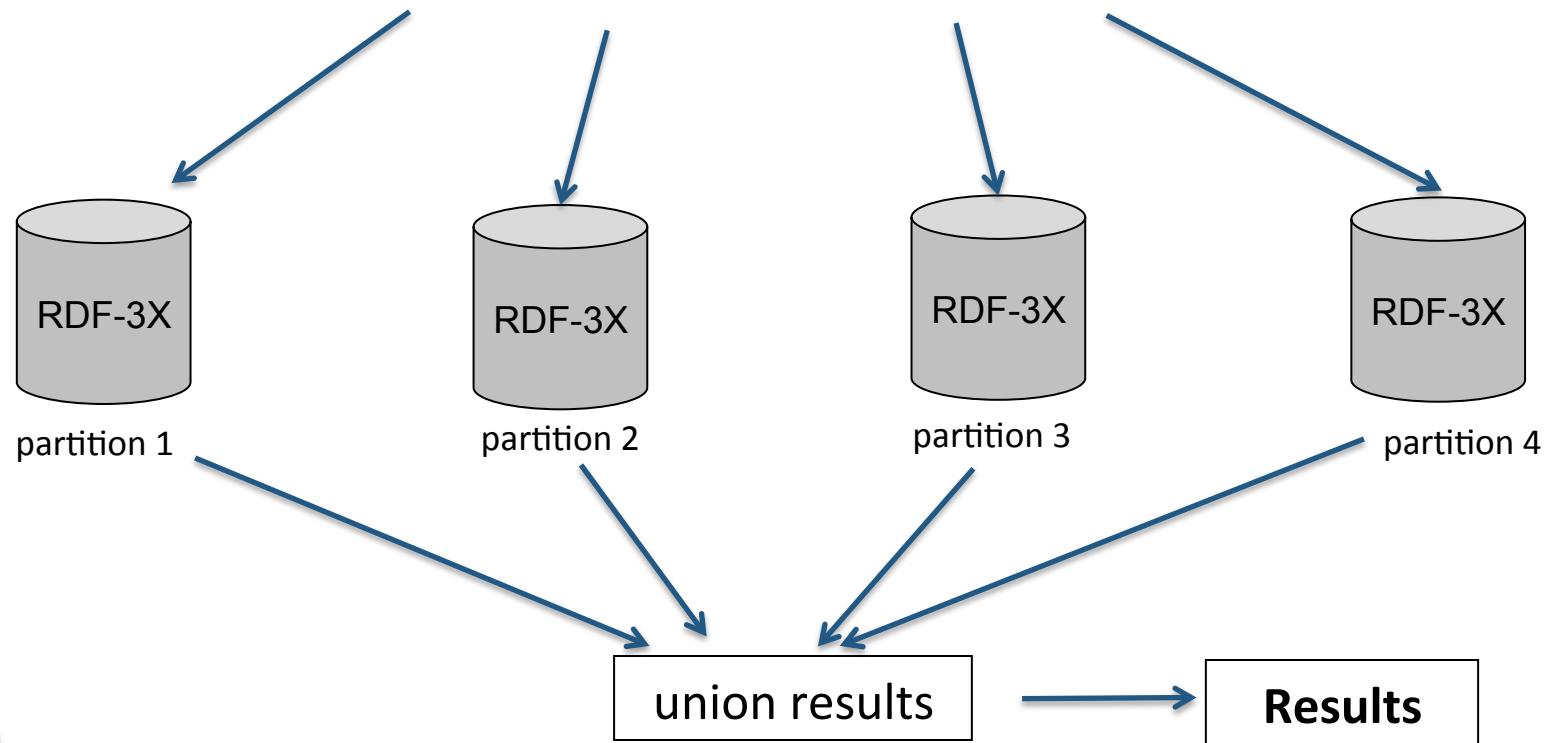


# Graph partitioning [Huang11] - querying

replication with  
1-hop guarantee

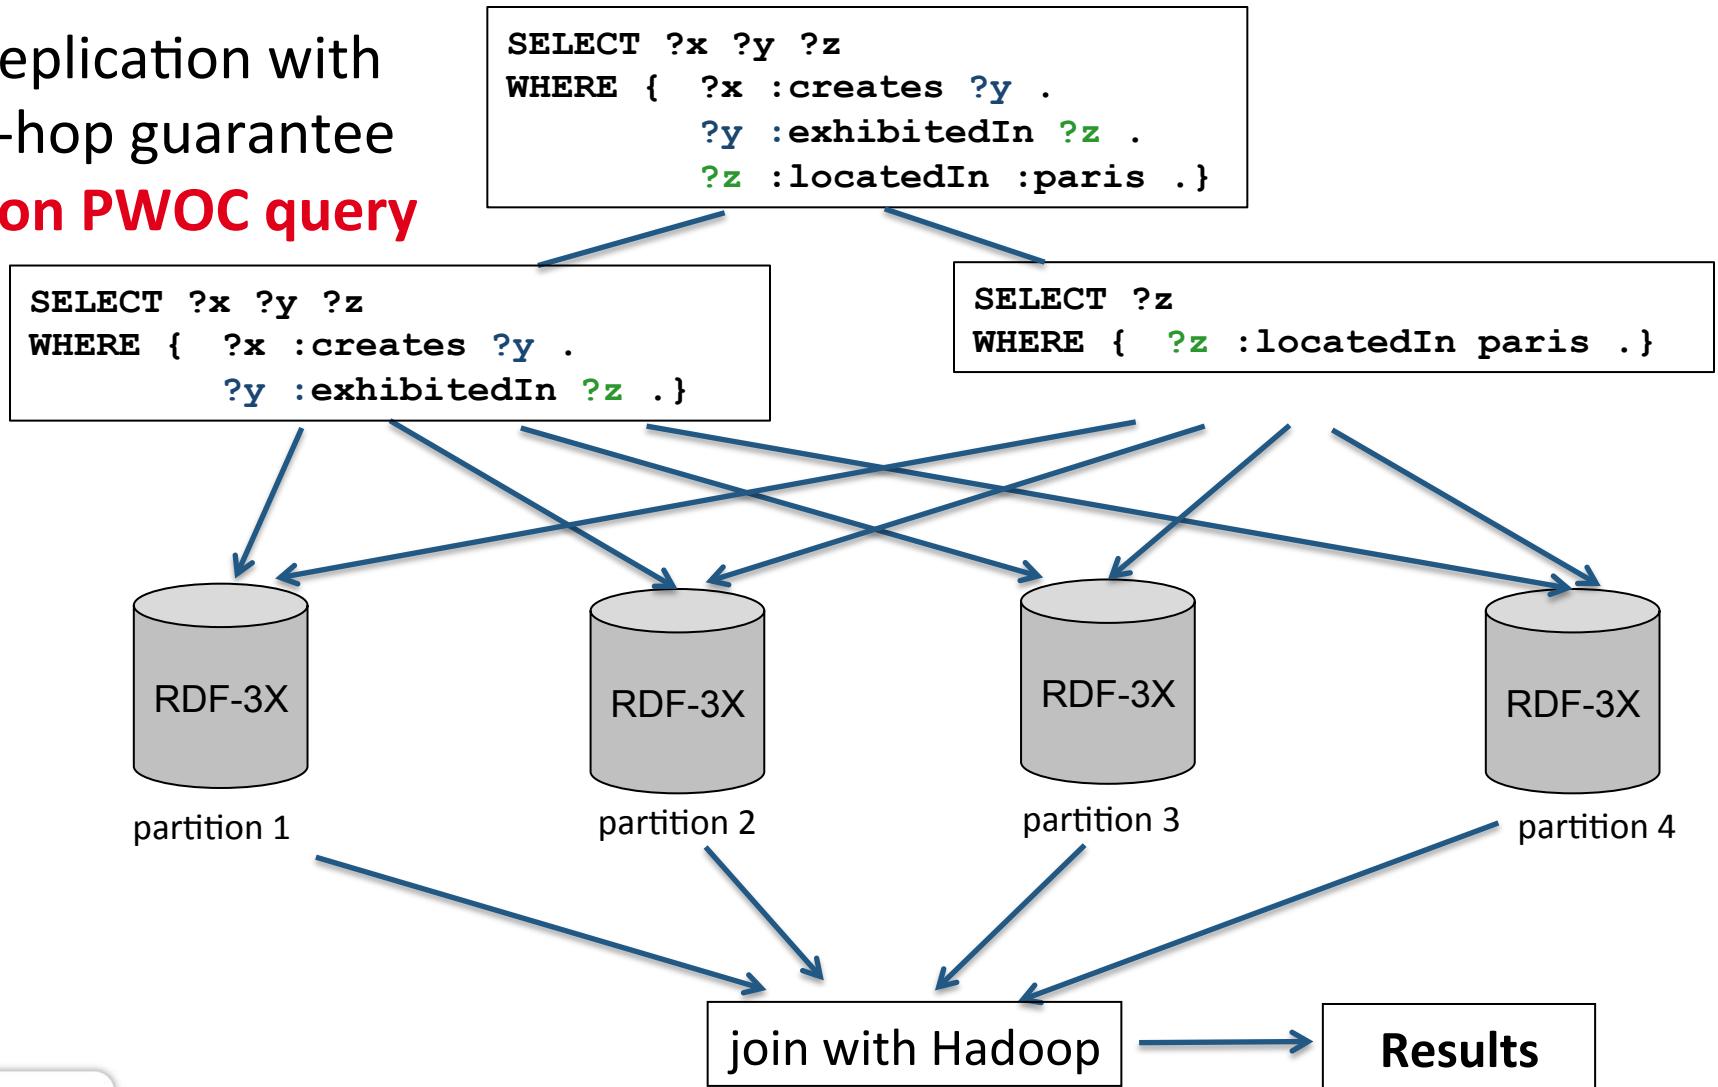
PWOC query

```
SELECT ?x ?y ?z
WHERE {
    ?x type :artist .
    ?x :firstName ?y .
    ?x :creates ?z . }
```



# Graph partitioning [Huang11] - querying

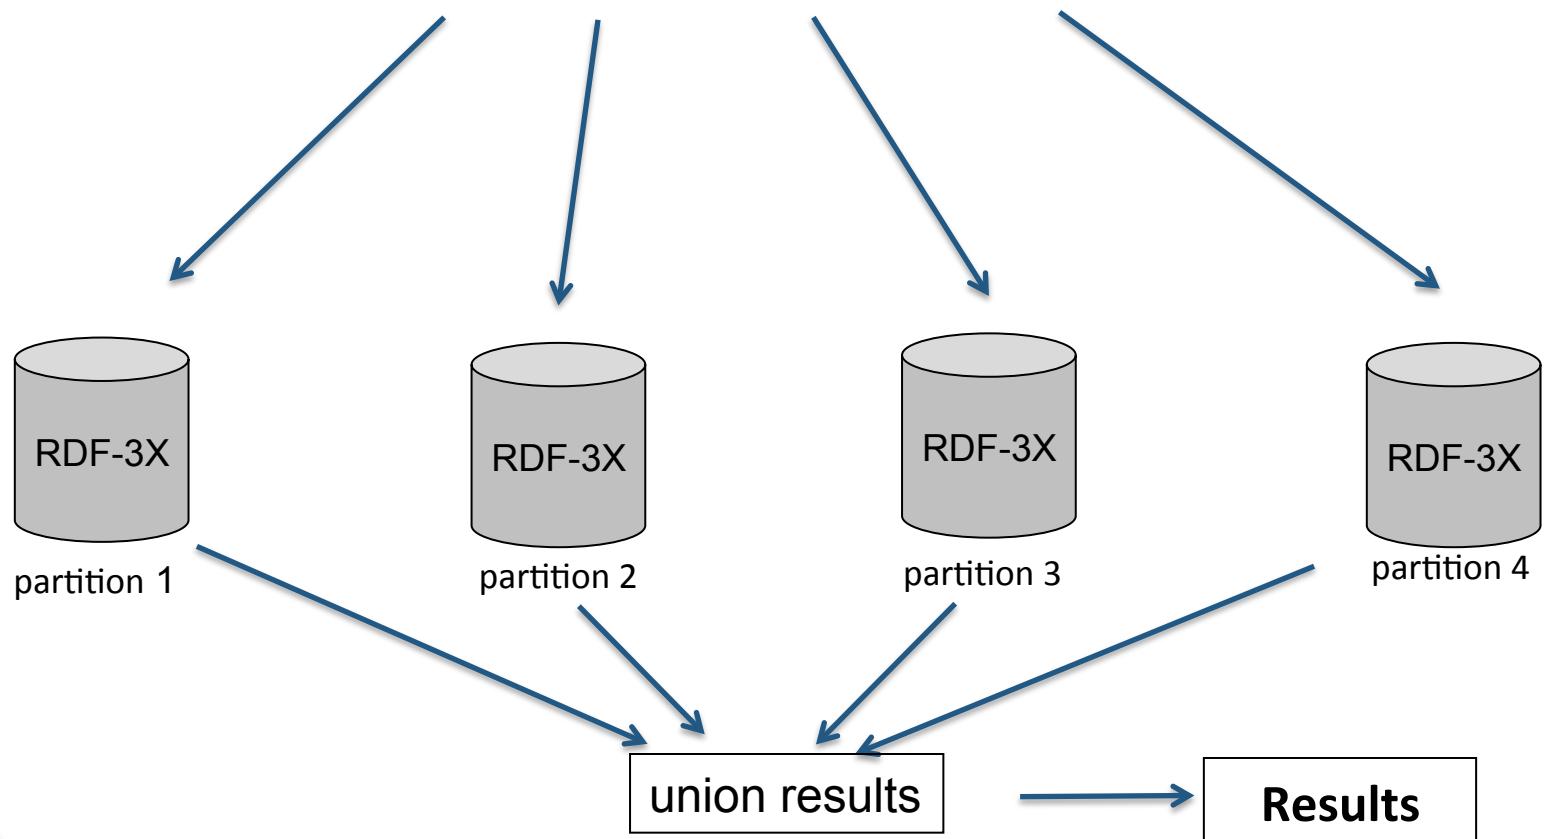
replication with  
1-hop guarantee  
**non PWOC query**



# Graph partitioning [Huang11] - querying

replication with  
2-hop guarantee  
**PWOC query**

```
SELECT ?x ?y ?z
WHERE { ?x :creates ?y .
        ?y :exhibitedIn ?z .
        ?z :locatedIn :paris . }
```



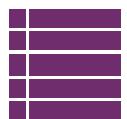
### 3. Graph-based approaches

Pros	Cons
Benefits from RDF structure	Not scalable for big datasets*
Data placement	Subqueries go to all machines
Shuffling is avoided	Uses a lot of resources

## 4. RDF stores based on cloud services

- Statustore [Stein10]
  - AWS SimpleDB
- AMADA [Bugiotti12, Aranda12]
  - Generic cloud-based architecture
  - Built on top of **Amazon** Web Services as a **SaaS**
  - **Route the queries** to a small **subset** of the data that are likely to have matches
  - Trade-offs between **efficiency** and **monetary costs**

# AMADA architecture



## Key-value store

Ideal for indexing and querying **small, structured data**



## Storage for raw data

Ideal for large files



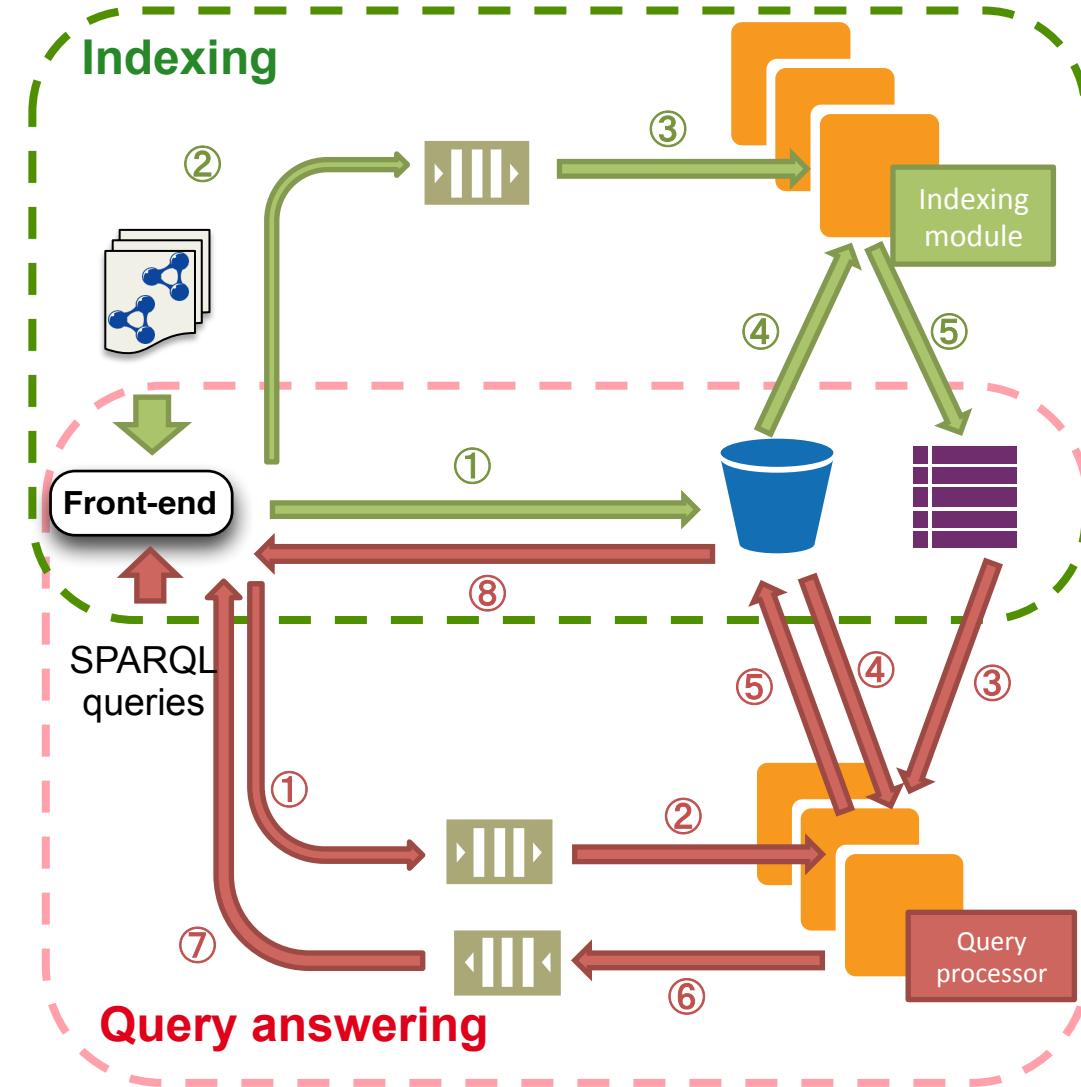
## Virtual machines

**Resizable computing capacity** in the cloud



## Messaging queues

**Queues** for communication between distributed components



# AMADA indexing and storage

**S table**

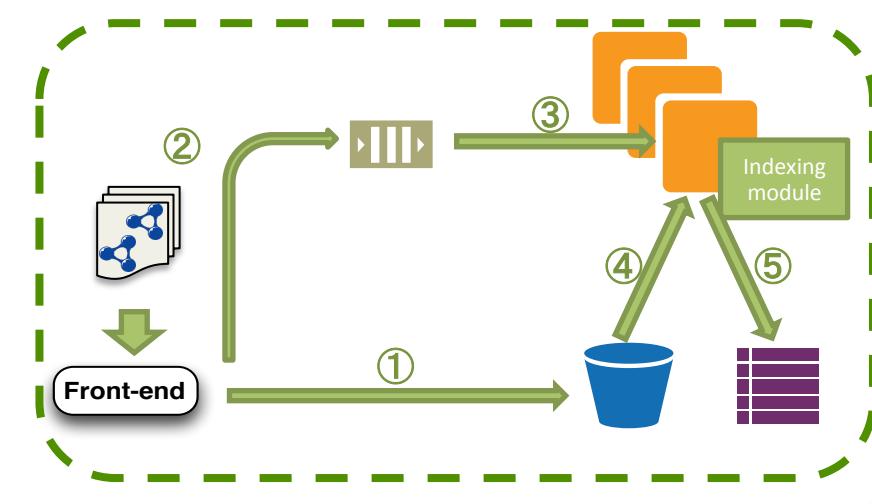
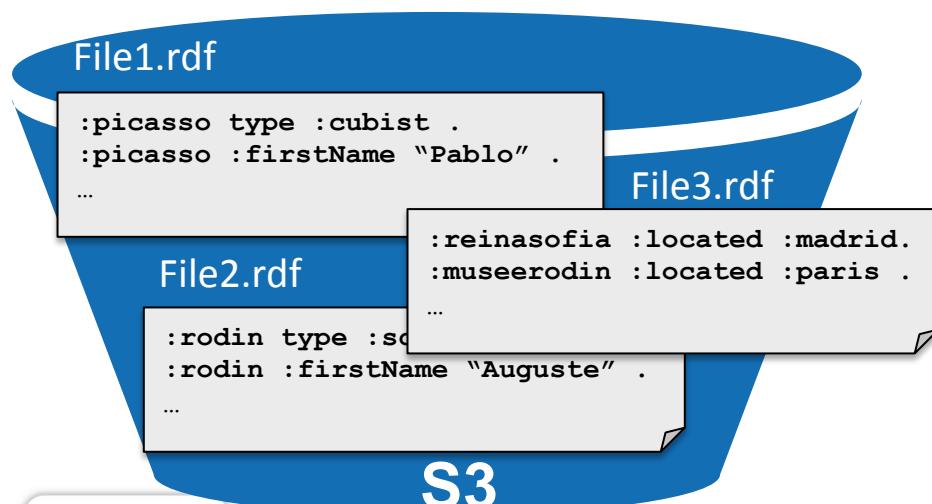
key	(attribute, value)
:picasso	(file1, -)
:guernica	(file1, -)
:reinasofia	(file1, -)
:rodin	(file2, -)
:thethinker	(file2, -)
:museerodin	(file2, -)
...	...

**P table**

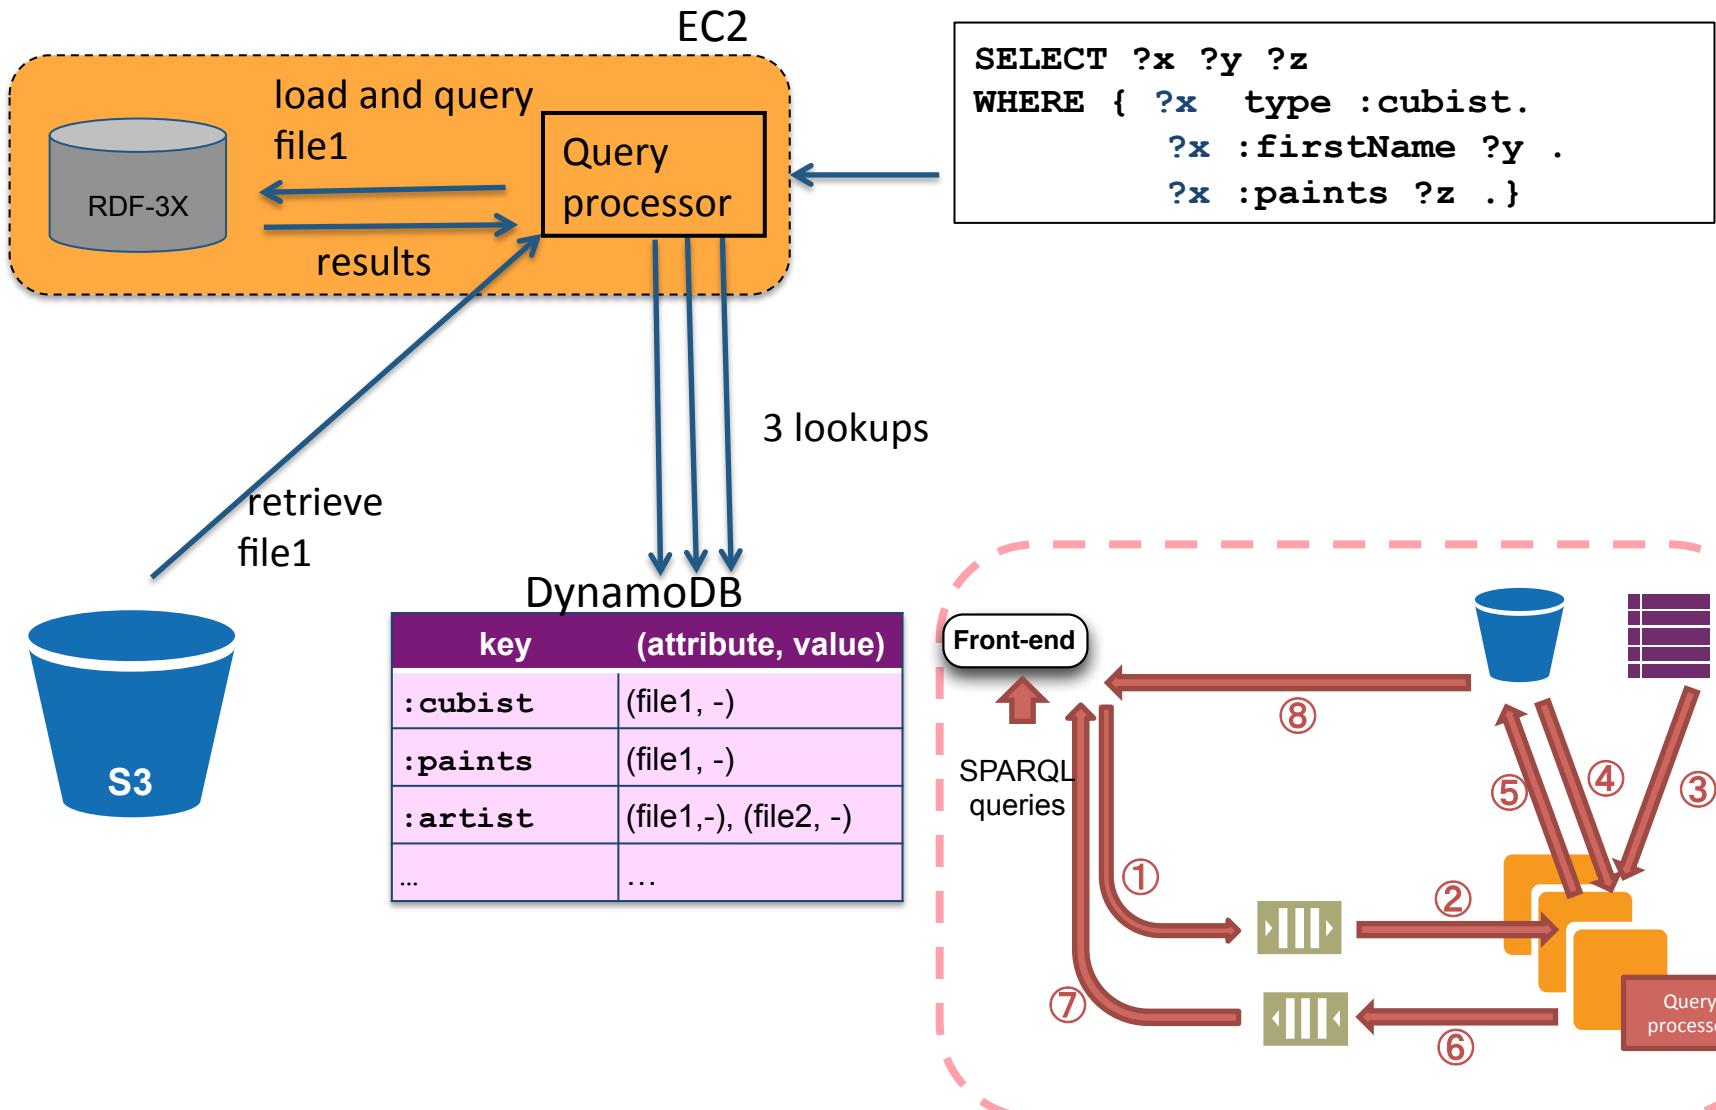
key	(attribute, value)
type	(file1,-), (file2, -)
:firstName	(file1,-), (file2, -)
:paints	(file1,-)
:exhibitedIn	(file1,-), (file2, -)
:locatedIn	(file1,-)
:creates	(file2,-)
...	...

**O table**

key	(attribute, value)
:cubist	(file1,-)
:guernica	(file1,-)
"Pablo"	(file1,-)
:sculptor	(file2,-)
:thethinker	(file2,-)
:museerodin	(file2,-)
...	...



# AMADA querying



## 4. RDF stores based on cloud services

Pros	Cons
Scalability	Depends on data partitioning
Fault-tolerance	Separation between storage and processing services
Elastic allocation of resources	Low latency for non-selective queries
Inter-query parallelism	No control on the services
Benefit from higher-level services (PaaS)	

# Classification of systems

Query processing

MapReduce

Locally

	MAPSIN	SHARD HadoopRDF RAPID+ PigSPARQL
Graph-partitioning	H2RDF	
	Rya AMADA Statustore CumulusRDF	

RDF-3X

key-value  
stores

DFS

Data storage

# Analysis dimensions

## Underlying data storage:

- Key-value stores
- DFS (distributed file system)
- Single-site RDF stores or data storage services supplied by cloud providers

## Query processing of conjunctive queries:

- MapReduce-based
- Local processing (joins are usually performed at a single site)

## RDFS entailment

# Analysis dimensions

## Underlying data storage:

- Key-value stores
- DFS (distributed file system)
- Single-site RDF stores or data storage services supplied by cloud providers

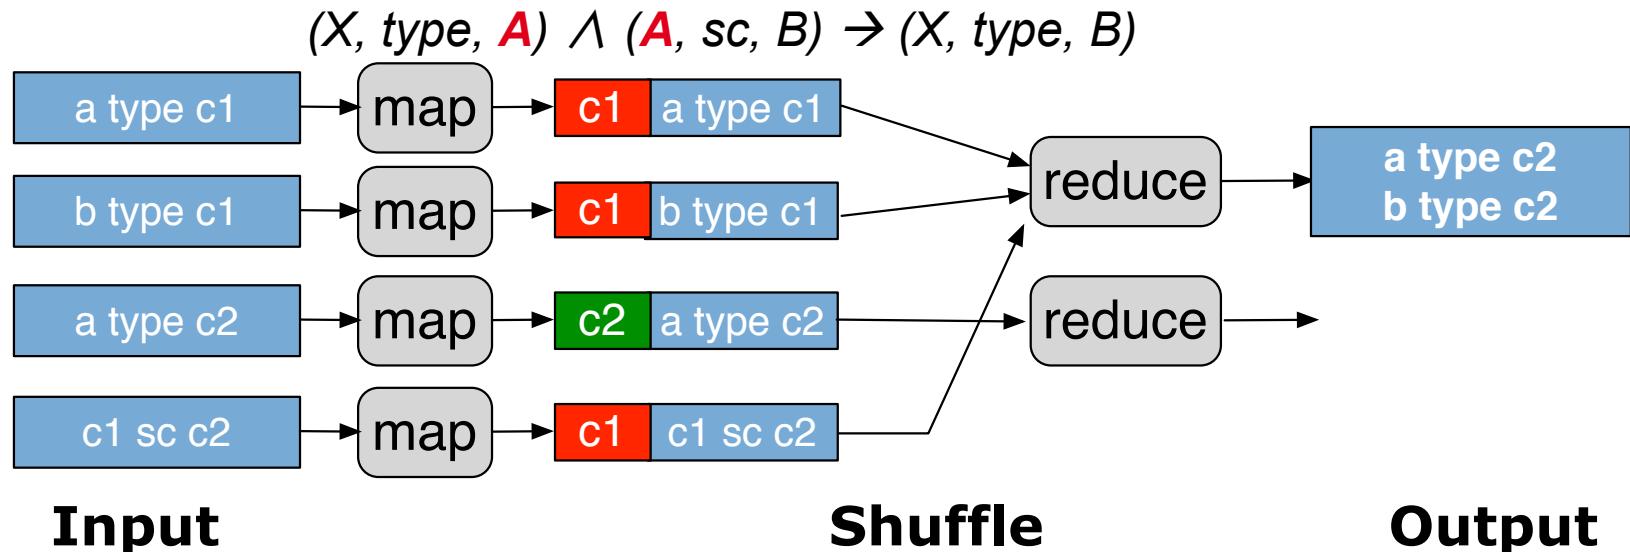
## Query processing of conjunctive queries:

- MapReduce-based
- Local processing (joins are usually performed at a single site)

## RDFS entailment

# RDFS closure in the cloud

- Forward chaining in MapReduce: **WebPie** [Urbani09]

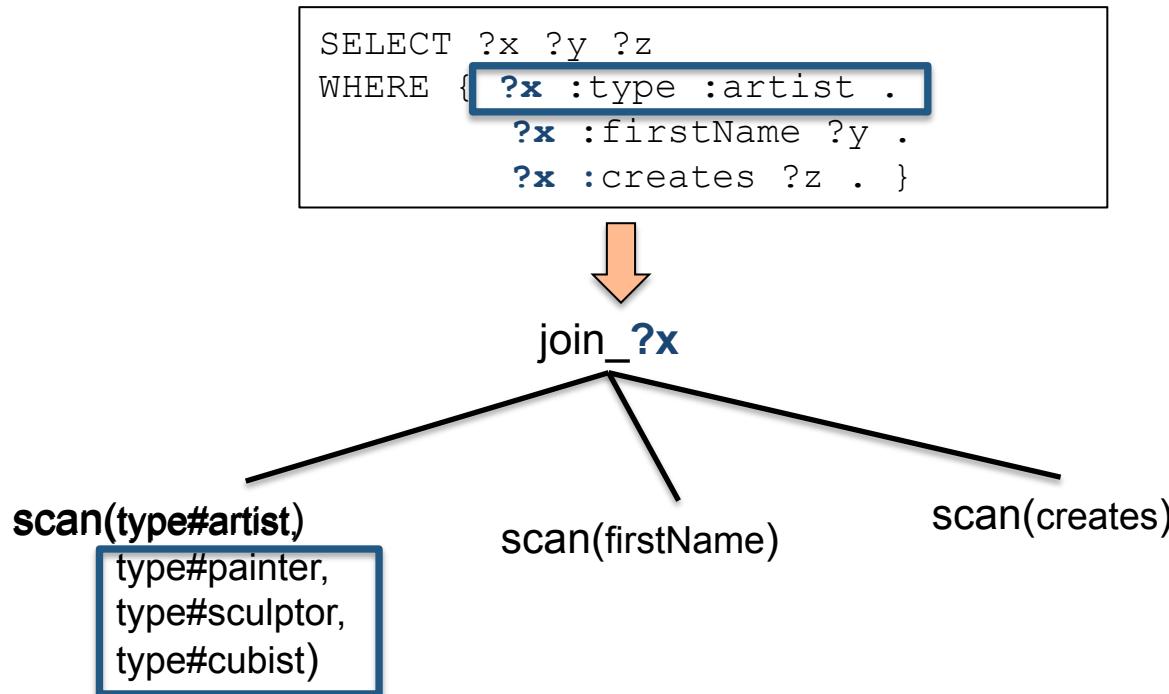


- RDFS triples are kept in memory in each node
- RDFS rules are executed in a specific order → the number of iterations is limited
- Full RDFS closure computation in parallel (MPI-based) [Weaver09]
  - embarrassingly parallel algorithm
  - RDFS triples are kept in memory in each process

# Query reformulation in the cloud

## HadoopRDF [Husain11]

- Query reformulation during the file selection of the query processing



- Transitive closure computation and recursive query processing in MapReduce [Afrati12, Bu12] may apply to RDFS reasoning

# Hybrid approaches

- QueryPie [Urbani11]
  - Precompute implicit triples of the schema
  - Reformulate the query (faster reformulations)
- Rya [Punnoose12]
  - 1 MapReduce job to precompute subclass/subproperty hierarchies
  - Query reformulation at query time

# Analysis dimensions

## Underlying data storage:

- Key-value stores
- DFS (distributed file system)
- Single-site RDF stores or data storage services supplied by cloud providers

## Query processing of conjunctive queries:

- MapReduce-based
- Local processing (joins are usually performed at a single site)

## RDFS entailment:

- Materialization of RDFS closure
- Query reformulation
- Hybrid (only the schema closure is precomputed)

# Do we have a winner?

- No, at least not yet 😊
- Performance depends on query workload
  - Star queries or path queries?
  - Selective or analytics-style queries?
- **Benchmarks** on how real SPARQL queries look:
  - [Duan11]
  - [Arias11]
  - [Picalausa11]
  - Linked Data benchmark council (ldbc.eu)

# 4

## OPEN ISSUES

# Where we go from here?

- A lot of room for optimization
  - Improve query response time in the first place
  - Can we apply traditional optimization techniques (join ordering, statistics, etc.)?
- RDF data partitioning
  - Do the existing graph partitioning algorithms fit?
- RDFS entailment
  - Query reformulation
- More SPARQL features
  - Blank nodes, optional, negation, property paths, etc.

# Where we go from here?

- RDF **views/indexes** in the cloud
- SPARQL multi-query optimization in the cloud
- RDF **updates** in the cloud
- RDF data **analytics**
- How does cloud **variability** affect RDF data management?

# Thank you



Slides can be found at:  
[http://pages.saclay.inria.fr/zoi.kaoudi/  
rdfcloud\\_icde13.pdf](http://pages.saclay.inria.fr/zoi.kaoudi/rdfcloud_icde13.pdf)

# References

- [Abadi07] D. J. Abadi, A. Marcus, S. Madden, K. J. Hollenbach, “Scalable Semantic Web Data Management Using Vertical Partitioning”, in VLDB 2007.
- [Afrati10] F. N. Afrati and J. D. Ullman, “Optimizing Joins in a Map-Reduce Environment,” in EDBT 2010.
- [Afrati11a] F. N. Afrati and J. D. Ullman, “Optimizing Multiway Joins in a Map-Reduce Environment,” IEEE Trans. Knowl. Data Eng., 2011.
- [Afrati11b] F. N. Afrati, V. R. Borkar, M. J. Carey, N. Polyzotis, and J. D. Ullman, “Map-Reduce Extensions and Recursive Queries,” in EDBT 2011.
- [Afrati12] F. N. Afrati, J. D. Ullman, “Transitive closure and recursive Datalog implemented on clusters” in EDBT 2012.
- [Alexaki01] S. Alexaki, V. Christophides, G. Karvounarakis, D. Plexousakis. “On Storing Voluminous RDF Descriptions: The Case of Web Portal Catalogs”, in WebDB 2001.
- [Aranda12] A. Aranda-Andujar, F. Bugiotti, J. Camacho-Rodriguez, D. Colazzo, F. Goasdoué, Z. Kaoudi, and I. Manolescu, “Amada: Web Data Repositories in the Amazon Cloud (demo)”, in CIKM 2012.
- [Arias11] M. Arias, J. D. Fernandez, M. A. Martínez-Prieto, “An Empirical Study of Real-World SPARQL Queries”, in USEWOD 2011.
- [Battre06] D. Battre, A. Hoing, F. Heine, O. Kao, “On Triple Dissemination, Forward-Chaining, and Load Balancing in DHT based RDF stores”, in DBISP2P 2006.
- [Blanas10] S. Blanas, J. M. Patel, V. Ercegovac, J. Rao, E. J. Shekita, and Y. Tian, “A Comparison of Join Algorithms for Log Processing in MapReduce,” in SIGMOD 2010.

# References

- [Bu12] Y. Bu, B. Howe, M. Balazinska, and M. D. Ernst, “The HaLoop Approach to Large-Scale Iterative Data Analysis,” VLDB J, 2012.
- [Bugiotti12] F. Bugiotti, F. Goasdoué, Z. Kaoudi, and I. Manolescu, “RDF Data Management in the Amazon Cloud,” in ICDT/EDBT Workshops 2012.
- [Cattell11] R. Cattell, “Scalable SQL and NoSQL data stores,” SIGMOD Record, May 2011.
- [Dean13] J. Dean, L. A. Barroso, “Tail at Scale”, communication of the ACM, February 2013.
- [Duan11] S. Duan, A. Kementsietsidis, K. Srinivas, O. Udrea, “Apples and oranges: a comparison of RDF benchmarks and real RDF datasets”, in SIGMOD 2011.
- [Goasdoué13] F. Goasdoué, I. Manolescu, and A. Roatis, “Efficient Query Answering against Dynamic RDF Databases”, in EDBT 2013.
- [Hayes04] P. Hayes, “RDF Semantics”, W3C Recommendation, February 2004, <http://www.w3.org/TR/rdf-mt/>.
- [Husain11] M. Husain, J. McGlothlin, M. M. Masud, L. Khan, and B. M. Thuraisingham, “Heuristics-Based Query Processing for Large RDF Graphs Using Cloud Computing,” IEEE Trans. on Knowl. and Data Eng., 2011.
- [Iosup11] A. Iosup, N. Yigitbasi, D.H. J. Epema, “On the Performance Variability of Production Cloud Services”, in CCGRID 2011.
- [Kaoudi13] Z. Kaoudi, M. Koubarakis, “Distributed RDFS Reasoning over Structured Overlay Networks”, Journal on Data Semantics, 2013.
- [Kotoulas10] S. Kotoulas, E. Oren, F. Harmelen, “Mind the data skew: distributed inferencing by speeddating in elastic regions”, in WWW 2010.

# References

- [Mallea11] A. Mallea, M. Arenas, A. Hogan, A. Polleres, “On Blank Nodes”, in ISWC 2011.
- [METIS] METIS - Serial Graph Partitioning and Fill-reducing Matrix Ordering:  
<http://glaros.dtc.umn.edu/gkhome/metis/metis/overview>
- [Neumann10] T. Neumann and G. Weikum, “The RDF-3X Engine for Scalable Management of RDF Data,” VLDBJ 2010.
- [Picalausa11] F. Picalausa, S. Vansumeren, “What are real SPARQL queries like?”, in SWIM 2011.
- [Punnoose12] R. Punnoose, A. Crainiceanu, and D. Rapp, “Rya: A Scalable RDF Triple Store for the Clouds,” in 1st International Workshop on Cloud Intelligence (in conjunction with VLDB), 2012.
- [Ravindra11] P. Ravindra, H. Kim, K. Anyanwu, “An Intermediate Algebra for Optimizing RDF Graph Pattern Matching on MapReduce”, in ESWC 2011.
- [Rohloff10] K. Rohloff and R. E. Schantz, “High-Performance, Massively Scalable Distributed Systems using the MapReduce Software Framework: the SHARD Triple-Store,” in Programming Support Innovations for Emerging Distributed Applications, 2010.
- [Schad10] J. Schad, J. Dittrich, J. Quiané-Ruiz, “Runtime Measurements in the Cloud: Observing, Analyzing, and Reducing Variance”, PVLDB 2010.
- [Schätzle12] A. Schätzle, M. Przyjacielski-Zablocki, C. Dorner, T. Hornung, G. Lausen, “Cascading Map-Side Joins over HBase for Scalable Join Processing”, in SSWS+HPCSW 2012.
- [Schätzle11] A. Schätzle, M. Przyjacielski-Zablocki, G. Lausen, ”PigSPARQL: mapping SPARQL to Pig Latin”, in SWIM 2011.

# References

- [Stein10] R. Stein and V. Zacharias, “RDF On Cloud Number Nine,” in 4<sup>th</sup> Workshop on New Forms of Reasoning for the Semantic Web: Scalable and Dynamic, May 2010.
- [Theoharis05] Y. Theoharis, V. Christophides, and G. Karvounarakis. “Benchmarking Database Representations of RDF/S Stores”, in ISWC 2005.
- [Urbani09] J. Urbani, S. Kotoulas, E. Oren, and F. van Harmelen, “Scalable Distributed Reasoning using MapReduce,” in ISWC 2009.
- [Urbani11] J. Urbani, F. van Harmelen, S. Schlobach, and H. Bal, “QueryPIE: Backward Reasoning for OWL Horst over Very Large Knowledge Bases,” in ISWC 2011.
- [Weaver09] J. Weaver and J. A. Hendler, “Parallel Materialization of the Finite RDFS Closure for Hundreds of Millions of Triples”, in ISWC 2009.
- [Weiss08] C. Weiss, P. Karras, A. Bernstein, “Hexastore: sextuple indexing for semantic web data management”, PVLDB 2008.
- [Wilkinson06] K. Wilkinson, “Jena Property Table Implementation”, in SSWS 2006.
- [Wu11] S. Wu, F. Li, S. Mehrotra, B. C. Ooi, “Query Optimization for Massively Parallel Data Processing”, in SOCC 2011.
- [Zeng13] K. Zeng, J. Yang, H. Wang, B. Shao, and Z. Wang, “A Distributed Graph Engine for Web Scale RDF Data”, in PVLDB 2013.